

Visual Computing I:

Interactive Computer Graphics and Vision



Features, Features Description and Feature Matching

Stefanie Zollmann and Tobias Langlotz

Last time..

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Linear Filters and Image Convolutions

Original



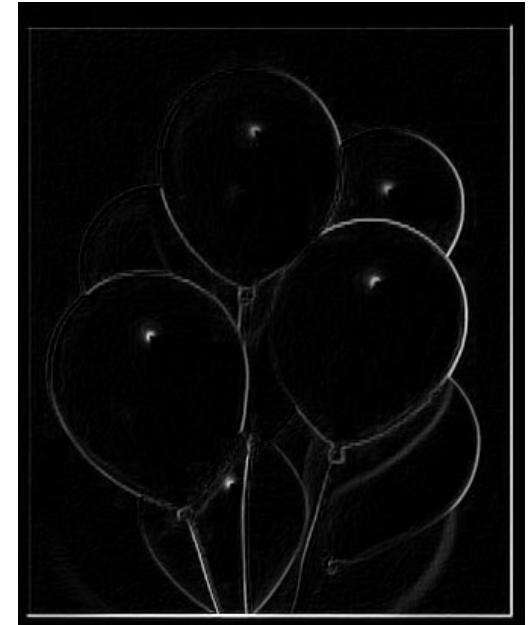
Edge horizontal



Edge vertical



Edge (Sobel Operator)



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$

$$\begin{matrix} -\frac{1}{2} & 0 & \frac{1}{2} \\ -1 & 0 & 1 \\ -\frac{1}{2} & 0 & \frac{1}{2} \end{matrix}$$

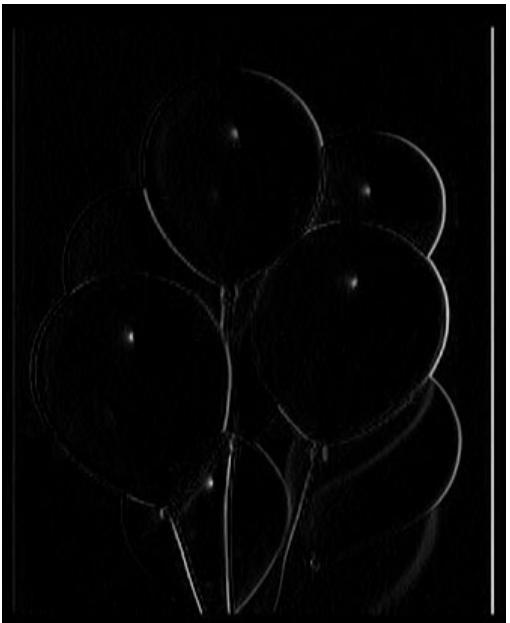
$$\begin{matrix} -\frac{1}{2} & -1 & -\frac{1}{2} \\ -0 & 0 & 0 \\ \frac{1}{2} & 1 & \frac{1}{2} \end{matrix}$$

Linear Filters and Image Convolutions

Original



Edge horizontal/
vertical



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$

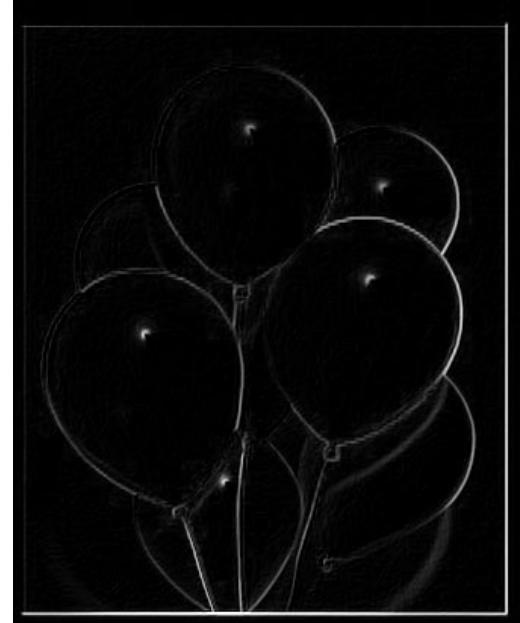
$$\begin{matrix} -\frac{1}{2} & 0 & \frac{1}{2} \\ -1 & 0 & 1 \\ -\frac{1}{2} & 0 & \frac{1}{2} \end{matrix}$$

Edge vertical/
horizontal



$$\begin{matrix} -\frac{1}{2} & -1 & -\frac{1}{2} \\ -0 & 0 & 0 \\ \frac{1}{2} & 1 & \frac{1}{2} \end{matrix}$$

Edge (Sobel Operator)



Linear Filters and Image Convolutions

- Flipping operator in convolutions?!
 - The short answer is: **Yes**, in pure mathematics and traditional signal processing, the definition of convolution involves flipping the kernel. However, also **No** because in computer vision, graphics, and deep learning, the operation we call "convolution" is actually cross-correlation, which does not flip the kernel.
 - There is a lot of subtlety and differences between the meanings of convolution and correlation in different fields (also Tobias did not know about that)

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

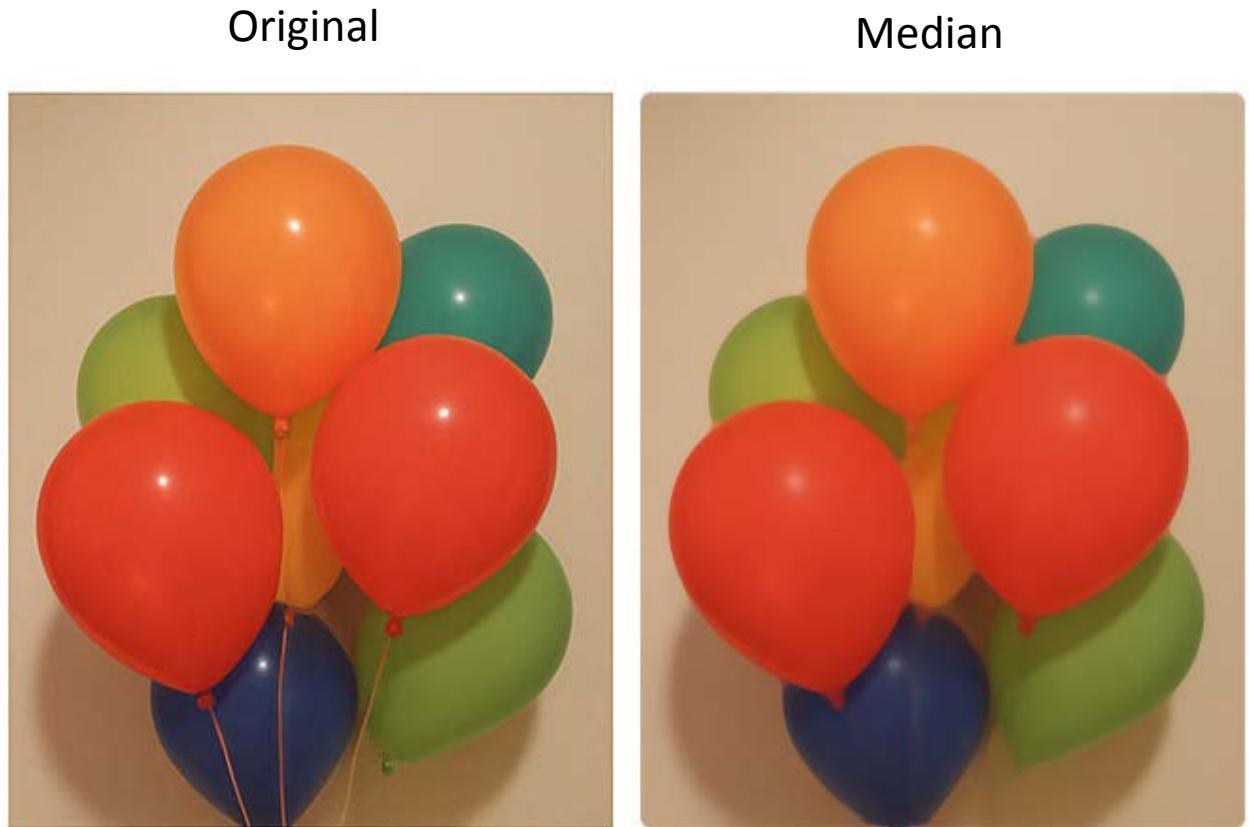
2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Image Filters - Non-Linear Filters

Non-Linear Filters

- Not all filters are implemented as convolution operations
- E.g. Median filter
 - Take a local neighbourhood (e.g. a 3×3 or 5×5 window)
 - Sort the pixel values
 - Use the middle value as the result
 - Has a smoothing effect, but preserves sharp edges



2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Image Features & Feature Detection

Image Features & Feature Detection

- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views



What are good image features?

Image Features & Feature Detection

- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views
- Feature descriptors
 - Used to tell features apart
 - Ideally robust to changes in lighting, viewpoint, contrast, etc.
- Three common types of features
 - Edges – gradient-based
 - Corners – gradient-based
 - Blobs – bright/dark centre/surround



What are good image features?

Image Features & Feature Detection

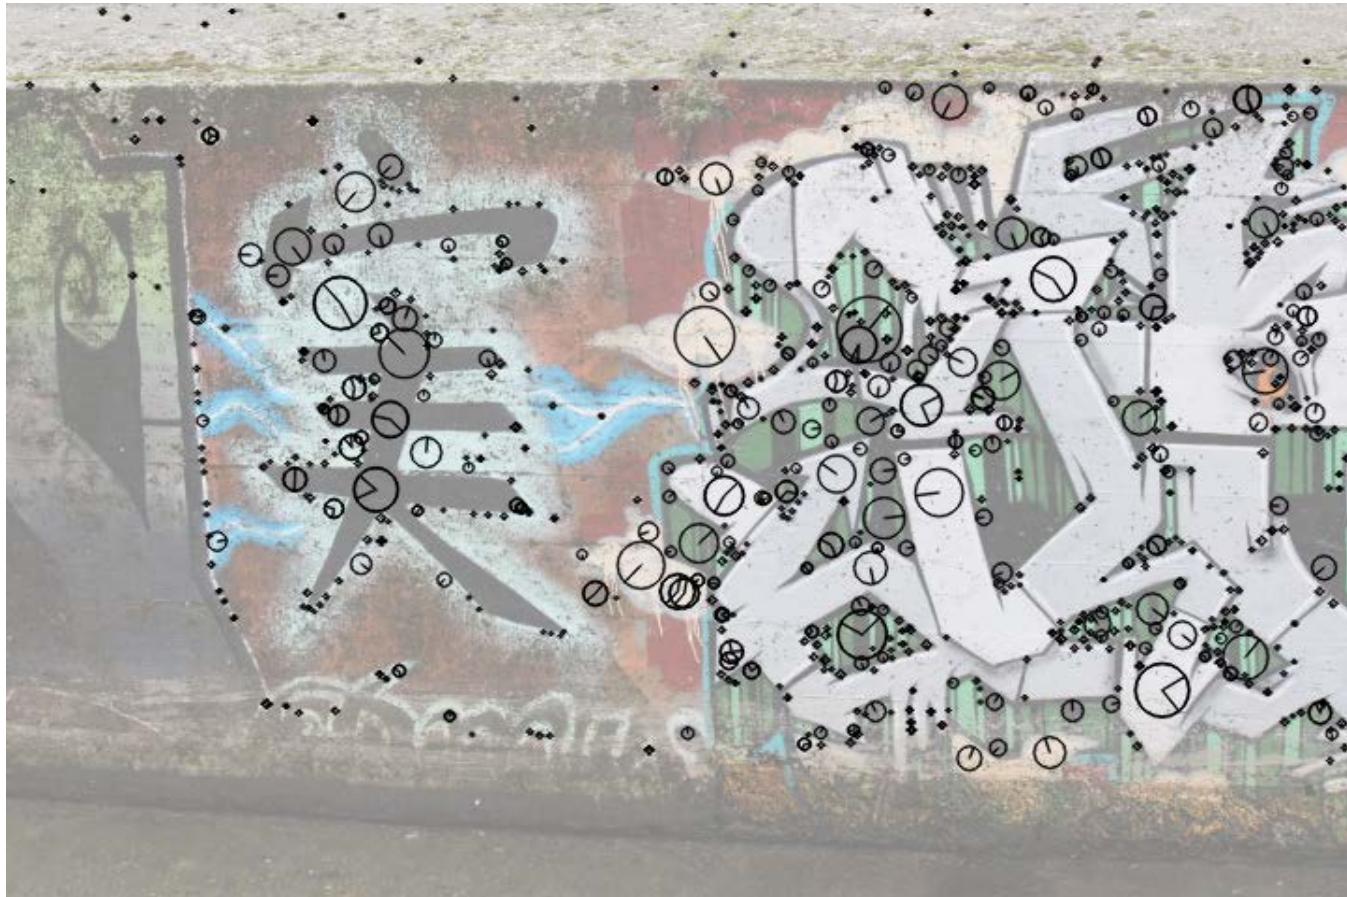
- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views
- Feature descriptors
 - Used to tell features apart
 - Ideally robust to changes in lighting, viewpoint, contrast, etc.
- Three common types of features
 - Edges – gradient-based
 - Corners – gradient-based
 - Blobs – bright/dark centre/surround



Example: Corners

Image Features & Feature Detection

- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views
- Feature descriptors
 - Used to tell features apart
 - Ideally robust to changes in lighting, viewpoint, contrast, etc.
- Three common types of features
 - Edges – gradient-based
 - Corners – gradient-based
 - Blobs – bright/dark centre/surround



Example: Blobs

Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)

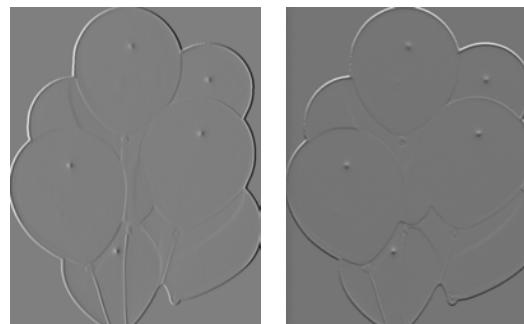
Original



Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)
 - Step 2: Compute gradients in x and y direction (> Sobel operators H_u, H_v), gradient is a vector
$$\mathbf{g} = [H_u \quad H_v]^T$$

Original

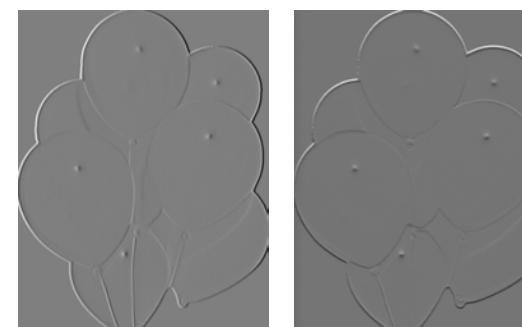


$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)
 - Step 2: Compute gradients in x and y direction (> Sobel operators H_u, H_v), gradient is a vector
$$\mathbf{g} = [H_u \quad H_v]^T$$
 - Step 3: compute gradient magnitudes (edge strength) and gradient directions (edge directions)

Original

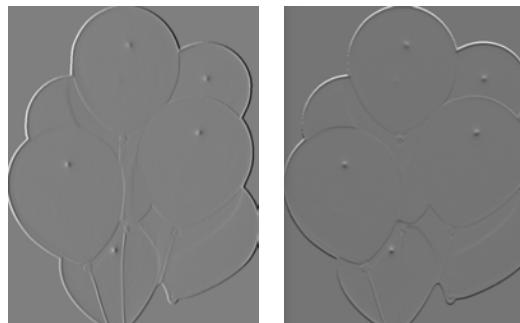


$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Magnitude

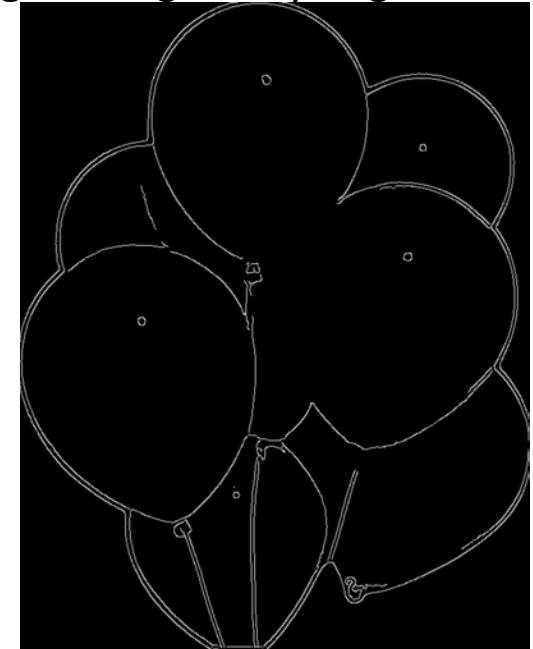
Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)
 - Step 2: Compute gradients in x and y direction (> Sobel operators H_u, H_v), gradient is a vector
$$\mathbf{g} = [H_u \quad H_v]^T$$
 - Step 3: compute gradient magnitudes (edge strength) and gradient directions (edge directions)
 - Step 4: select points with maximum magnitudes (above threshold T1)
 - Step 5: non-maximum suppression with hysteresis:
 - For all points with magnitudes above high-threshold T1
 - Trace along the edge (in edge direction) suppress all pixels that are not on traced line (pixel by pixel) and mark as visited
 - Stop tracing when point's magnitude is below low-threshold T2 or already visited



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Edges using Canny Edge Detector



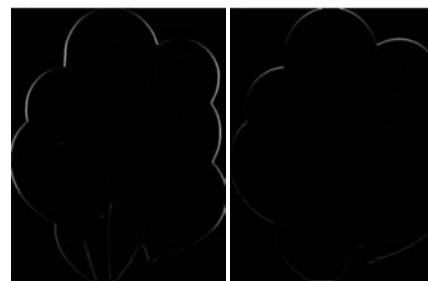
Magnitude

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

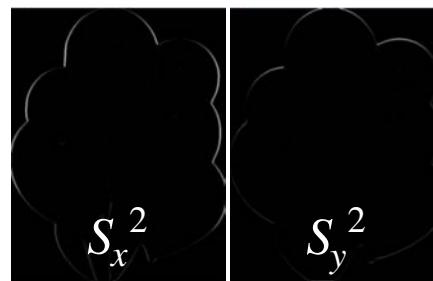


Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



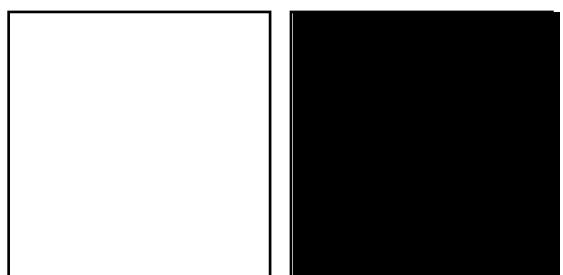
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



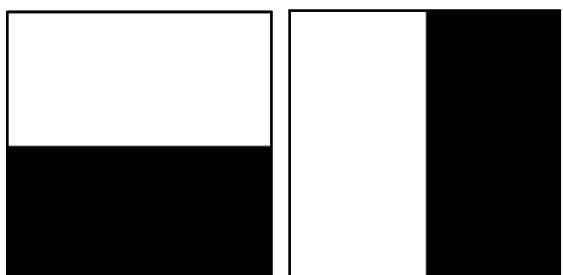
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low



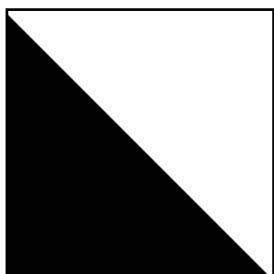
$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)



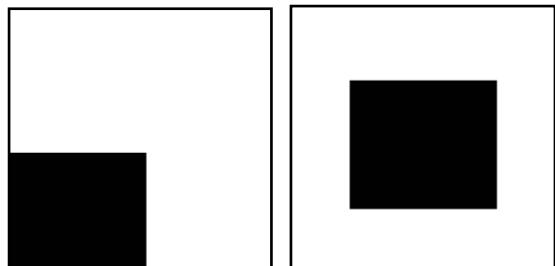
$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



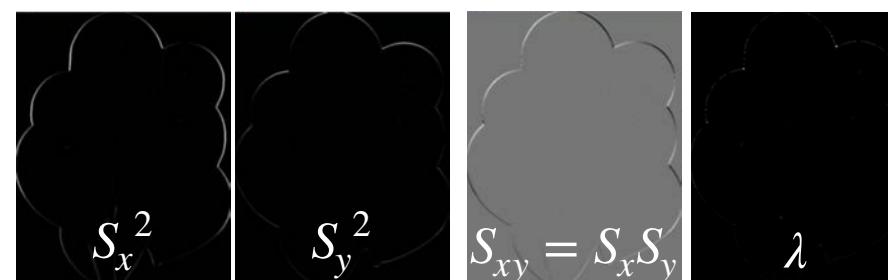
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
 - Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



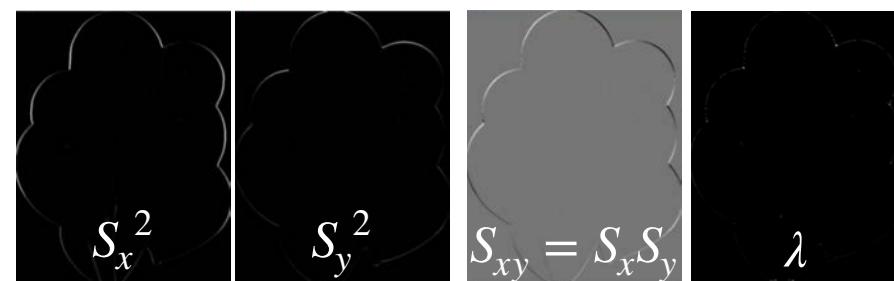
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
 - Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge
 - Step 4 (Corner Detection): A point is identified if the eigenvalues or their approximation are large (meaning the local image patch changes significantly in all directions)
 - 6. Post-processing:
Non-maximal suppression to keep only the strongest corners



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



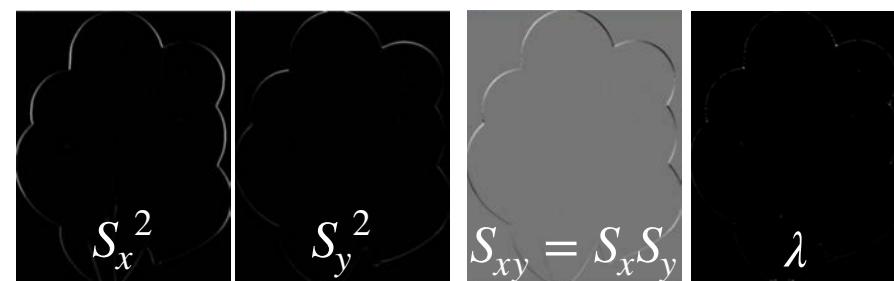
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_xS_y \\ S_xS_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_xS_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
 - Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge
 - Step 4 (Corner Detection): A point is identified if the eigenvalues or their approximation are large (meaning the local image patch changes significantly in all directions)
 - 6. Post-processing:
Non-maximal suppression to keep only the strongest corners



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector



Original



Shi-Tomasi

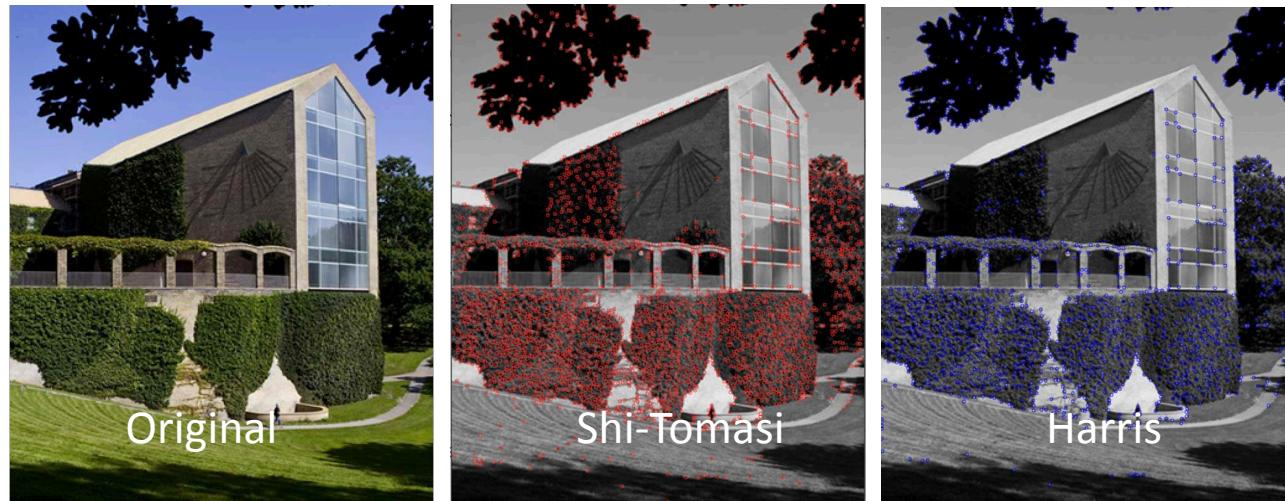


Harris

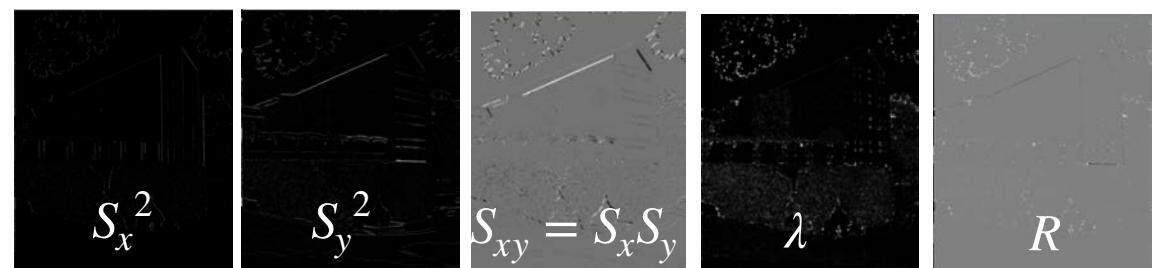
Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection

- Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
- Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
- Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge
- Step 4 (Corner Detection): A point is identified if the eigenvalues or their approximation are large (meaning the local image patch changes significantly in all directions)
- 6. Post-processing:
Non-maximal suppression to keep only the strongest corners



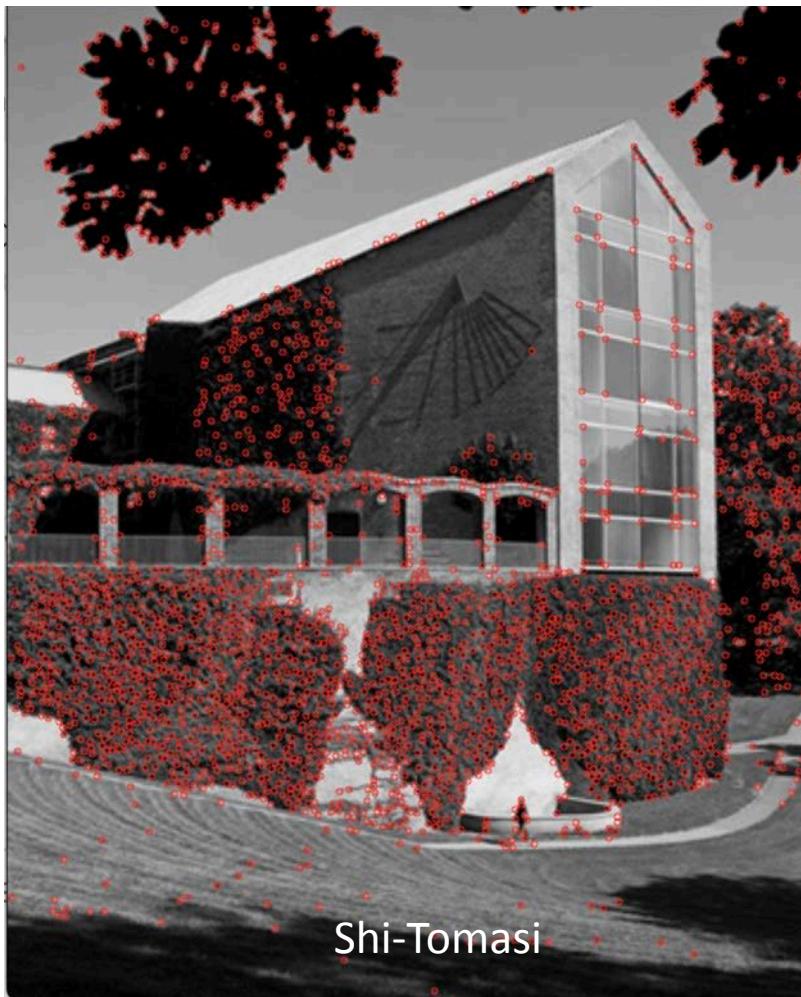
$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

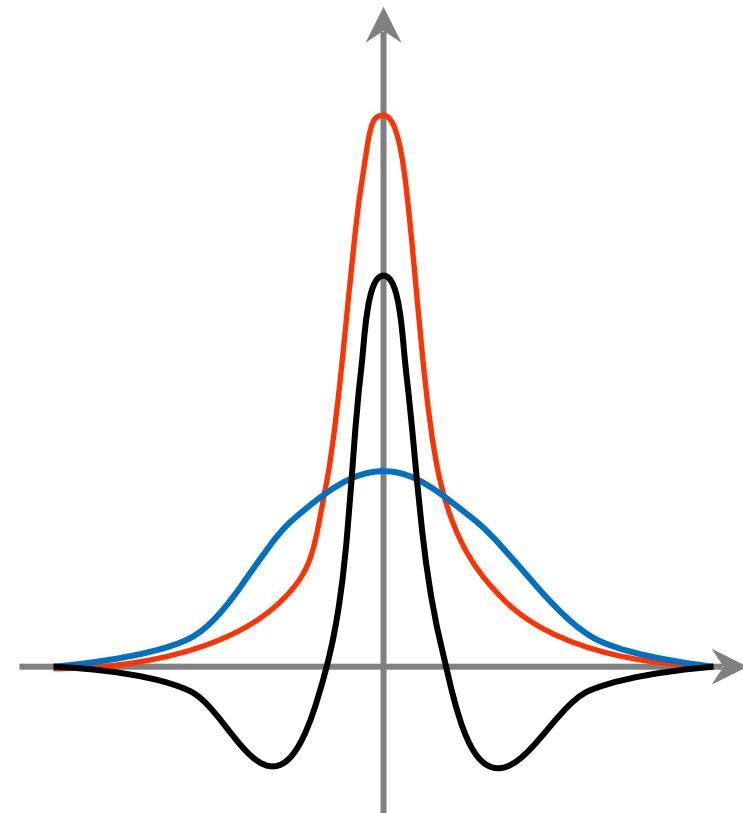
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector



Blob Detection - Difference of Gaussians (DoG)

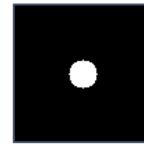
- ‘Blobs’ are alternative features
 - A blob is a dark region surrounded by a bright region, or vice-versa
 - Blobs have a scale as well as a location
- Difference of Gaussian filter
 - Two Gaussians of different width
 - Subtract wide from narrow
 - Bright blobs – high positive response
 - Dark blobs – high negative response



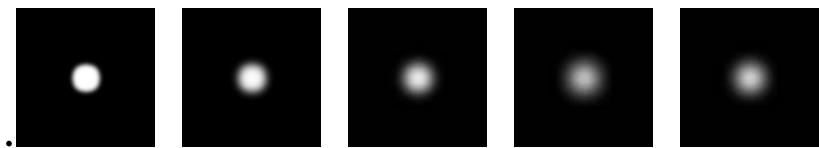
Blob Detection - Difference of Gaussians (DoG)

- ‘Blobs’ are alternative features
 - A blob is a dark region surrounded by a bright region, or vice-versa
 - Blobs have a scale as well as a location
- Difference of Gaussian filter
 - Two Gaussians of different width
 - Subtract wide from narrow
 - Bright blobs – high positive response
 - Dark blobs – high negative response

Original Image



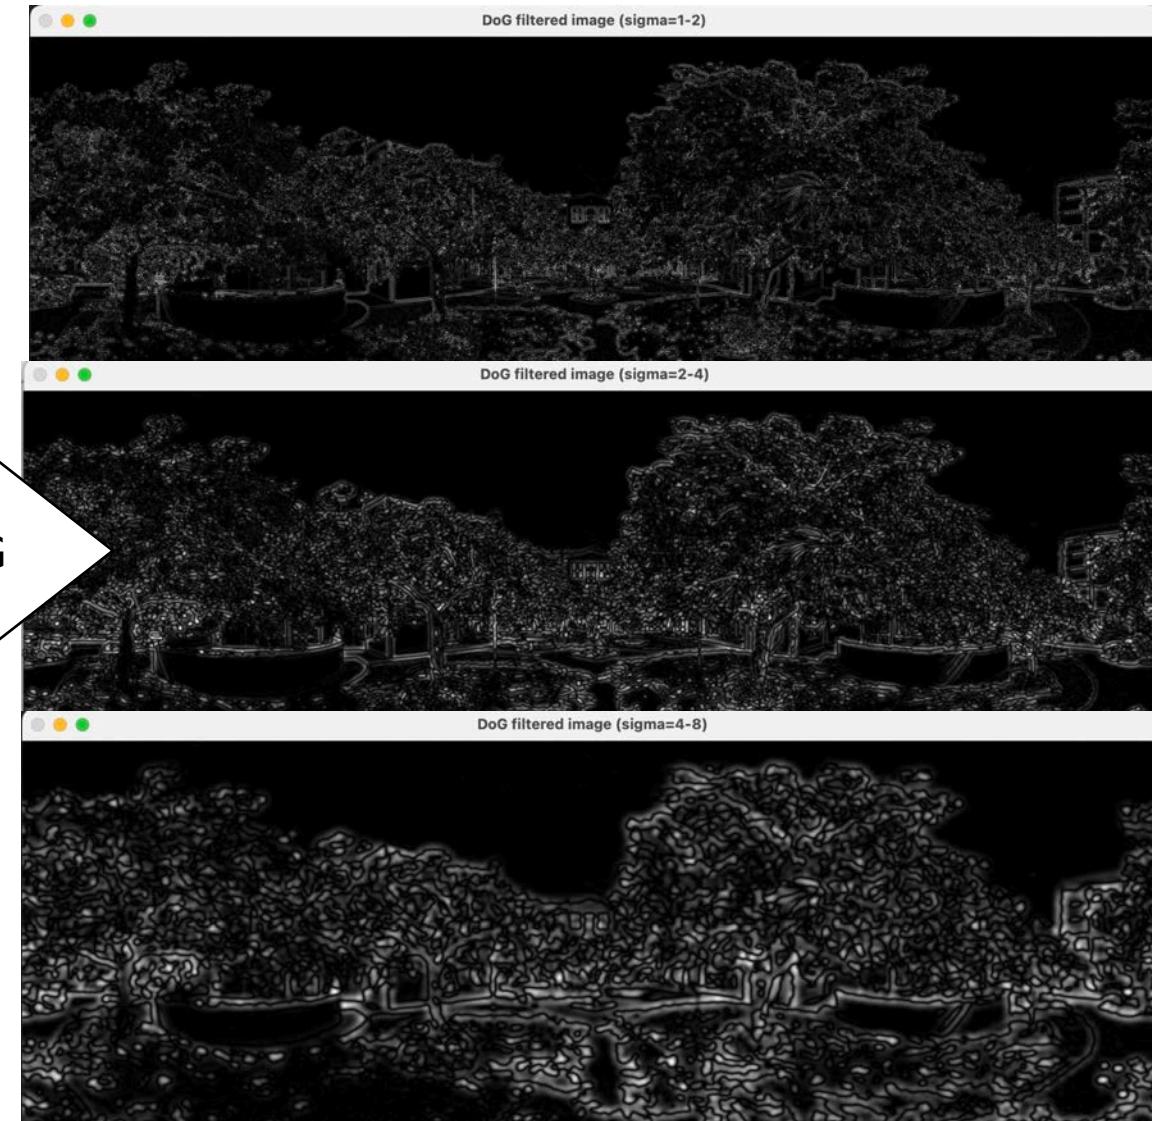
Gaussian blur
with
 $\sigma = 1.5, 3, 4.5, 6, \dots$



Difference of
Gaussians



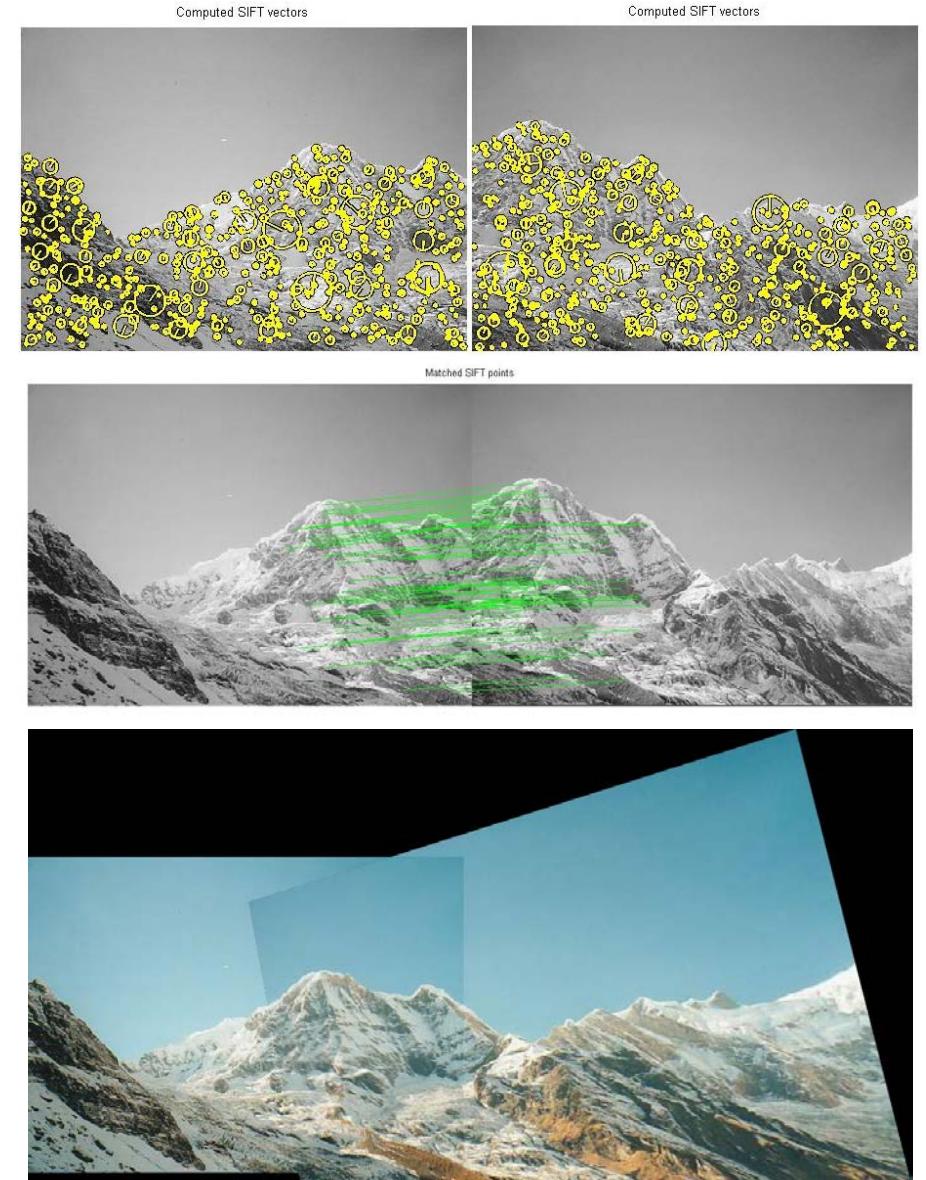
Blob Detection



**Lets put it all together -
Scale Invariant Features Transform (SIFT)**

Scale Invariant Features Transform (SIFT)

- Common problems:
 - Panorama stitching
 - Image Search
 - Object recognition
 - Tracking (covered a different time)
- Scale Invariant Features Transform (SIFT):
 - Computer vision algorithm to
 - Detect,
 - Describe,
 - And match local features in images



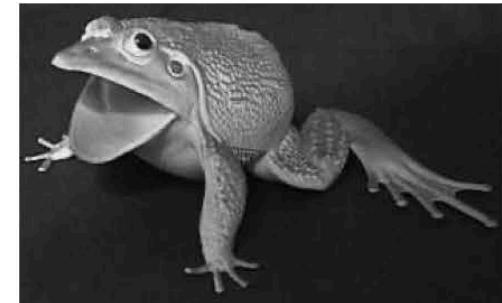
Scale Invariant Features Transform (SIFT)

- Common problems:
 - Panorama stitching
 - Image Search
 - Object recognition
 - Tracking (covered a different time)
- Scale Invariant Features Transform (SIFT):
 - Computer vision algorithm to
 - Detect,
 - Describe,
 - And match local features in images



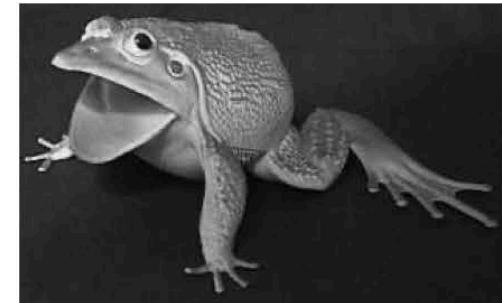
Scale Invariant Features Transform (SIFT)

- Common problems:
 - Panorama stitching
 - Image Search
 - Object recognition
 - Tracking (covered a different time)
- Scale Invariant Features Transform (SIFT):
 - Computer vision algorithm to
 - Detect,
 - Describe,
 - And match local features in images



Scale Invariant Features Transform (SIFT)

- Common problems:
 - Panorama stitching
 - Image Search
 - Object recognition
 - Tracking (covered a different time)
- Scale Invariant Features Transform (SIFT):
 - Computer vision algorithm to
 - Detect,
 - Describe,
 - And match local features in images



Scale Invariant Features Transform (SIFT)

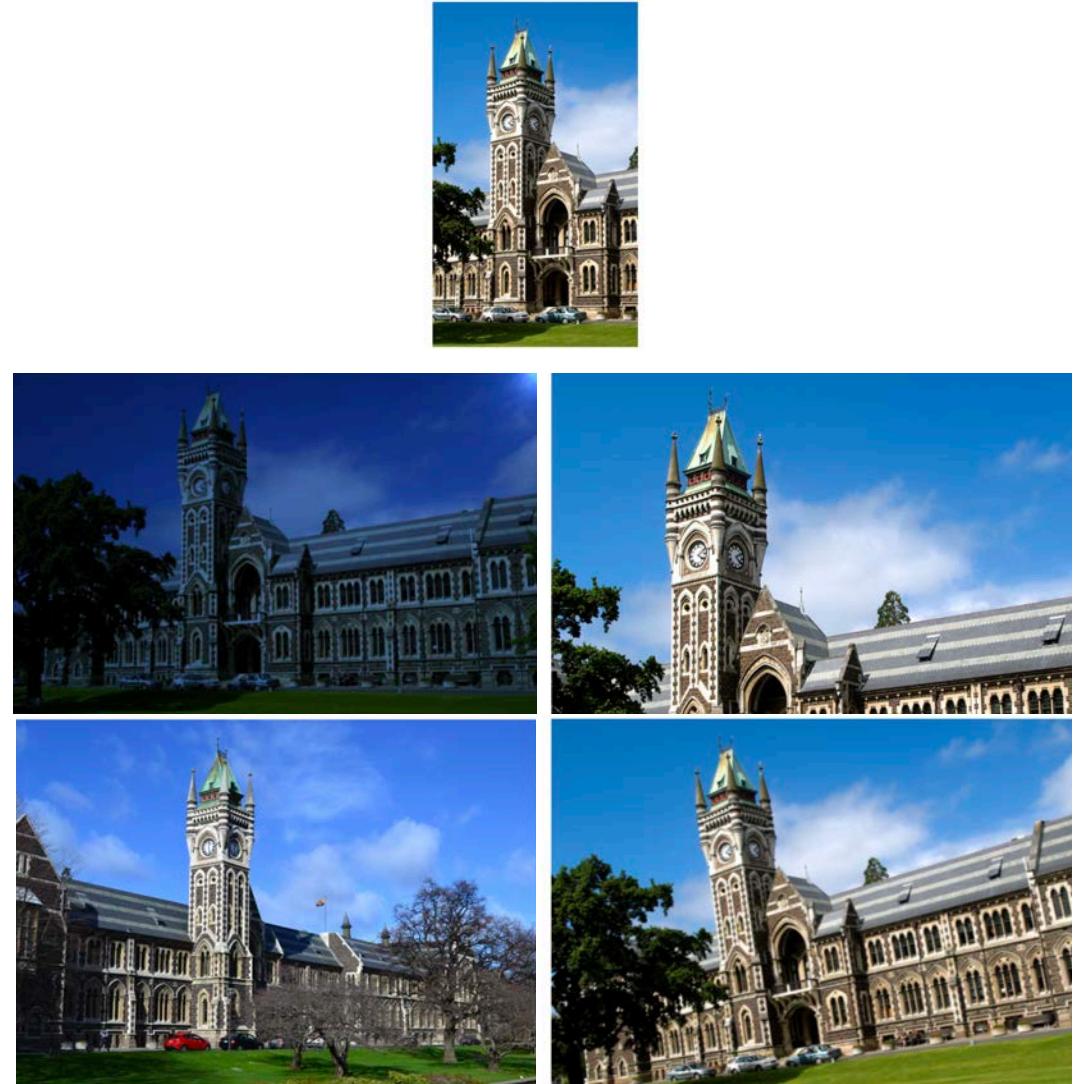
- Common problems:
 - Panorama stitching
 - Image Search
 - Object recognition
 - Tracking (covered a different time)
- Scale Invariant Features Transform (SIFT):
 - Computer vision algorithm to
 - Detect,
 - Describe,
 - And match local features in images



Scale Invariant Features Transform (SIFT)

- Properties of SIFT

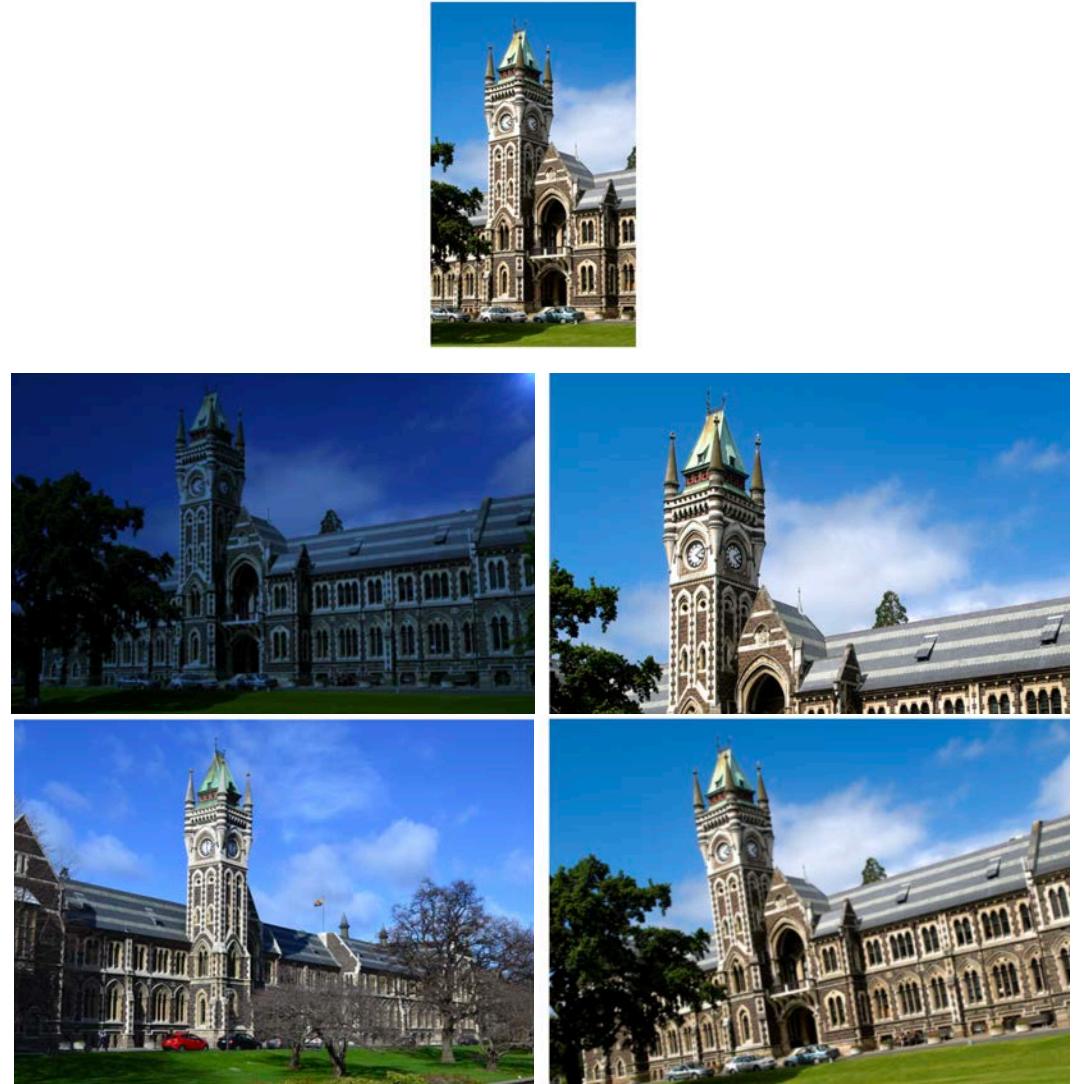
- Robustness:** Invariance to changes in illumination, scale, rotation, affine, perspective.
- Locality:** robustness to occlusion and clutter.
- Distinctiveness:** Easy to match to a large database of objects (e.g. different images).
- Quantity:** Many features (points) can be generated for even small objects.
- Efficiency:** Computationally “cheap”, real-time performance.



Scale Invariant Features Transform (SIFT)

- Properties of SIFT

- Robustness:** Invariance to changes in illumination, scale, rotation, affine, perspective.
- Locality:** robustness to occlusion and clutter.
- Distinctiveness:** Easy to match to a large database of objects (e.g. different images).
- Quantity:** Many features (points) can be generated for even small objects.
- Efficiency:** Computationally “cheap”, real-time performance.



SIFT - Feature detection (blobs)

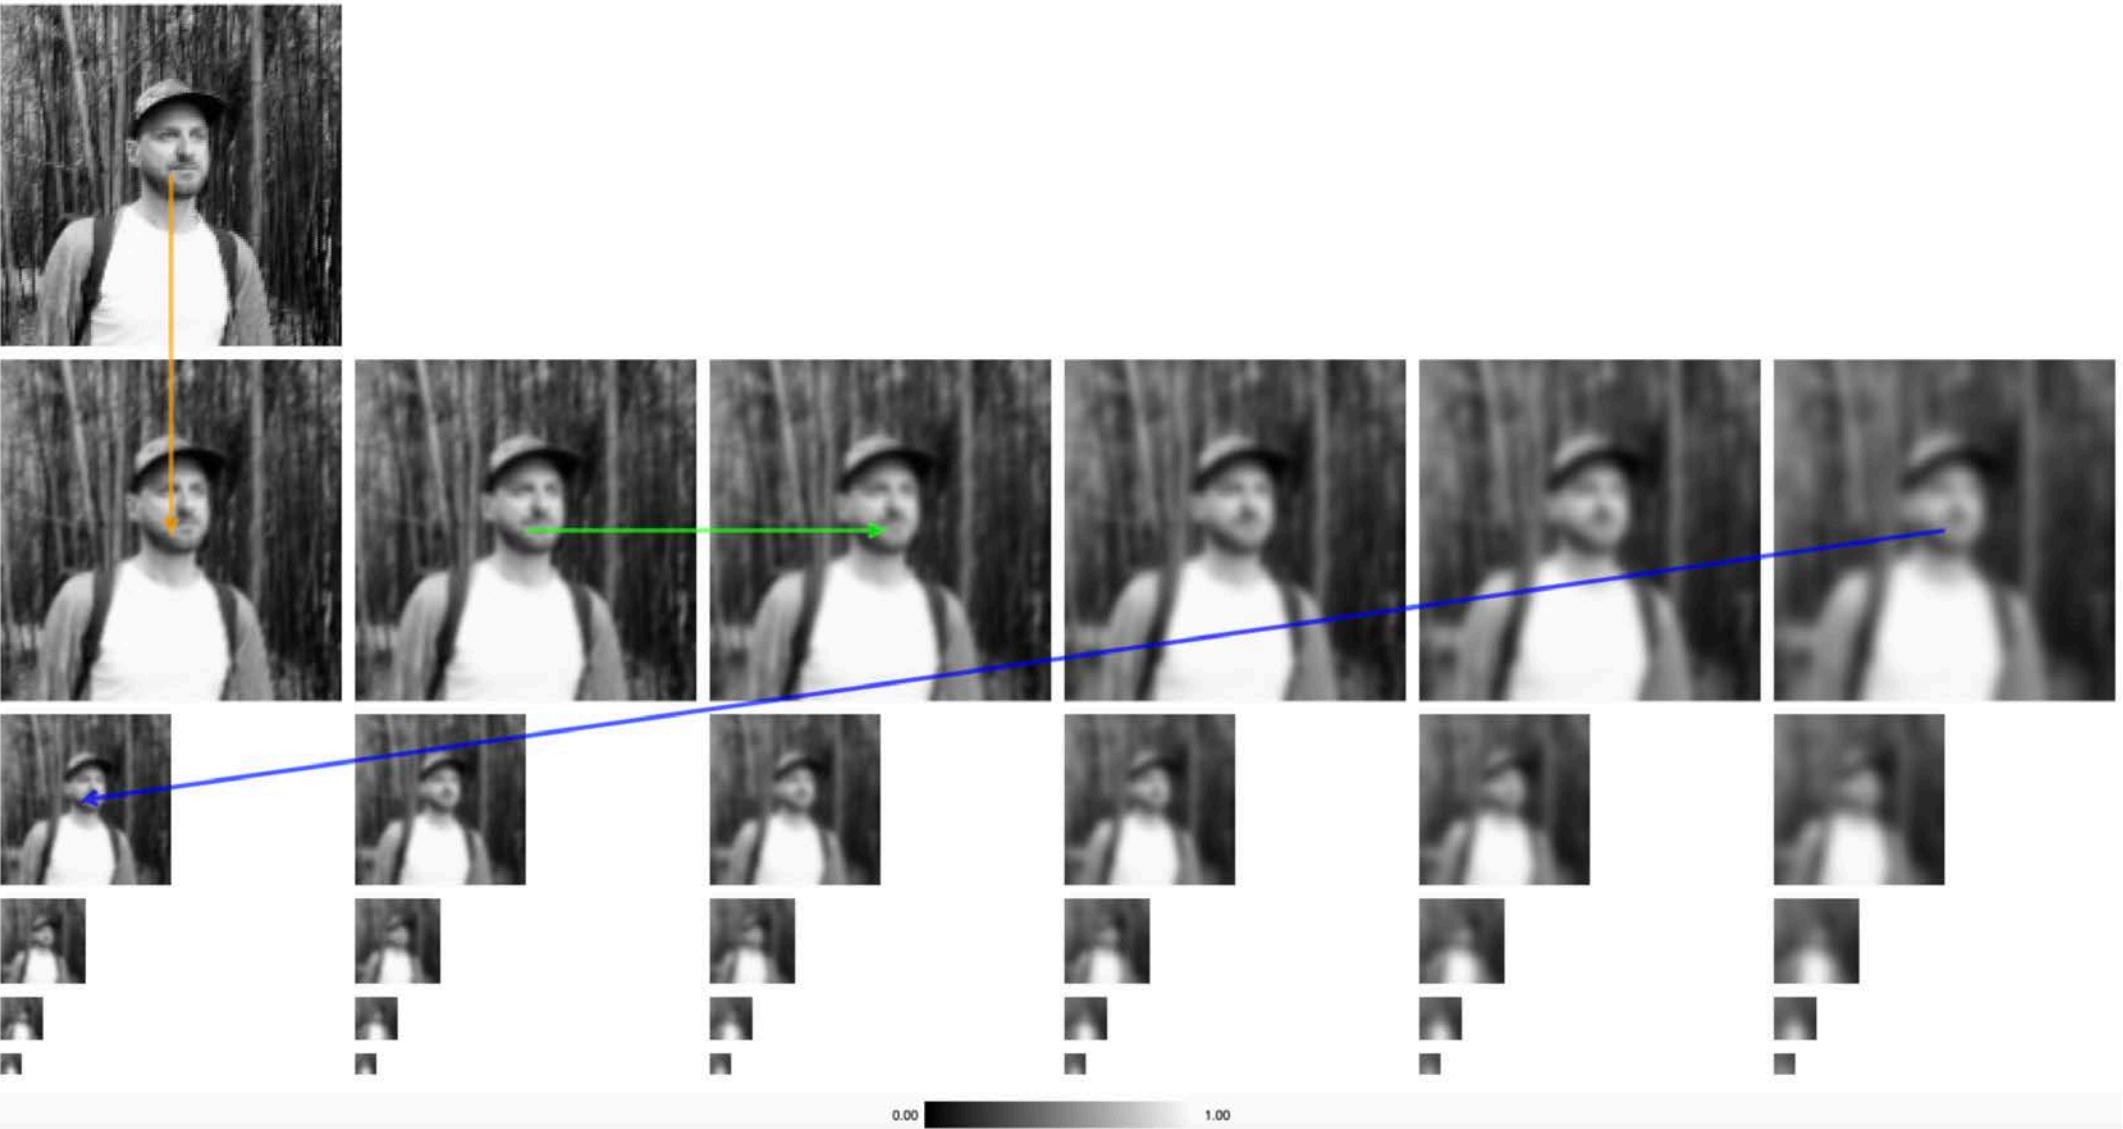
Scale Invariant Features Transform (SIFT)

- Properties of SIFT
 - **Robustness:** Invariance to changes in illumination, scale, rotation, affine, perspective.
 - **Locality:** robustness to occlusion and clutter.
 - **Distinctiveness:** Easy to match to a large database of objects (e.g. different images).
 - **Quantity:** Many features (points) can be generated for even small objects.
 - **Efficiency:** Computationally “cheap”, real-time performance.

SIFT - Scalespace creation

Input image
128x128

Normalised
Doubled input size
using
bilinear_interpolation



Uses some material from : <https://weitz.de/sift/>

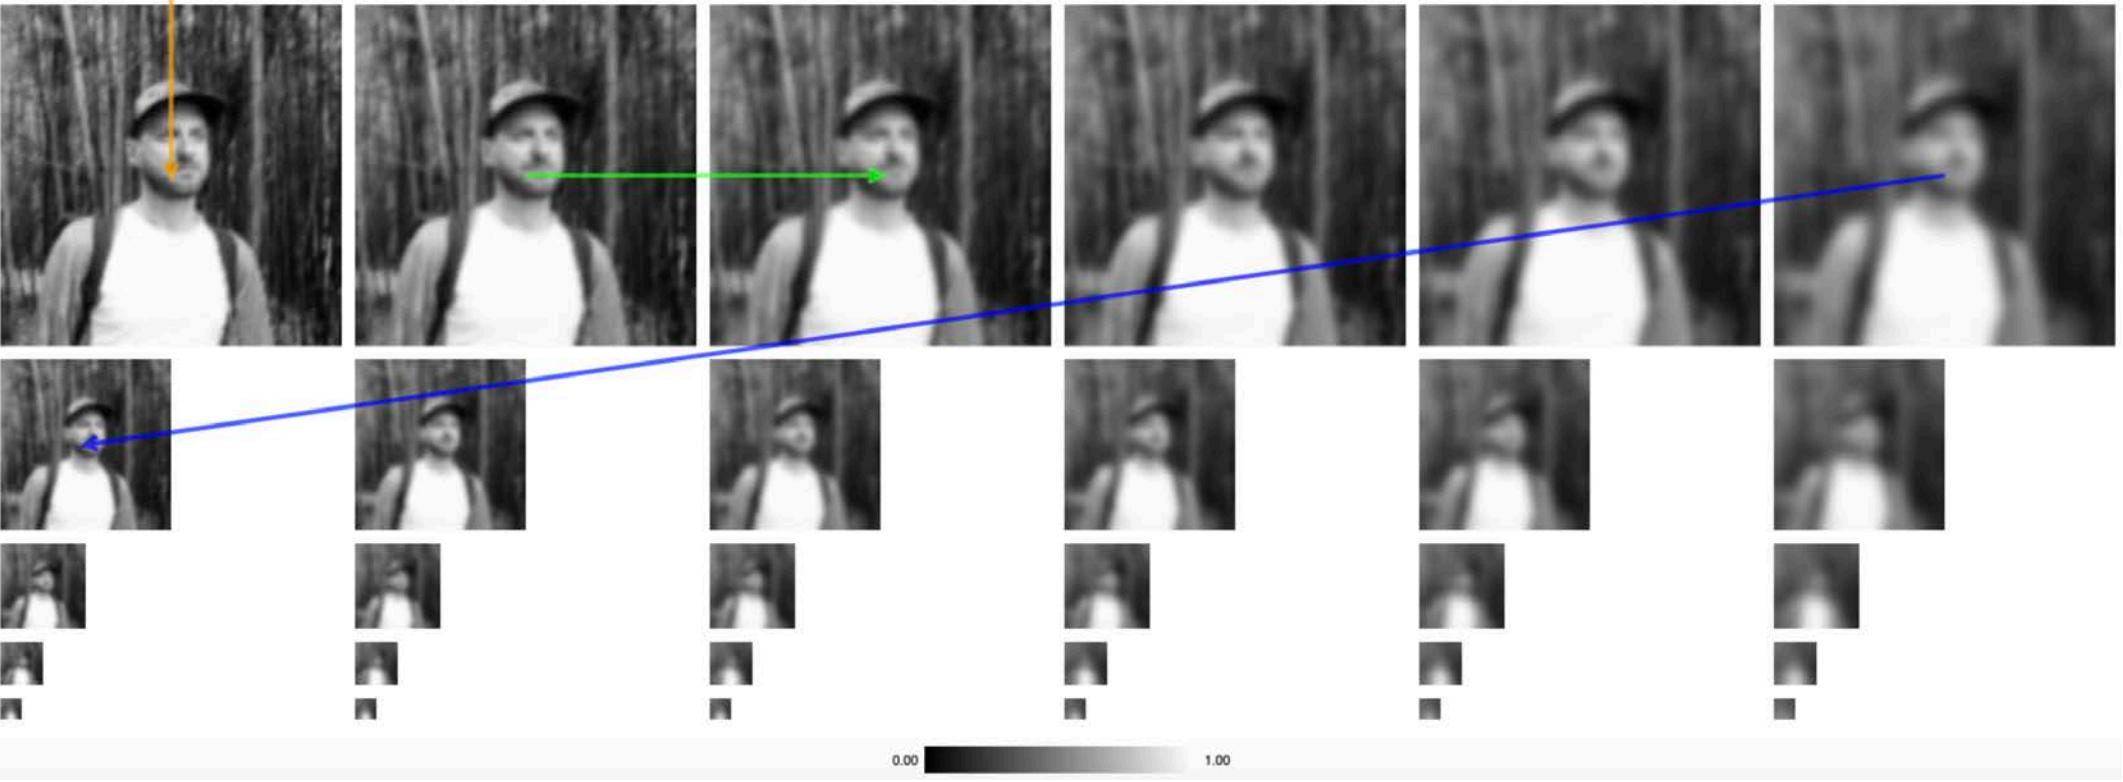
SIFT - Scalespace creation

Input image
128x128

Normalised
Doubled input size
using
bilinear_interpolation



Blurring using a
Gaussian convolution



SIFT - Scalespace creation

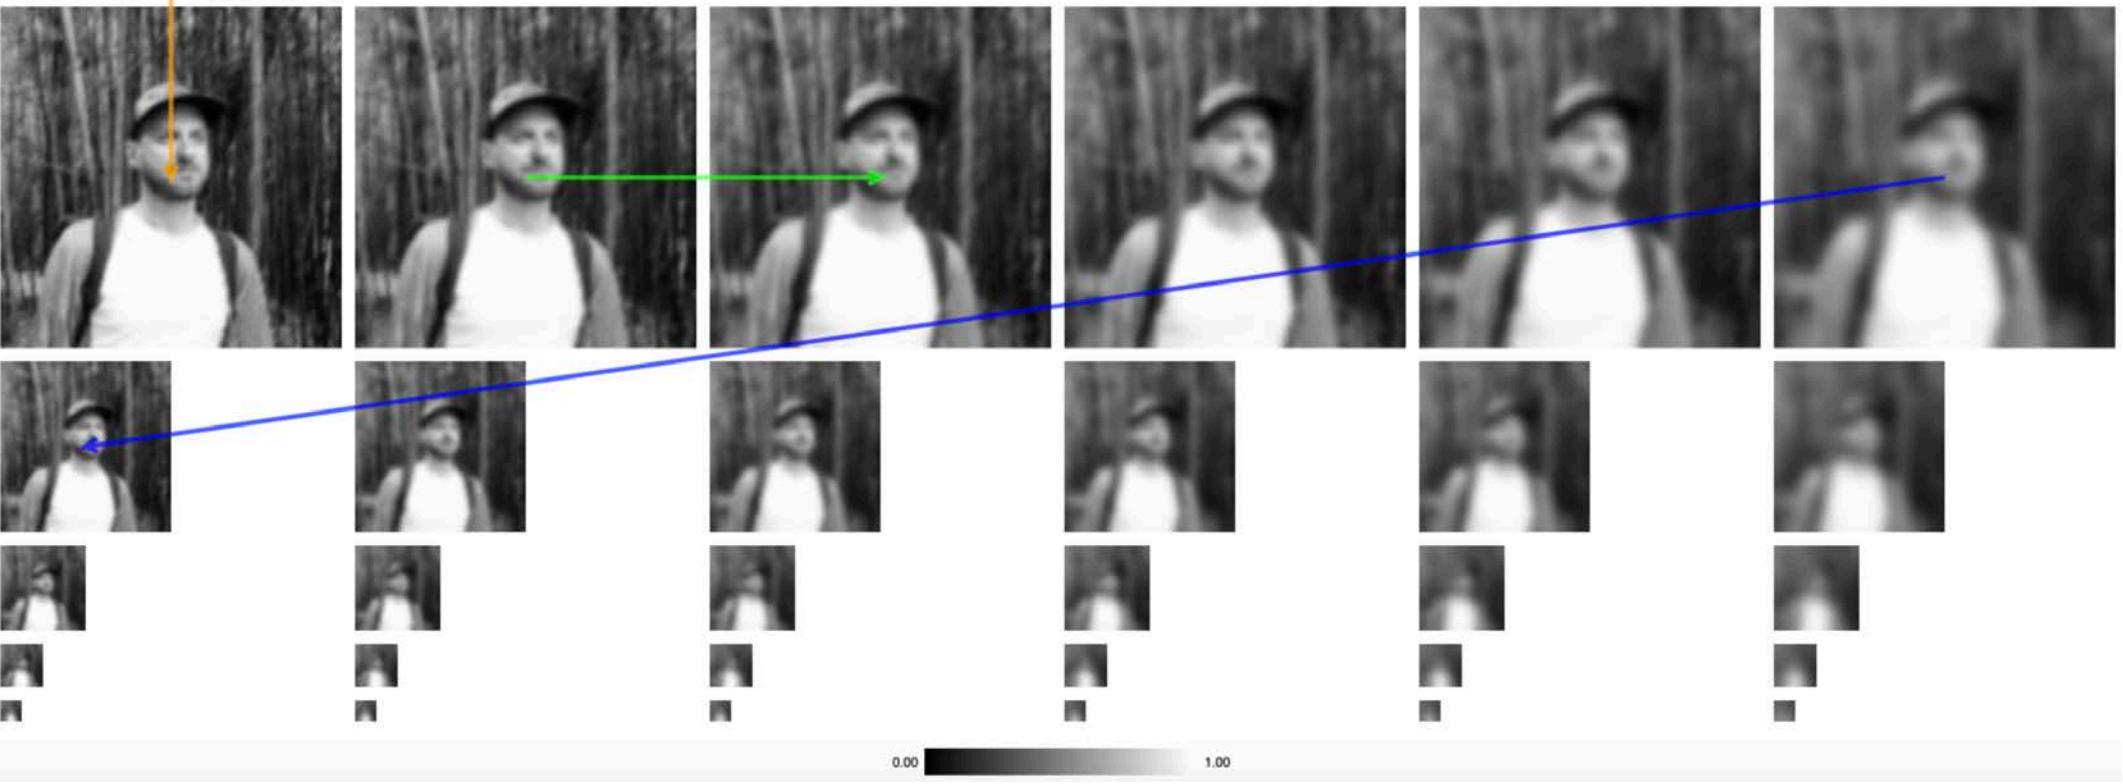
Input image
128x128

Normalised
Doubled input size
using
bilinear_interpolation



Blurring using a
Gaussian convolution

Blurring using a
Gaussian convolution
with increasing
standard deviation
(more blur)



SIFT - Scalespace creation

Input image
128x128

Normalised
Doubled input size
using
bilinear_interpolation



Blurring using a
Gaussian convolution



Blurring using a
Gaussian convolution
with increasing
standard deviation
(more blur)



Downsampling (size
reduction)



0.00 1.00

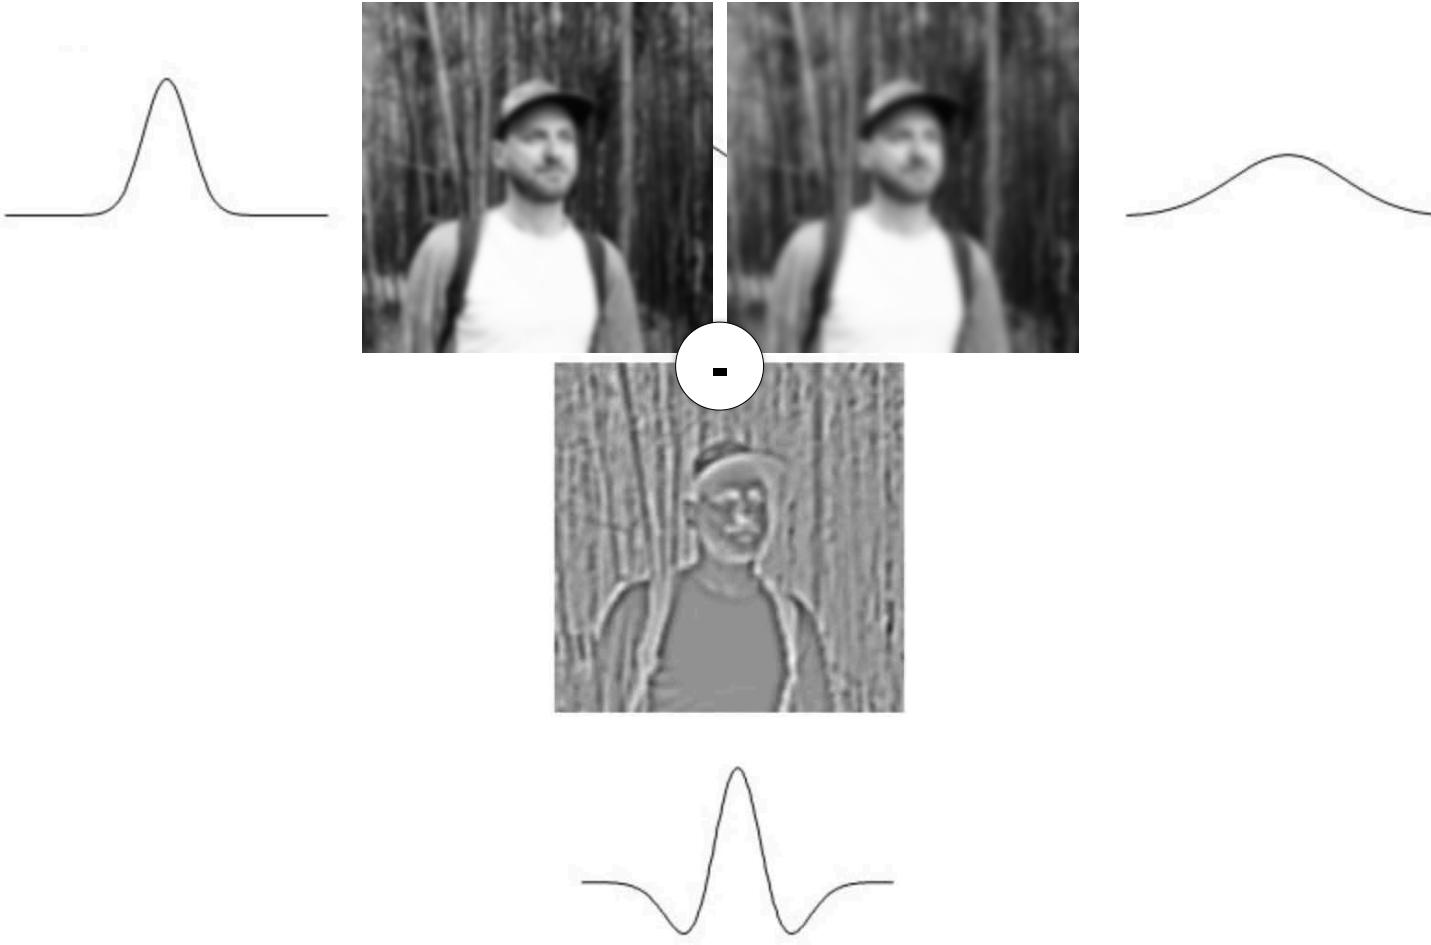
SIFT - Difference of Gaussians



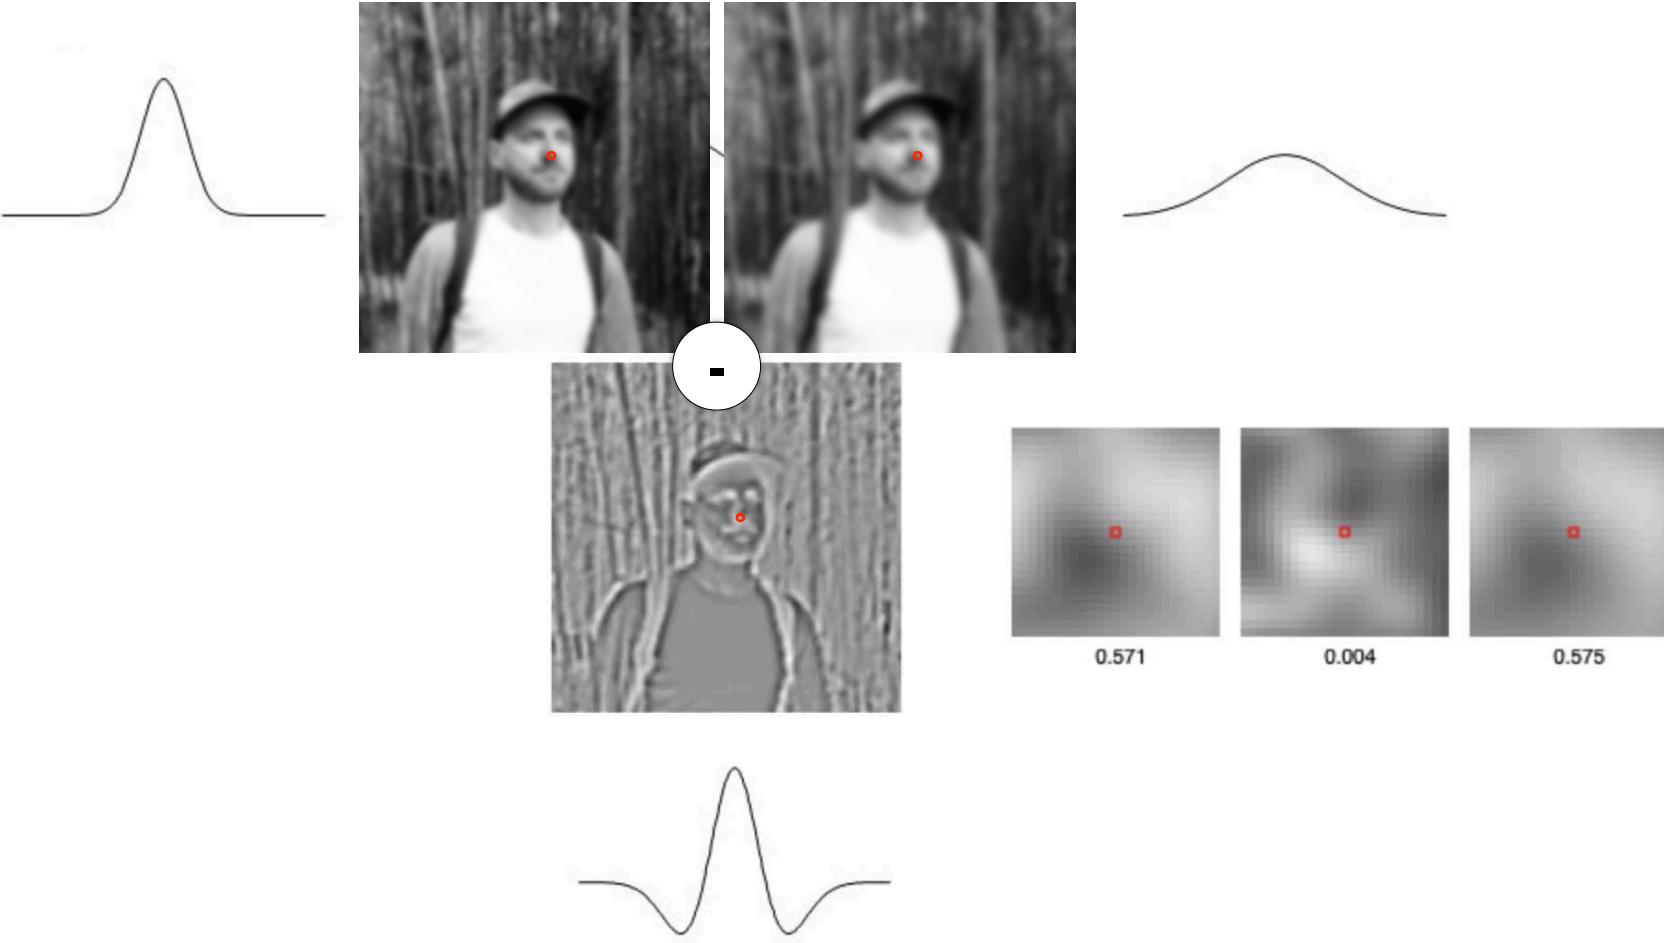
SIFT - Difference of Gaussians



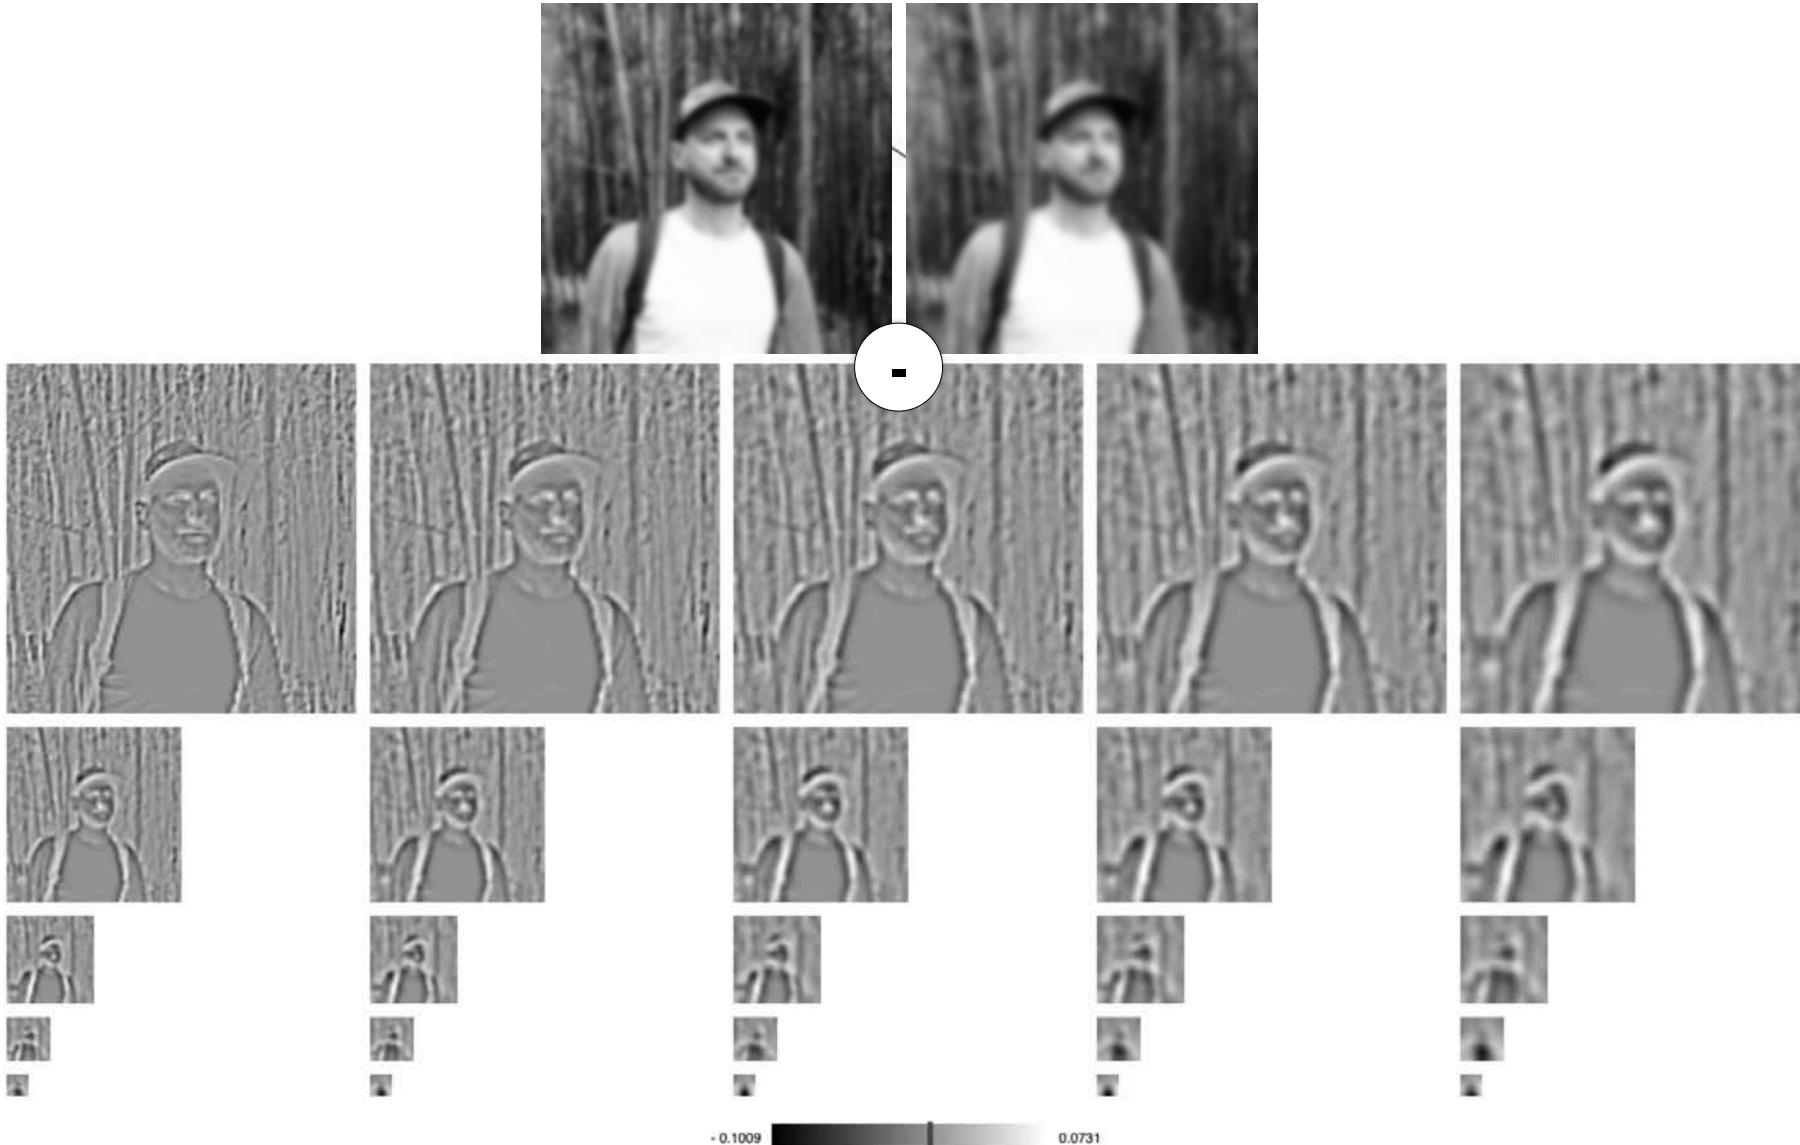
SIFT - Difference of Gaussians



SIFT - Difference of Gaussians

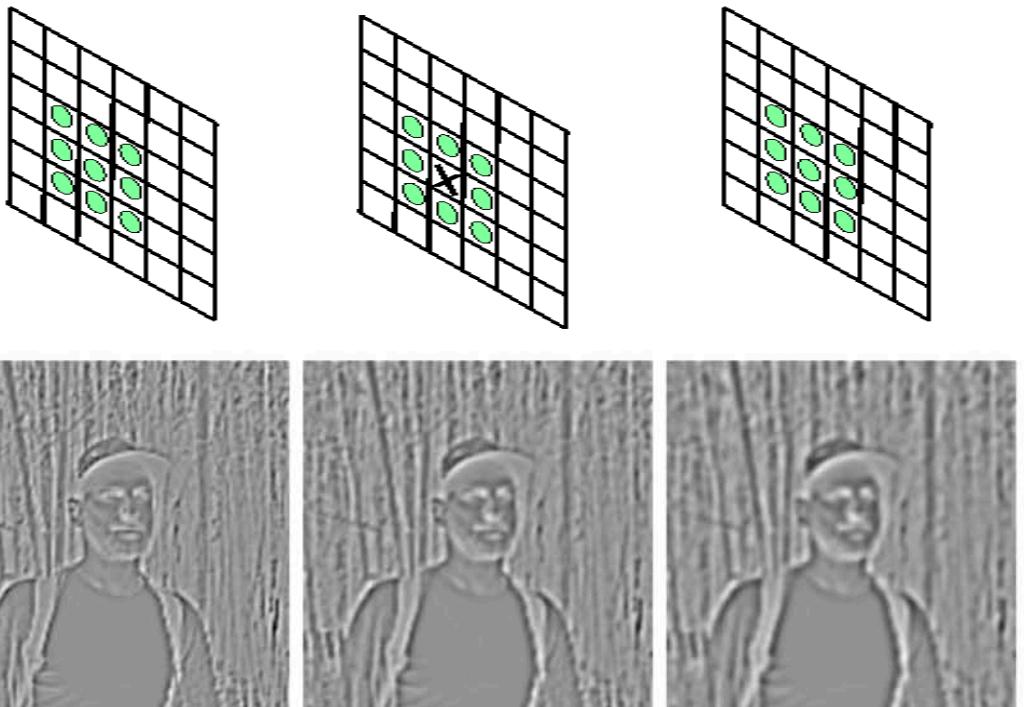


SIFT - Difference of Gaussians

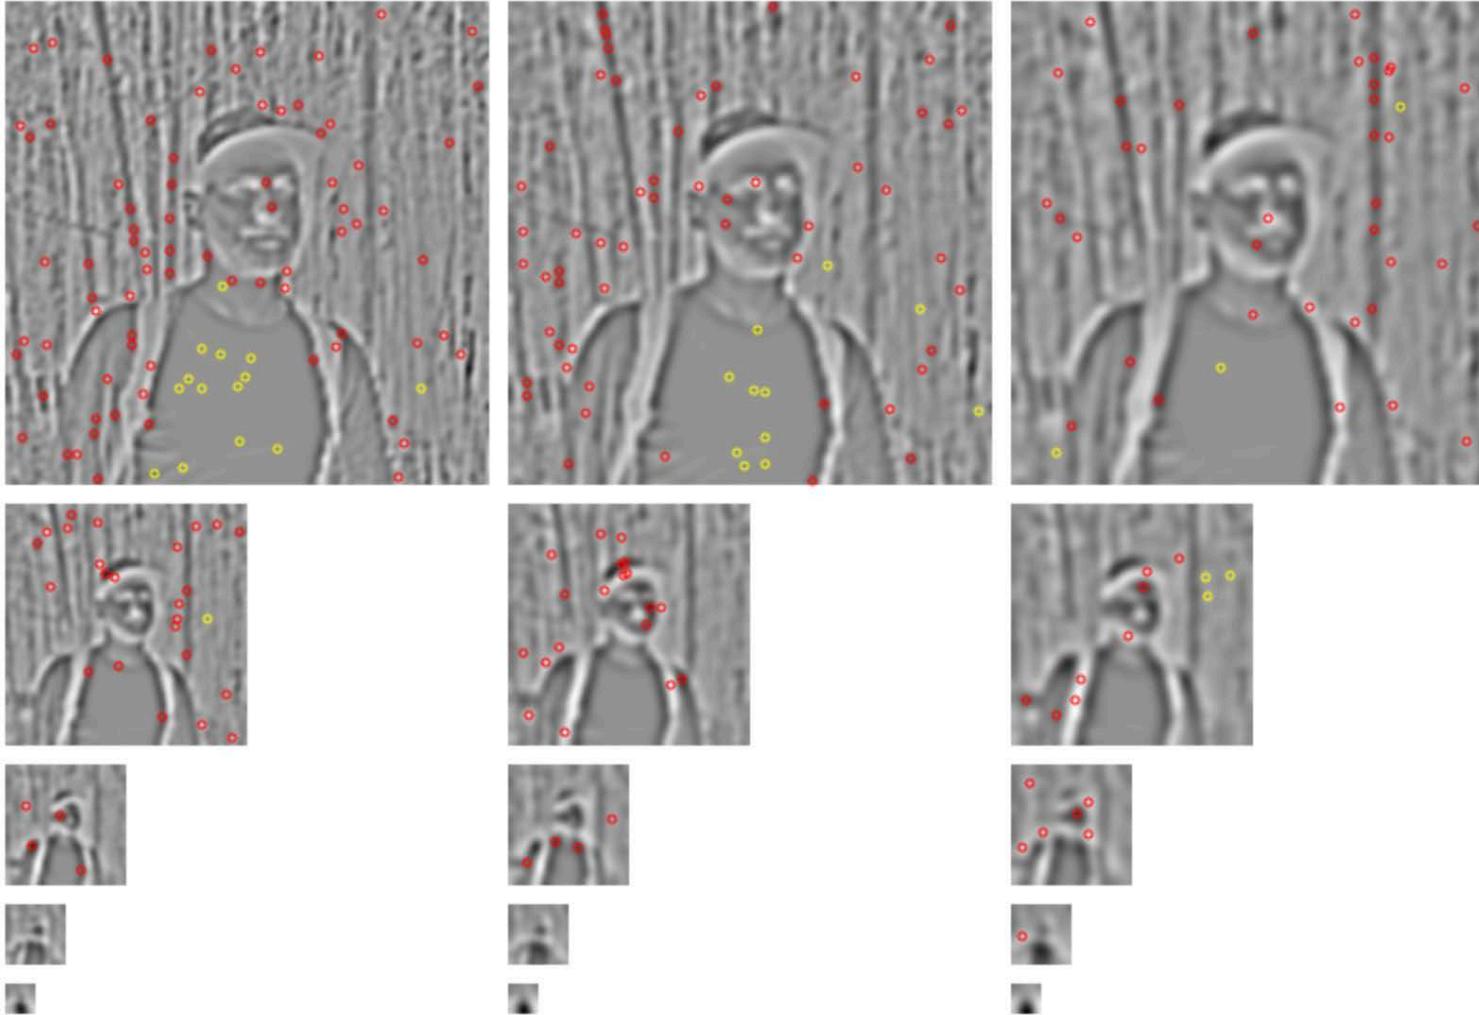


SIFT - Discrete extrema

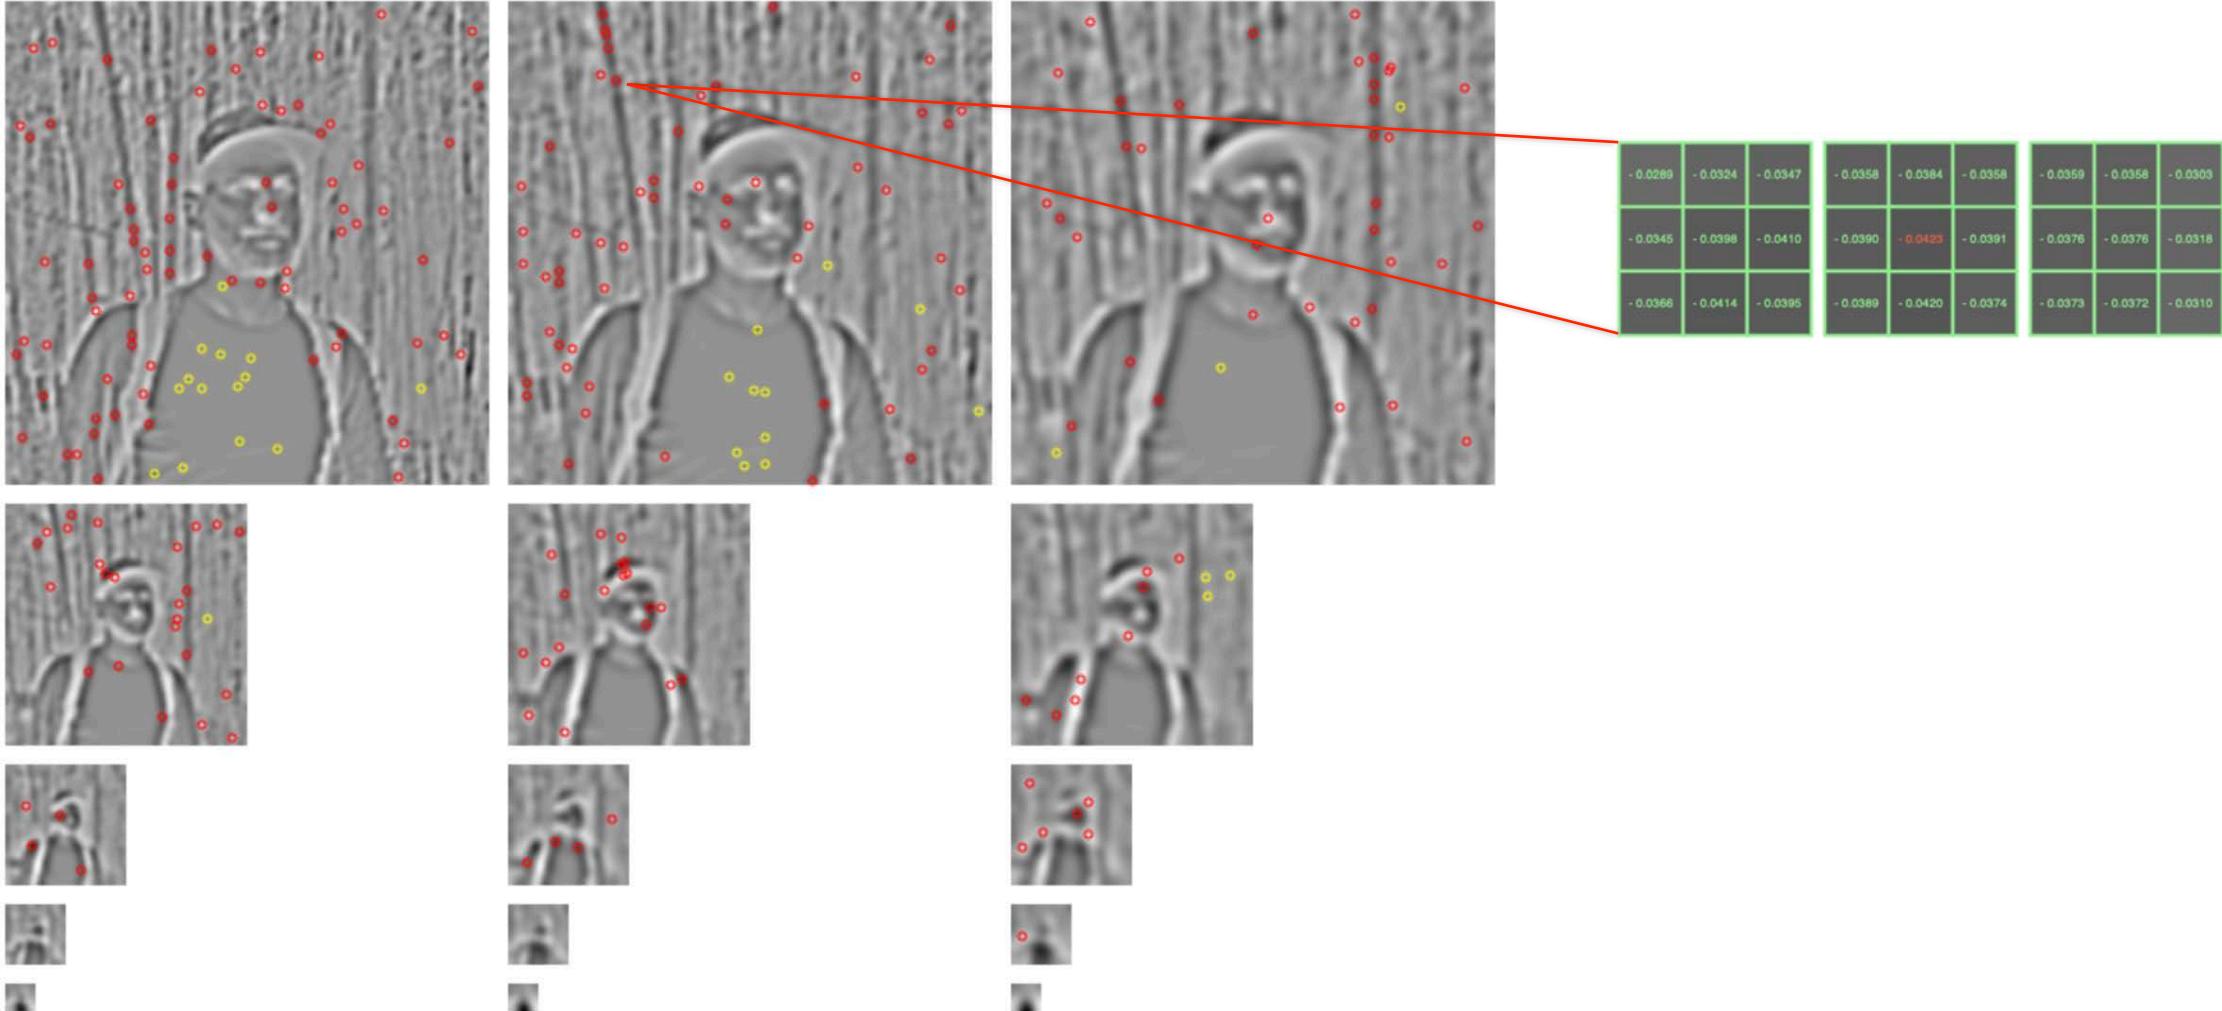
- Pixel whose gray value is larger than those of all of its 26 neighbor pixel
(Discrete Maximum)
- Pixel whose gray value is smaller than those of all of its 26 neighbor pixel
(Discrete minimum)



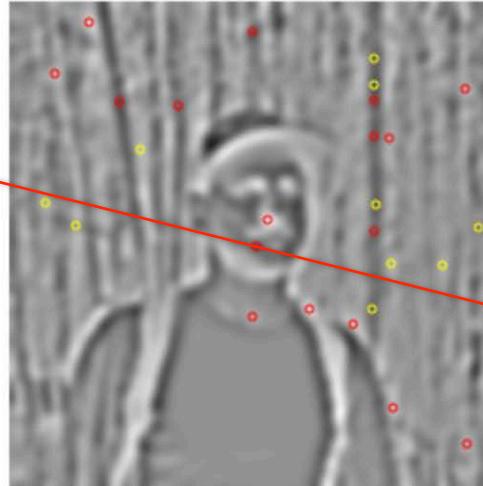
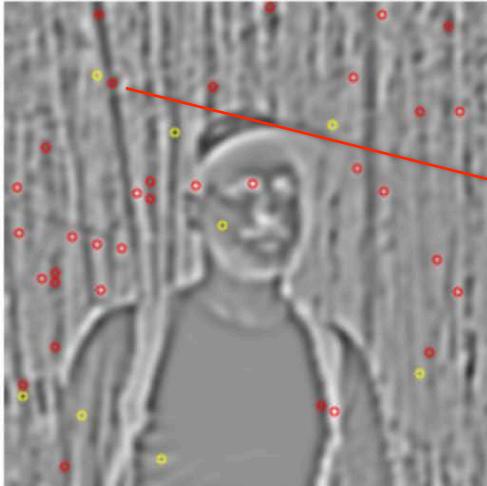
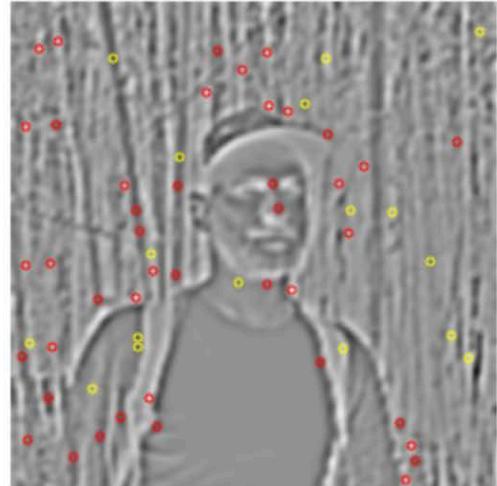
SIFT - Discrete extrema



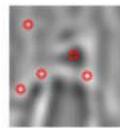
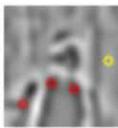
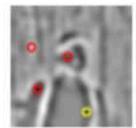
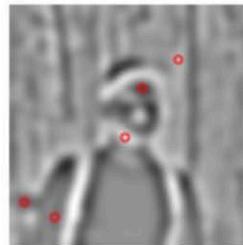
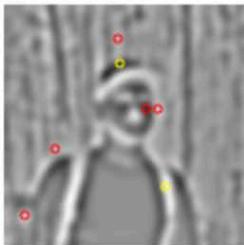
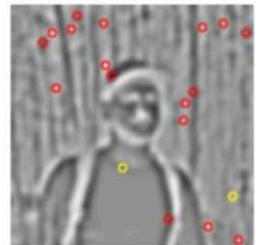
SIFT - Discrete extrema



SIFT - Discrete extrema



	discrete	interpolated
x	28.5	28.571
y	21.0	21.278
<i>scale</i>	2	1.641

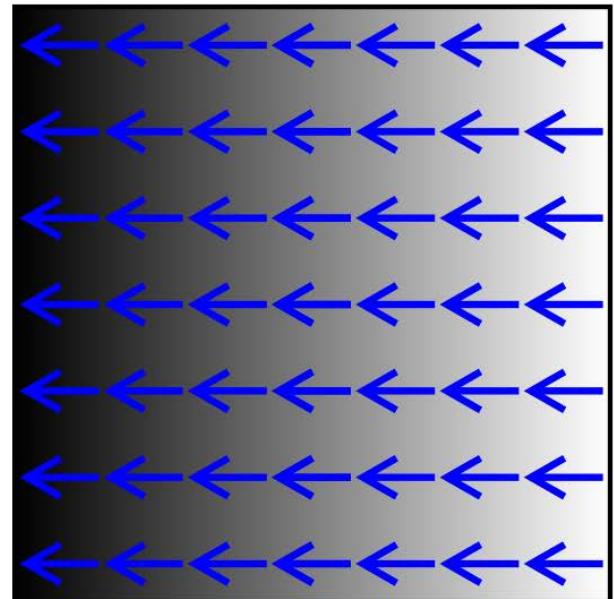
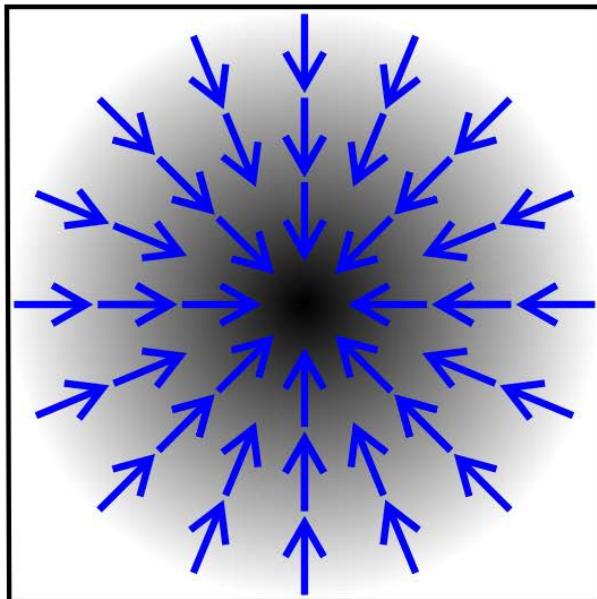


- We discard points (yellow) with small absolute values (noise) or extrema on an edge

SIFT - Feature description

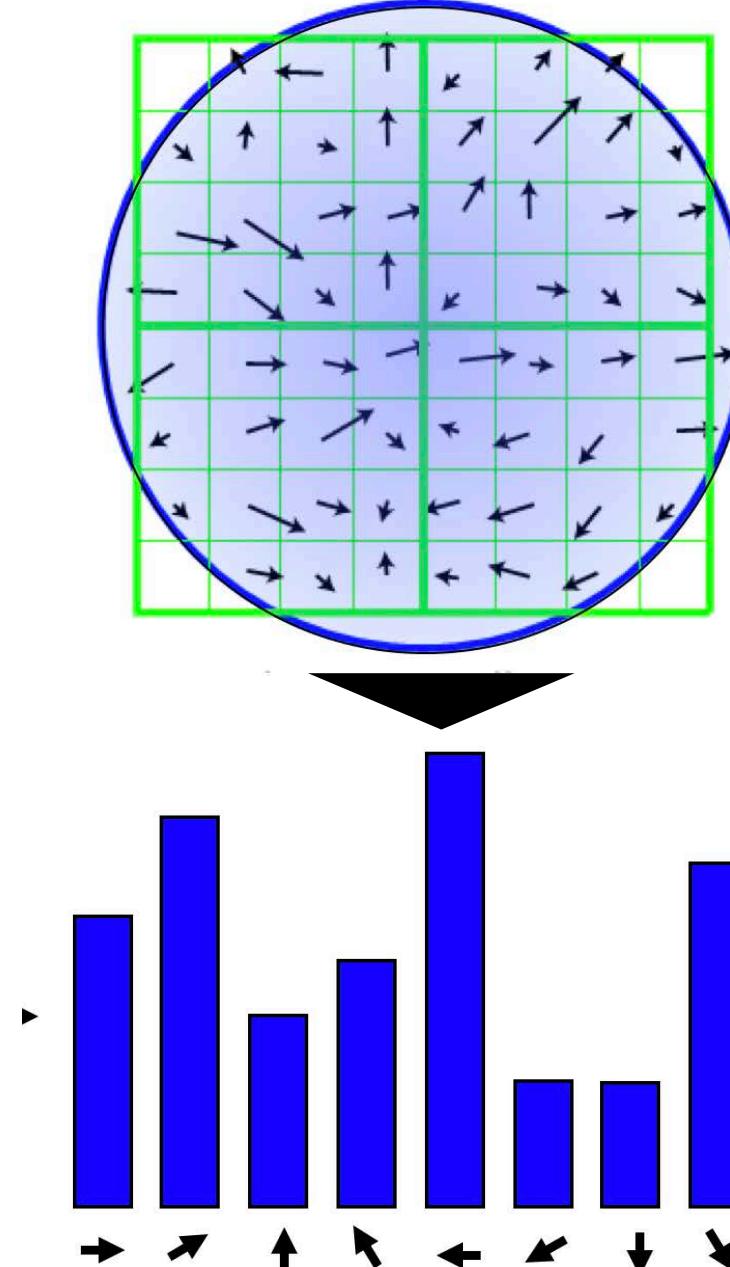
SIFT - Feature description

- Feature description (orientation and gradient)
- For each feature (corner) we compute gradient's magnitude and orientation
- Image gradients:



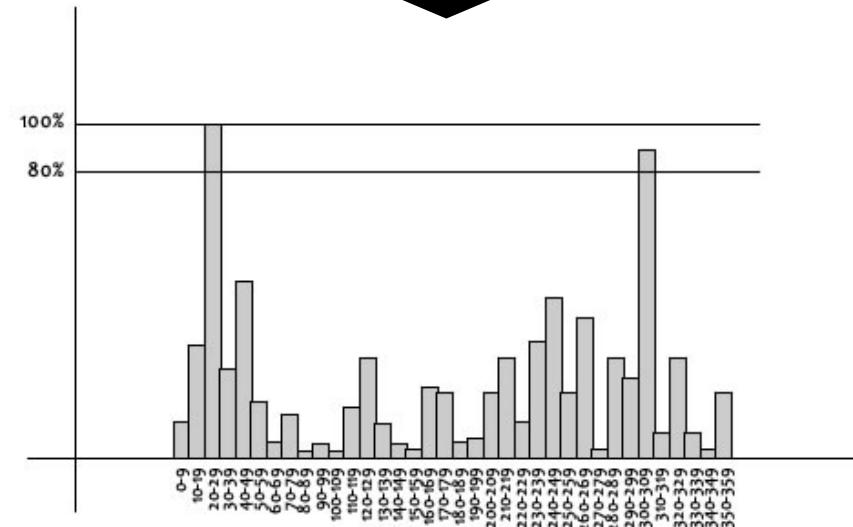
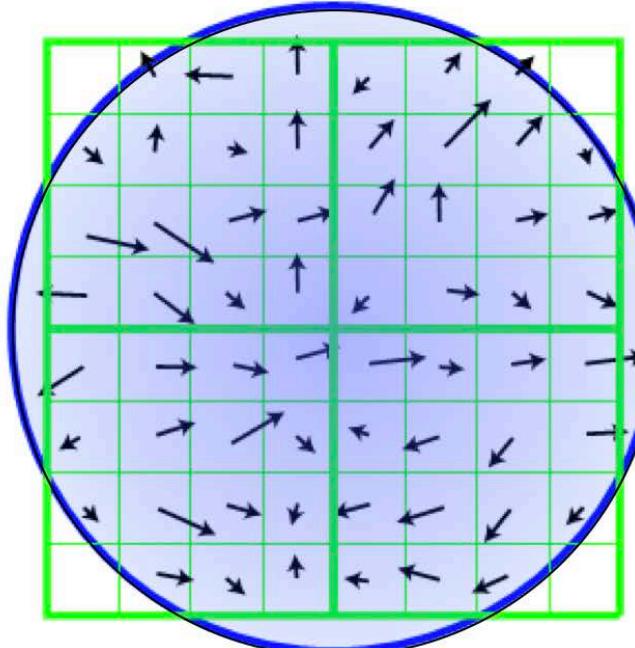
SIFT - Feature description

- Feature description (orientation and gradient)
- Create gradient histogram (8 bins, simplified) weighted by magnitude and Gaussian window
- Any histogram peak within 80% of highest peak is assigned to keypoint (multiple assignments possible)
- Result is keypoints/features main orientation

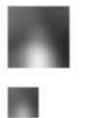
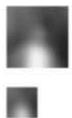
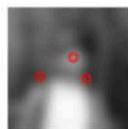
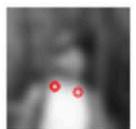
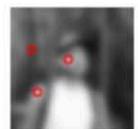
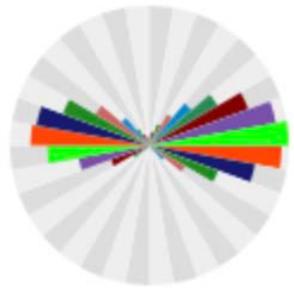
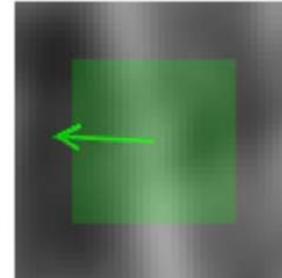
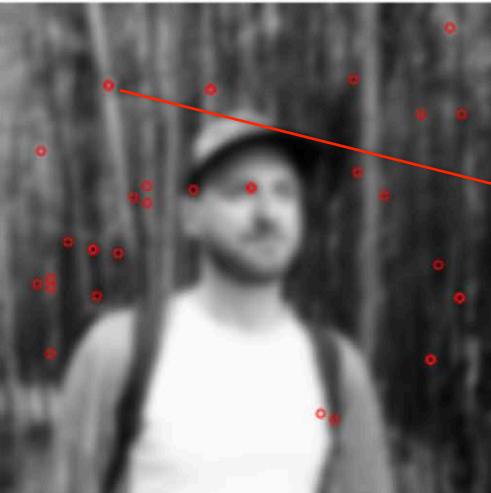


SIFT - Feature description

- Feature description (orientation and gradient)
- Create gradient histogram (8 bins, simplified) weighted by magnitude and Gaussian window
- Any histogram peak within 80% of highest peak is assigned to keypoint (multiple assignments possible)
- Result is keypoints/features main orientation

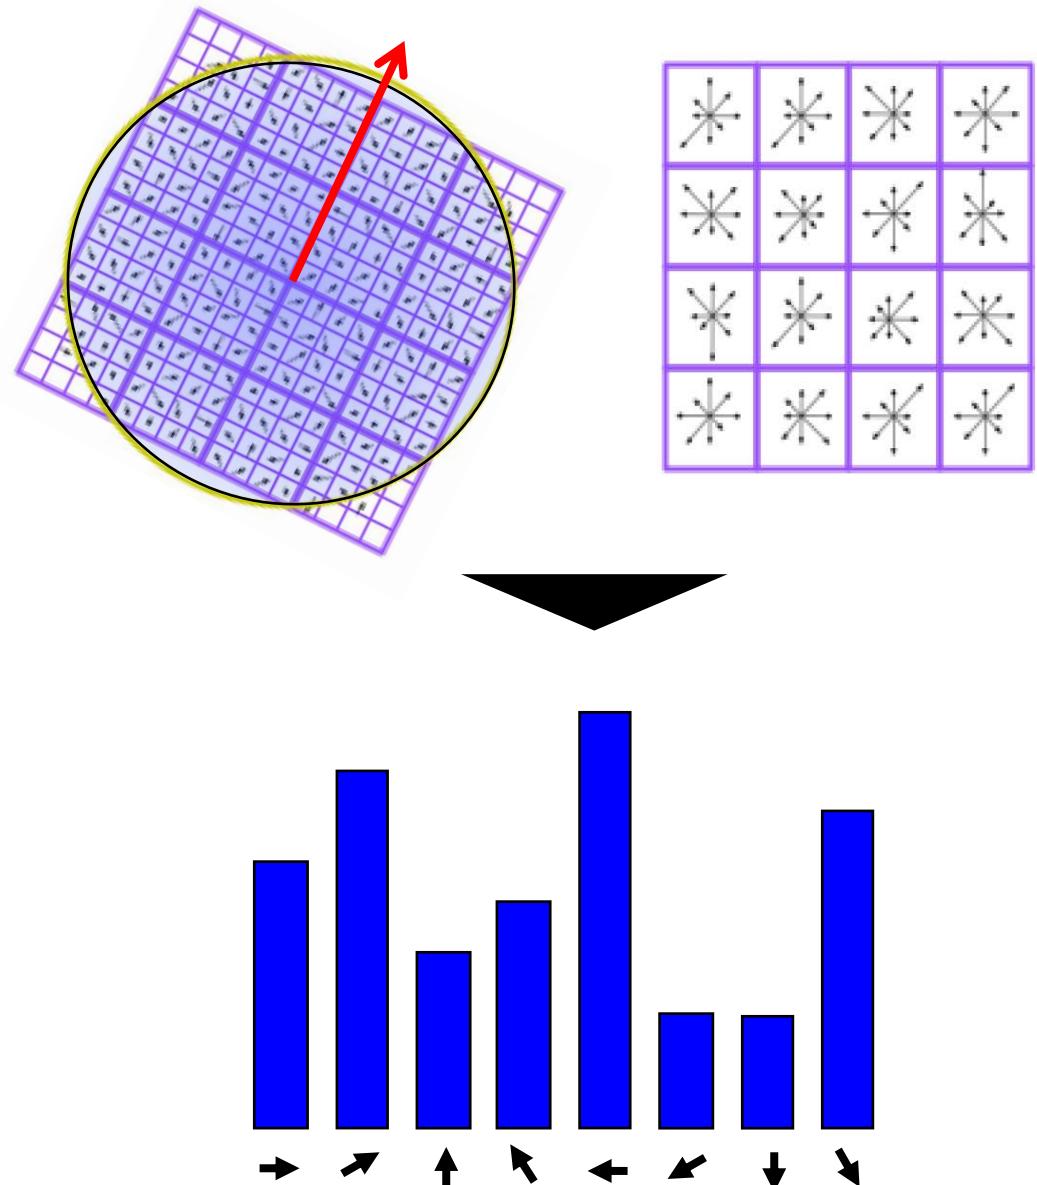


SIFT - Feature description



SIFT - Feature description from gradients

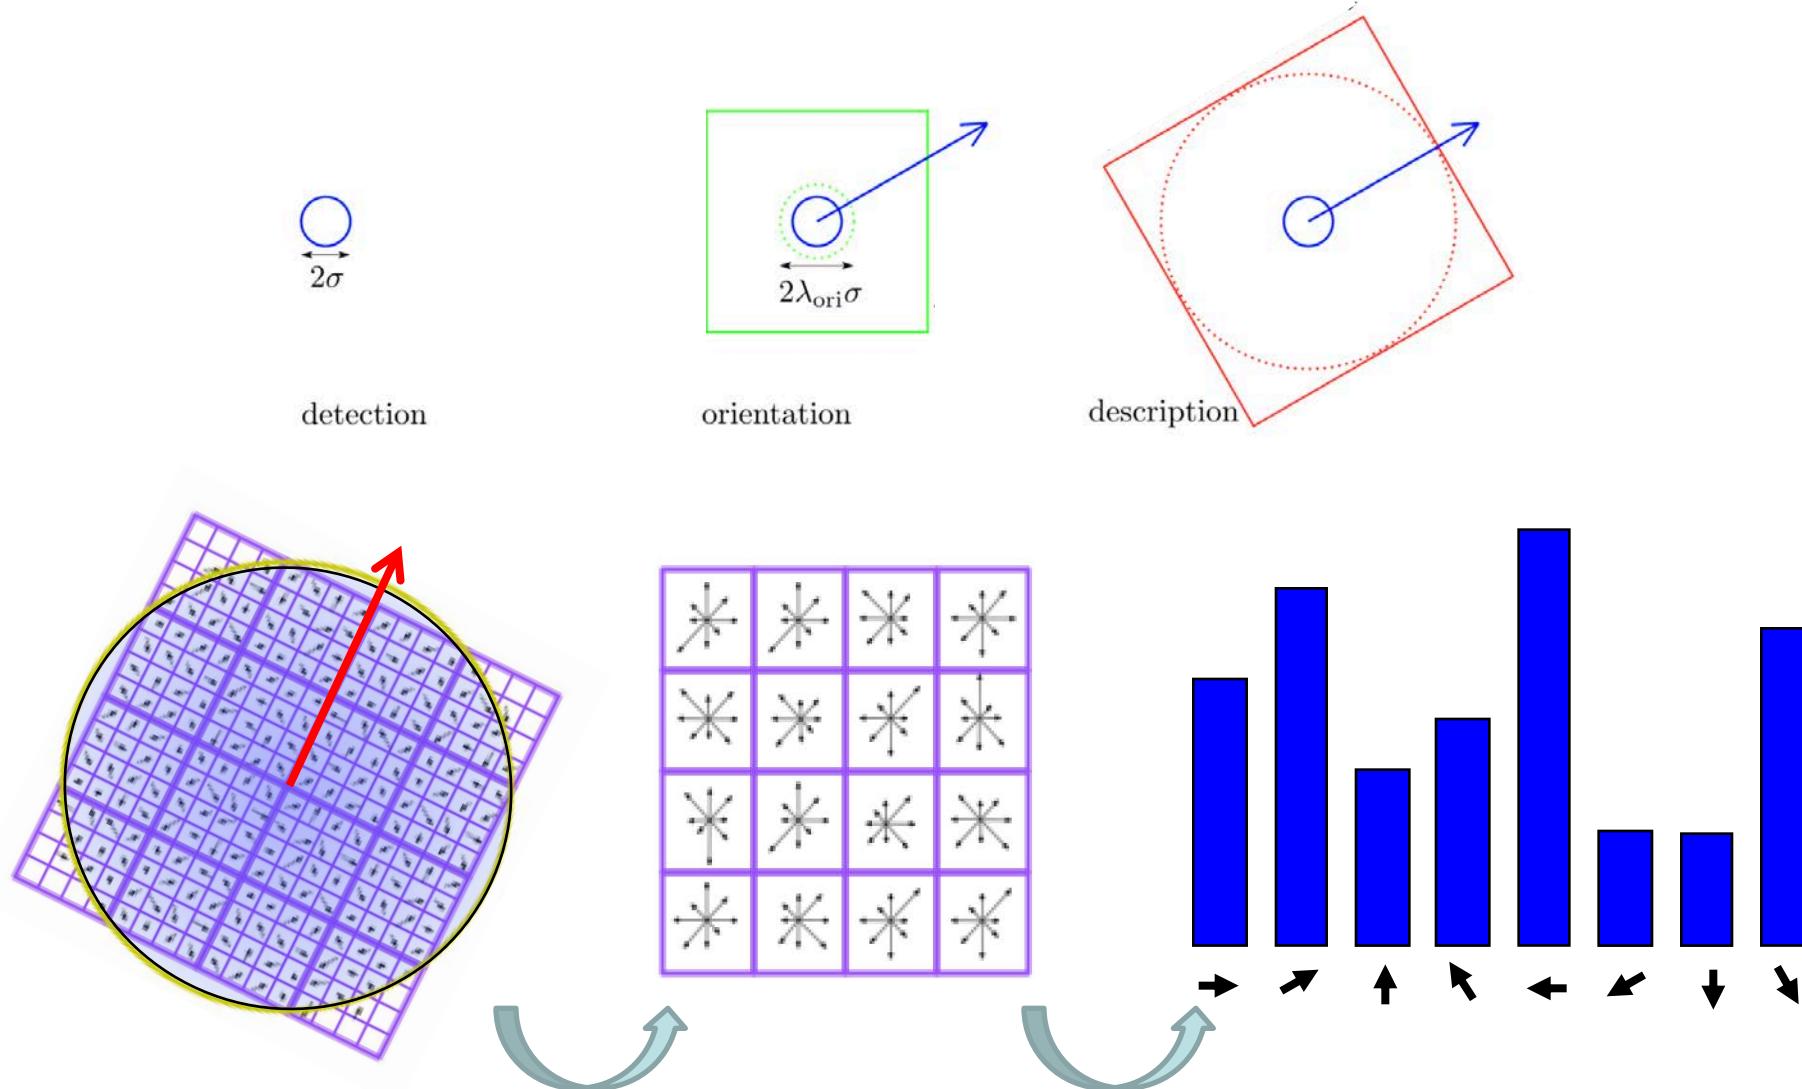
- Feature description (orientation and gradient)
 - Reorient image patch using main orientation
 - Create 16 gradient histograms (8 bins) weighted by magnitude and Gaussian window
 - Final keypoint Descriptor - 128 (4x4x8) element vector



SIFT - Feature description from gradients

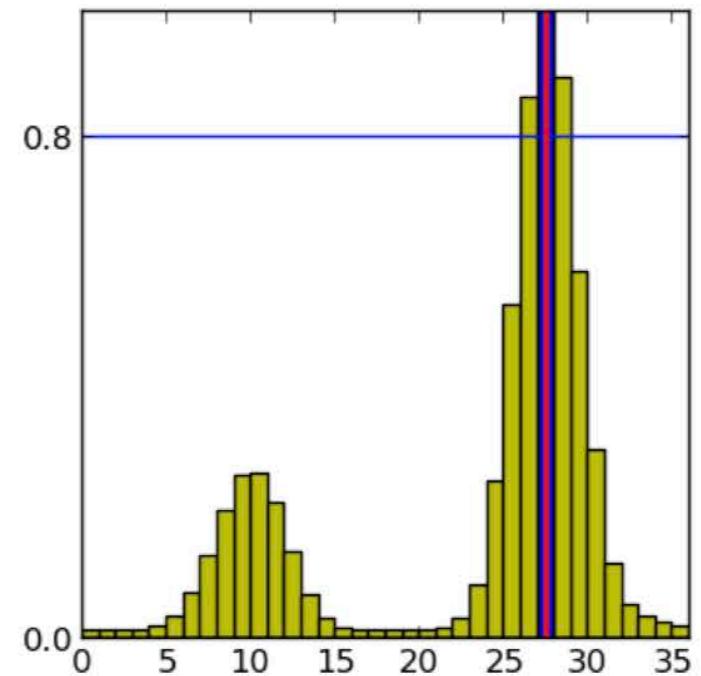
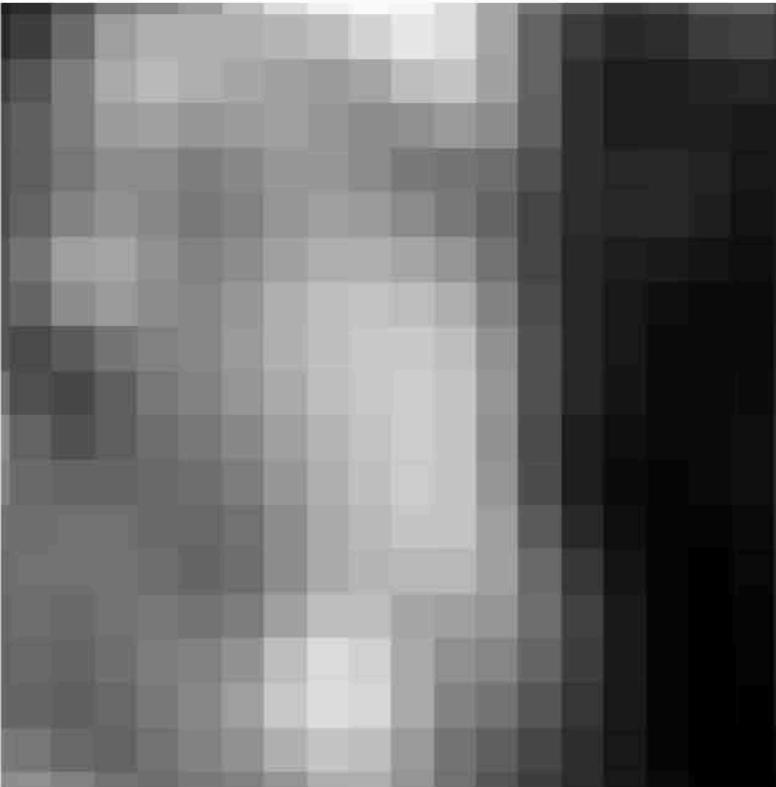
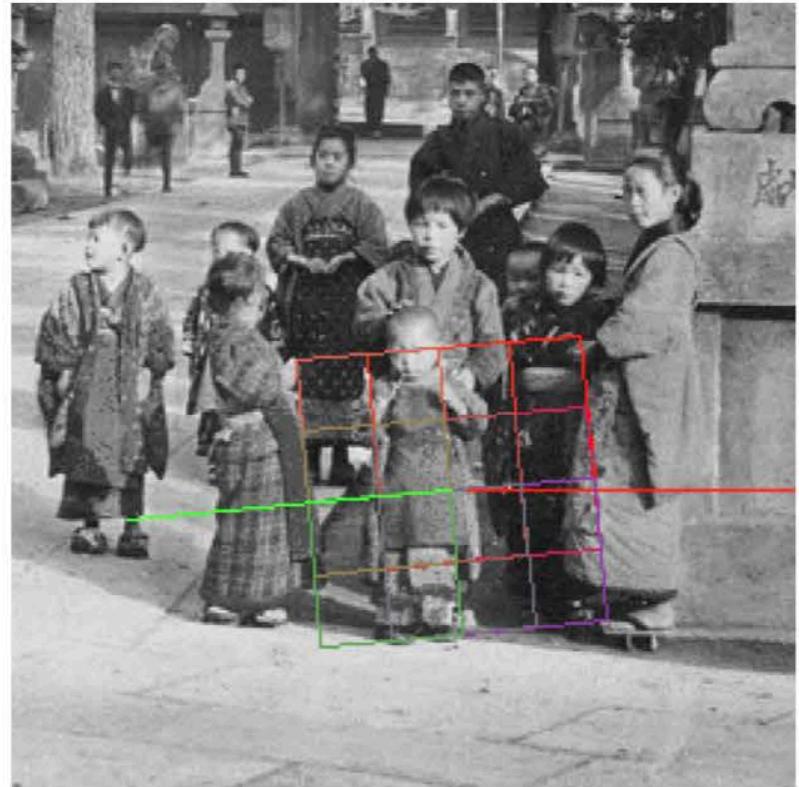


SIFT

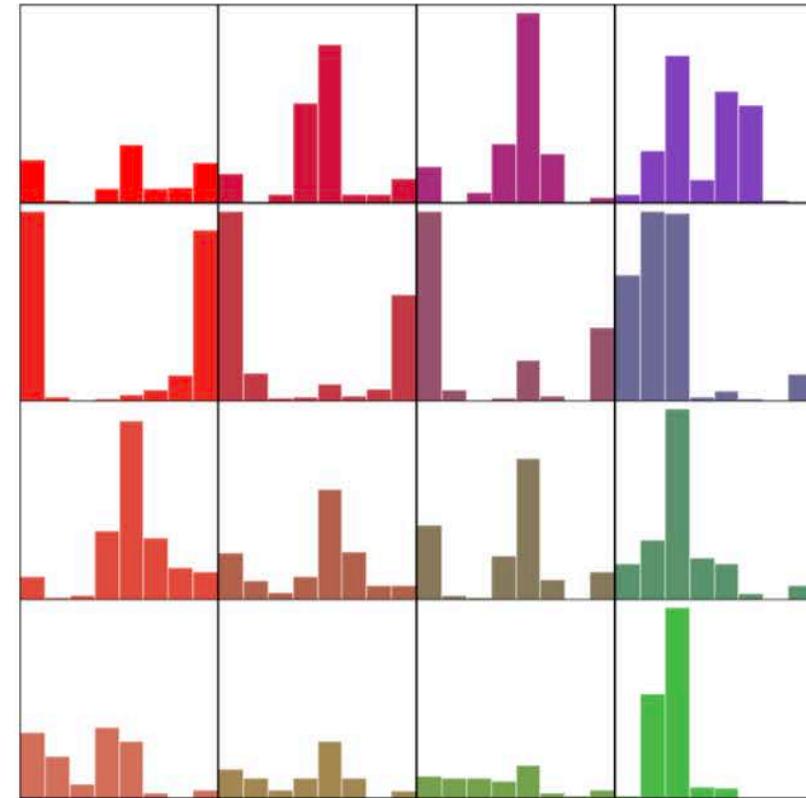
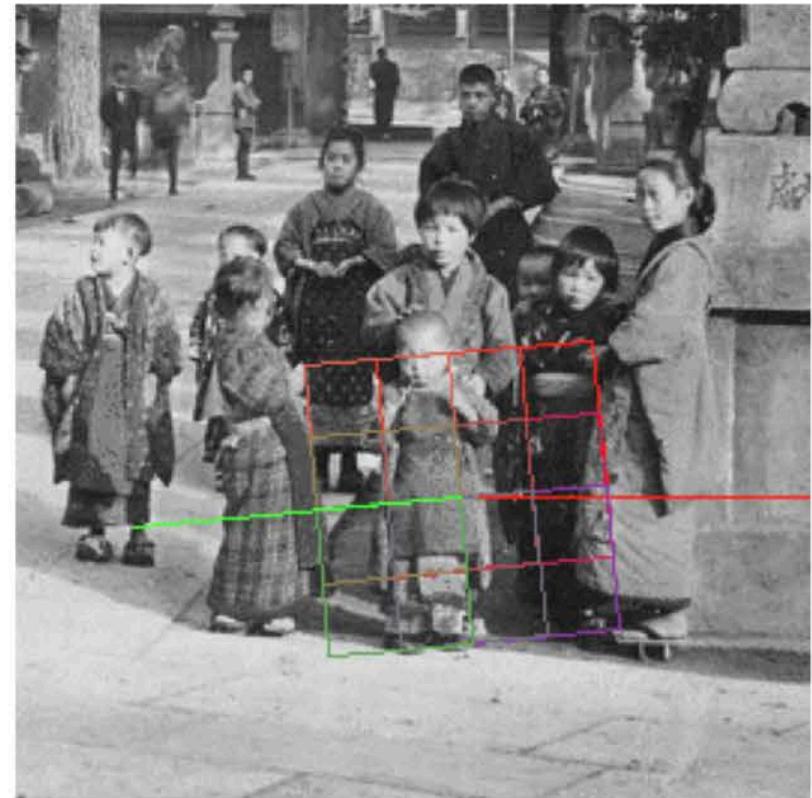




SIFT

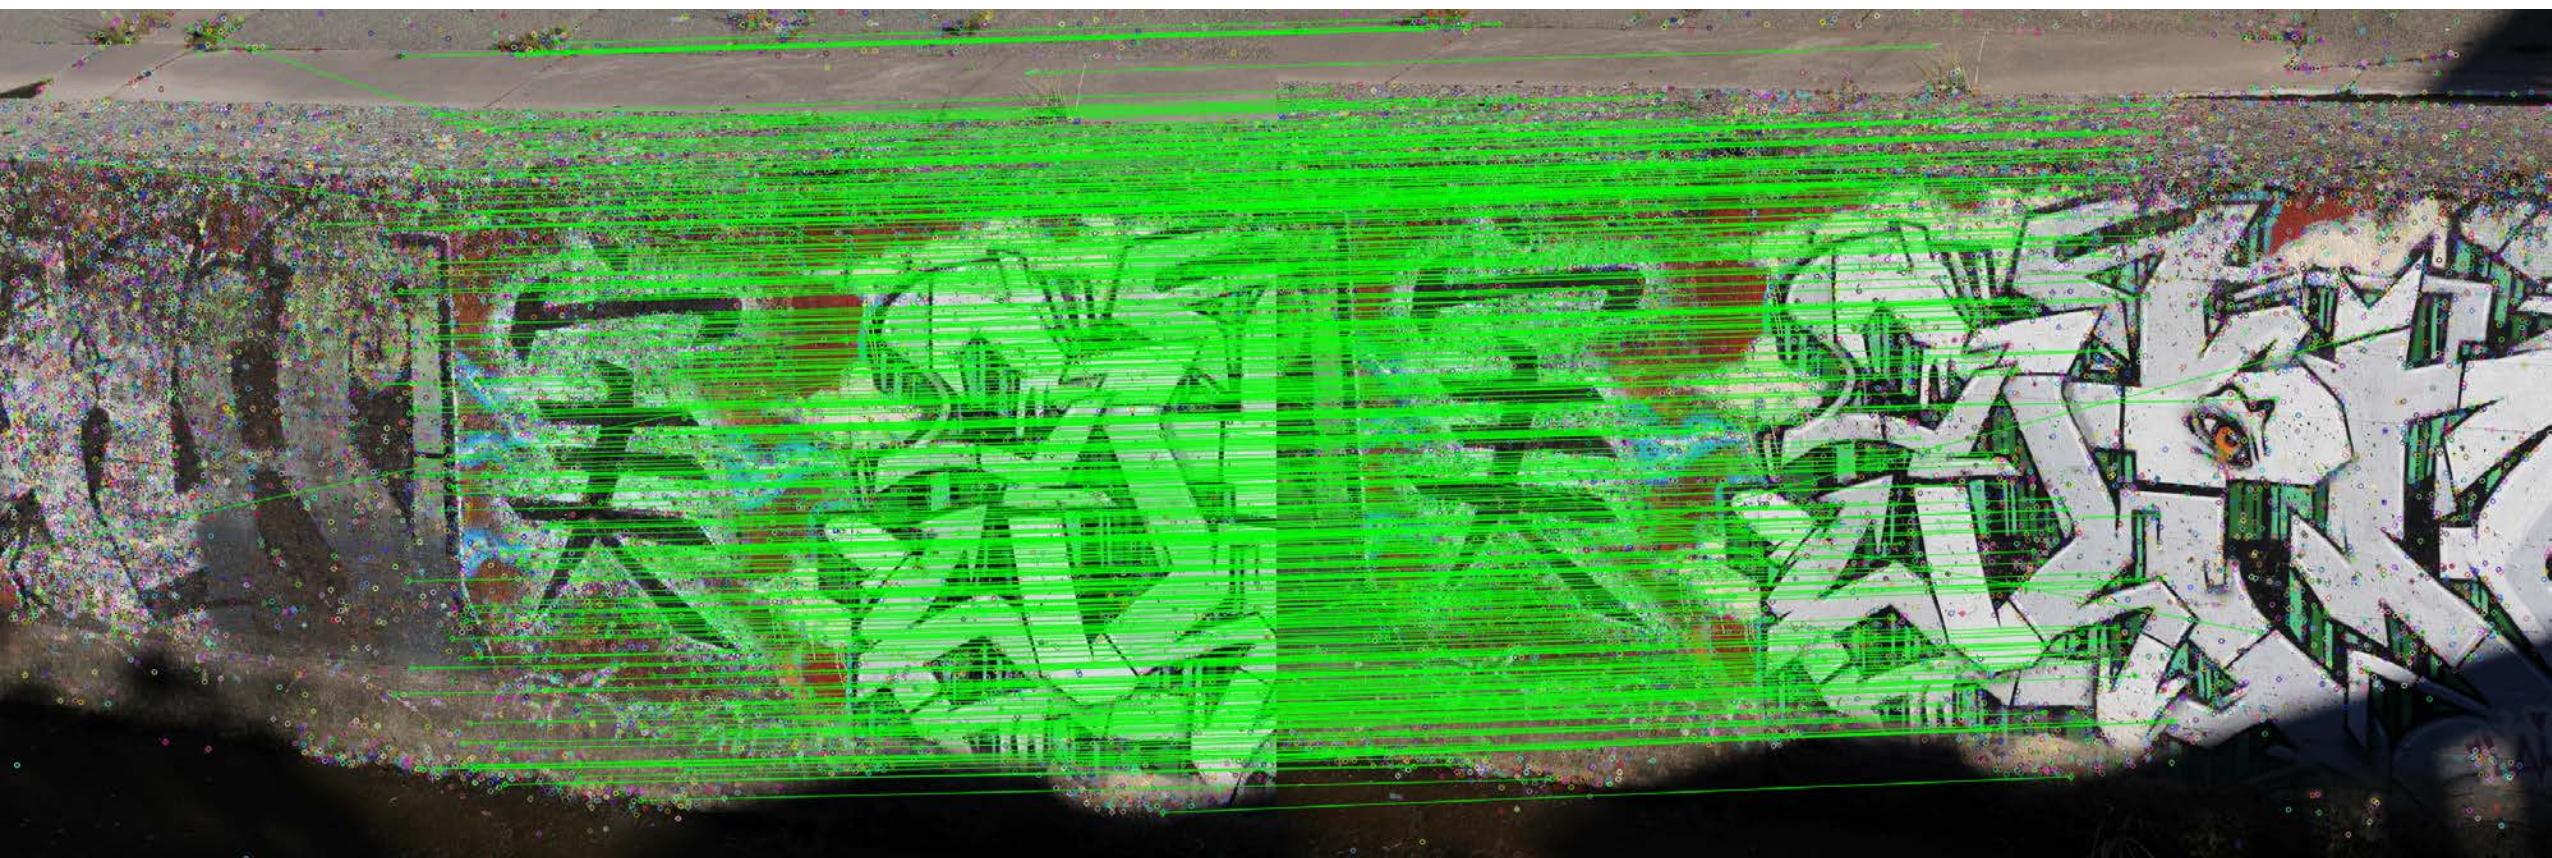


SIFT



SIFT - Feature matching

SIFT - Matching



SIFT - Matching

- SIFT descriptor 128 values
 - Usually these are bytes
- Computing the distance between features:

$$d = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots}$$

- 128 subtractions
- 128 multiplications
- 127 additions
- 1 square root (optional)

SIFT - Matching

- SIFT descriptor 128 values
 - Usually these are bytes
- Computing the distance between features:

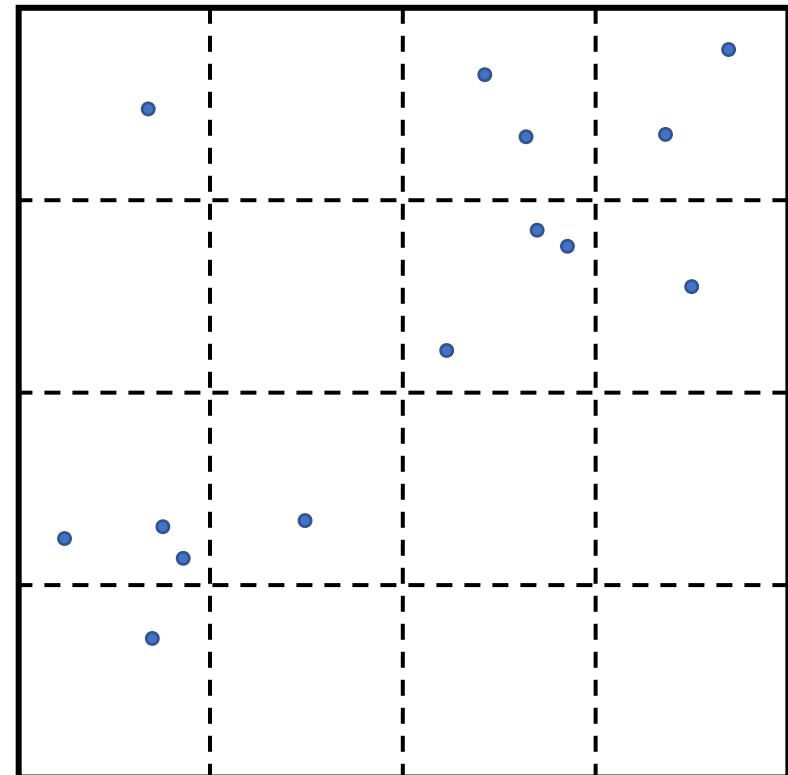
$$d = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2 + \dots}$$

- 128 subtractions
- 128 multiplications
- 127 additions
- 1 square root (optional)

- Suppose we have 10,000 features in each image
 - Brute force method – find distance for each pair
 - Match is smallest distance
 - Matching one feature:
3,840,000 operations
 - Matching all features:
38,400,000,000 ops

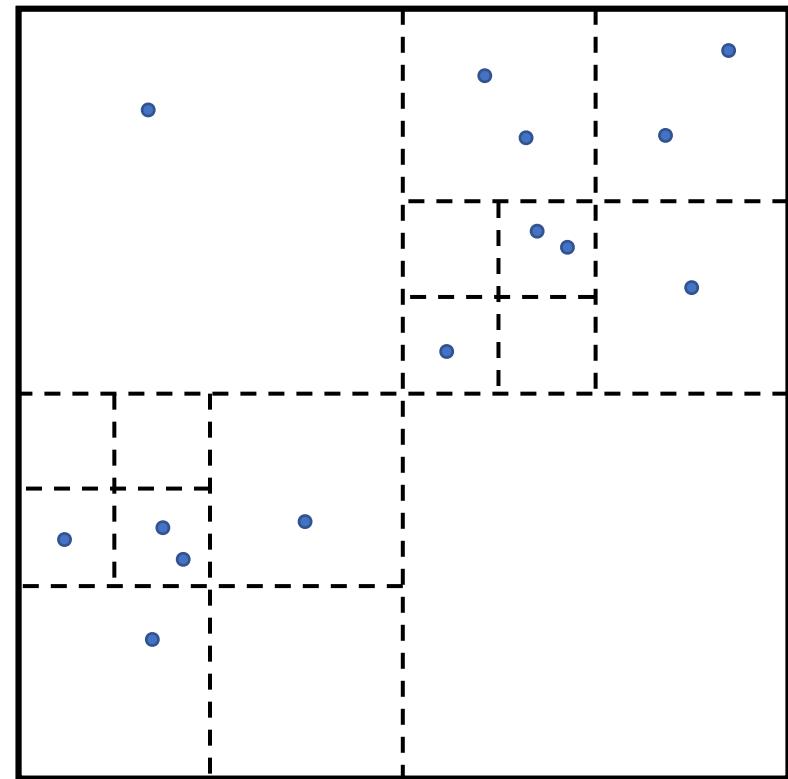
SIFT - Approximate Nearest Neighbours

- Split space into smaller regions
 - Each feature lies in one region
 - Only look for matches in same region
 - Might not find nearest match (why?)
- Uniform subdivision:
 - Divide space into a regular grid
 - 10 divisions on 2-D grid – 100 regions
 - 10 divisions on 128-D grid – 10128



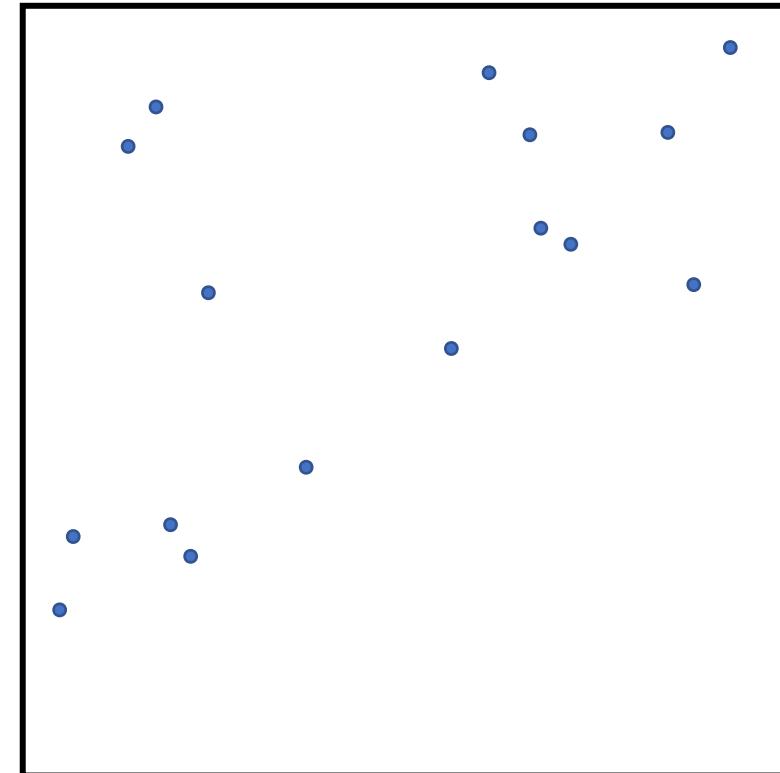
SIFT - Quadtrees, Octrees, etc.

- Recursively split space in half along each dimension
- Stop splitting when:
 - Some max depth is reached
 - Only a few items in the cell
- 2-D gives a quadtree
- 3-D gives an octree
- 128-D: : still too big
- 10 divisions on 128-D grid – 10128

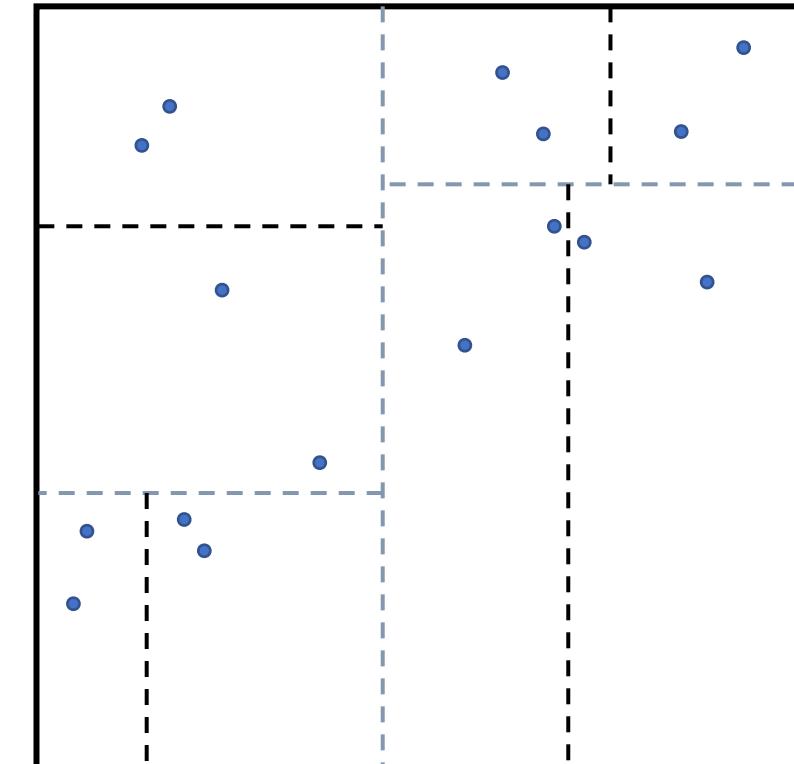
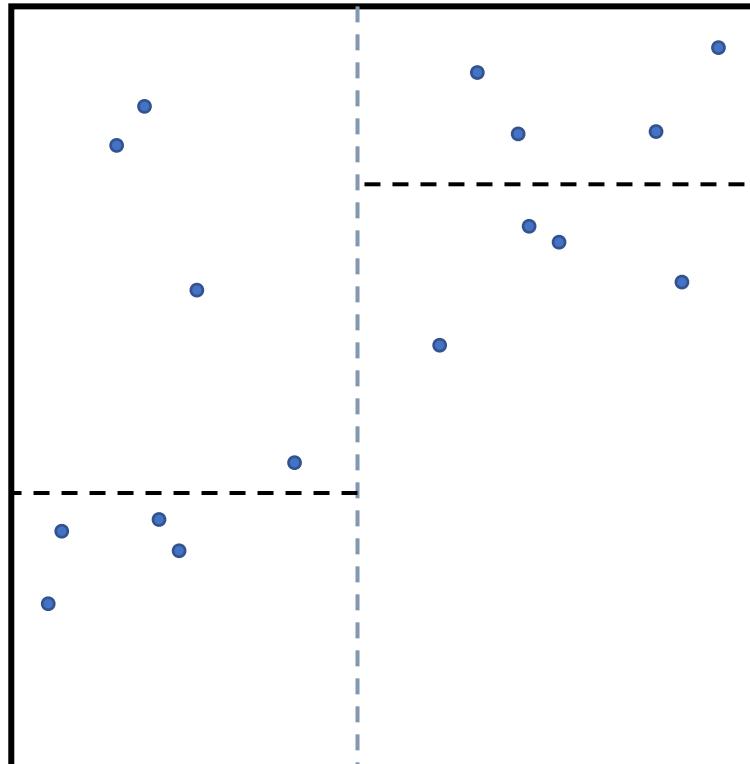
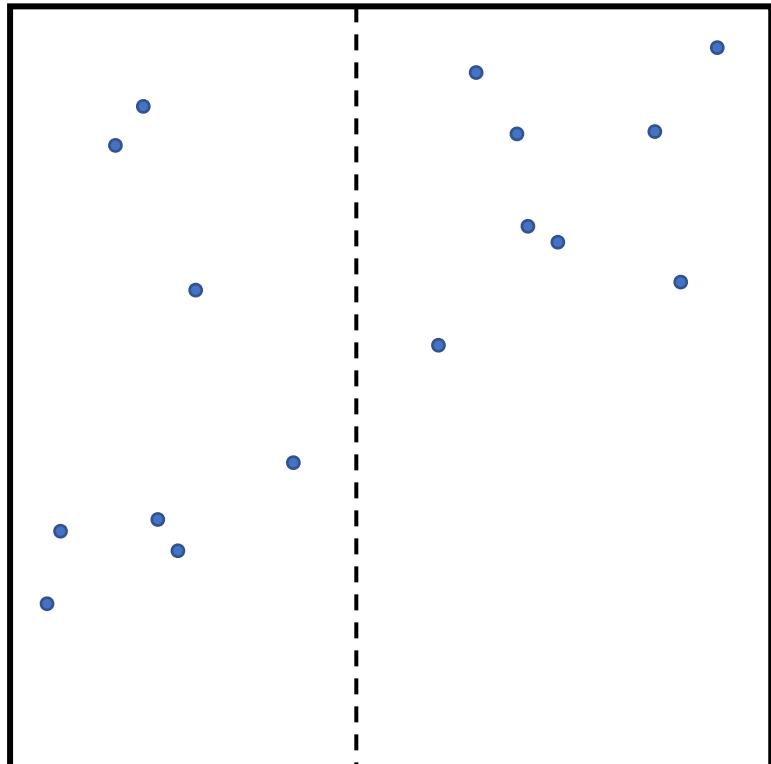


SIFT - K-D Trees

- Choose an axis and split along it
- Then take each half and repeat
- Stop when:
 - Only a few items in the region
 - Reached some maximum depth
- How to choose splits?
 - Axis – next, random, greatest spread?
 - Splitting point – middle or even split?



SIFT - K-D Trees



SIFT - K-D Trees and Feature Matching

- Put the feature descriptors from one image in a k -d tree
 - This has some overhead, but worth it if many features
- Given a feature from the other image:
 - Find what cell in the k -d tree contains that feature's descriptor
 - Compute the distance to all features in that cell
 - The nearest one is (probably) the best match
- For a tree with n layers and F features, this requires:
 - n comparisons to find the appropriate cell
 - $F/2^n$ descriptor-distance computations to find the best match

SIFT - Matching SIFT Features

- Even with brute-force, most SIFT matches are wrong
 - Lack of texture, repeating features, etc. -> ambiguous matches
 - And SIFT is the best we have
- With k -d trees this gets worse, but not much
- Solution – find best two matches
 - Ambiguous matches will have two similar distances
 - Keep only matches with a big difference between the two
 - Find two nearest matches $d(f, m_1) < \alpha d(f, m_2)$, $d(f, m_i)$ is distance to i th match, $\alpha < 1$ is a parameter, e.g.: 0.8
- Makes things better, but still some wrong matches

SIFT - Matching SIFT Features

- Even with brute-force, most SIFT matches are wrong
 - Lack of texture, repeating features, etc. -> ambiguous matches
 - And SIFT is the best we have
- With k -d trees this gets worse, but not much
- Solution – find best two matches
 - Ambiguous matches will have two similar distances
 - Keep only matches with a big difference between the two
 - Find two nearest matches $d(f, m_1) < \alpha d(f, m_2)$, $d(f, m_i)$ is distance to i th match, $\alpha < 1$ is a parameter, e.g.: 0.8
- Makes things better, but still some wrong matches

SIFT - Summary

SIFT - Matching SIFT Features

- Why does this work?
 - Key points are extracted at different scales and blur levels and all subsequent computations are performed within the scale space framework
 - **Invariance to image scaling and small changes in perspective**
 - Descriptor computation is relative to a reference (main) orientation
 - **Descriptors robust against rotation**
 - Descriptor information is stored relative to the key point position
 - **Invariance against translations**
 - Many potential key points are discarded if they are deemed unstable or hard to locate precisely
 - **Remaining key points should thus be relatively immune to image noise**
 - Histograms are normalized at the end which means the descriptors will not store the magnitudes of the gradients, only their relations to each other
 - **Descriptors are invariant against global, uniform illumination changes**
 - The histogram values are also thresholded to reduce the influence of large gradients. This will make the information partly immune to local, non-uniform changes in illumination.

The end!