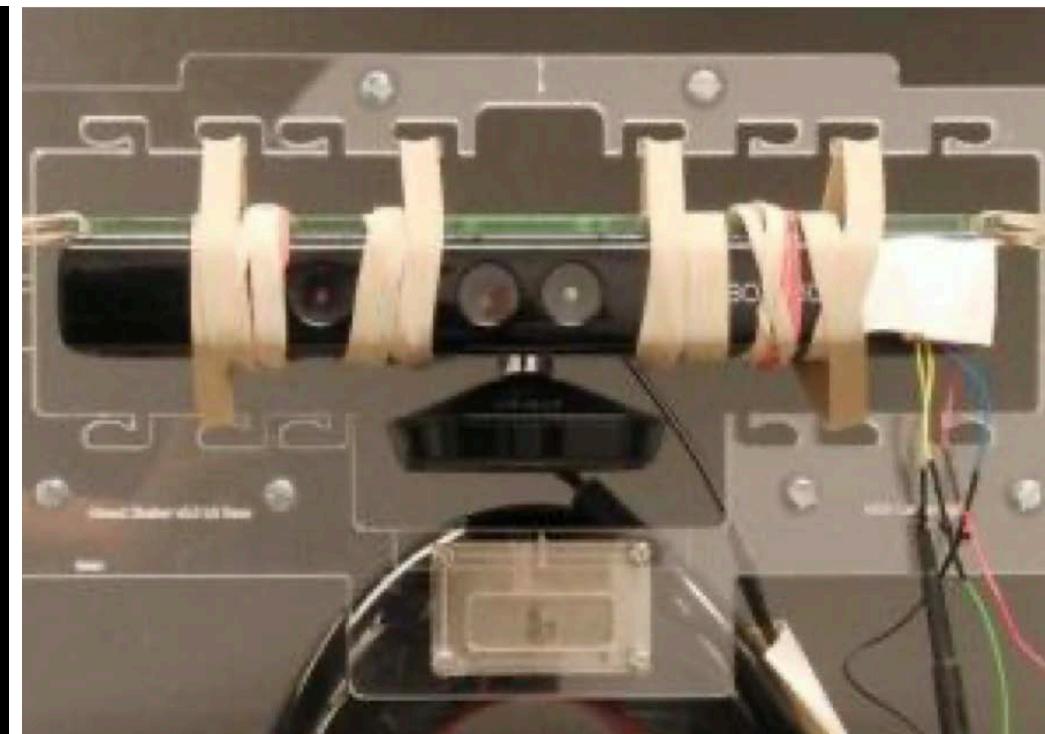
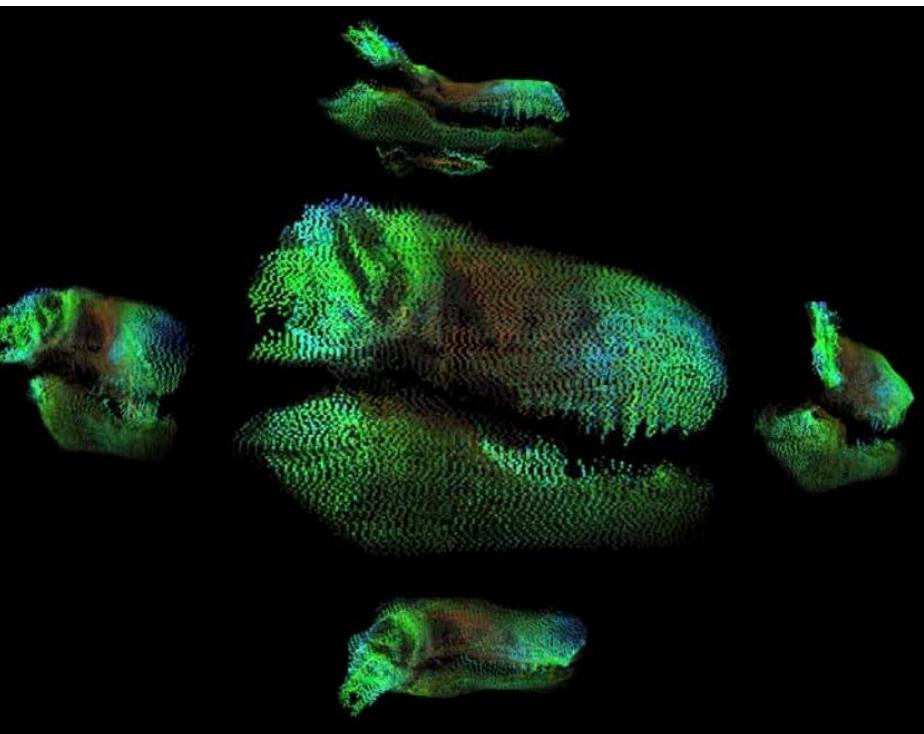
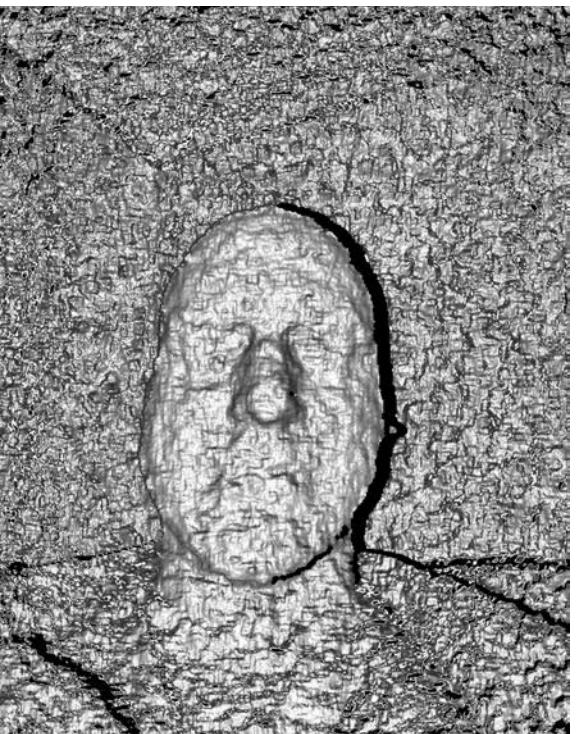


Visual Computing I:

Interactive Computer Graphics and Vision



Practical Depth Sensing and Working with Range Data

Stefanie Zollmann and Tobias Langlotz

Last time..

The Fundamental Matrix

- One way to express F is

$$F = K_2^{-T} [t]_x R K_1^{-1}$$

- Here
 - K_1 and K_2 are the calibration matrices of the two cameras
 - R and t are the rotation and translation between them
 - $[t]_x$ is the matrix form of the cross product with t :

$$[t]_x = \begin{bmatrix} 0 & -t_z & t_y \\ t_z & 0 & -t_x \\ -t_y & t_x & 0 \end{bmatrix}$$

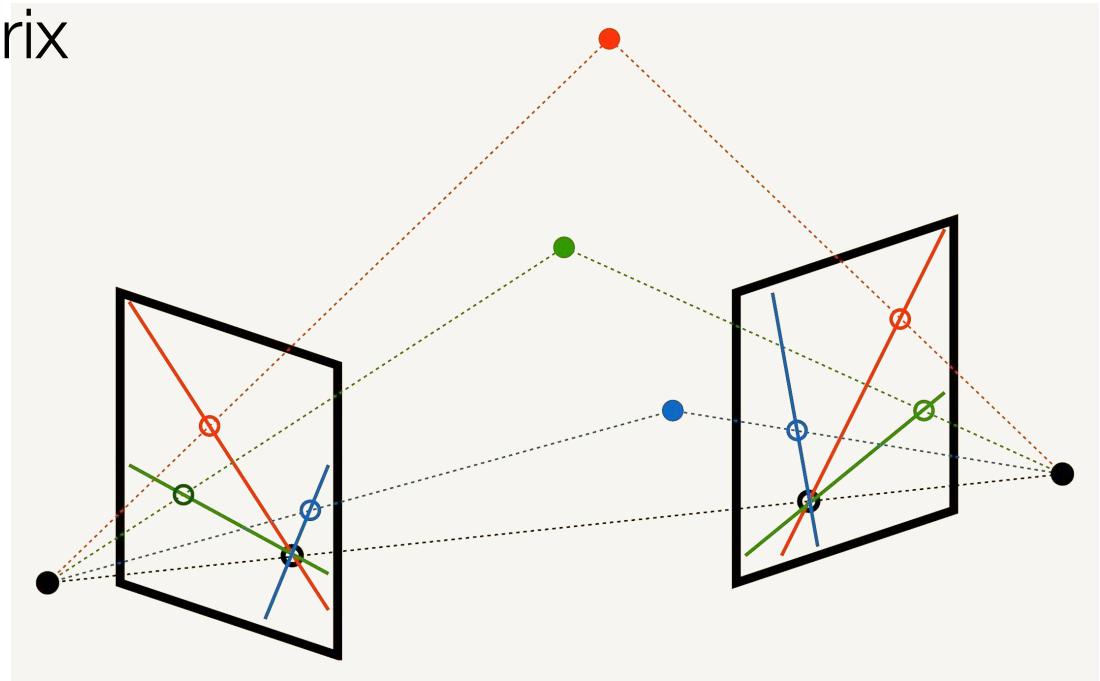
Notes:

$$p_1 = K_1[I \mid 0]x$$

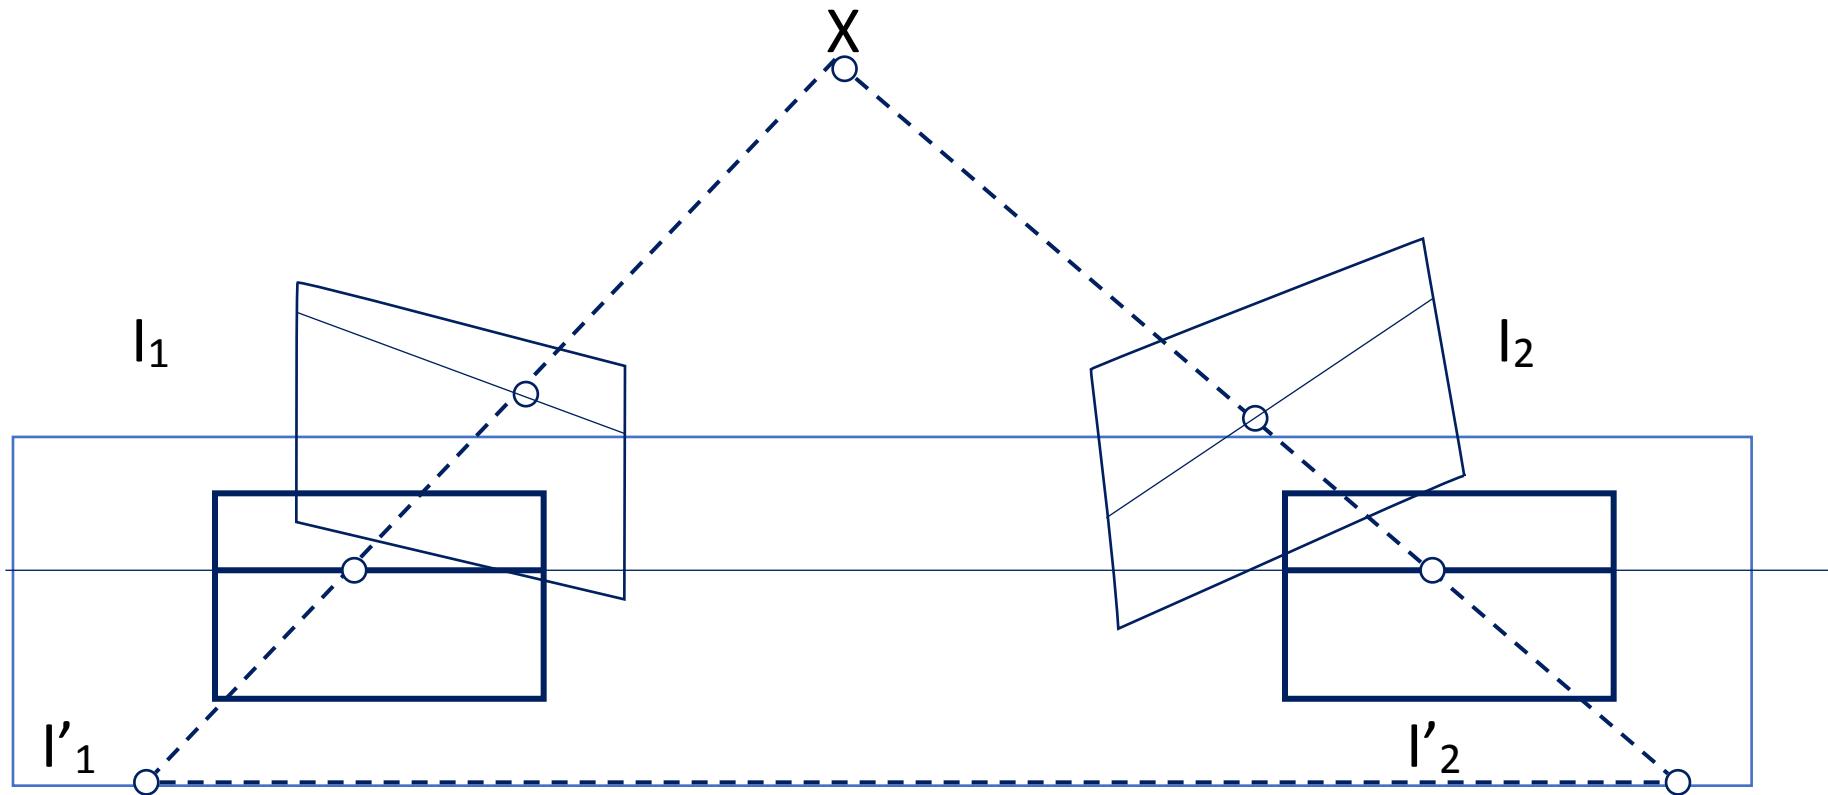
$$p_2 = K_2[R \mid t]x$$

The Essential Matrix

- For calibrated cameras
 - We can factor out \mathbf{K}_1 and \mathbf{K}_2 in $\mathbf{F} = \mathbf{K}_2^{-T}[\mathbf{t}]_{\times}\mathbf{R}\mathbf{K}_1^{-1}$
 - This gives us the essential matrix
 - $\mathbf{E} = \mathbf{K}_2^T\mathbf{F}\mathbf{K}_1 = [\mathbf{t}]_{\times}\mathbf{R}$
 - \mathbf{E} has only five DoF:
 - 3 for 3D rotation
 - 3 for 3D translation
 - Less 1 for unknown scale



Stereo Rectification

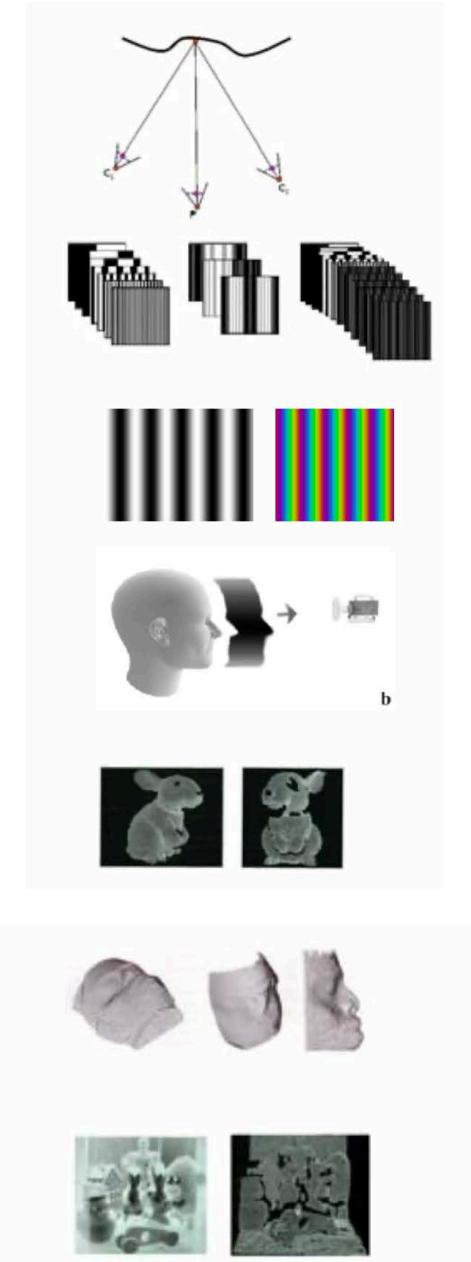


Today:

Practical Depth Sensing and Range Data

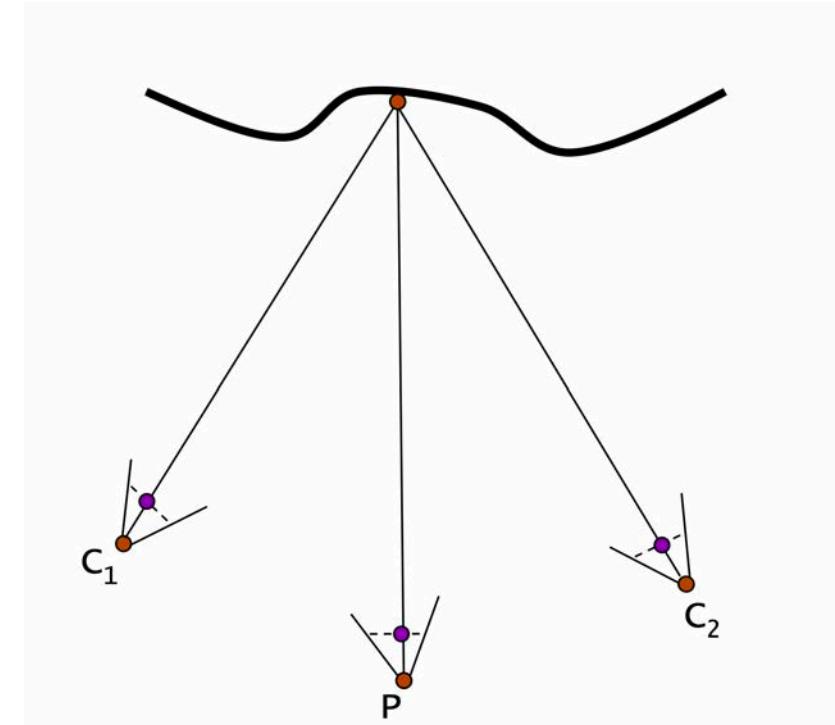
Overview

- Range Scanning
 - Passive:
 - Stereo or Multi-Camera
 - Active:
 - Structured Light Projection
 - Optical Time-of-Flight
 - Direct ToF
 - Indirect ToF
- Registration of Depth Maps
- Processing Depth Maps

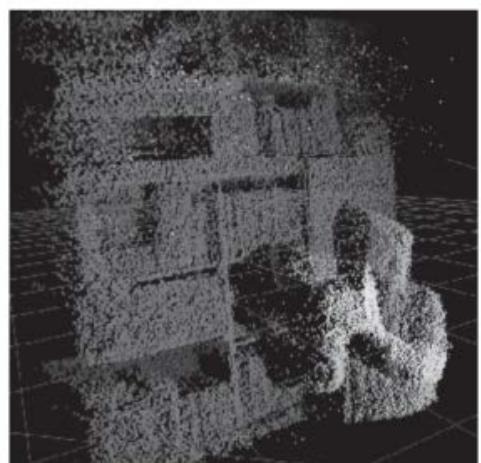


Range Scanning

- Range sensors:
 - Passive: ≥ 2 cameras (stereovision)
 - Active: ≥ 1 cameras + ≥ 1 projectors (structured light)
- Result is depth information given as:
 - Range image / depth map store depth at pixel positions from a specific view (RGBD image)
 - A point cloud is a 3D dataset of discrete (X,Y,Z) surface points



Depth Map

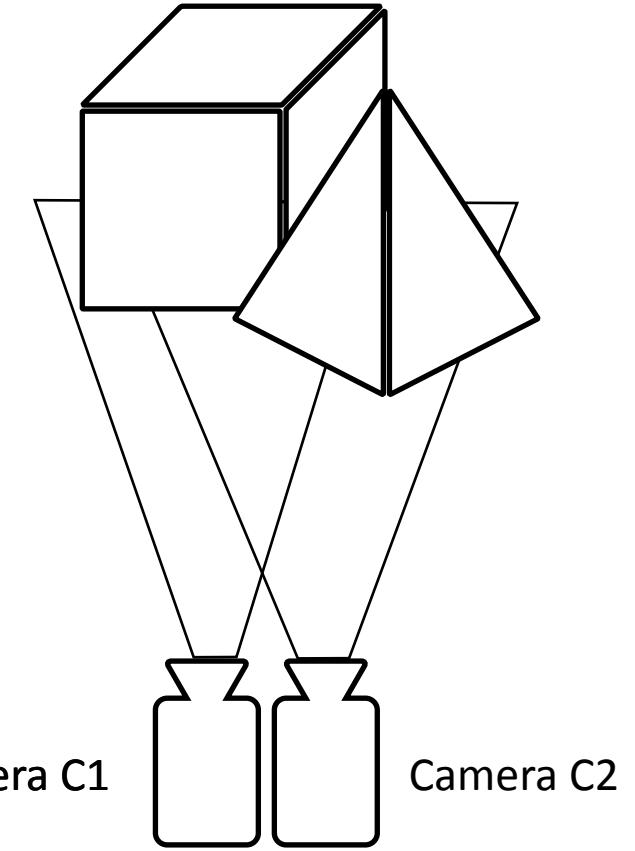


Point cloud

Passive Range Scanning

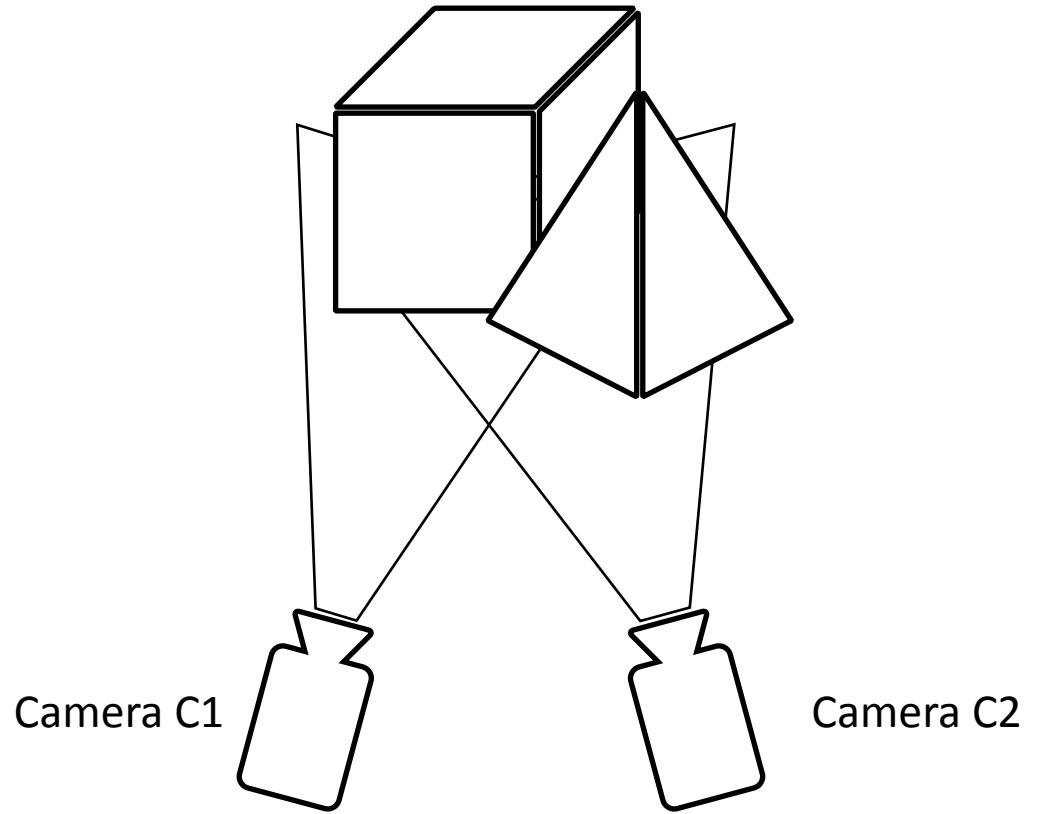
Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?



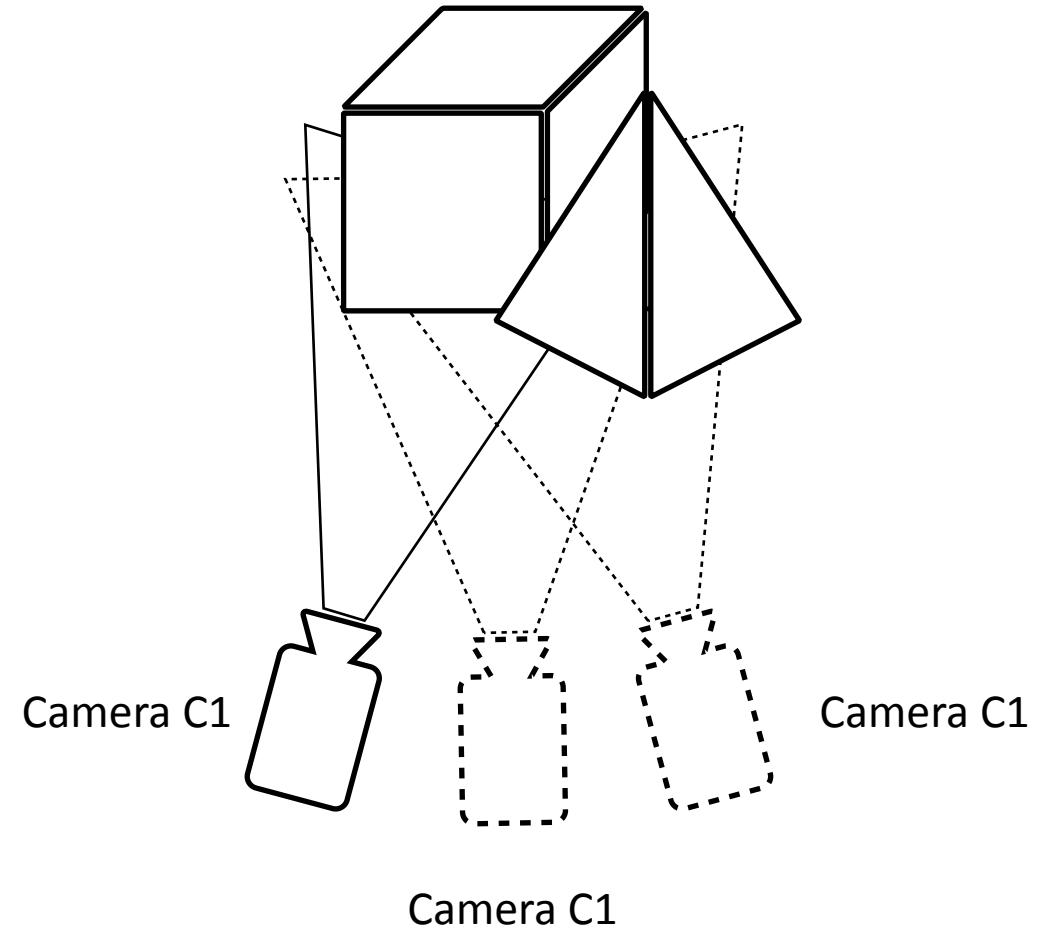
Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?



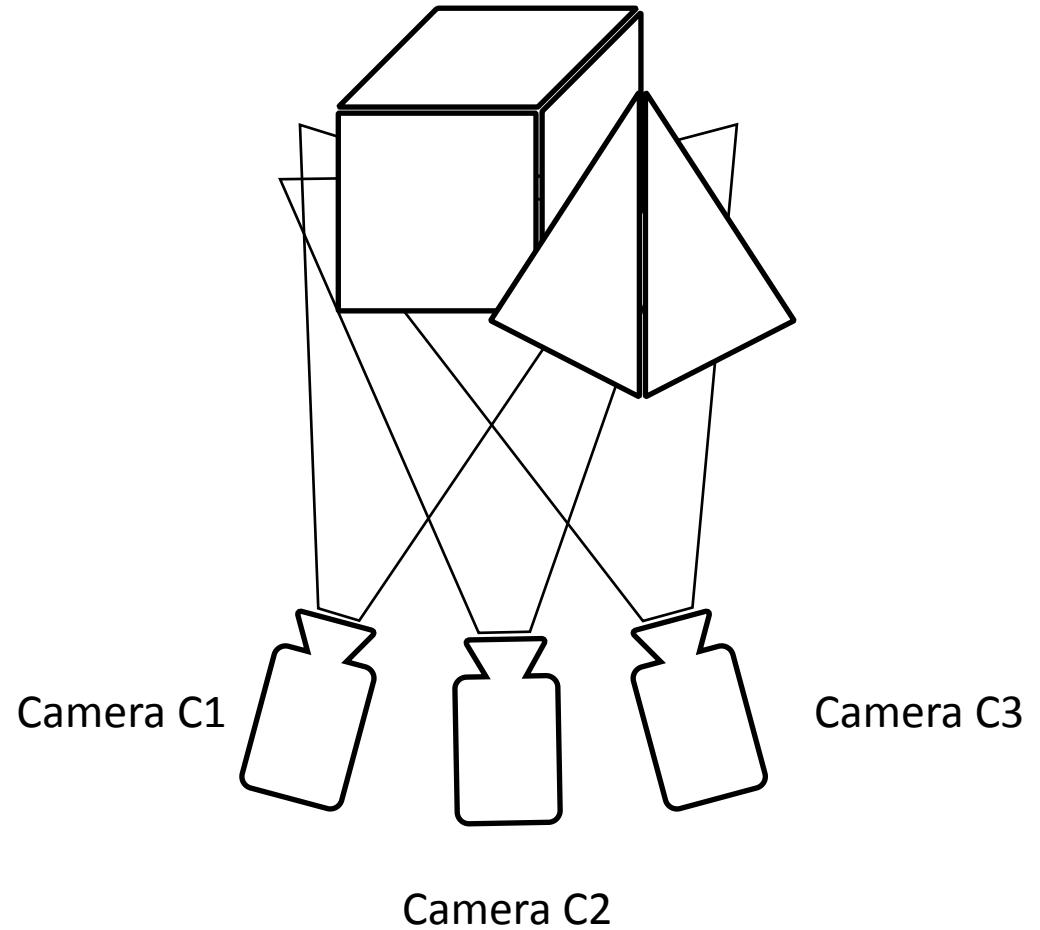
Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?



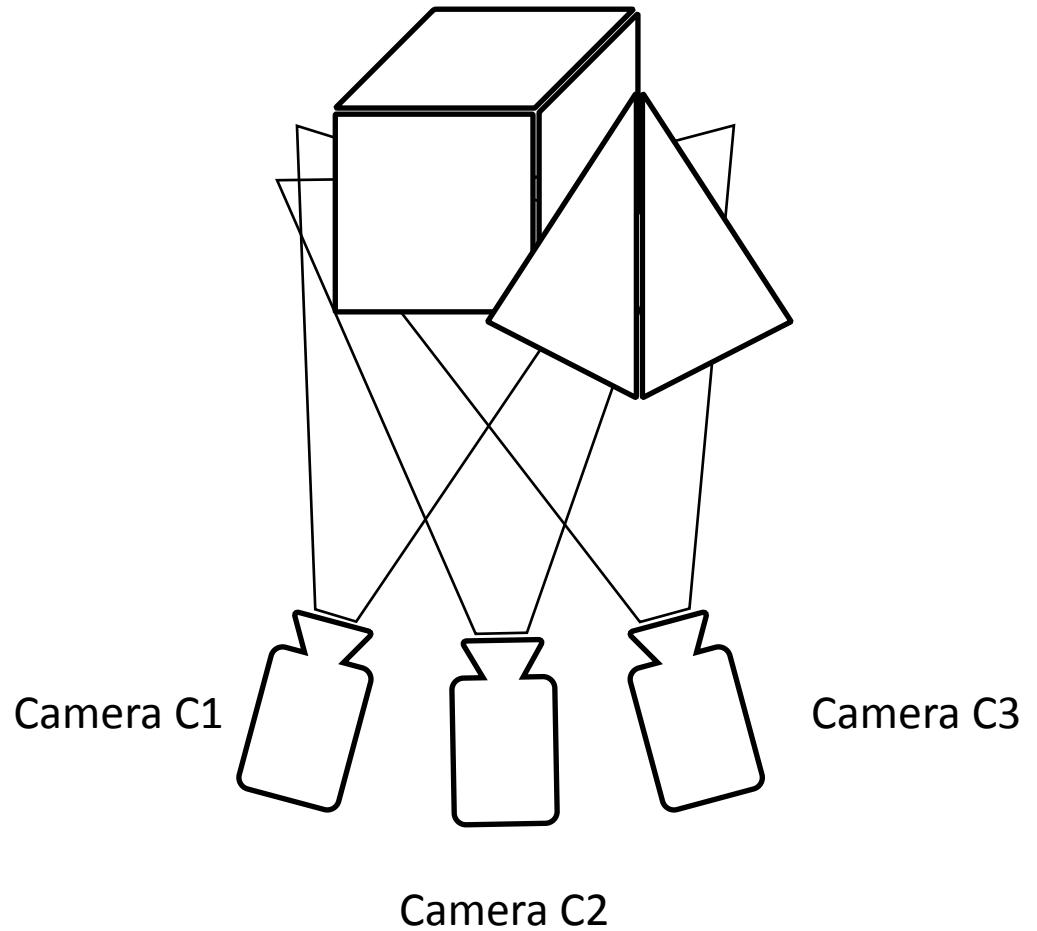
Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?



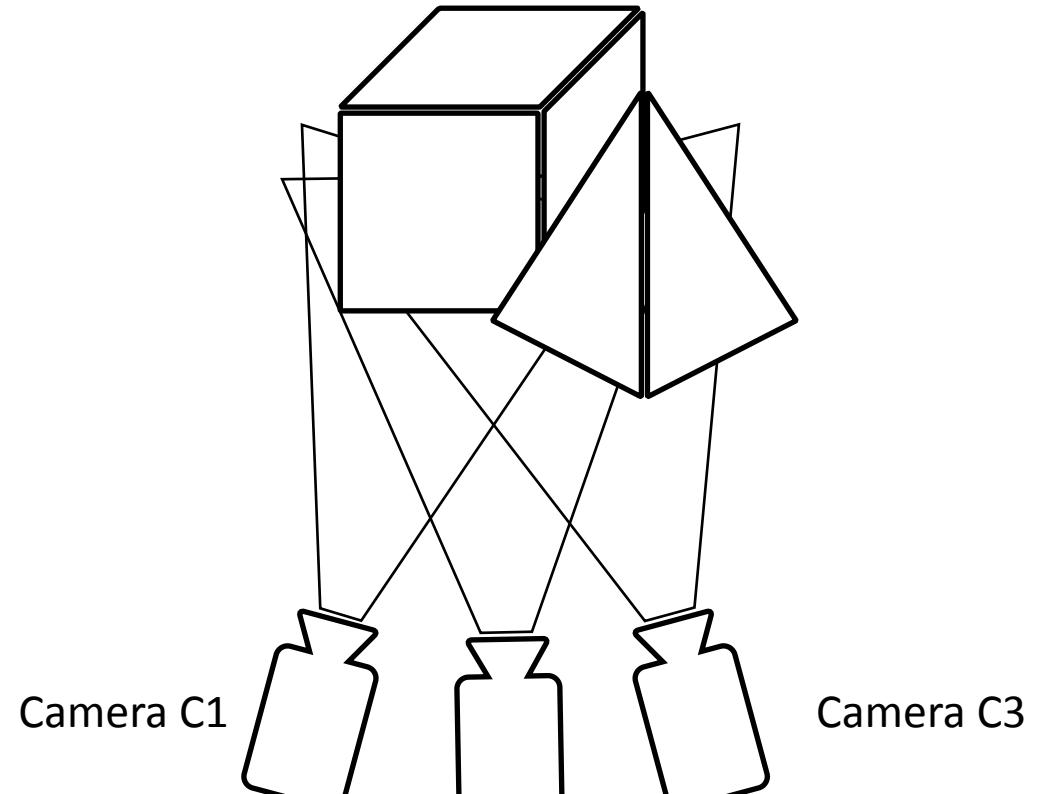
Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?
- What is the main challenge?



Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?
- What is the main challenge?



PointGrey Bumble Bee



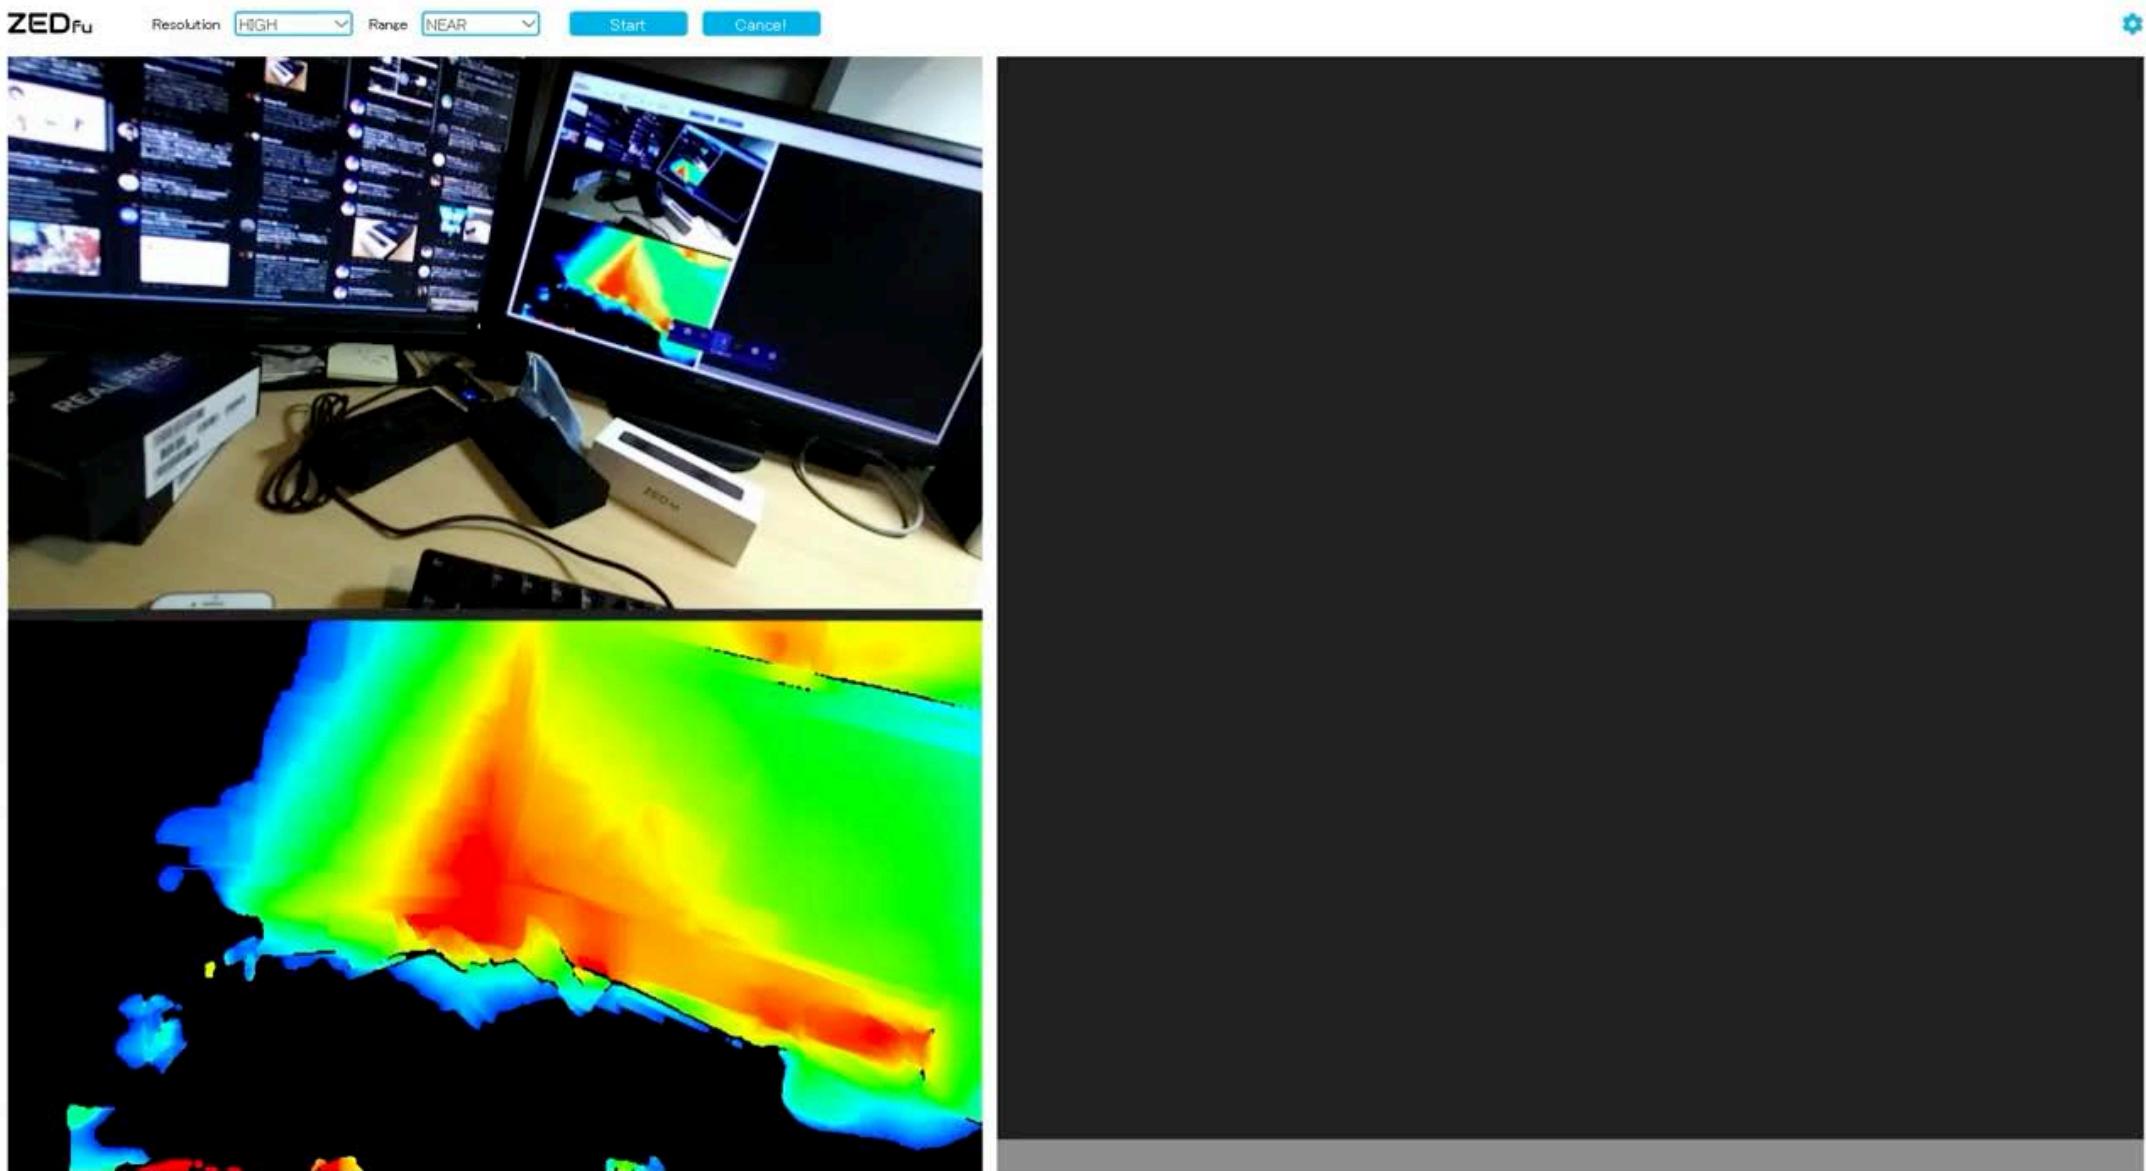
Zed Mini

Passive range scanning



Youtube Taeyeon Kim

Passive range scanning



Youtube Atsizumi

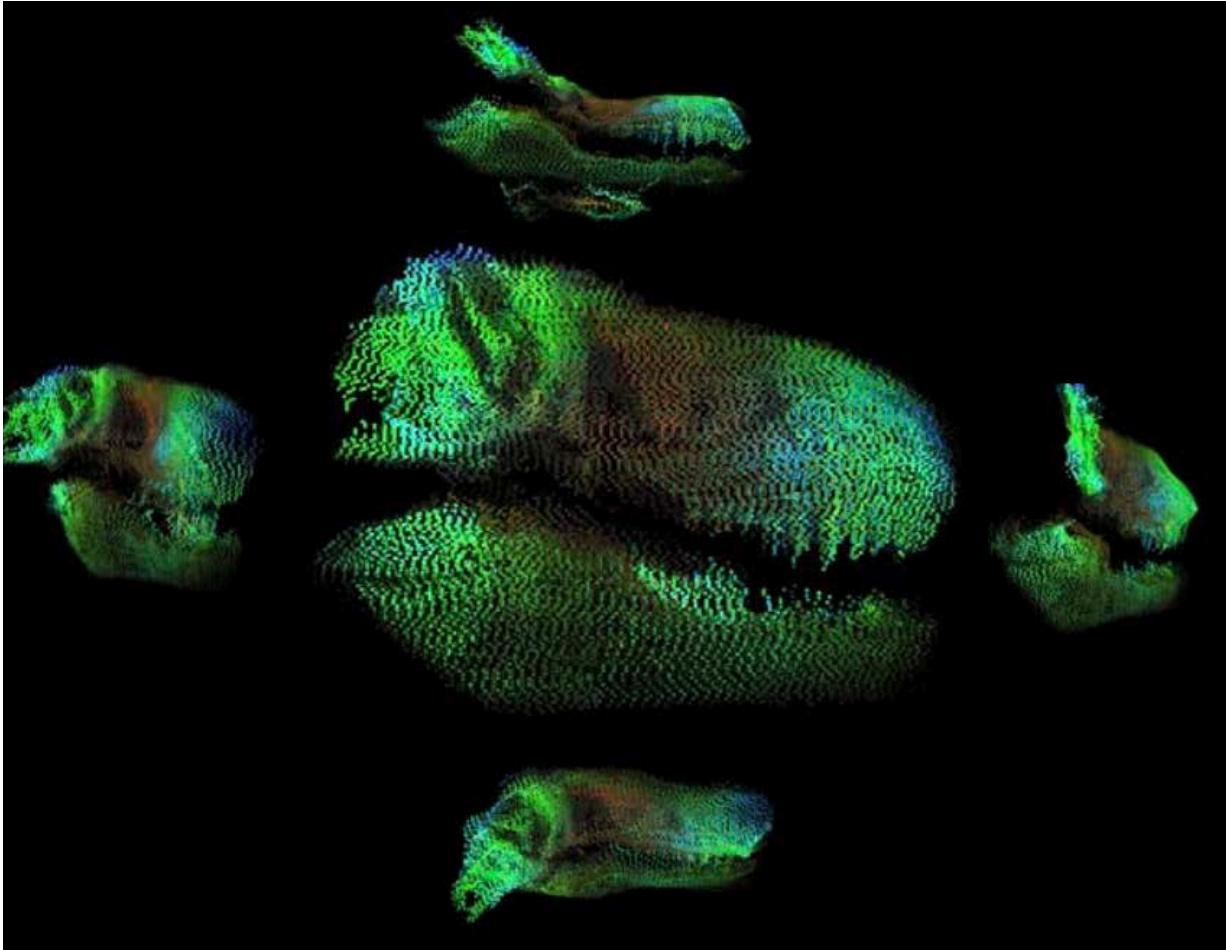
Passive range scanning



Vexcel (YouTube Video: Skyline Software Systems)

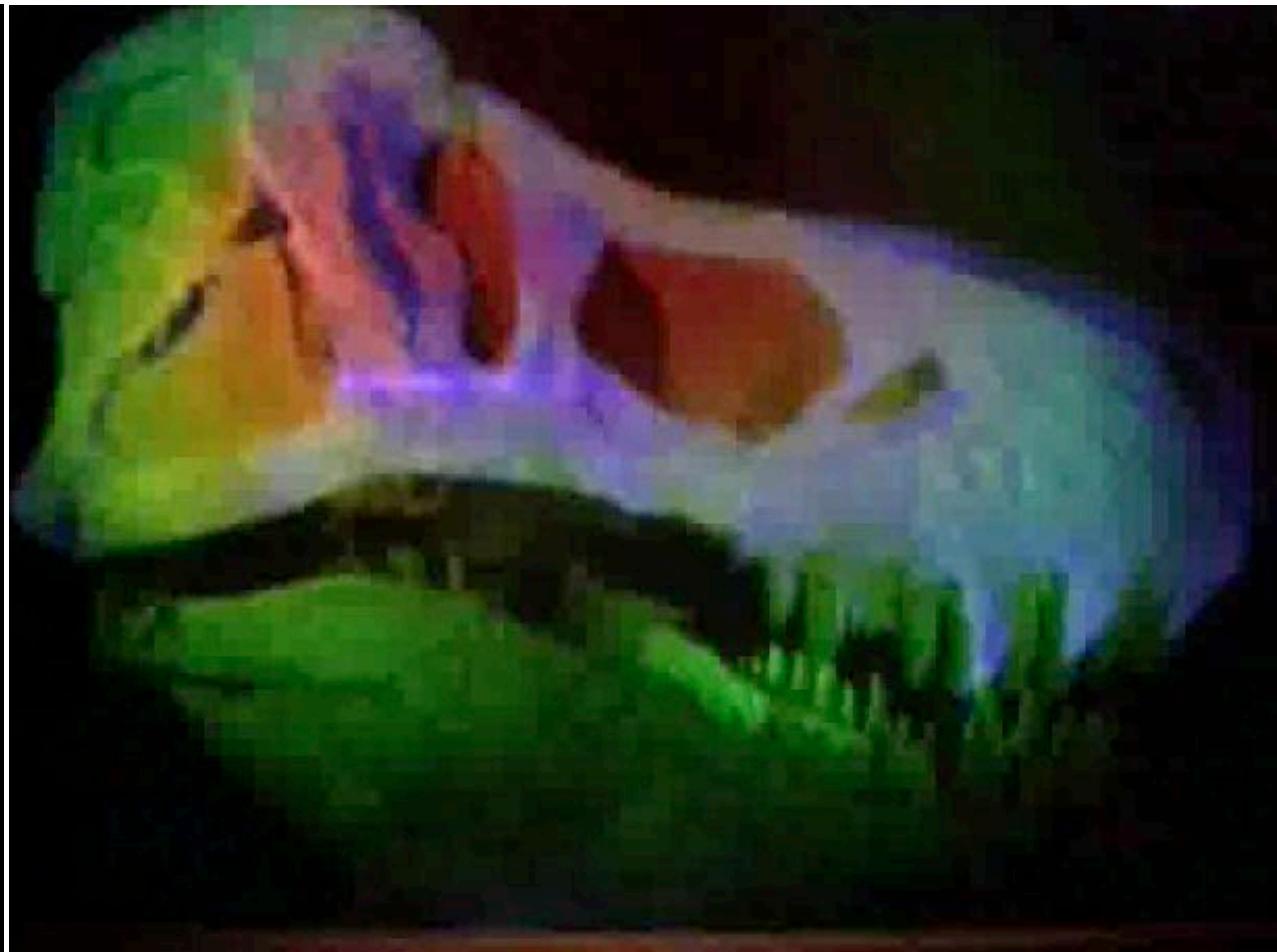
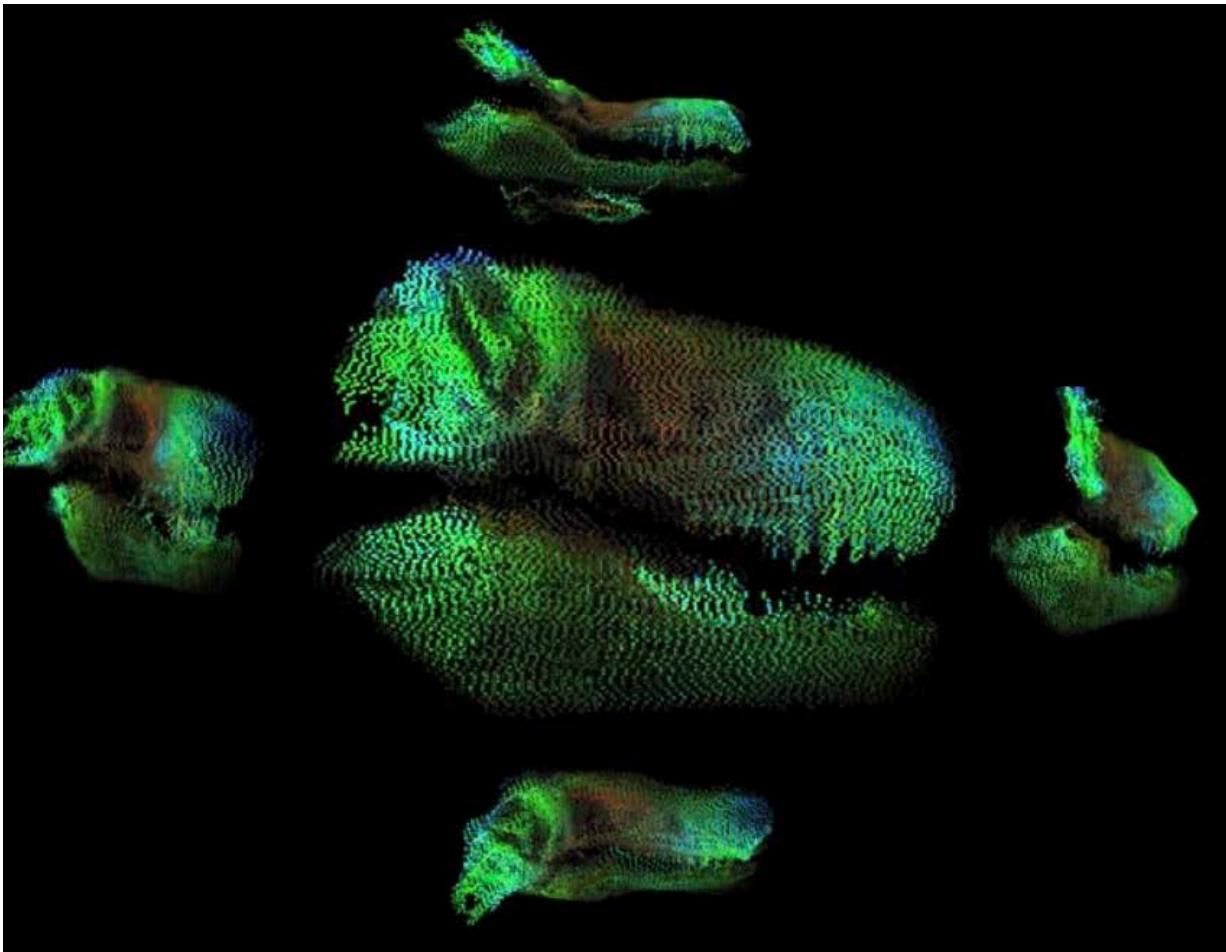
Skyline VEXCEL
IMAGING

Passive range scanning



Holographics

Passive range scanning

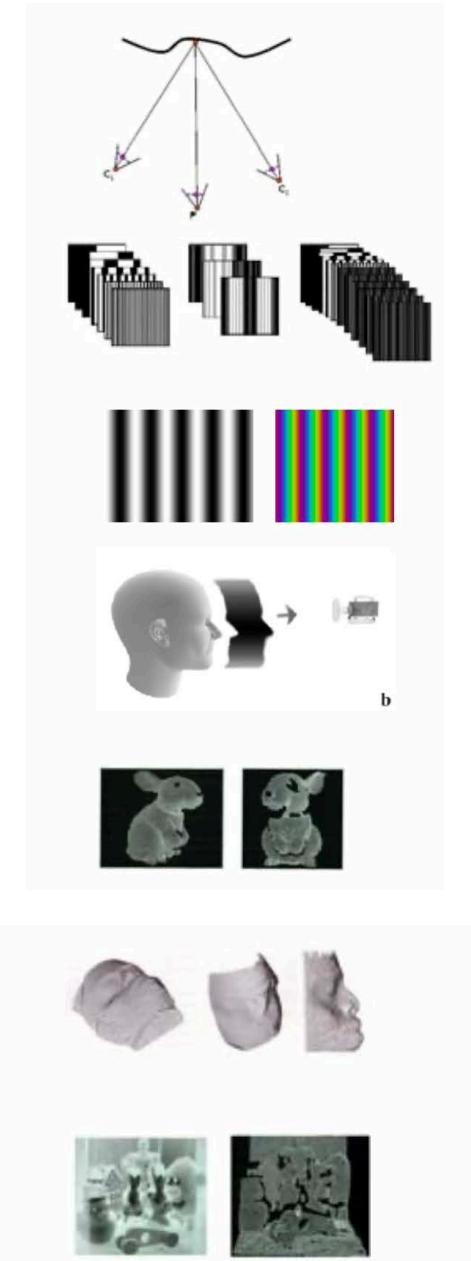


Holographics

Active Range Scanning

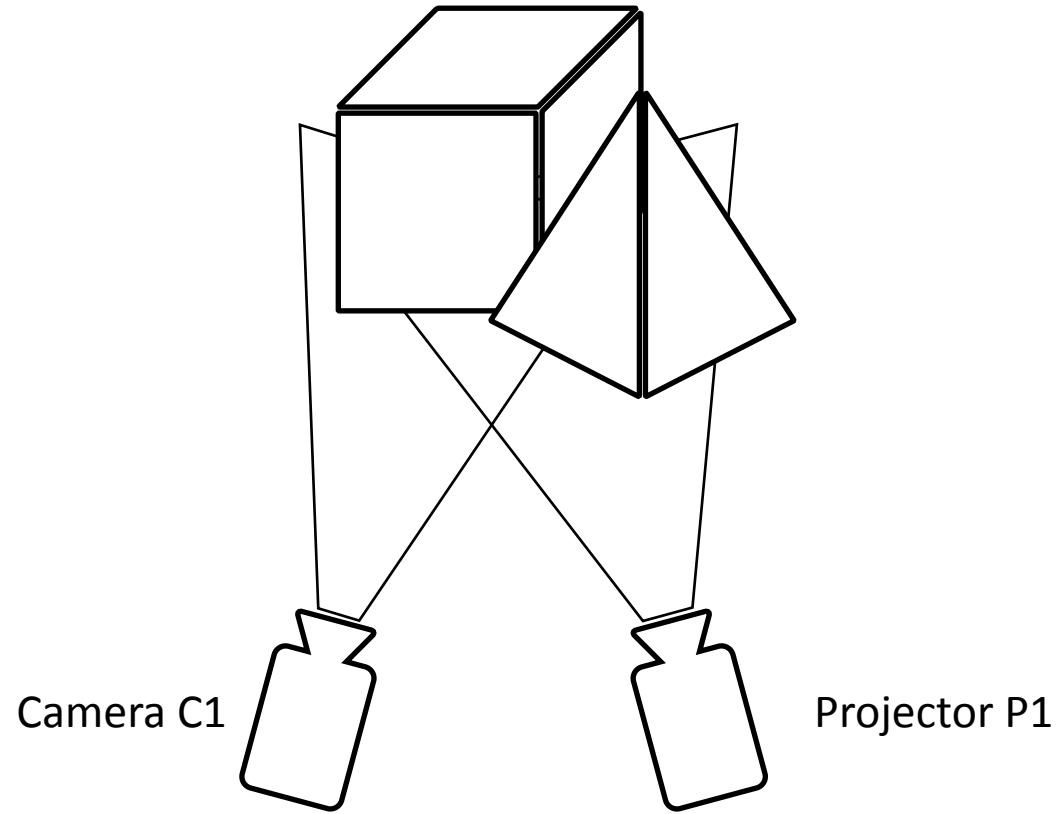
Overview

- Range Scanning
 - Passive:
 - Stereo or Multi-Camera
 - Active:
 - Structured Light Projection
 - Optical Time-of-Flight
 - Direct ToF
 - Indirect ToF
- Registration of Depth Maps
- Processing Depth Maps



Active Range Scanning

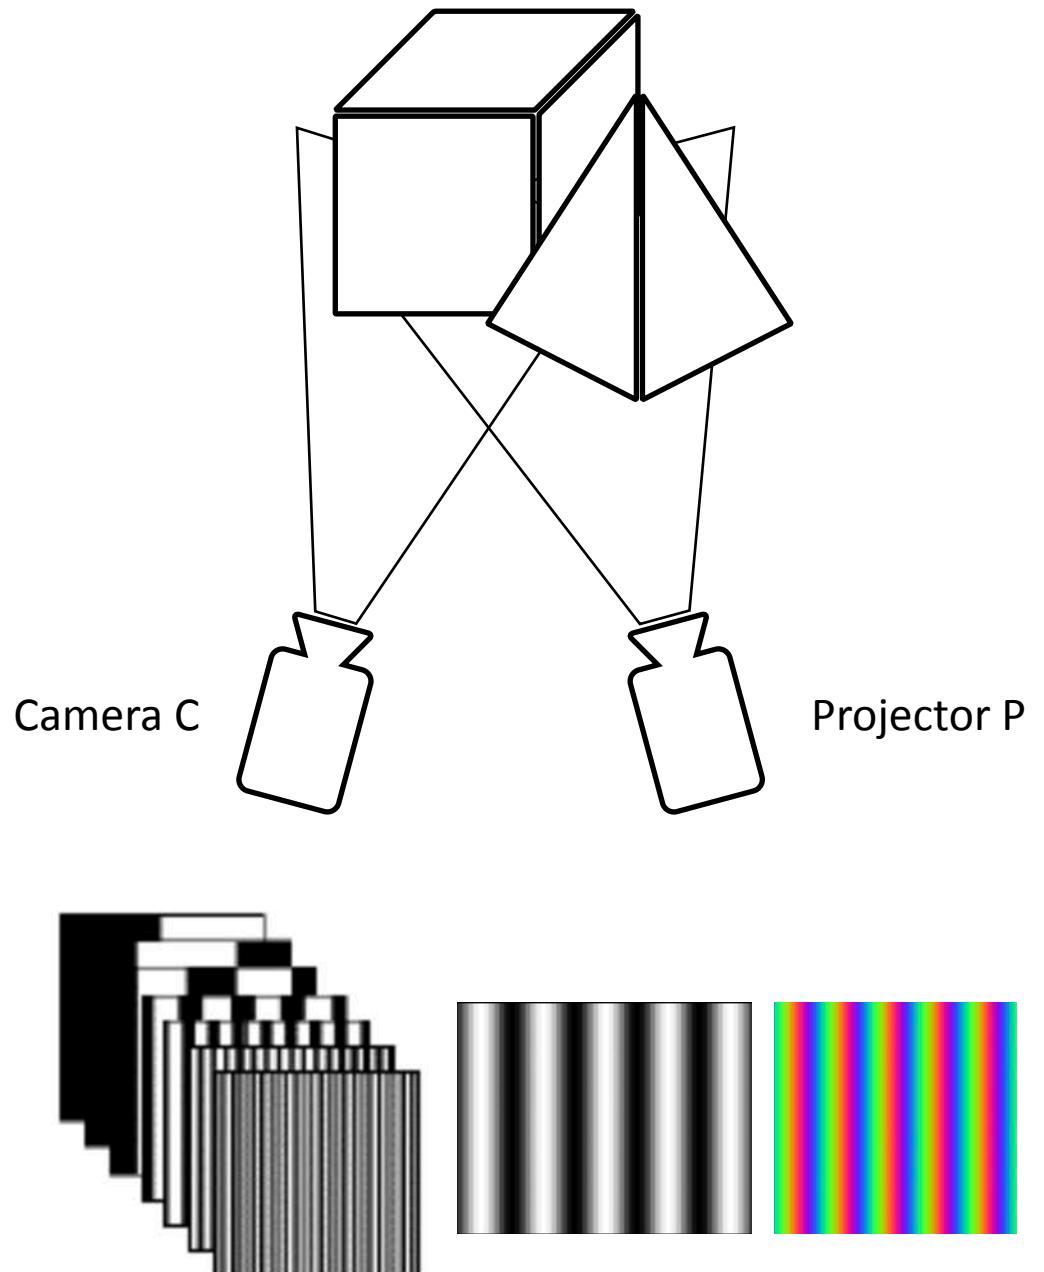
- Why active range scanning?
 - It avoids the correlation problem
- Replace one camera with a projector / light emitter
 - Structured Light Projection
 - Optical Time-of-Flight
 - Direct ToF
 - Indirect ToF



Structured Light Projection

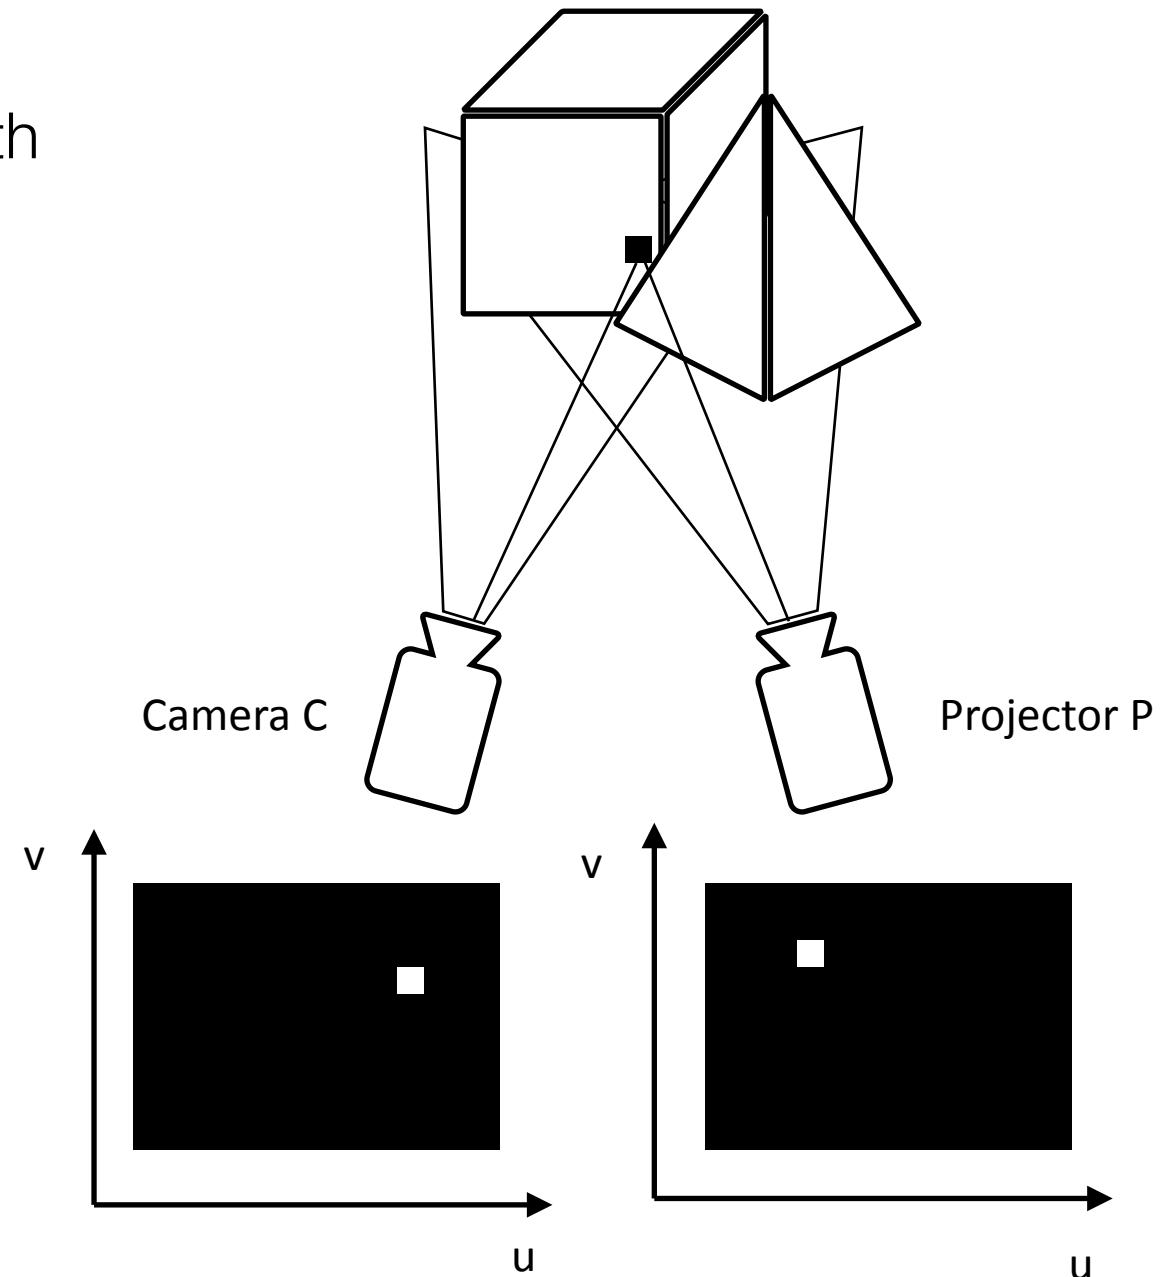
Structured Light Projection

- Many structured light techniques exist that follow the two objectives:
 - Determine projector-camera correspondences with as few images as possible
 - Be robust against surface modulation and noise
- Codes include:
 - Binary codes, such as the Gray code (most common): <100 images, quite robust
 - Intensity codes that apply cosine patterns and phase shifts: 6 images, but are not robust against strong absorptions
 - Color coding: <3 images, but very fragile (does not work on coloured surfaces)



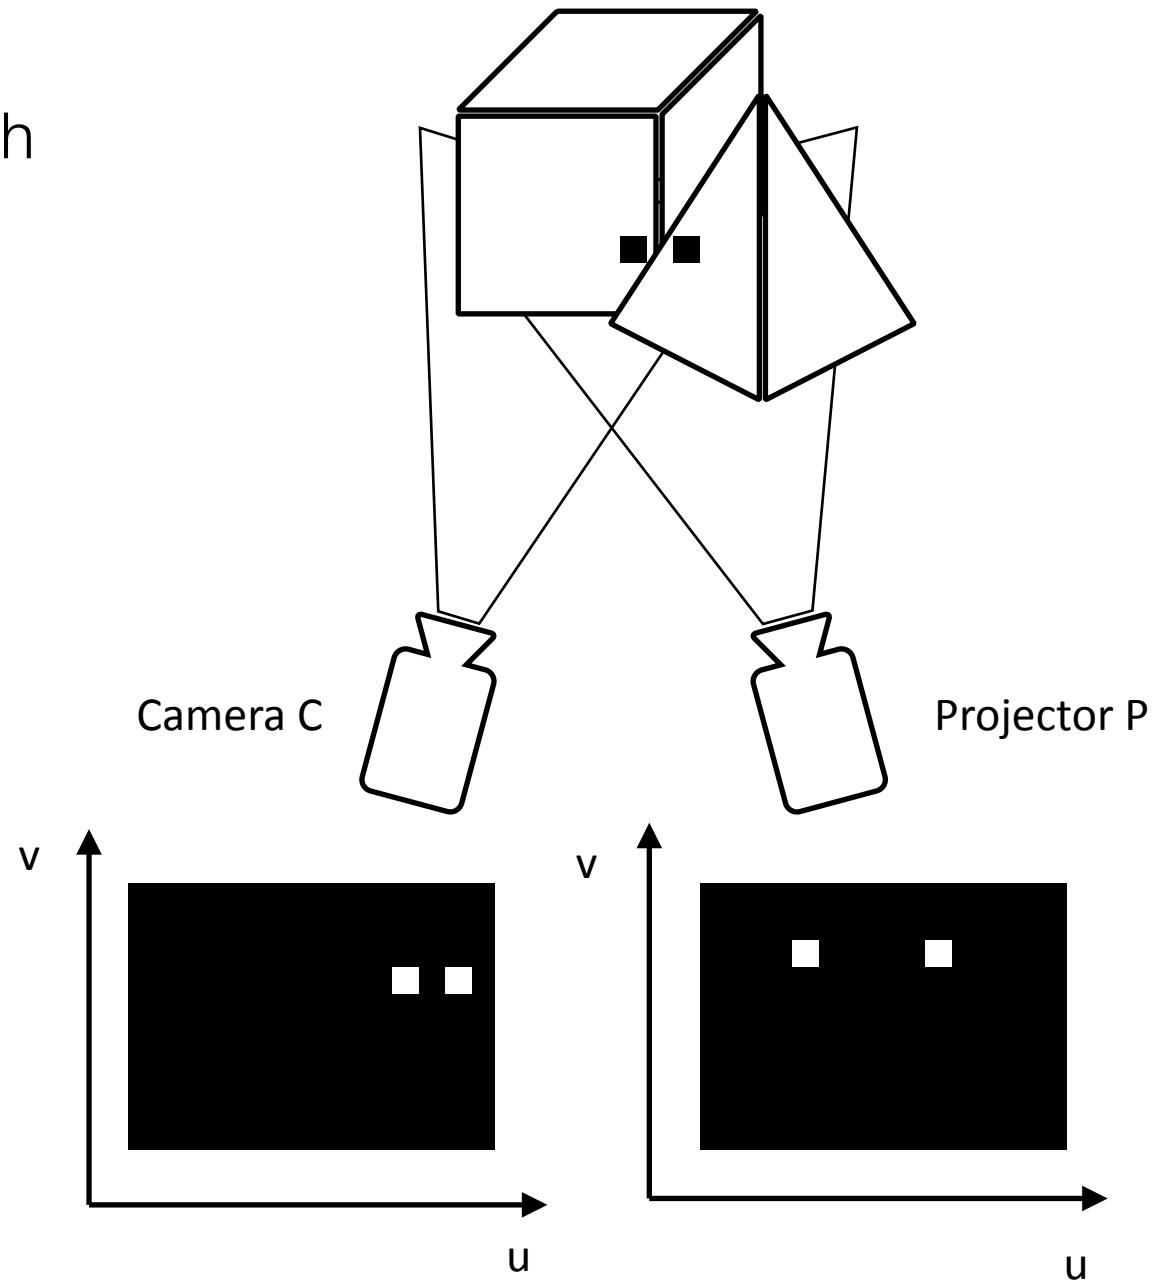
Structured Light Projection

- Let's assume 1 projector + 1 camera (both are calibrated)
- The simplest light pattern is projecting a point at $[u_P, v_P]$ that can be found at $[u_C, v_C]$. This has to be repeated for all points
- What is the problem?
 - Too slow!
 - What is a better approach?



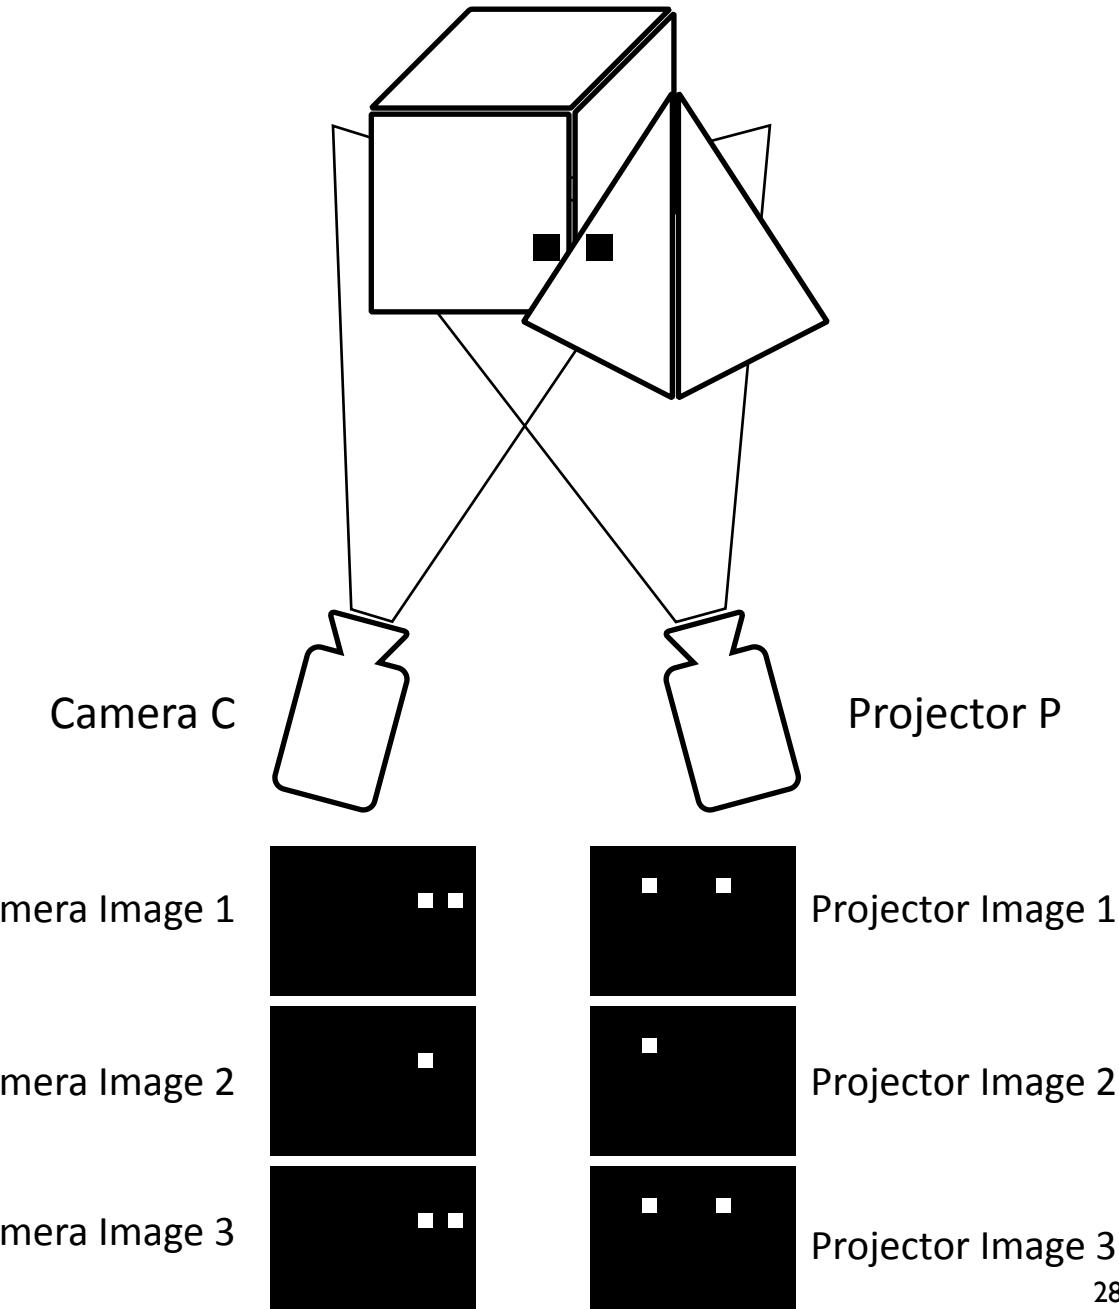
Structured Light Projection

- Let's assume 1 projector + 1 camera (both are calibrated)
- The simplest light pattern is projecting a point at $[u_P, v_P]$ that can be found at $[u_C, v_C]$. This has to be repeated for all points
- What is the problem?
 - Too slow!
 - Projecting two points!
 - But how to distinguish?



Structured Light Projection

- Let's assume 1 projector + 1 camera (both are calibrated)
- The simplest light pattern is projecting a point at $[u_P, v_P]$ that can be found at $[u_C, v_C]$. This has to be repeated for all points
- What is the problem?
 - Too slow!
 - Projecting two points!
 - Project unique “code”
 - But how to scale up?



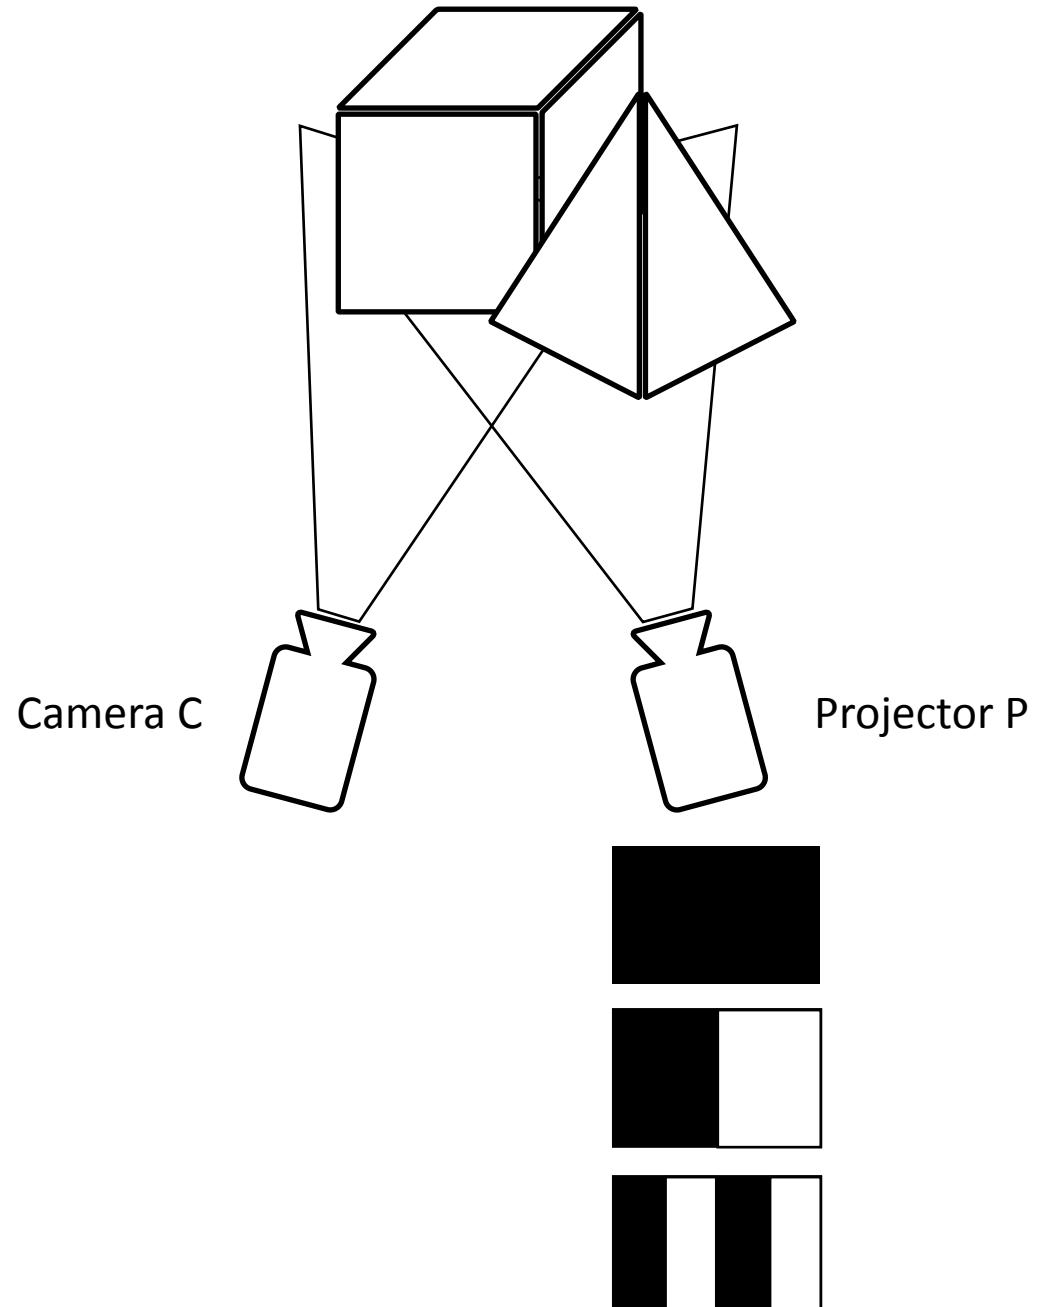
Structured Light Projection



Binary calibration

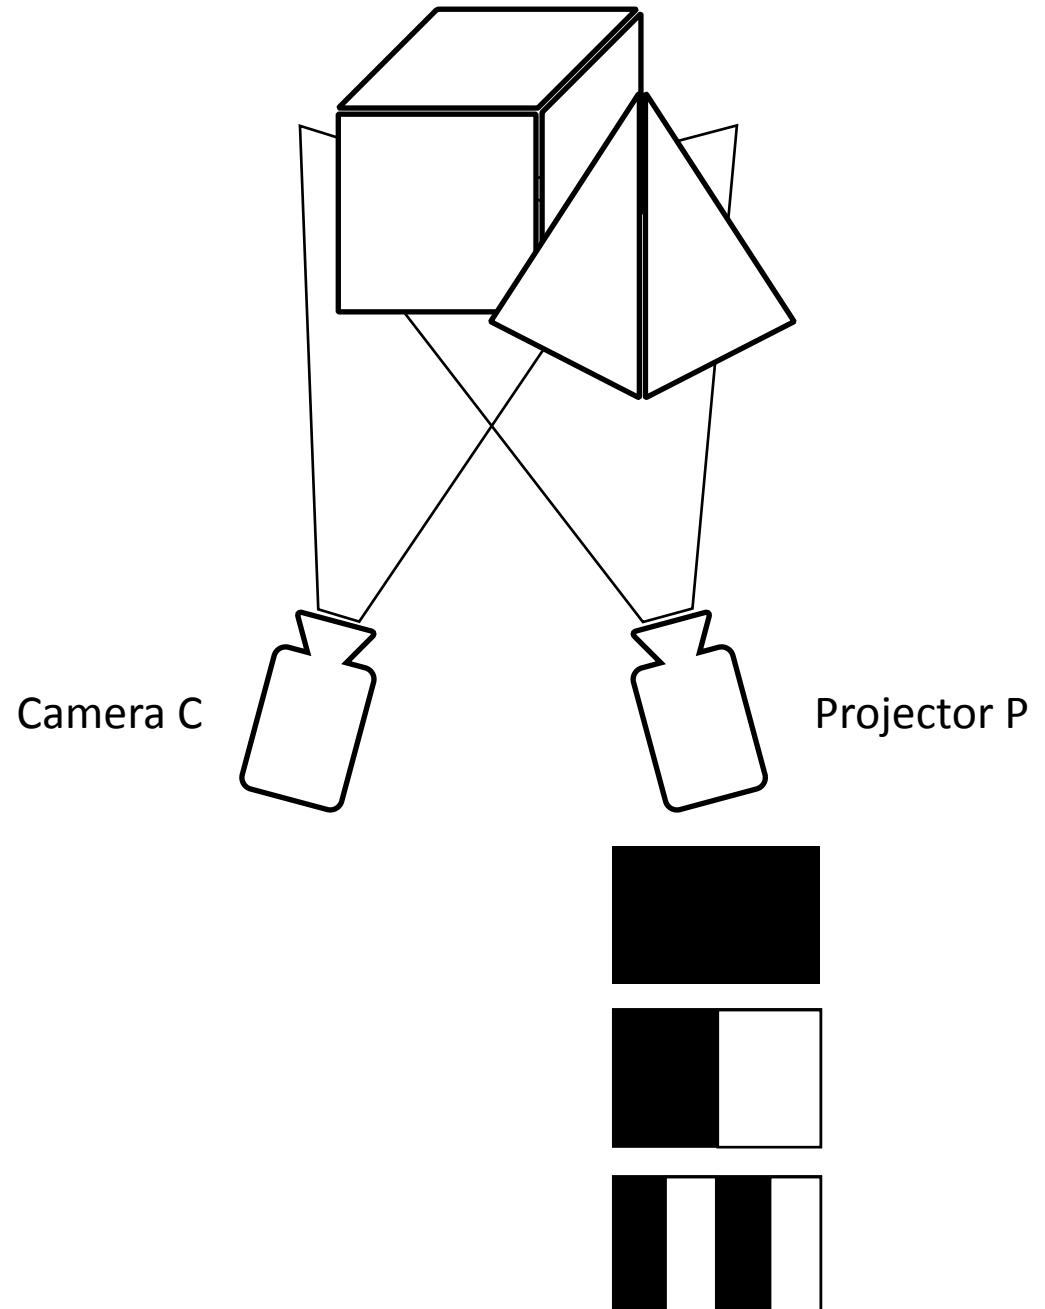
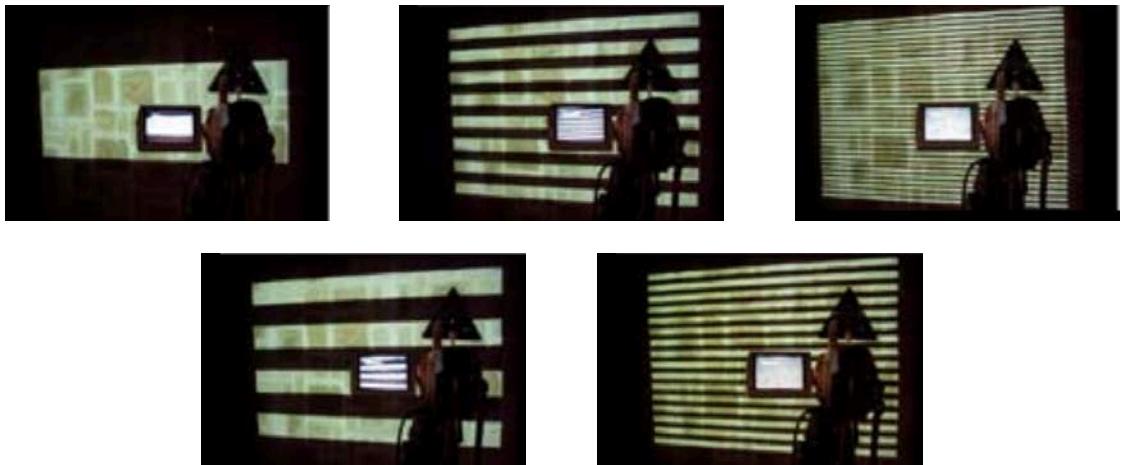
Structured Light Projection

- Let's further optimise:
 - Time multiplexing with line strip scanning
 - Each projected pixel codes a unique ID (e.g. Binary pattern)
 - How many images for 1024px resolution?

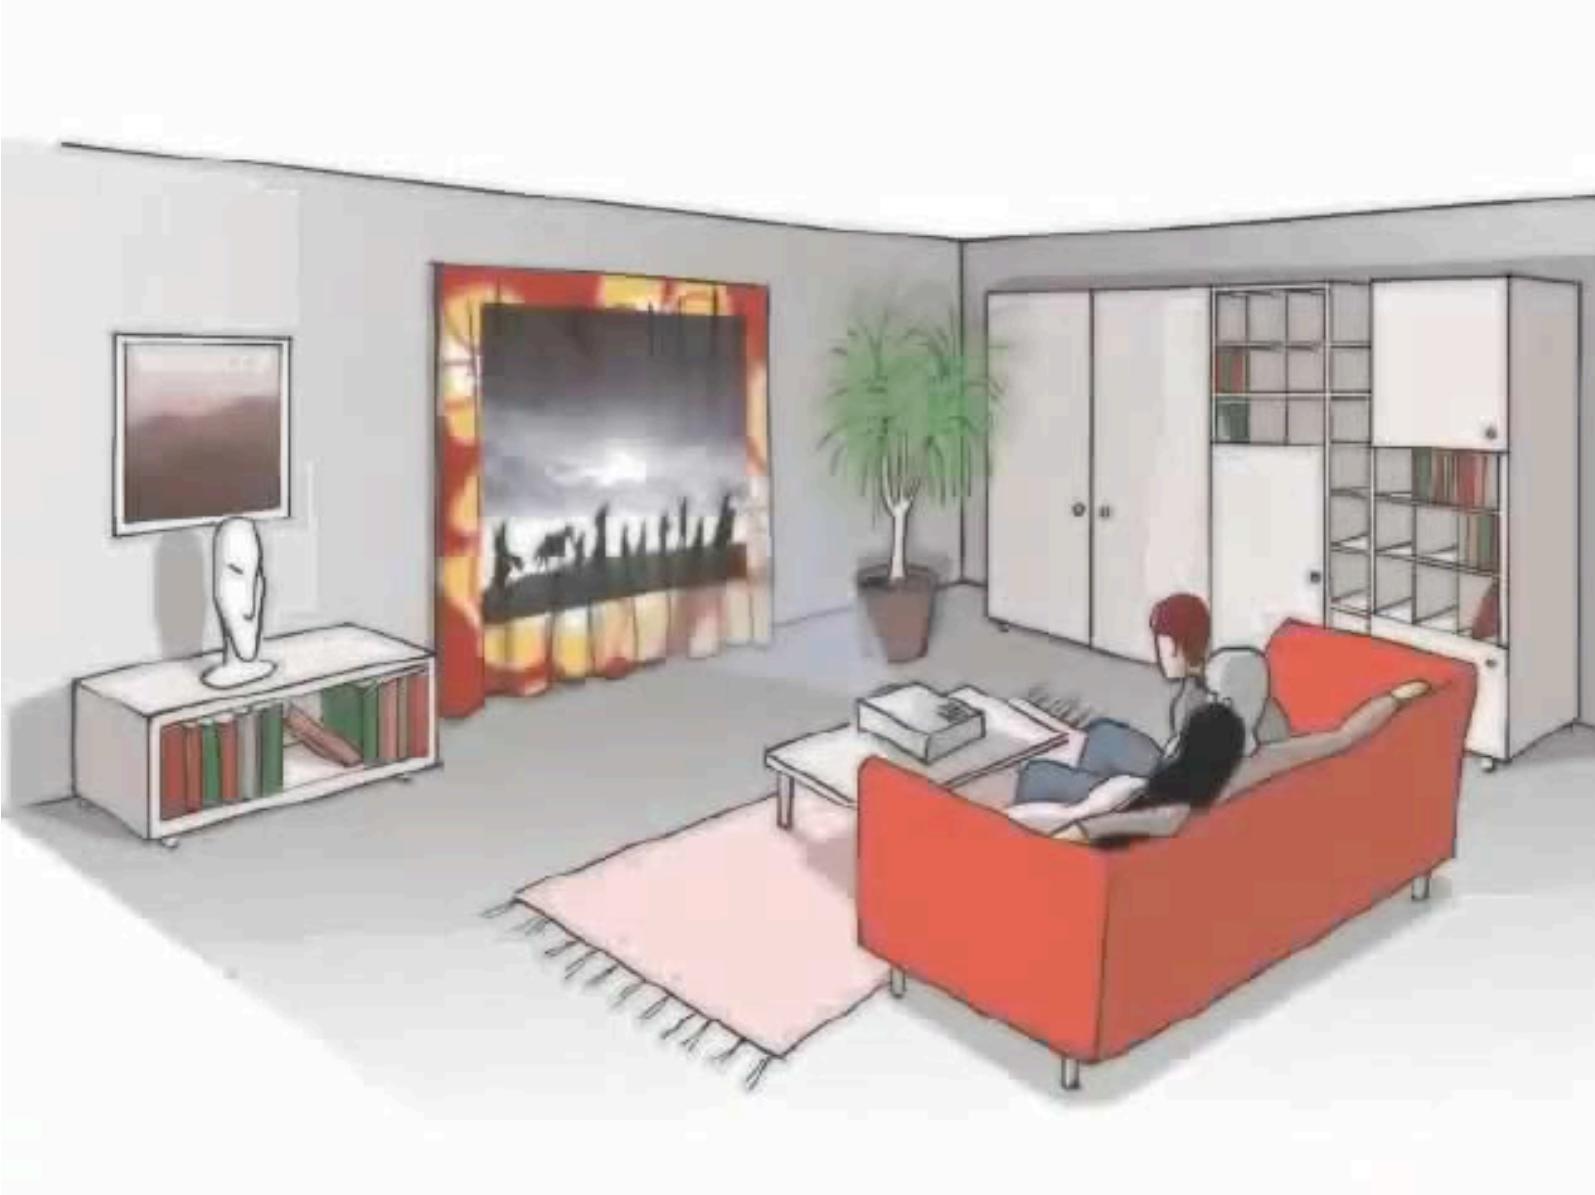


Structured Light Projection

- Let's further optimise:
 - Time multiplexing with line strip scanning
 - Each projected pixel codes a unique ID (e.g. Binary pattern)
 - How many images for 1024px resolution?

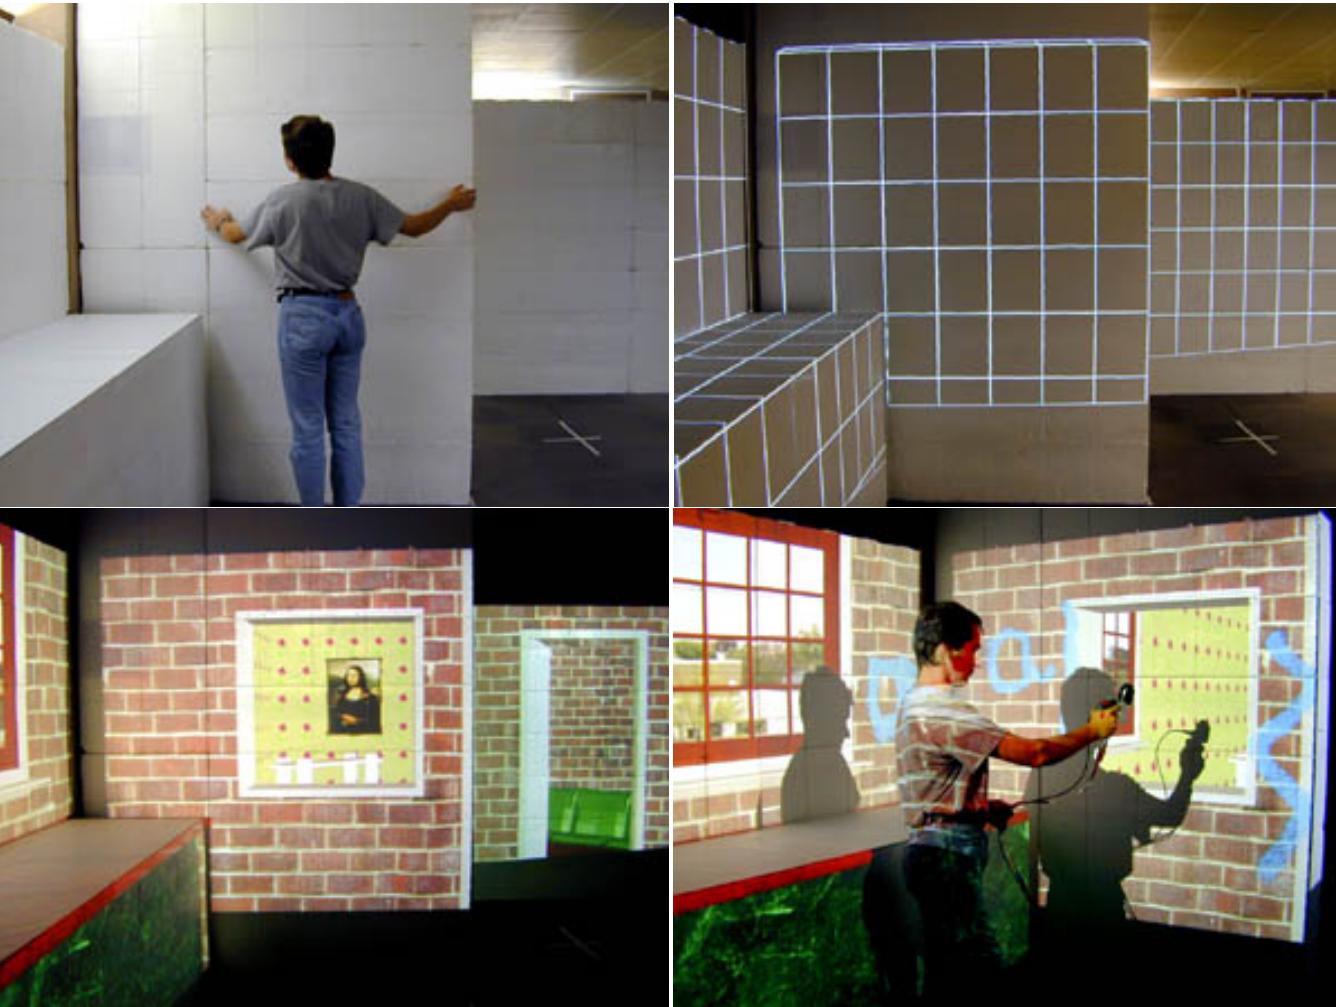


Structured Light Projection



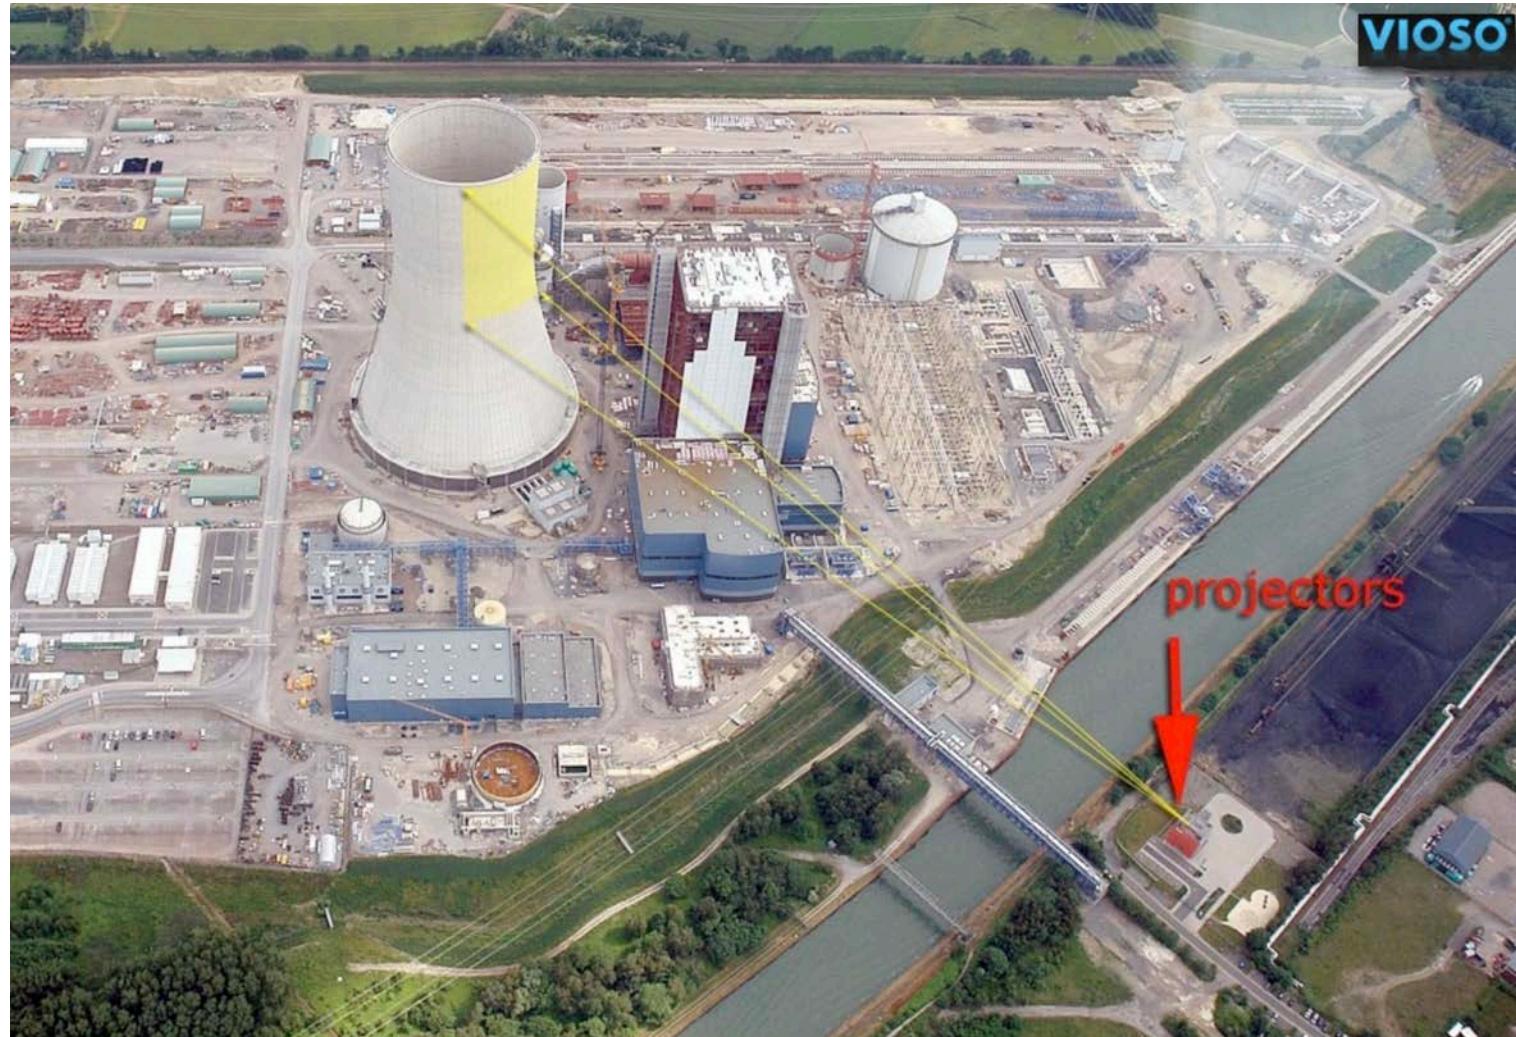
Smart projectors

Structured Light Projection



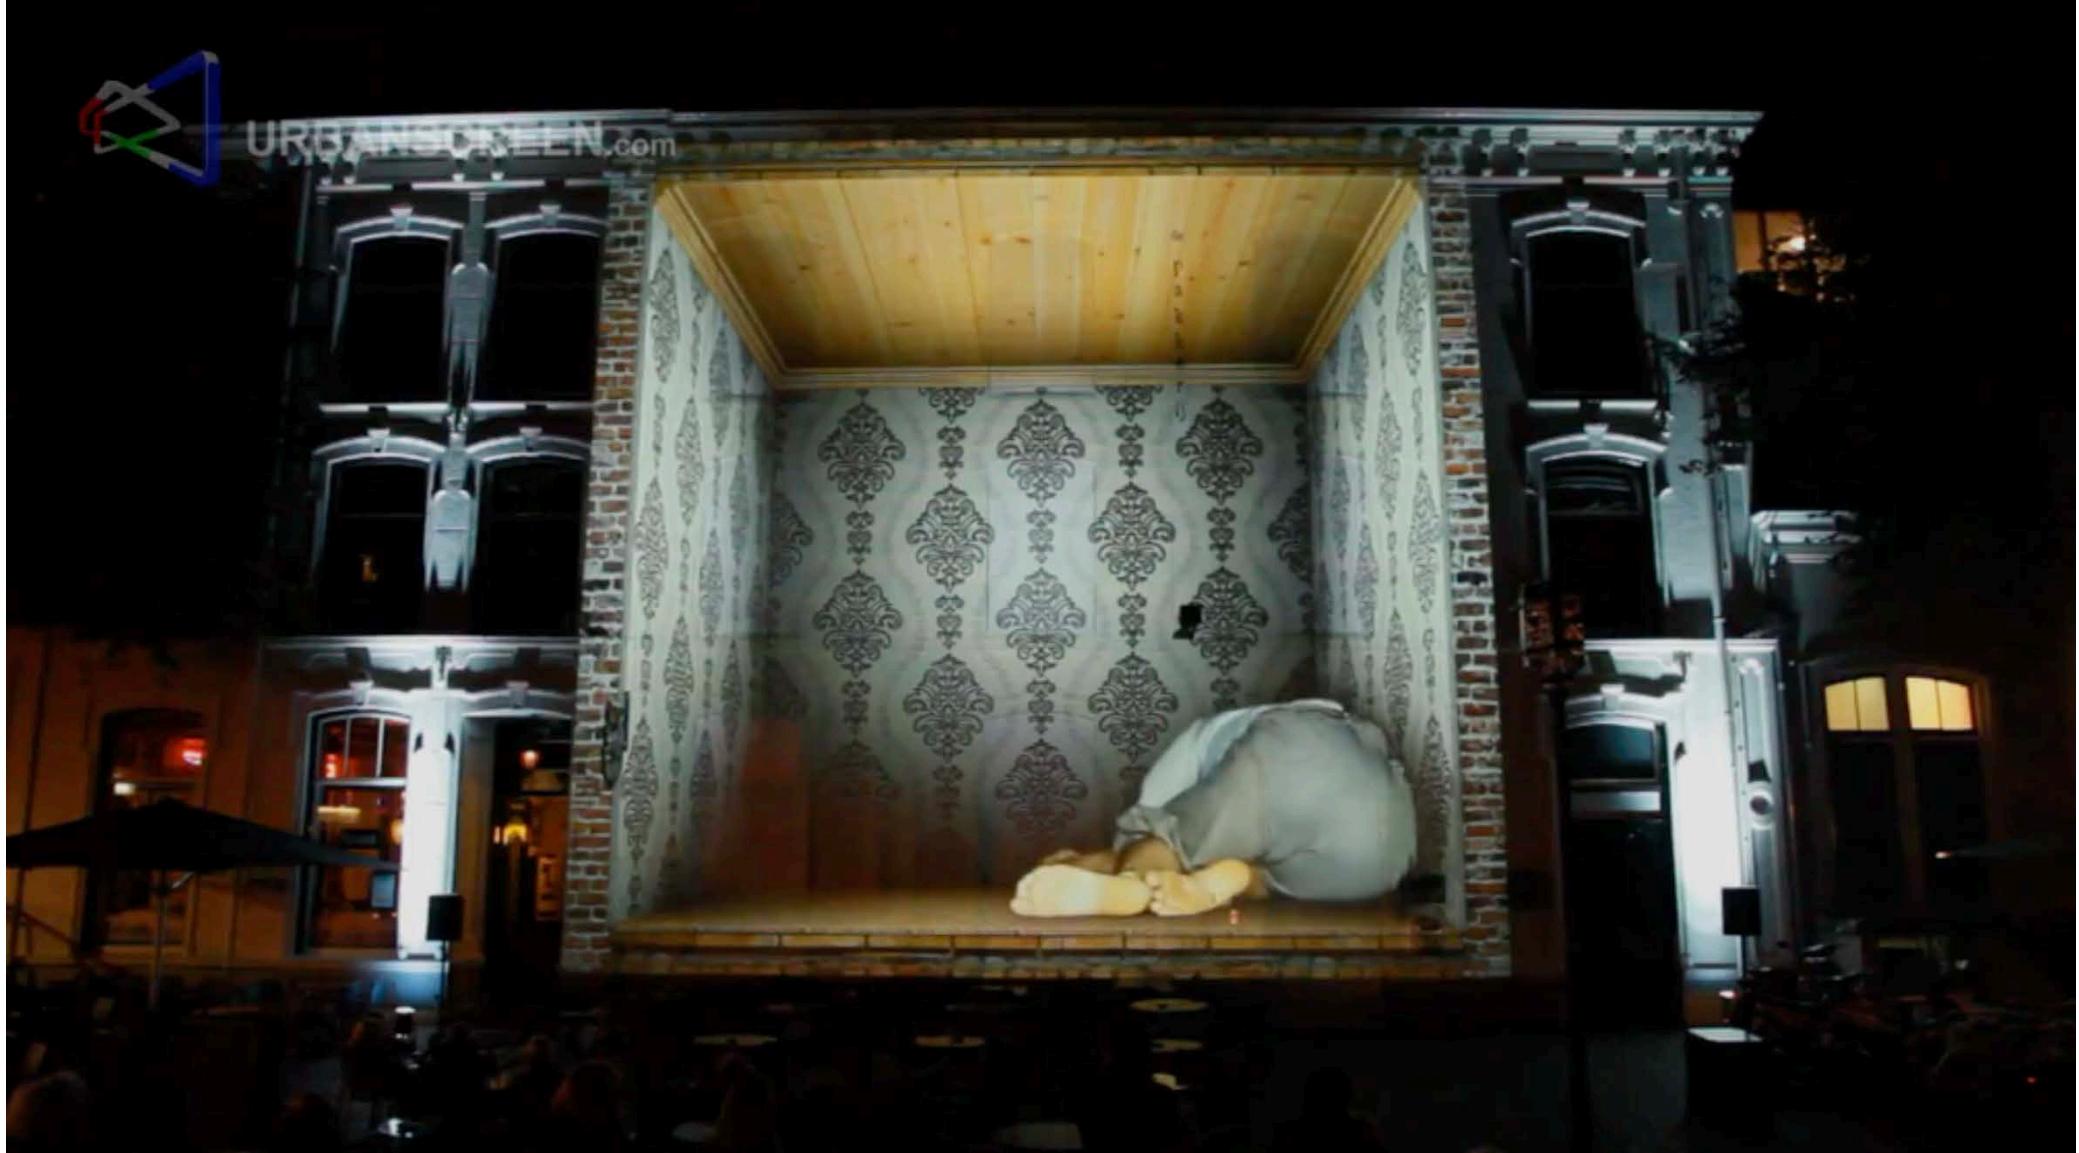
Raskat et al. "Shaderlamps"

Structured Light Projection



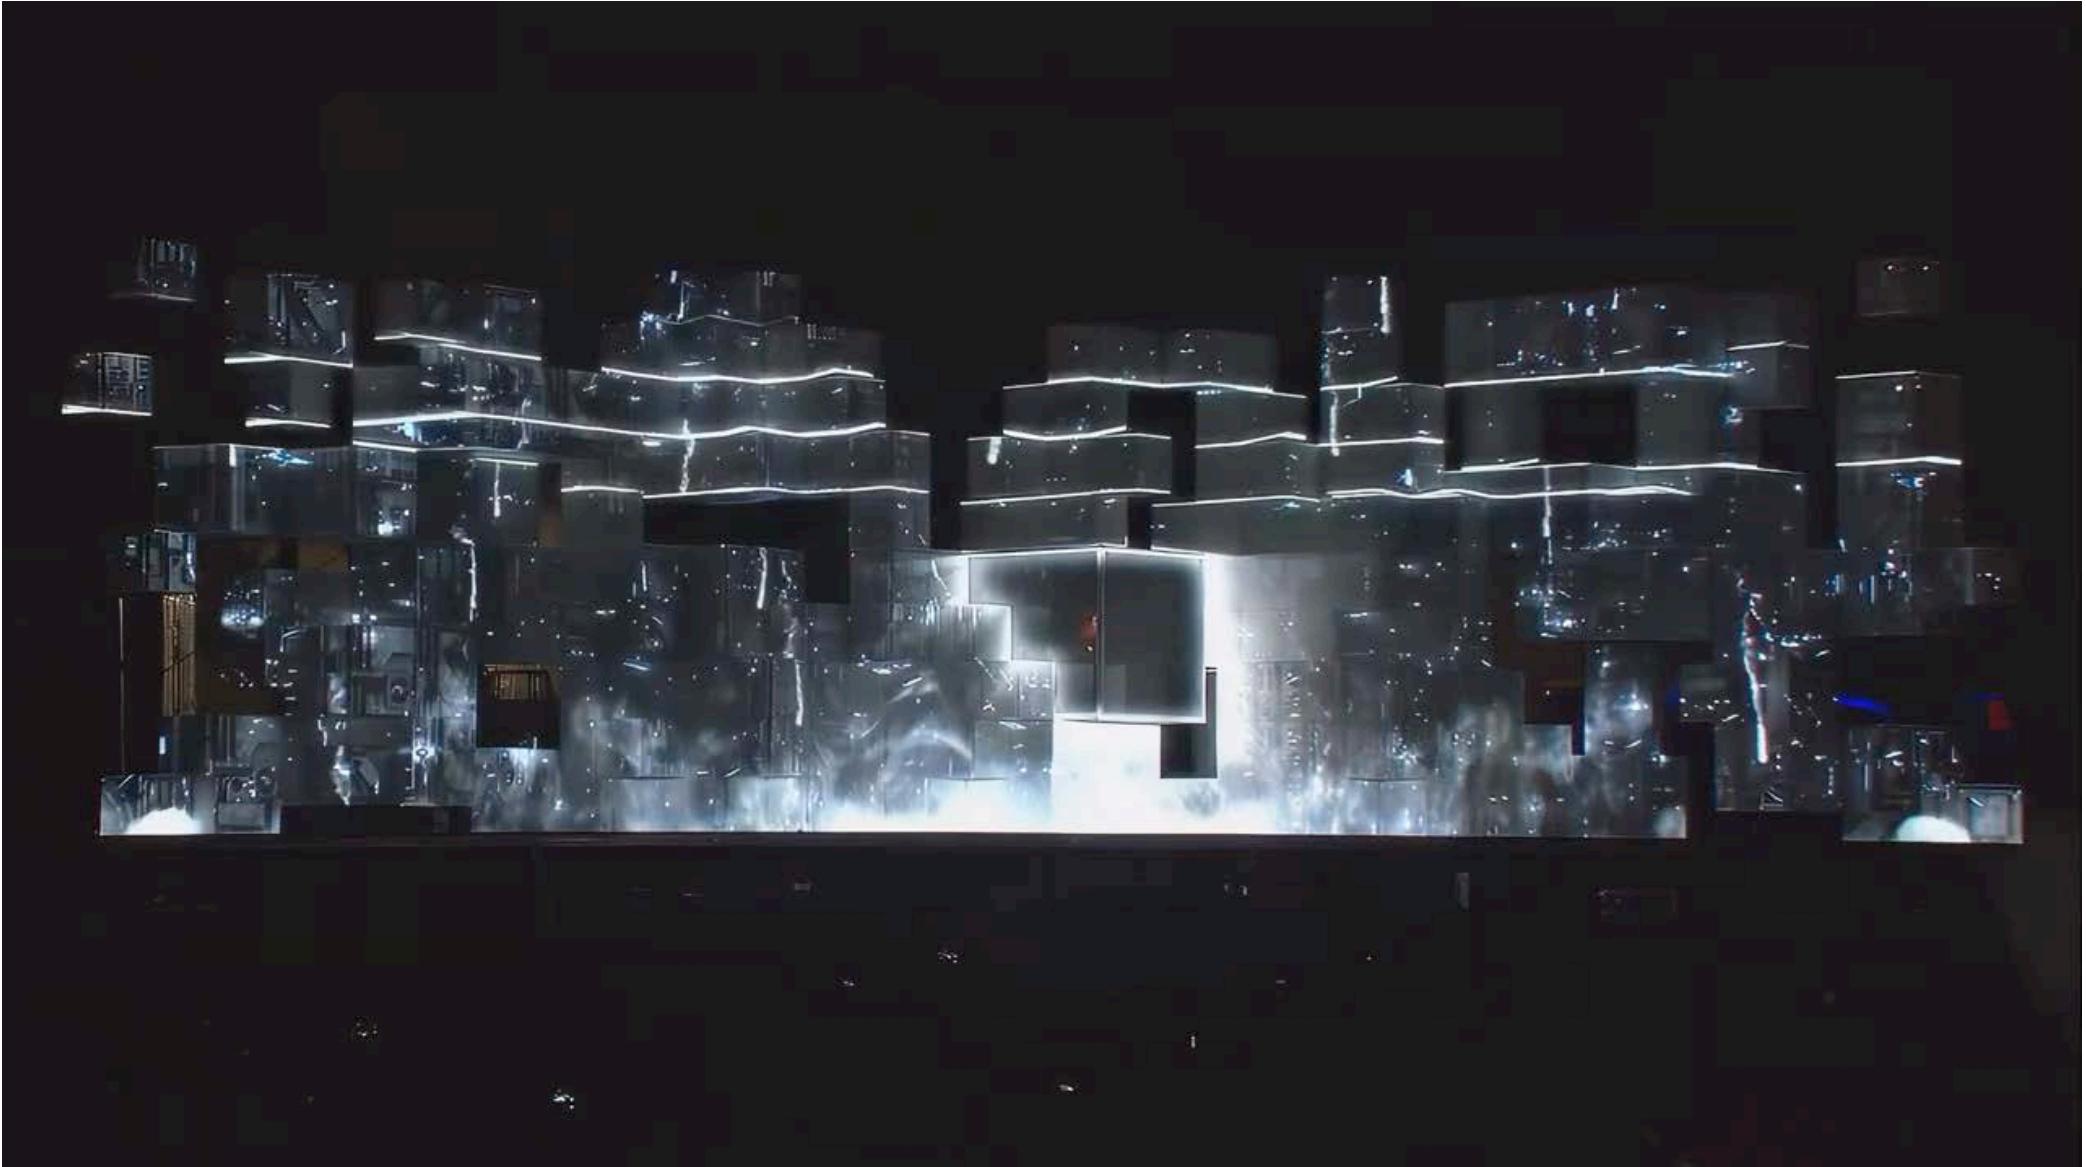
Smart projectors / Vioso

Structured Light Projection



Smart projectors / Vioso

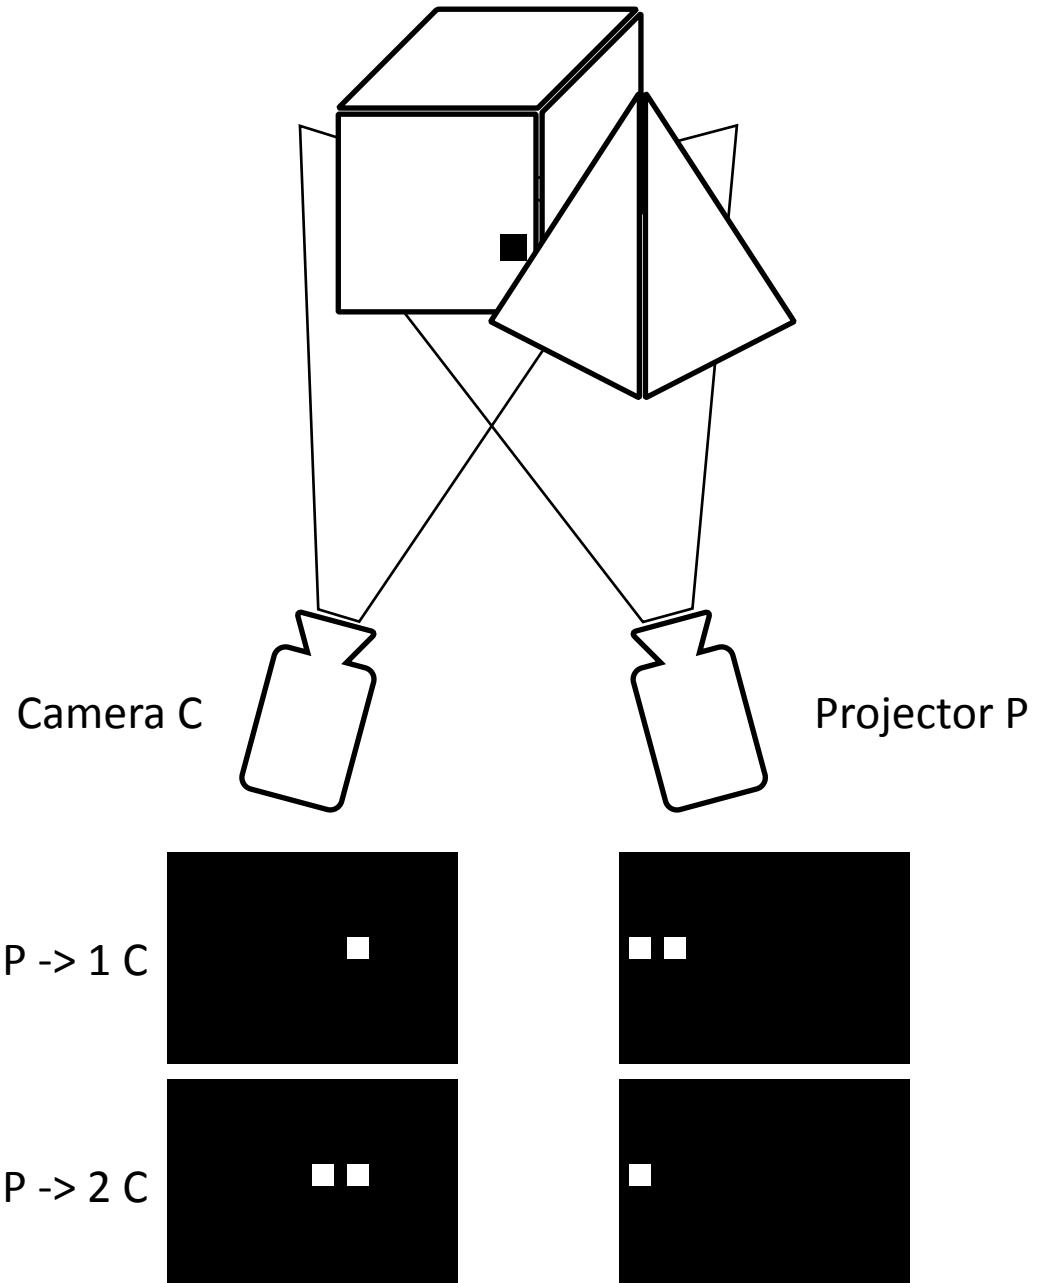
Structured Light Projection



Amob Tobin

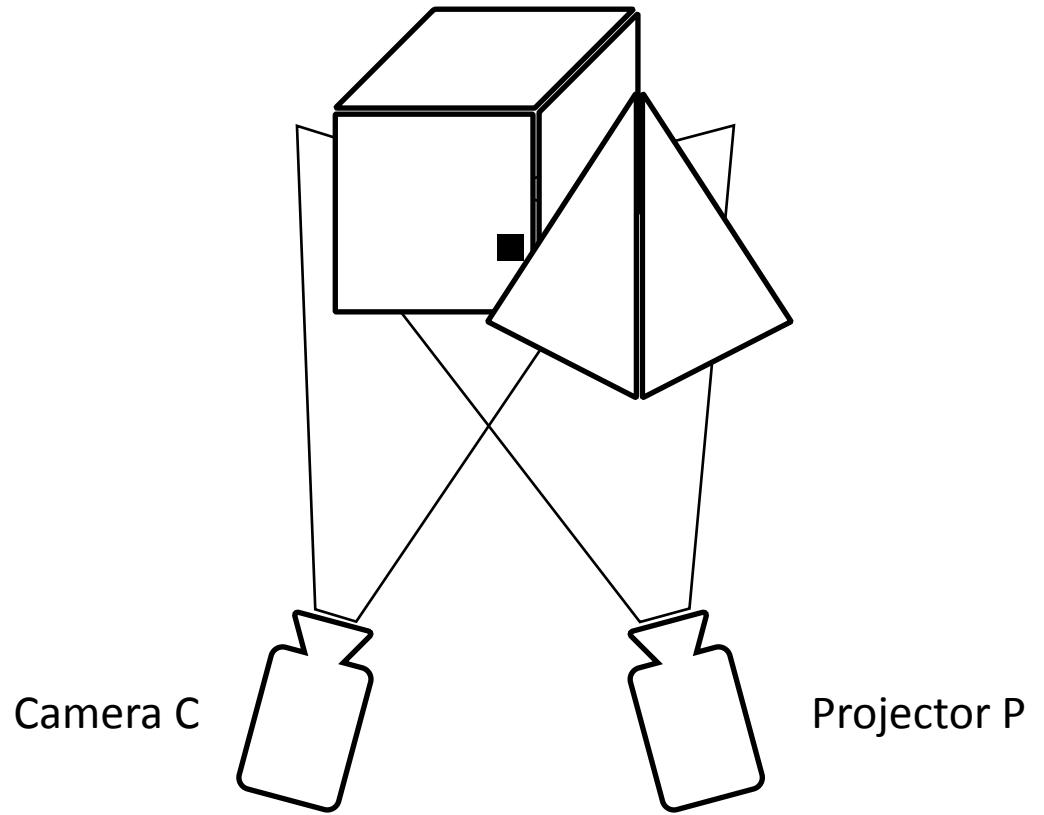
Structured Light Projection

- Other issues:
 - Results usually not in 1-to-1 mappings!
 - Multiple projector pixels can map on same camera pixels / multiple camera pixels can be addressed by single projector pixel
 - Solution?
 - Multiple entries must be averaged



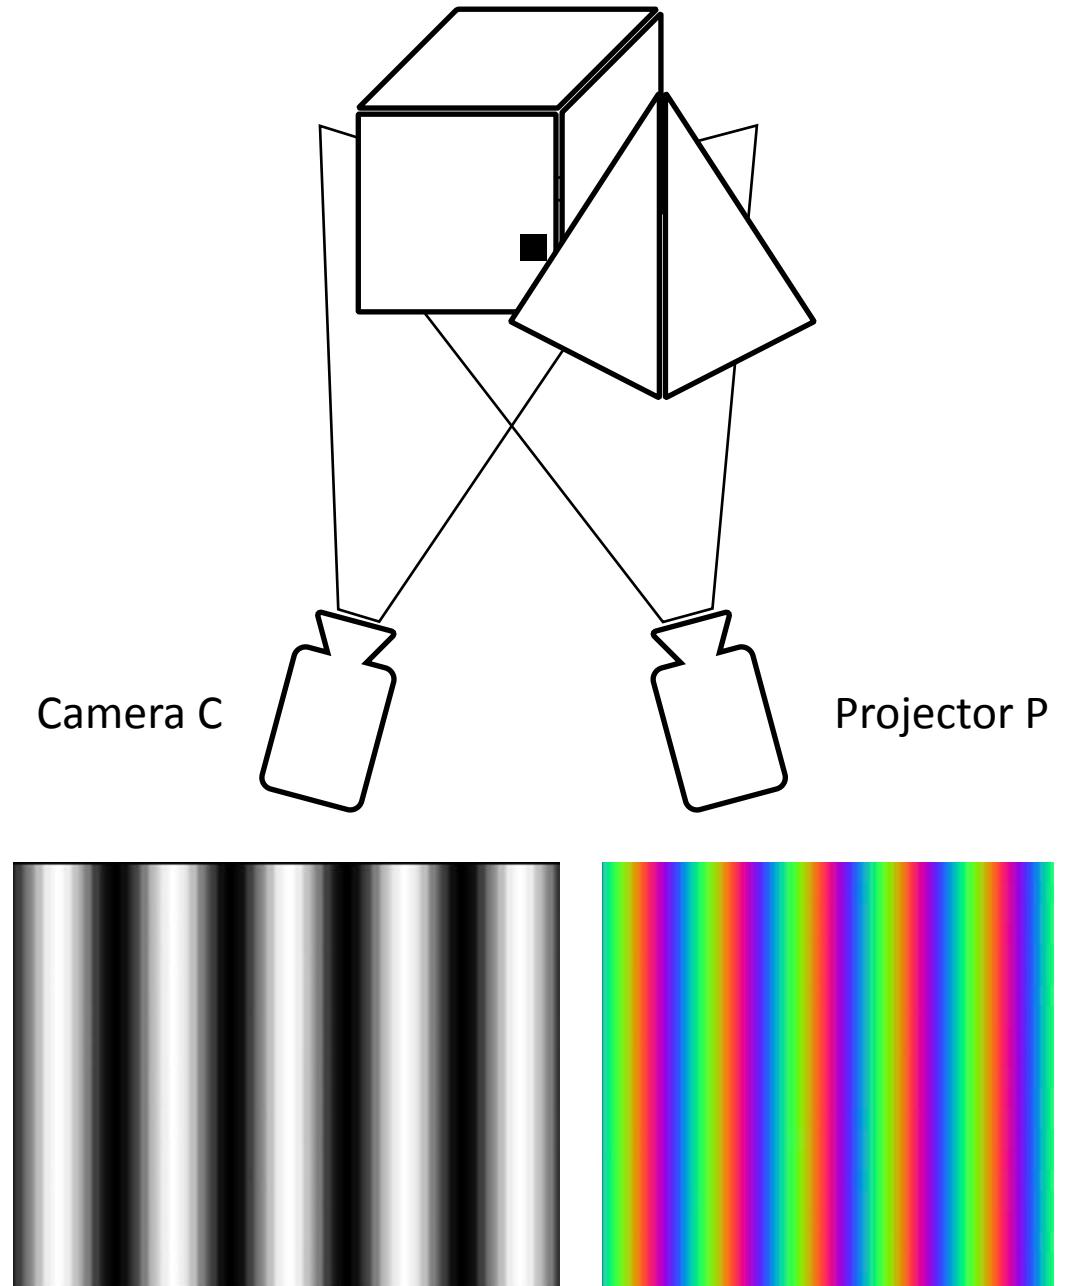
Structured Light Projection

- Other issues:
 - Results usually not in 1-to-1 mappings!
 - Multiple projector pixels can map on same camera pixels / multiple camera pixels can be addressed by single projector pixel
 - Solution?
 - Multiple entries must be averaged
- Any other ideas for improvements?



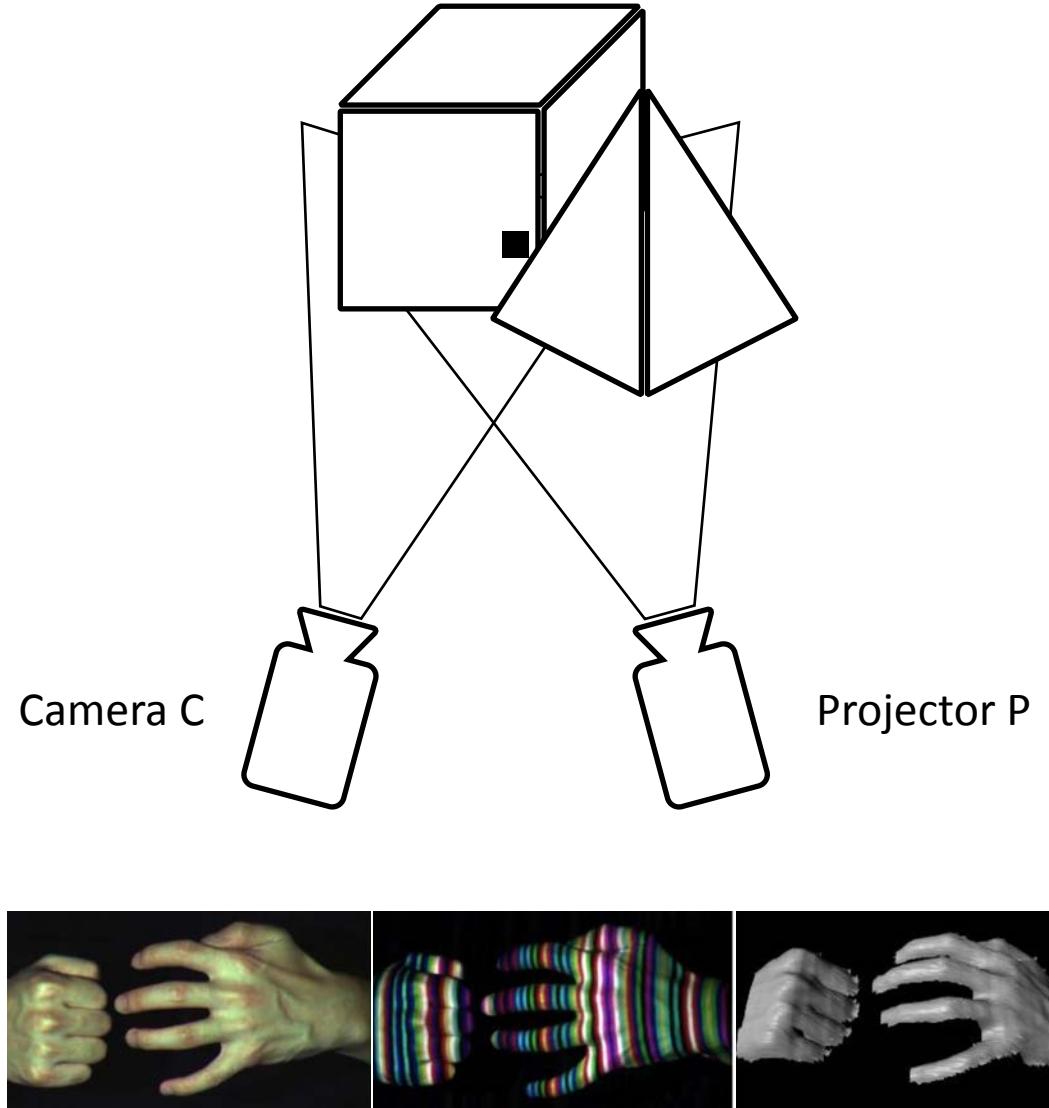
Structured Light Projection

- Other issues:
 - Results usually not in 1-to-1 mappings!
 - Multiple projector pixels can map on same camera pixels / multiple camera pixels can be addressed by single projector pixel
 - Solution?
 - Multiple entries must be averaged
- Any other ideas for improvements?
 - Encode more information into structured projection
 - For example: colours and intensities



Structured Light Projection

- Other issues:
 - Results usually not in 1-to-1 mappings!
 - Multiple projector pixels can map on same camera pixels / multiple camera pixels can be addressed by single projector pixel
 - Solution?
 - Multiple entries must be averaged
- Any other ideas for improvements?
 - Encode more information into structured projection
 - For example: colours and intensities

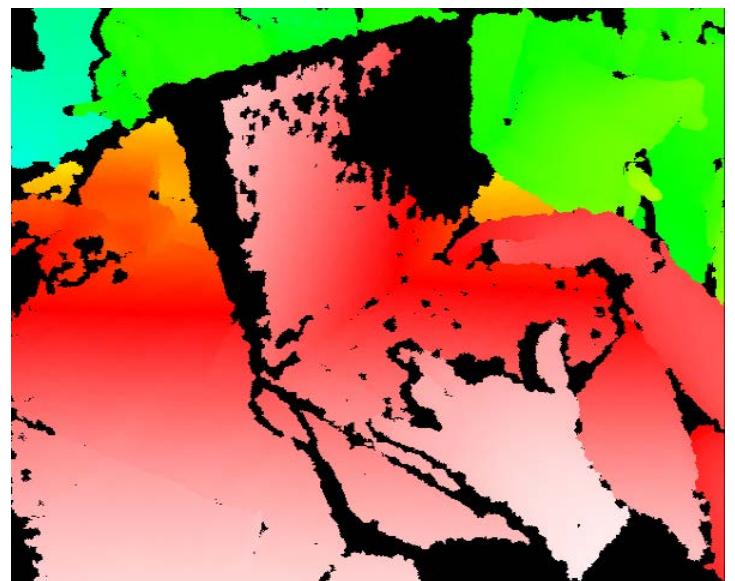


Structured Light Projection



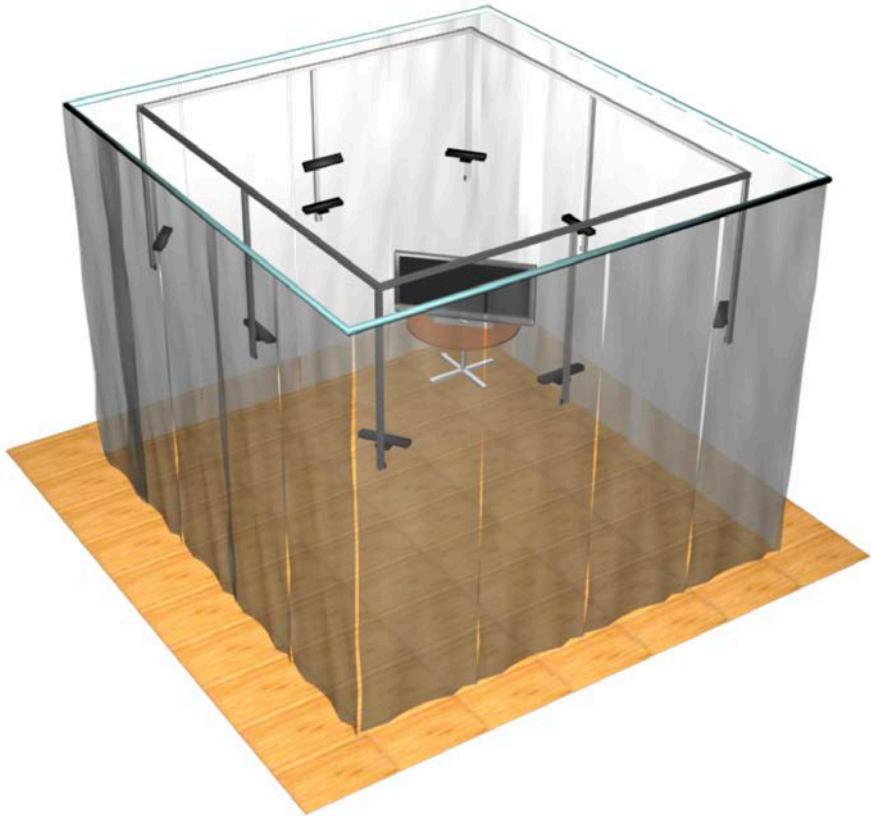
<https://fts.gr/structured-light-3d-scanner/>

Structured Light Projection



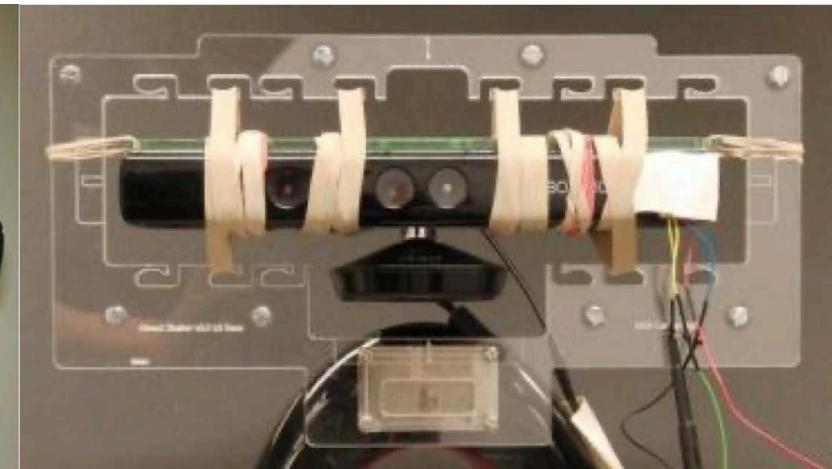
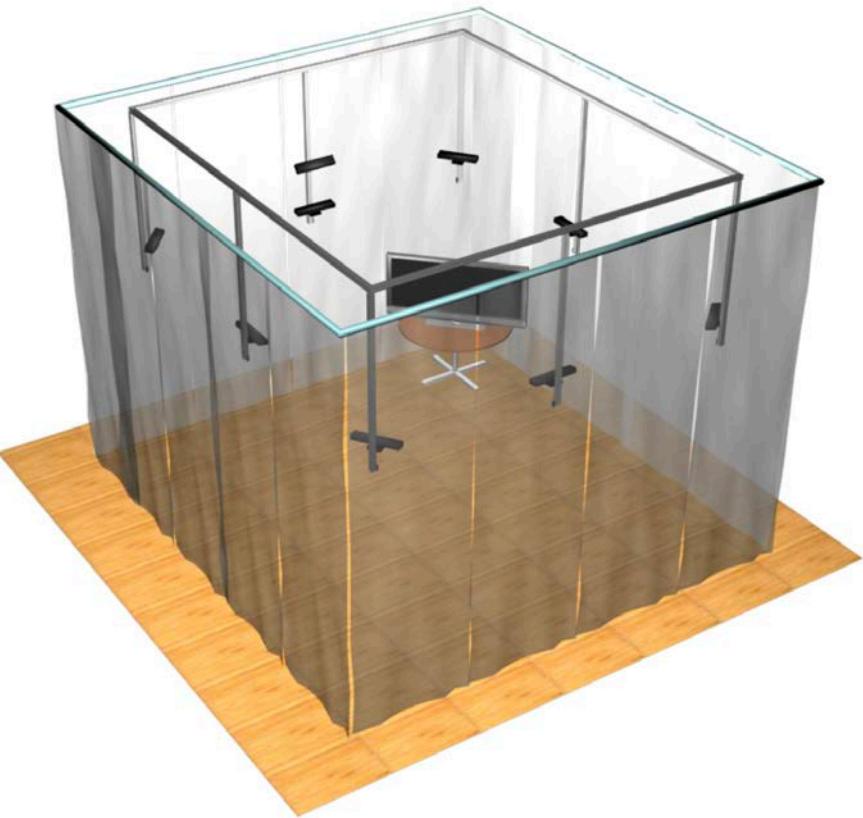
MS Kinect 1 (PrimeSense sensor, not MS Kinect 2 uses ToF discussed later)

Structured Light Projection



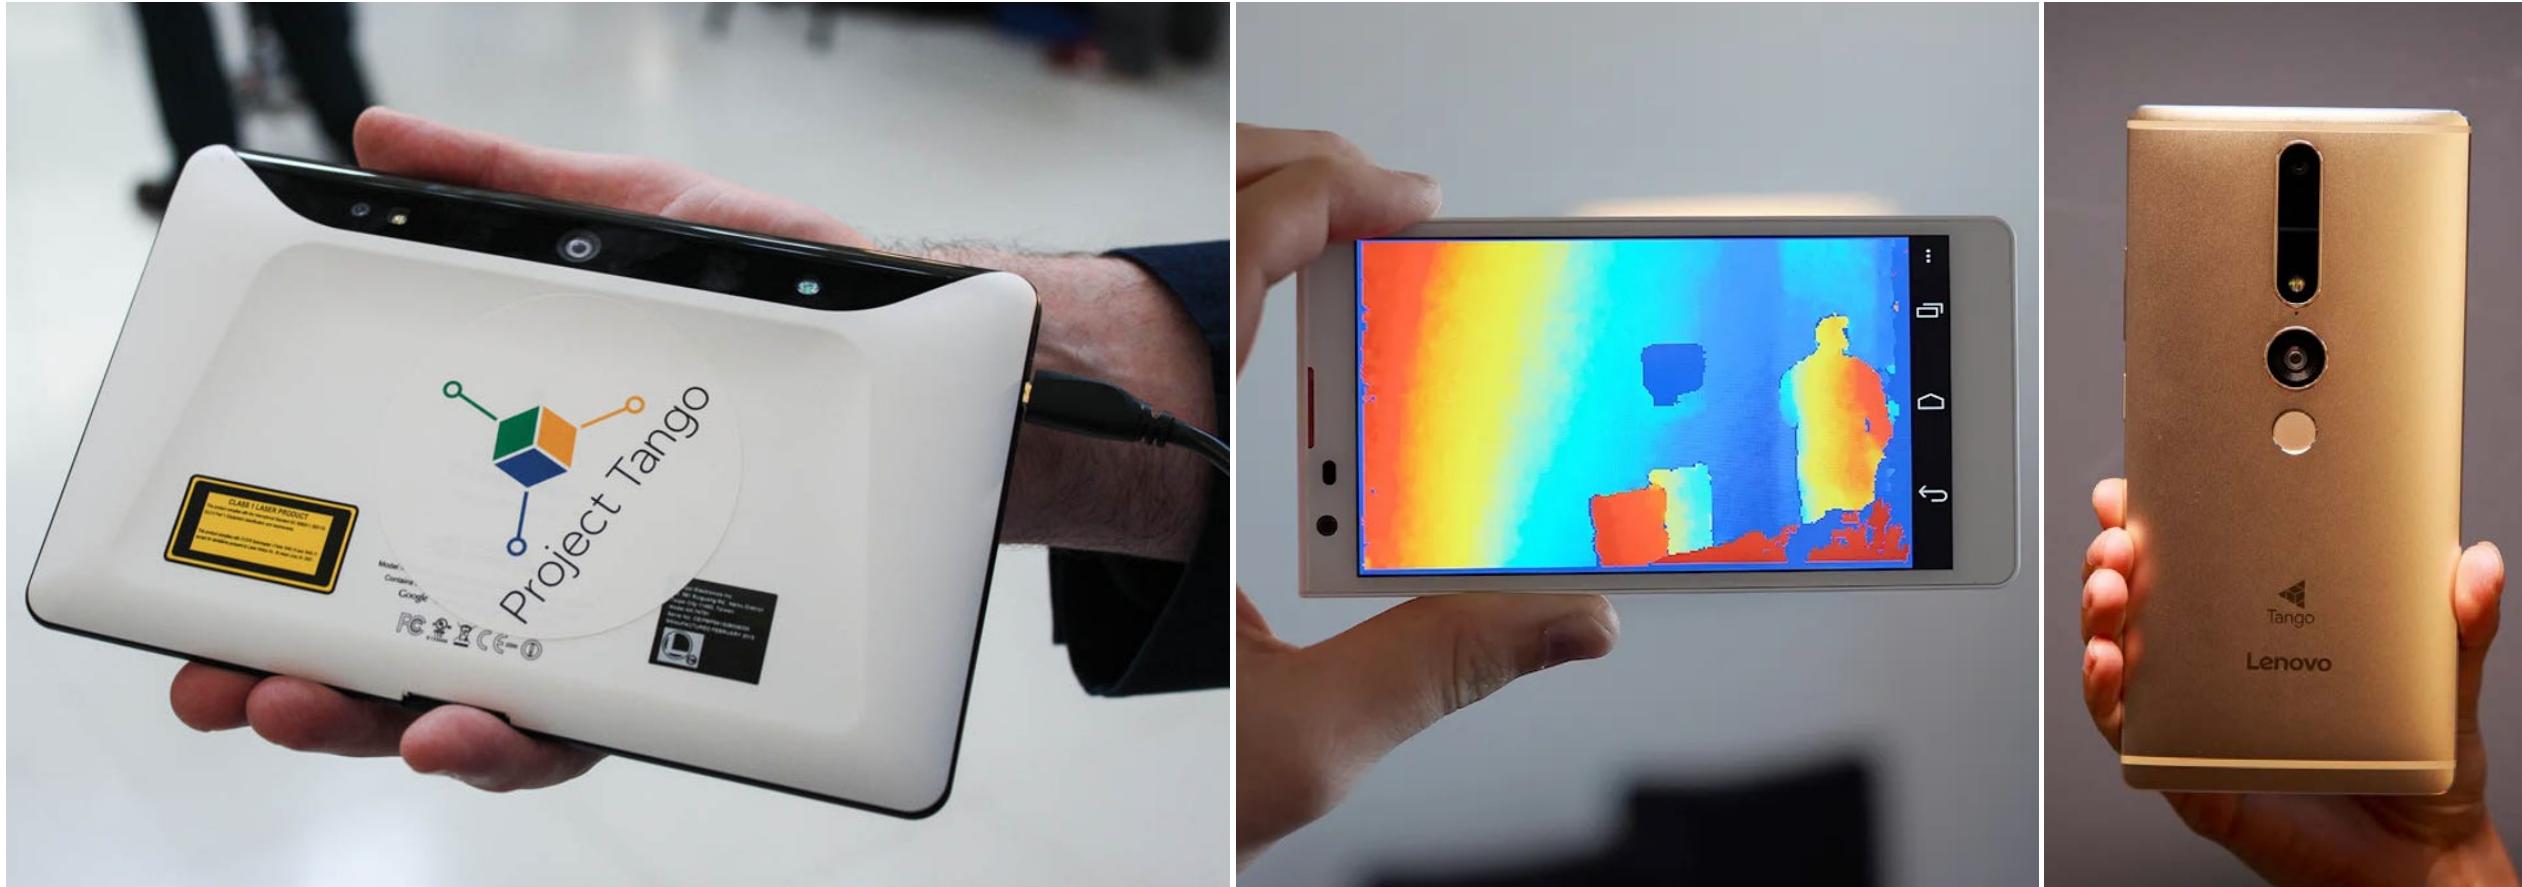
Maimone and Fuchs “Reducing Interference Between Multiple Structured Light Depth Sensors Using Motion”

Structured Light Projection



Maimone and Fuchs “Reducing Interference Between Multiple Structured Light Depth Sensors Using Motion”

Structured Light Projection



Google Project Tango (also using ToF discussed later)

Structured Light Projection



Mobile Portation

Structured Light Projection

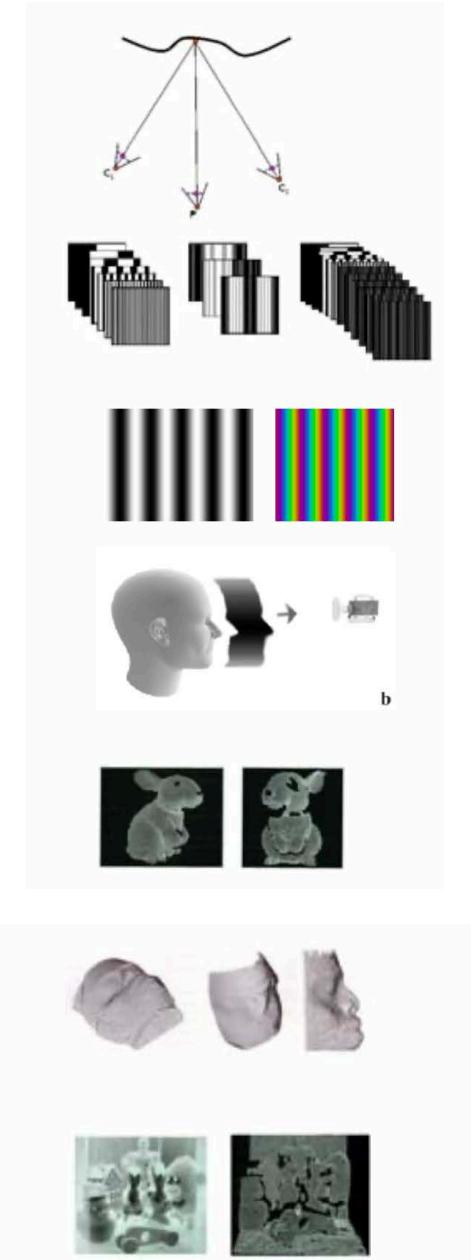


Mobile Portation

Time of Flight (ToF)

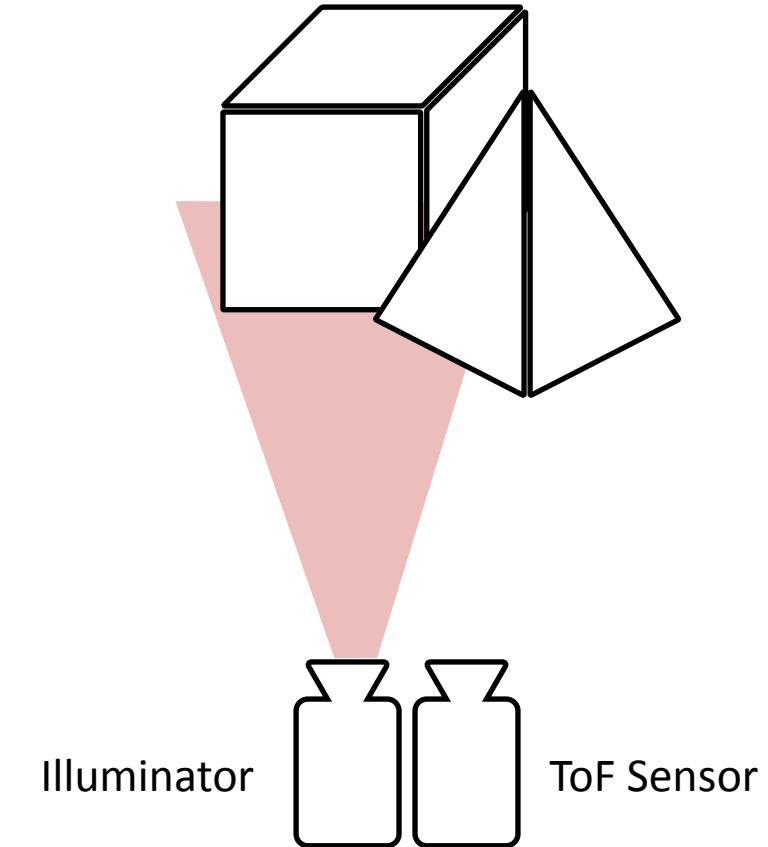
Overview

- Range Scanning
 - Passive:
 - Stereo or Multi-Camera
 - Active:
 - Structured Light Projection
 - Optical Time-of-Flight
 - Direct ToF
 - Indirect ToF
- Registration of Depth Maps
- Processing Depth Maps



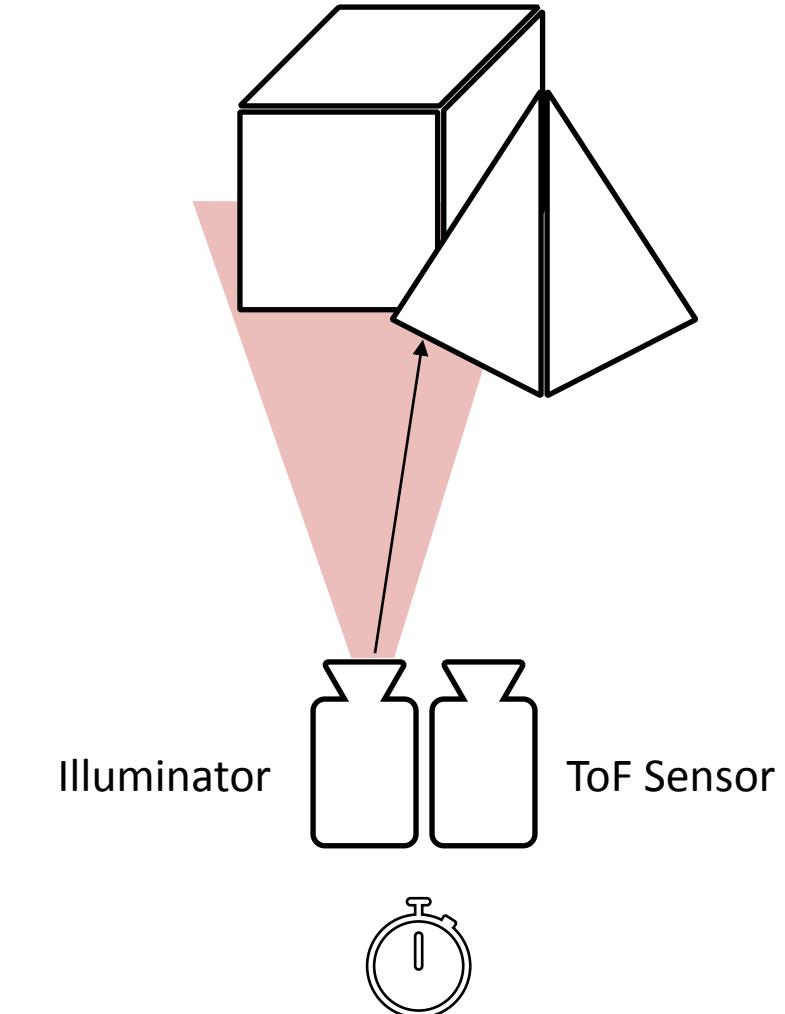
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t/2$
- Use low-power IR illumination for pulsed illumination
 - Can be switched on/off with <1ns
 - Generates a “light wall” that is reflected back by objects



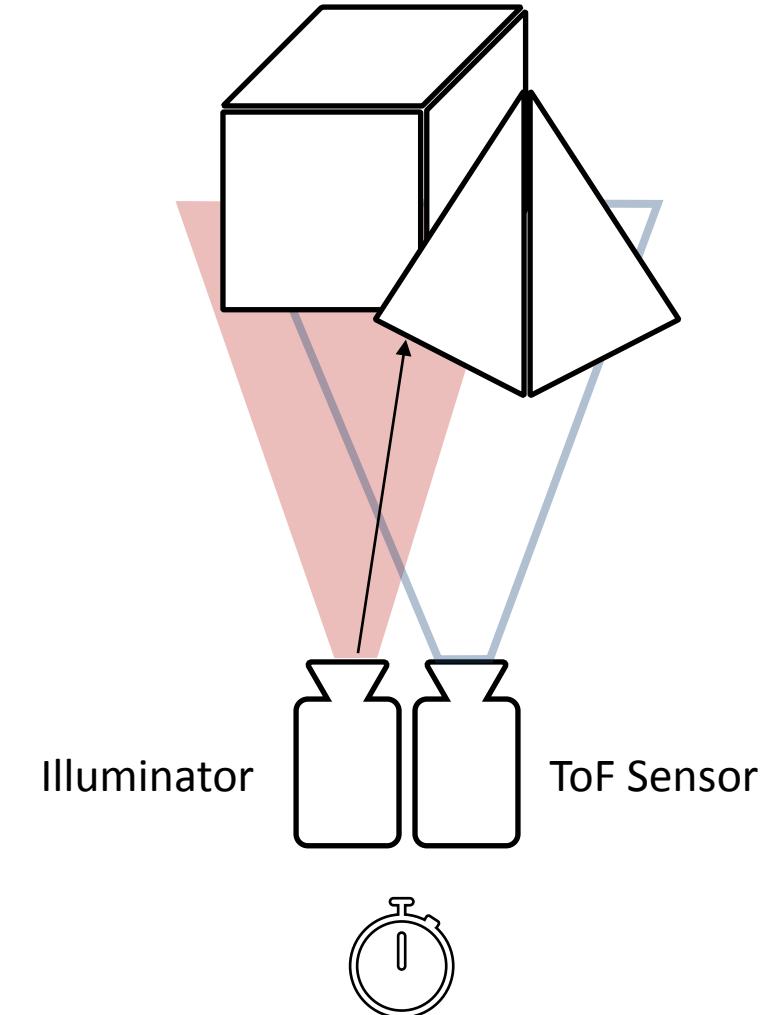
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t/2$
- Use low-power IR illumination for pulsed illumination
 - Can be switched on/off with <1ns
 - Generates a “light wall” that is reflected back by objects



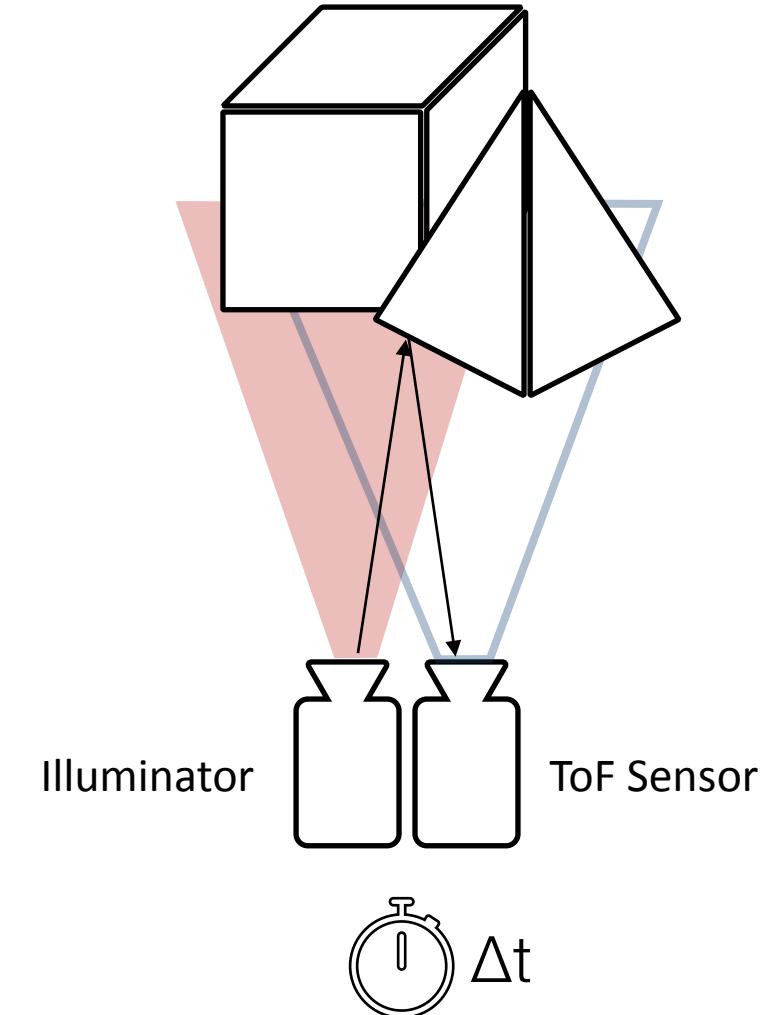
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t/2$
- Use low-power IR illumination for pulsed illumination
 - Can be switched on/off with <1ns
 - Generates a “light wall” that is reflected back by objects



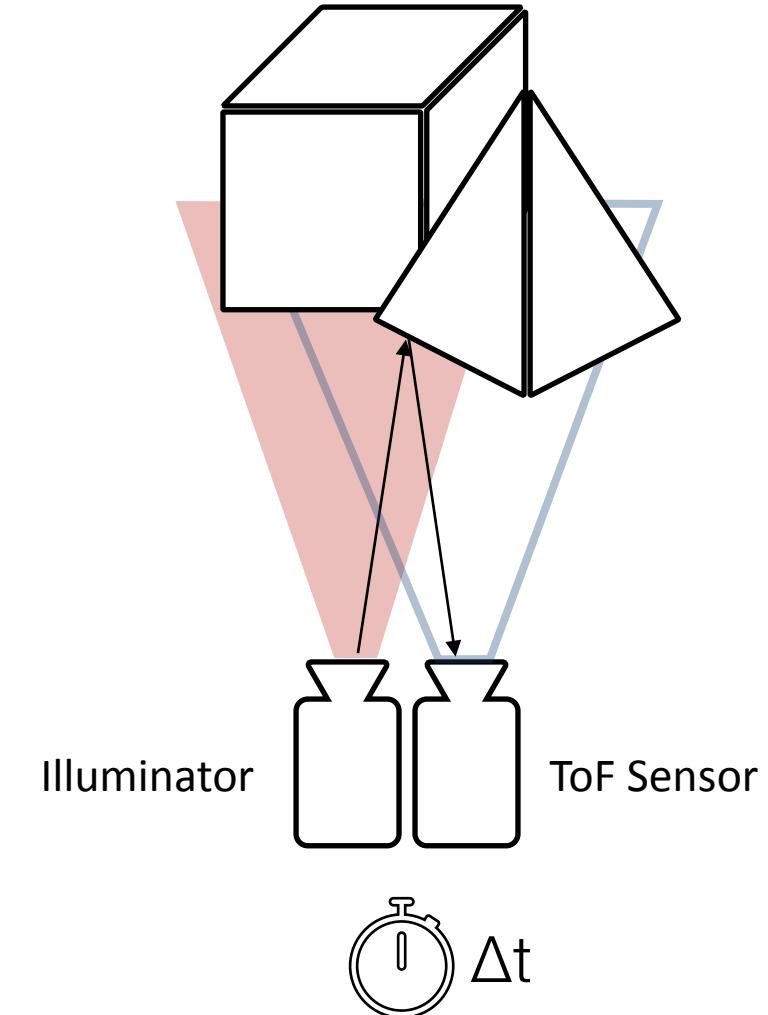
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t / 2$
- Use low-power IR illumination for pulsed illumination
 - Can be switched on/off with <1ns
 - Generates a “light wall” that is reflected back by objects



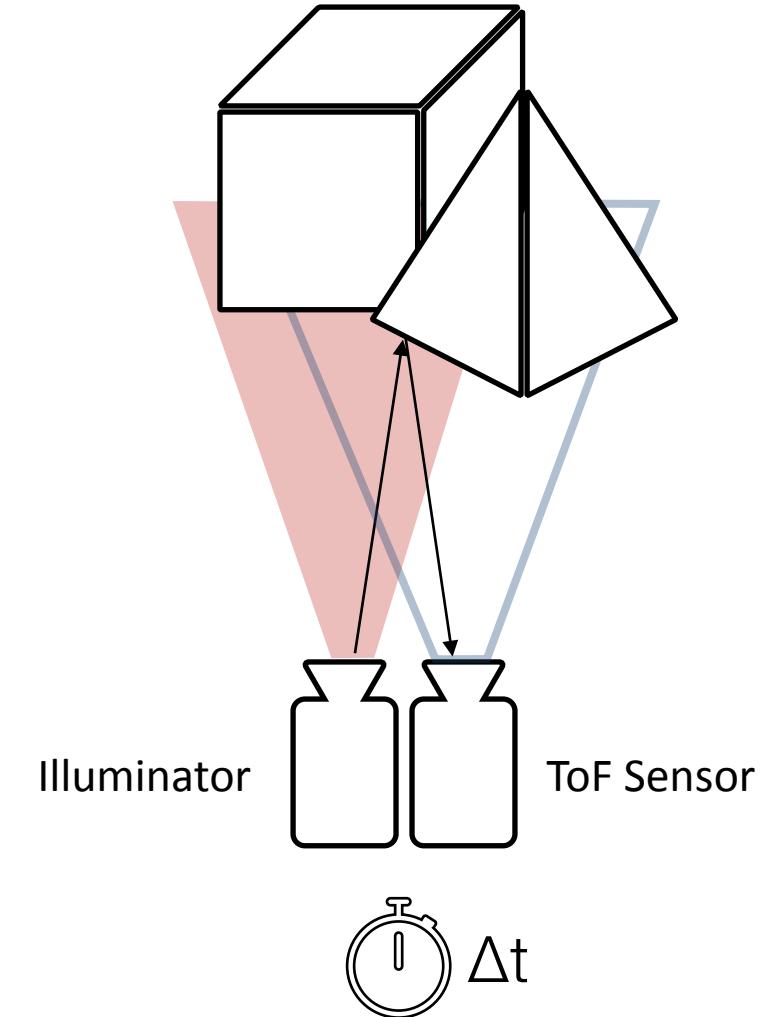
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t / 2$
 - $\Delta t = 2D/C$
 - $D=1\text{m} \Delta t = 6.6 \cdot 10^{-9} \text{ s}$
 - $D=0.01\text{m} \Delta t = 6.6 \cdot 10^{-11} \text{ s}$
 - $D=0.001\text{m} \Delta t = 6.6 \cdot 10^{-12} \text{ s}$ (pico seconds)



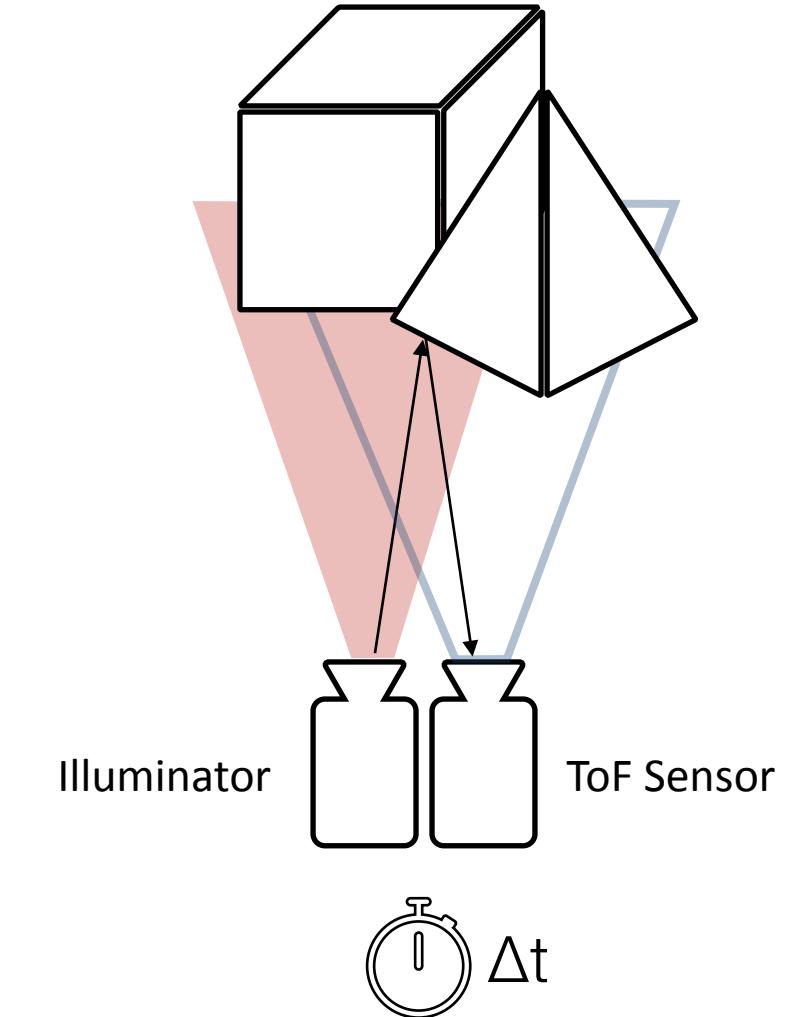
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t / 2$
 - $\Delta t = 2D/C$
 - $D=1\text{m} \Delta t = 6.6 \cdot 10^{-9} \text{ s}$
 - $D=0.01\text{m} \Delta t = 6.6 \cdot 10^{-11} \text{ s}$
 - $D=0.001\text{m} \Delta t = 6.6 \cdot 10^{-12} \text{ s}$ (pico seconds)
 - 3.3 picoseconds for light to travel 1 millimeter



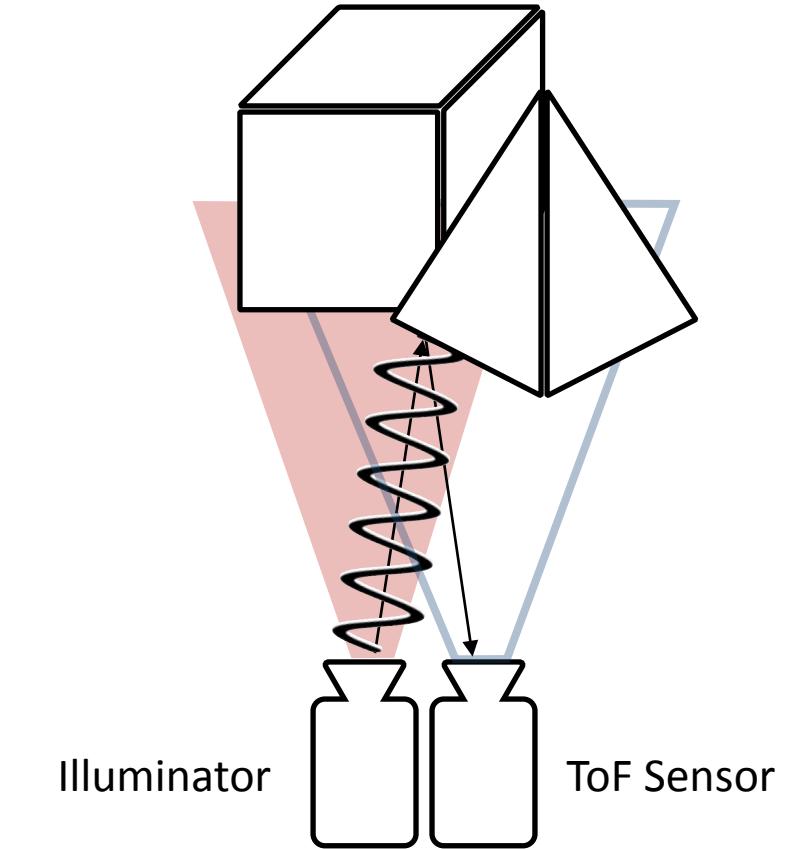
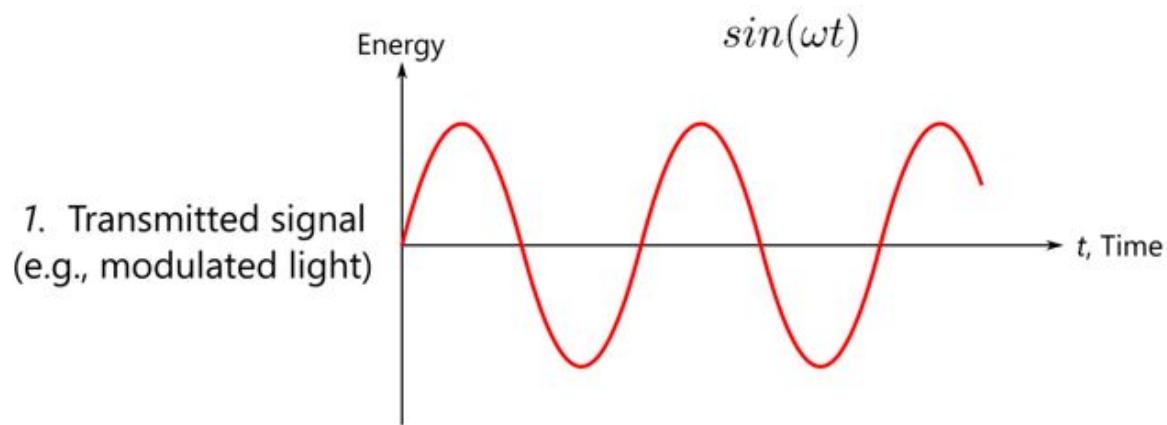
Optical Time-of-Flight (ToF)

- Practical realisation:
 - Direct Time of Flight
 - Measuring actual round-trip of light pulse
 - Technical limits to measuring time affect accuracy
 - Indirect Time of Flight
 - Indirectly measuring round-trip of light pulse
 - Usually utilising phase information



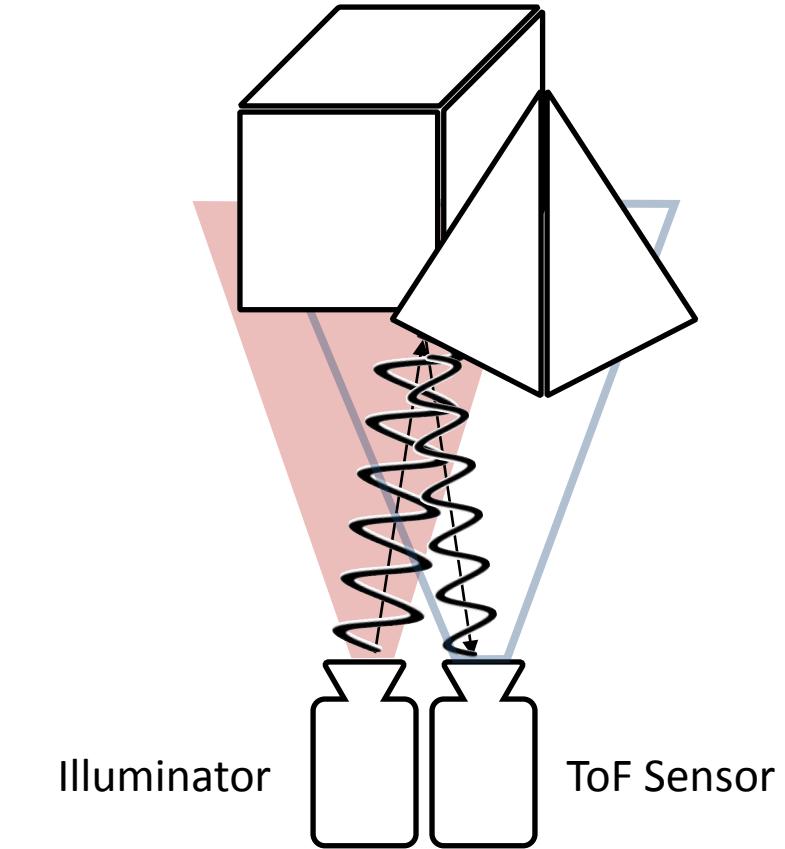
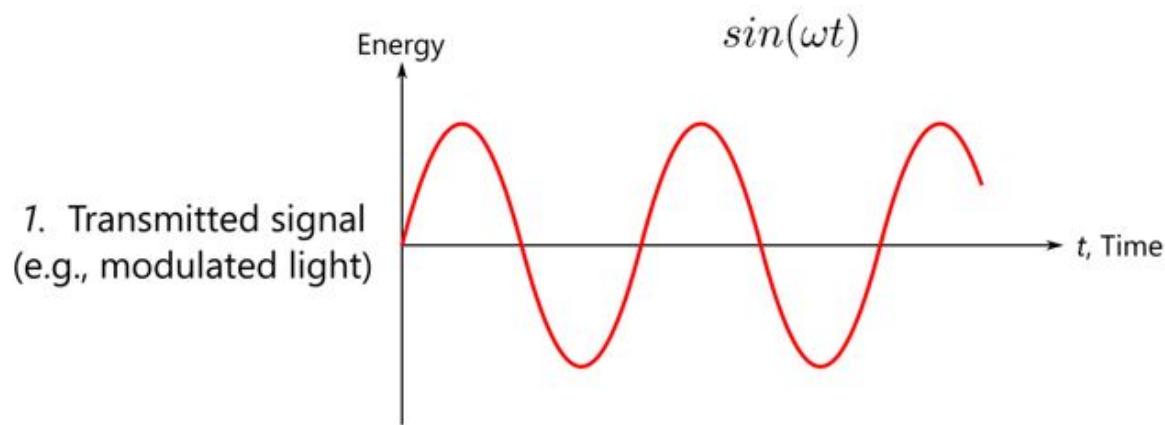
Indirect Time-of-Flight (ToF)

- We emit light with a periodic pattern, then calculate the phase shift in the return measurement.



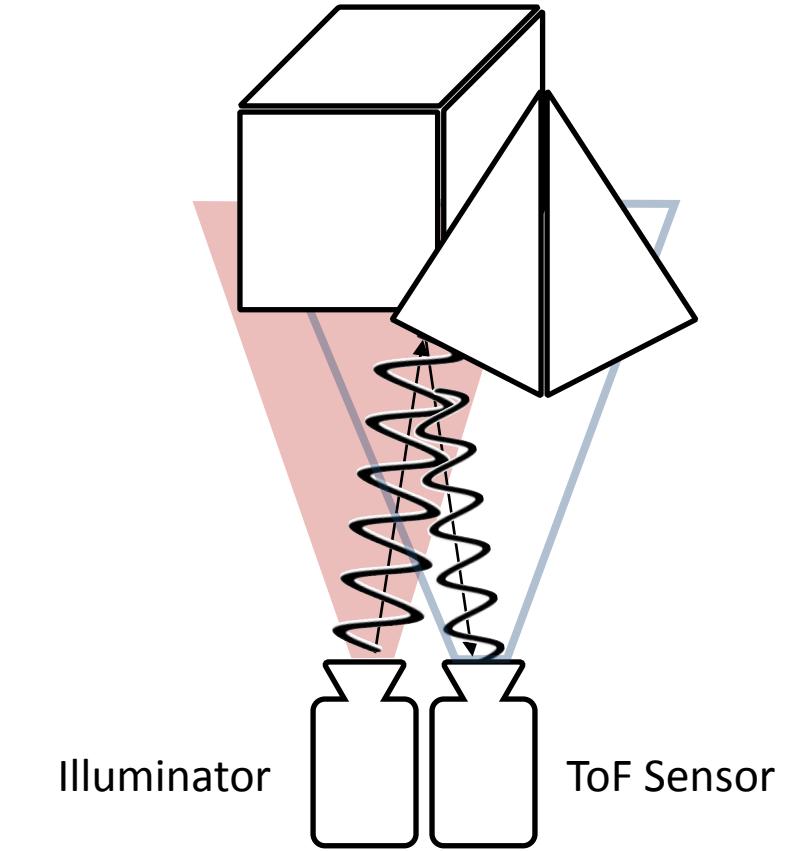
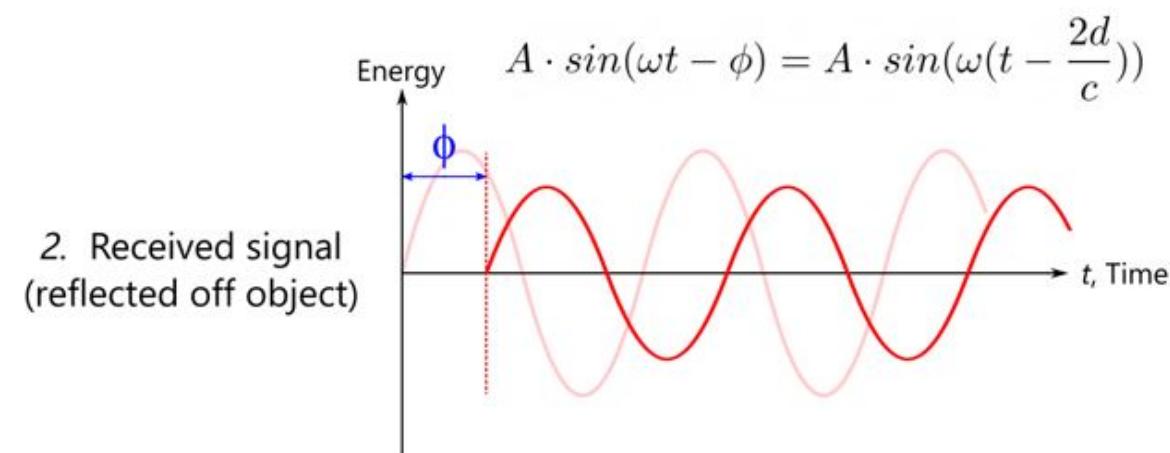
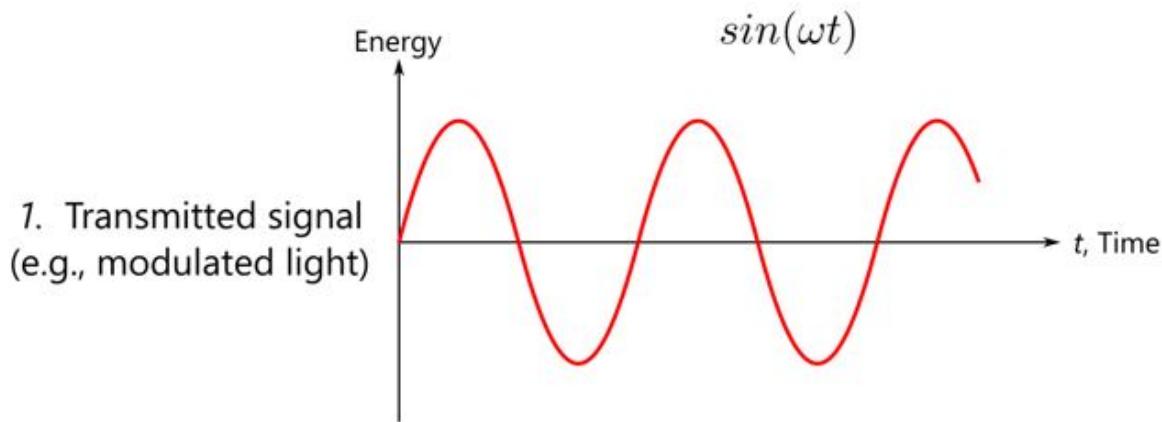
Indirect Time-of-Flight (ToF)

- We emit light with a periodic pattern, then calculate the phase shift in the return measurement.



Indirect Time-of-Flight (ToF)

- We emit light with a periodic pattern, then calculate the phase shift in the return measurement.



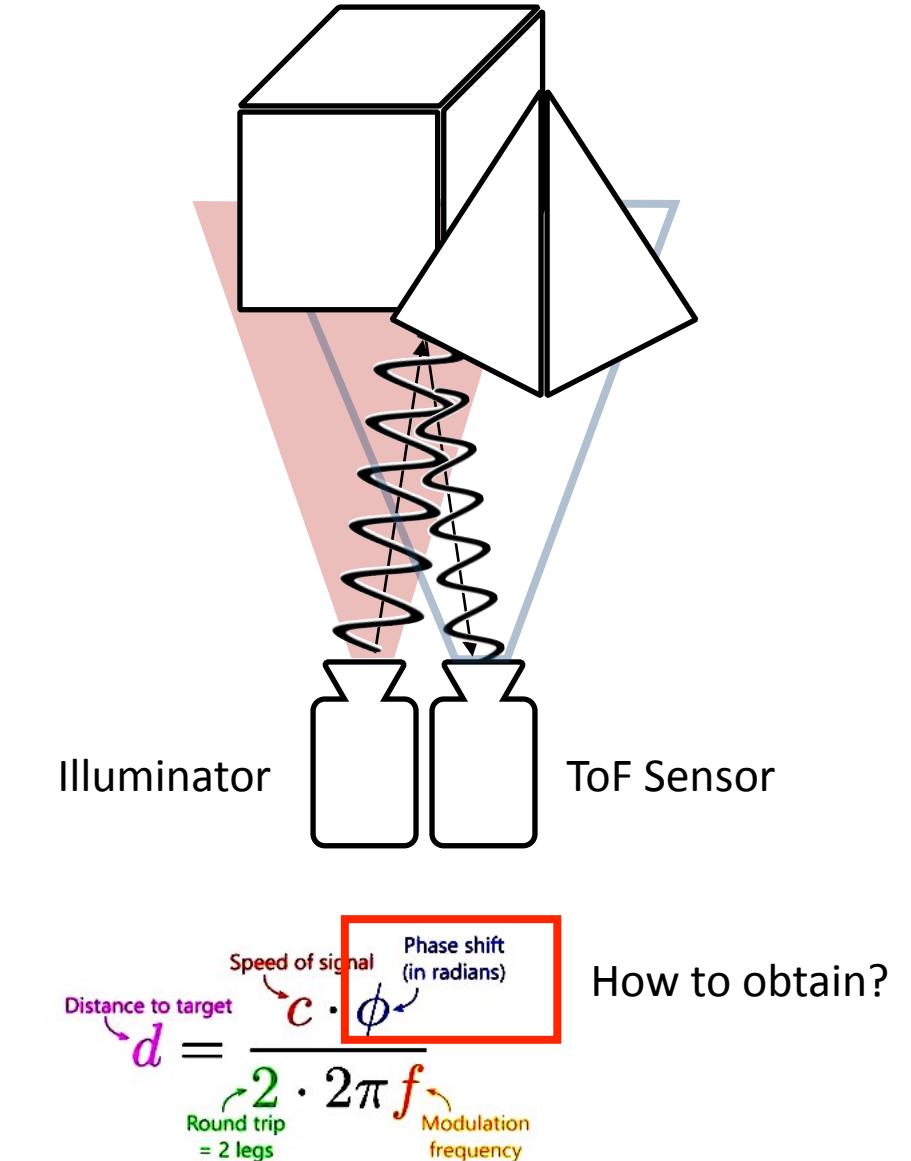
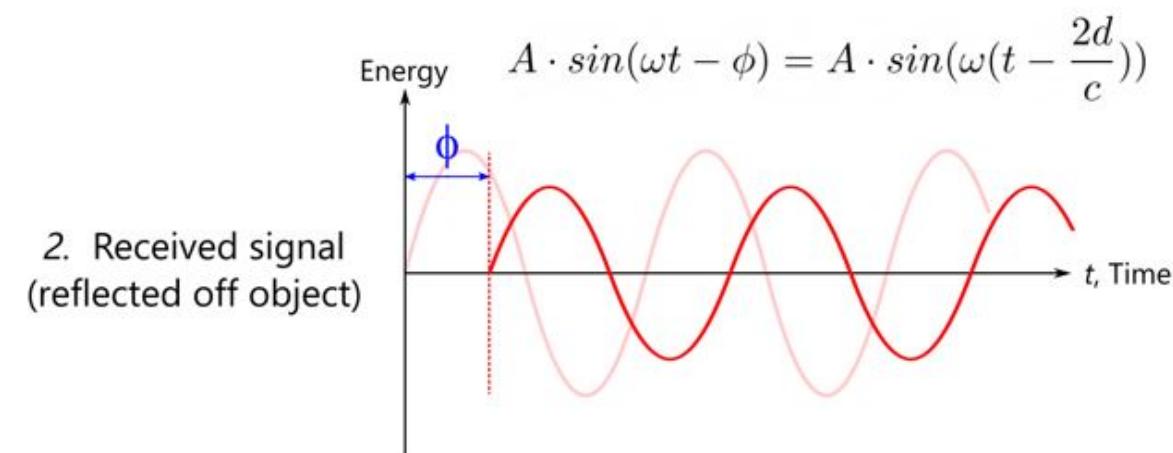
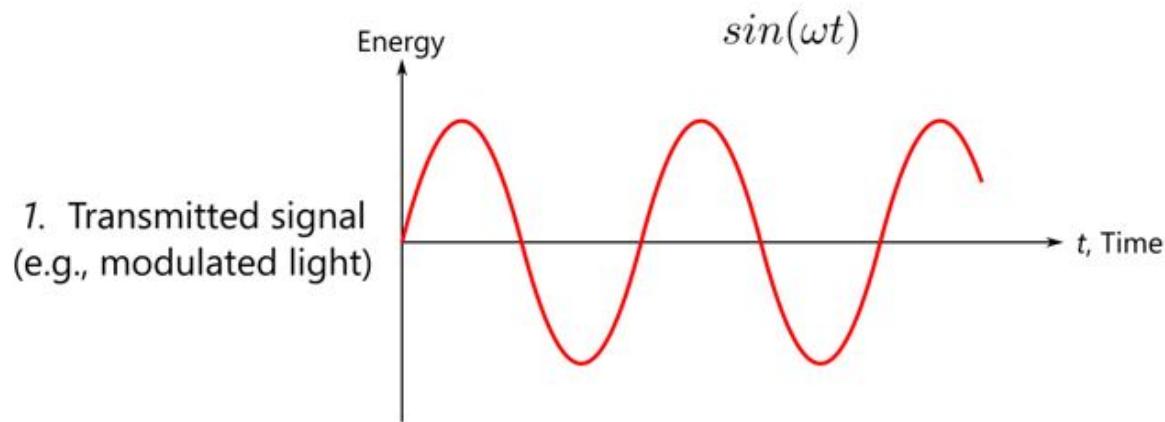
$$d = \frac{c \cdot \phi}{2 \cdot 2\pi f}$$

Annotations explain the variables:

- Distance to target: d
- Speed of signal: c
- Phase shift (in radians): ϕ
- Modulation frequency: f
- Round trip = 2 legs

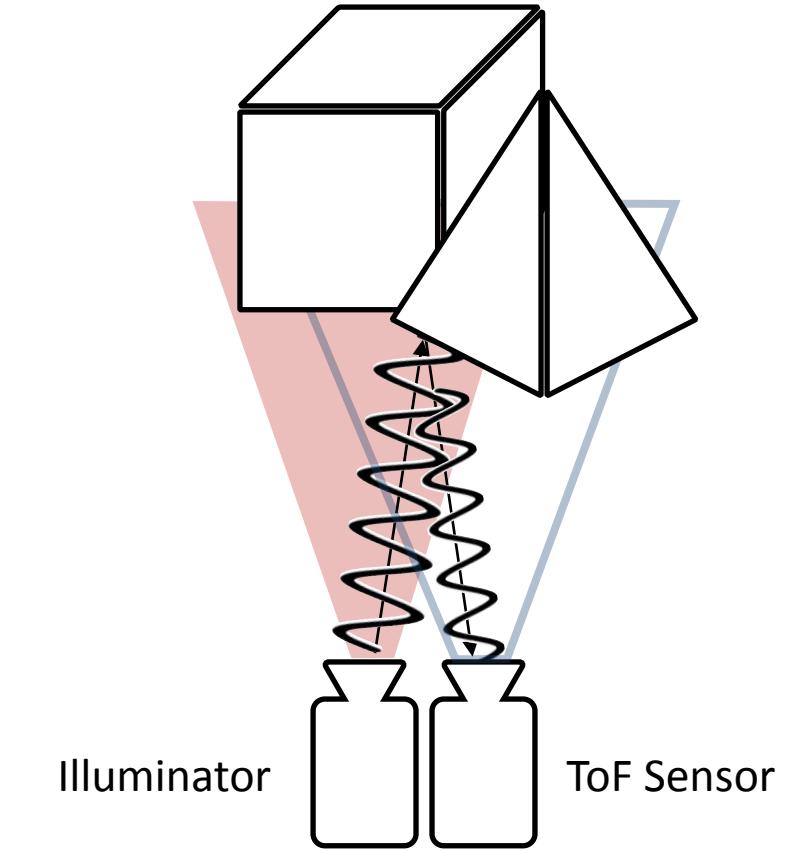
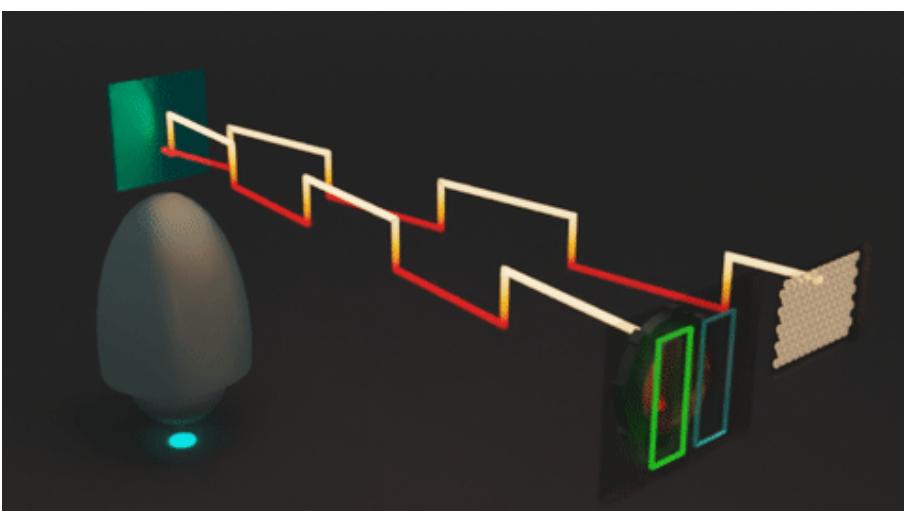
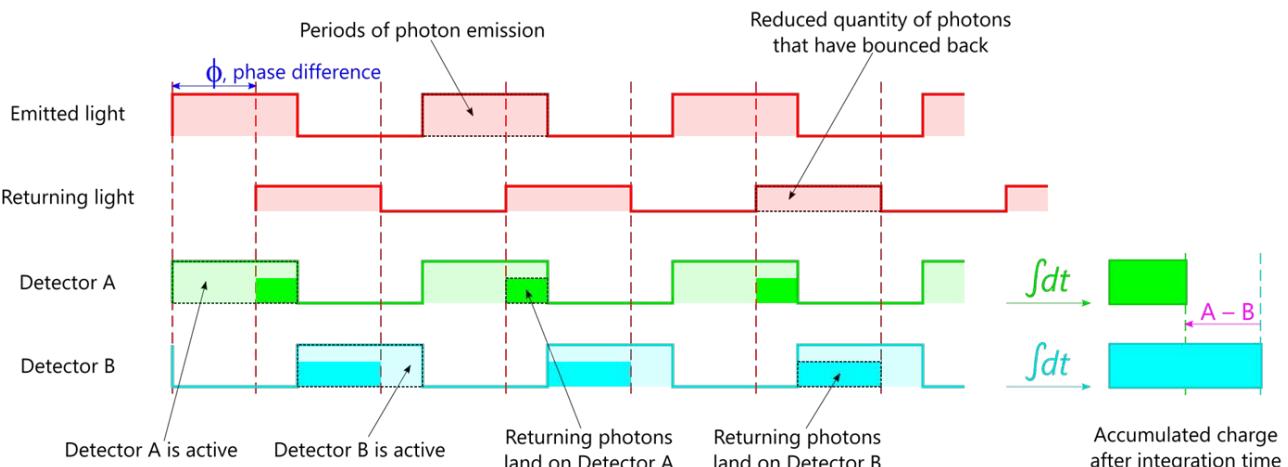
Indirect Time-of-Flight (ToF)

- We emit light with a periodic pattern, then calculate the phase shift in the return measurement.



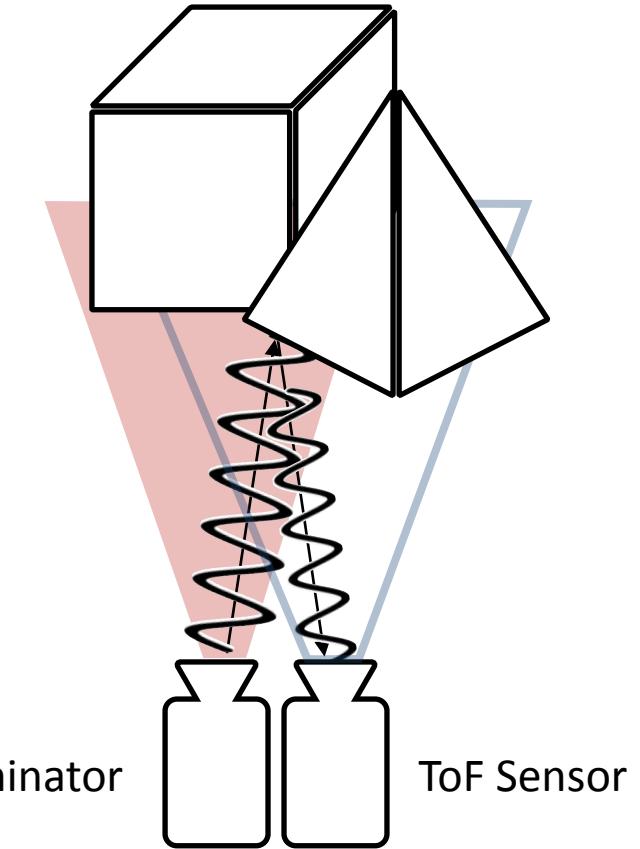
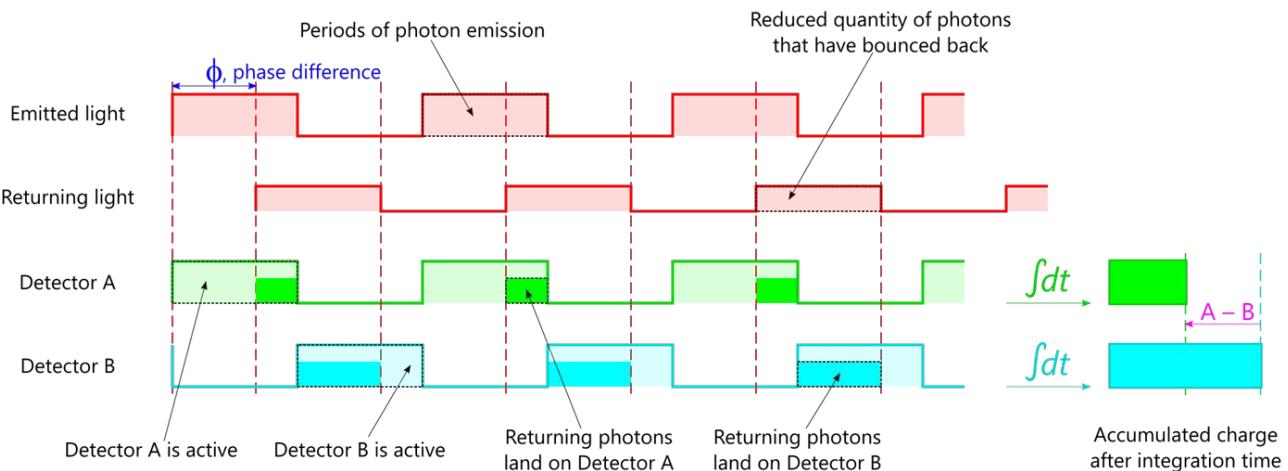
Indirect Time-of-Flight (ToF)

- Two detectors synchronised with the light emitter



Indirect Time-of-Flight (ToF)

- Two detectors synchronised with the light emitter



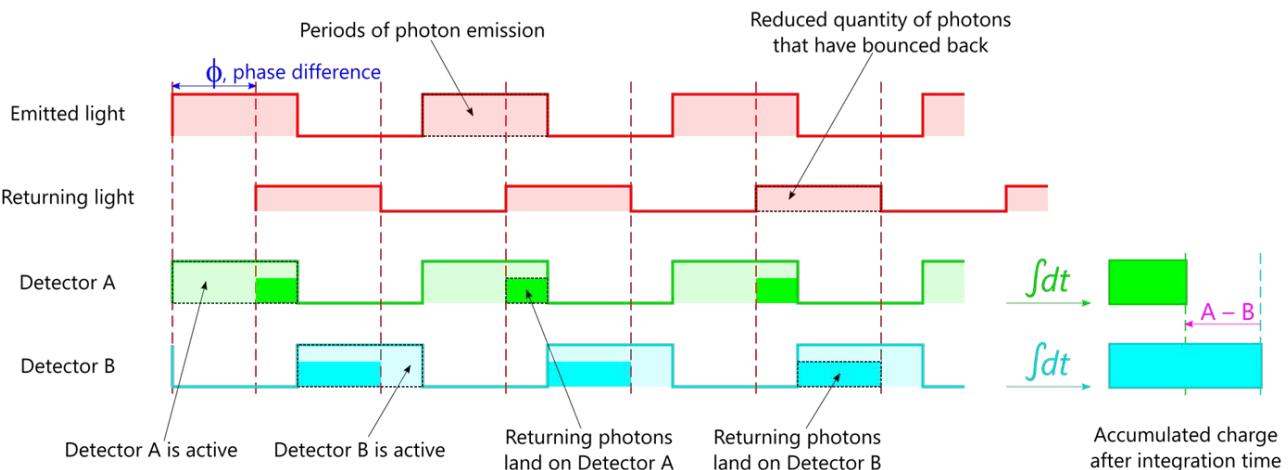
$$d = \frac{c \cdot \phi}{2 \cdot 2\pi f}$$

Annotations for the equation:

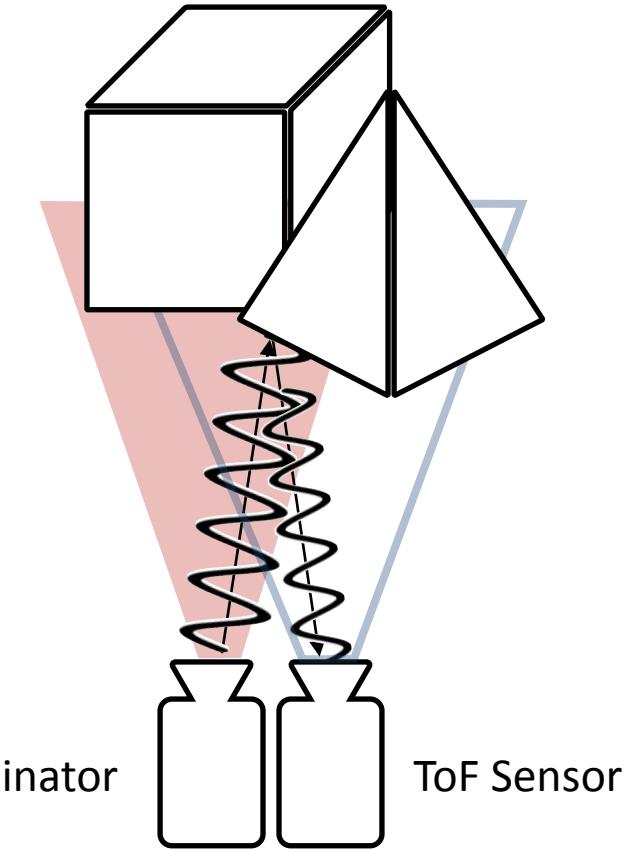
- Distance to target** (d) is highlighted in purple.
- Speed of signal** (c) is highlighted in red.
- Phase shift (in radians)** (ϕ) is highlighted in blue.
- Modulation frequency** (f) is highlighted in orange.
- Round trip = 2 legs** is highlighted in green.

Indirect Time-of-Flight (ToF)

- Two detectors synchronised with the light emitter



- What can go wrong?
 - Ambient light?
 - Ambiguity?

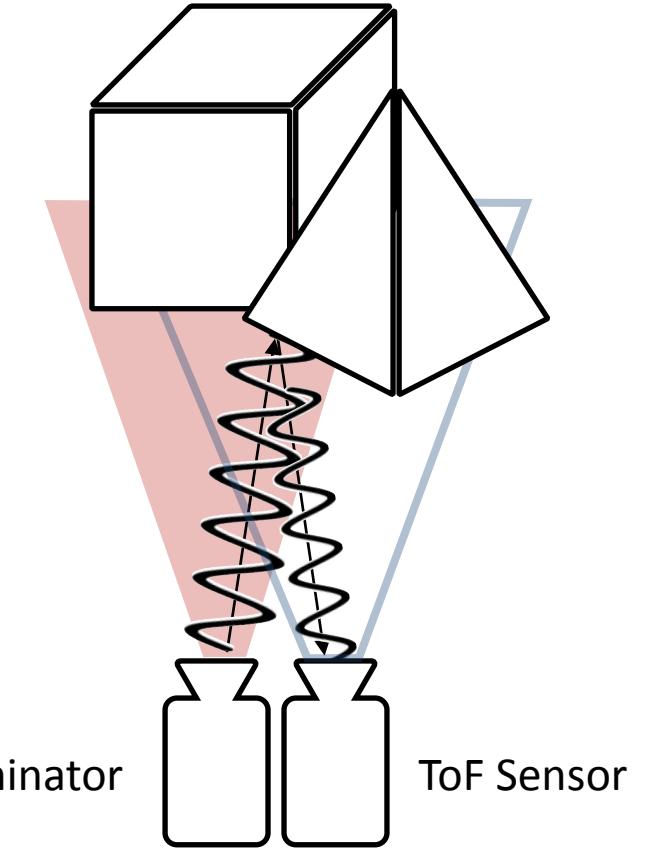
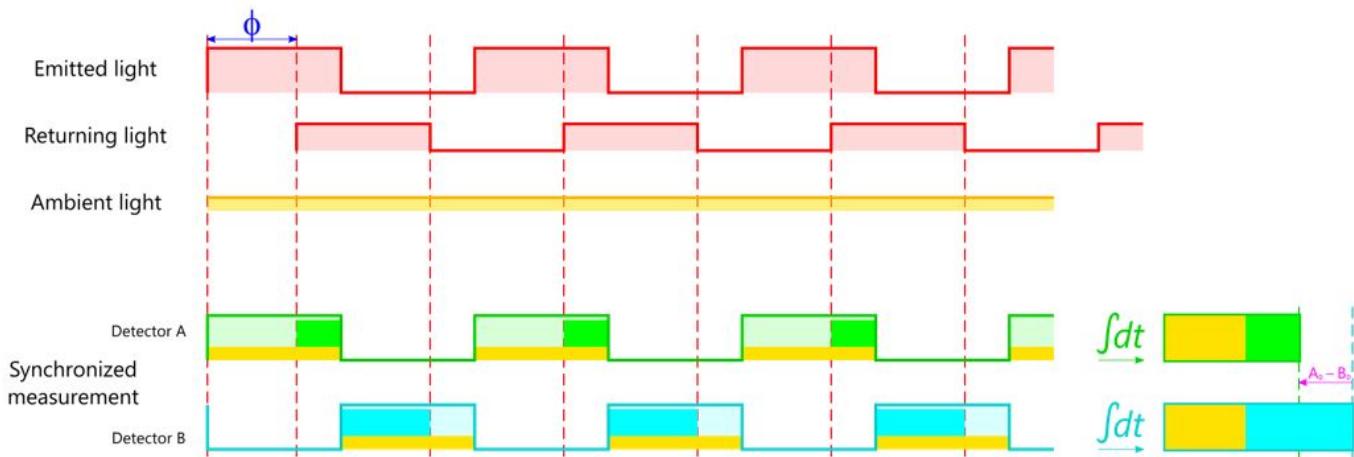


$$d = \frac{c \cdot \phi}{2 \cdot 2\pi f}$$

Annotations explain the variables:
Distance to target: d
Speed of signal: c
Phase shift (in radians): ϕ
Modulation frequency: f
Round trip = 2 legs

Indirect Time-of-Flight (ToF)

- Ambient light:
 - Usually not a problem



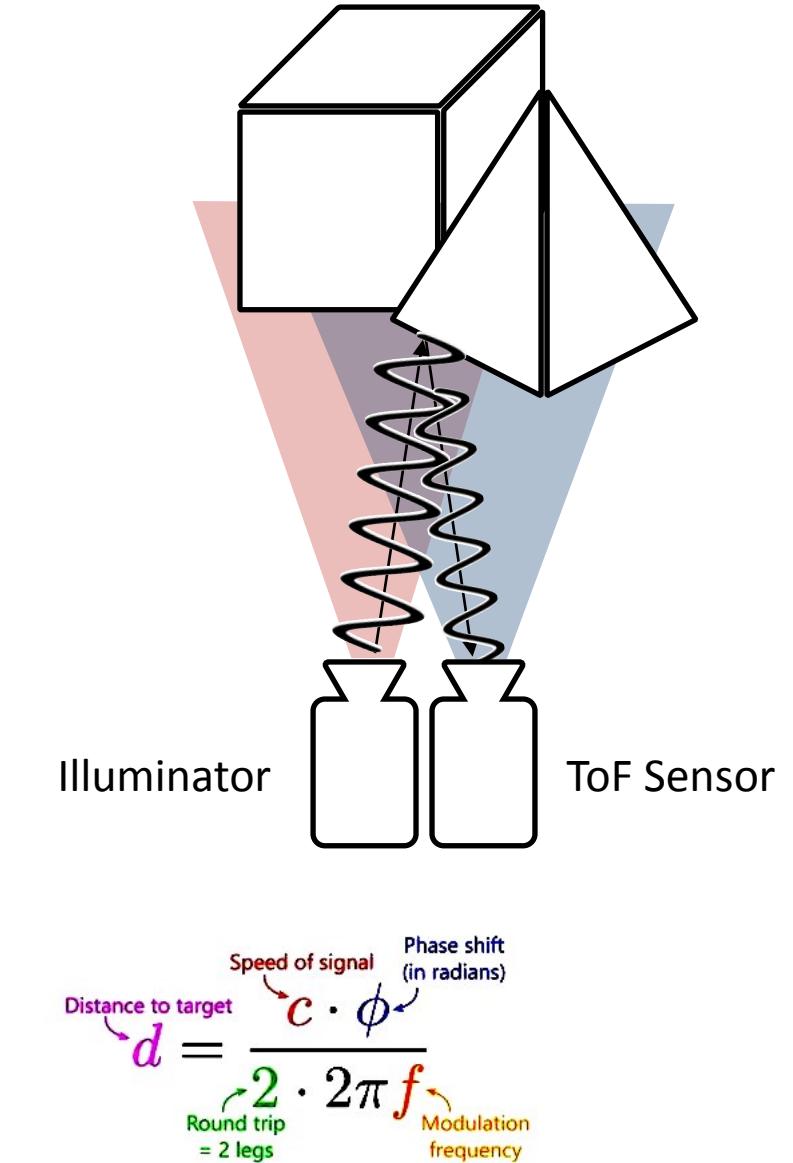
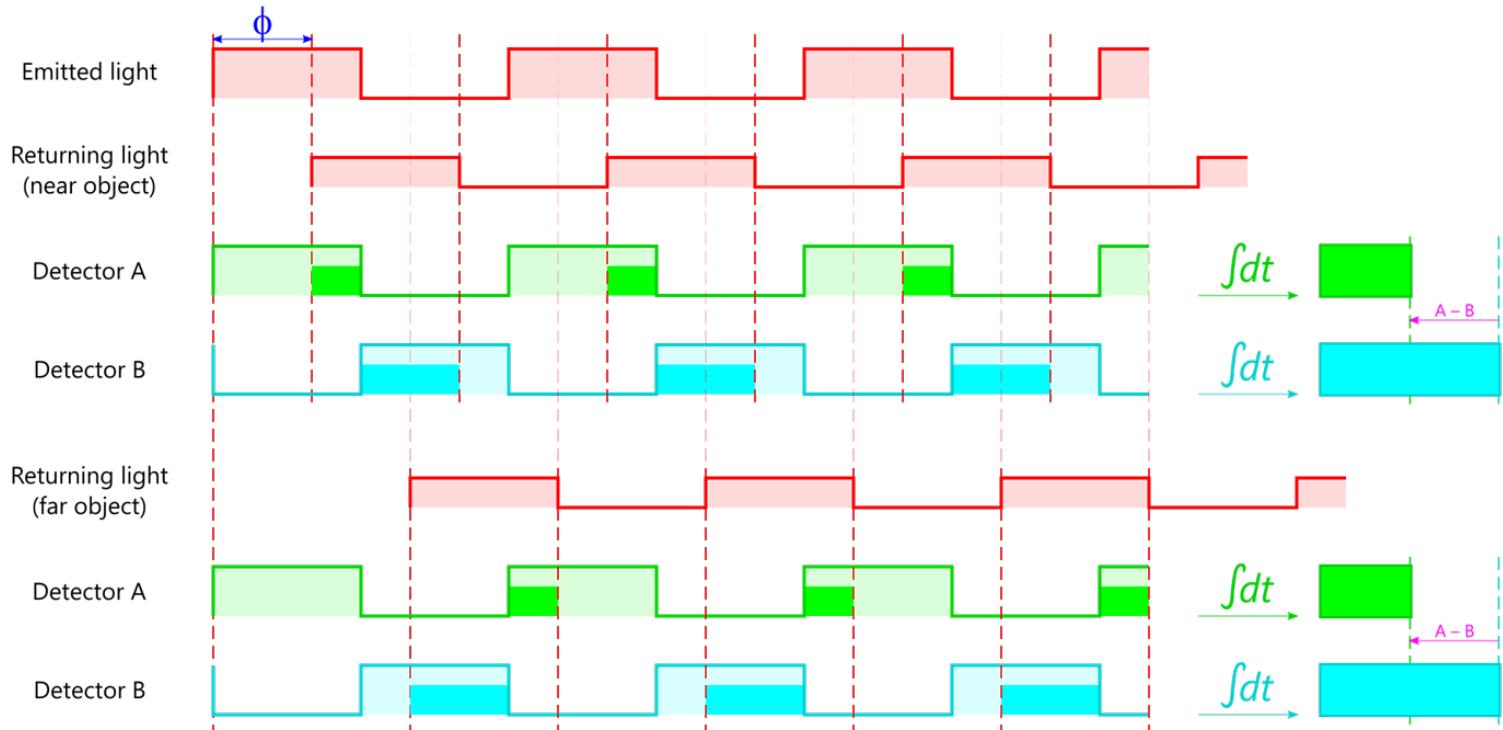
$$d = \frac{c \cdot \phi}{2 \cdot 2\pi f}$$

Annotations explain the variables:

- Distance to target: d
- Speed of signal: c
- Phase shift (in radians): ϕ
- Modulation frequency: f
- Round trip = 2 legs

Indirect Time-of-Flight (ToF)

- Ambiguity
 - Usually a problem (e.g. object with another distance creating same phase measurement)



Active range scanner / depth cameras work in infrared spectrum

We cannot perceive the infrared spectrum...



Active range scanner / depth cameras work in infrared spectrum

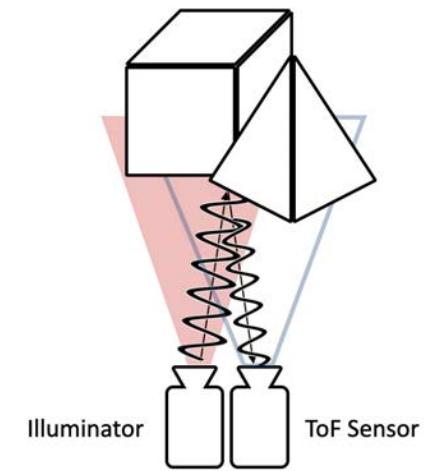
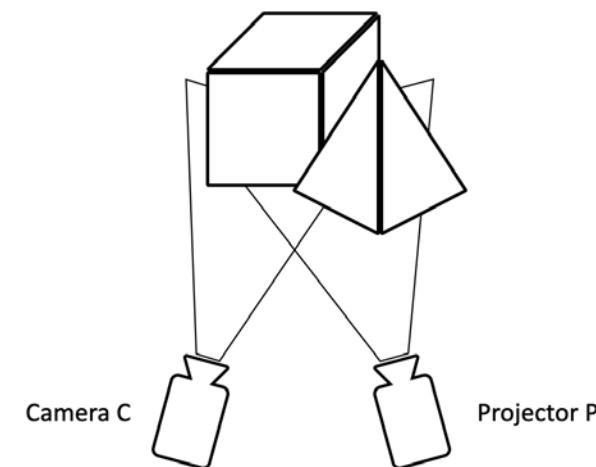
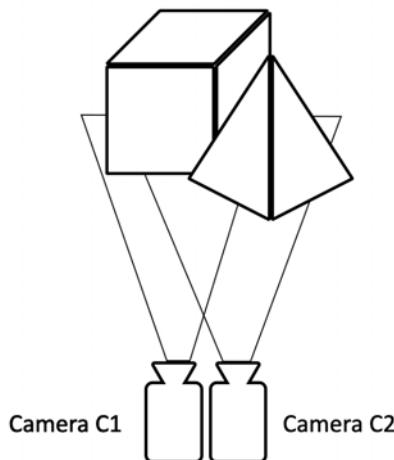
We cannot perceive the infrared spectrum...



But they can...

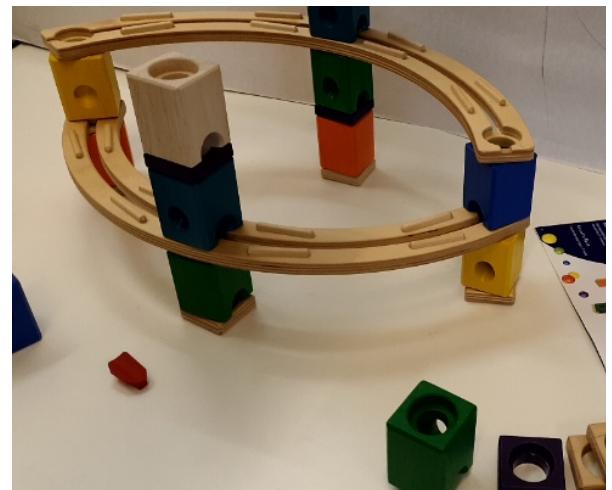
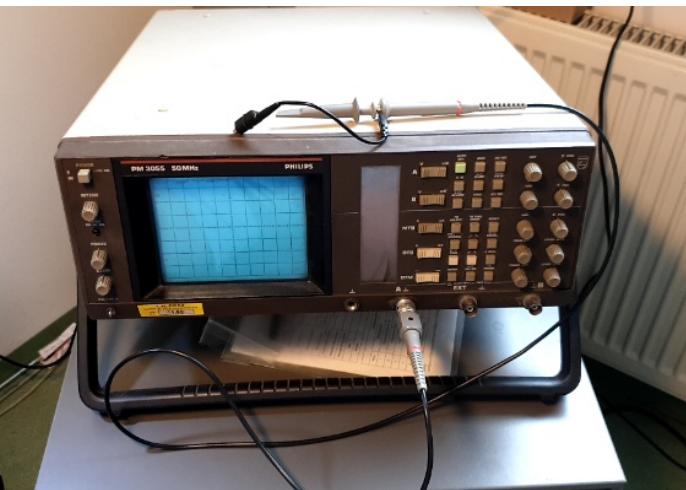
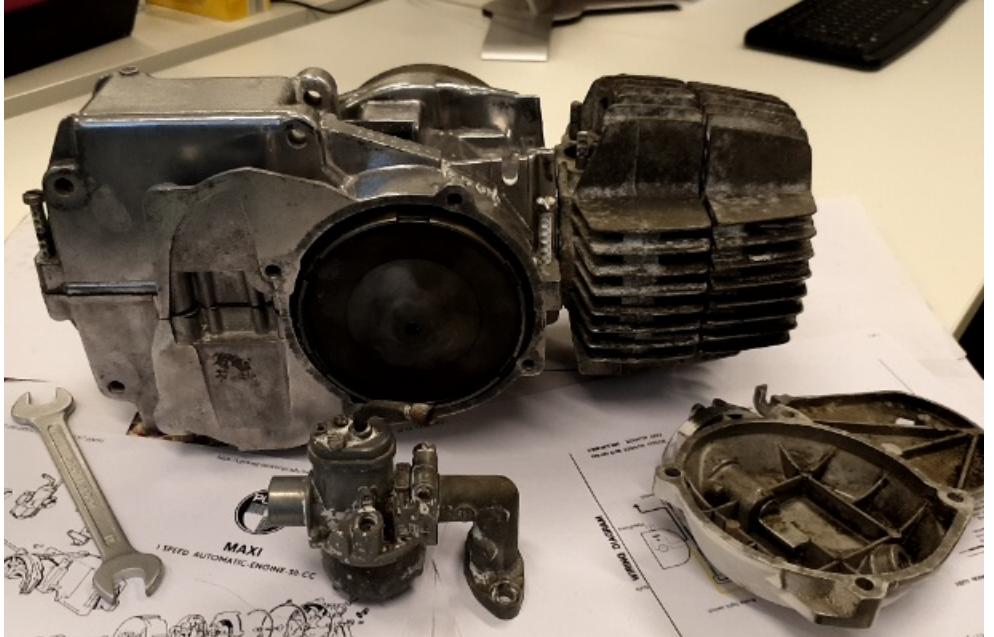
Summary Range Scanning

Considerations	Stereo Vision (Passive)	Structured Light (Active)	Time of Flight (Active)
Response time / Speed			
Depth Accuracy			
Low light performance			
Bright light performance			



Summary Range Scanning

Any problems when scanning these objects?



No Depth Cam?

Depth from Mono / Depth Estimation using ML

- Convolutional neural networks (CNNs) have been applied to monocular depth estimation
- Common architectures include U-Net, ResNet, and DenseNet
- During training:
 - Network is given pairs of images and their corresponding ground-truth depth maps
 - Minimize the difference between predicted depth and ground truth depth



<https://github.com/yuhsuanyeh/BiFuse>

Depth from Mono: Depth Anything V2



<https://huggingface.co/spaces/depth-anything/Depth-Anything-V2>

Depth from Mono: Depth Anything V2

Spaces | depth-anything / Depth-Anything-V2 | like 542 | Running on ZERO | App | Files | Community 14 | [Edit](#)

Depth Anything V2

Official demo for Depth Anything V2. Please refer to our [paper](#), [project page](#), and [github](#) for more details.

Depth Prediction demo

Input Image



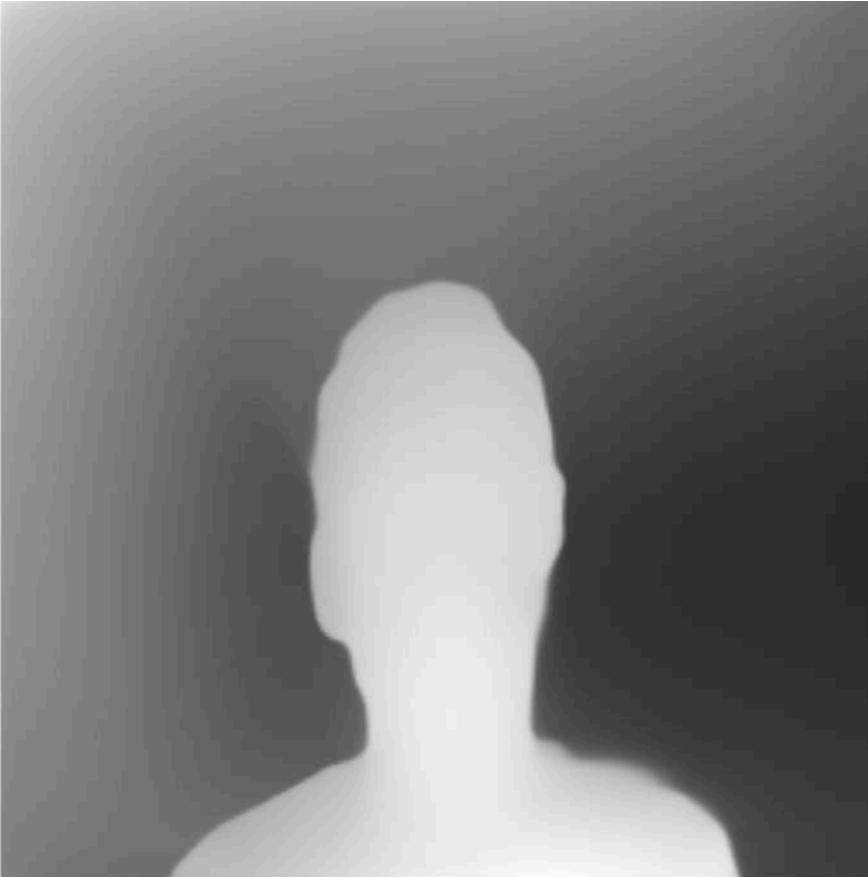
Depth Map with Slider View



Compute Depth

<https://huggingface.co/spaces/depth-anything/Depth-Anything-V2>

Depth from Mono: Tensorflow JS

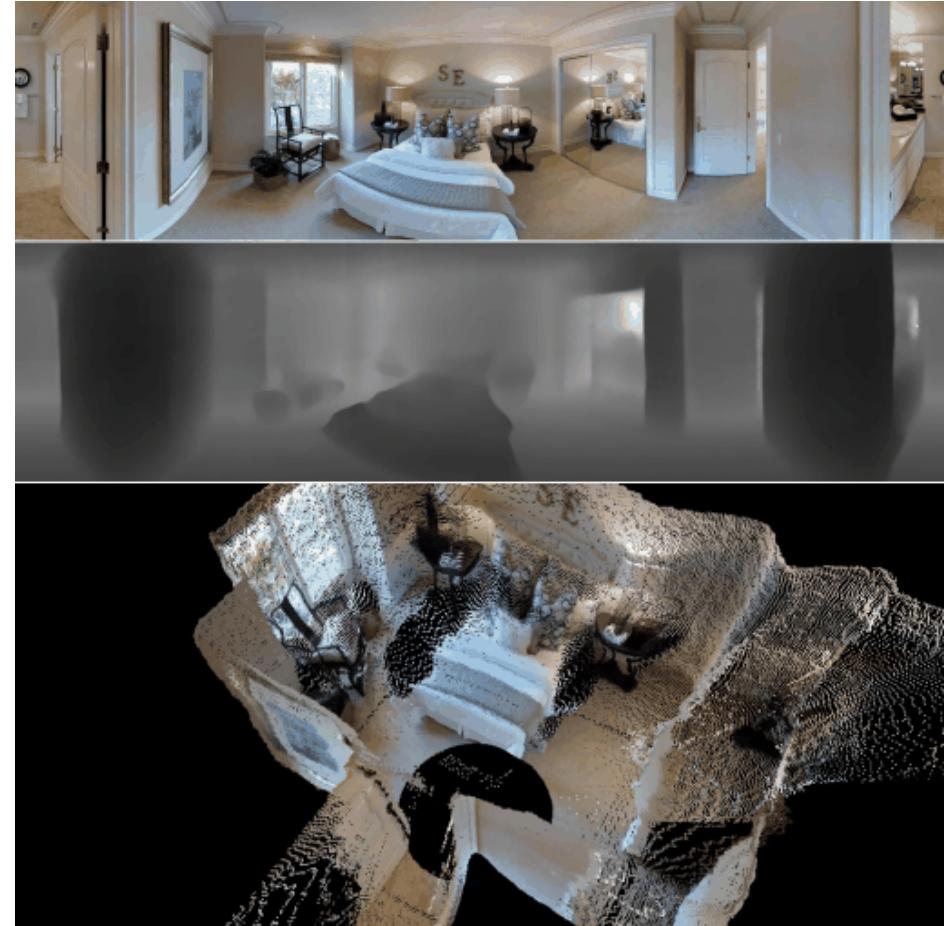


Average Depth: 126.56

<https://kuromadara.github.io/live-depth-prediction-using-tfjs/>

Reminder: Depth from Mono / Depth Estimation using ML

- Convolutional neural networks (CNNs) have been applied to monocular depth estimation
- Common architectures include U-Net, ResNet, and DenseNet
- During training:
 - Network is given pairs of images and their corresponding ground-truth depth maps
 - Minimize the difference between predicted depth and ground truth depth

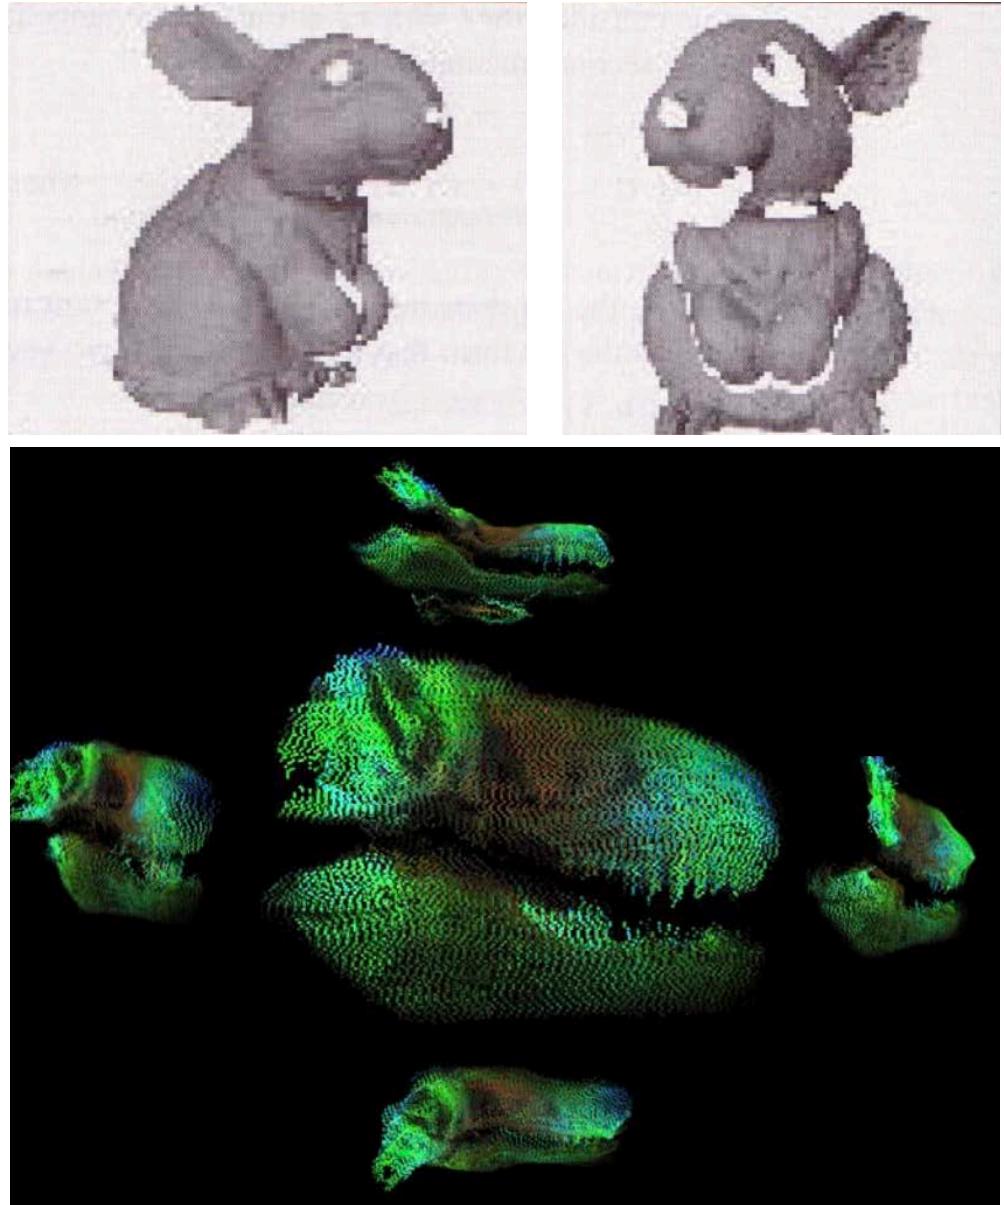


<https://github.com/yuhuanyeh/BiFuse>

Registration of Depth Maps

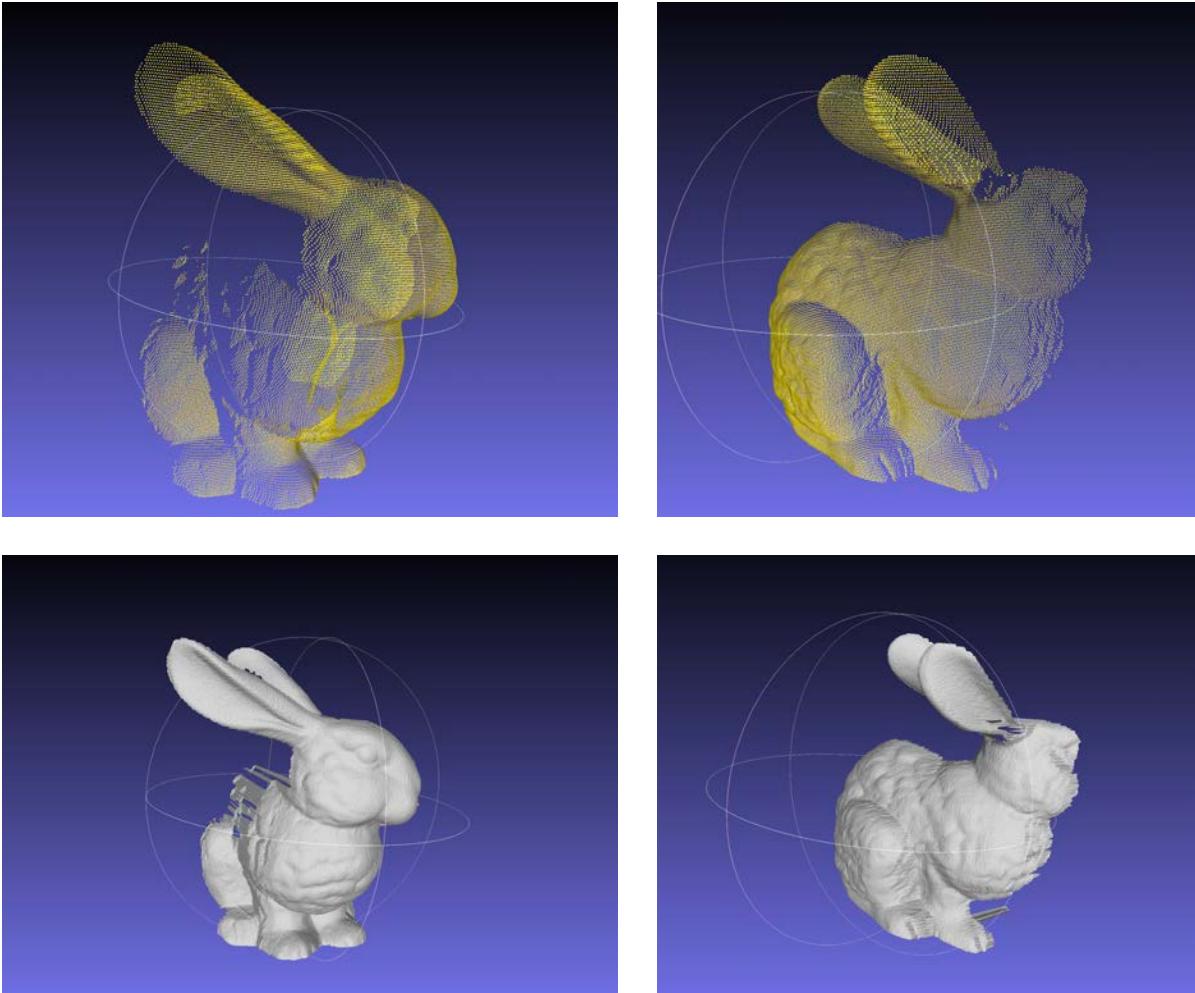
Registration of Depth Maps

- If correspondences are known, the depth (distance from camera / projector to scene point) can be estimated for every pixel (depth map)
- Having multiple depth maps of the same scene, how can they be merged?
- This process is known as registration



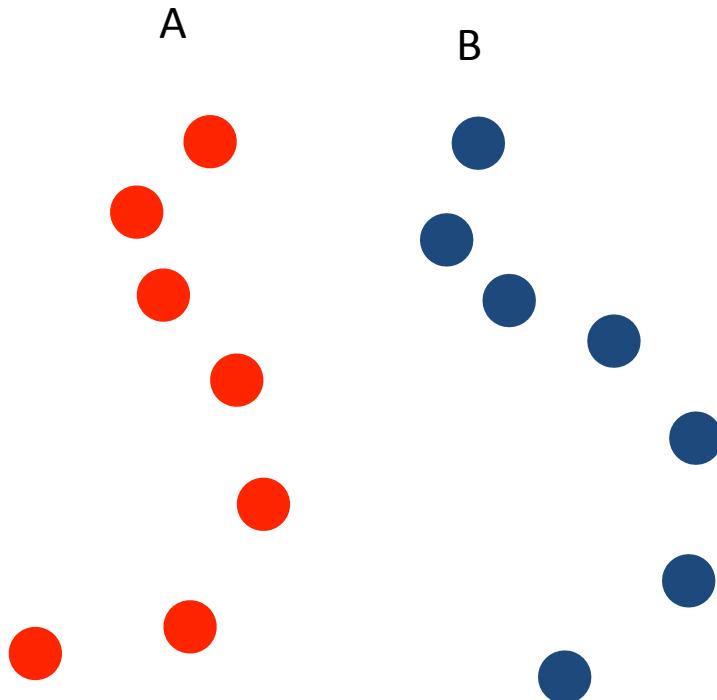
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



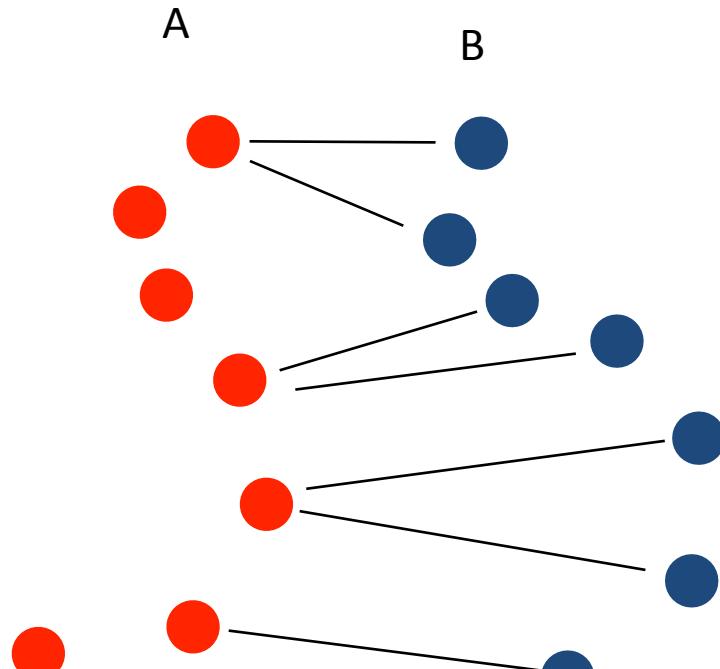
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



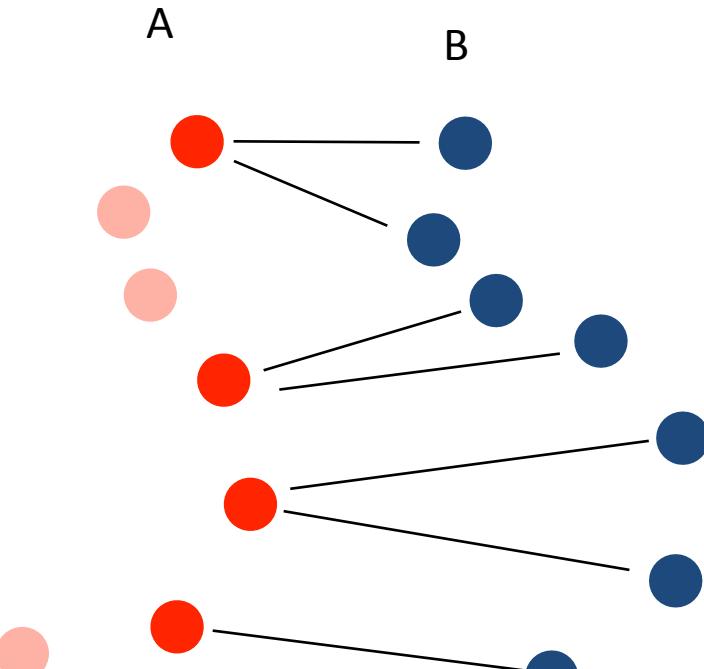
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



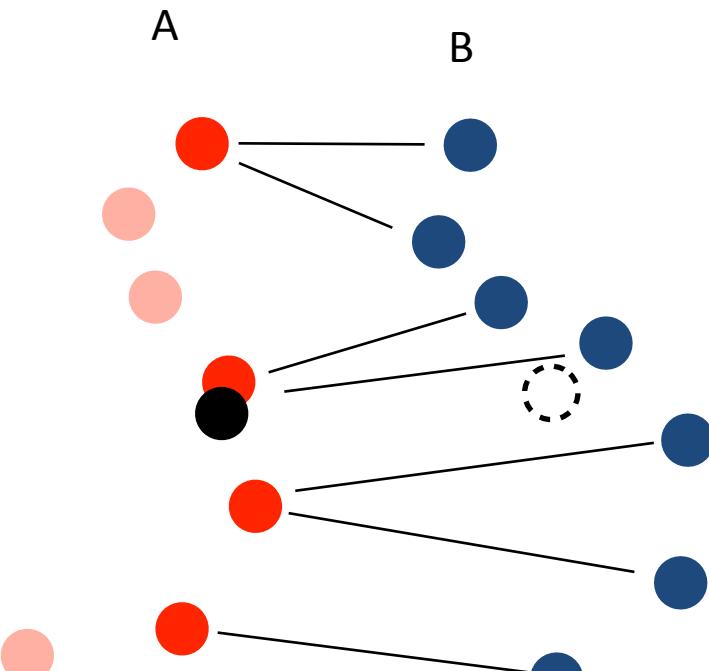
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



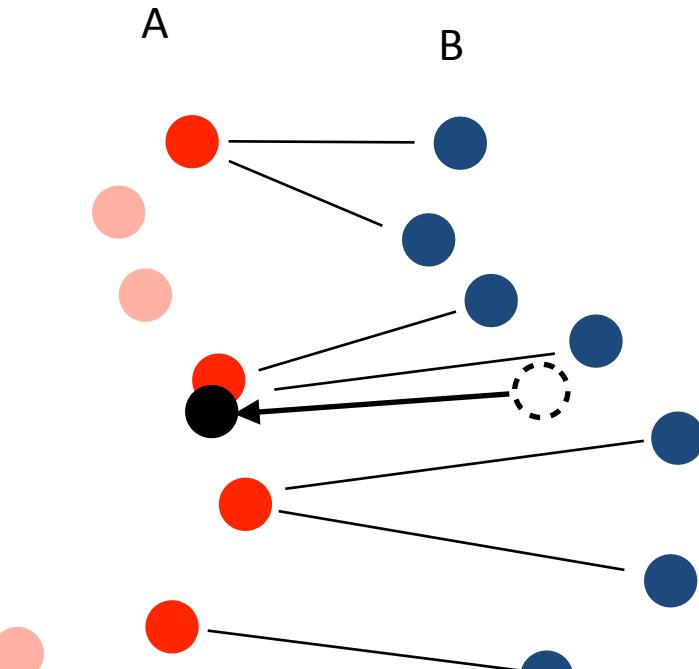
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



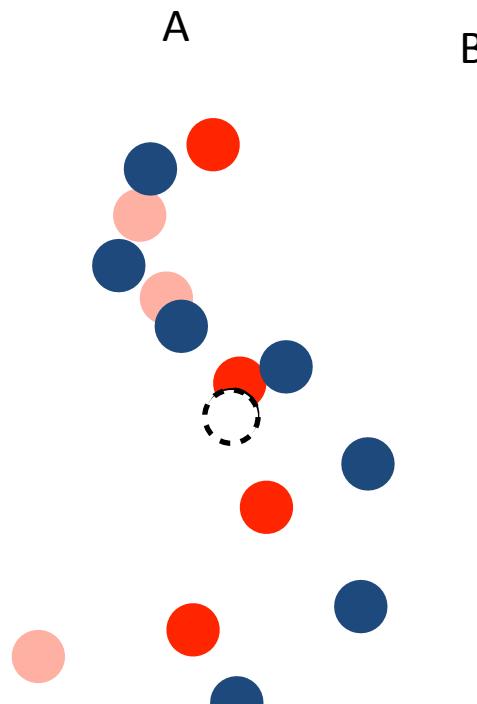
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



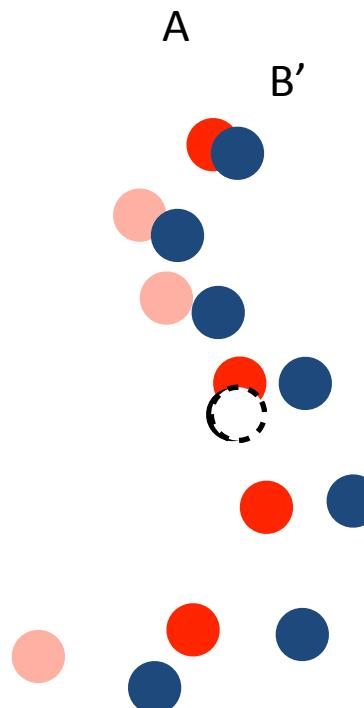
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



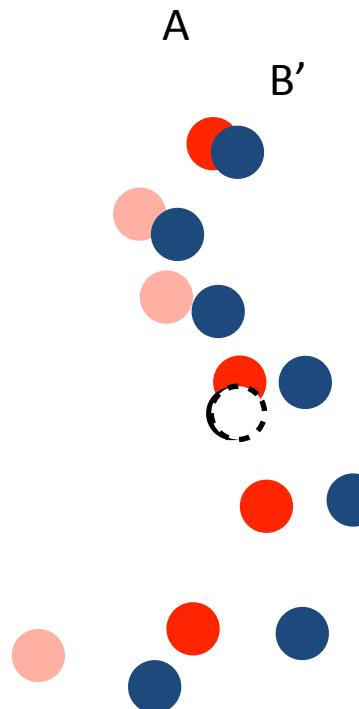
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



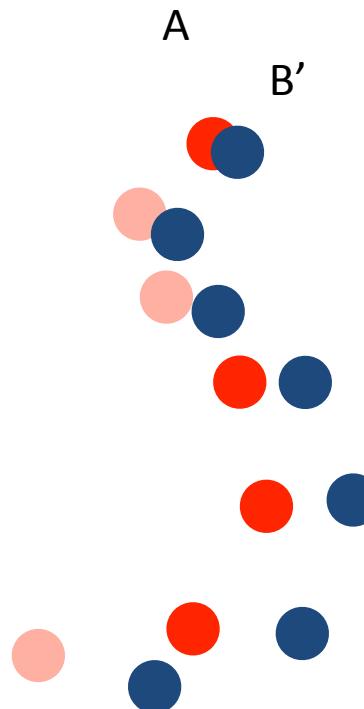
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



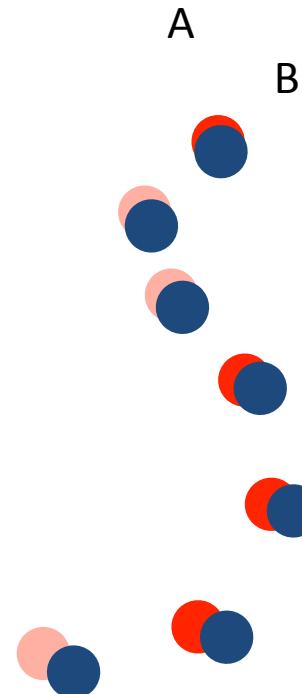
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



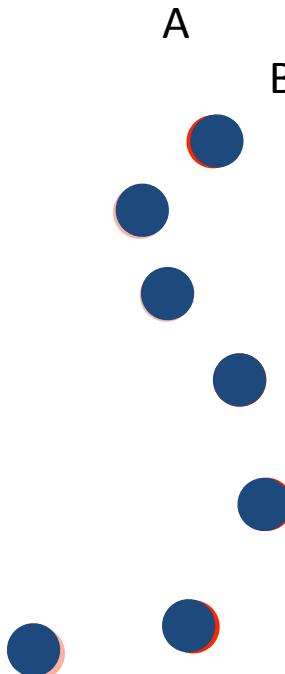
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



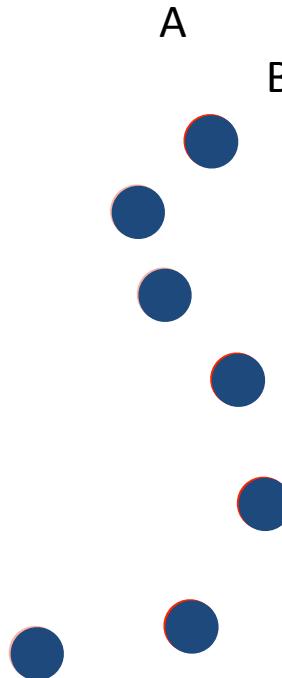
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



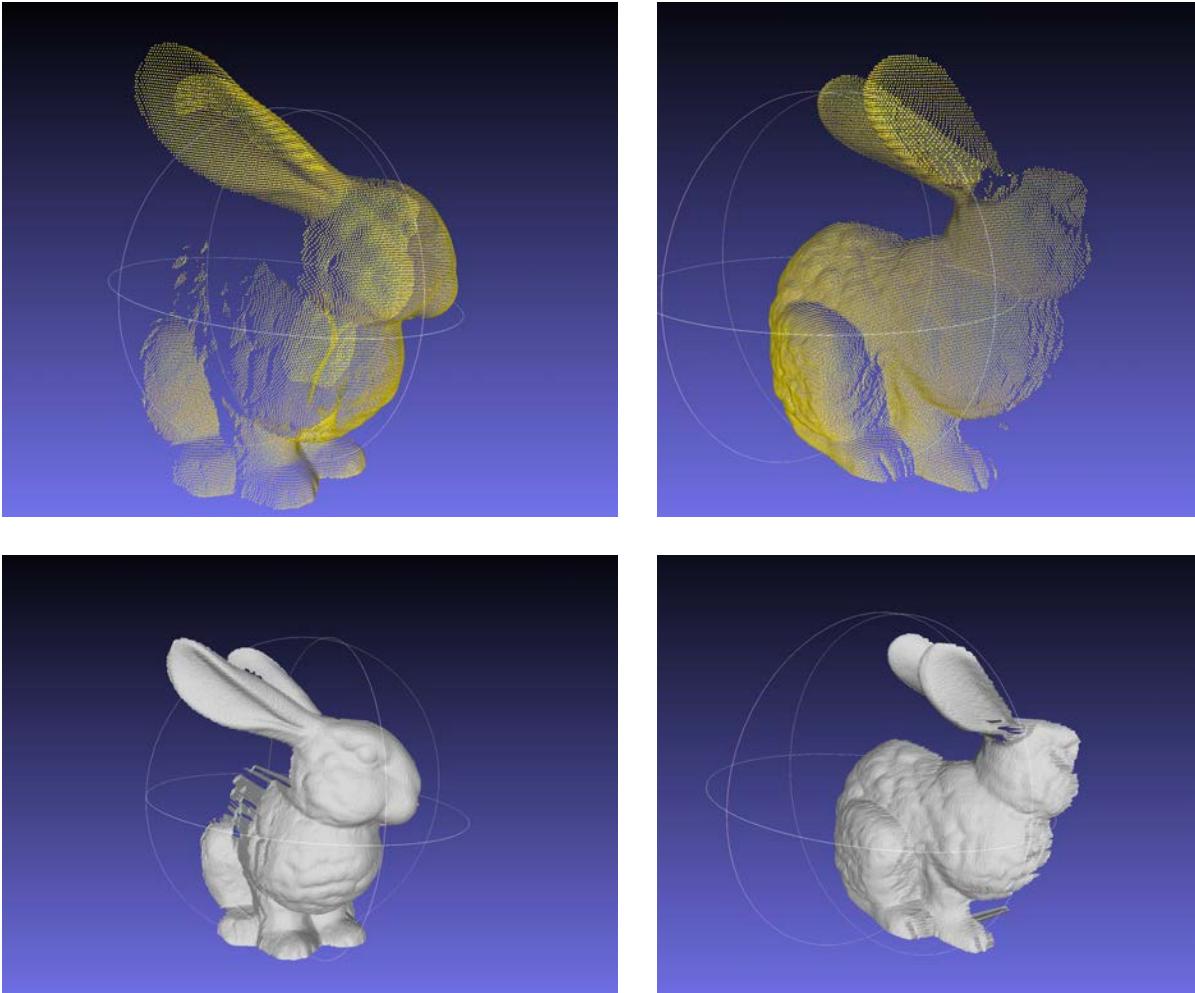
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



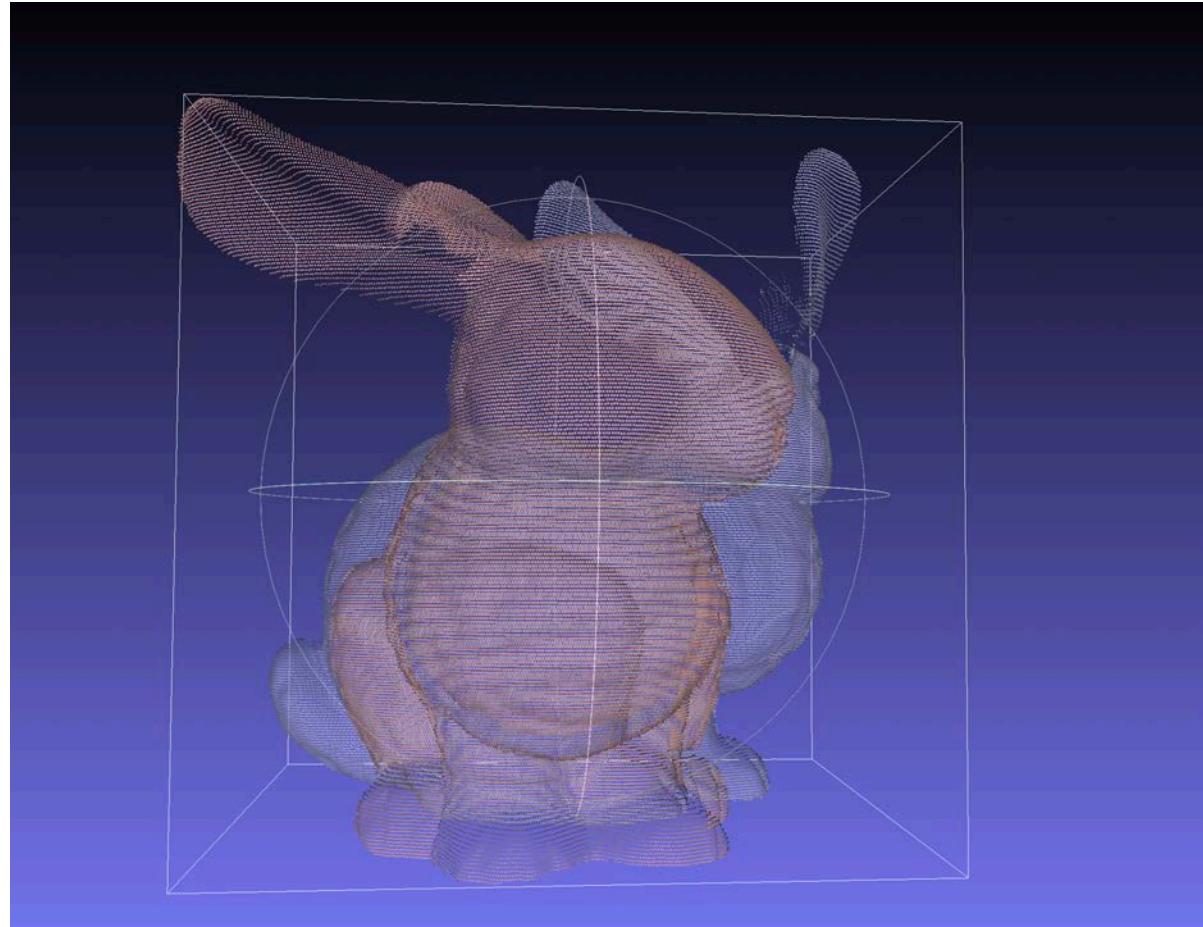
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



Iterative Closest-Point Method

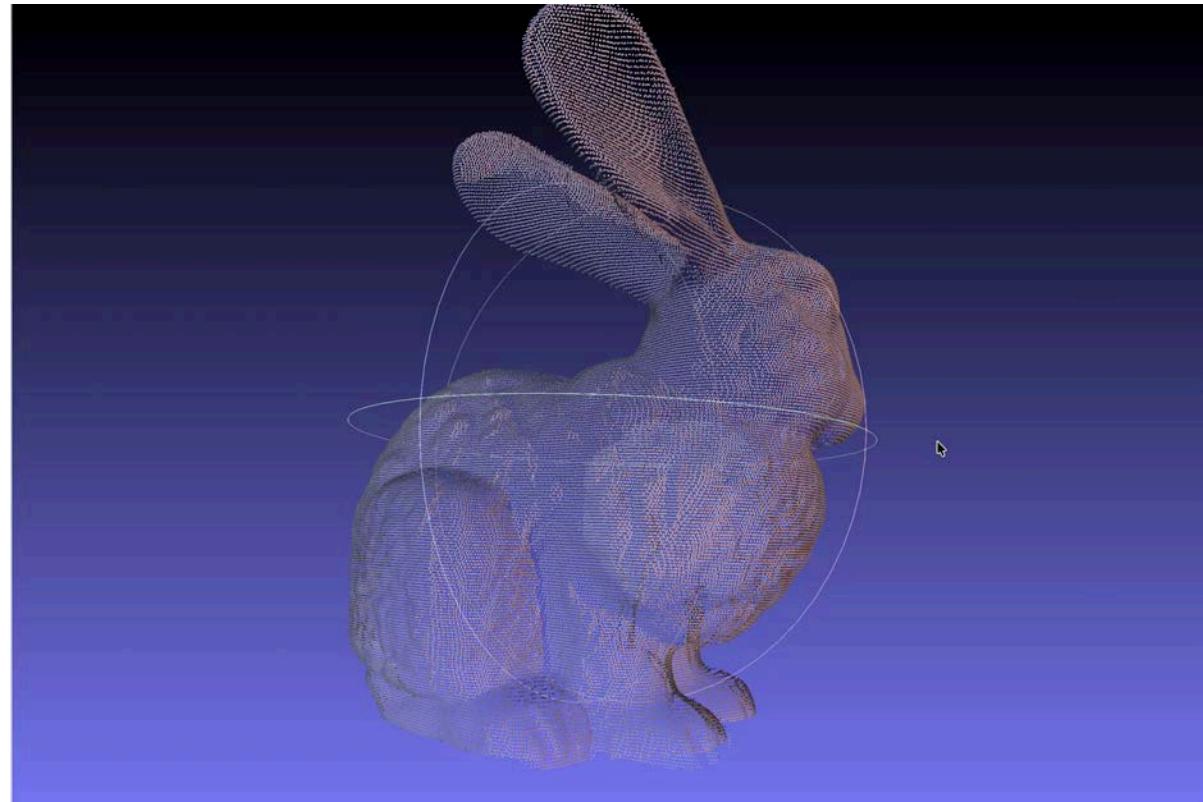
- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



ICP in Meshlab

Iterative Closest-Point Method

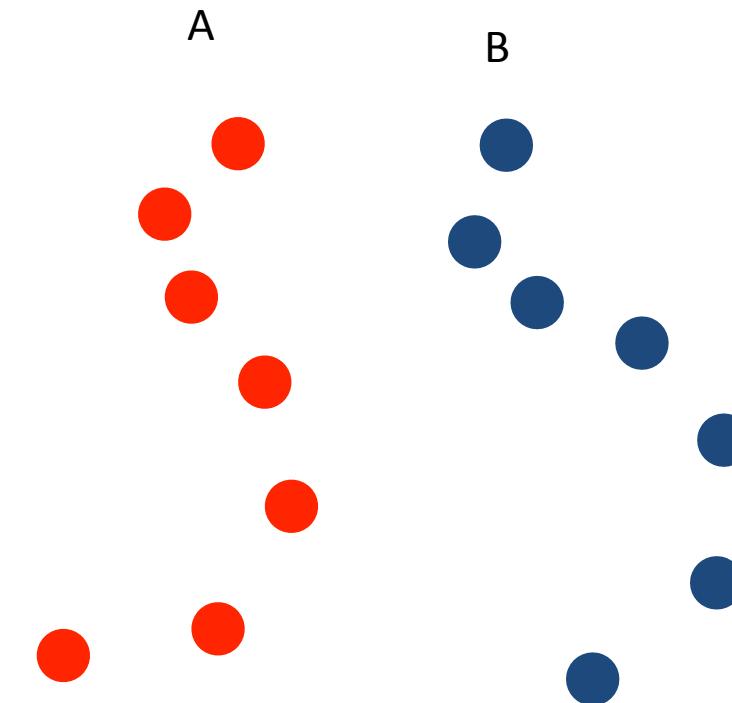
- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



ICP in Meshlab

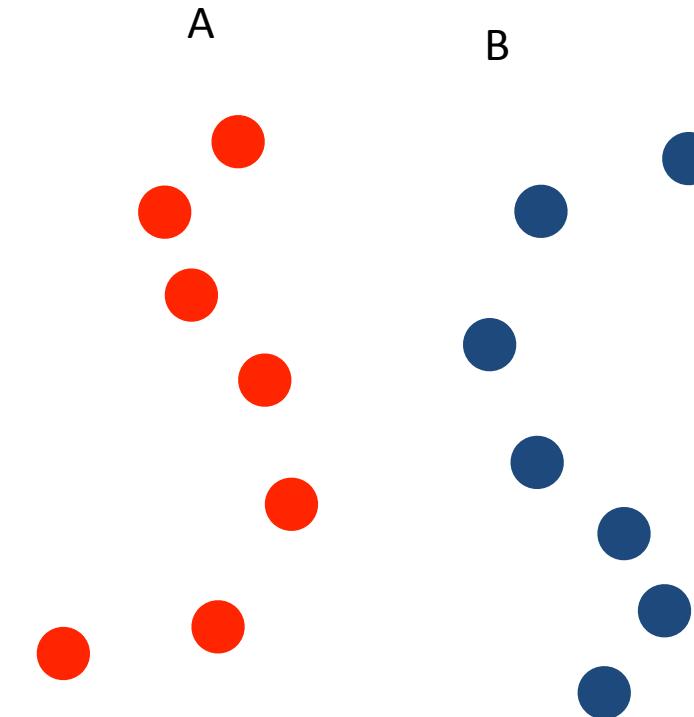
Iterative Closest-Point Method

- How to find a reasonable initial guess?



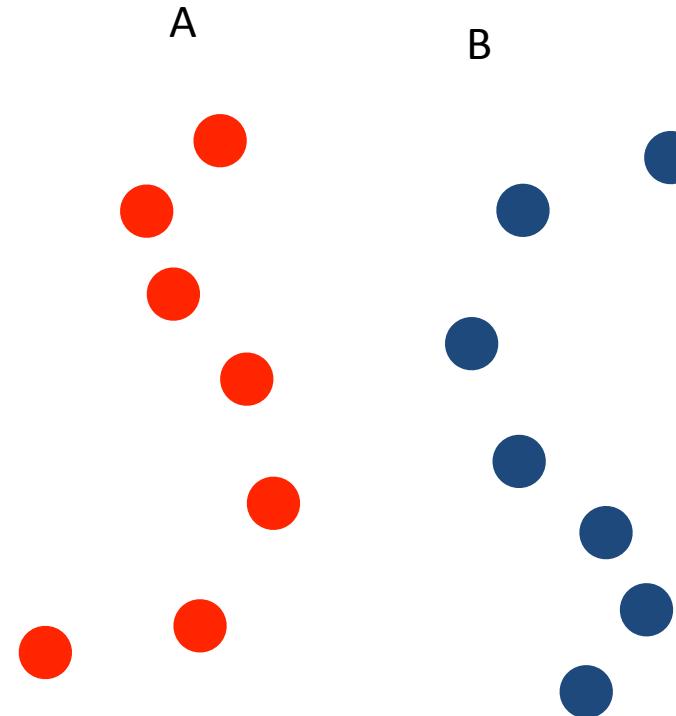
Iterative Closest-Point Method

- How to find a reasonable initial guess?

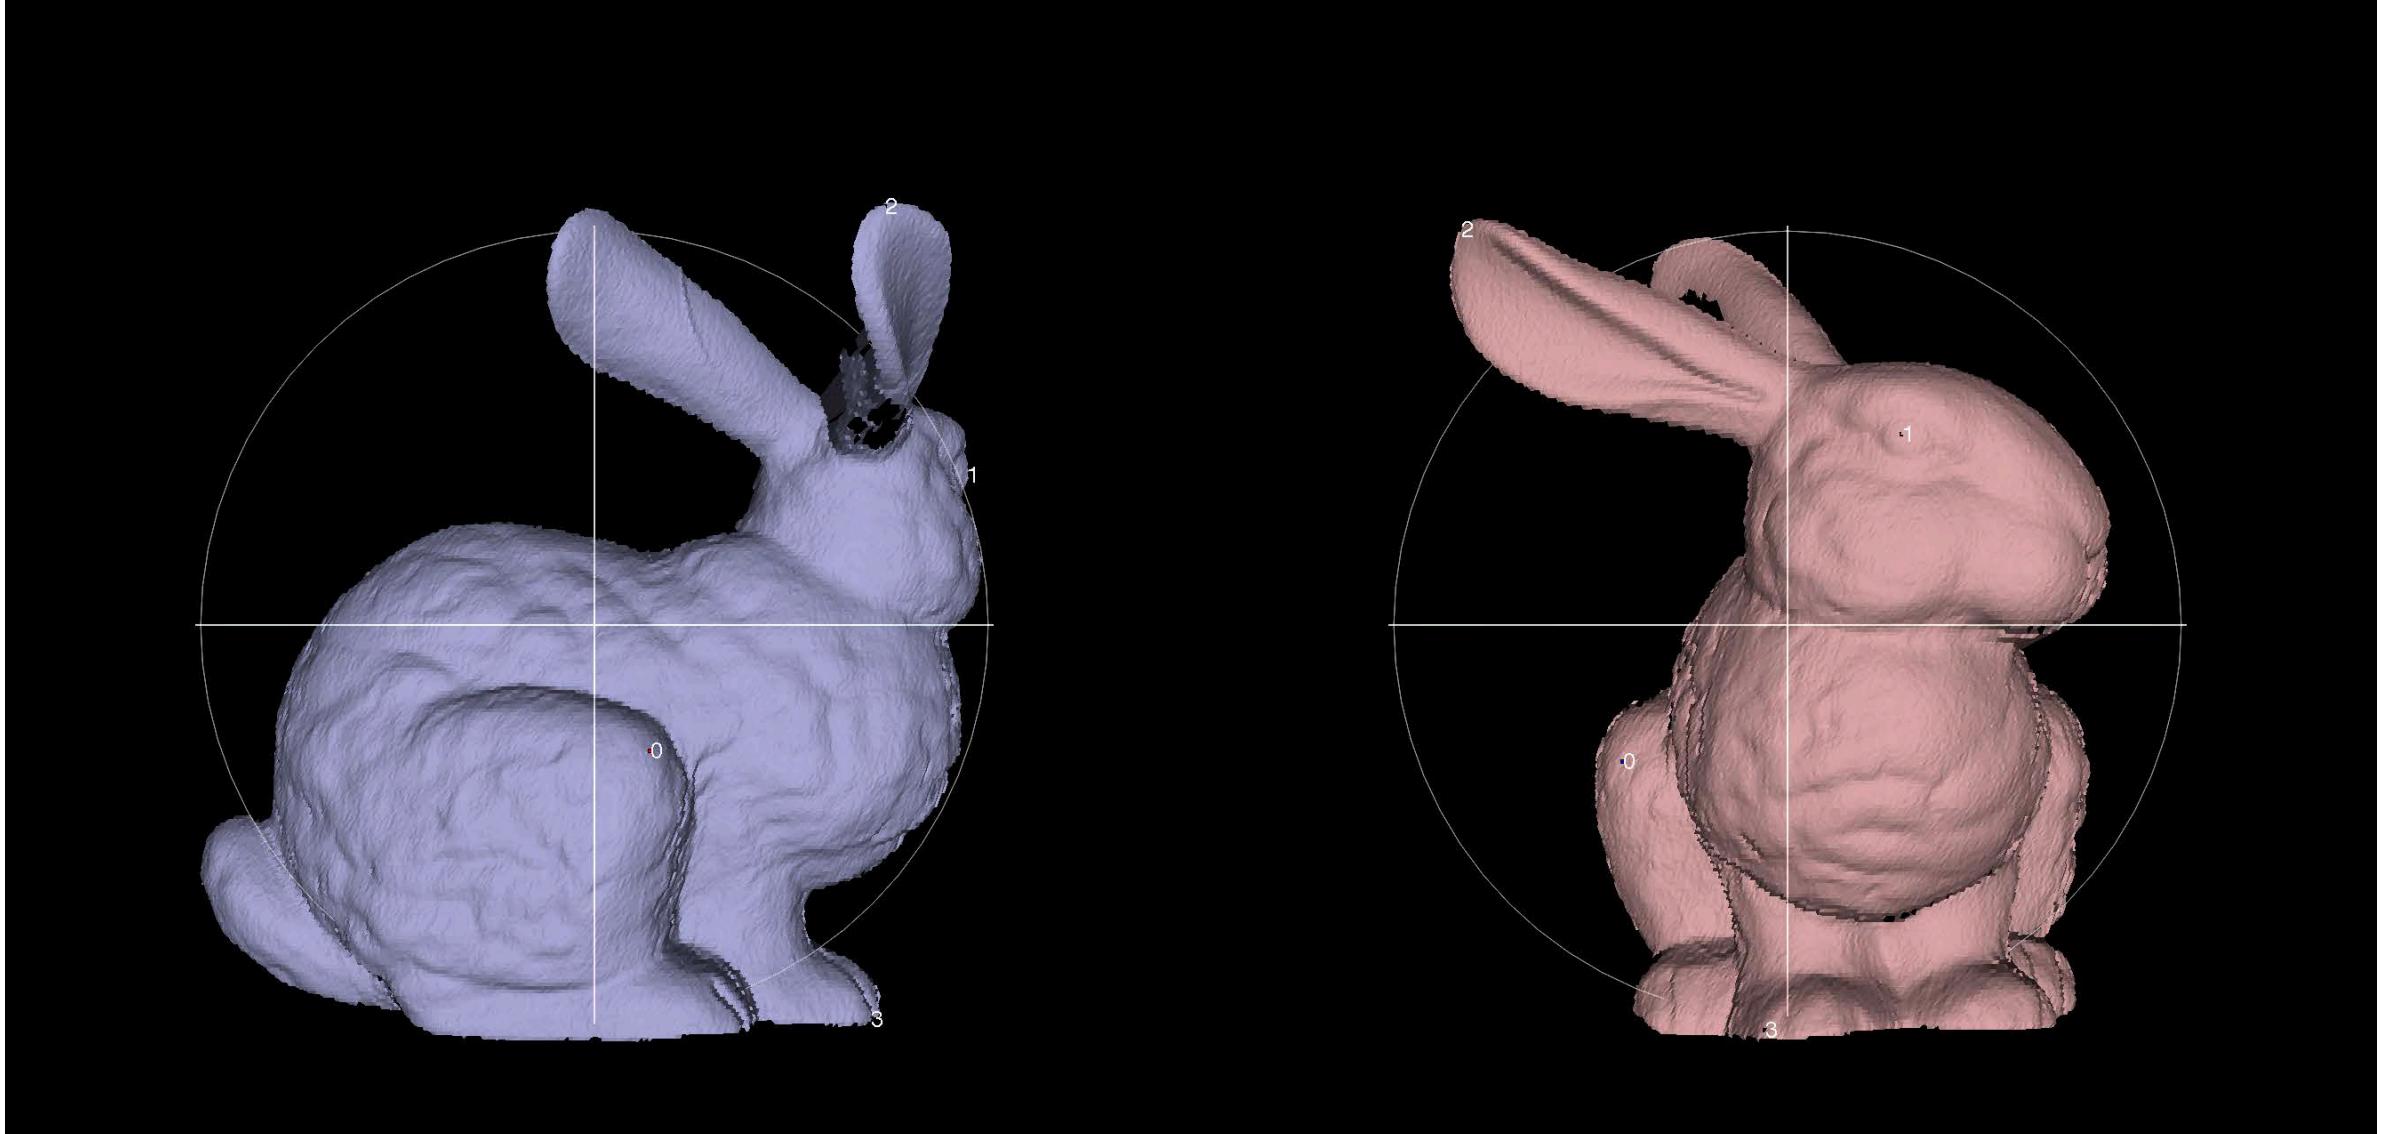


Iterative Closest-Point Method

- How to find a reasonable initial guess?
 - Manually
 - Matching distinctive features between clouds
 - Using sensor data
 - Global registration algorithms
(algorithms not needing initial guess)
- At every iteration finding the closest point M in A to a given point S in B takes (for n points) how long?
 - $O(n^2)$ brute-force
- Can this be done faster?
 - $O(\log n)$ for nearest-neighbour matches with k-d trees



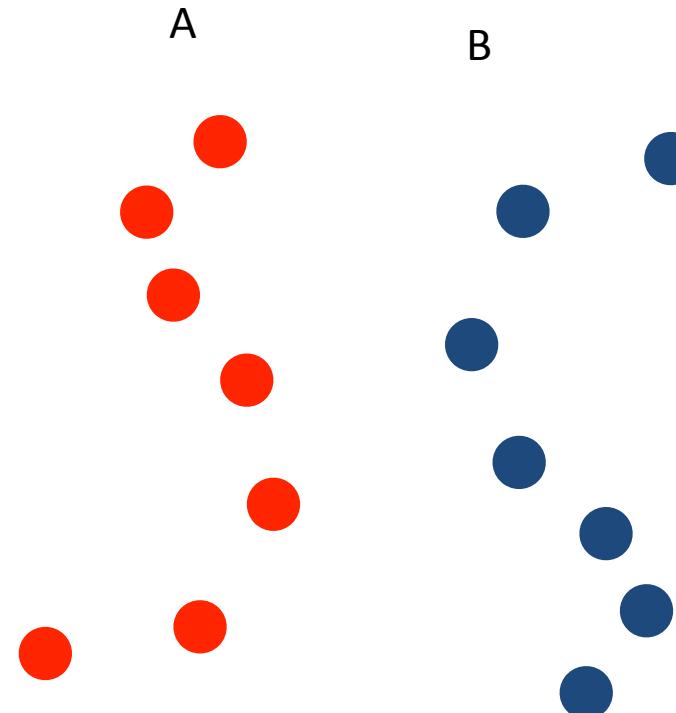
Iterative Closest-Point Method



Manual alignment in Meshlab to initialise ICP

Iterative Closest-Point Method

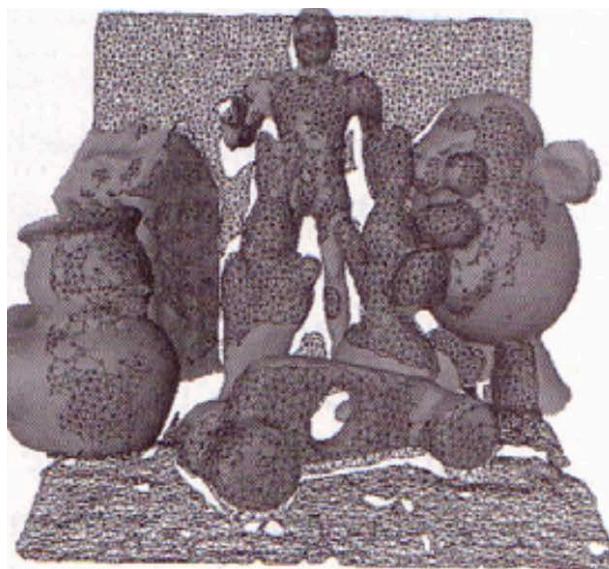
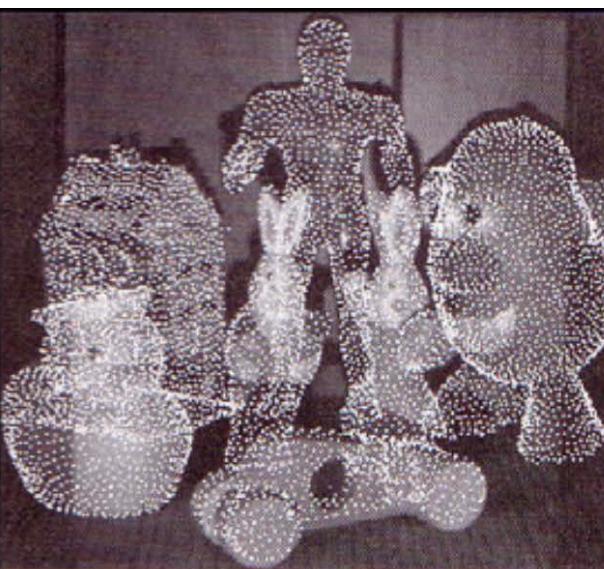
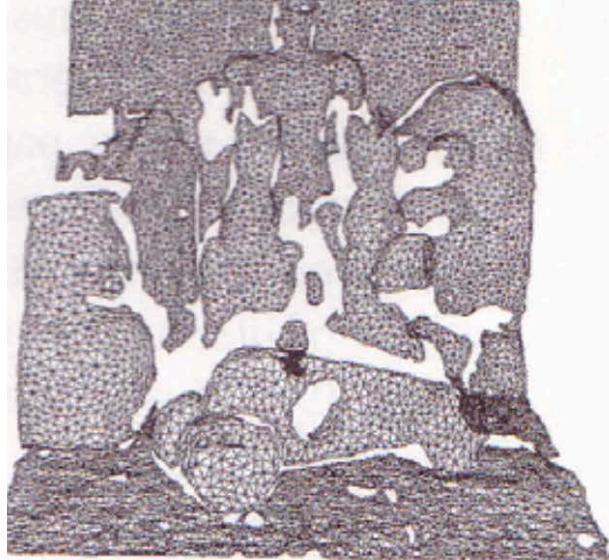
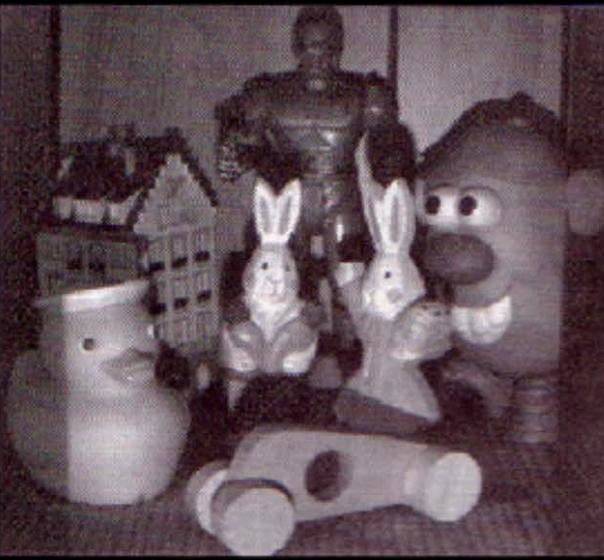
- How to find a reasonable initial guess?
 - Manually
 - Matching distinctive features between clouds
 - Using sensor data
 - Global registration algorithms
(algorithms not needing initial guess)
- At every iteration finding the closest point M in A to a given point S in B takes (for n points) how long?
 - $O(n^2)$ brute-force
- Can this be done faster?
 - $O(\log n)$ for nearest-neighbour matches with k-d trees



Processing Depth Maps

Processing Depth Maps

- Depth maps can be processed like images
- Image operators, such as filters, can be applied to:
 - Handle discontinuities and overlaps
 - Fill missing portions
 - Smoothen geometric noise
 - Analyse depth images
 - Find features such as edges
 - Segmentation of objects
 - Etc.



The end!