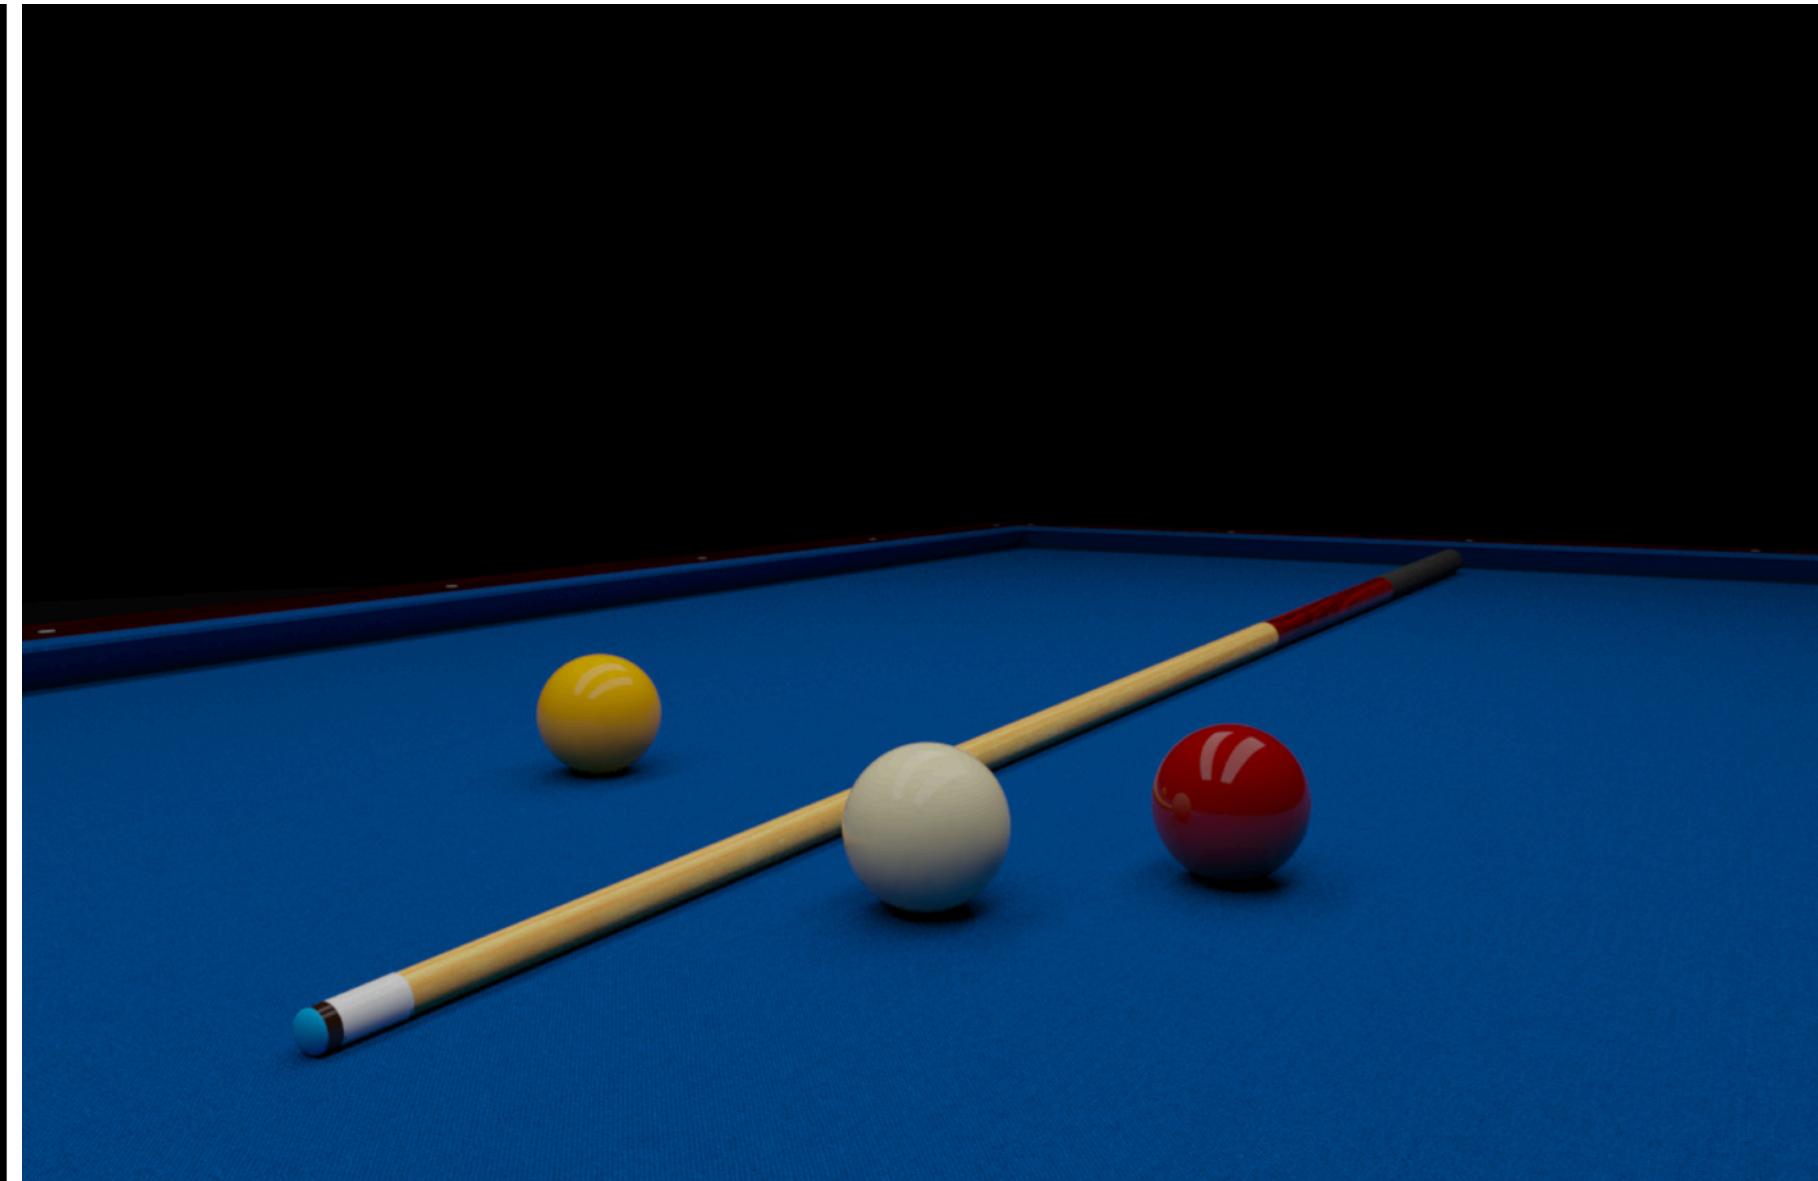
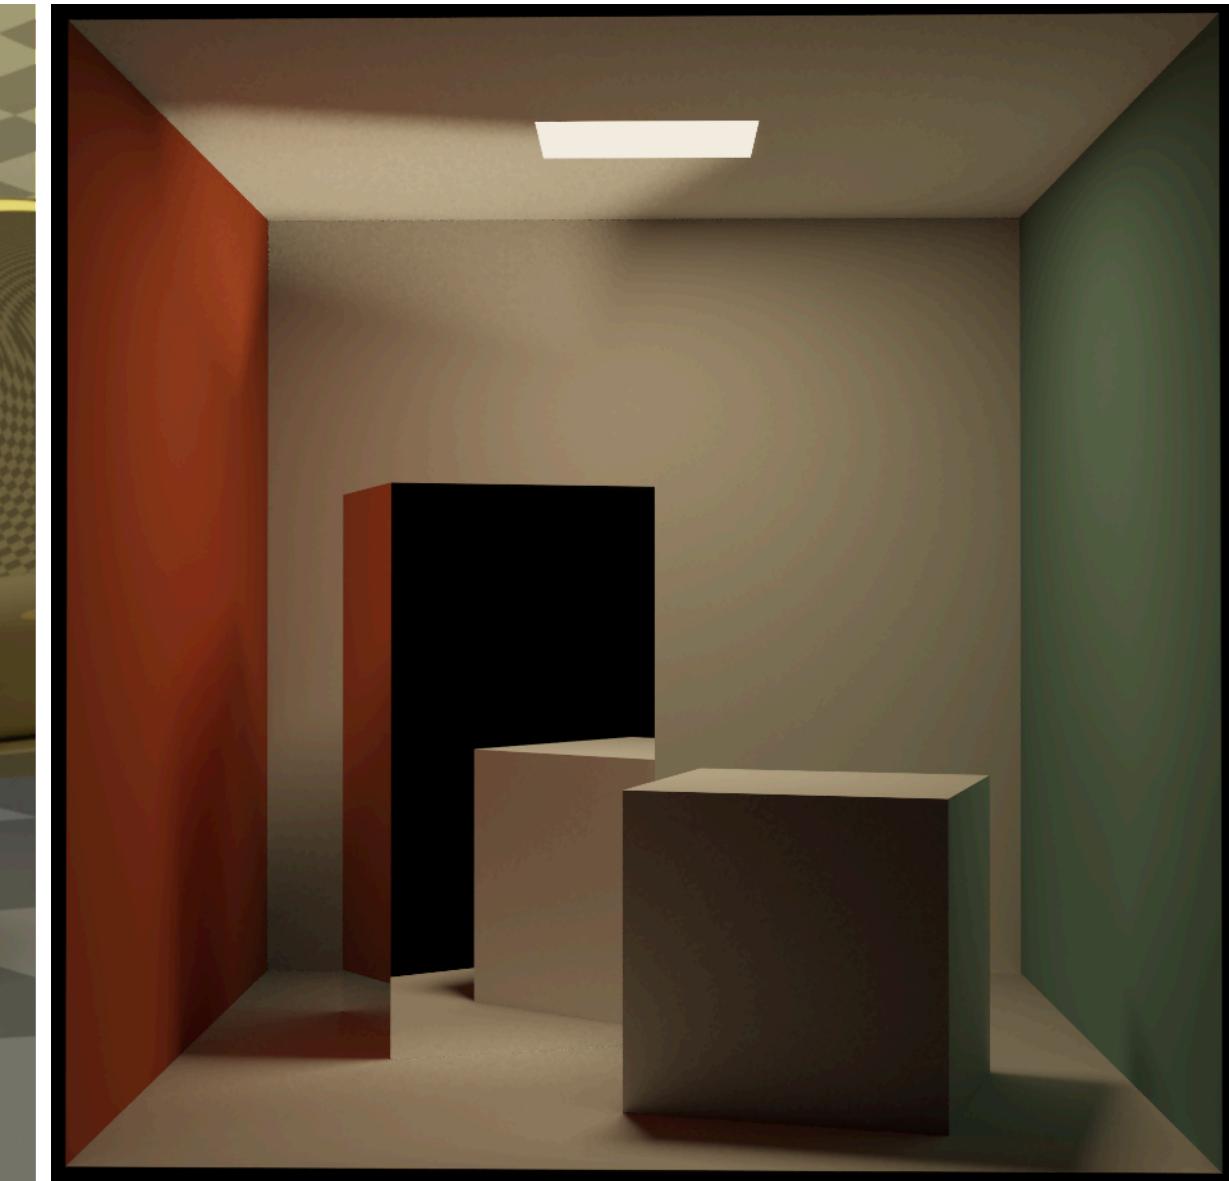
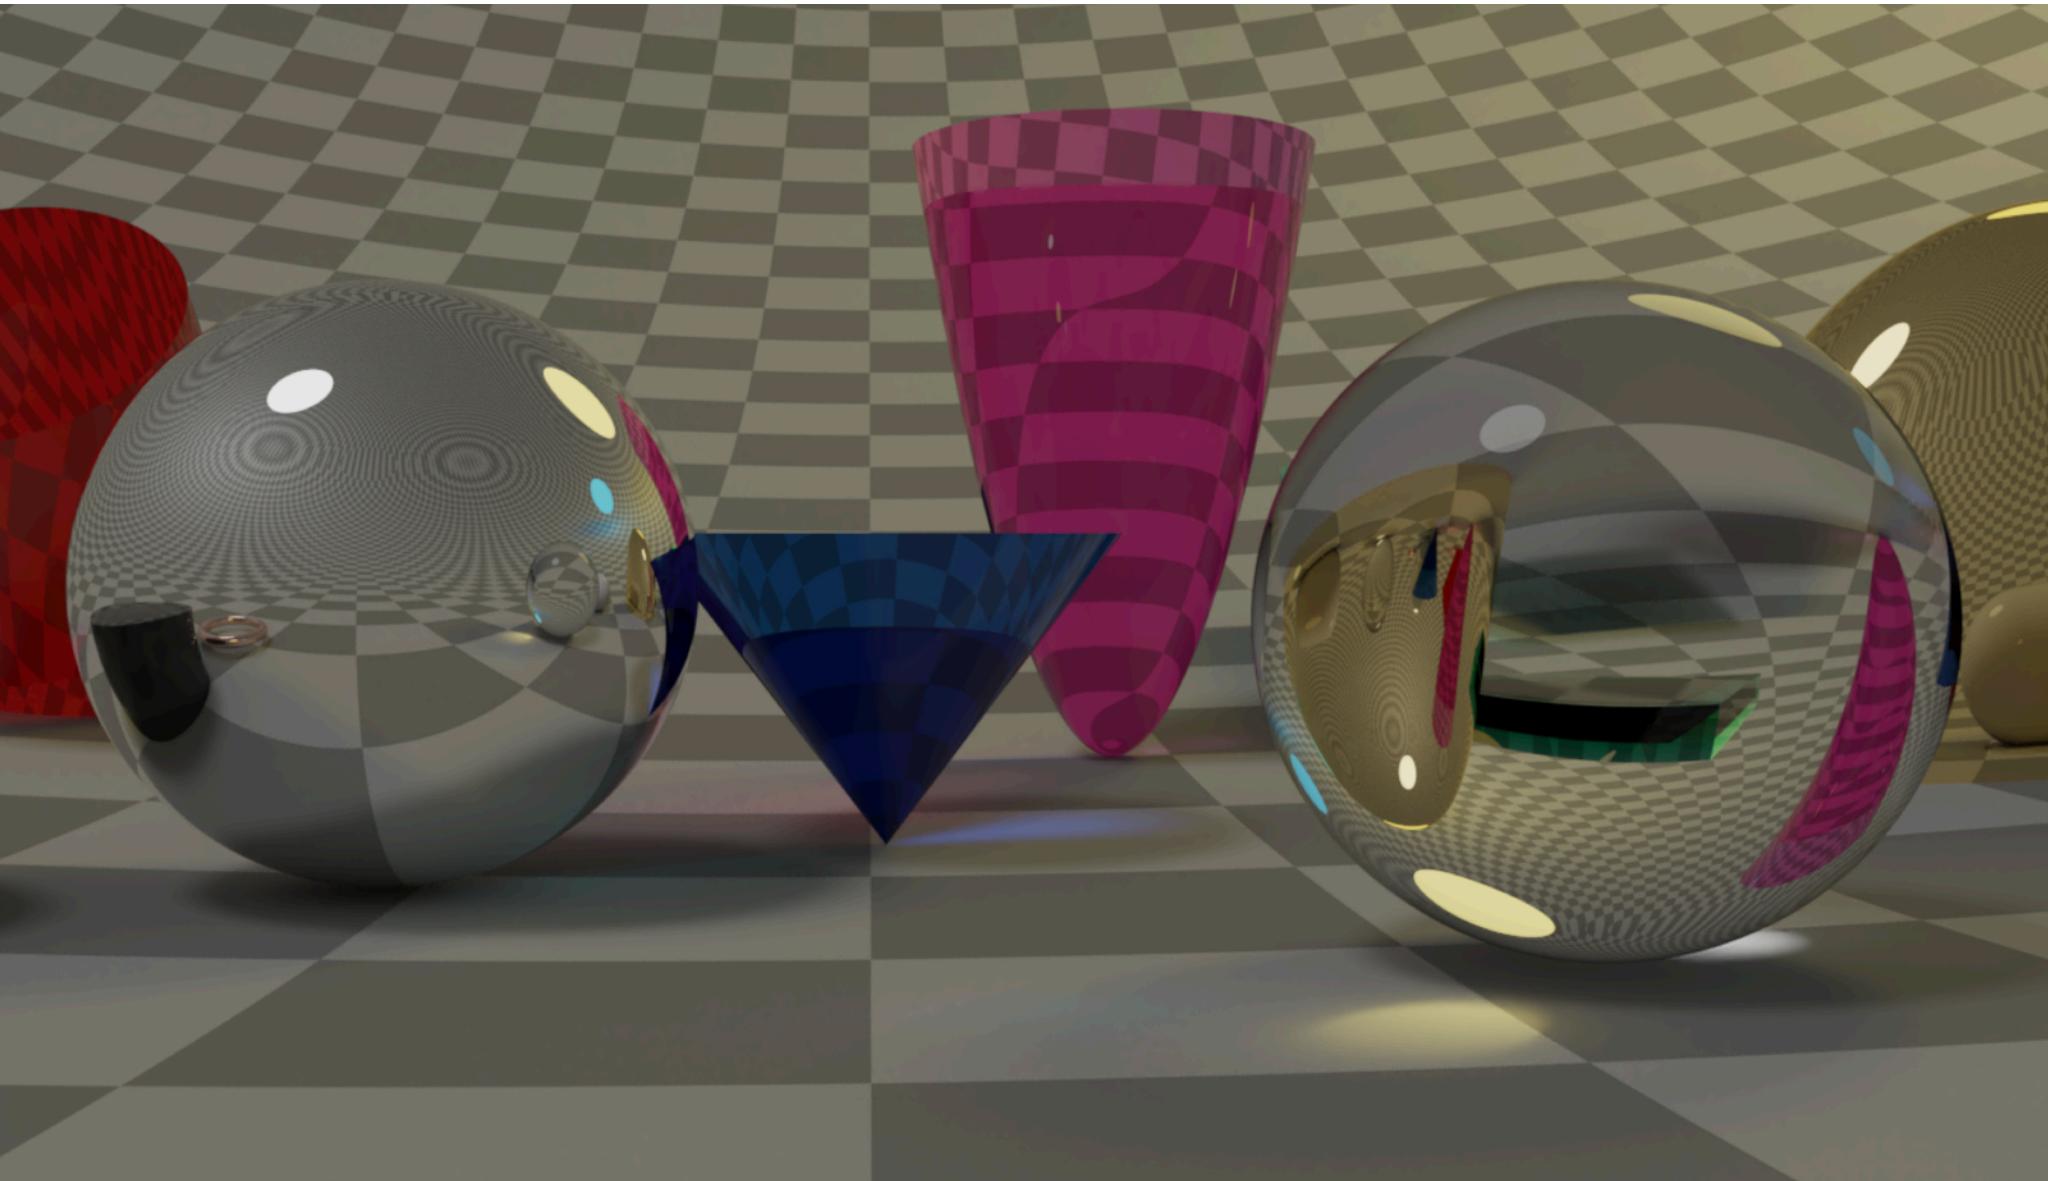


Visual Computing I:

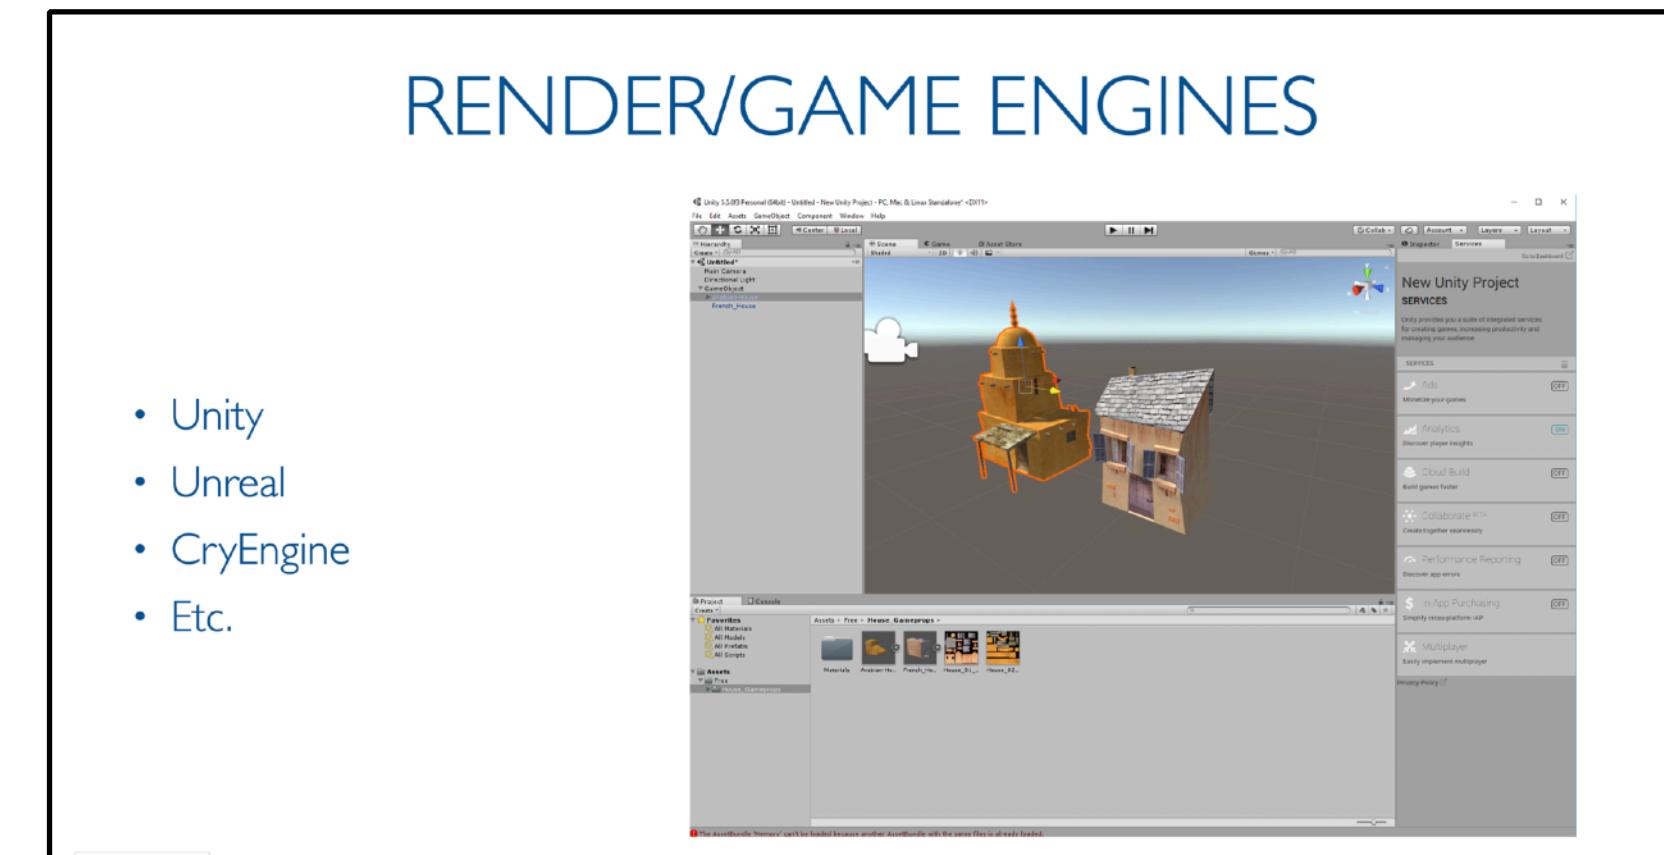
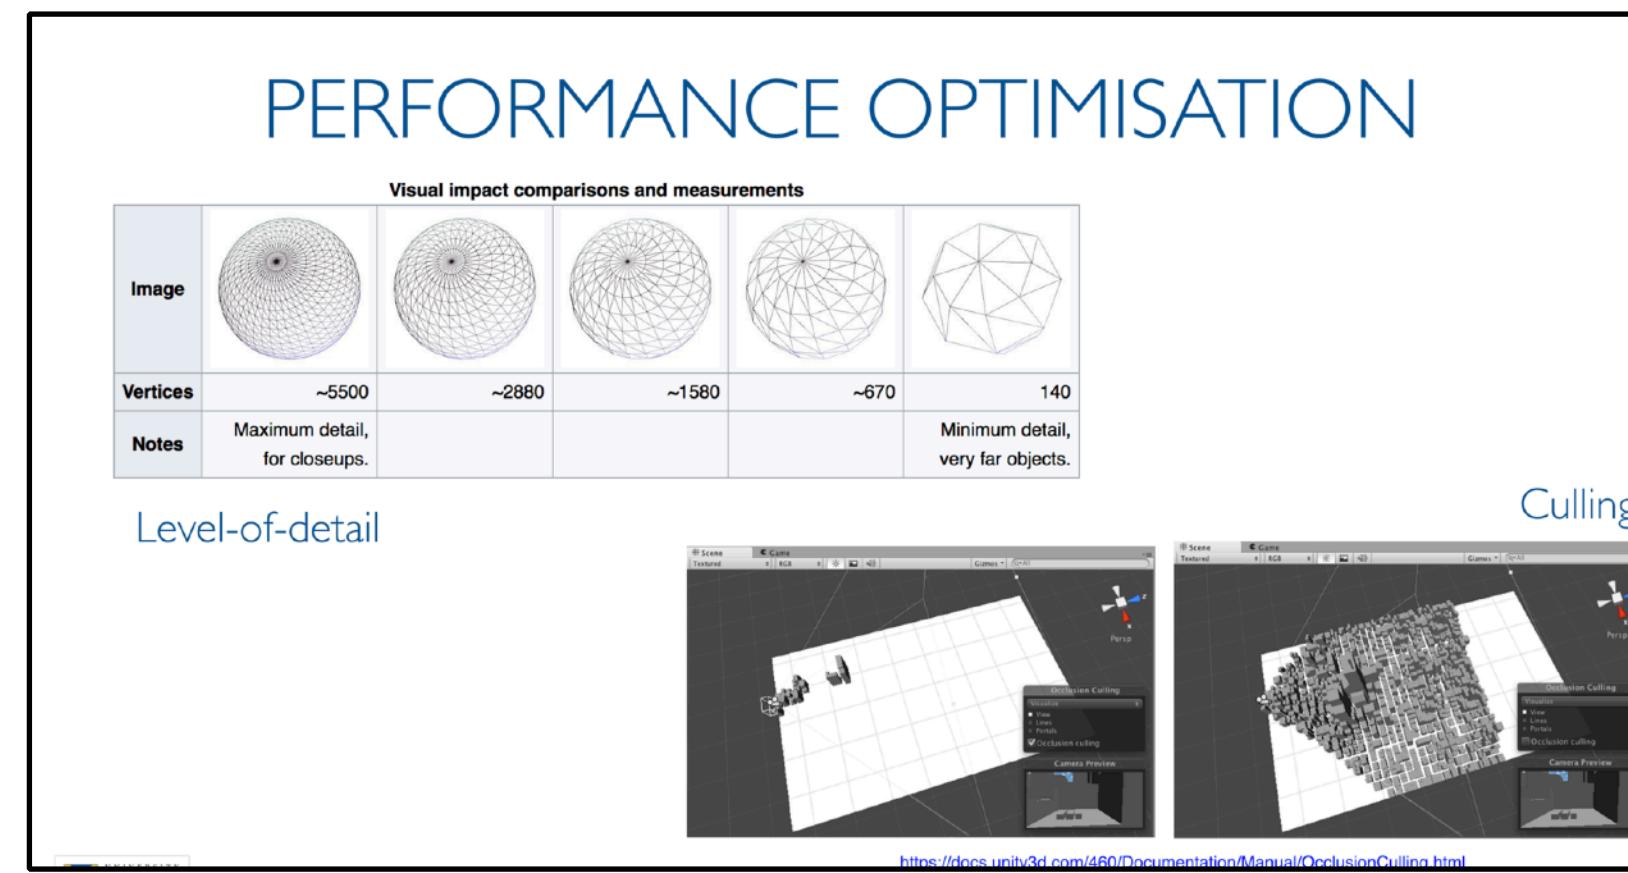
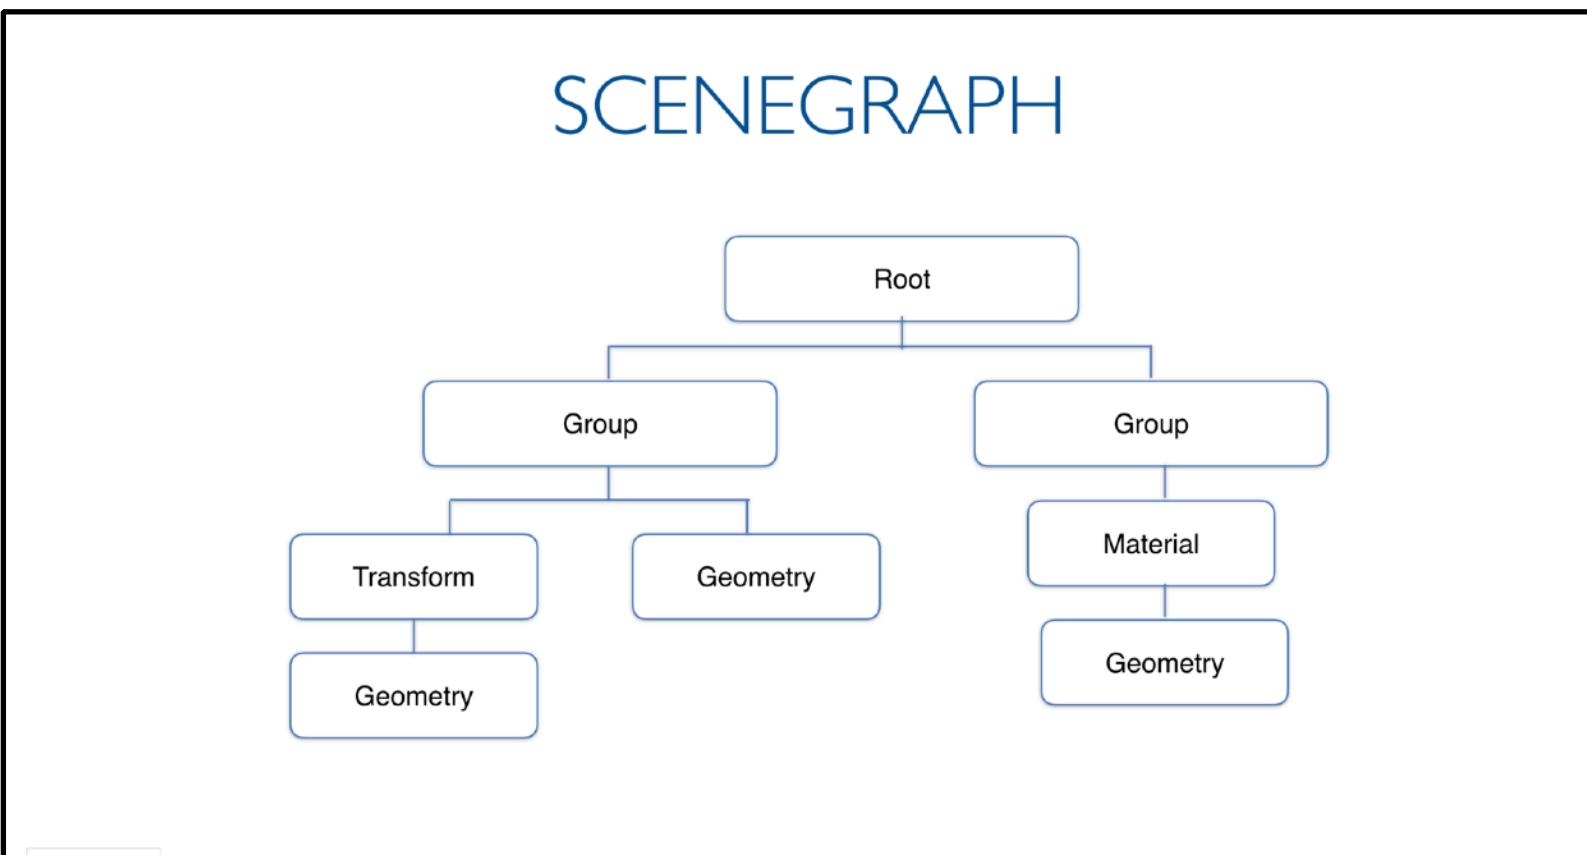
Interactive Computer Graphics and Vision



Raytracing

Stefanie Zollmann and Tobias Langlotz

Last time



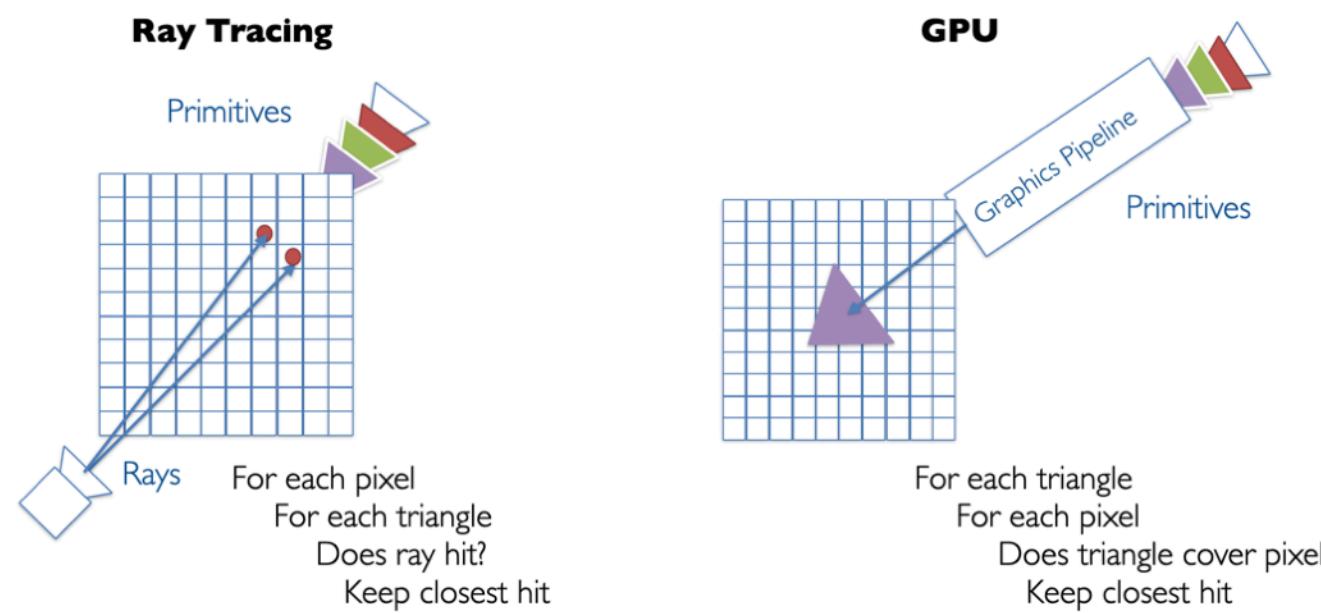
SCENEGRAPH

OPTIMISATIONS

GAME ENGINES

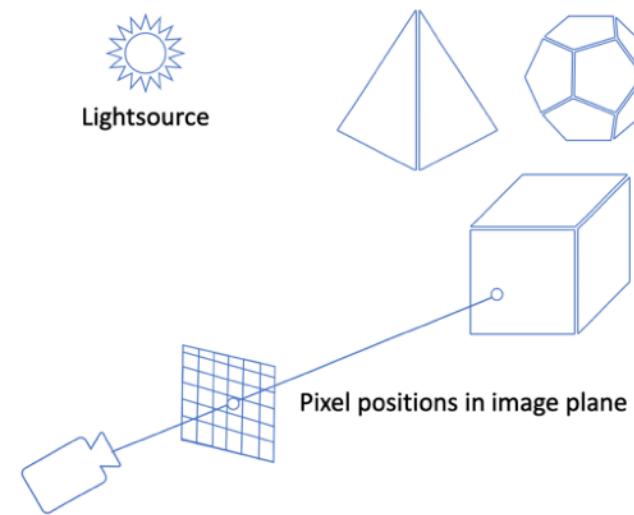
Today

Raytracing vs Rasterised Graphics



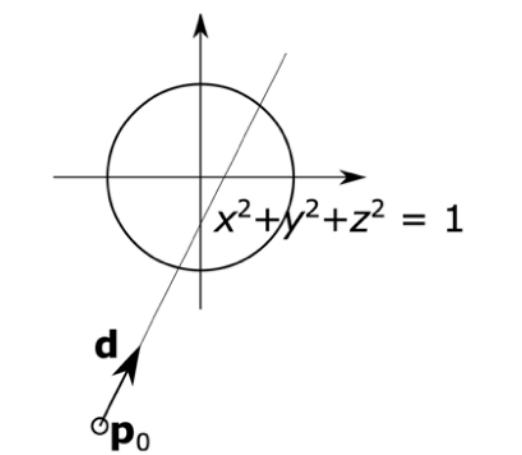
Raytracing Process

- Create an image plane and viewpoint
- For each pixel trace a 'ray' from the eye through a corresponding point in the image plane
- For each ray, return the colour of the object at the hit point (closest to the camera)



Ray-Sphere Intersections

- This is a quadratic equation in $a\lambda^2 + b\lambda + c = 0$
- So the solutions are $\lambda = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$
- 0, 1, or 2 real solutions (what do they mean)



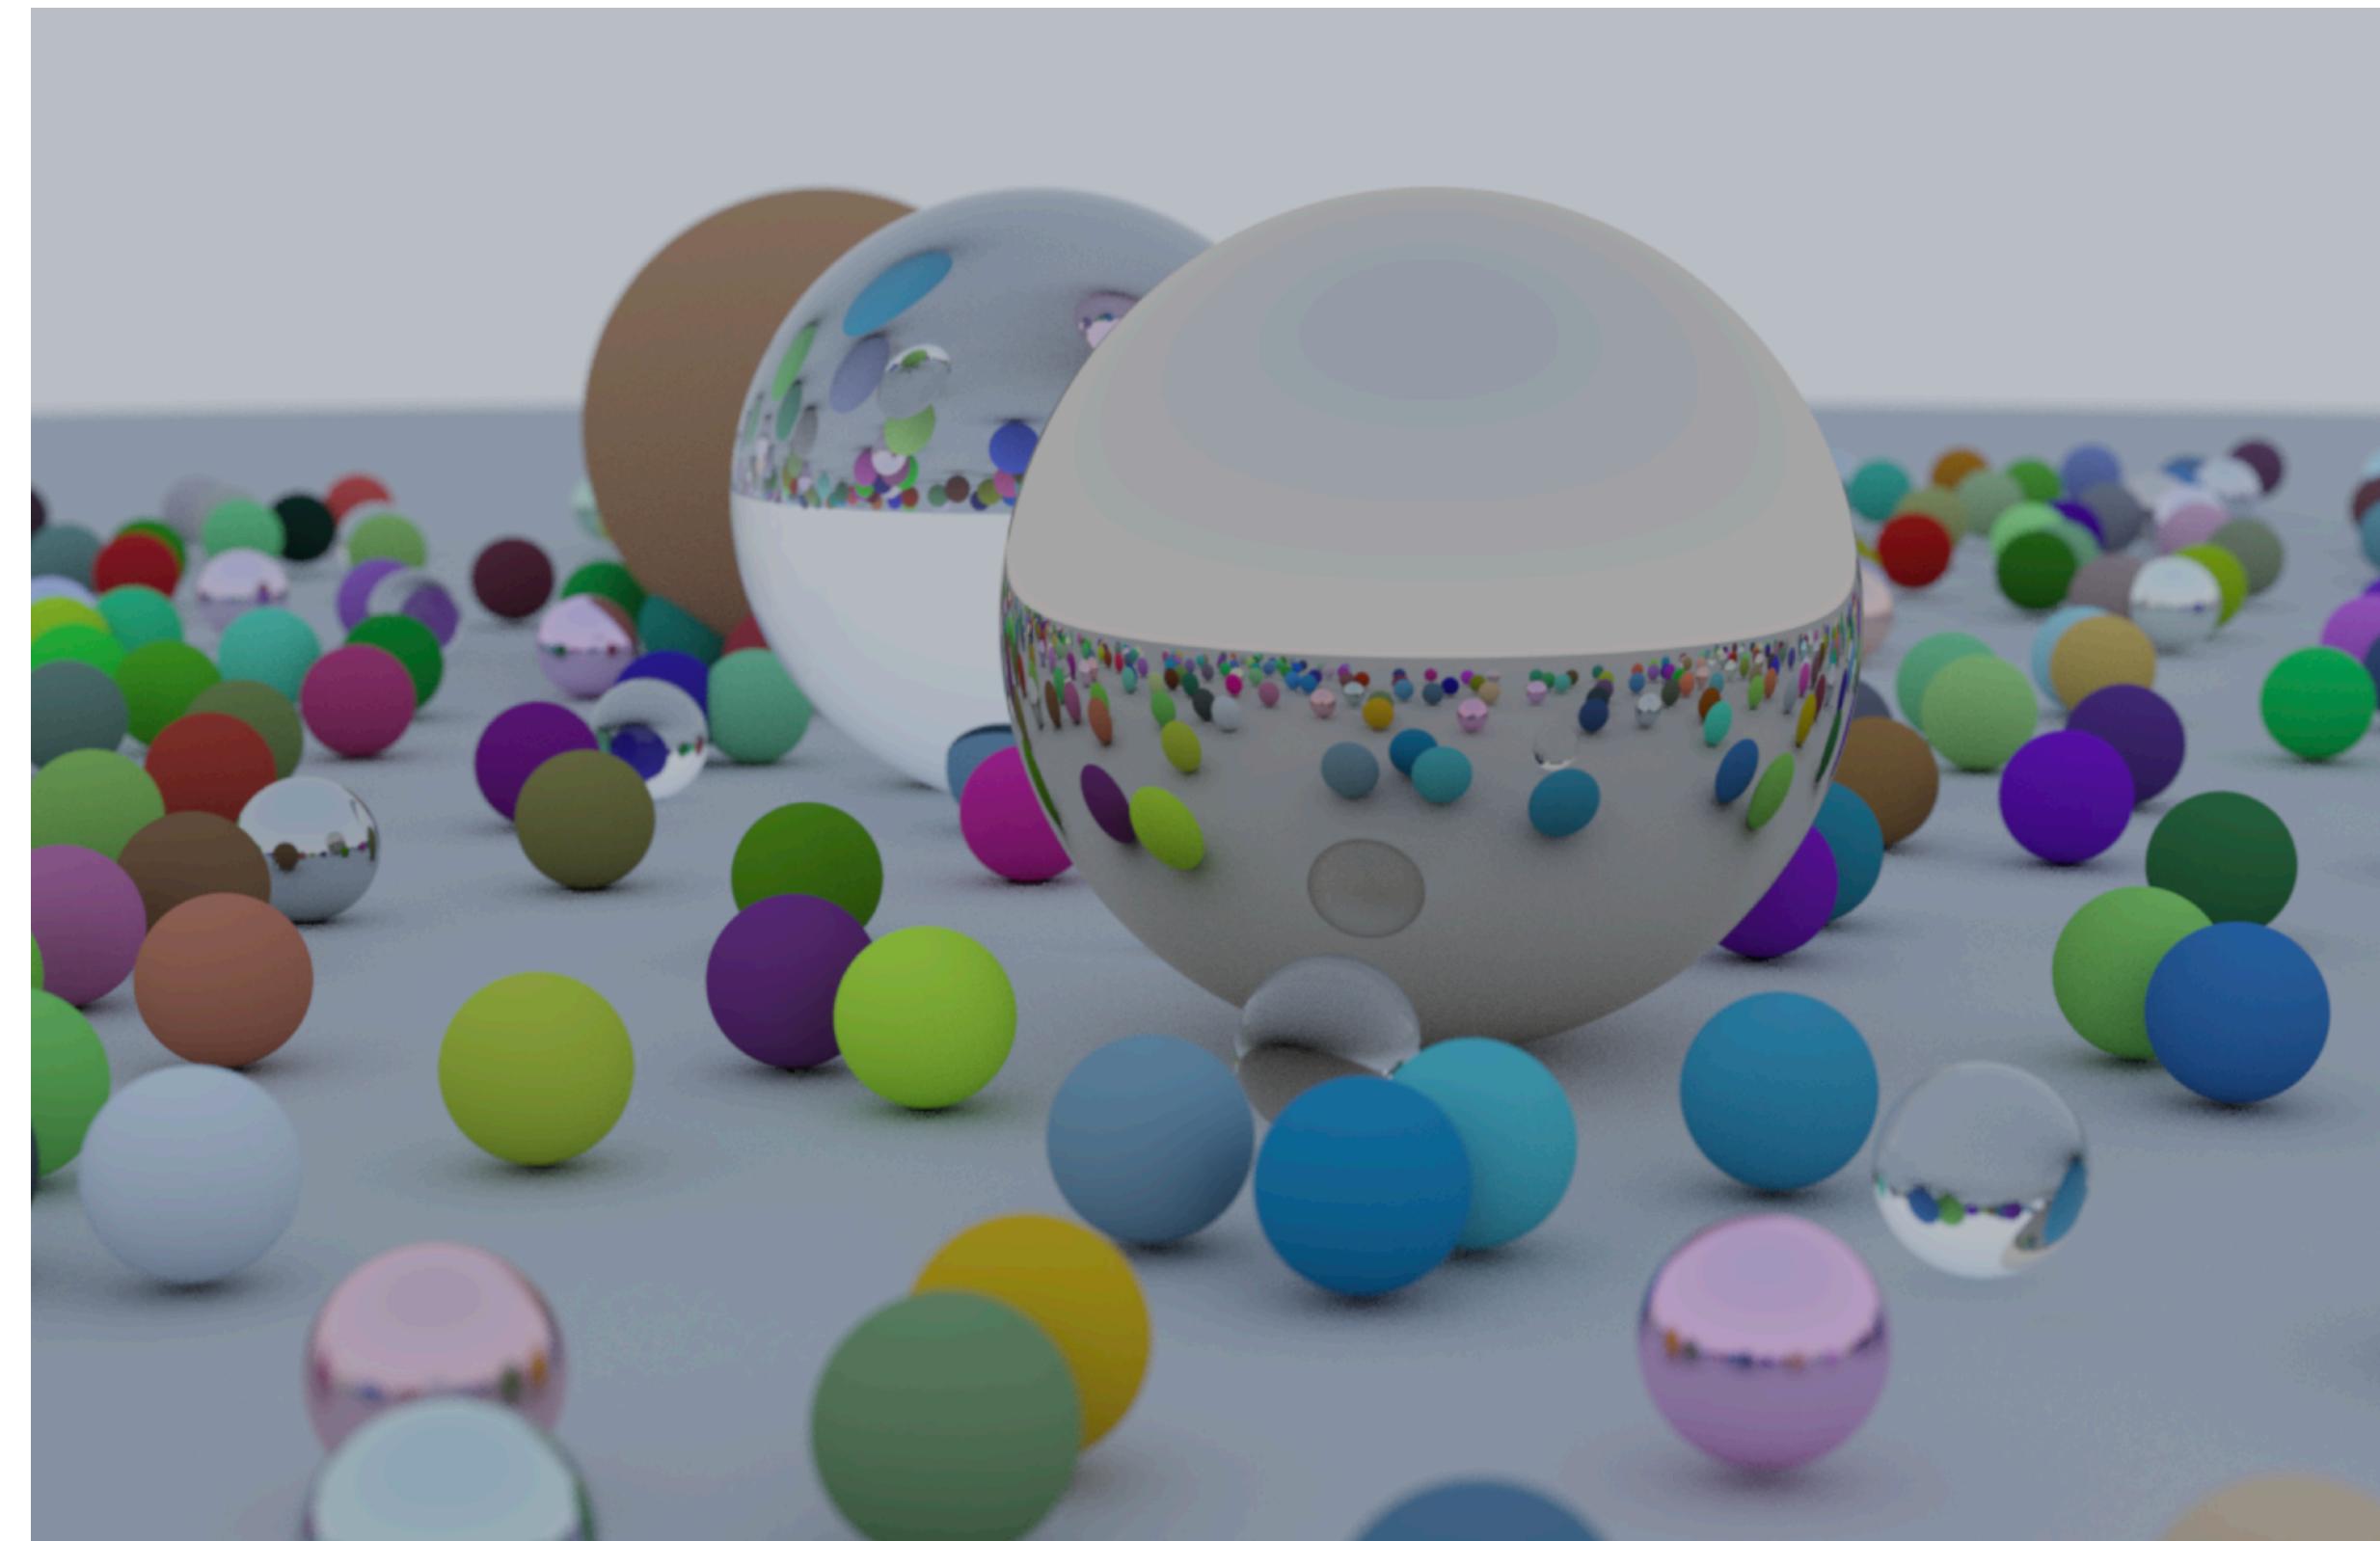
Raytracing/Rasterised

Raytracing Process

Intersections

Motivation

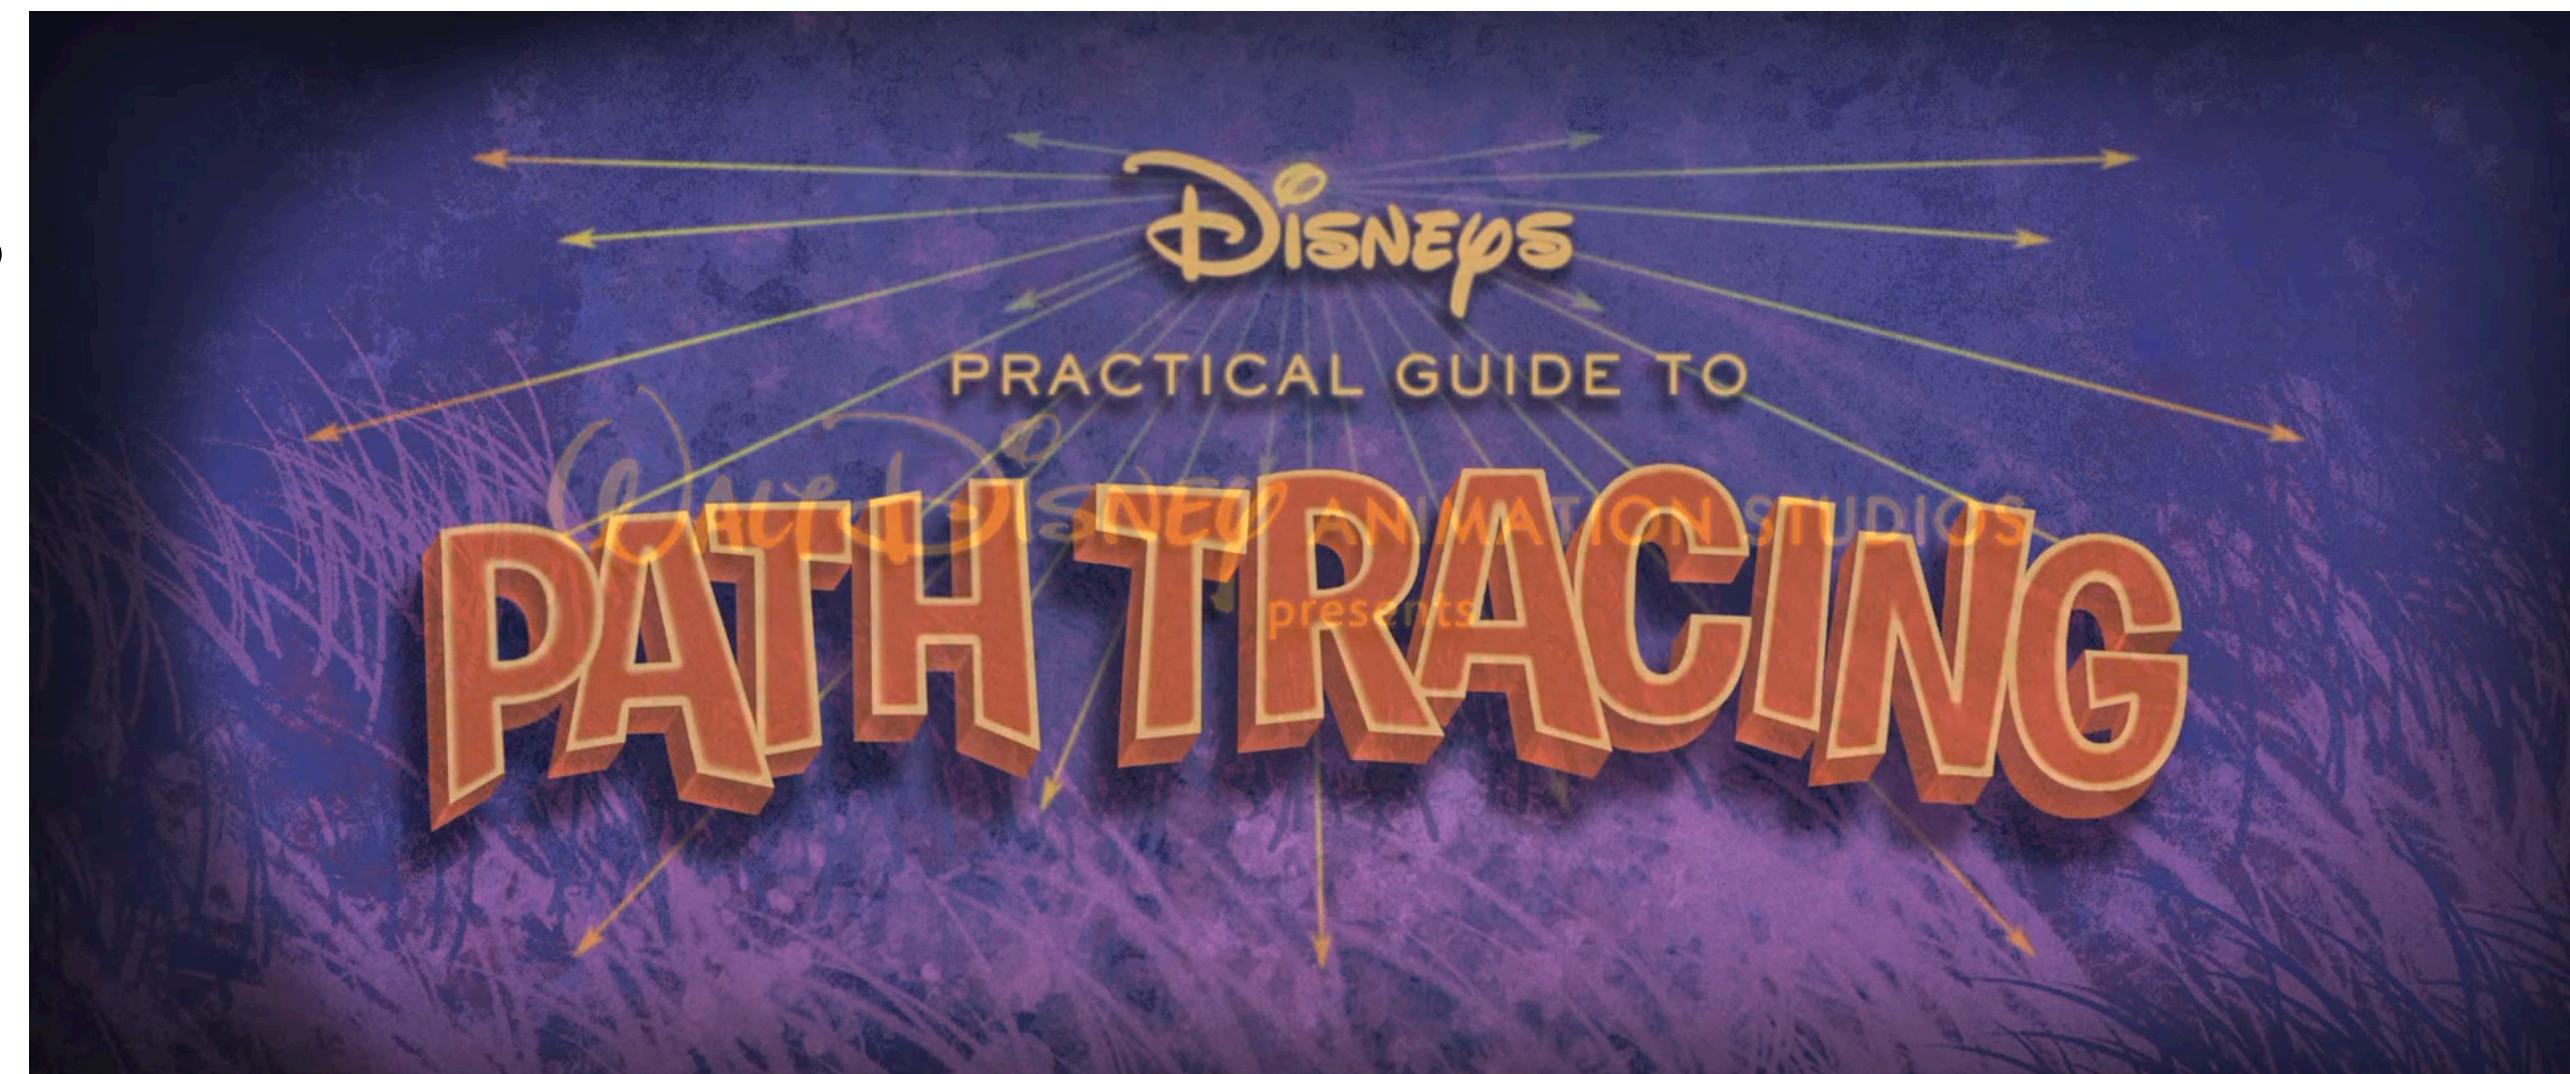
- Rasterisation provides high-performance image synthesis but is limited to local lighting models
- Global illumination phenomena (e.g., indirect lighting, soft shadows, caustics) are difficult to approximate in rasterisation pipelines
- Ray tracing approximates physical light transport, enabling more accurate optical effects



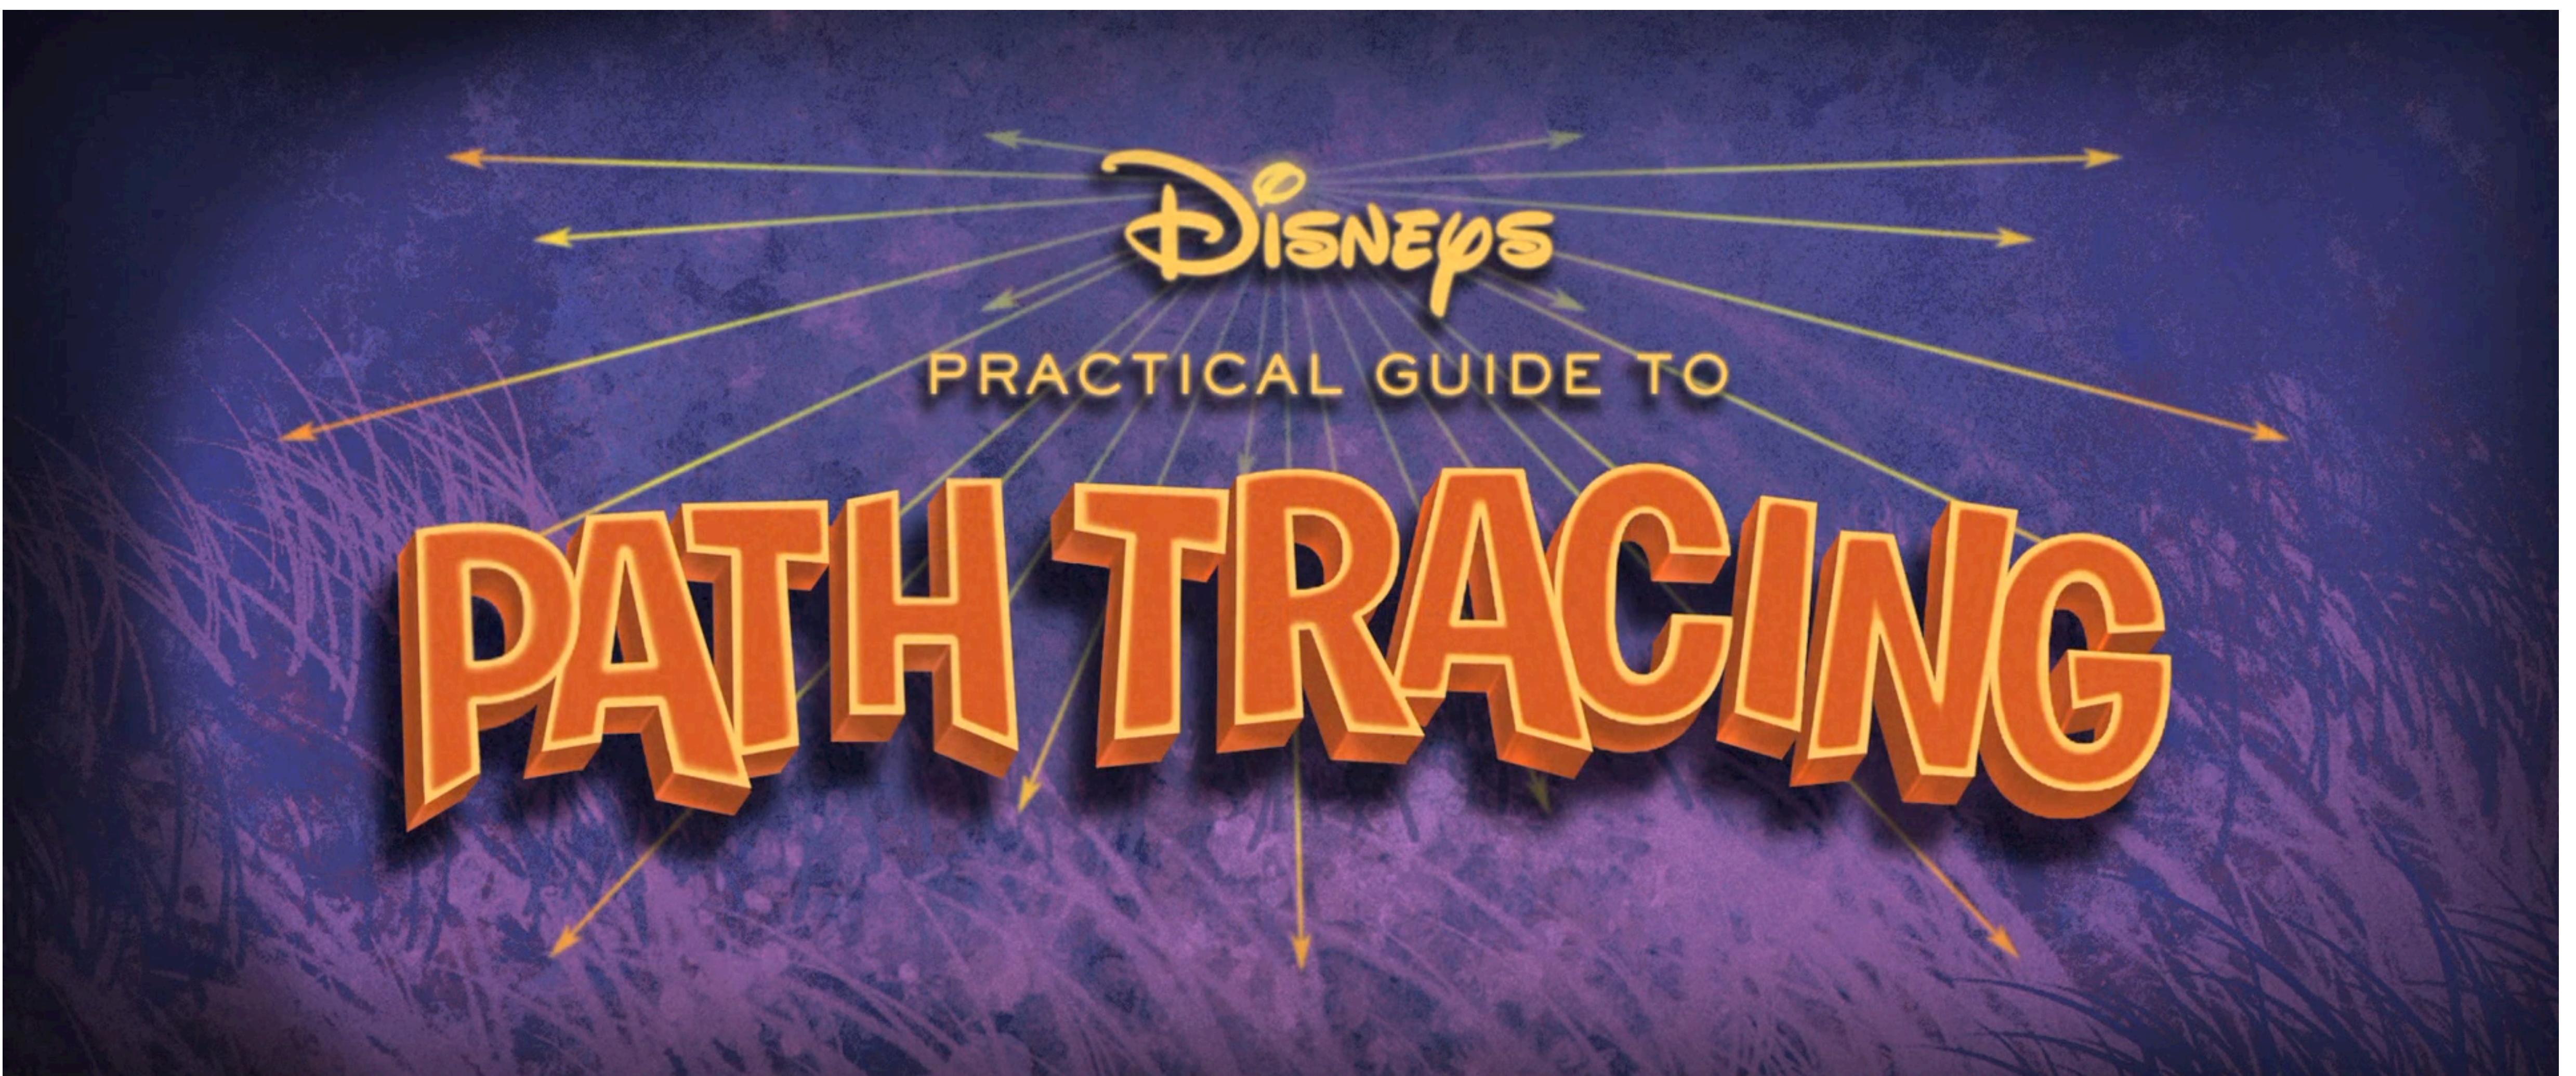
https://erichlof.github.io/THREE.js-PathTracing-Renderer/RayTracing_In_One_Weekend.html

Fundamentals

- Treat light propagation as geometric rays
- For each pixel, generate one or more camera rays
- Determine the nearest surface intersection for each ray
- Evaluate material response at the intersection
- Emit secondary rays as required to account for light interactions

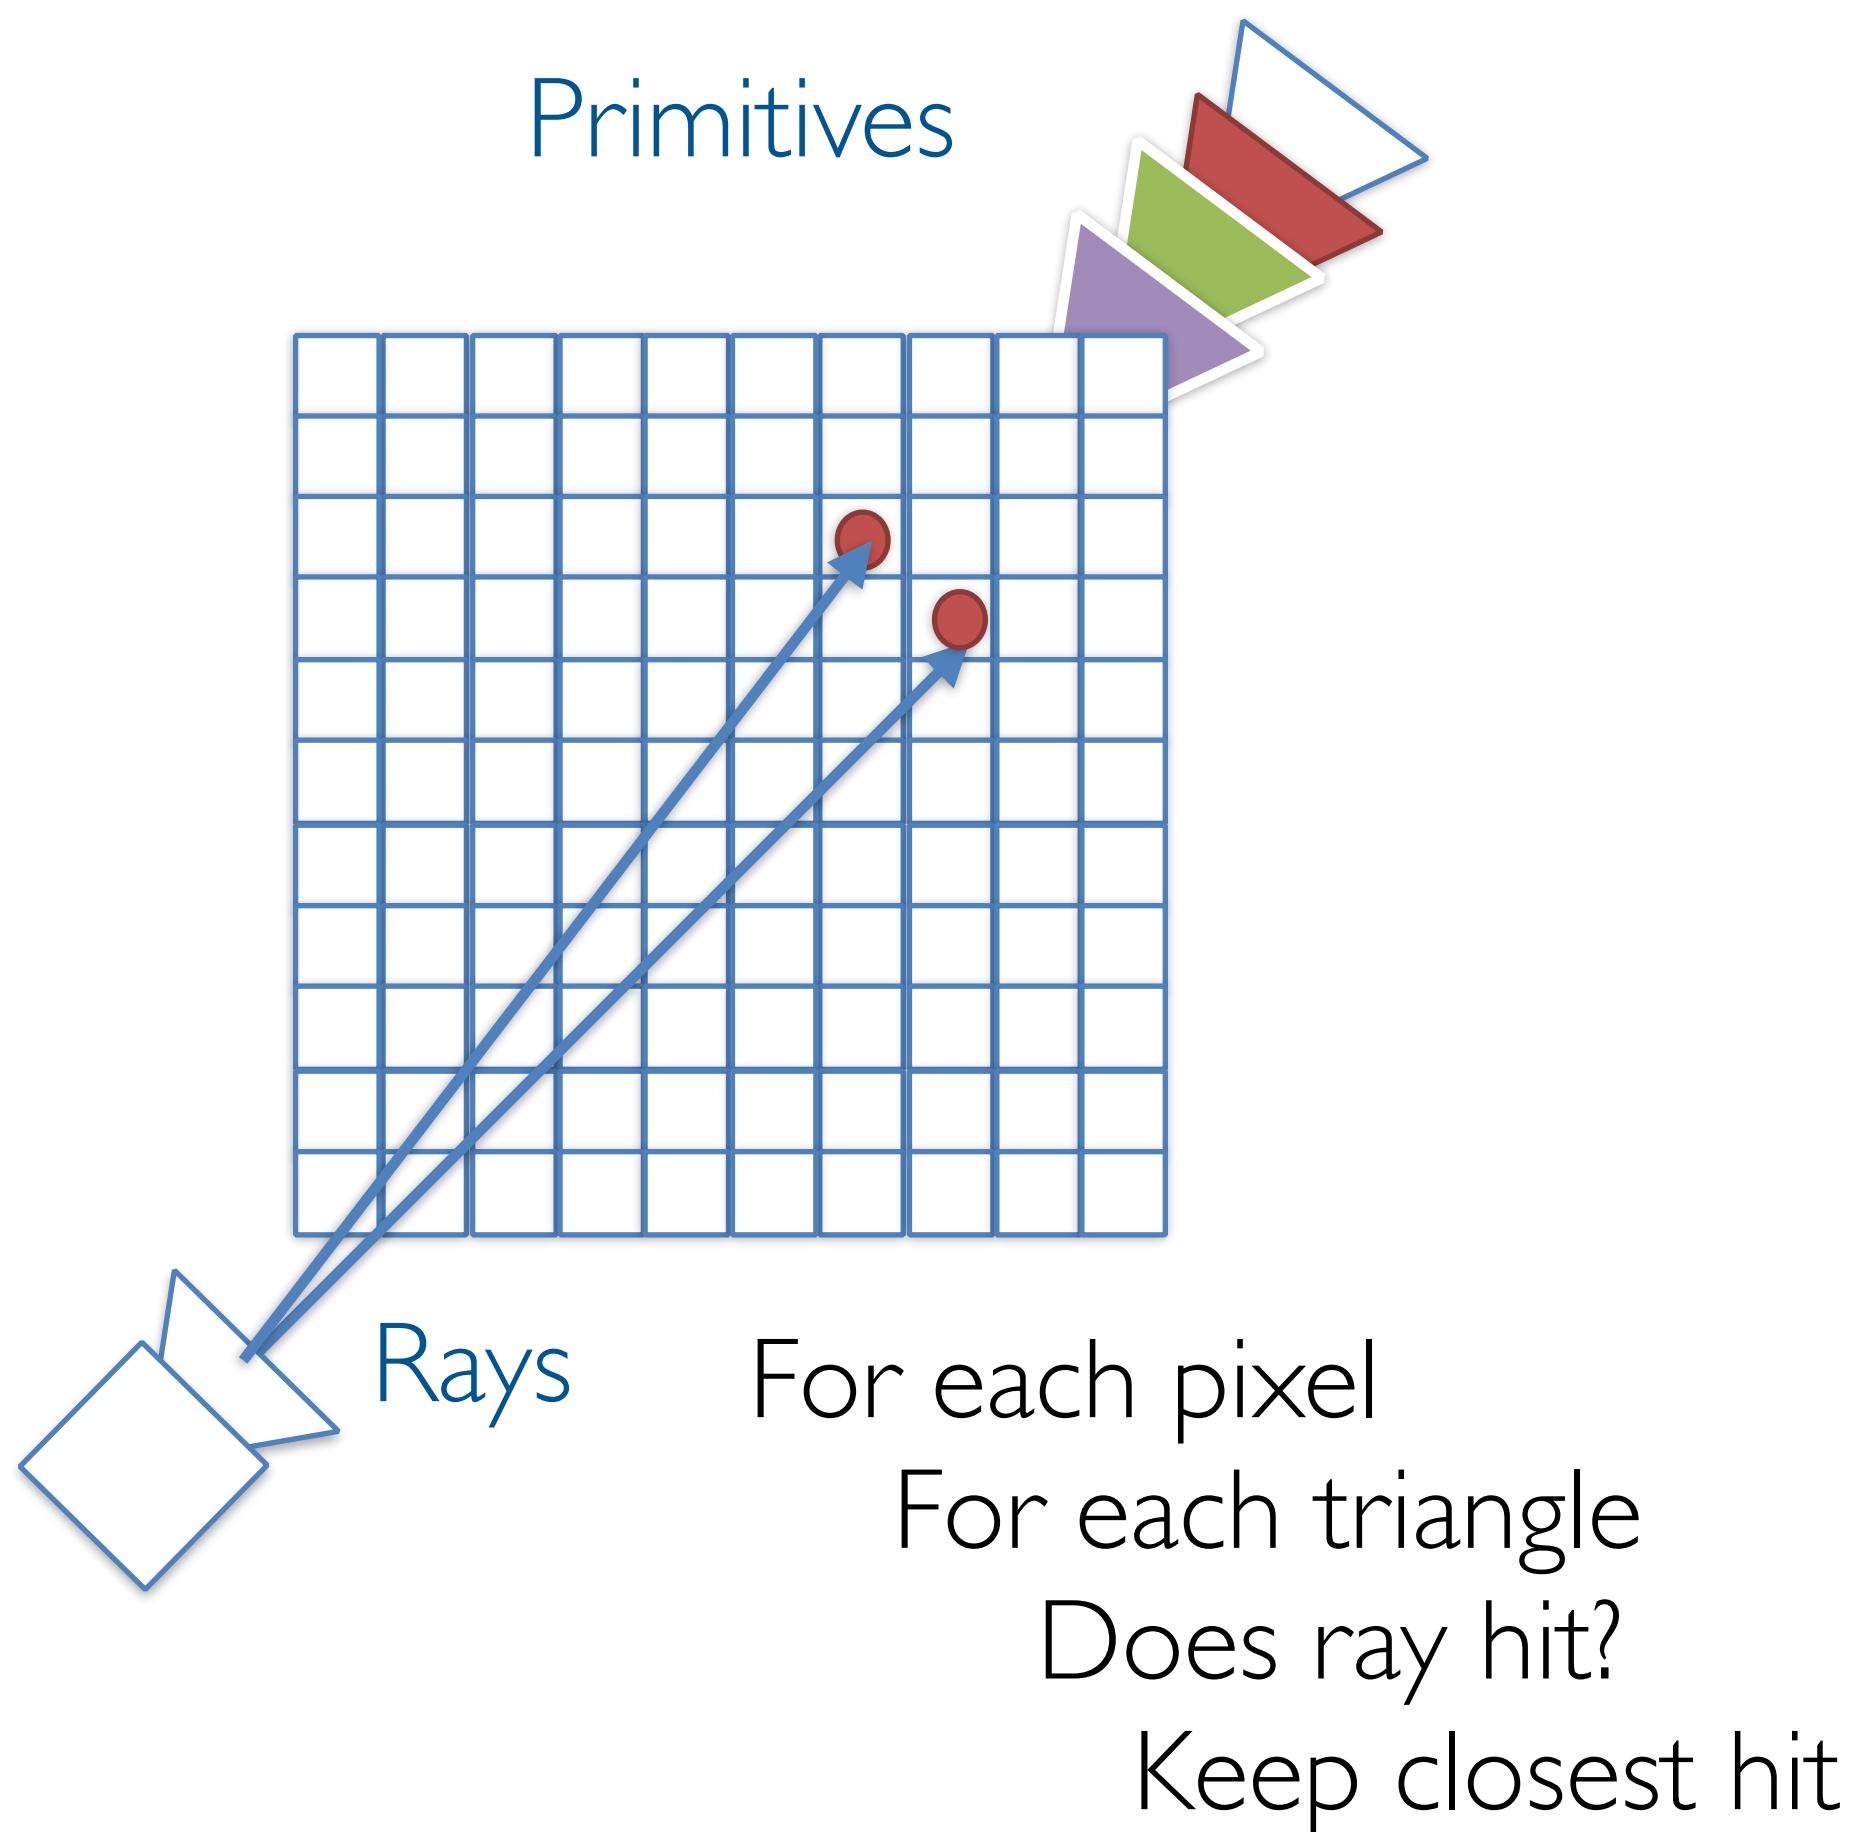


Fundamentals

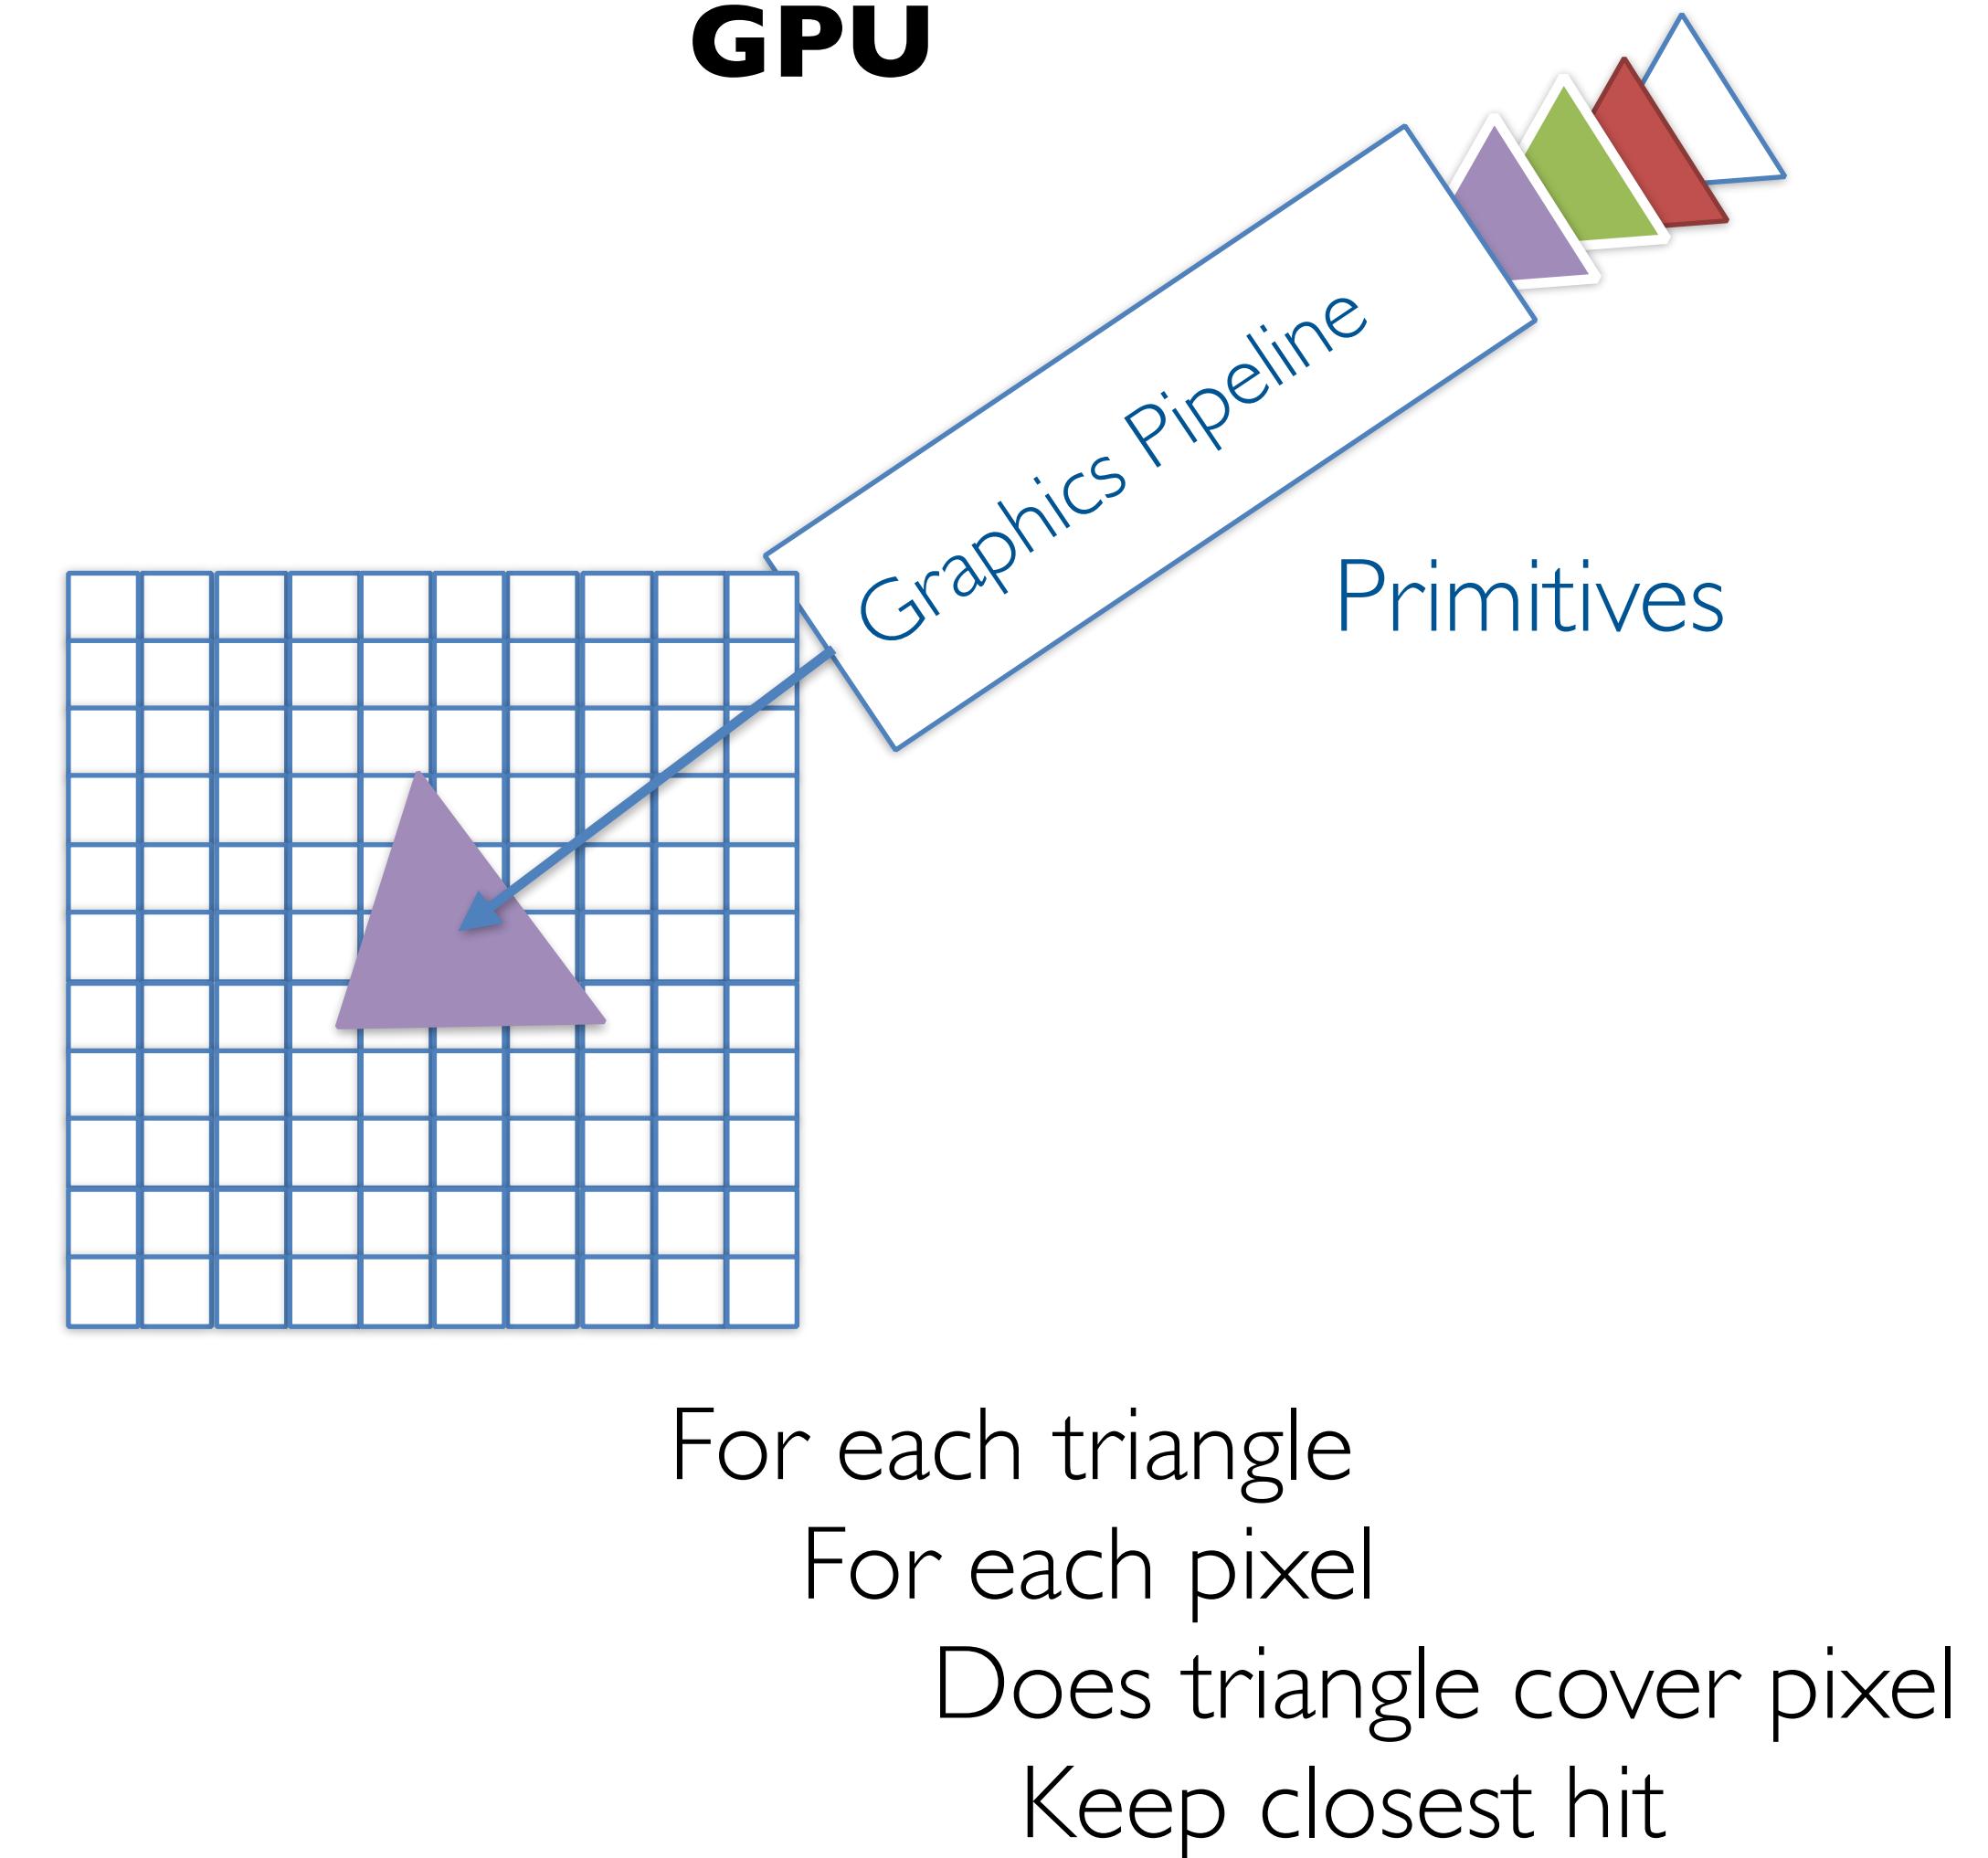


Raytracing vs Rasterised Graphics

Ray Tracing

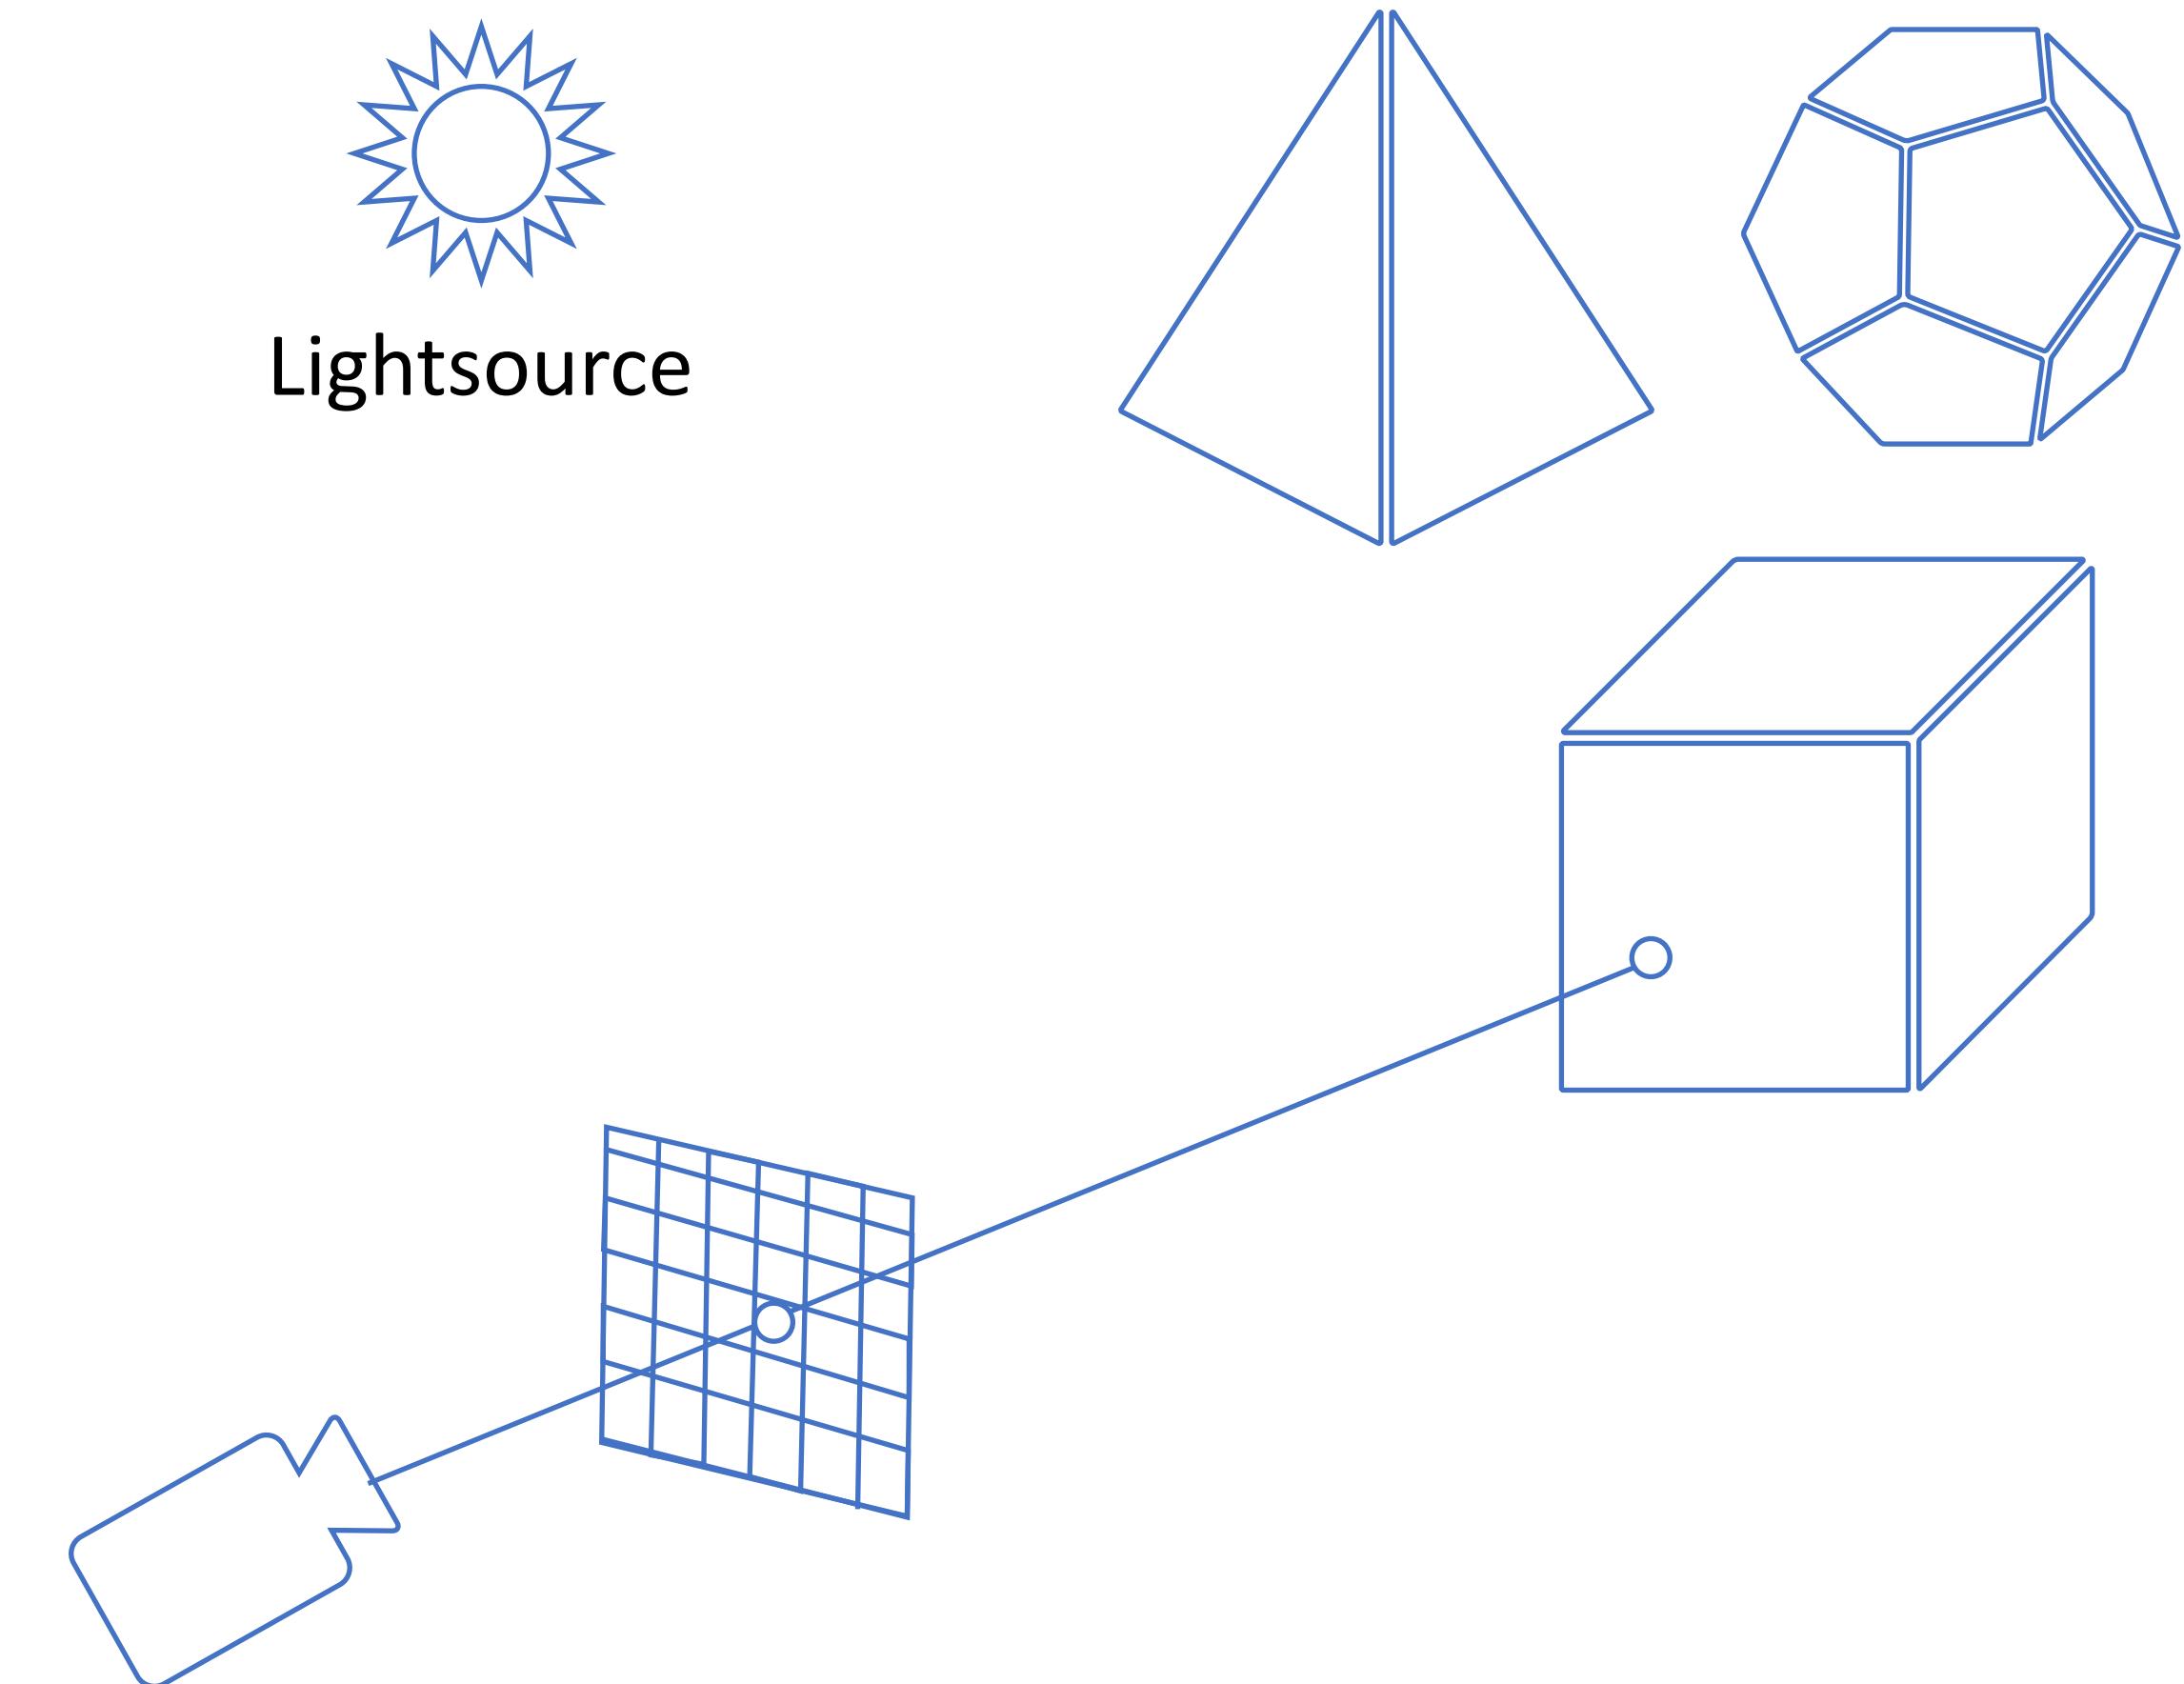


GPU



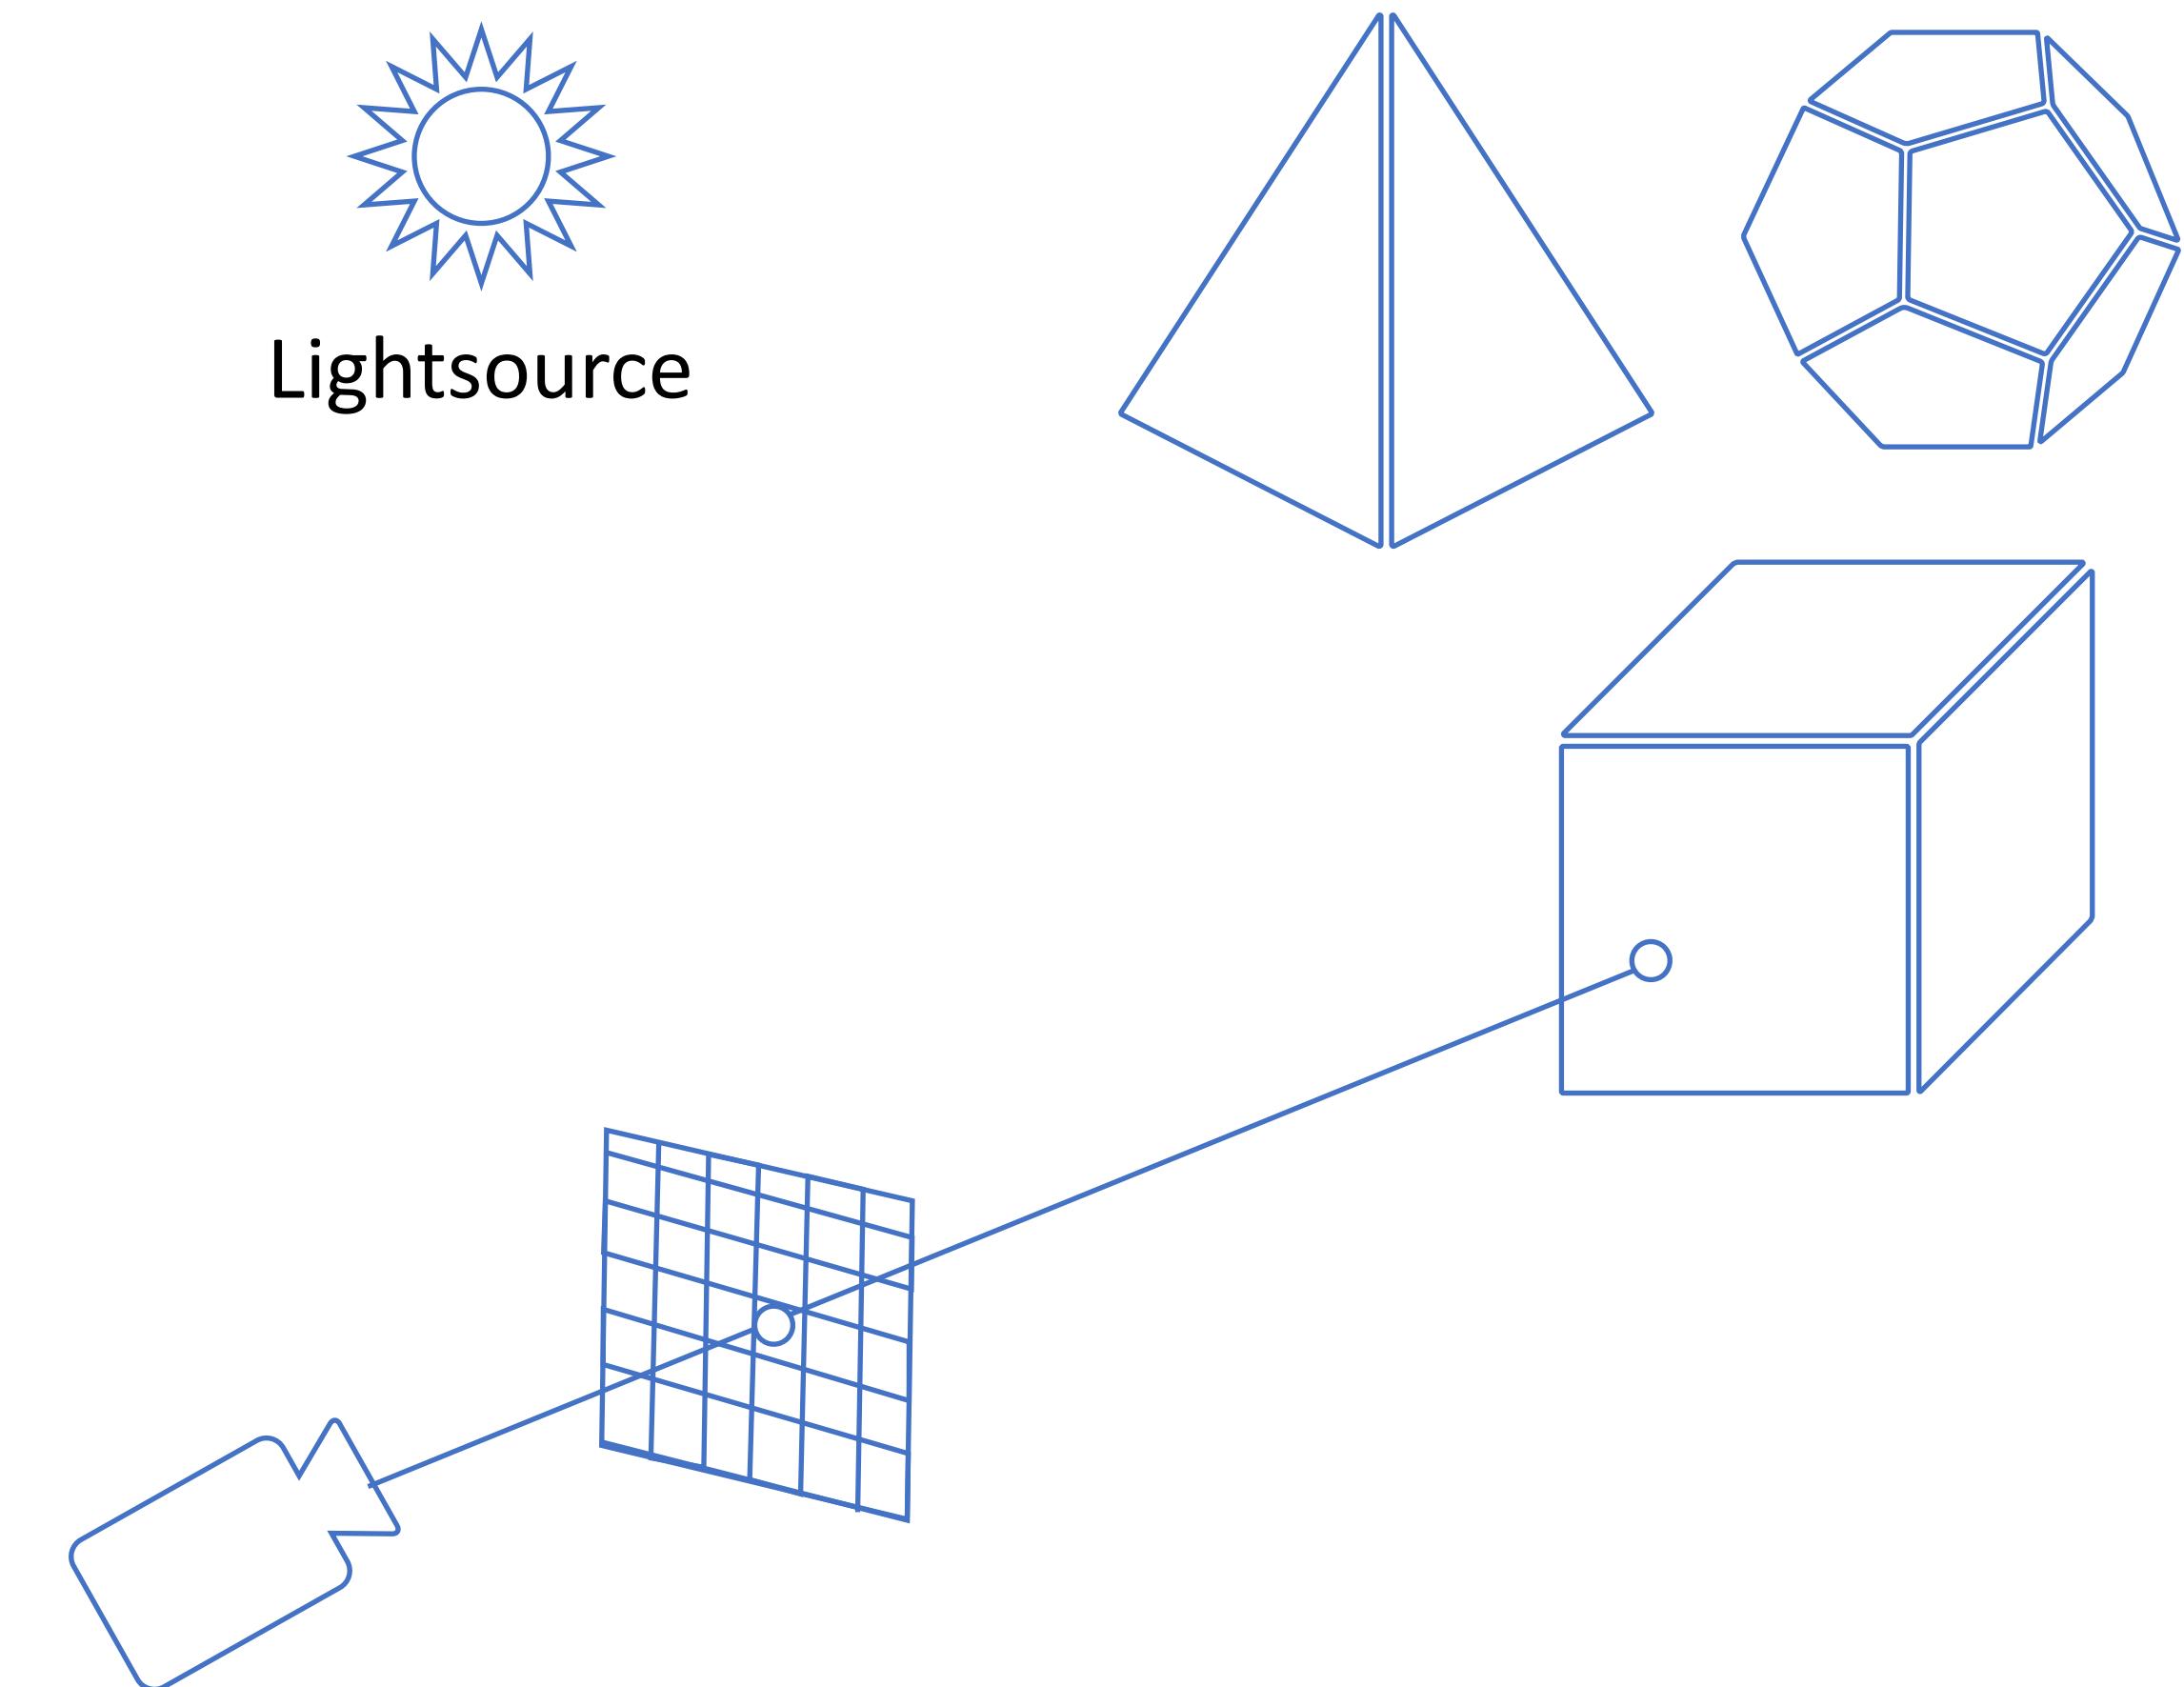
A Simple Raytracer

- Supports:
 - Spotlights and Shadows
 - Reflections
 - Refraction



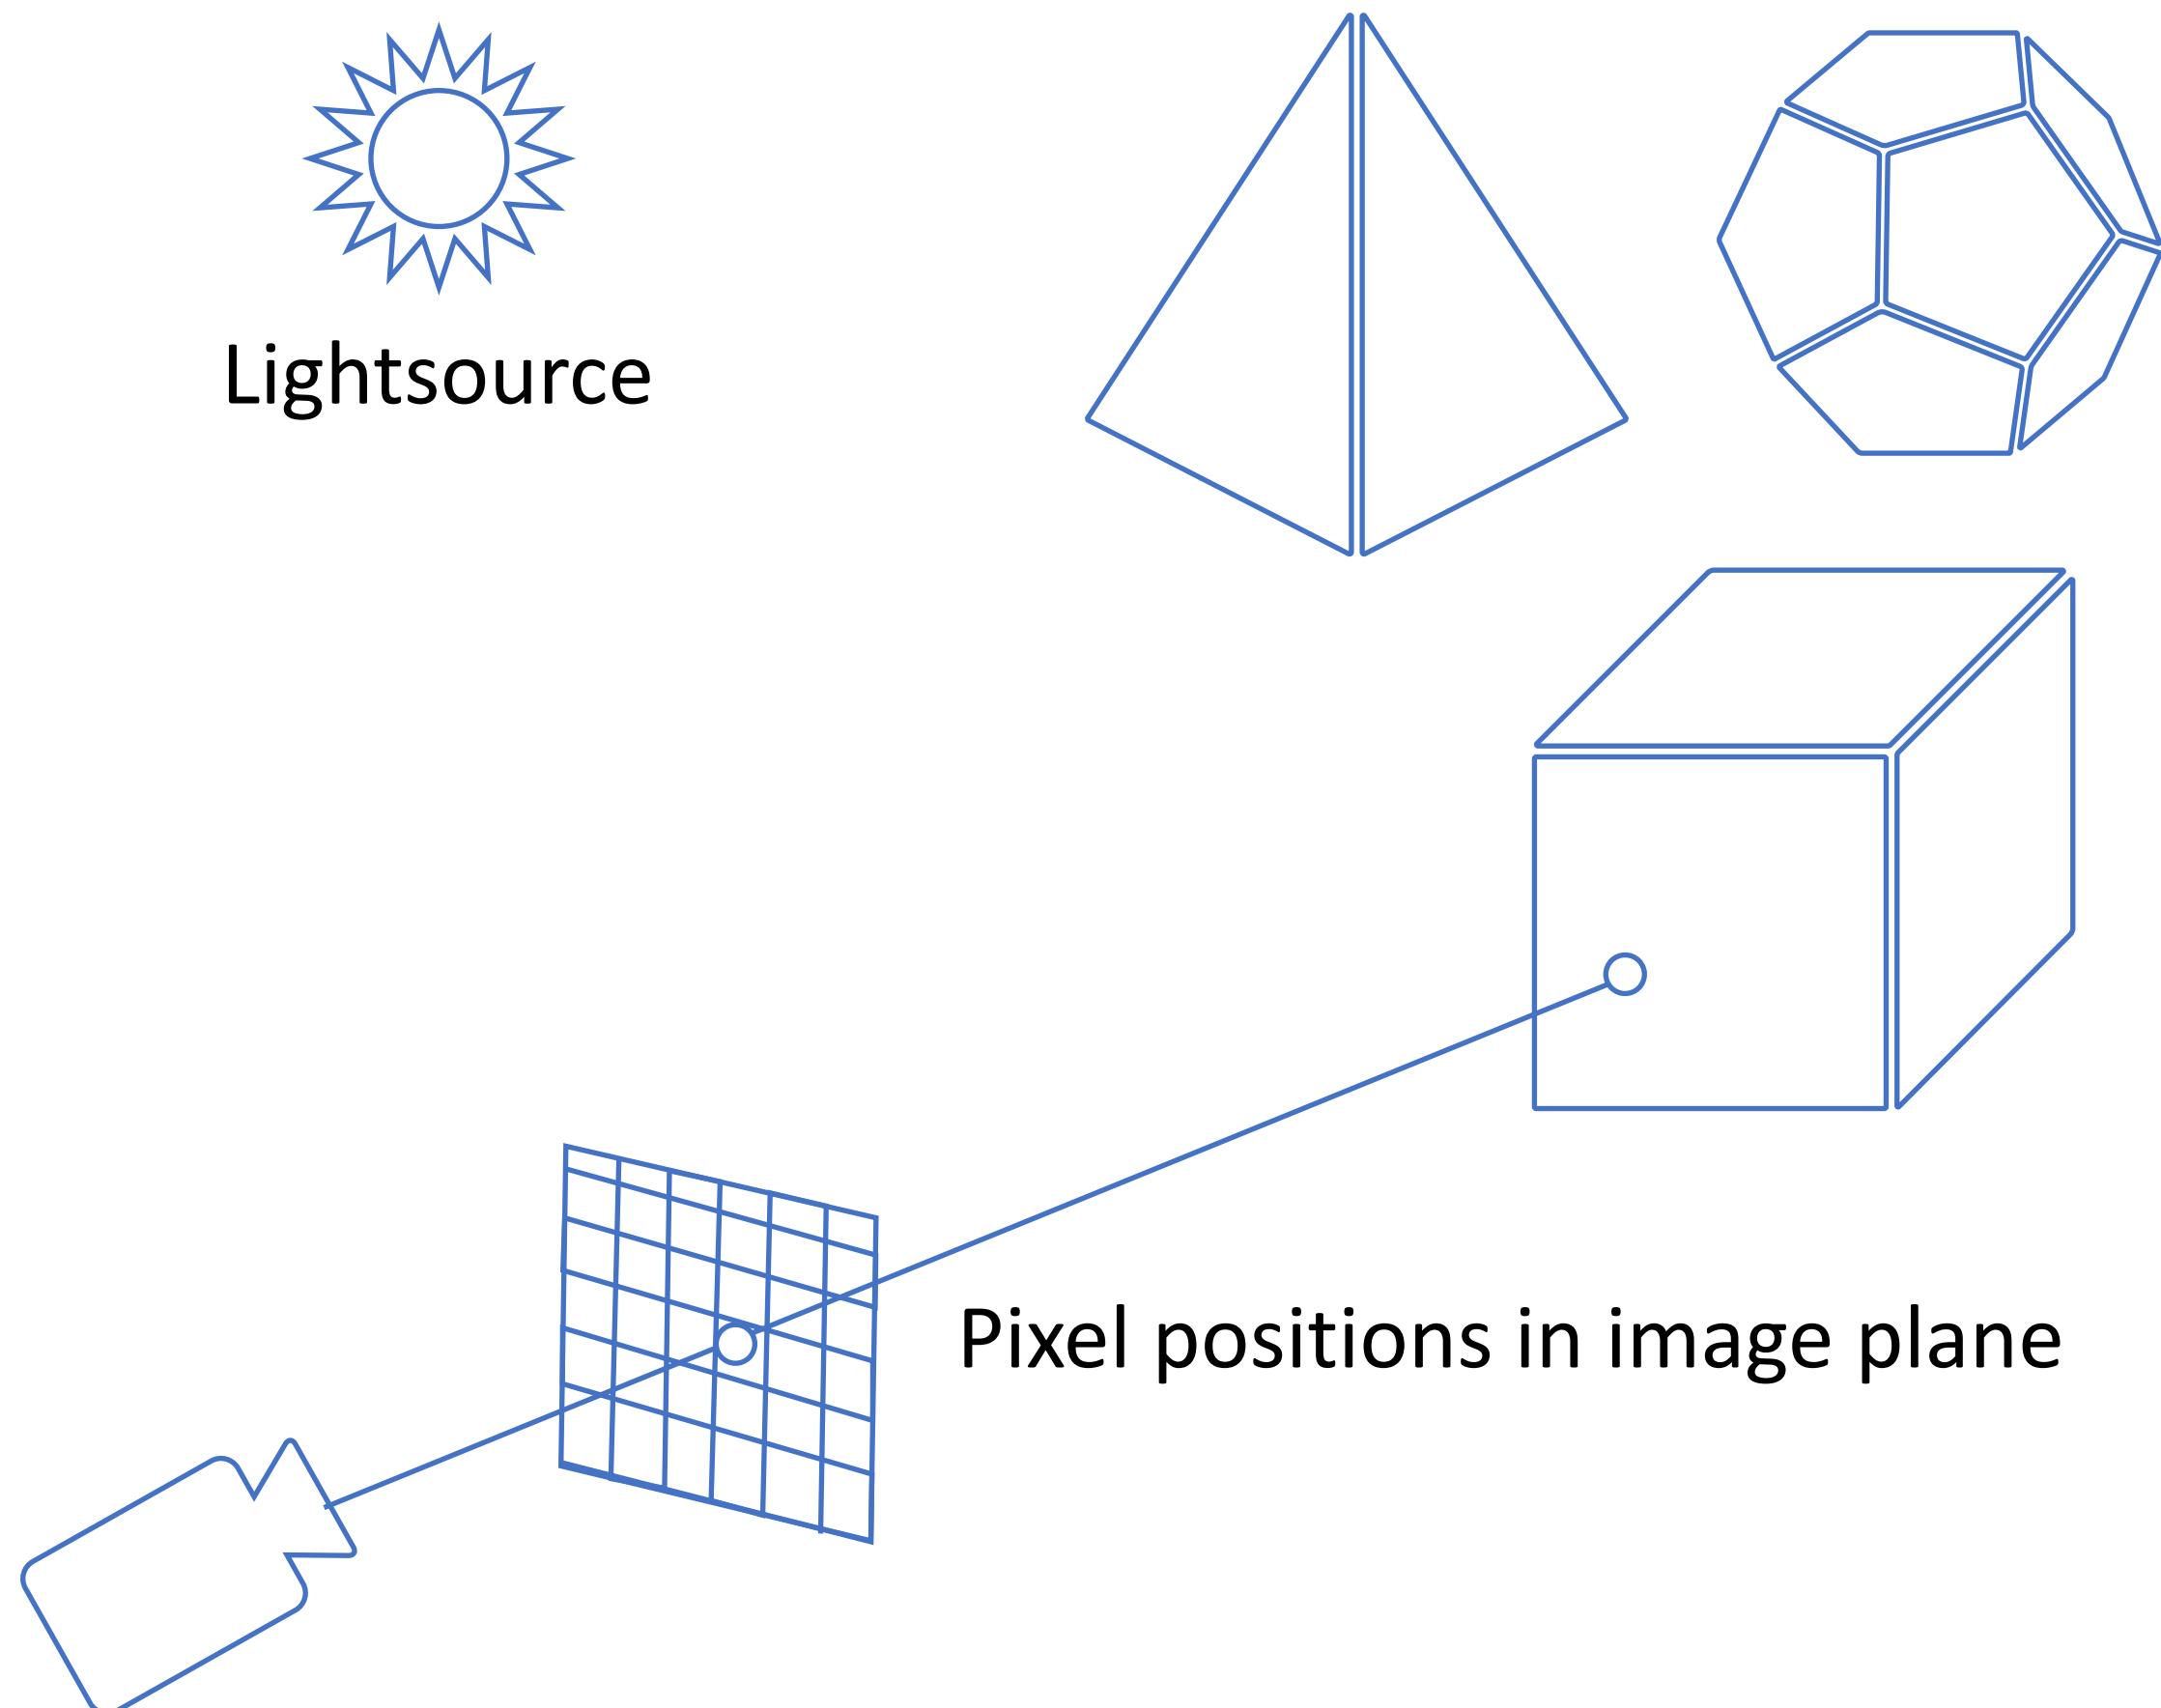
Local Illumination Models

- Compute outgoing radiance at an intersection
- Classic models:
 - Lambertian diffuse
 - Phong / Blinn–Phong specular components
- Requires shadow rays to determine direct-light visibility



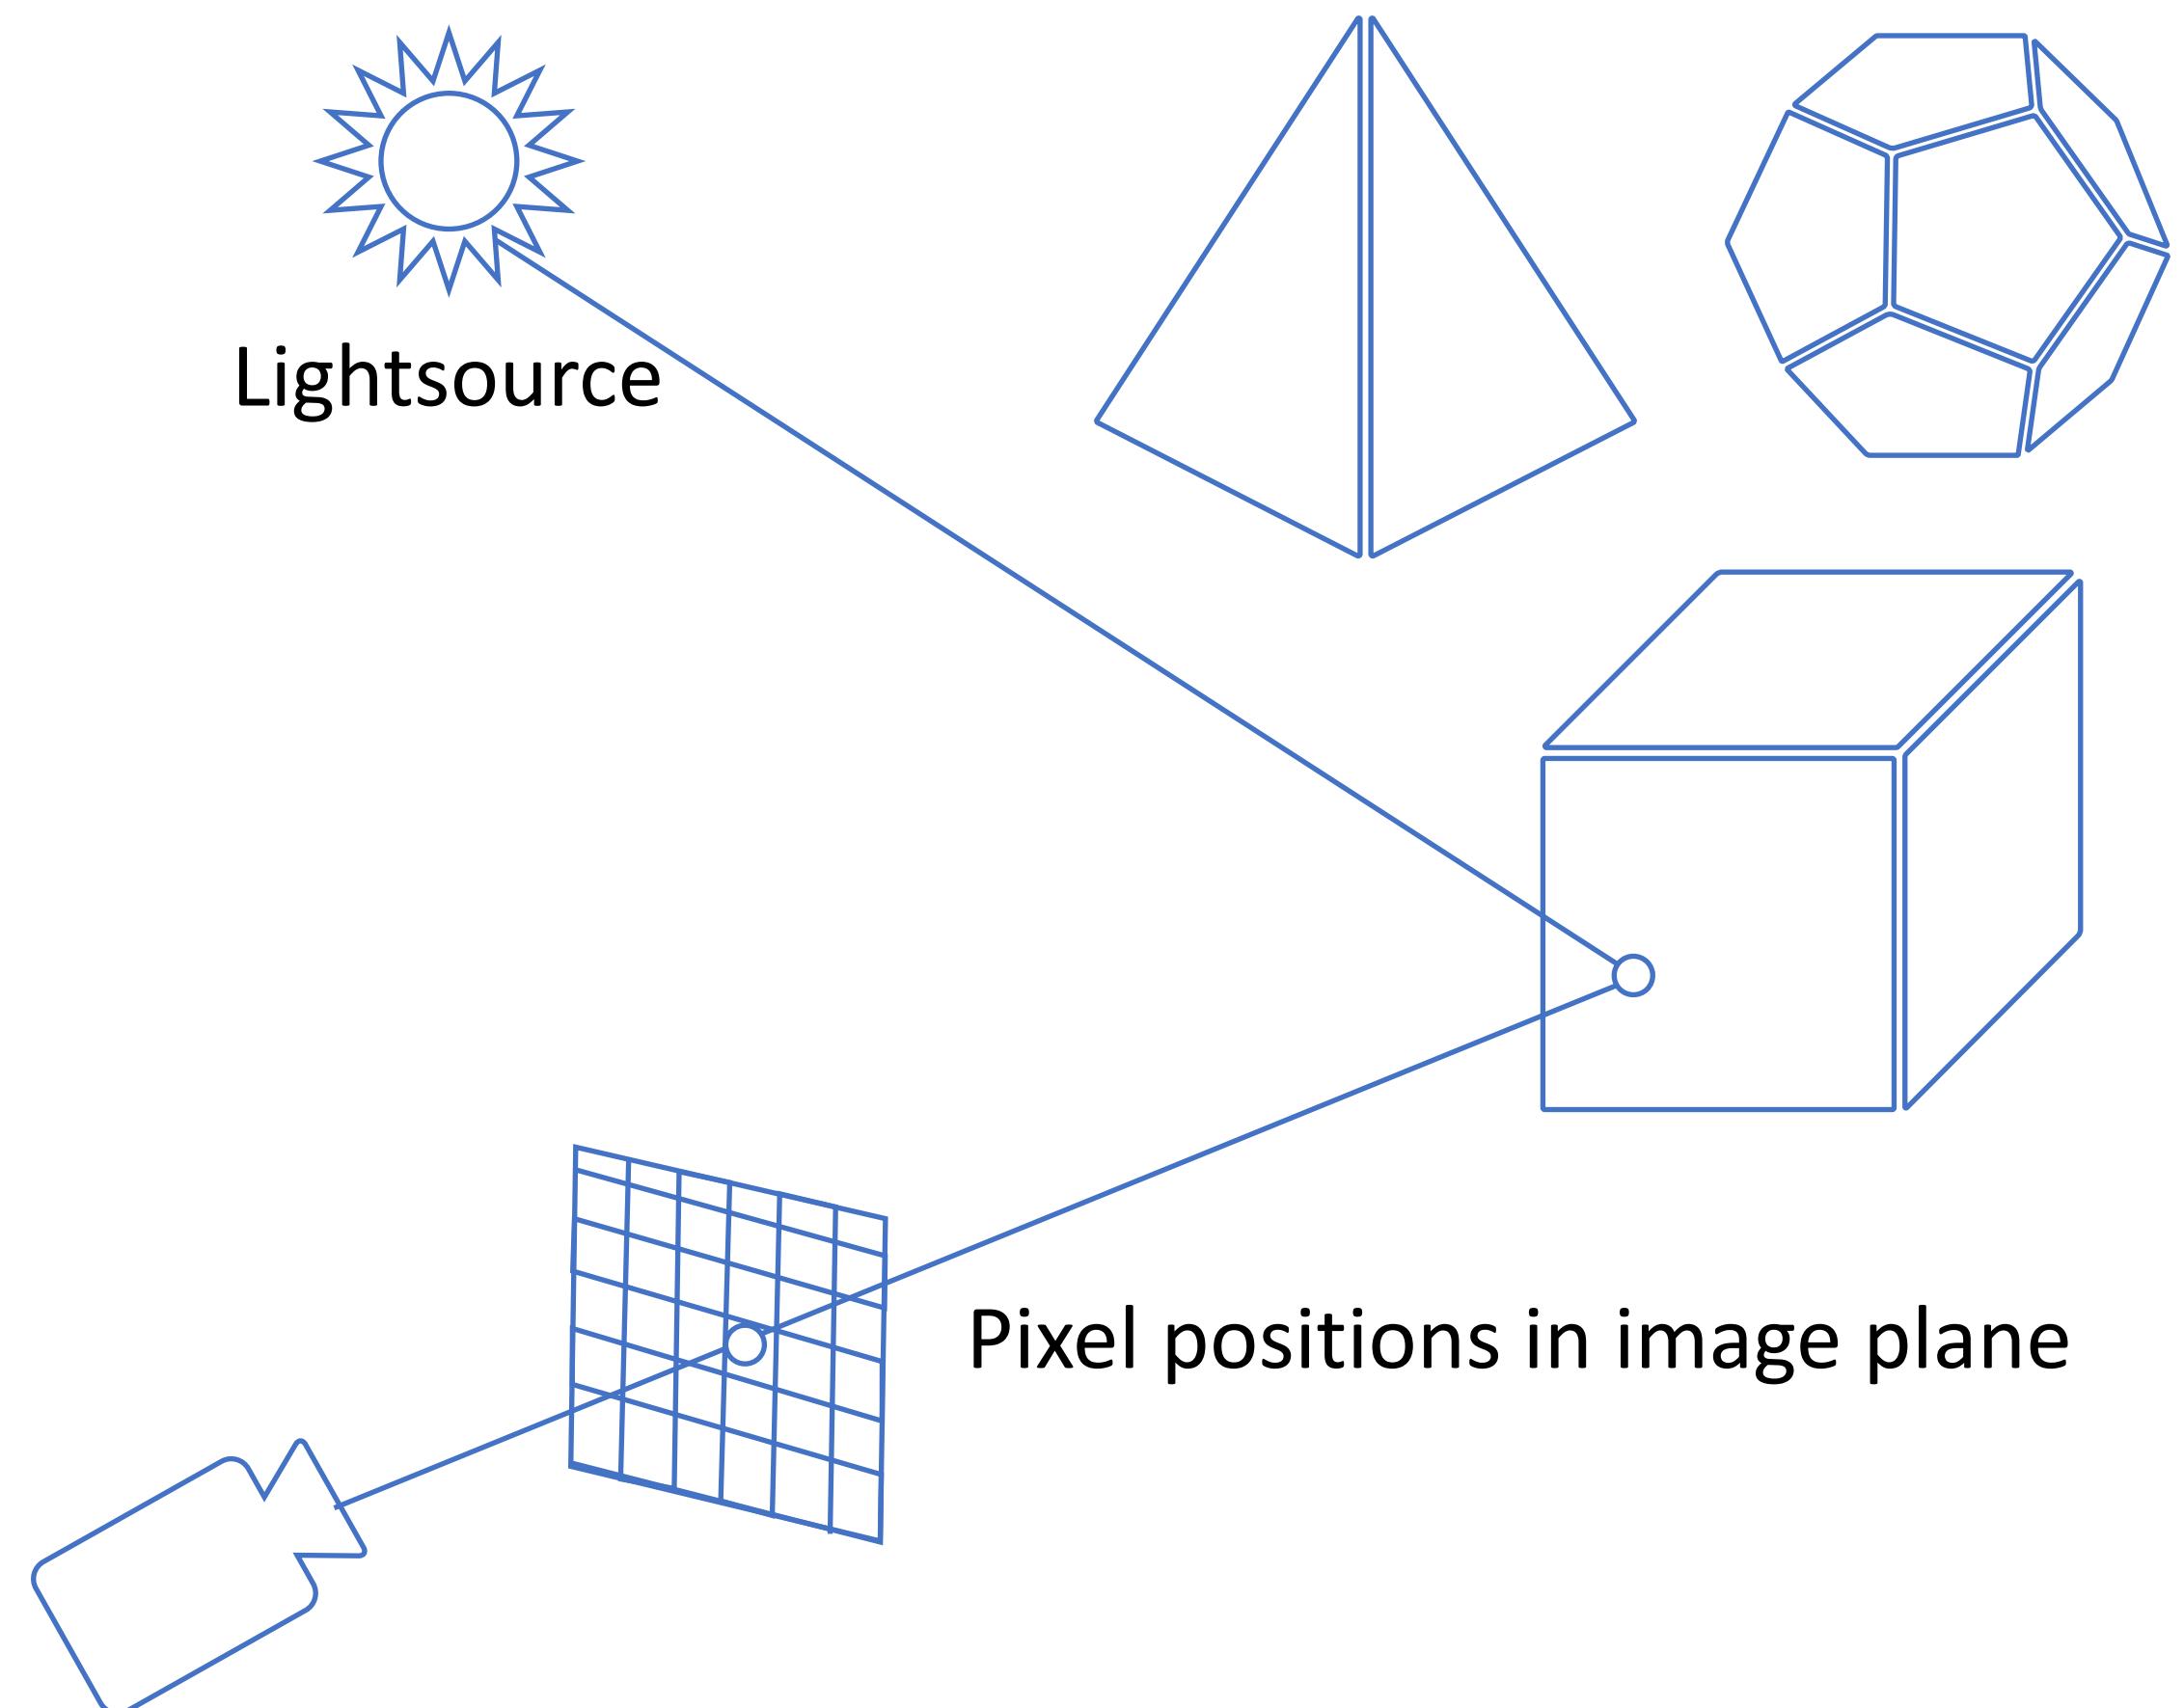
Raytracing Process

- Create an image plane and viewpoint
- For each pixel trace a ‘ray’ from the eye through a corresponding point in the image plane
- For each ray, return the colour of the object at the hit point (closest to the camera)



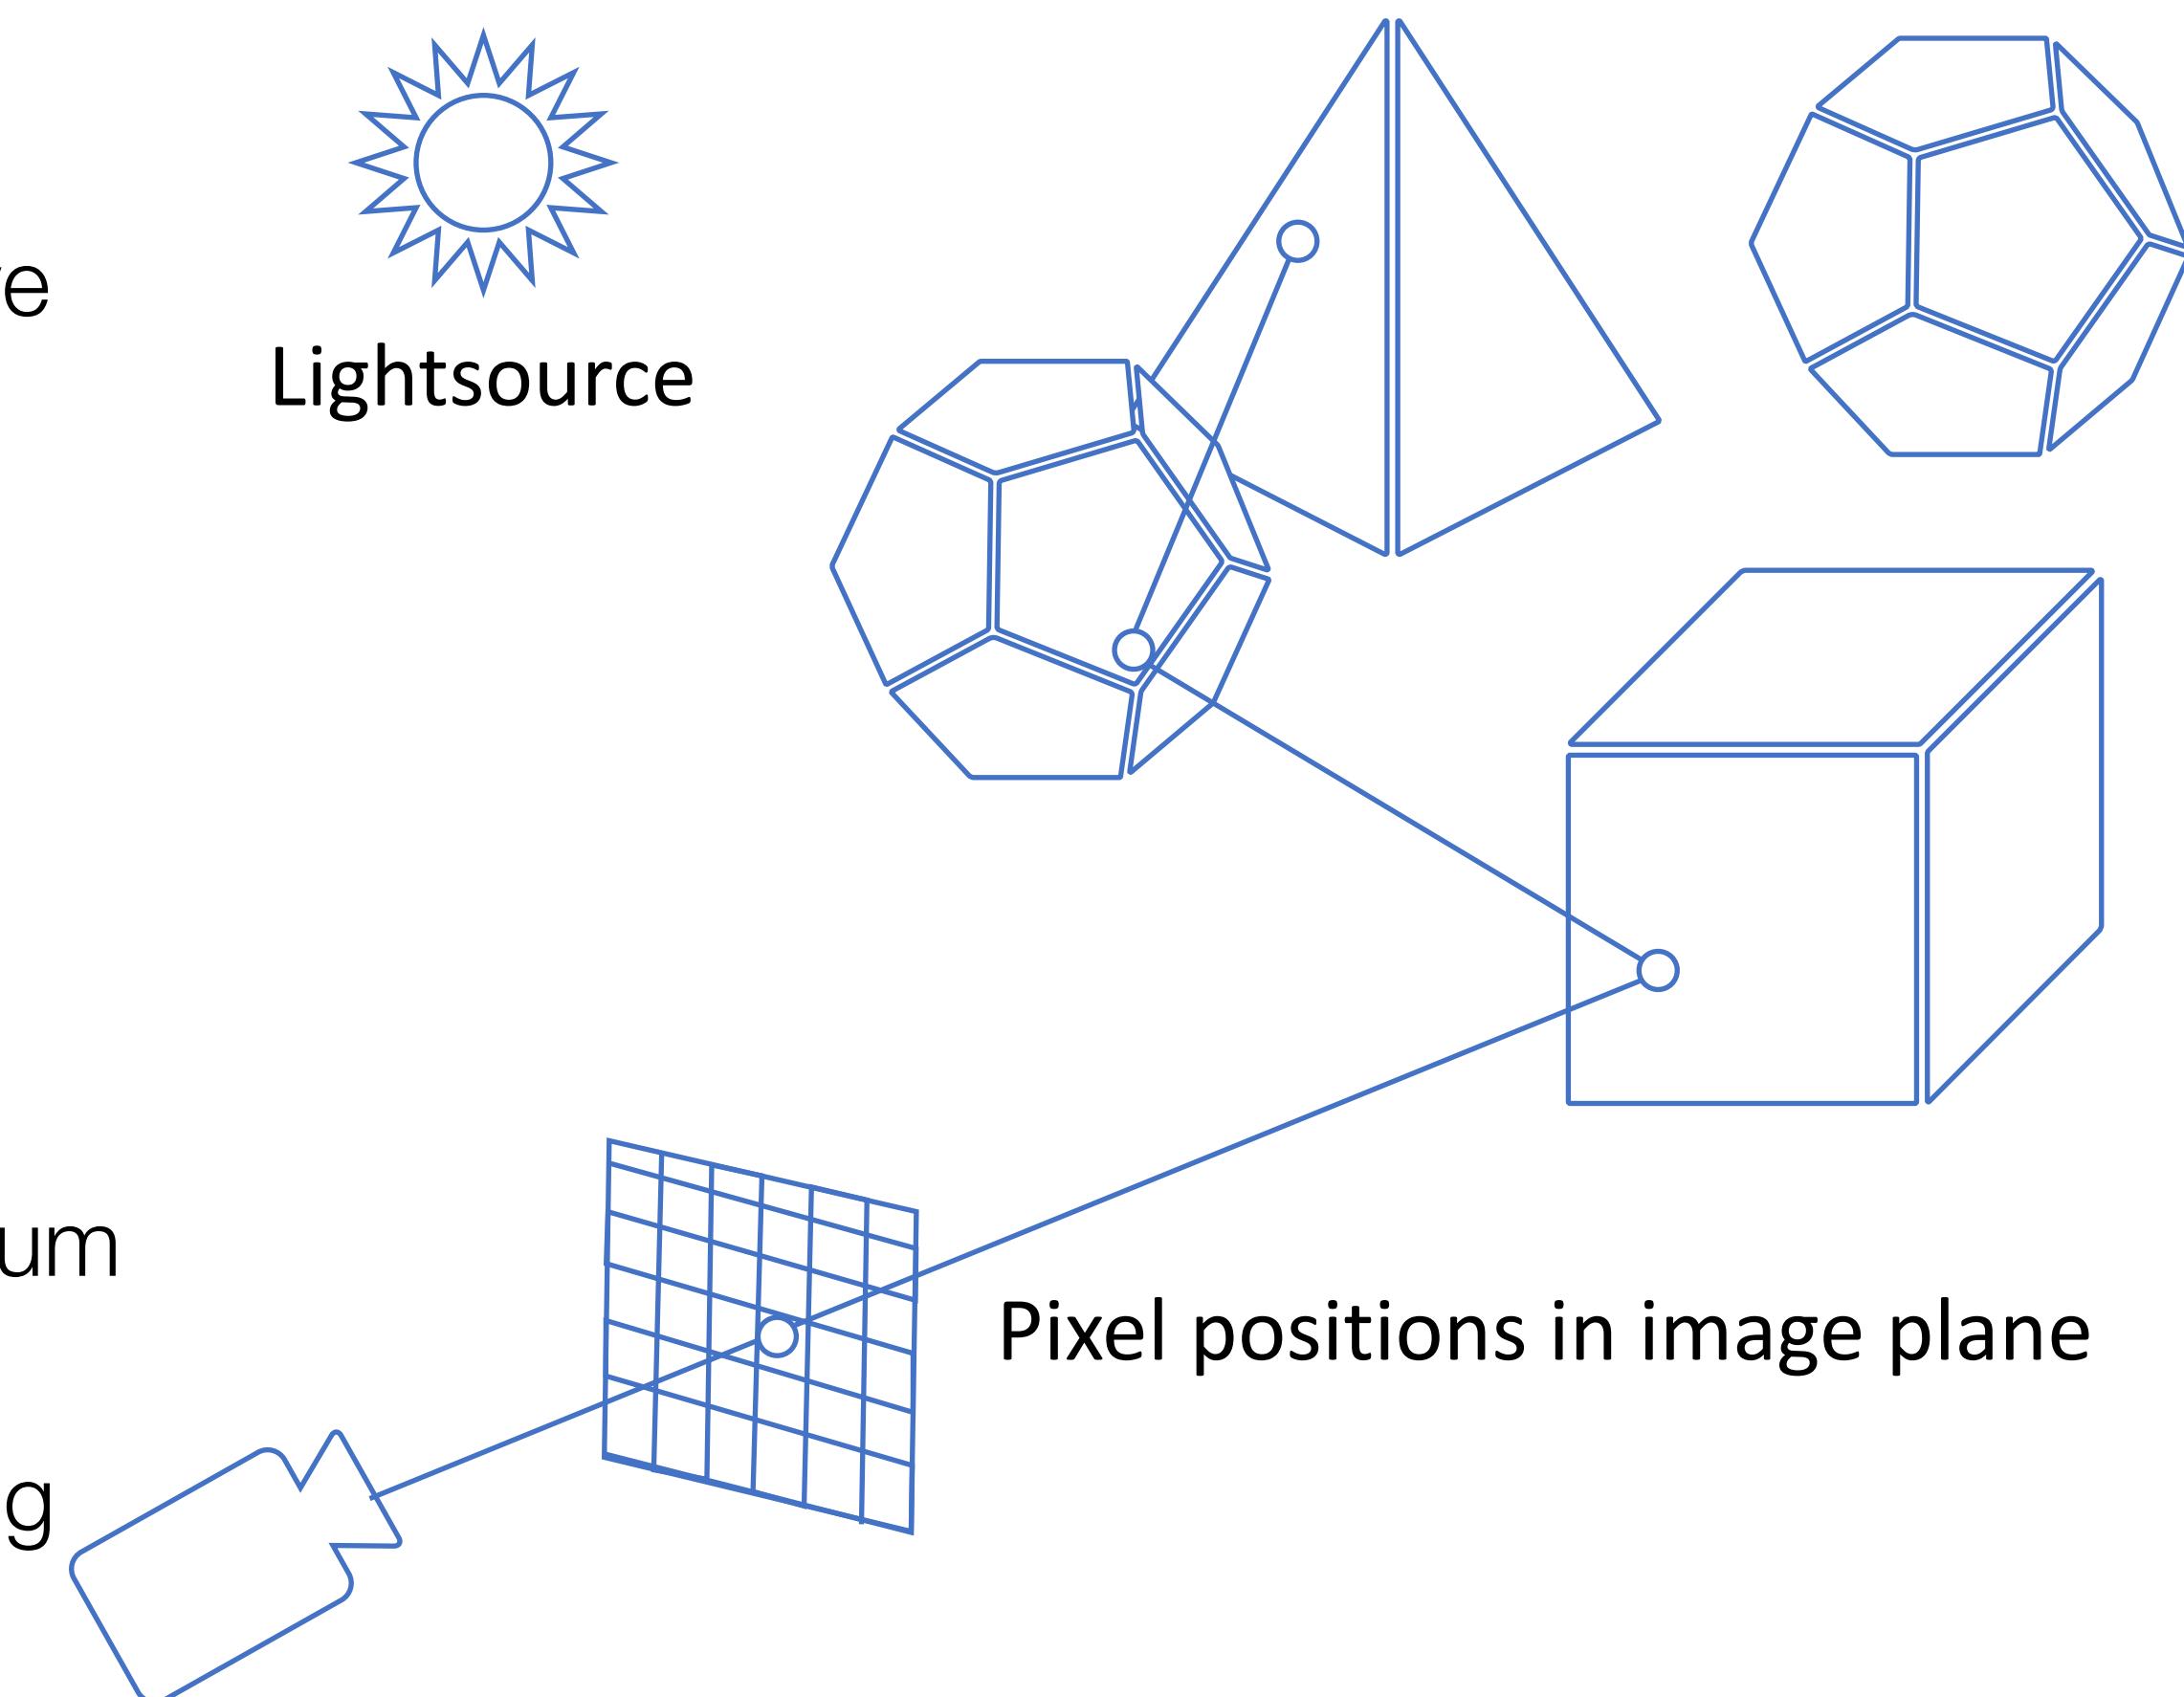
Raytracing Process - Shadows

- Shadows:
 - Ray is casted towards light source
 - If an object is located between hit point (surface point) and light source then the hit point is in shadow
 - If not, there is no shadow
 - Point lights: hard binary shadows
 - Area lights: sample positions on the emitter to produce soft penumbras



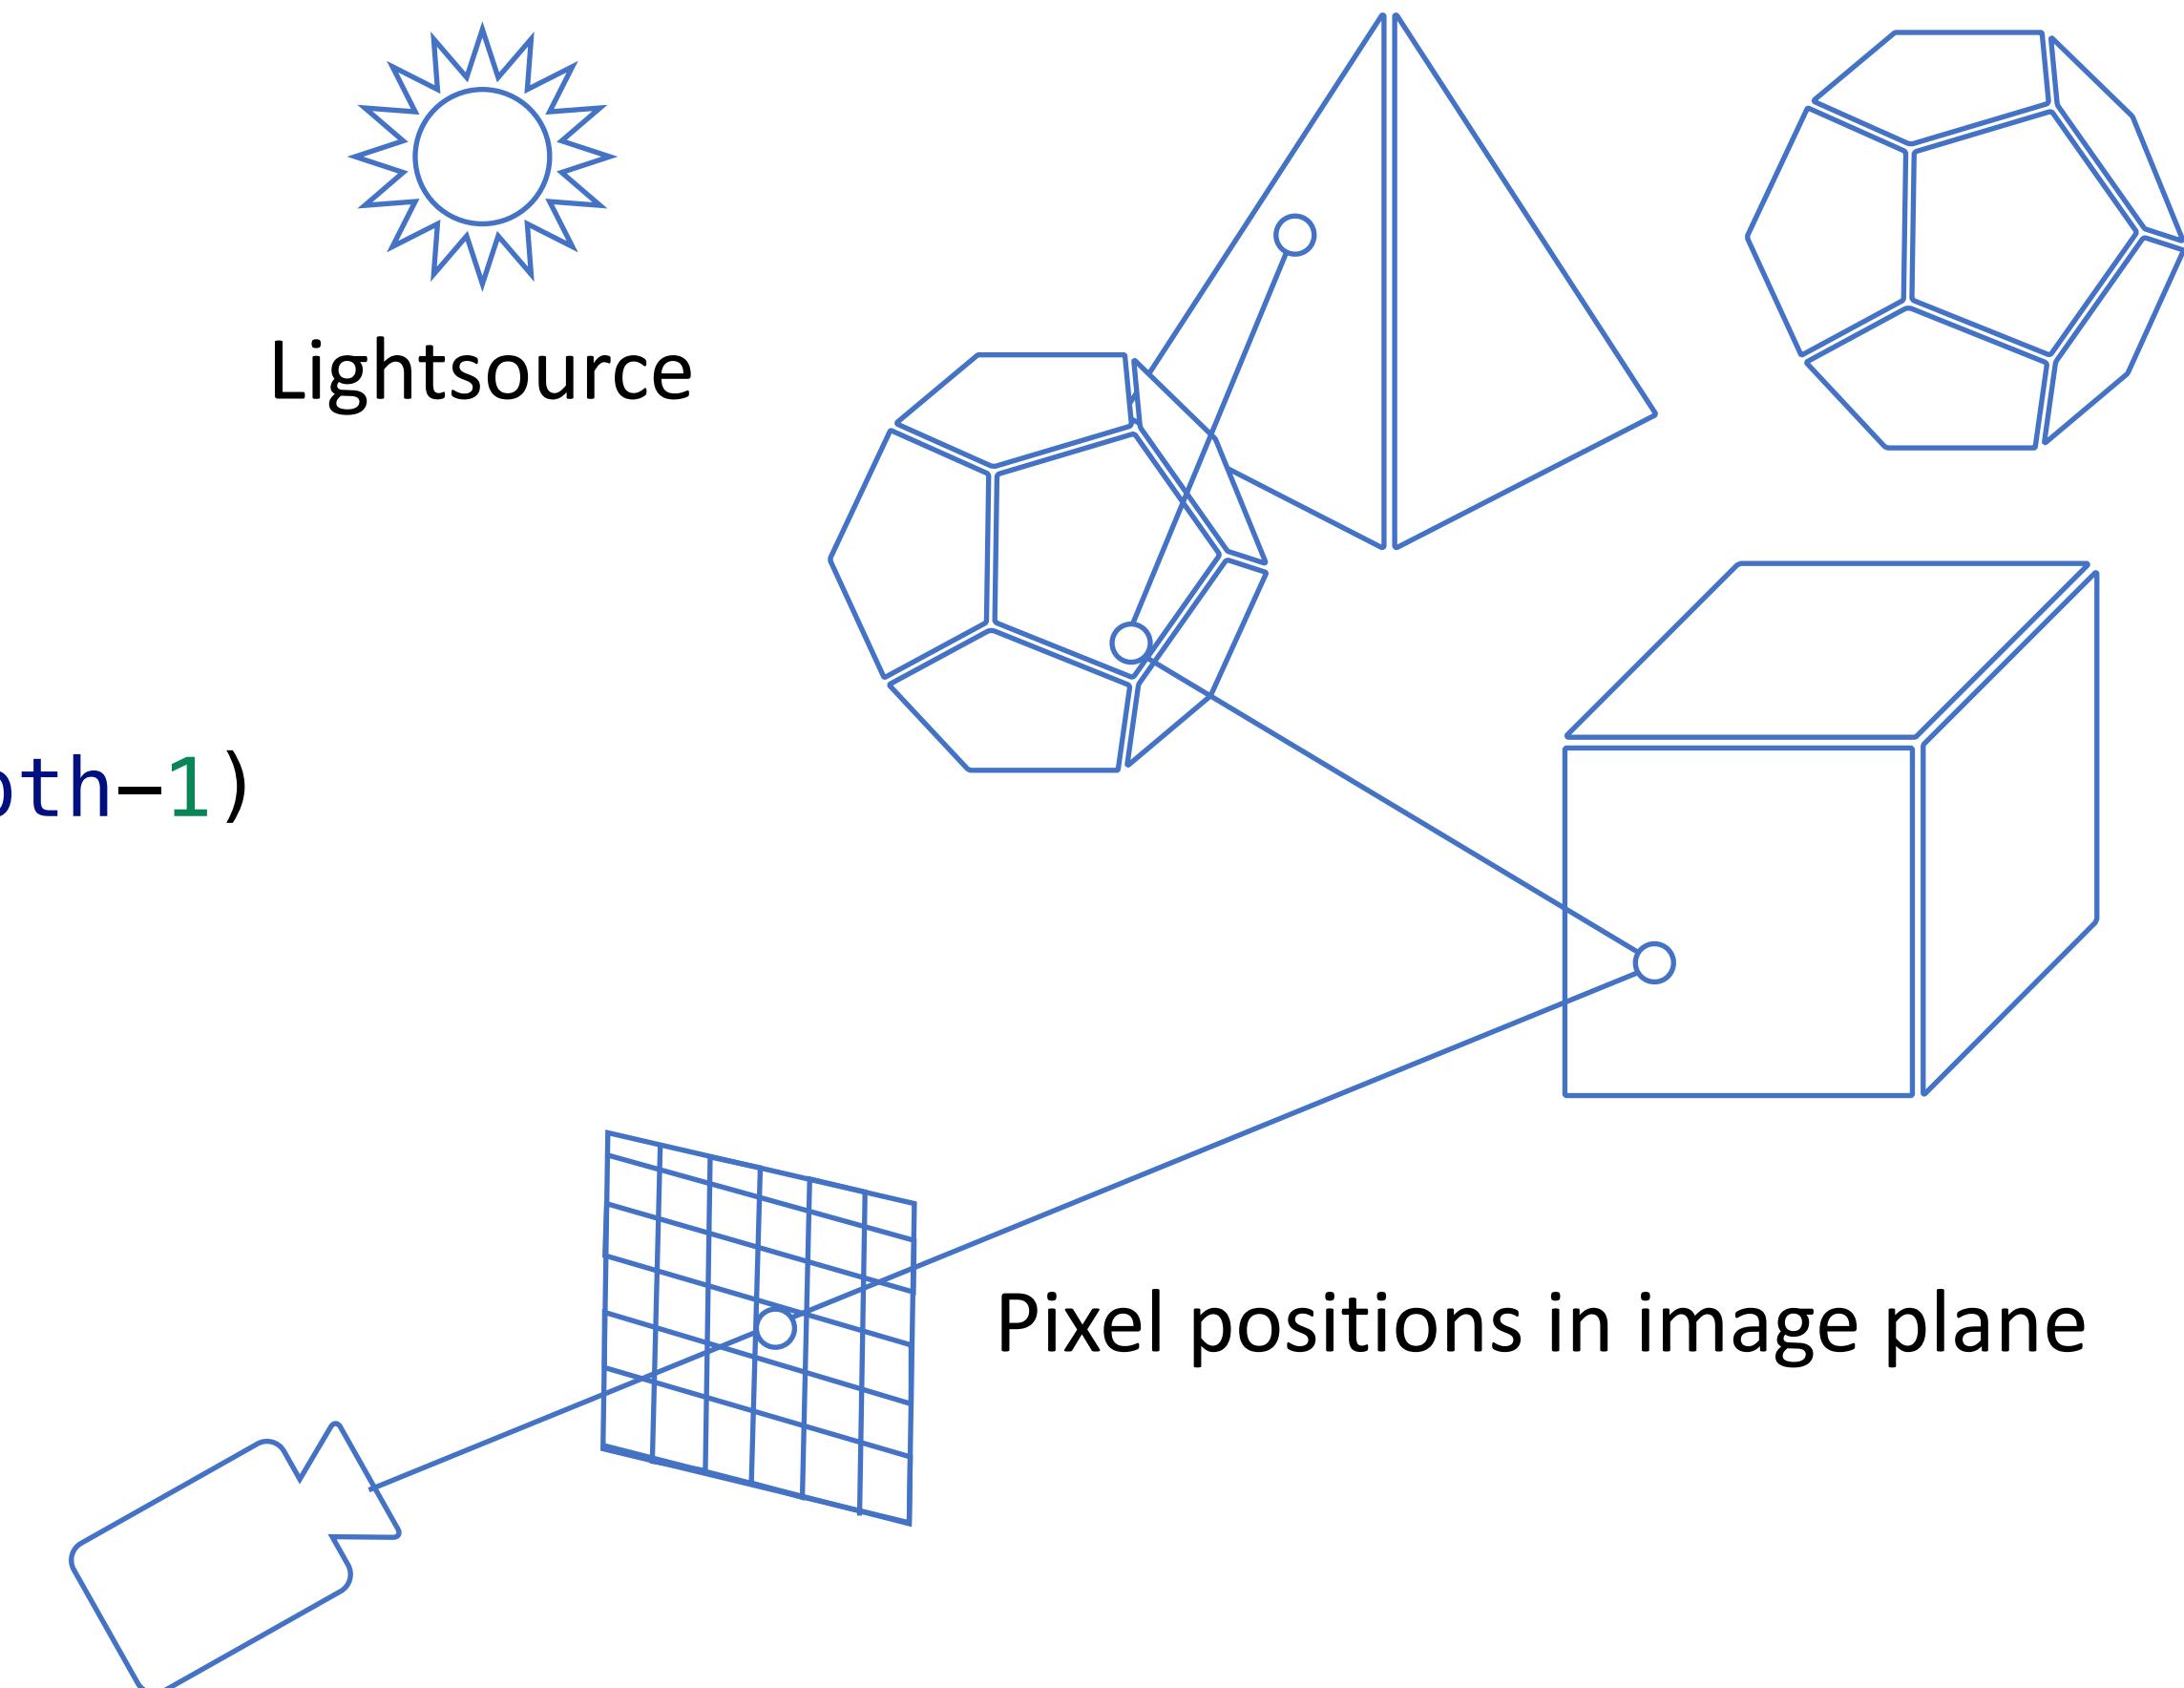
Raytracing Process - Reflection

- Method:
 - If ray is reflected on the surface (reflective material), a new (secondary) ray is created and cast in the reflective direction
 - Compute reflected direction using the law of reflection
$$\mathbf{r} = \mathbf{v} - 2(\mathbf{v} \cdot \mathbf{n})\mathbf{n}$$
 - Spawn secondary ray from hit point
 - Modulate the reflected contribution by the material's reflectivity ρ_r
 - Trace the reflected ray recursively, respecting a maximum depth or probabilistic termination
 - Finally: Add returned reflected radiance to local shading
$$L = L_{\text{local}} + \rho_r L_{\text{reflected}}$$



Raytracing Process - Reflection

```
colour render(ray, scene, fraction, depth)
hit = intersect(ray, scene)
if (depth == 0 || hit.kr * fraction < TINY)
    thisColour = hit.colour
else
    thisColour = (1 - hit.kr) * hit.colour +
        hit.kr * render(reflection, scene,
                          hit.kr*fraction, depth-1)
return fraction * thisColour
```



*hit.kr stands for “coefficient of reflection” (sometimes called reflectivity).

It is a value between 0 and 1:

0: the surface is non-reflective (all light is absorbed or diffused)

1: the surface is a perfect mirror (all light is reflected)

Values in between → partial reflection (some light is reflected, some is diffuse).

Raytracing Process - Refraction

- Transparent objects
- Light bends at the surface
- Light ‘bends’ between materials
- Material has a refractive index
- Ratio of refractive indices determines the bending (Snell’s law): $n_1 \sin \theta_1 = n_2 \sin \theta_2$
 - Vacum: $\eta = 1$
 - Air: $\eta = 1.0003$
 - Water : $\eta = 1.333$
- Refracted ray:

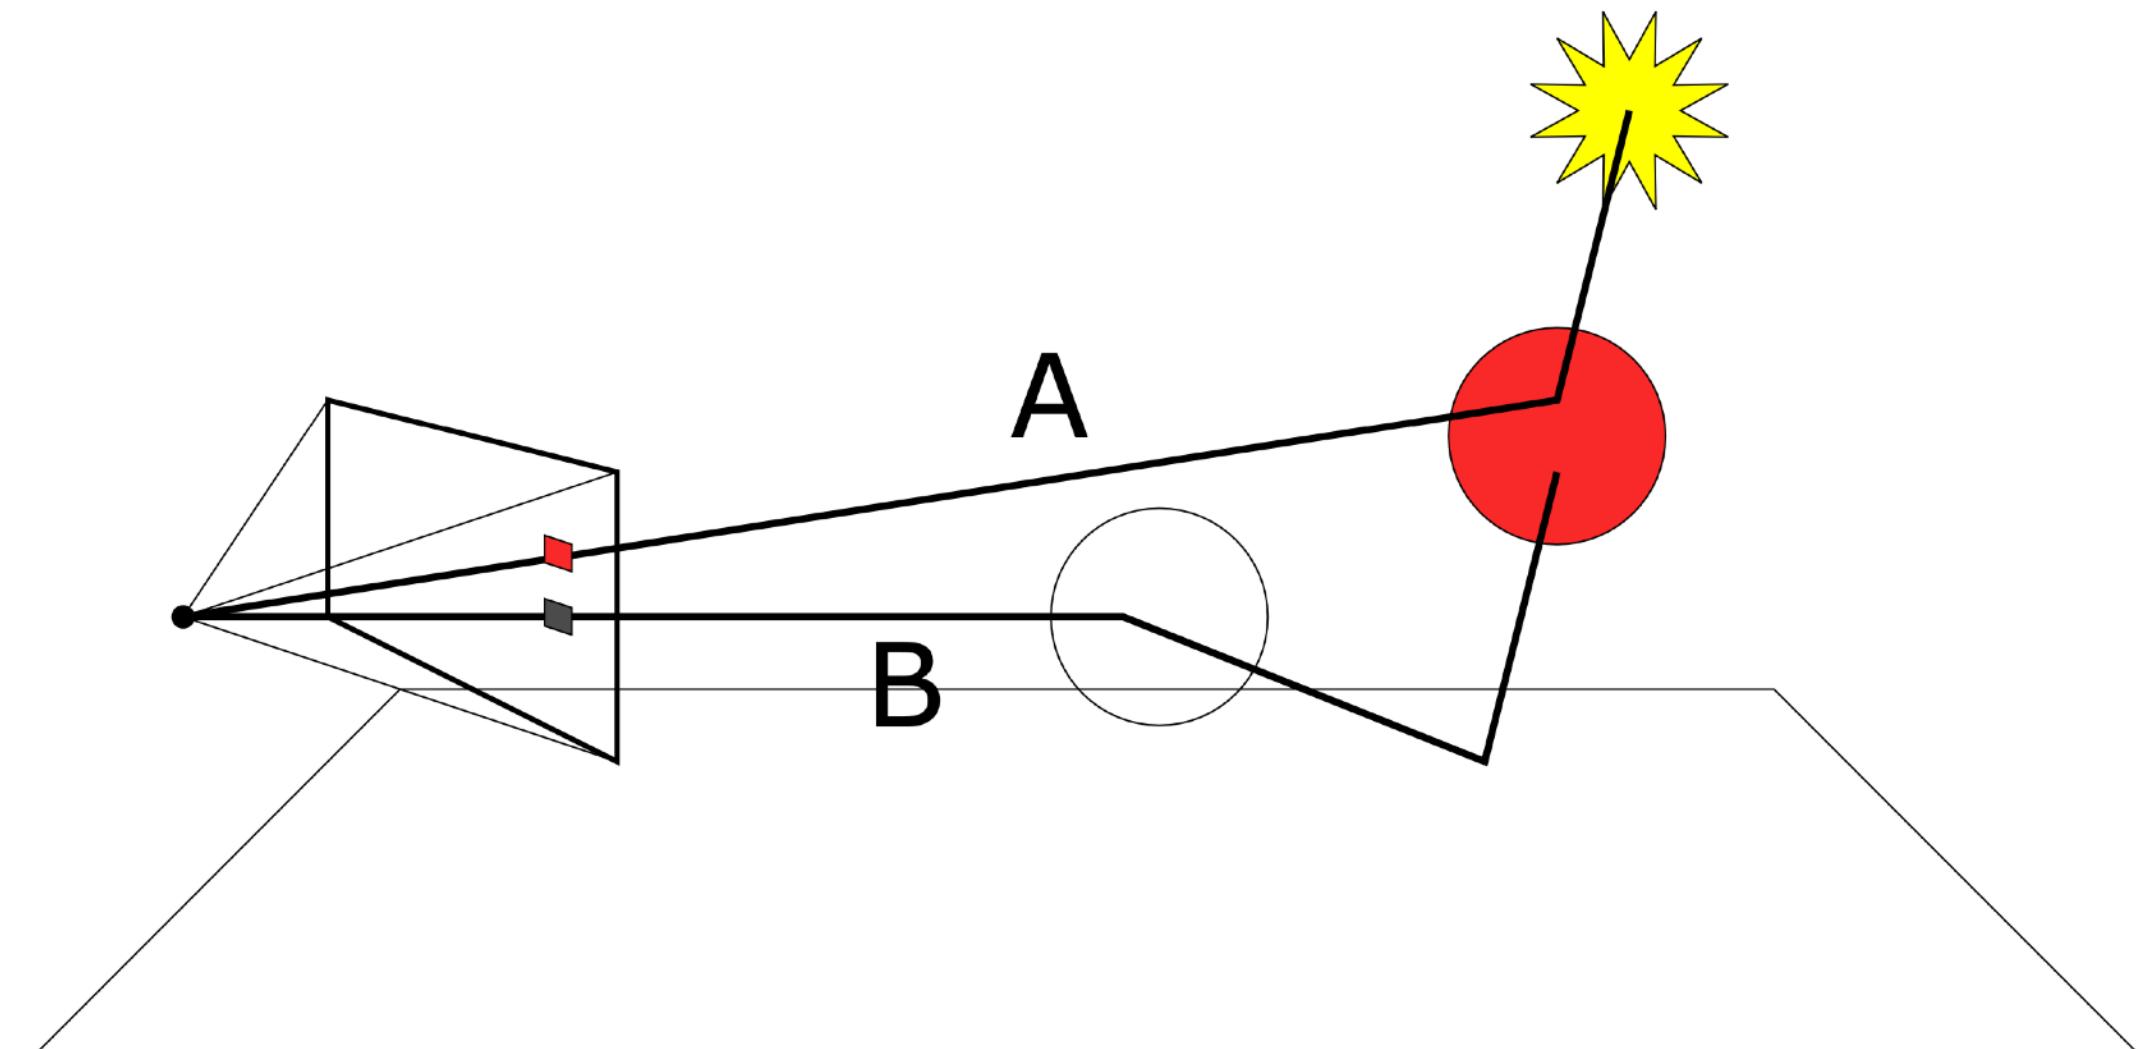
$$\mathbf{T} = \frac{n_1}{n_2} \mathbf{I} + \left(\frac{n_1}{n_2} \cos \theta_i - \cos \theta_t \right) \mathbf{N}$$



By Ulflund - Own work, CC0, <https://commons.wikimedia.org/w/index.php?curid=73684026>

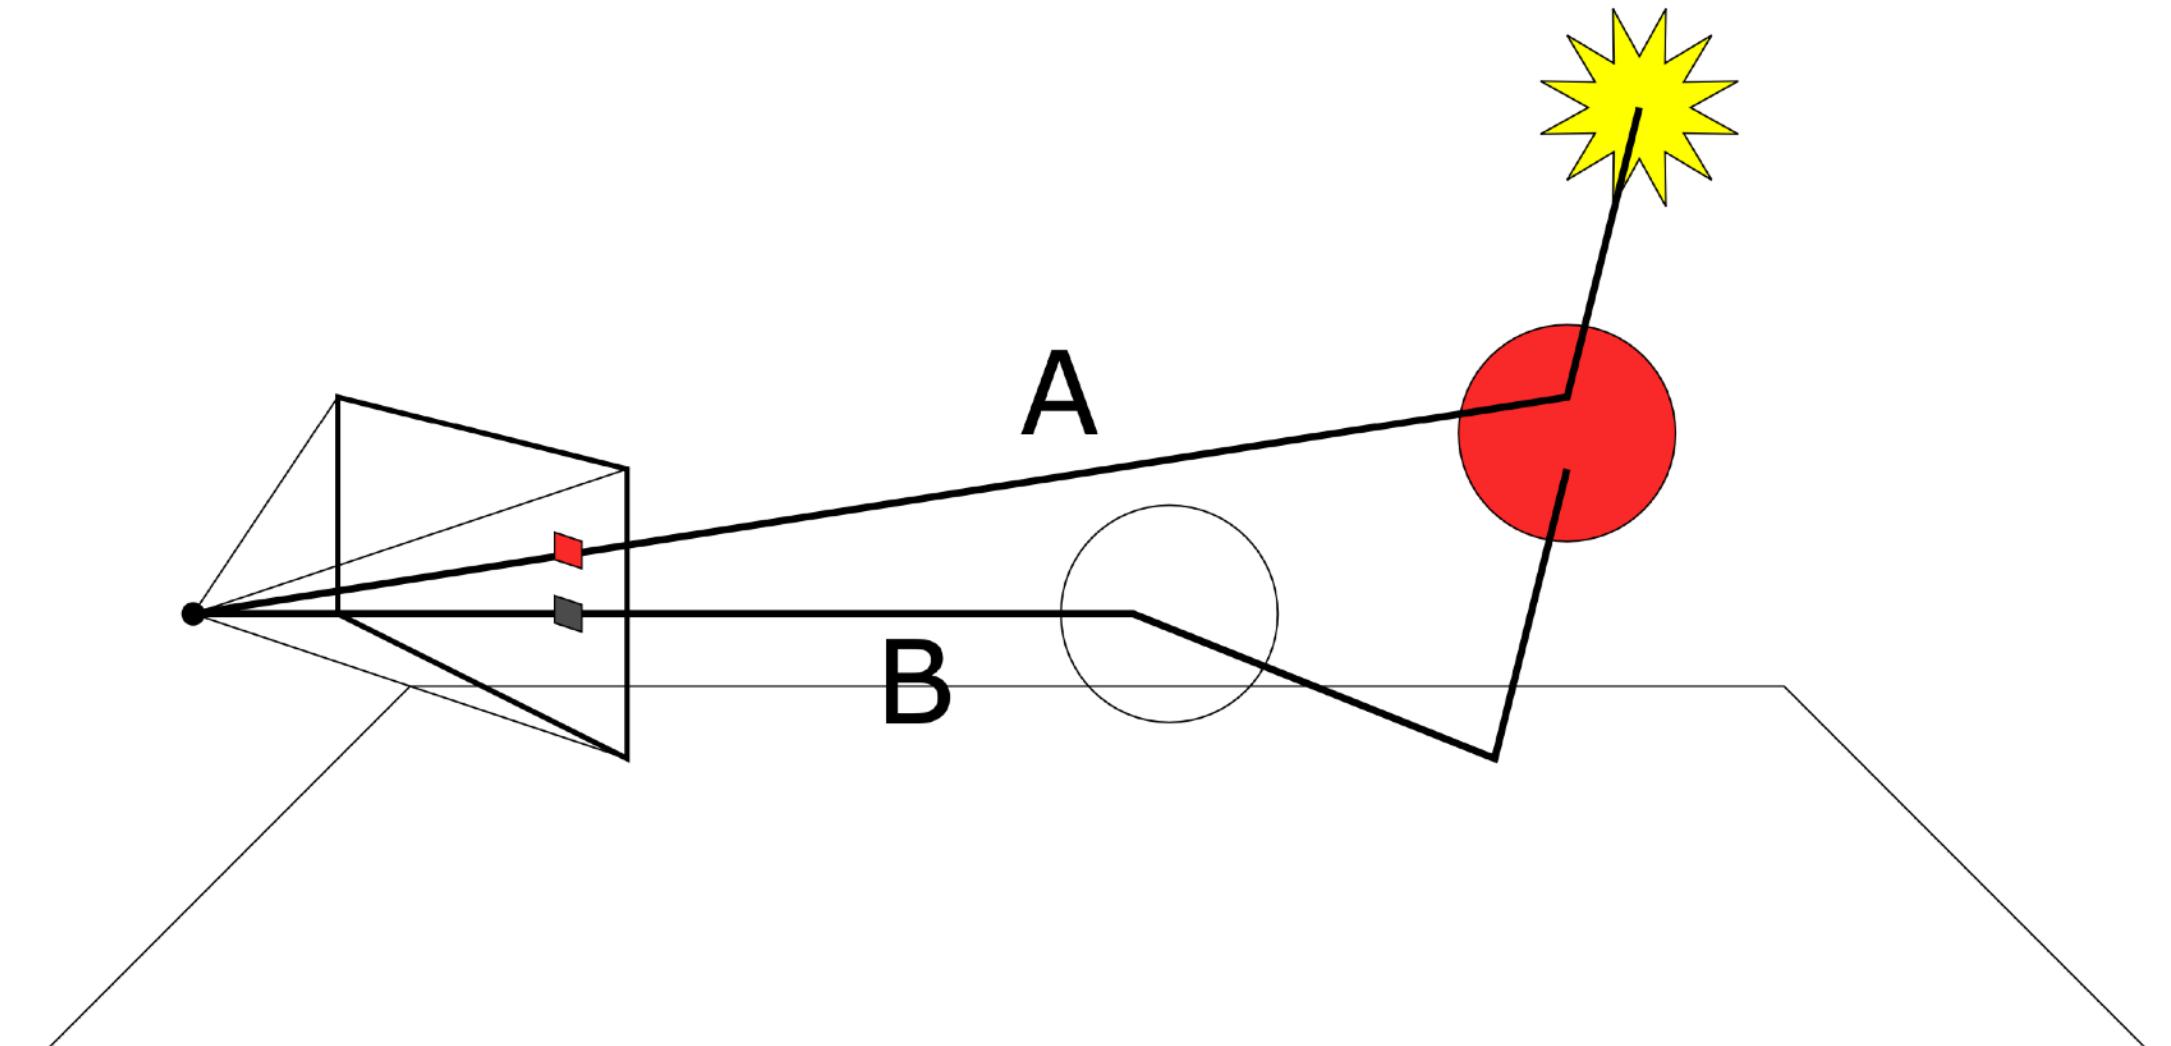
Raytracing Process - Examples

- Ray A is the ‘typical’ case
 - Ray hits the coloured sphere
 - A secondary ray is cast towards the light
 - That secondary ray doesn’t hit anything, so there is no shadow
 - The colour of the hit point on the sphere is used for the pixel
- Ray B is more complex
 - The ray hits a glass sphere which bends the ray, so a new ray is cast
 - The ray bends again as it leaves the sphere, then hits the ground plane
 - A ray is then cast to check for shadows
 - This hits the coloured sphere so the point is in shadow
 - The pixel is, therefore, black



Summary Raytracing Process

- The basic ray tracing algorithm:
 - Set up a camera – a projection point and an image plane
 - For each pixel in the plane:
 - Cast a ray from the projection point through the pixel
 - Determine the first object hit by that ray
 - Cast additional ray(s) to determine lighting
 - Additional rays can be used for reflection, refraction, etc.



Raytracing Process (Pseudo Code)

```
function raycast():
    for all pixels (x,y):
        ray = compute_eye_ray(x, y)
        image(x,y) = trace(ray)

function trace(ray r):
    closest_t = ∞
    closest_o = null

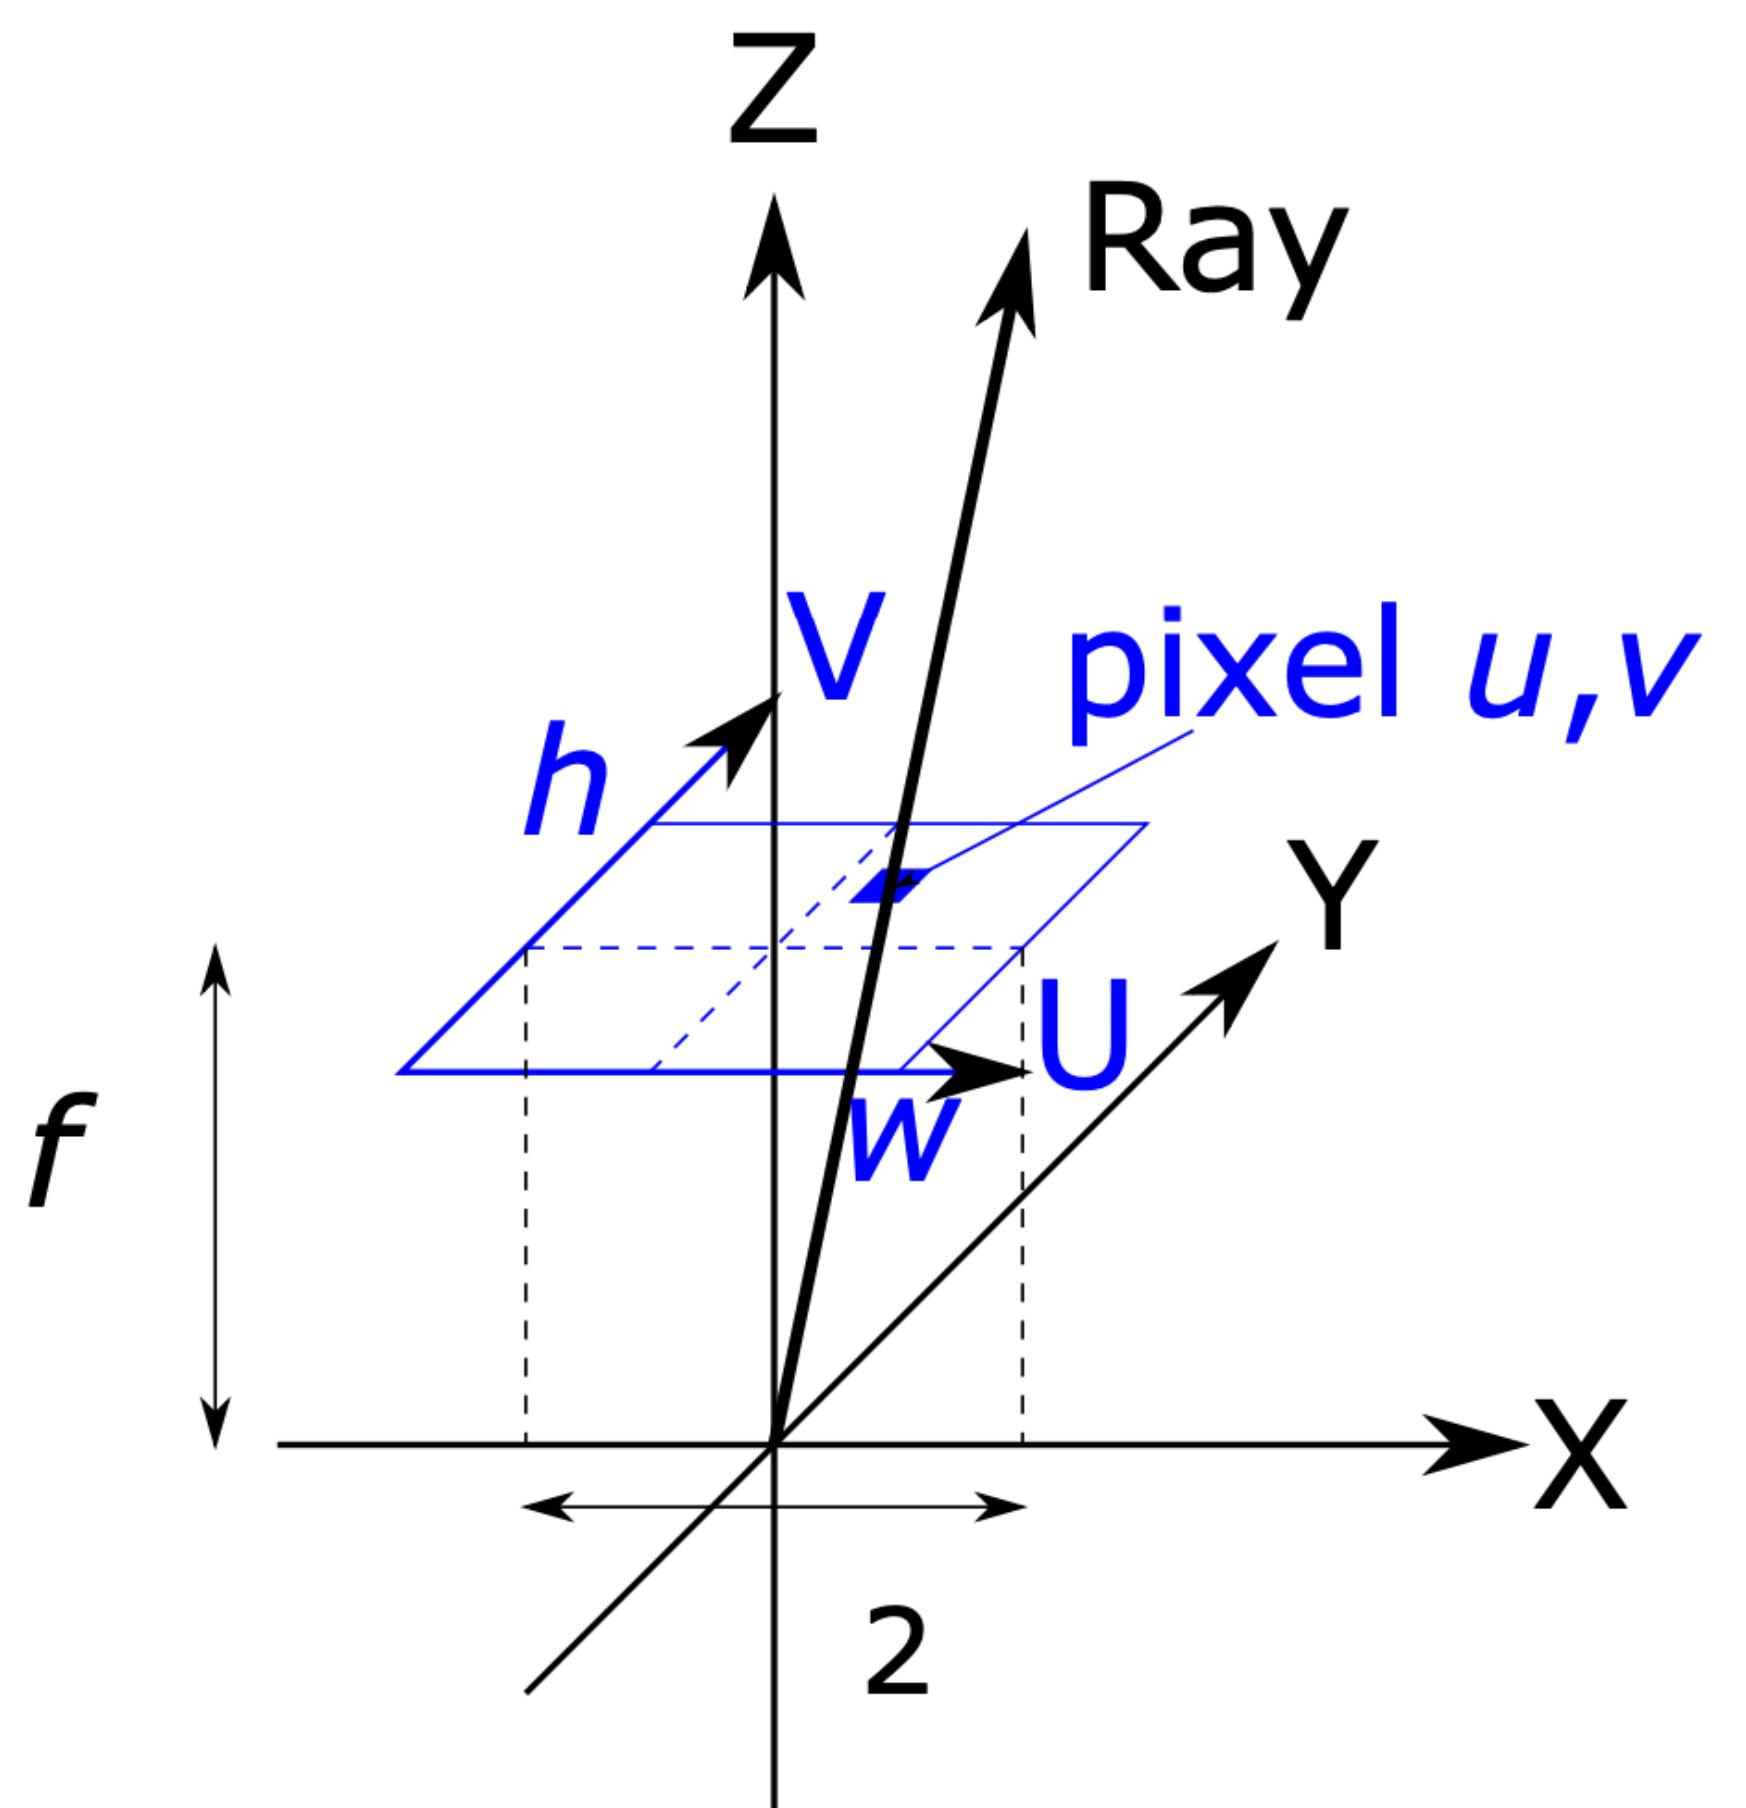
    for all objects o:
        t = compute_intersection(r, o)
        if t > 0 and t < closest_t:
            closest_t = t
            closest_o = o

    if closest_o != null:
        return shade(closest_o, r, closest_t)
    else
        return background_color

function shade(object o, ray r, t):
    x = r(t)          // intersection point
    // evaluate Phong illumination model at x
    return compute_phong(o, x, r)
```

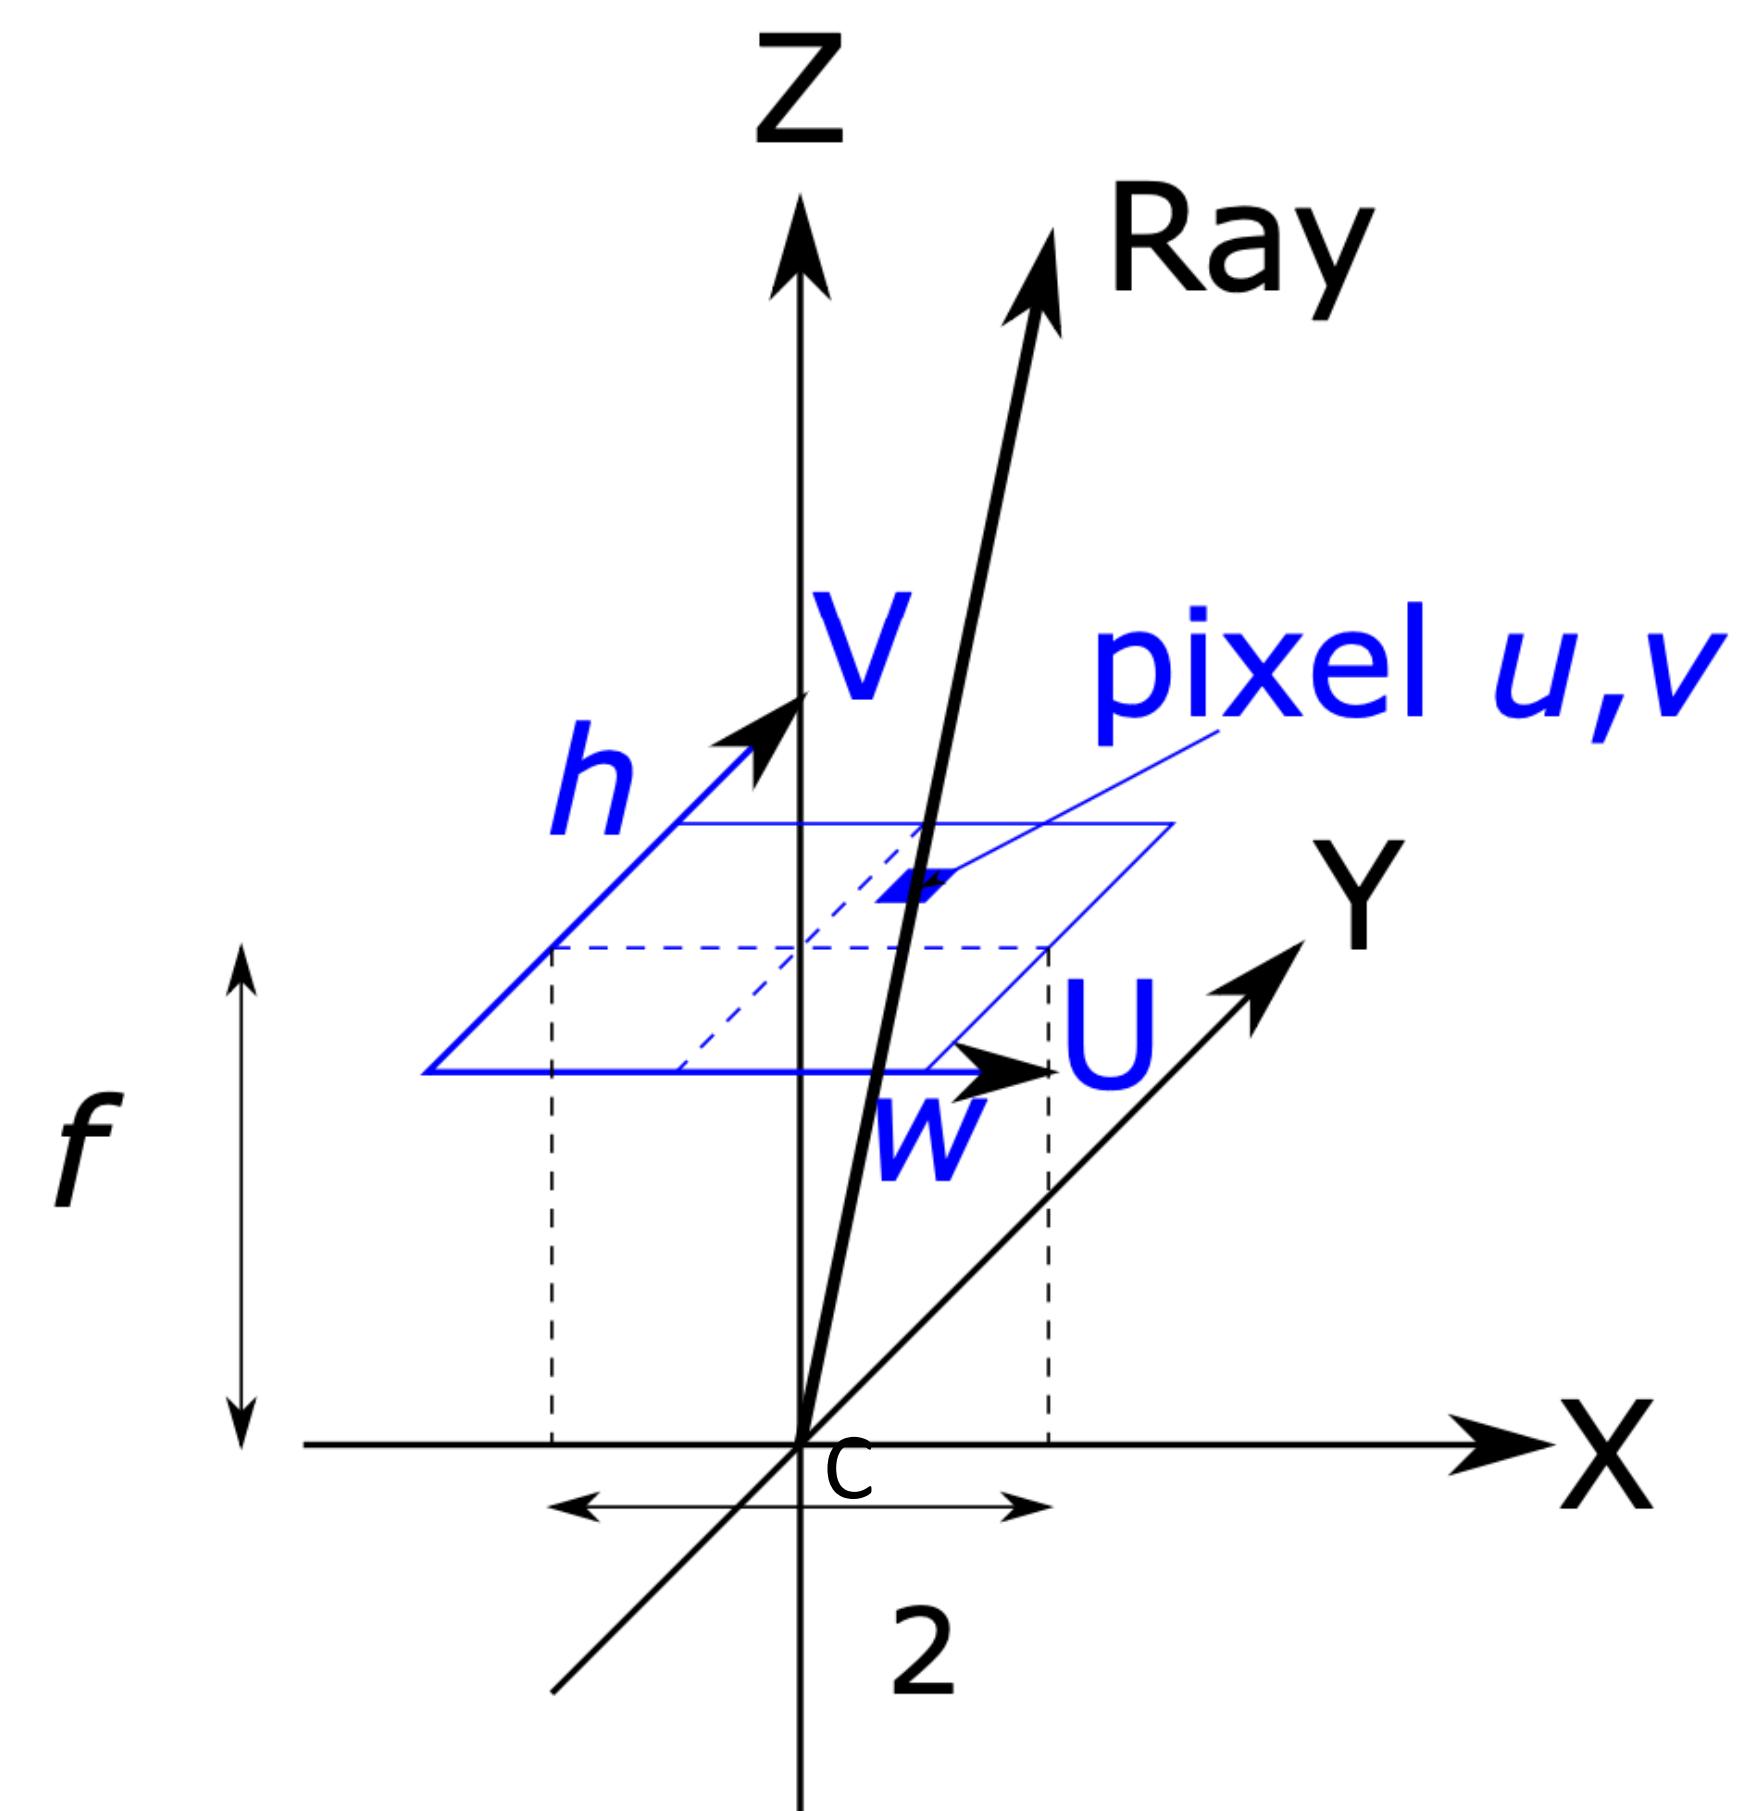
Primary Ray and Camera Model

- Primary Ray
 - Defined by:
 - Origin at the camera center
 - Direction through a pixel on the image plane
 - Intersection tests determine visible surfaces
- Camera models:
 - Pinhole (standard for educational ray tracers)
 - Thin-lens (focus, depth of field)



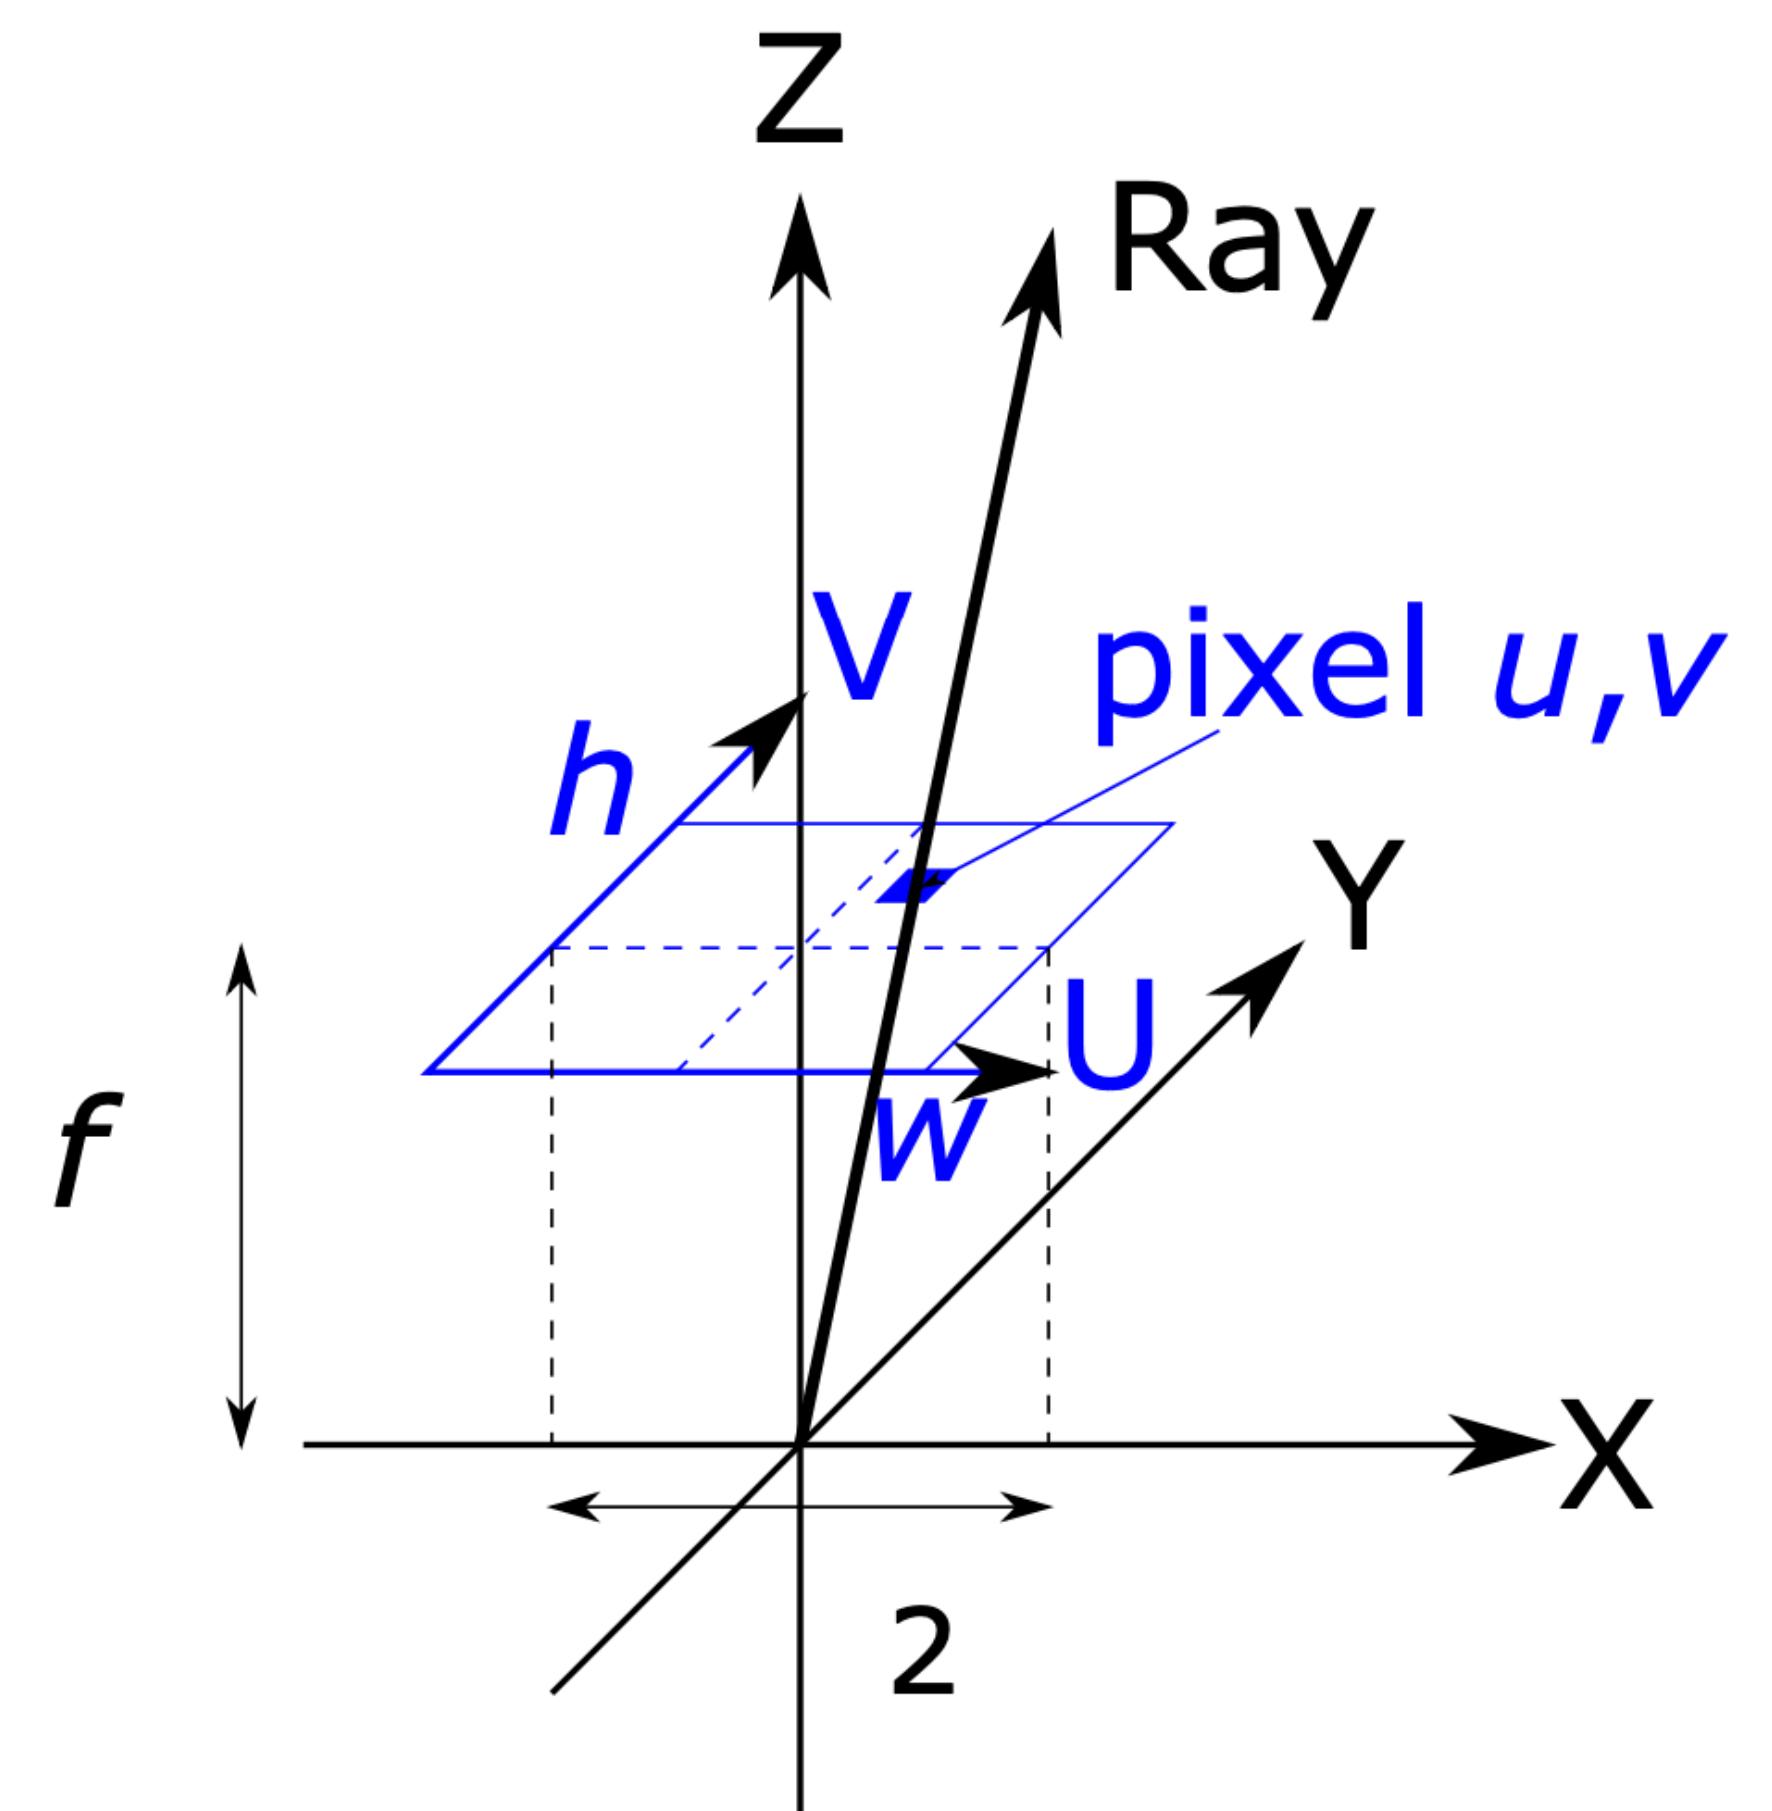
Primary Ray and Camera Model

- The primary ray for each pixel
 - Starts at the camera center, c
 - Goes through the middle of the pixel
- We make some simplifying assumptions
 - The camera is at the origin, looking along $+Z$
 - The (image) U and V axes are aligned with the (world) X and Y axes
 - The image plane is two ‘world units’ wide
 - We are given a focal length, f , and a rendered image size, $w \times h$
 - Our pixels are given integer coordinates (u, v) starting from $(0, 0)$



Primary Ray and Camera Model

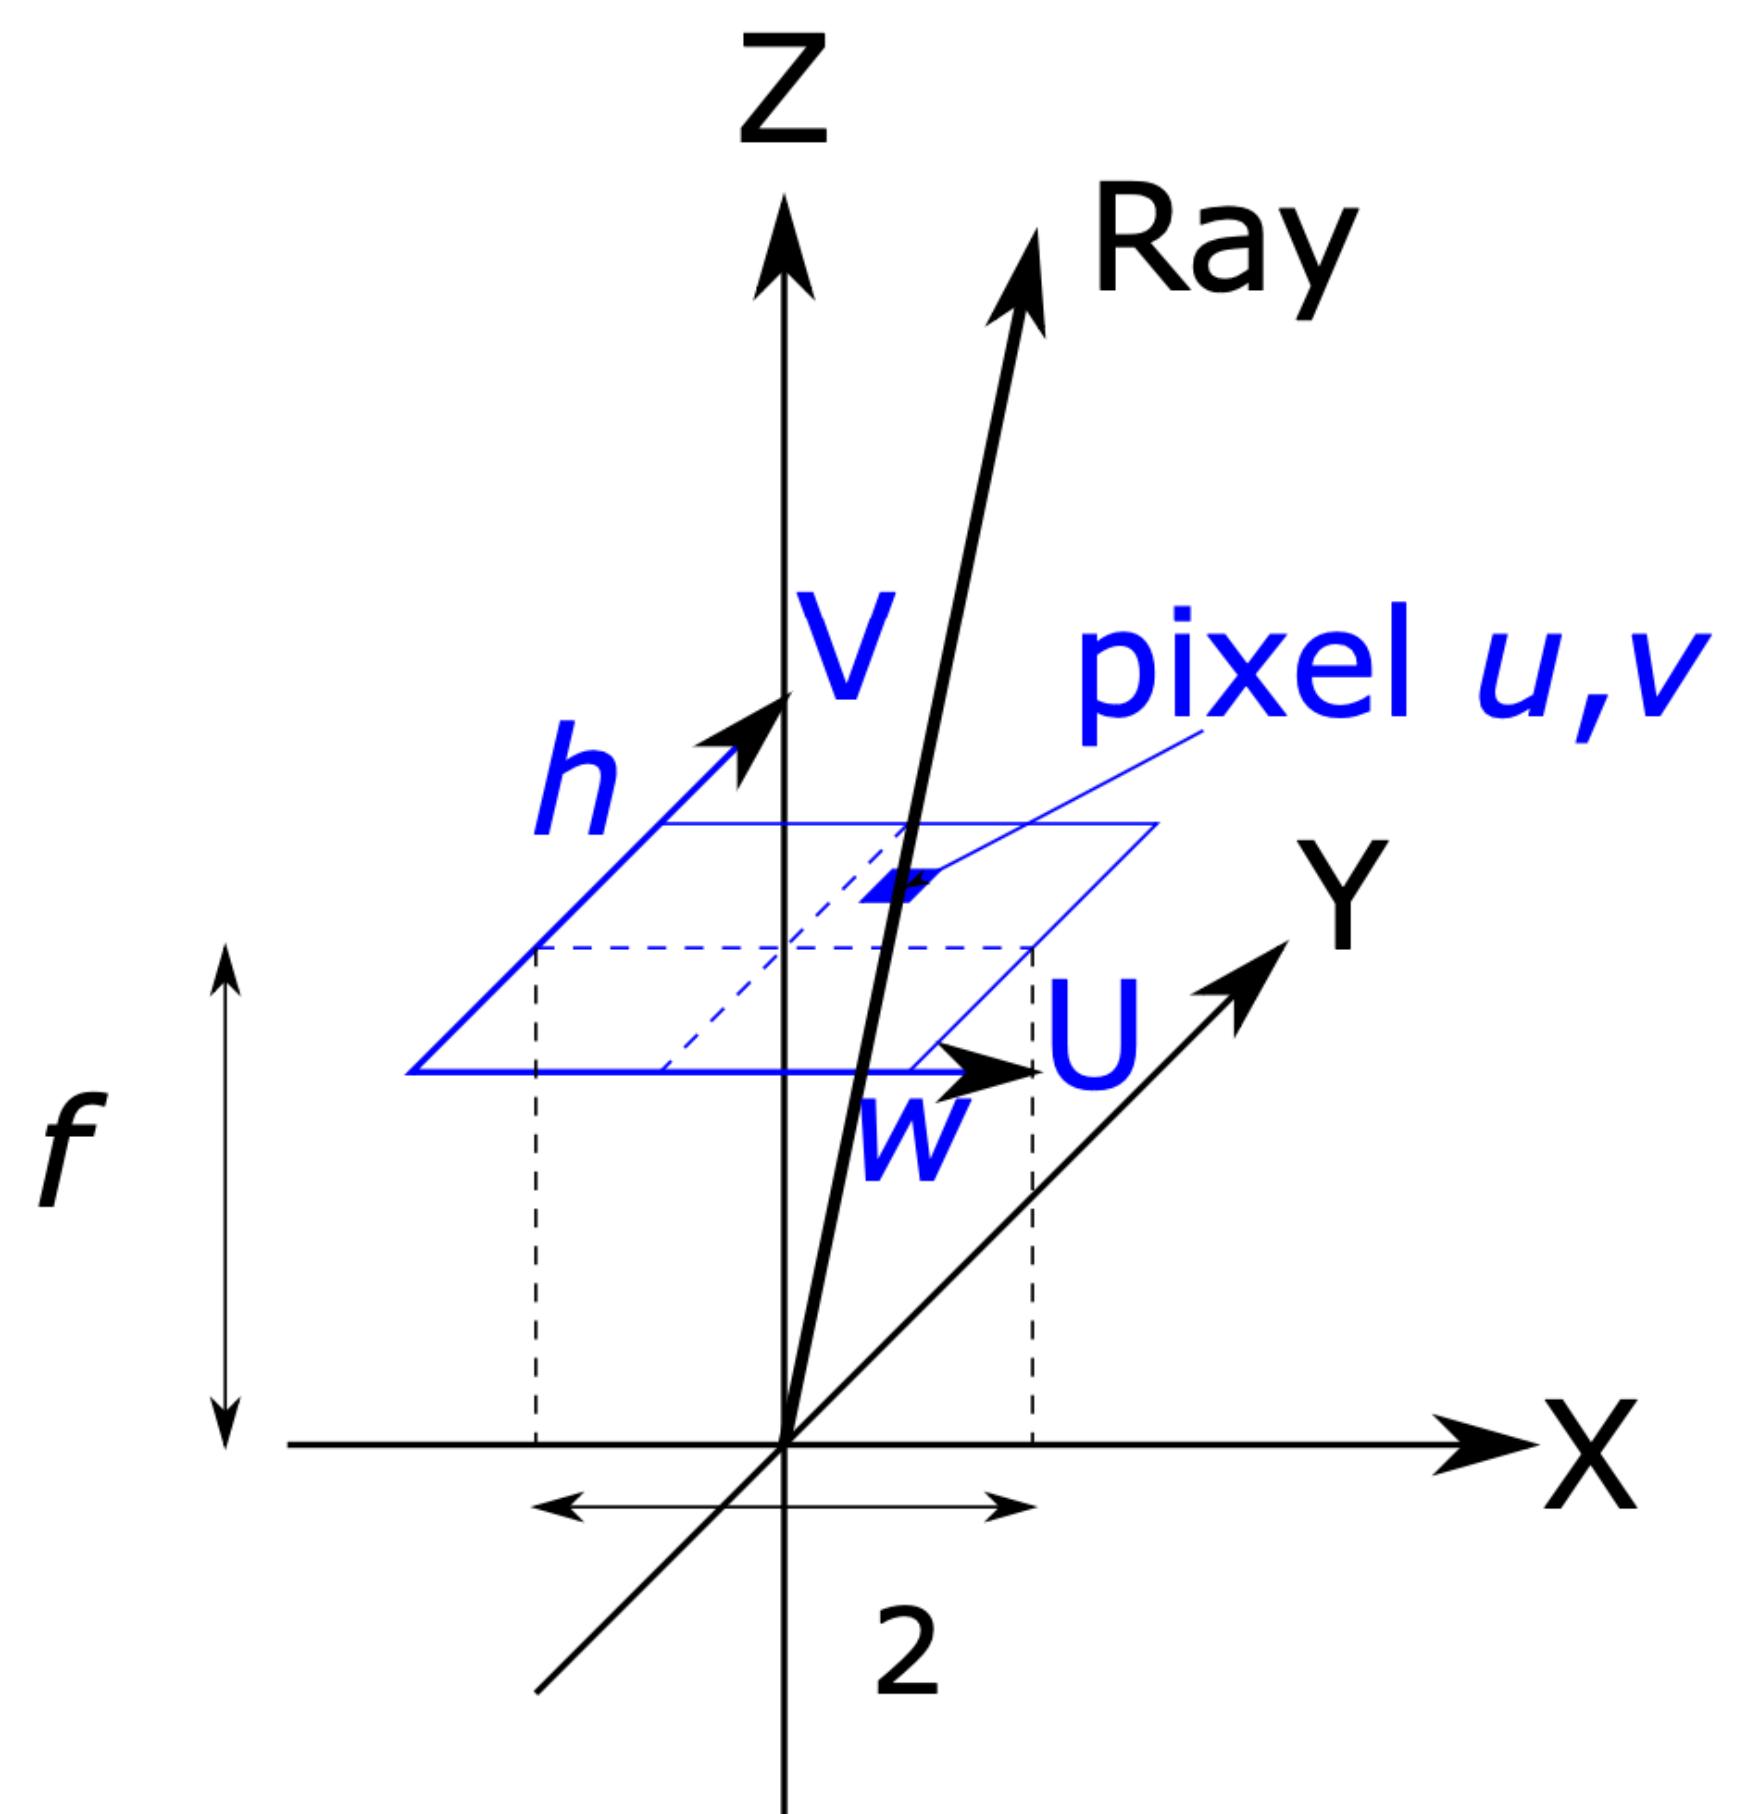
- The ray is defined by two points
 - 1) Camera centre [0,0,0]
 - 2) The centre of the pixel (u, v)
- Pixel/image coordinates are (U, V)
- Need world (X, Y, Z) coordinates
- Where is the image origin?
- How big is each pixel?



Primary Ray and Camera Model

- We can define the image origin:
 - Z value is focal length, f
 - X value is -1 (x values range from -1 to 1)
 - Y value is $\frac{h}{w}$ (y values range from $-\frac{h}{w}$ to $\frac{h}{w}$)
 - Size of each pixel in world units:
 - Width of image is 2
 - So pixels are $\frac{2}{w}$ on a side
 - Coordinates of pixel (u, v) :
$$[-1 + (u + \frac{1}{2})\frac{2}{w}, \frac{-h}{w} + (v + \frac{1}{2})\frac{2}{w}, f]^T$$
 - Our ray: $\mathbf{r}(t) = \mathbf{o} + t \mathbf{d}(u, v)$:

$$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} + \lambda \begin{bmatrix} -1 + \left(u + \frac{1}{2}\right) \frac{2}{w} \\ -\frac{h}{w} + \left(v + \frac{1}{2}\right) \frac{2}{w} \\ f \\ 0 \end{bmatrix}$$



Primary Ray and Camera Model

- We can define the image origin:

- Z value is focal length, f

- X value is -1 (x values range from -1 to 1)

- Y value is $\frac{h}{w}$ (y values range from $-\frac{h}{w}$ to $\frac{h}{w}$)

- Size of each pixel in world units:

- Width of image is 2

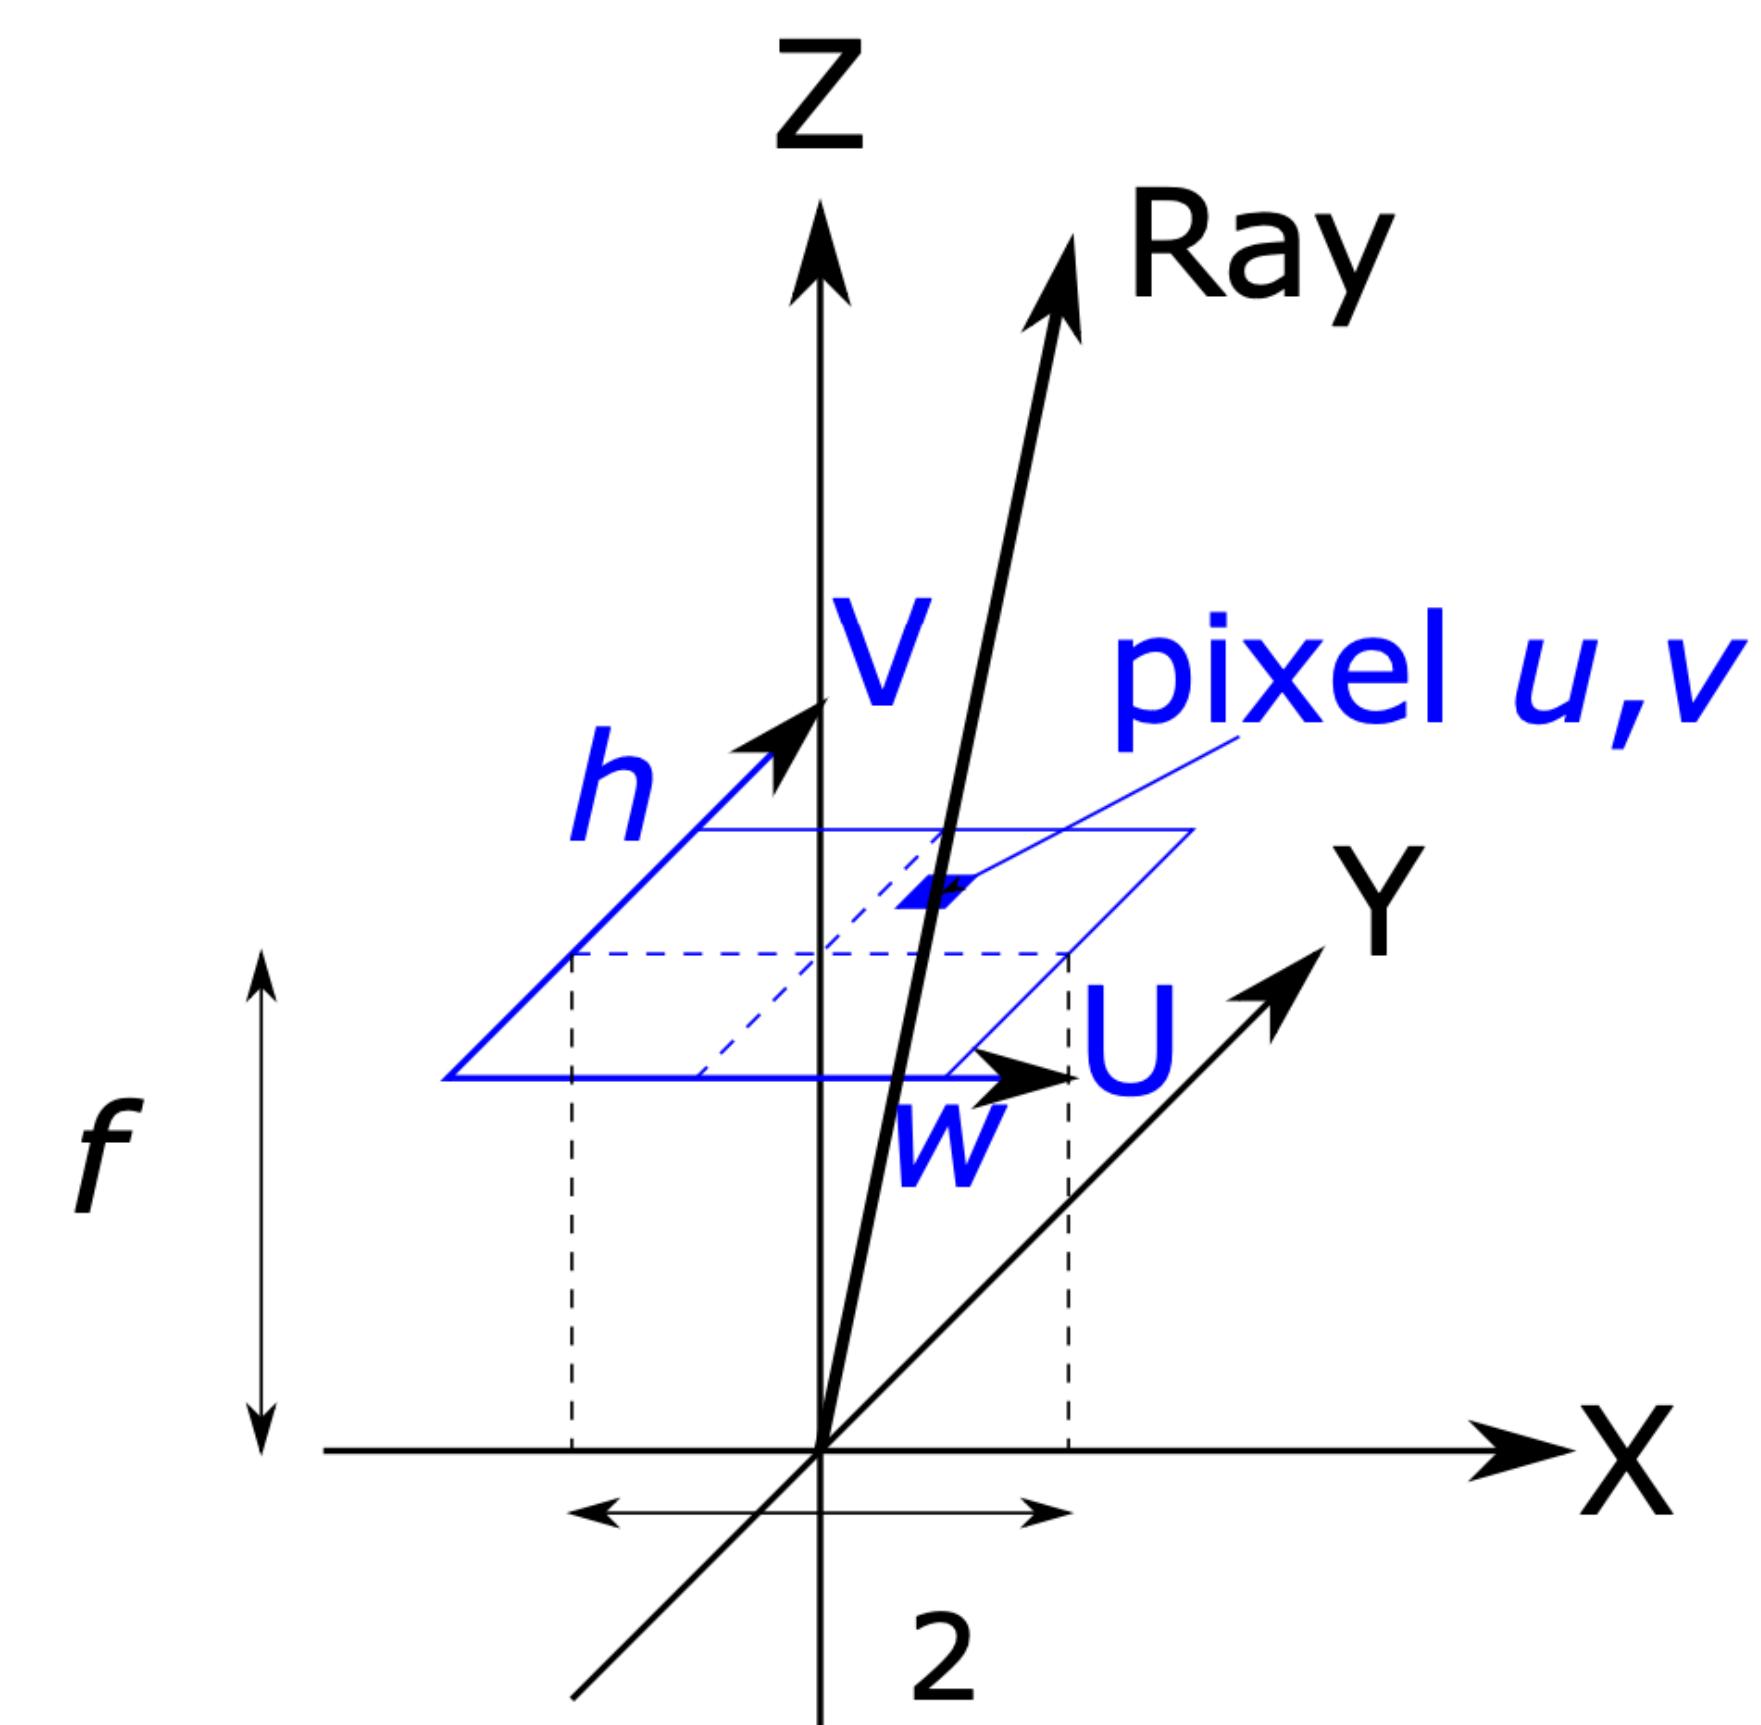
- So pixels are $\frac{2}{w}$ on a side

- Coordinates of pixel (u, v) :

$$[-1 + (u + \frac{1}{2})\frac{2}{w}, \frac{-h}{w} + (v + \frac{1}{2})\frac{2}{w}, f]^T$$

- Our ray: $\mathbf{r}(t) = \mathbf{o} + t \mathbf{d}(u, v)$:

$$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} + \lambda \begin{bmatrix} -1 + (u + \frac{1}{2})\frac{2}{w} \\ \frac{-h}{w} + (v + \frac{1}{2})\frac{2}{w} \\ f \\ 0 \end{bmatrix}$$

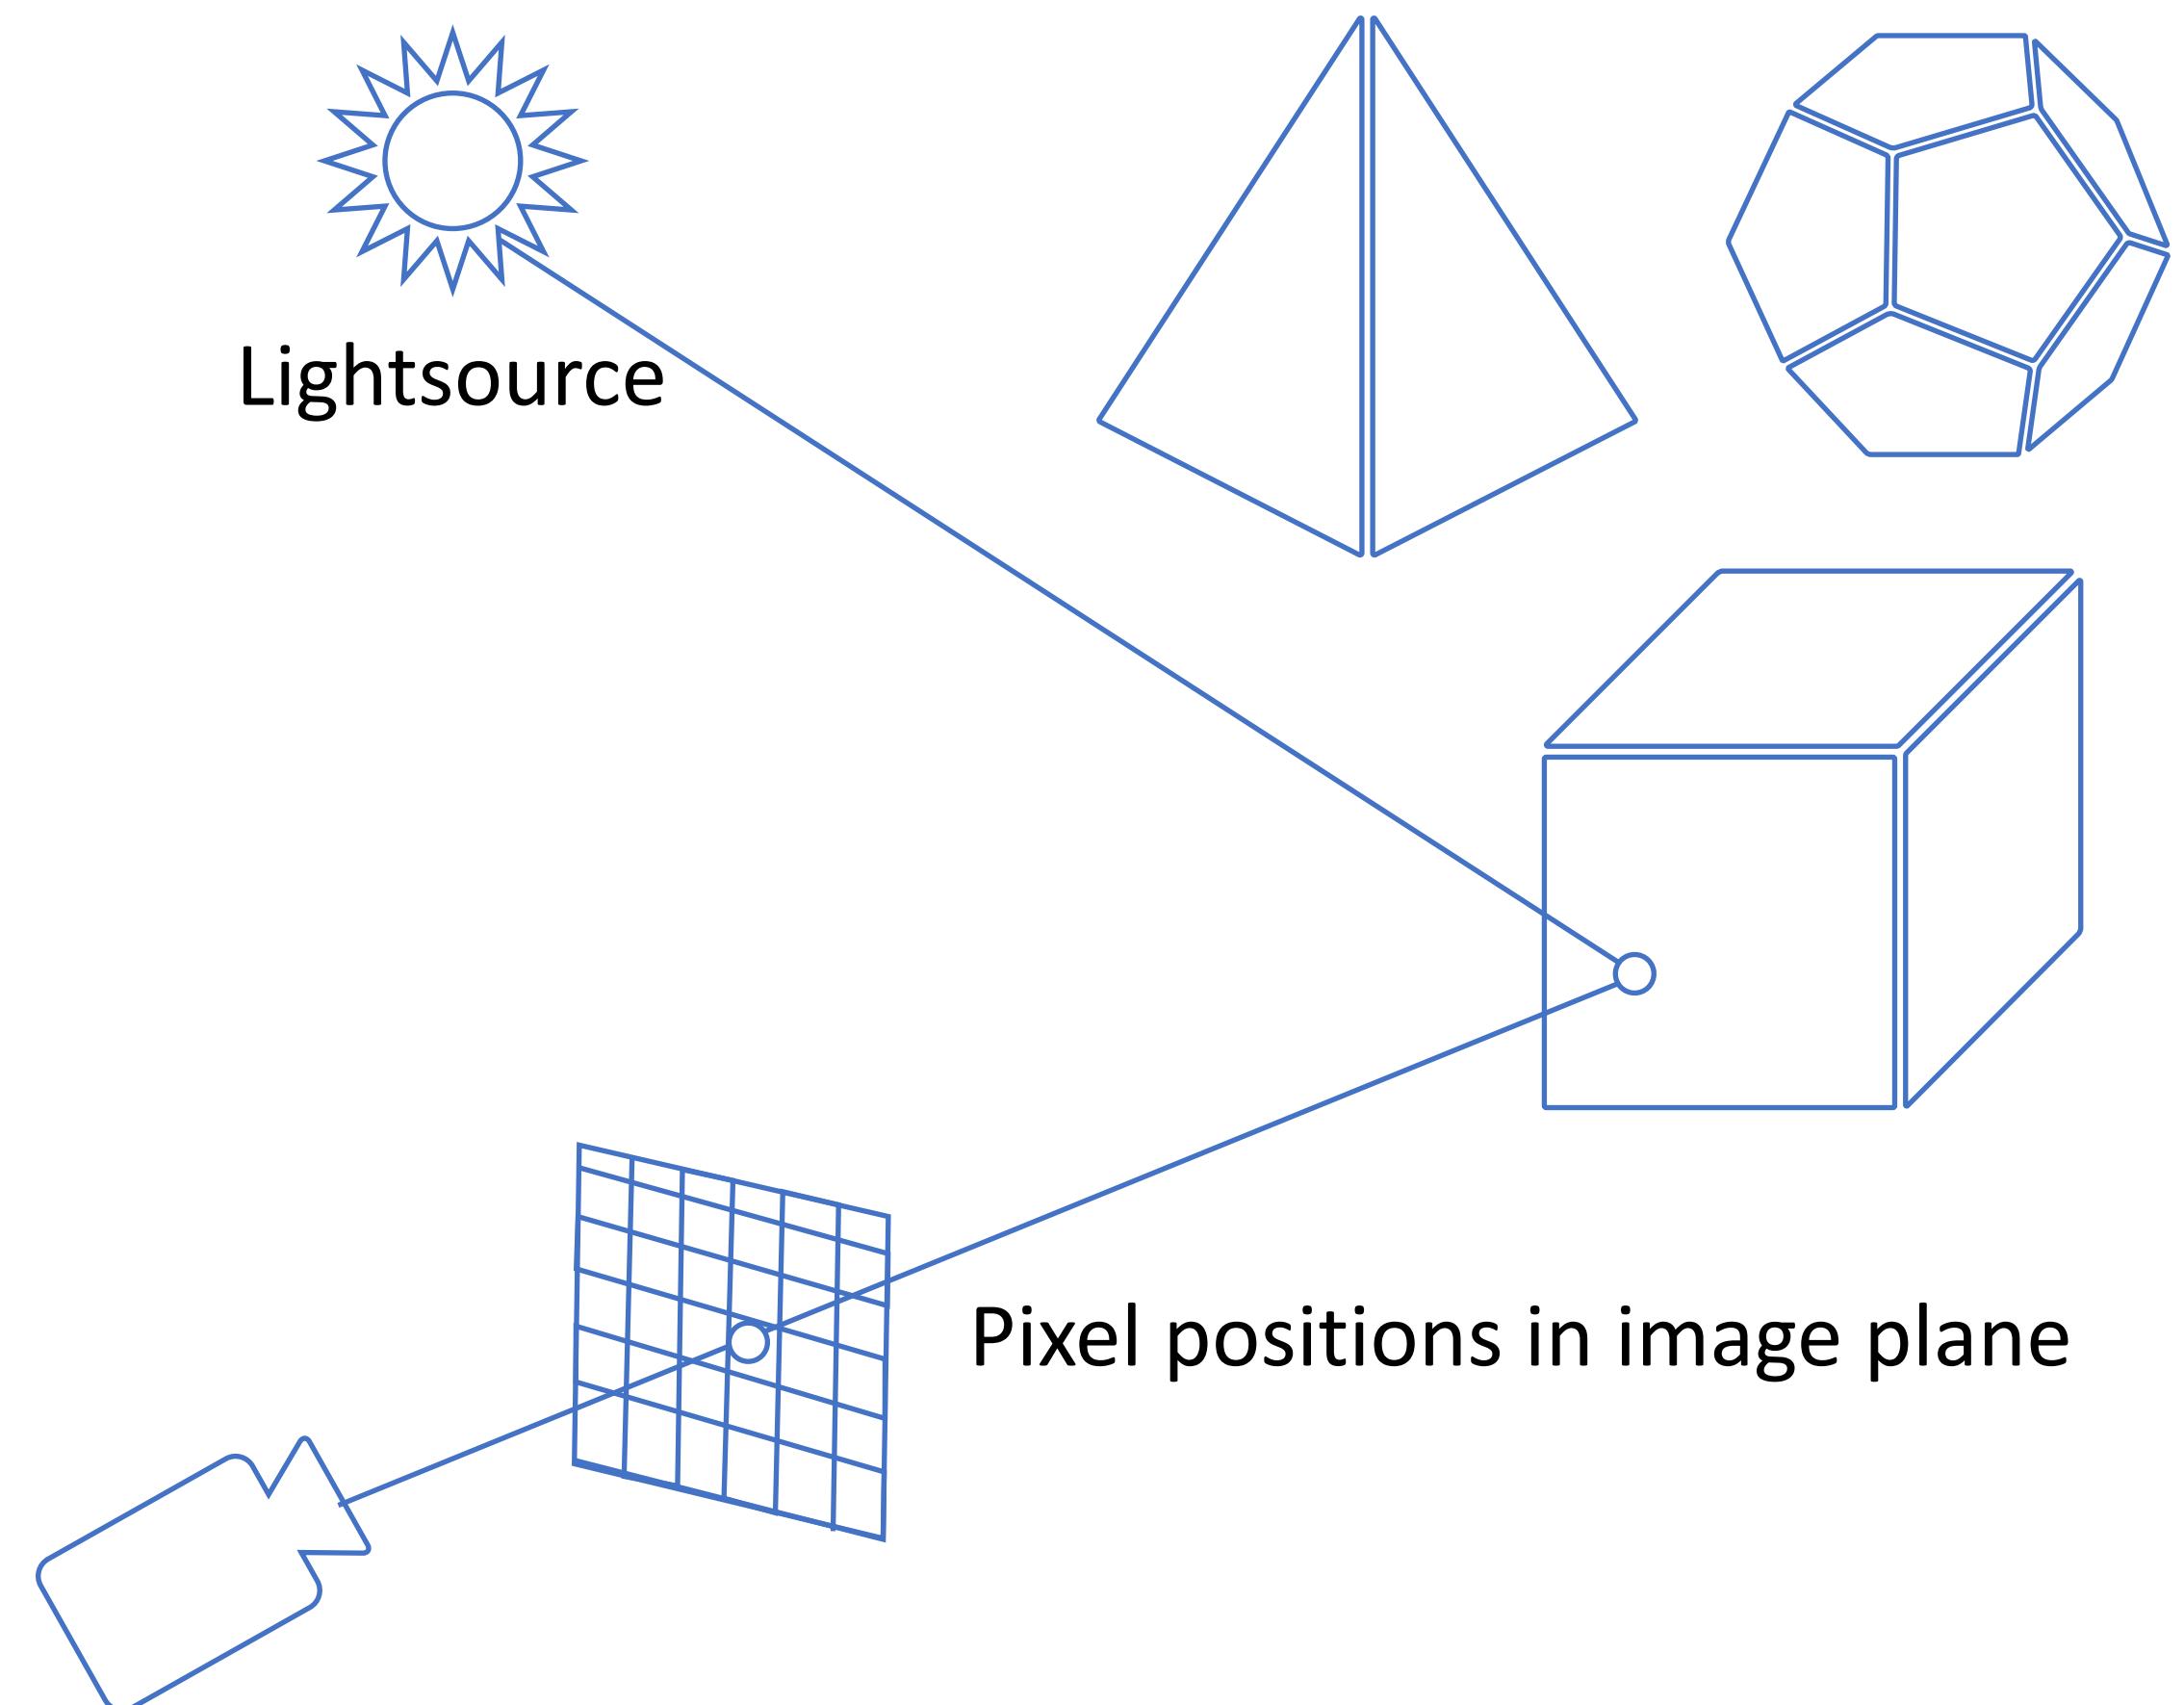


$$\mathbf{p}(0, 0) = \begin{bmatrix} -\frac{1}{2} \\ -\frac{1}{2} \\ f \end{bmatrix}, \quad \mathbf{p}(1, 0) = \begin{bmatrix} \frac{1}{2} \\ -\frac{1}{2} \\ f \end{bmatrix},$$

$$\mathbf{p}(0, 1) = \begin{bmatrix} -\frac{1}{2} \\ \frac{1}{2} \\ f \end{bmatrix}, \quad \mathbf{p}(1, 1) = \begin{bmatrix} \frac{1}{2} \\ \frac{1}{2} \\ f \end{bmatrix}.$$

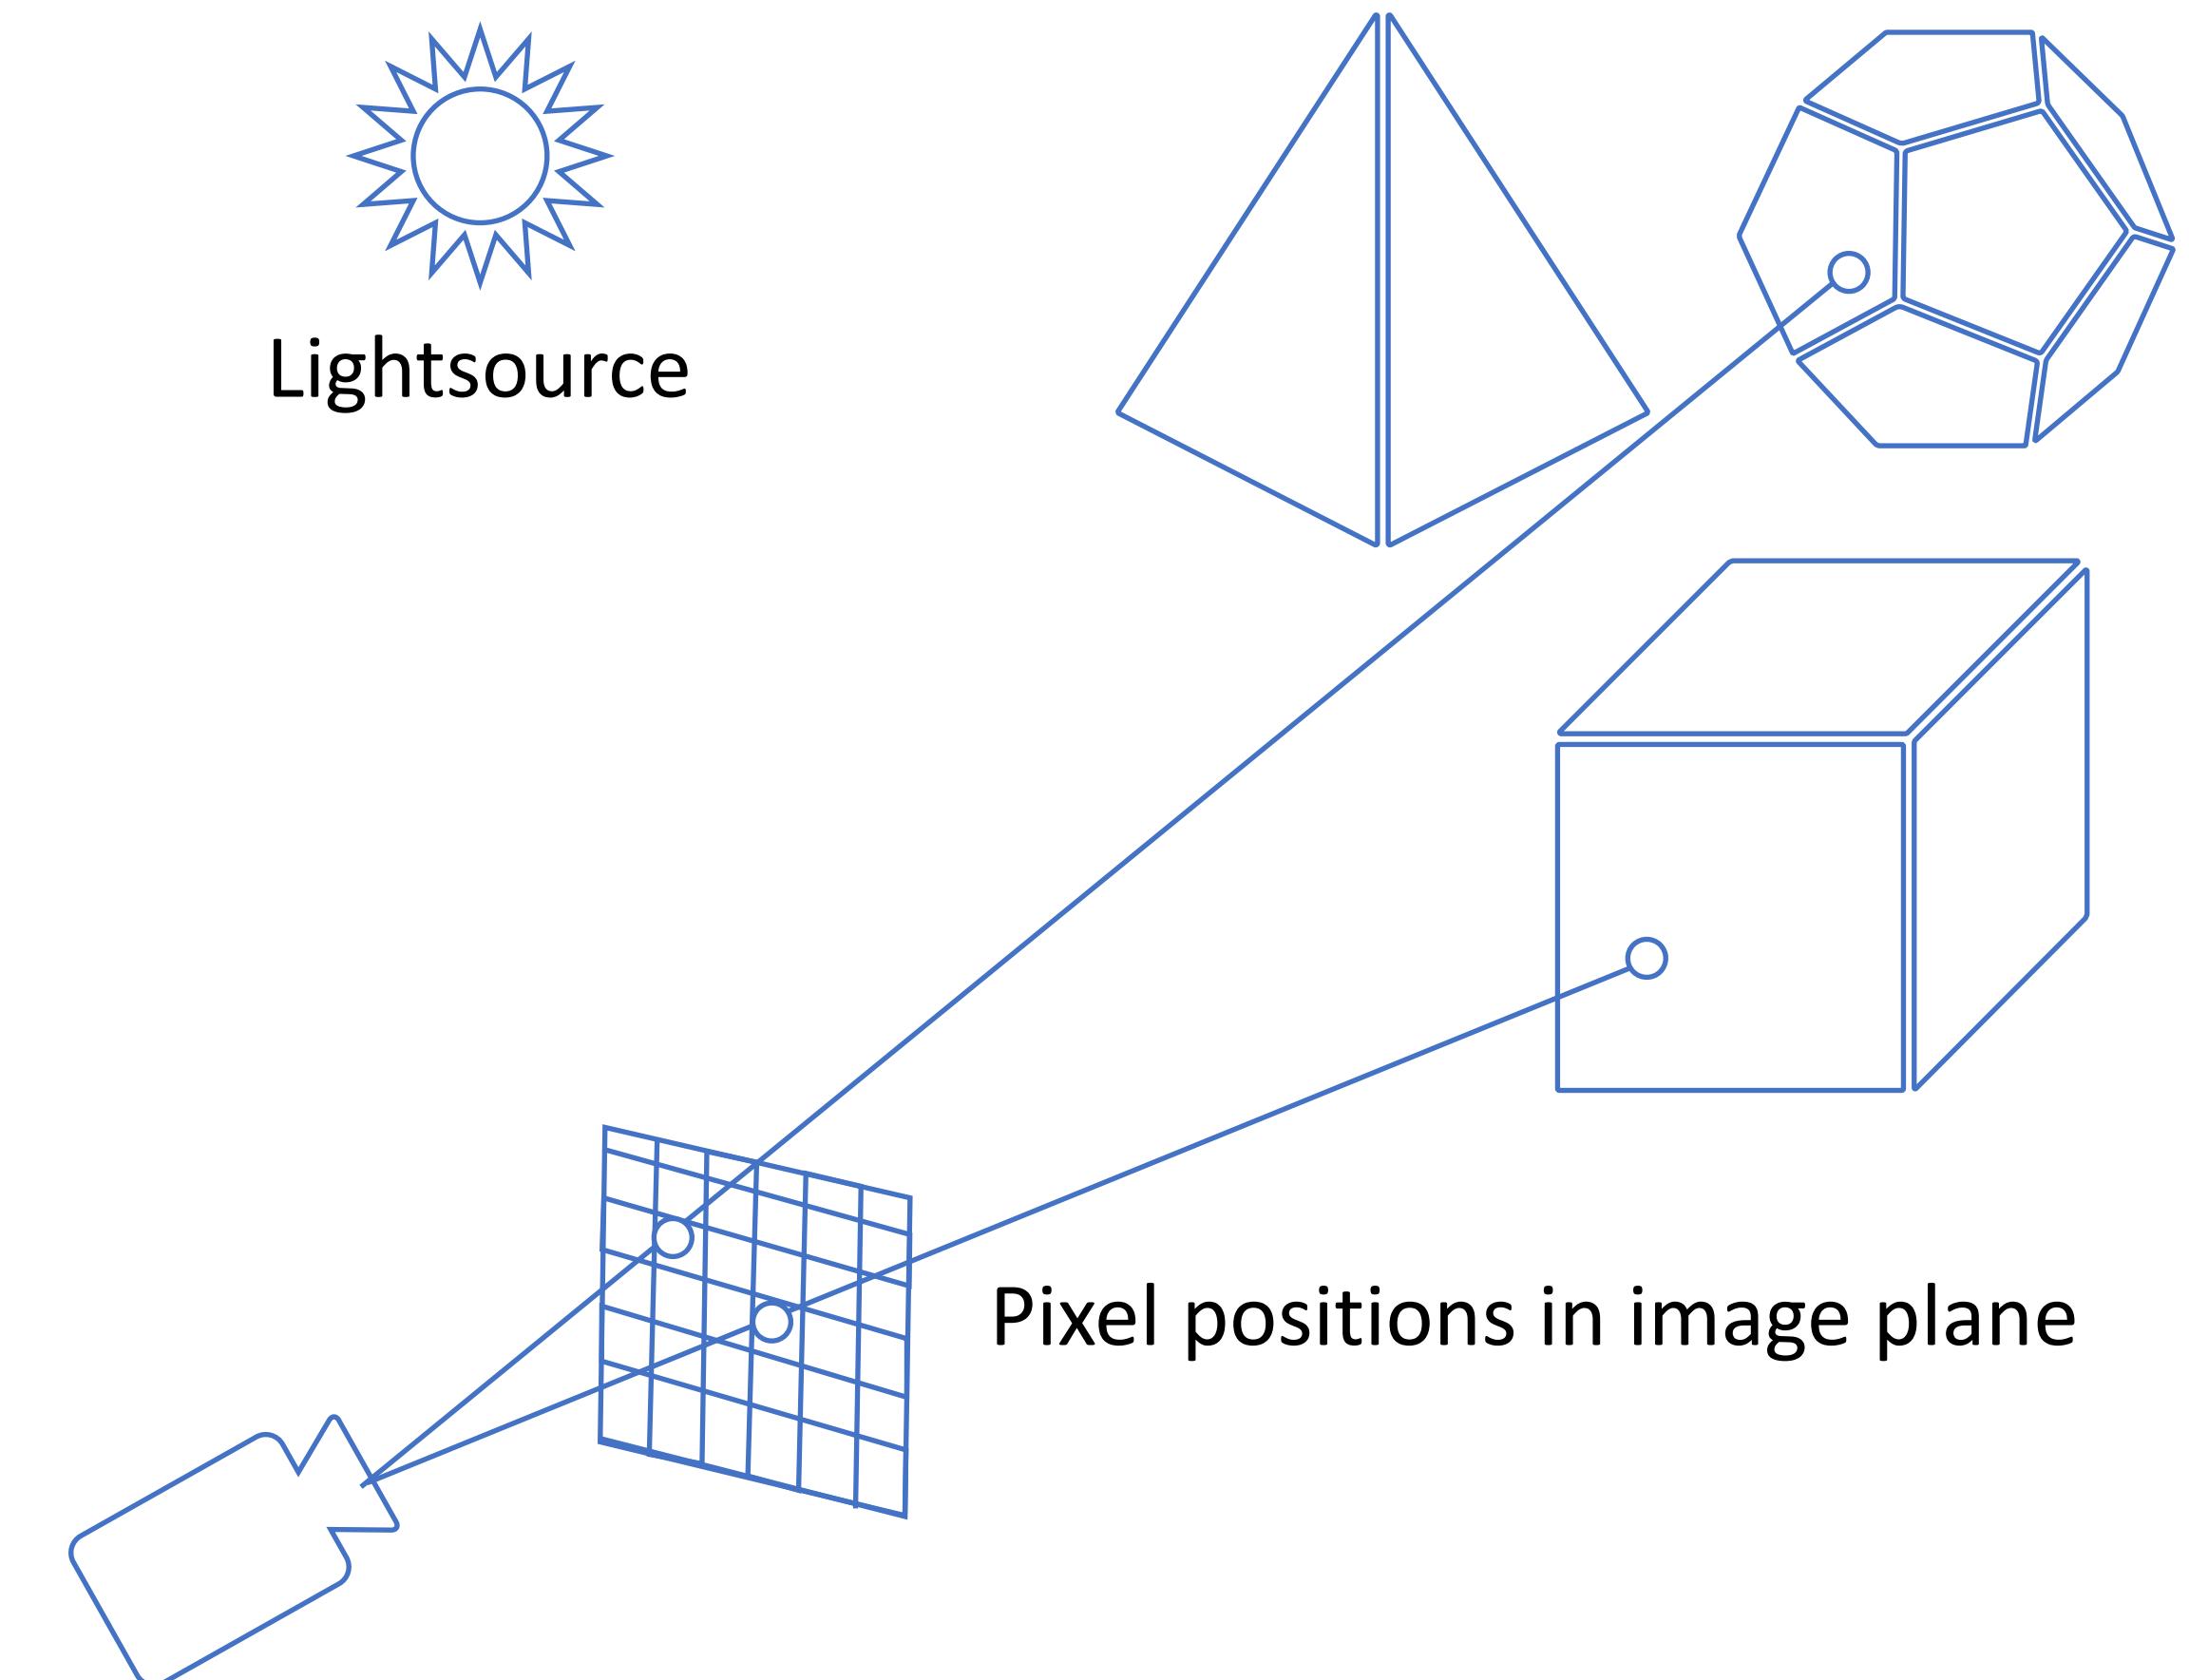
Secondary Rays

- Reflection rays to model specular and mirror-like materials
- Refraction rays for transparent materials:
 - Snell's law and Fresnel terms
 - Recursion depth controls complexity and runtime
- Enables physically plausible mirror reflections, transparency, and glossy effects



Geometric Intersection Tests

- Ray–sphere: closed-form quadratic equation
- Ray–box: testing each side
- Intersection cost dominates rendering -> acceleration required



Ray-Sphere Intersections

- Spheres have a simple mathematical form
 - Given a point $p = (x, y, z)$
 - Its squared distance from the origin is
$$\|p\|^2 = p \cdot p = x^2 + y^2 + z^2$$
 - So a unit sphere at the origin is defined by
$$p^2 = x^2 + y^2 + z^2 = 1$$

Ray-Sphere Intersections

- We have the equations:

$$p = p_0 + \lambda d$$

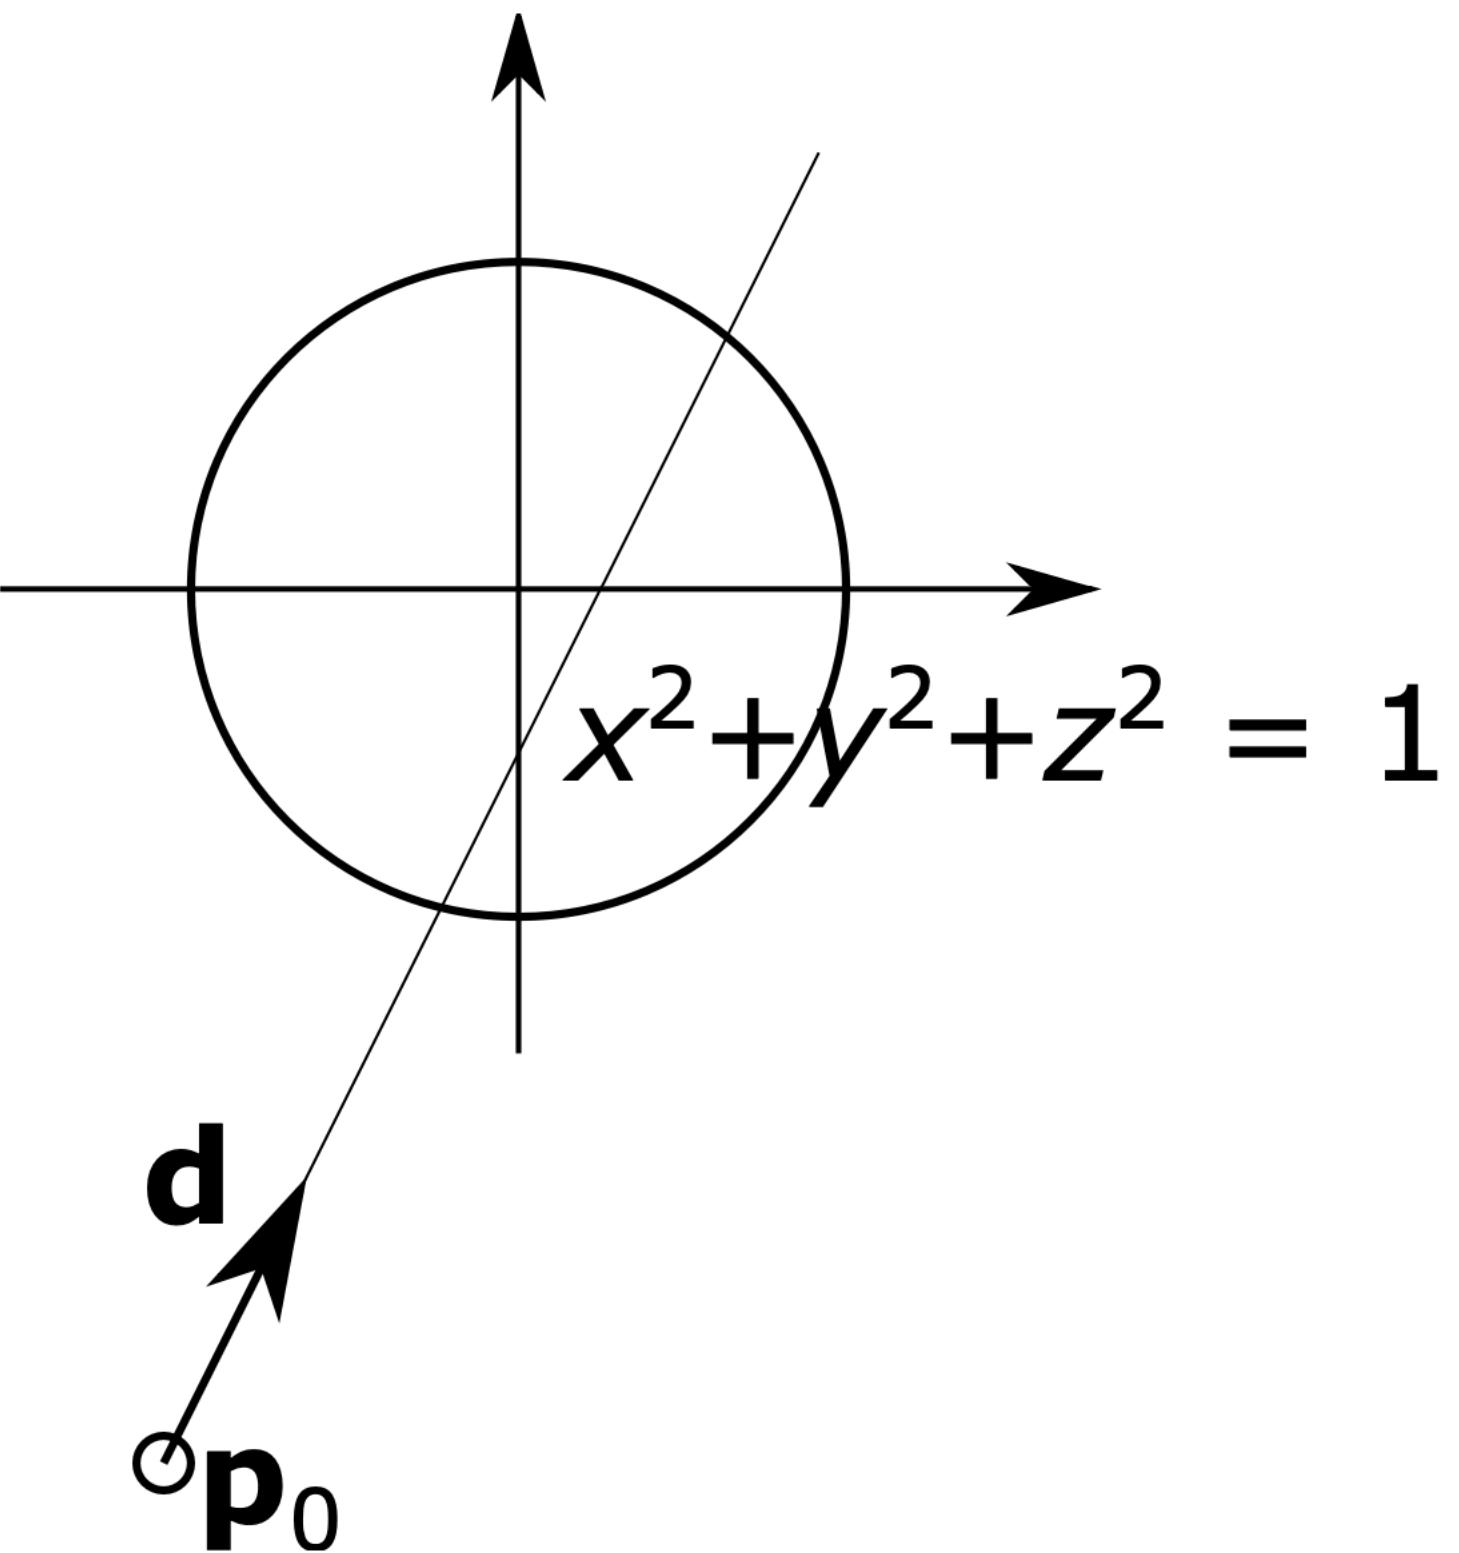
and

$$p \cdot p = 1$$

So

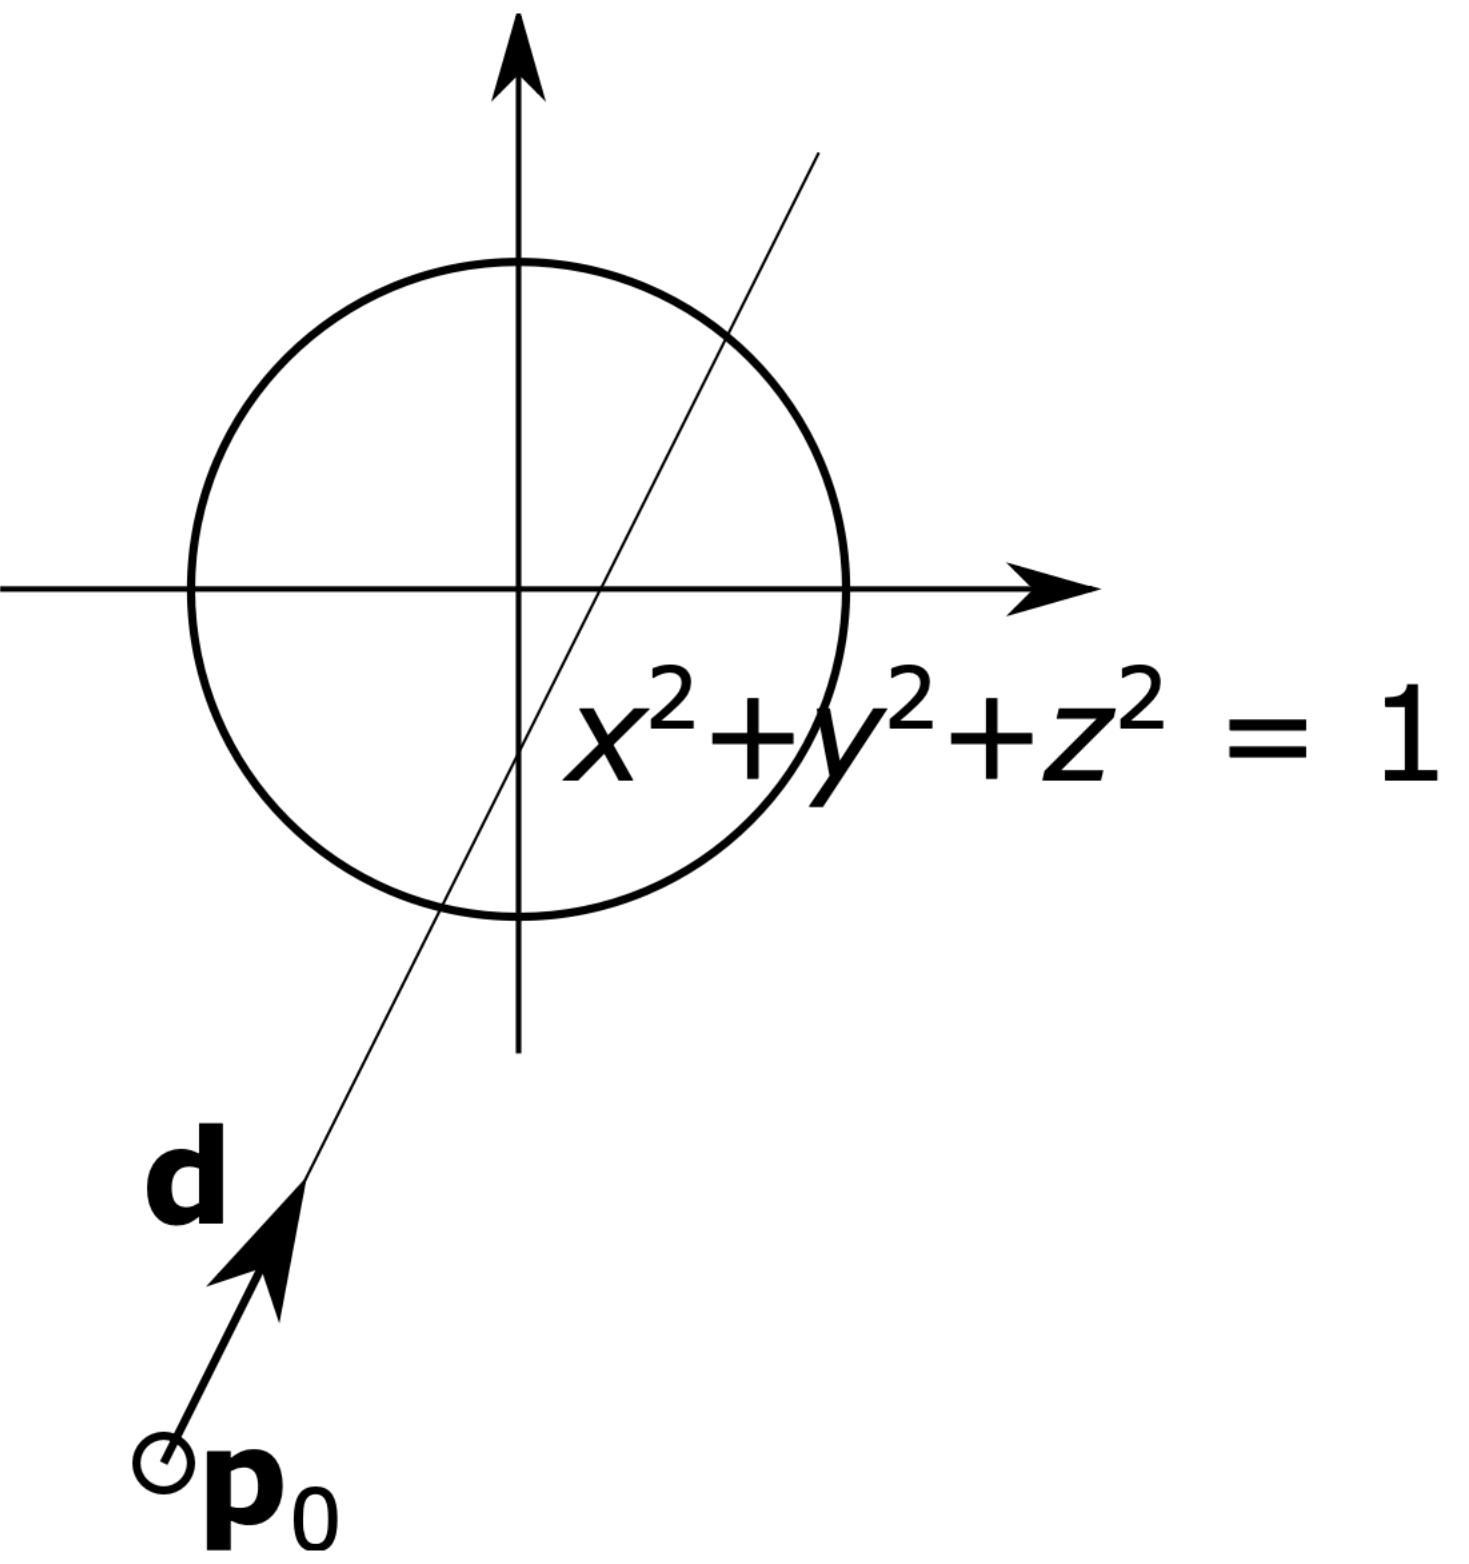
$$(p_0 + \lambda d) \cdot (p_0 + \lambda d) = 1$$

$$p_0 \cdot p_0 + 2\lambda p_0 \cdot d + \lambda^2 d \cdot d = 1$$



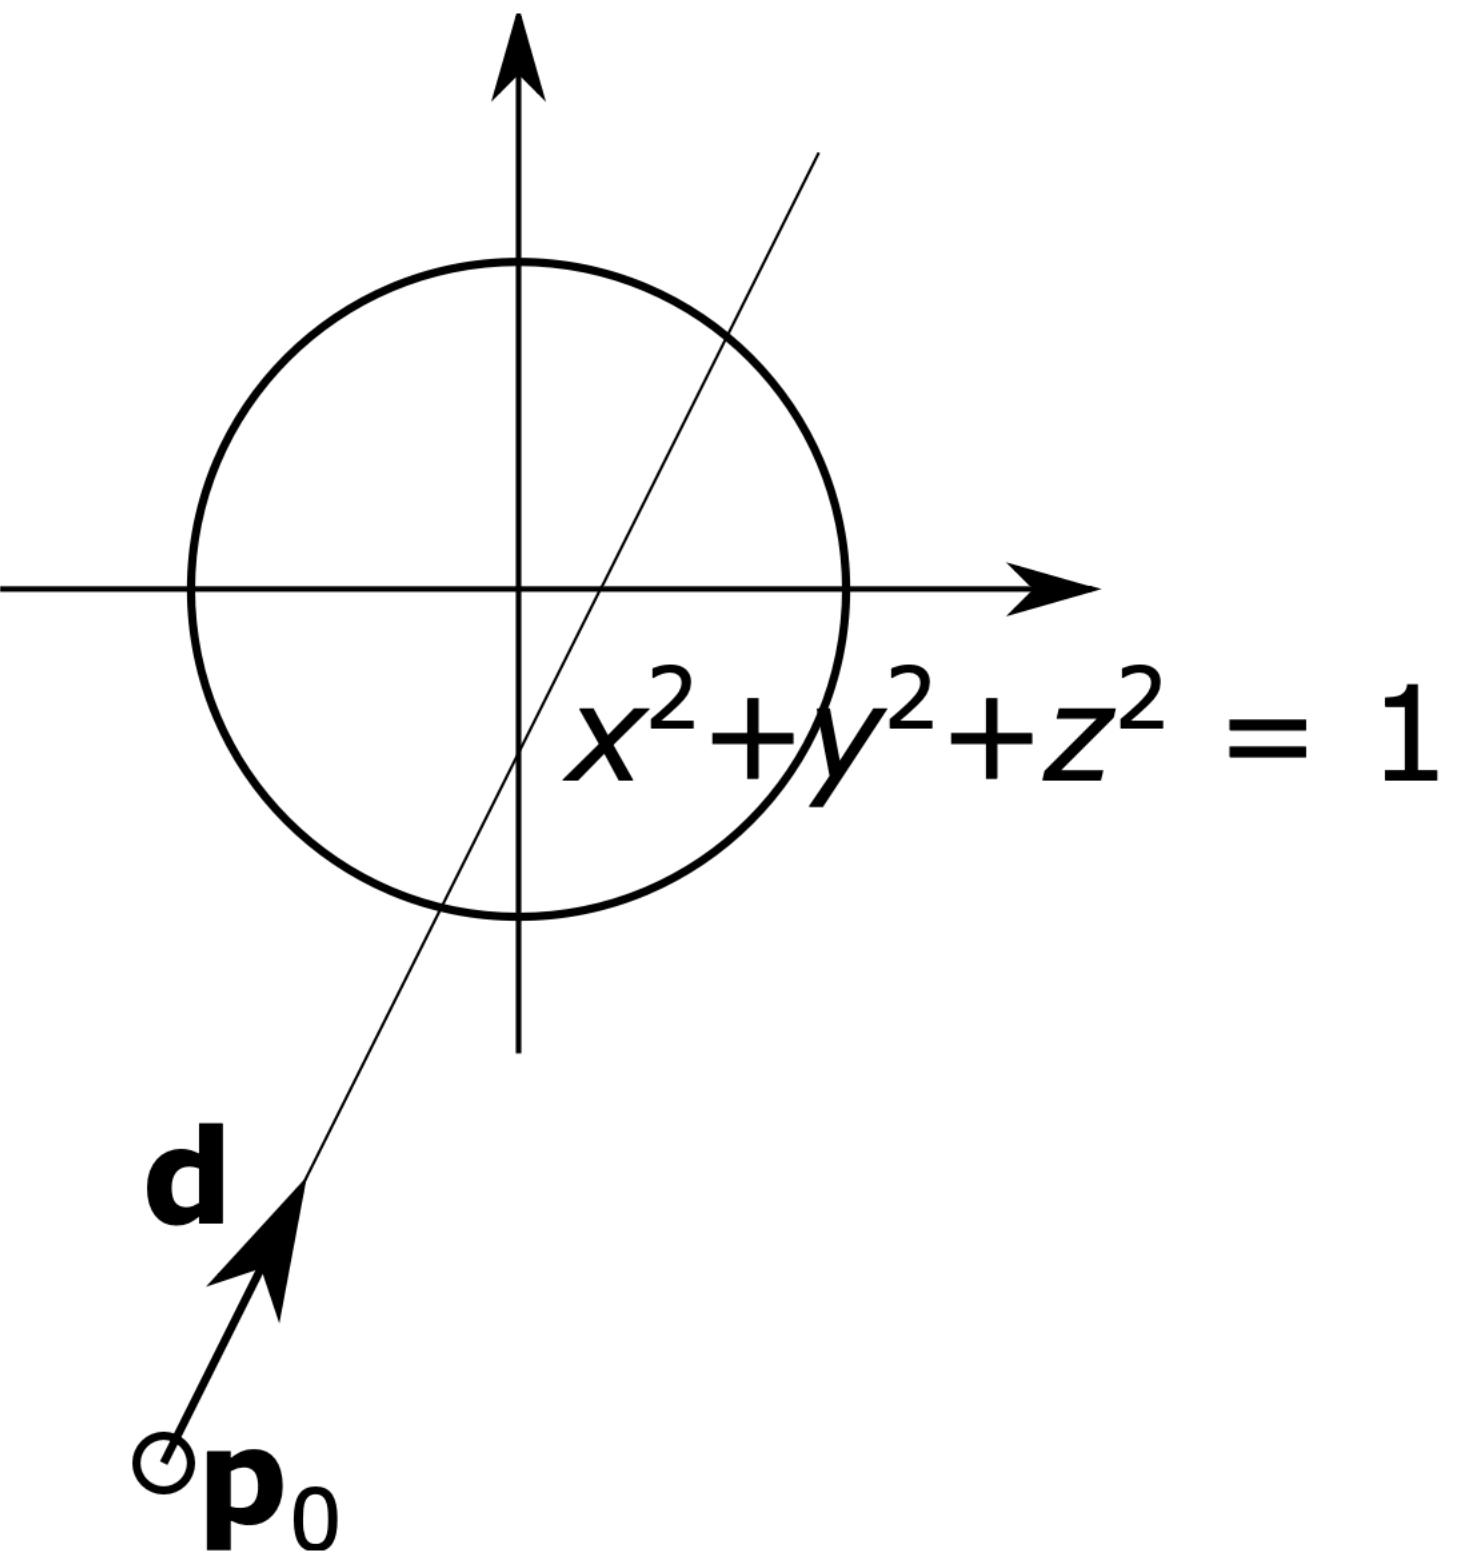
Ray-Sphere Intersections

- This is a quadratic equation in
 $a\lambda^2 + b\lambda + c = 0$
- So the solutions are
$$\lambda = \frac{-b \pm \sqrt{b^2 - 4ac}}{2a}$$
- 0, 1, or 2 real solutions (what do they mean)



Ray-Sphere Intersections

- 0 solutions:
 - The ray passes next to the sphere
 - No intersection
 - No shading at this pixel
 - This is the no-hit case
- 1 solution:
 - The ray just touches the sphere at exactly one point
 - This point is called a tangent point
- 2 solutions:
 - The ray hits the sphere on the way in → first intersection: t_1
 - The ray hits again on the way out → second intersection: t_2
 - These two points are where the ray cuts through the sphere



Ray-Sphere Intersections

- What if the sphere is not at the origin?
 - We can represent this by a translation
 - More generally we can have any transformation, T
 - Represented as a 4×4 homogeneous matrix
- Easier to apply the transform to the ray
 - Apply the inverse transform, T^{-1} to the ray

Ray-Cube Intersections

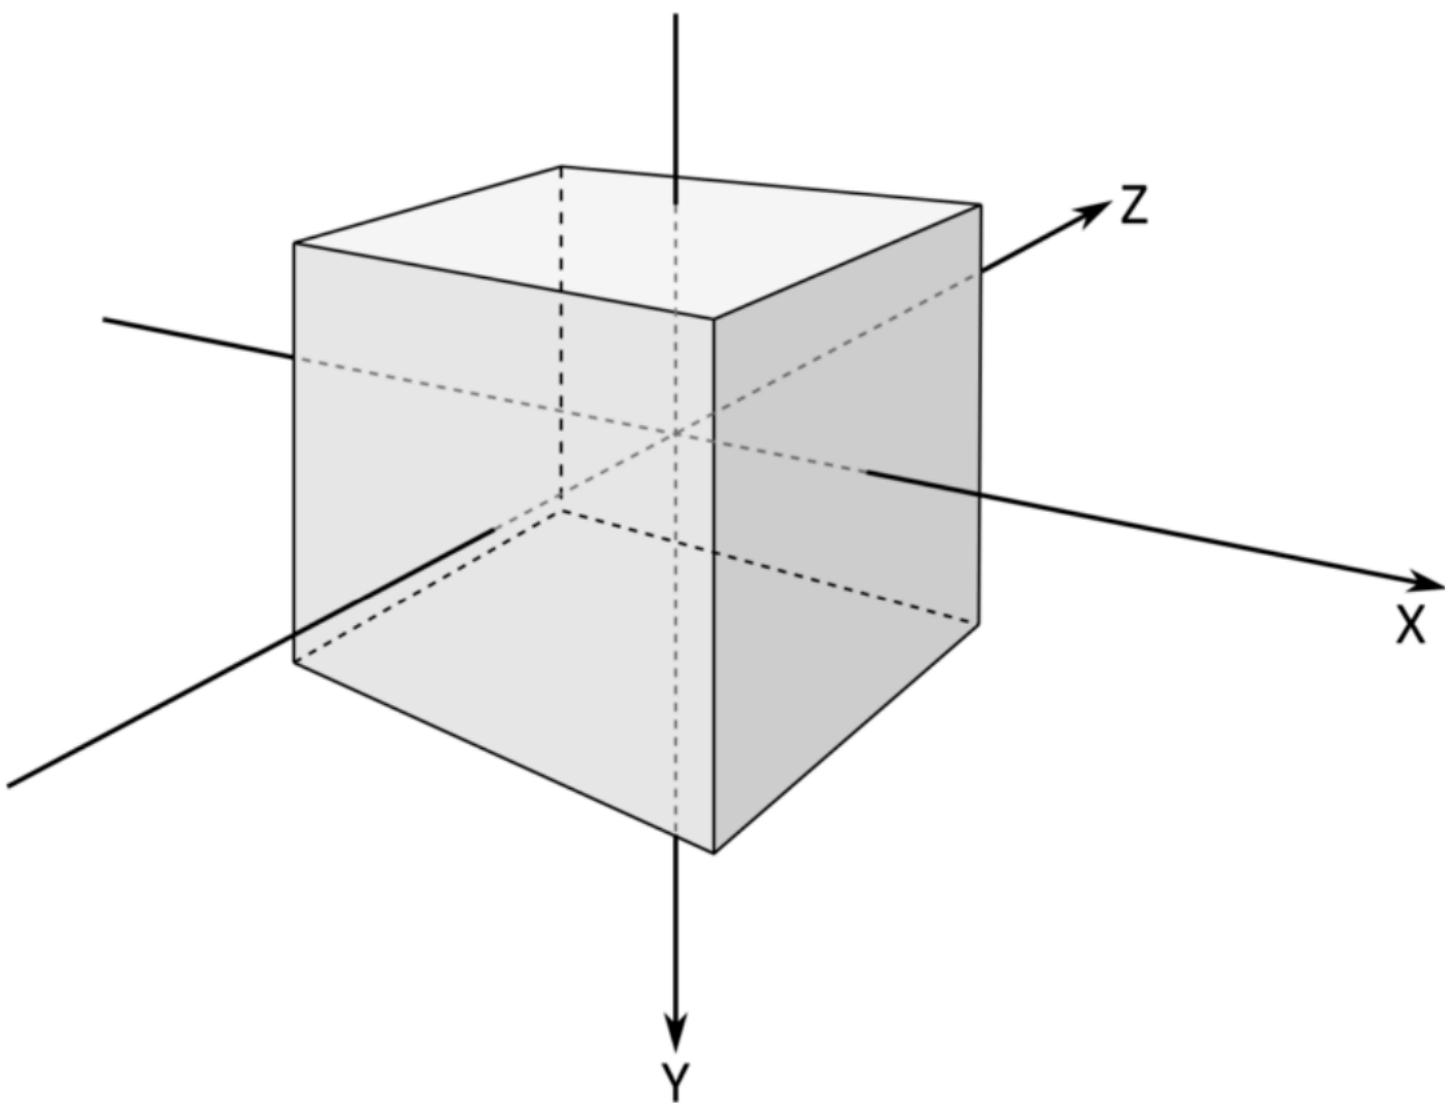
- A unit cube
 - From -1 to +1 on each axis
 - Intersect ray with each face
- X coordinate of a ray is

$$x_0 + \lambda \Delta x$$

- So at the plane $X = 1$ we have

$$1 = x_0 + \lambda \Delta x$$

$$\rightarrow \text{Solve for } \lambda = \frac{1 - x_0}{\Delta x}$$



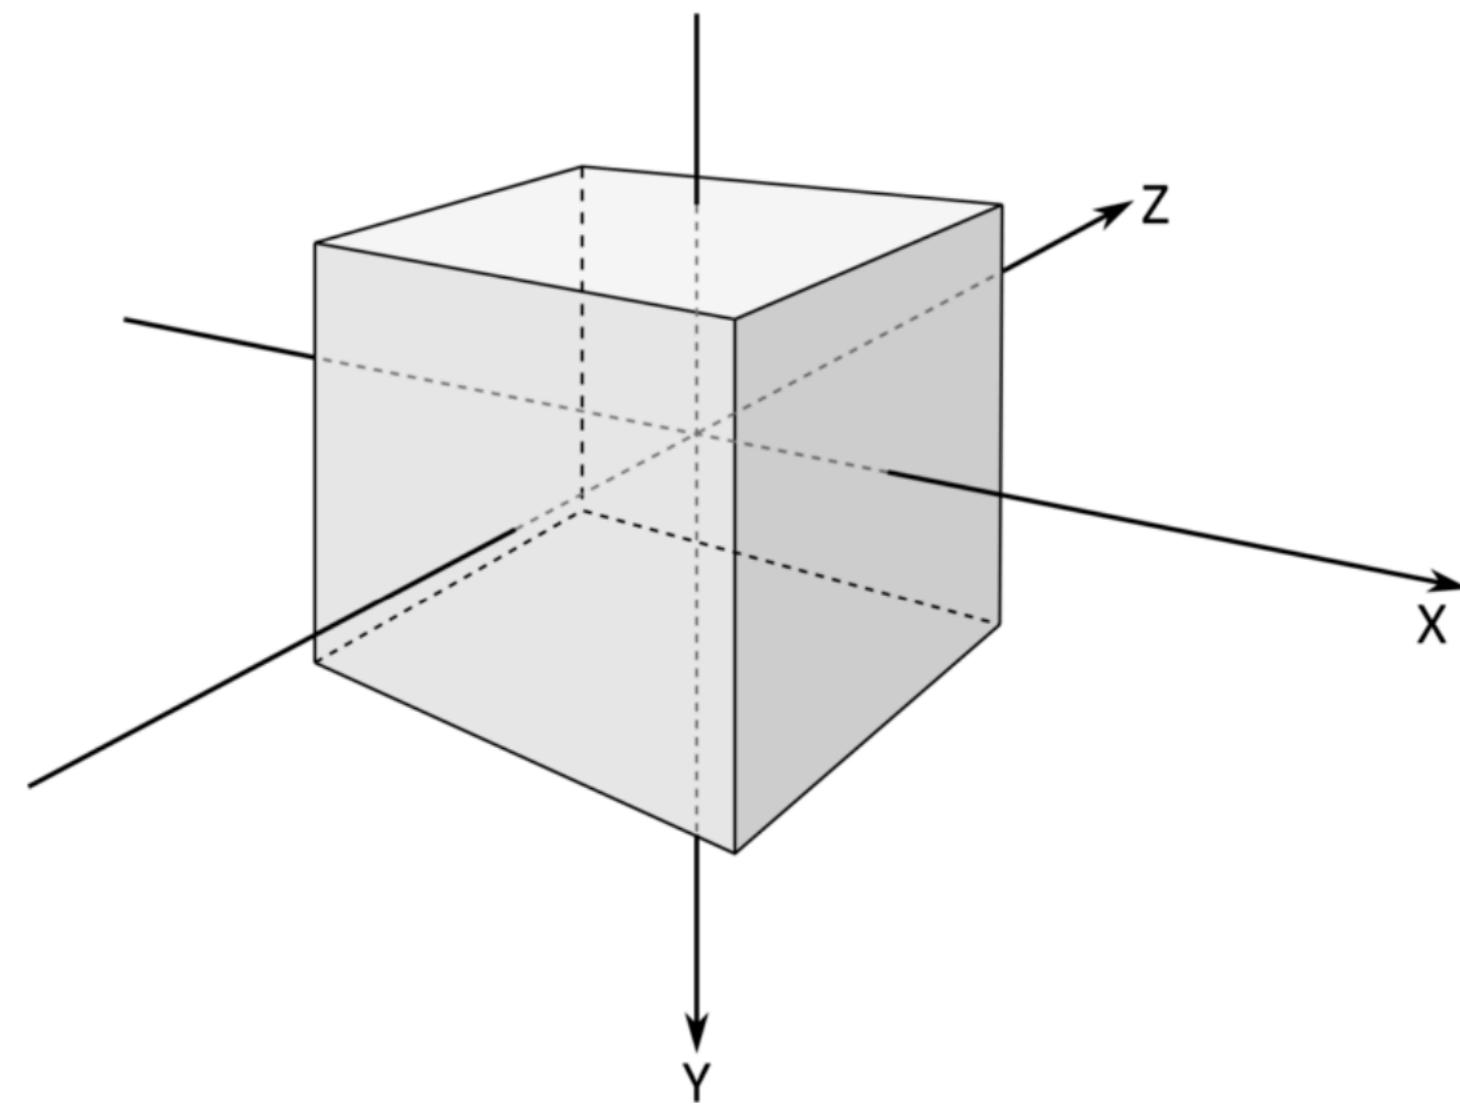
Ray-Cube Intersections

- Once we have λ , check whether we are in the face:

$$y_0 + \lambda \Delta y \in [-1, 1]$$

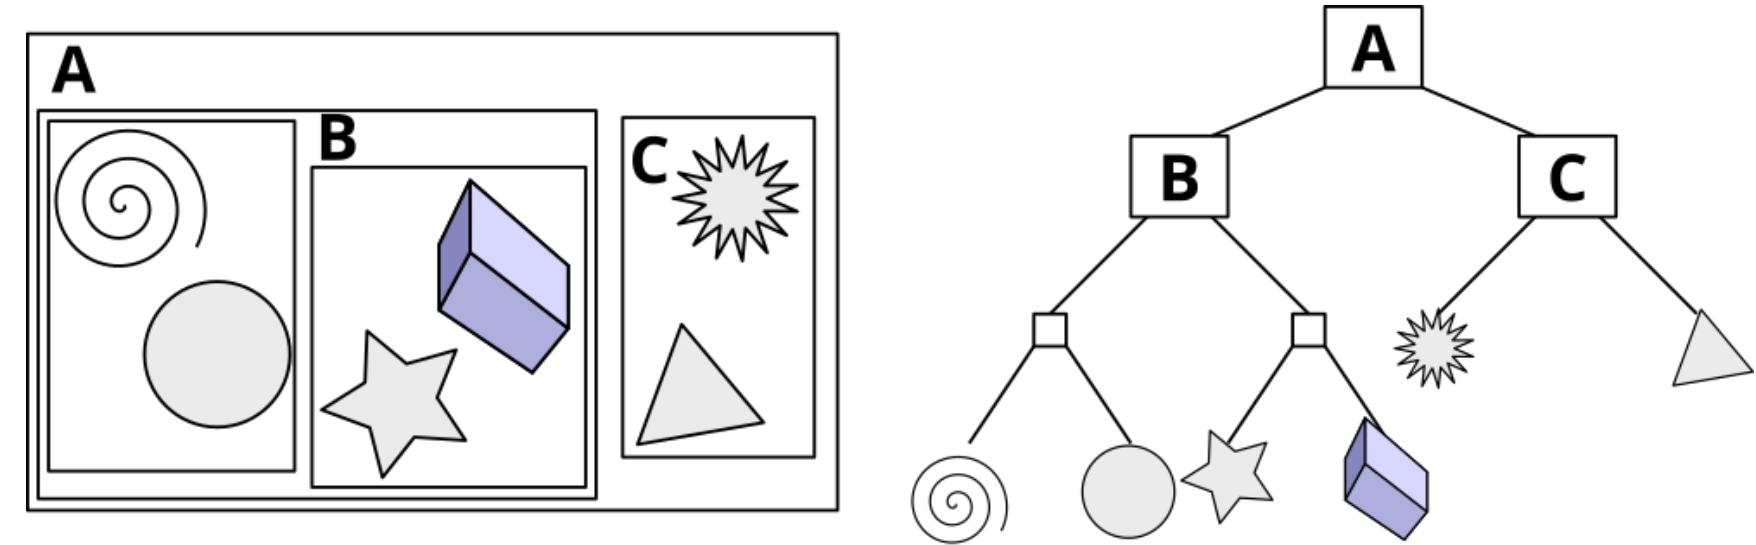
$$z_0 + \lambda \Delta z \in [-1, 1]$$

- Problems if $\Delta x = 0$:
-> parallel
- Repeat for all 6 faces

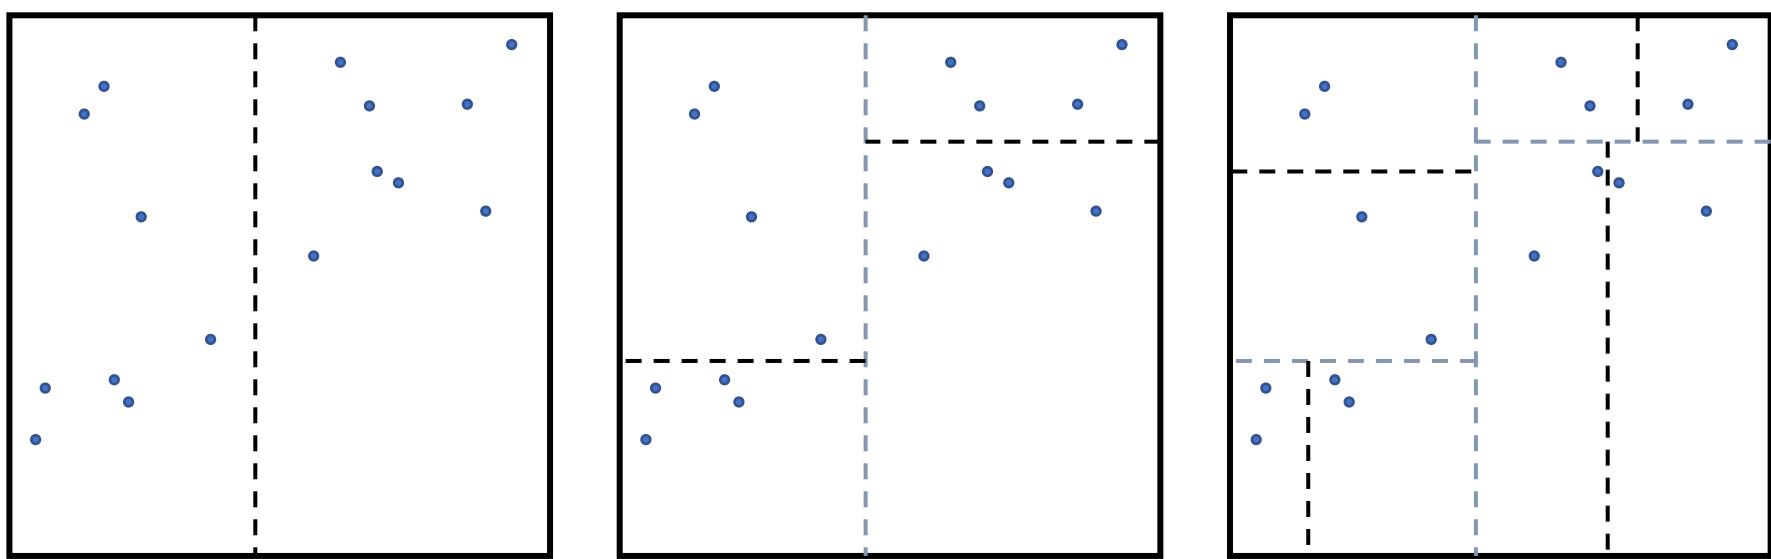


Acceleration Structures

- Reduce the number of ray-primitive tests
- Common structures:
 - BVH (Bounding Volume Hierarchy): widely used due to efficient construction and traversal.
 - KD-tree: effective for static scenes requiring high traversal performance
 - Uniform grids / octrees: beneficial for scenes with known spatial regularity
- Key trade-off: build time vs. traversal efficiency



Schreiberx, CC BY-SA 3.0 <<https://creativecommons.org/licenses/by-sa/3.0>>, via Wikimedia Commons



Real-time Raytracing

- Hardware acceleration (e.g., dedicated RT cores) -> significantly improved ray-tracing performance, enabling more efficient execution of complex lighting and visibility calculations
- AI-driven techniques are increasingly integrated into the ray-tracing pipeline, supporting tasks such as shading, denoising, and reconstruction with higher quality and stability
- Machine-learning-based denoisers have begun to replace traditional handcrafted approaches, providing cleaner images with fewer artefacts at lower sample counts
- Contemporary rendering systems now adopt hybrid pipelines, combining rasterisation, ray tracing, neural rendering, and volumetric methods to balance visual fidelity and performance:
 - Ray-traced shadows
 - Ray-traced reflections
 - Ray-traced ambient occlusion



RadFoam: Volumetric mesh ray tracing algorithm

The end!