

# Visual Computing I:

Interactive Computer Graphics and Vision

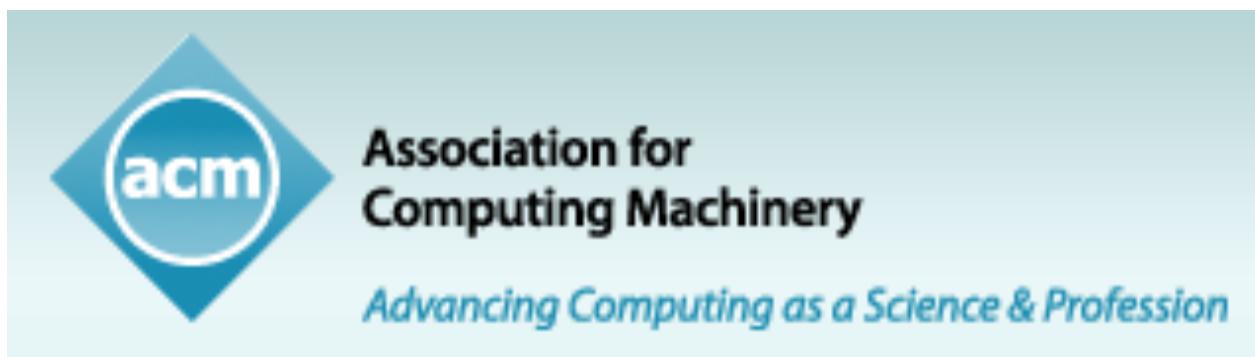


Digital Images and Colour Representation

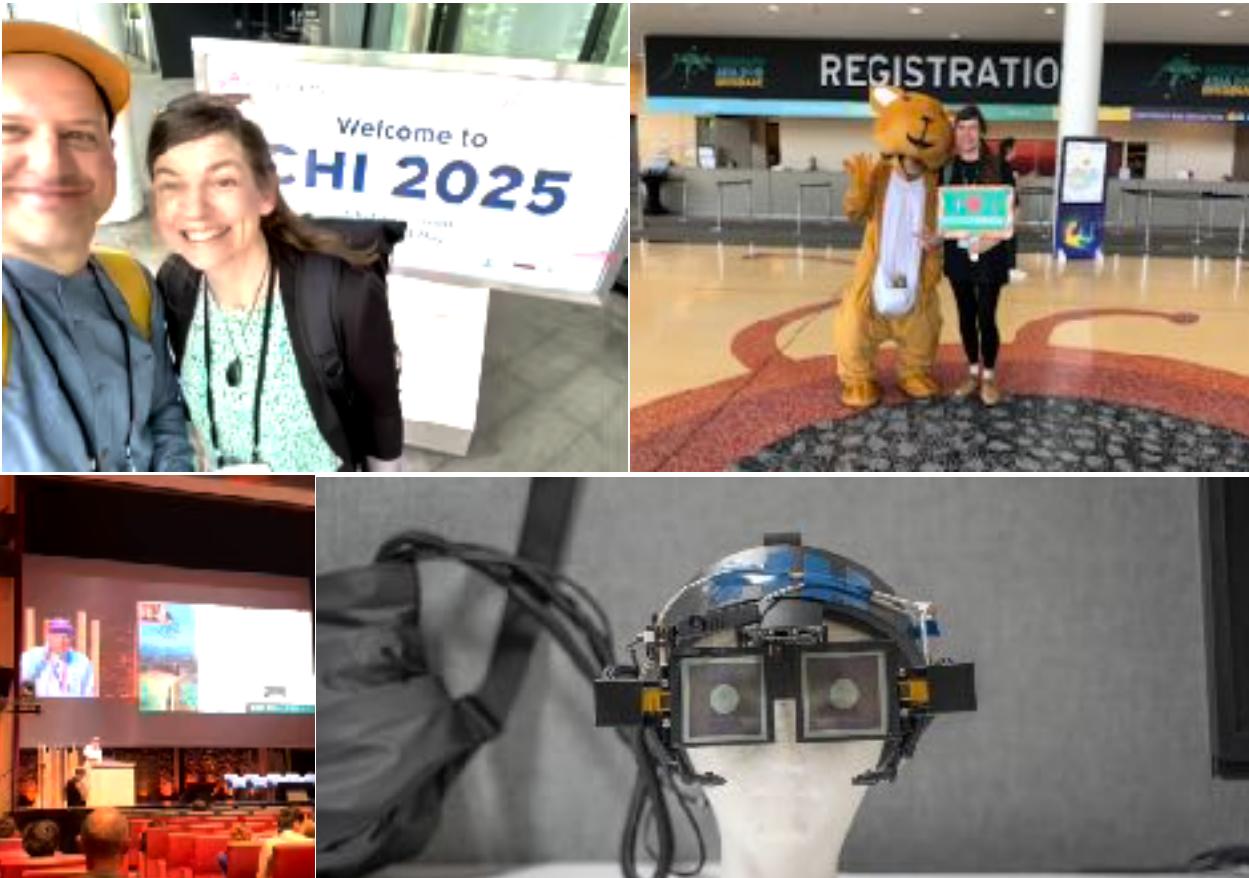
Tobias Langlotz

**Last time**

# Design of the course



Expert Societies



Academic Background



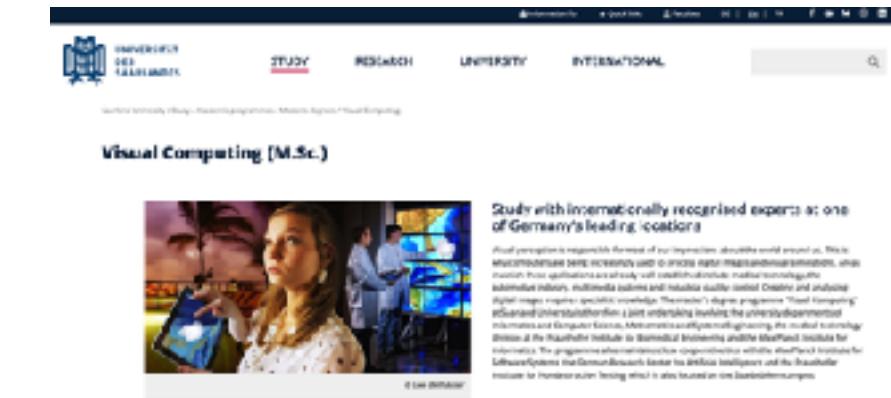
Industry Experience and Examples



Visual Computing

Semester: Autumn 2024

Catalogue Link: [252-0209-00L](#)

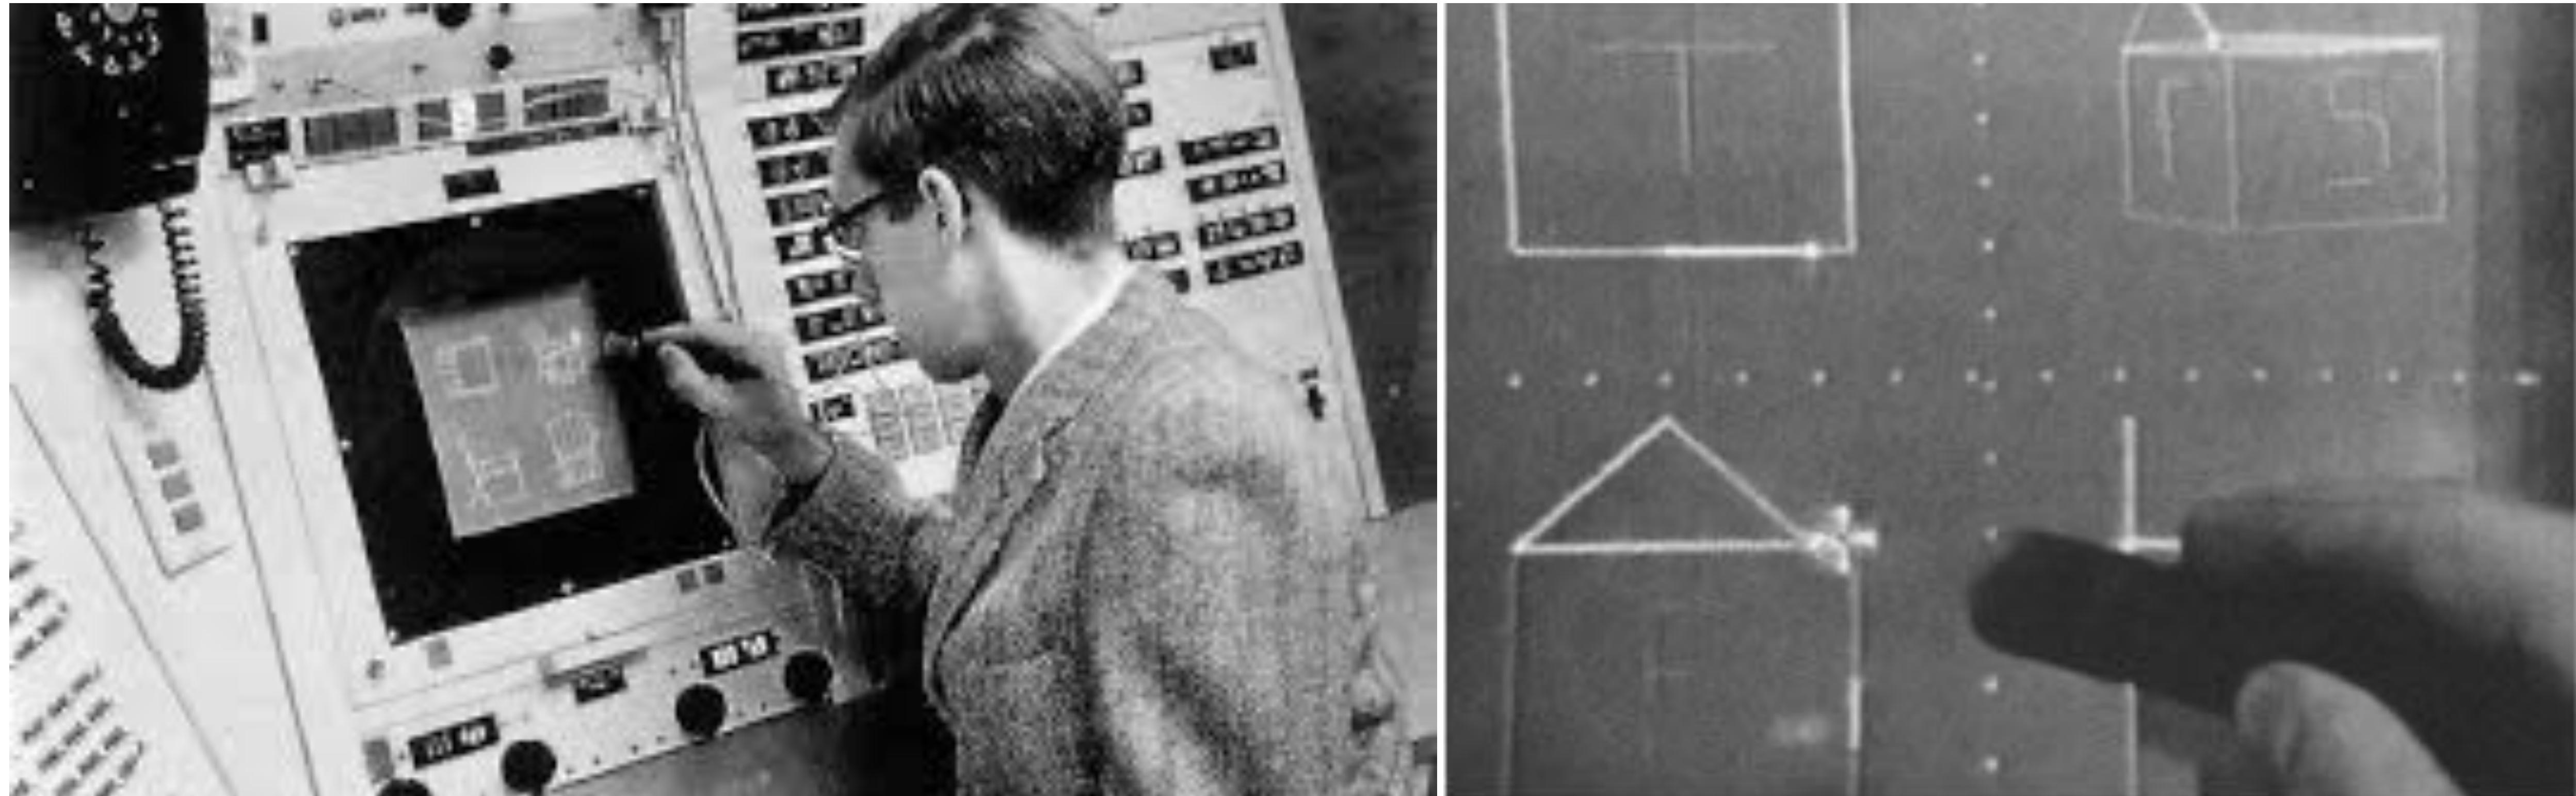


Similar Courses



Literature

# A Short History of Visual Computing

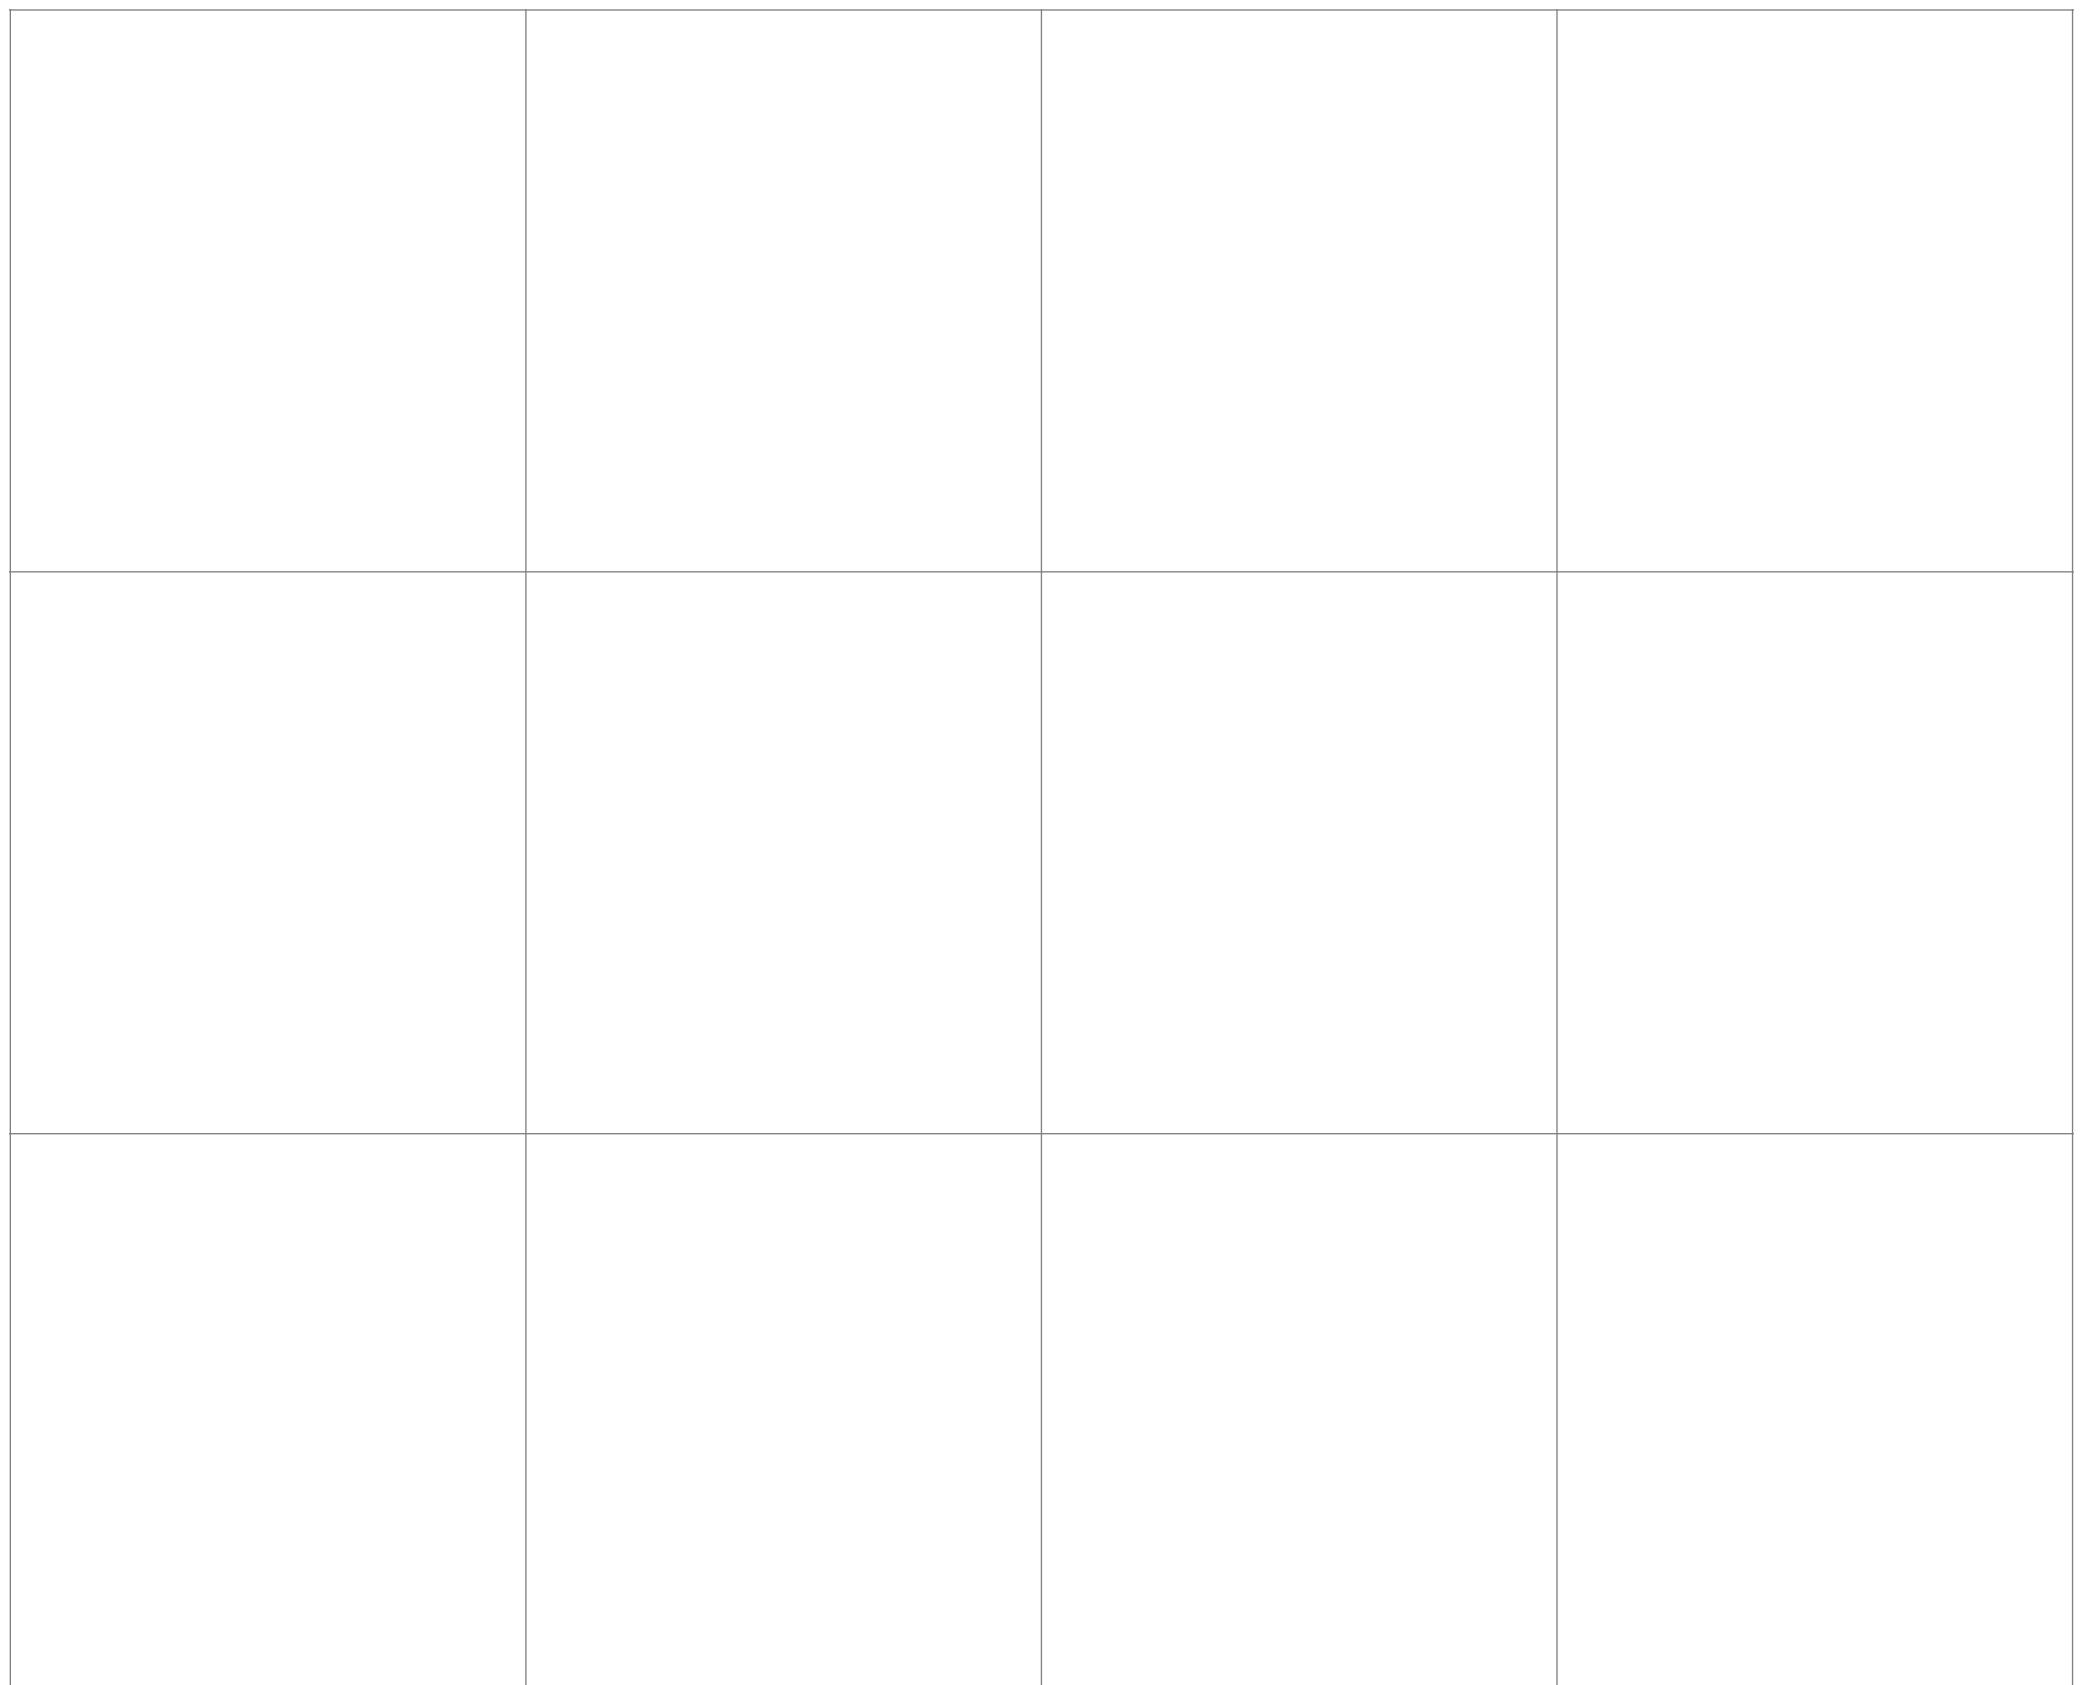


Ivan Sutherland, *Sketchpad, a Man–Machine Graphical Communication System*, 1963

# **Digital Images and 2D Image Representation**

# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

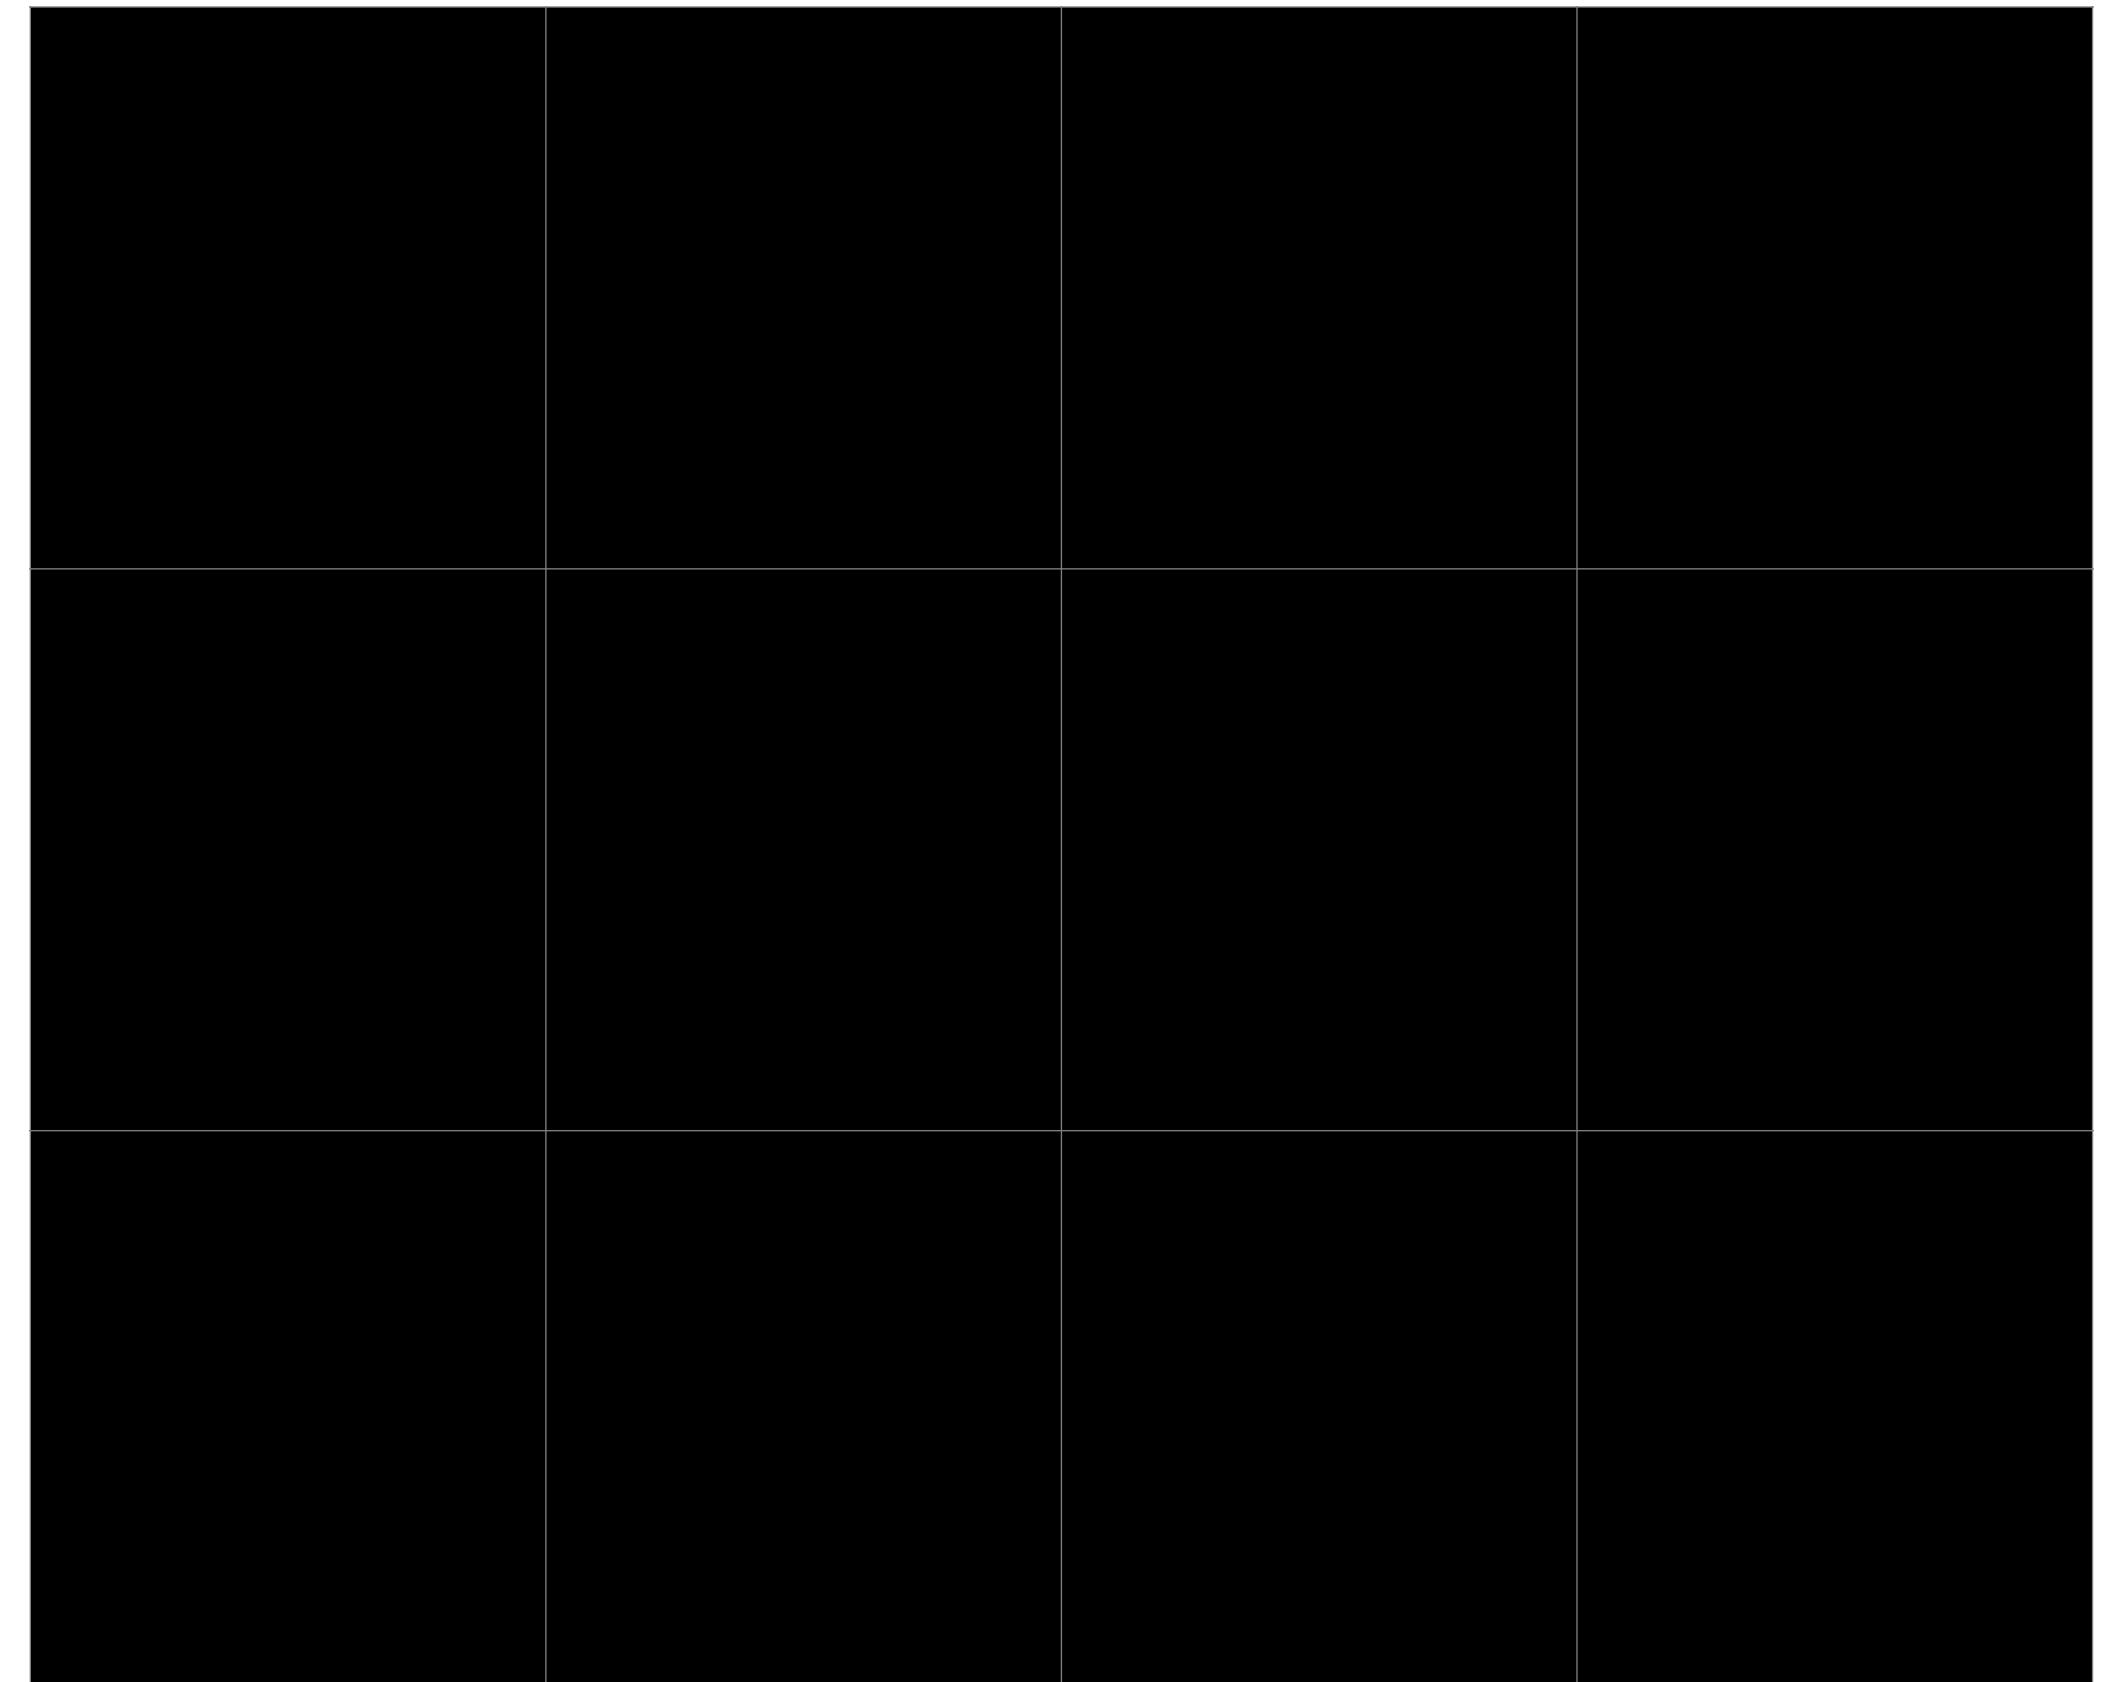


# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125
```



# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125
```



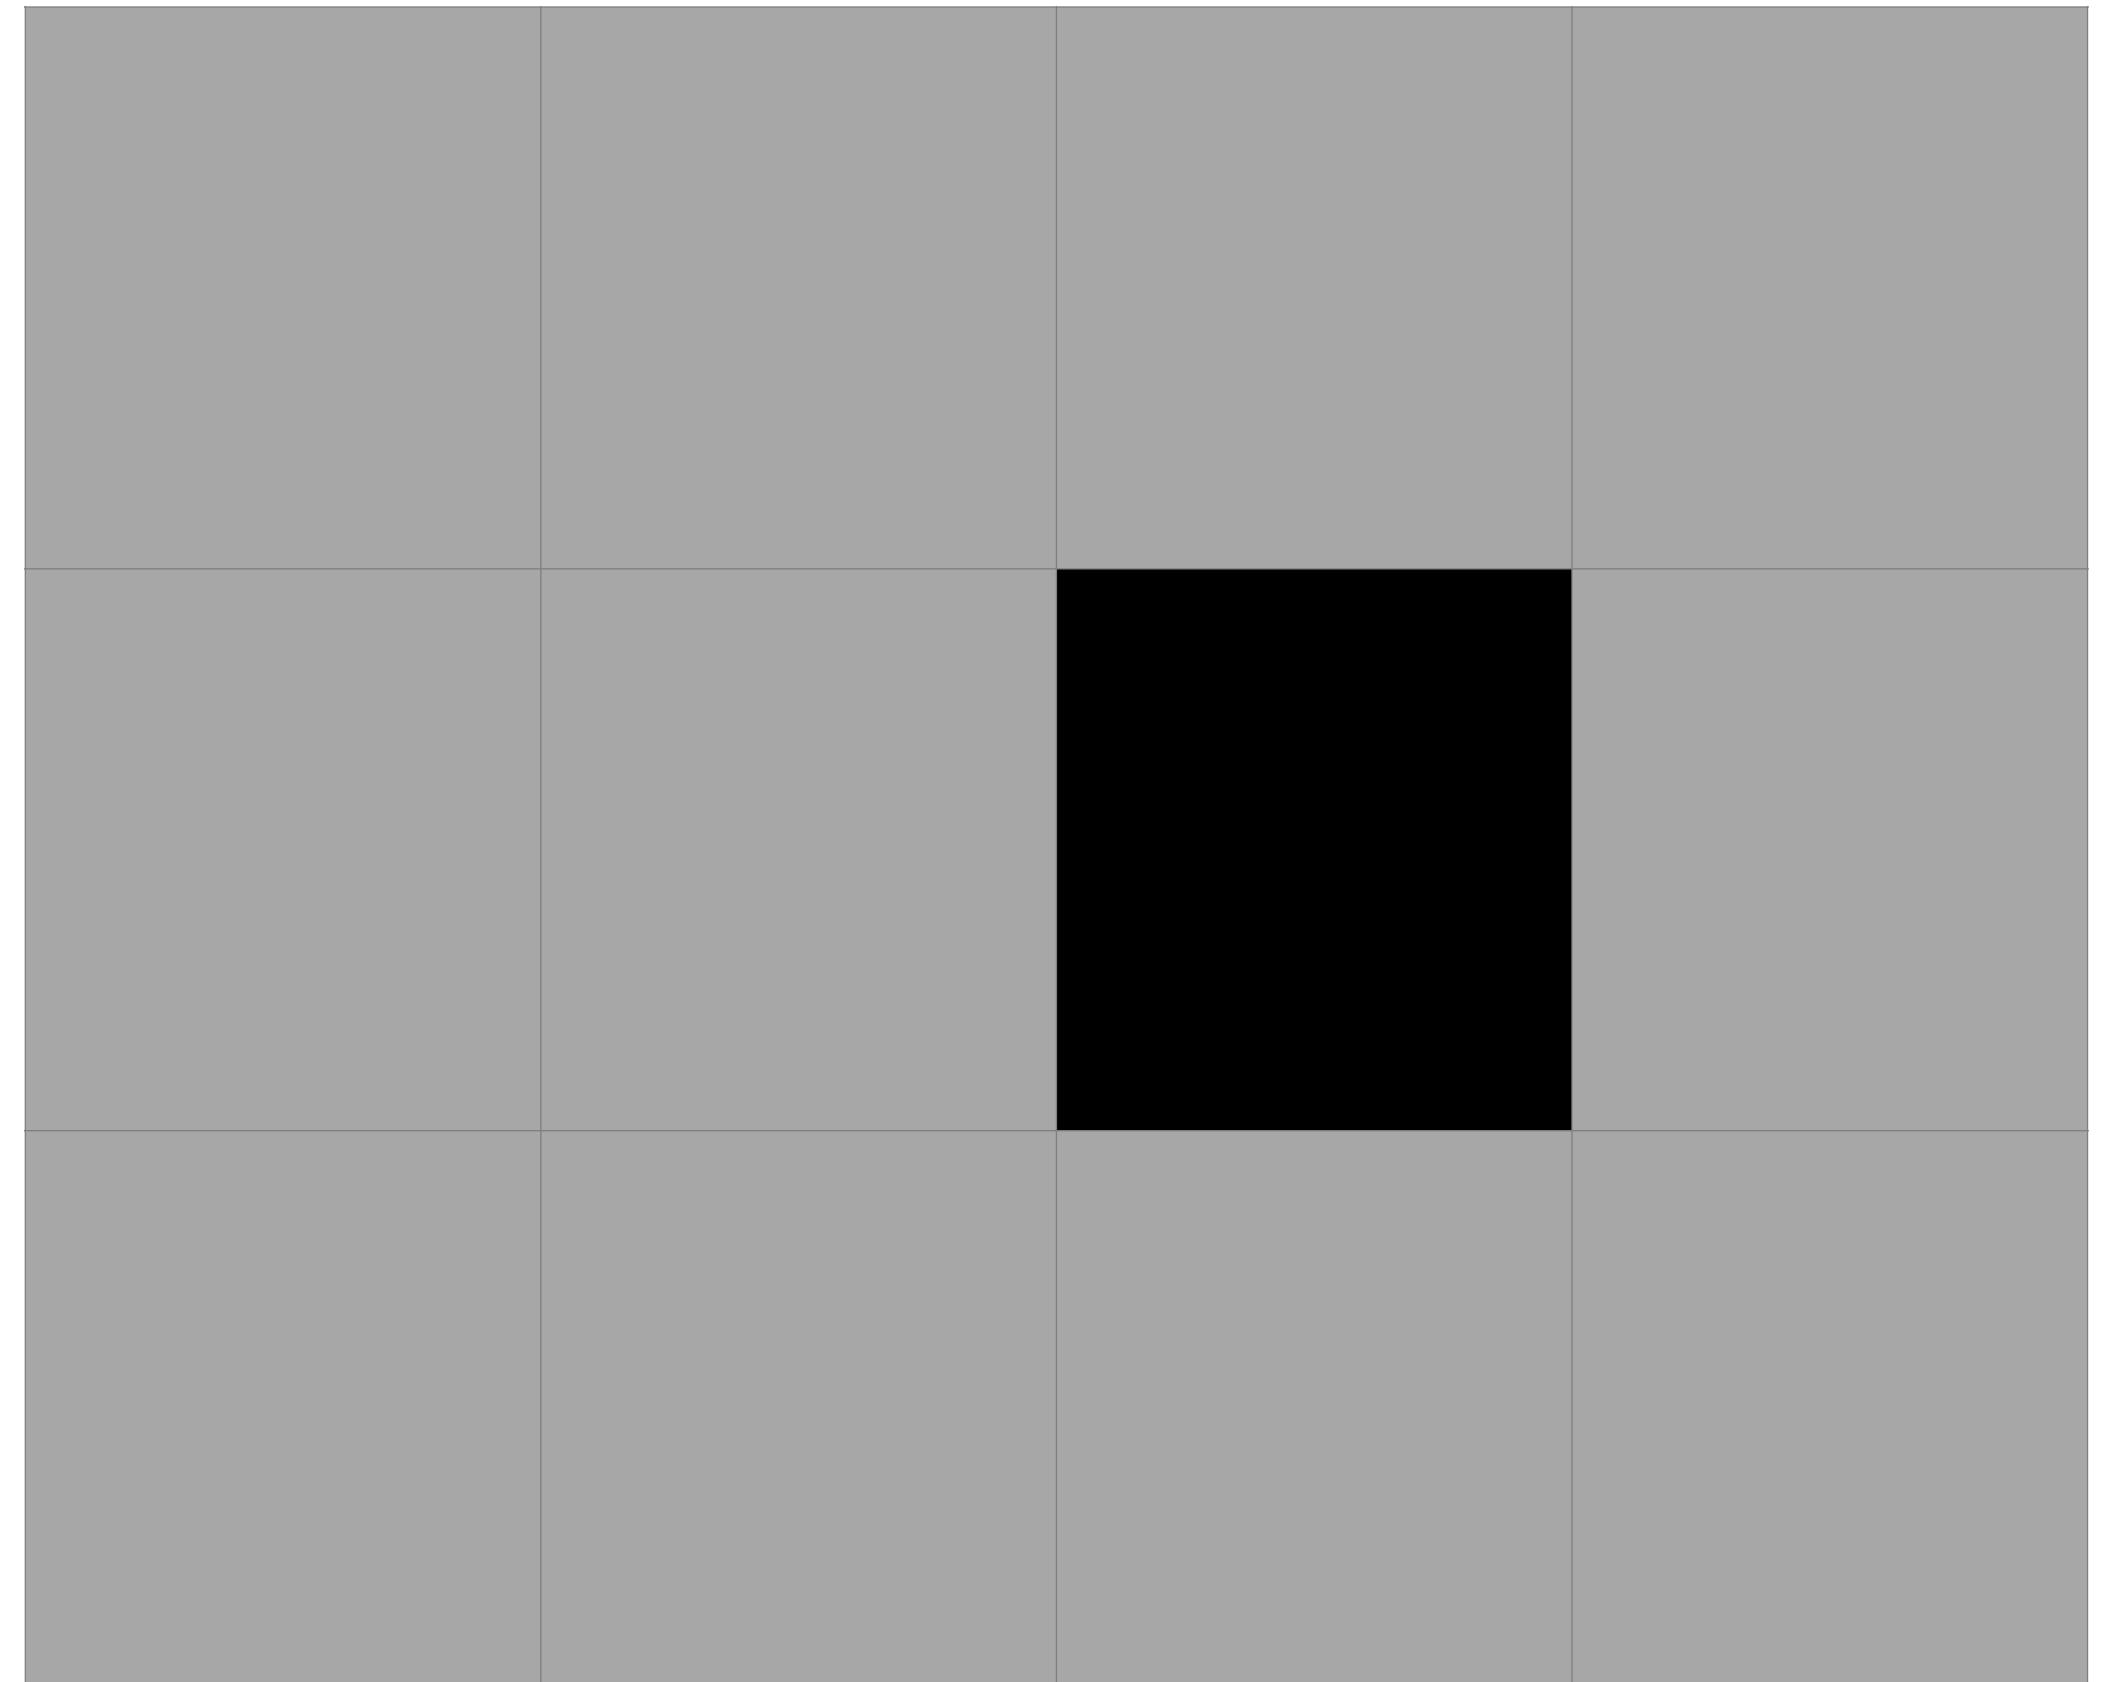
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125

# Access pixel at (row=2, col=3)
image2D[1][2]= 0
```



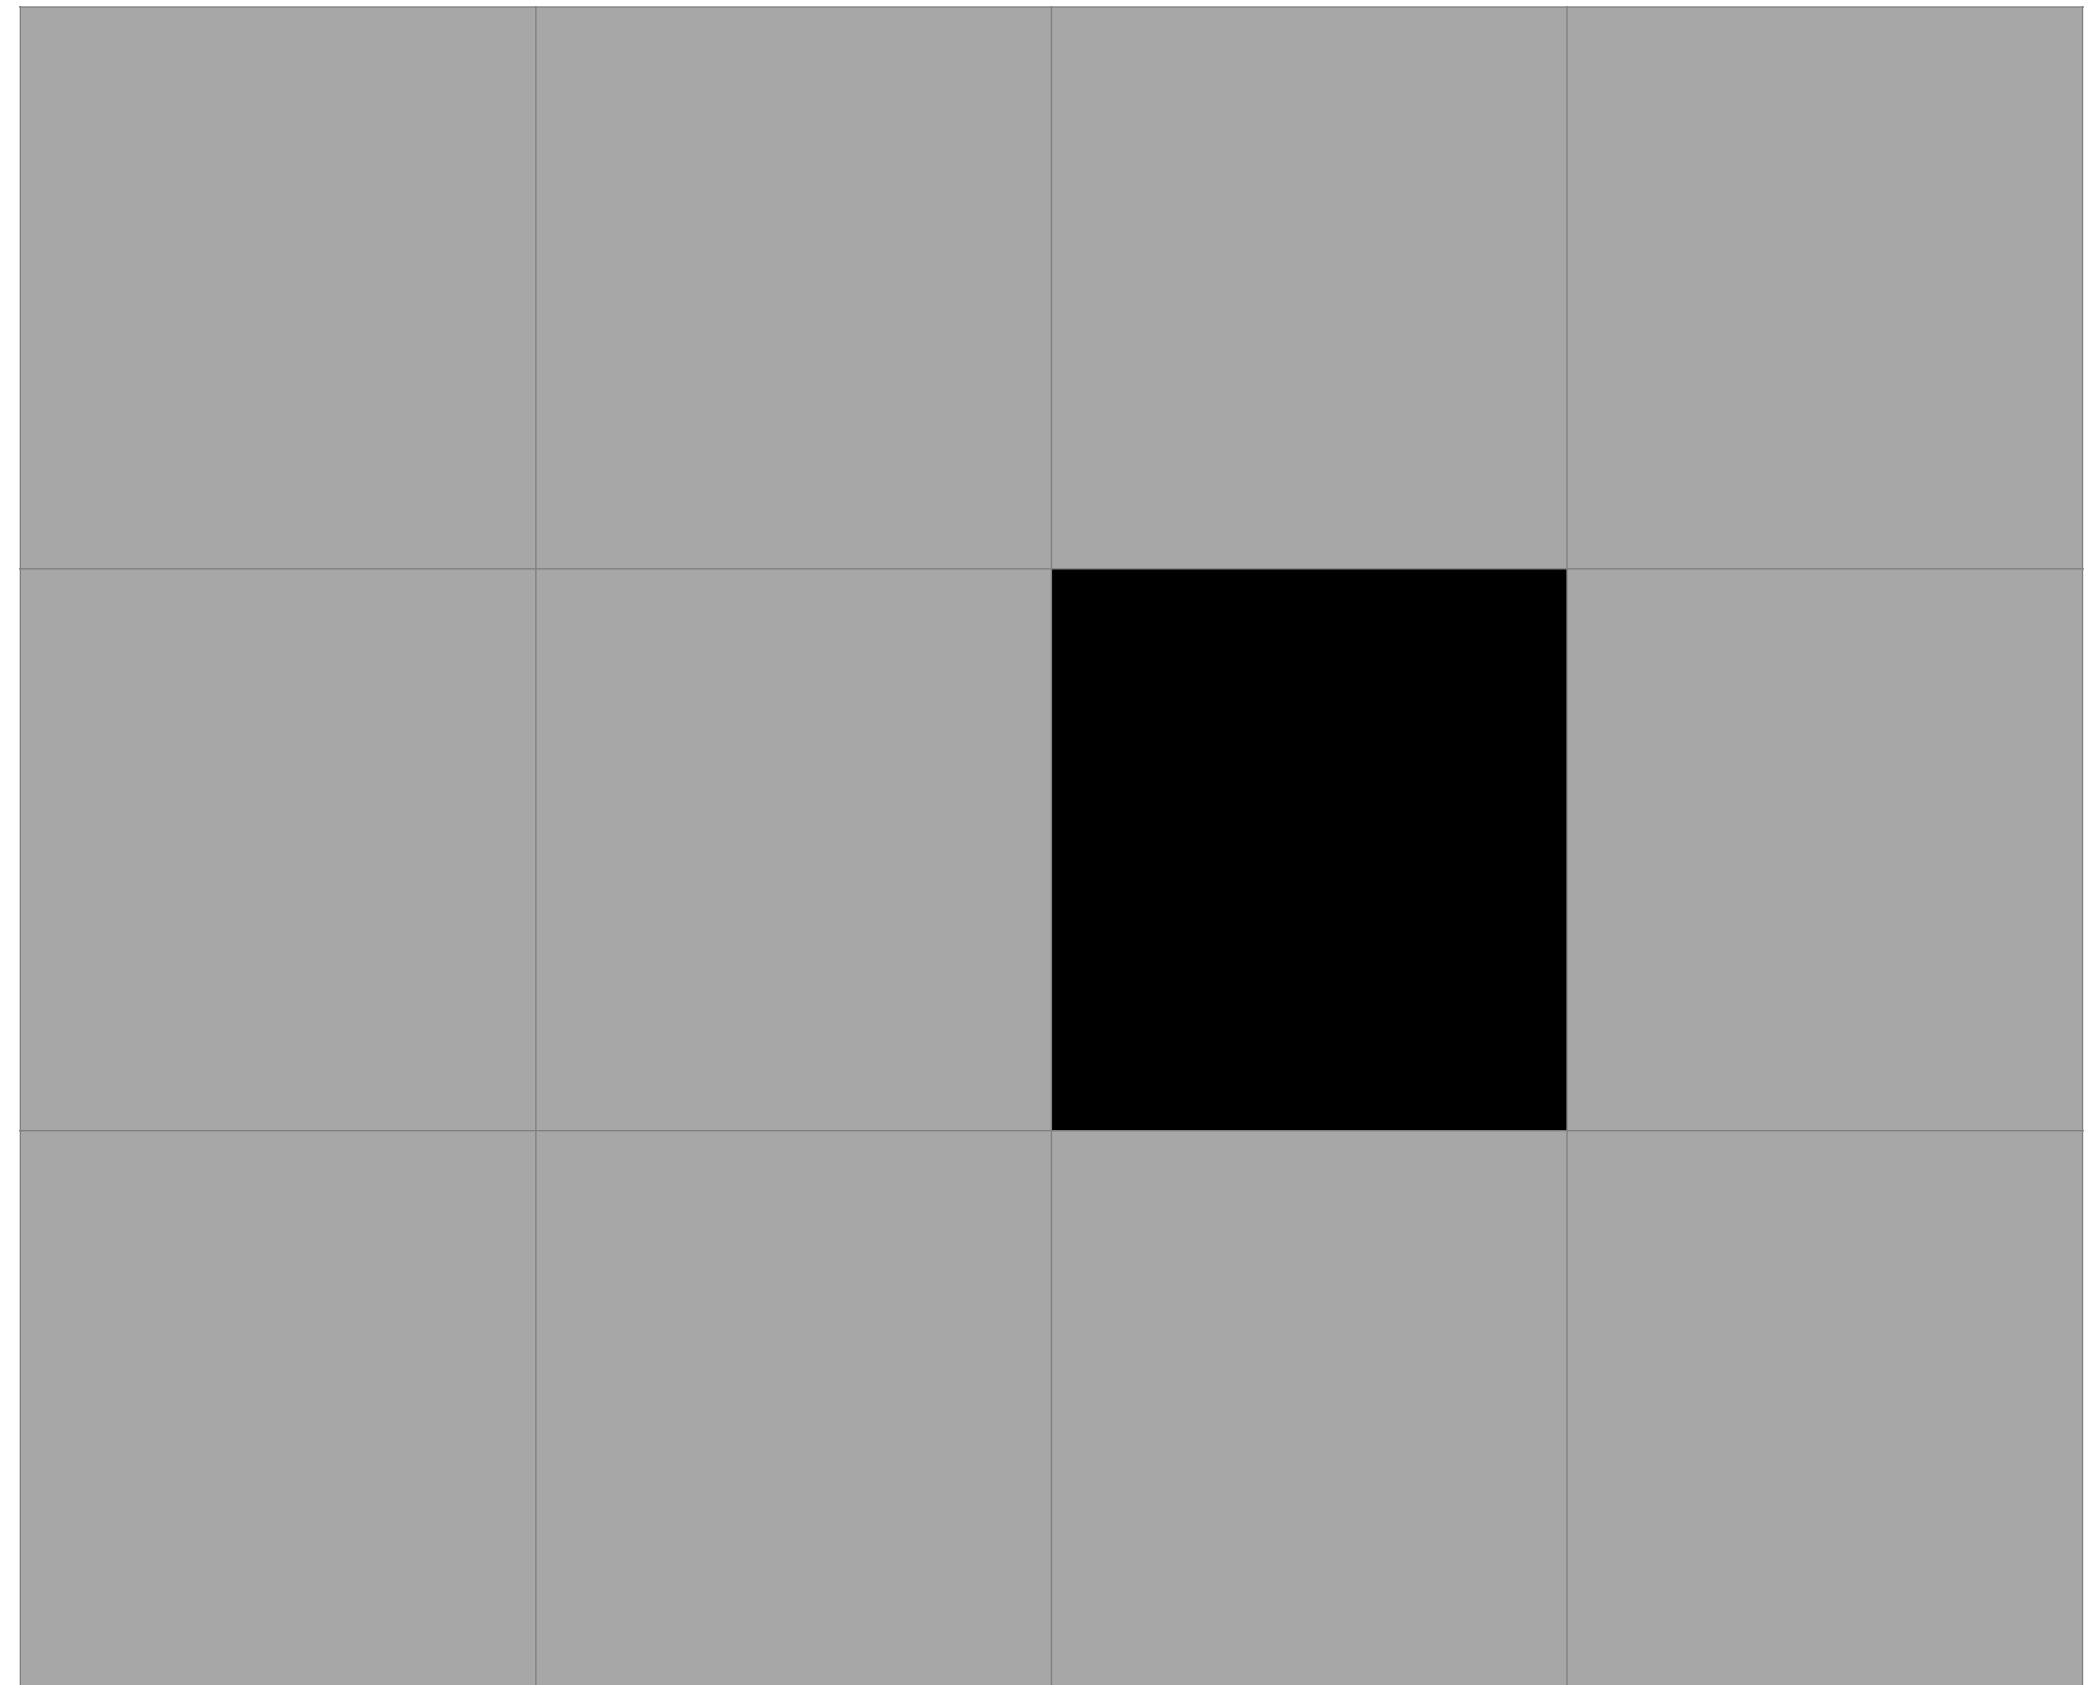
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125

image2D =
[
    [ 0, 1, 2, 3 ],      # row 0
    [ 0, 1, 2, 3 ],      # row 1
    [ 0, 1, 2, 3 ]       # row 2
]
```



# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 1D array
image1D = create_array(width * height)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        index = (y * width) + x
        image1D[index] = 125
```



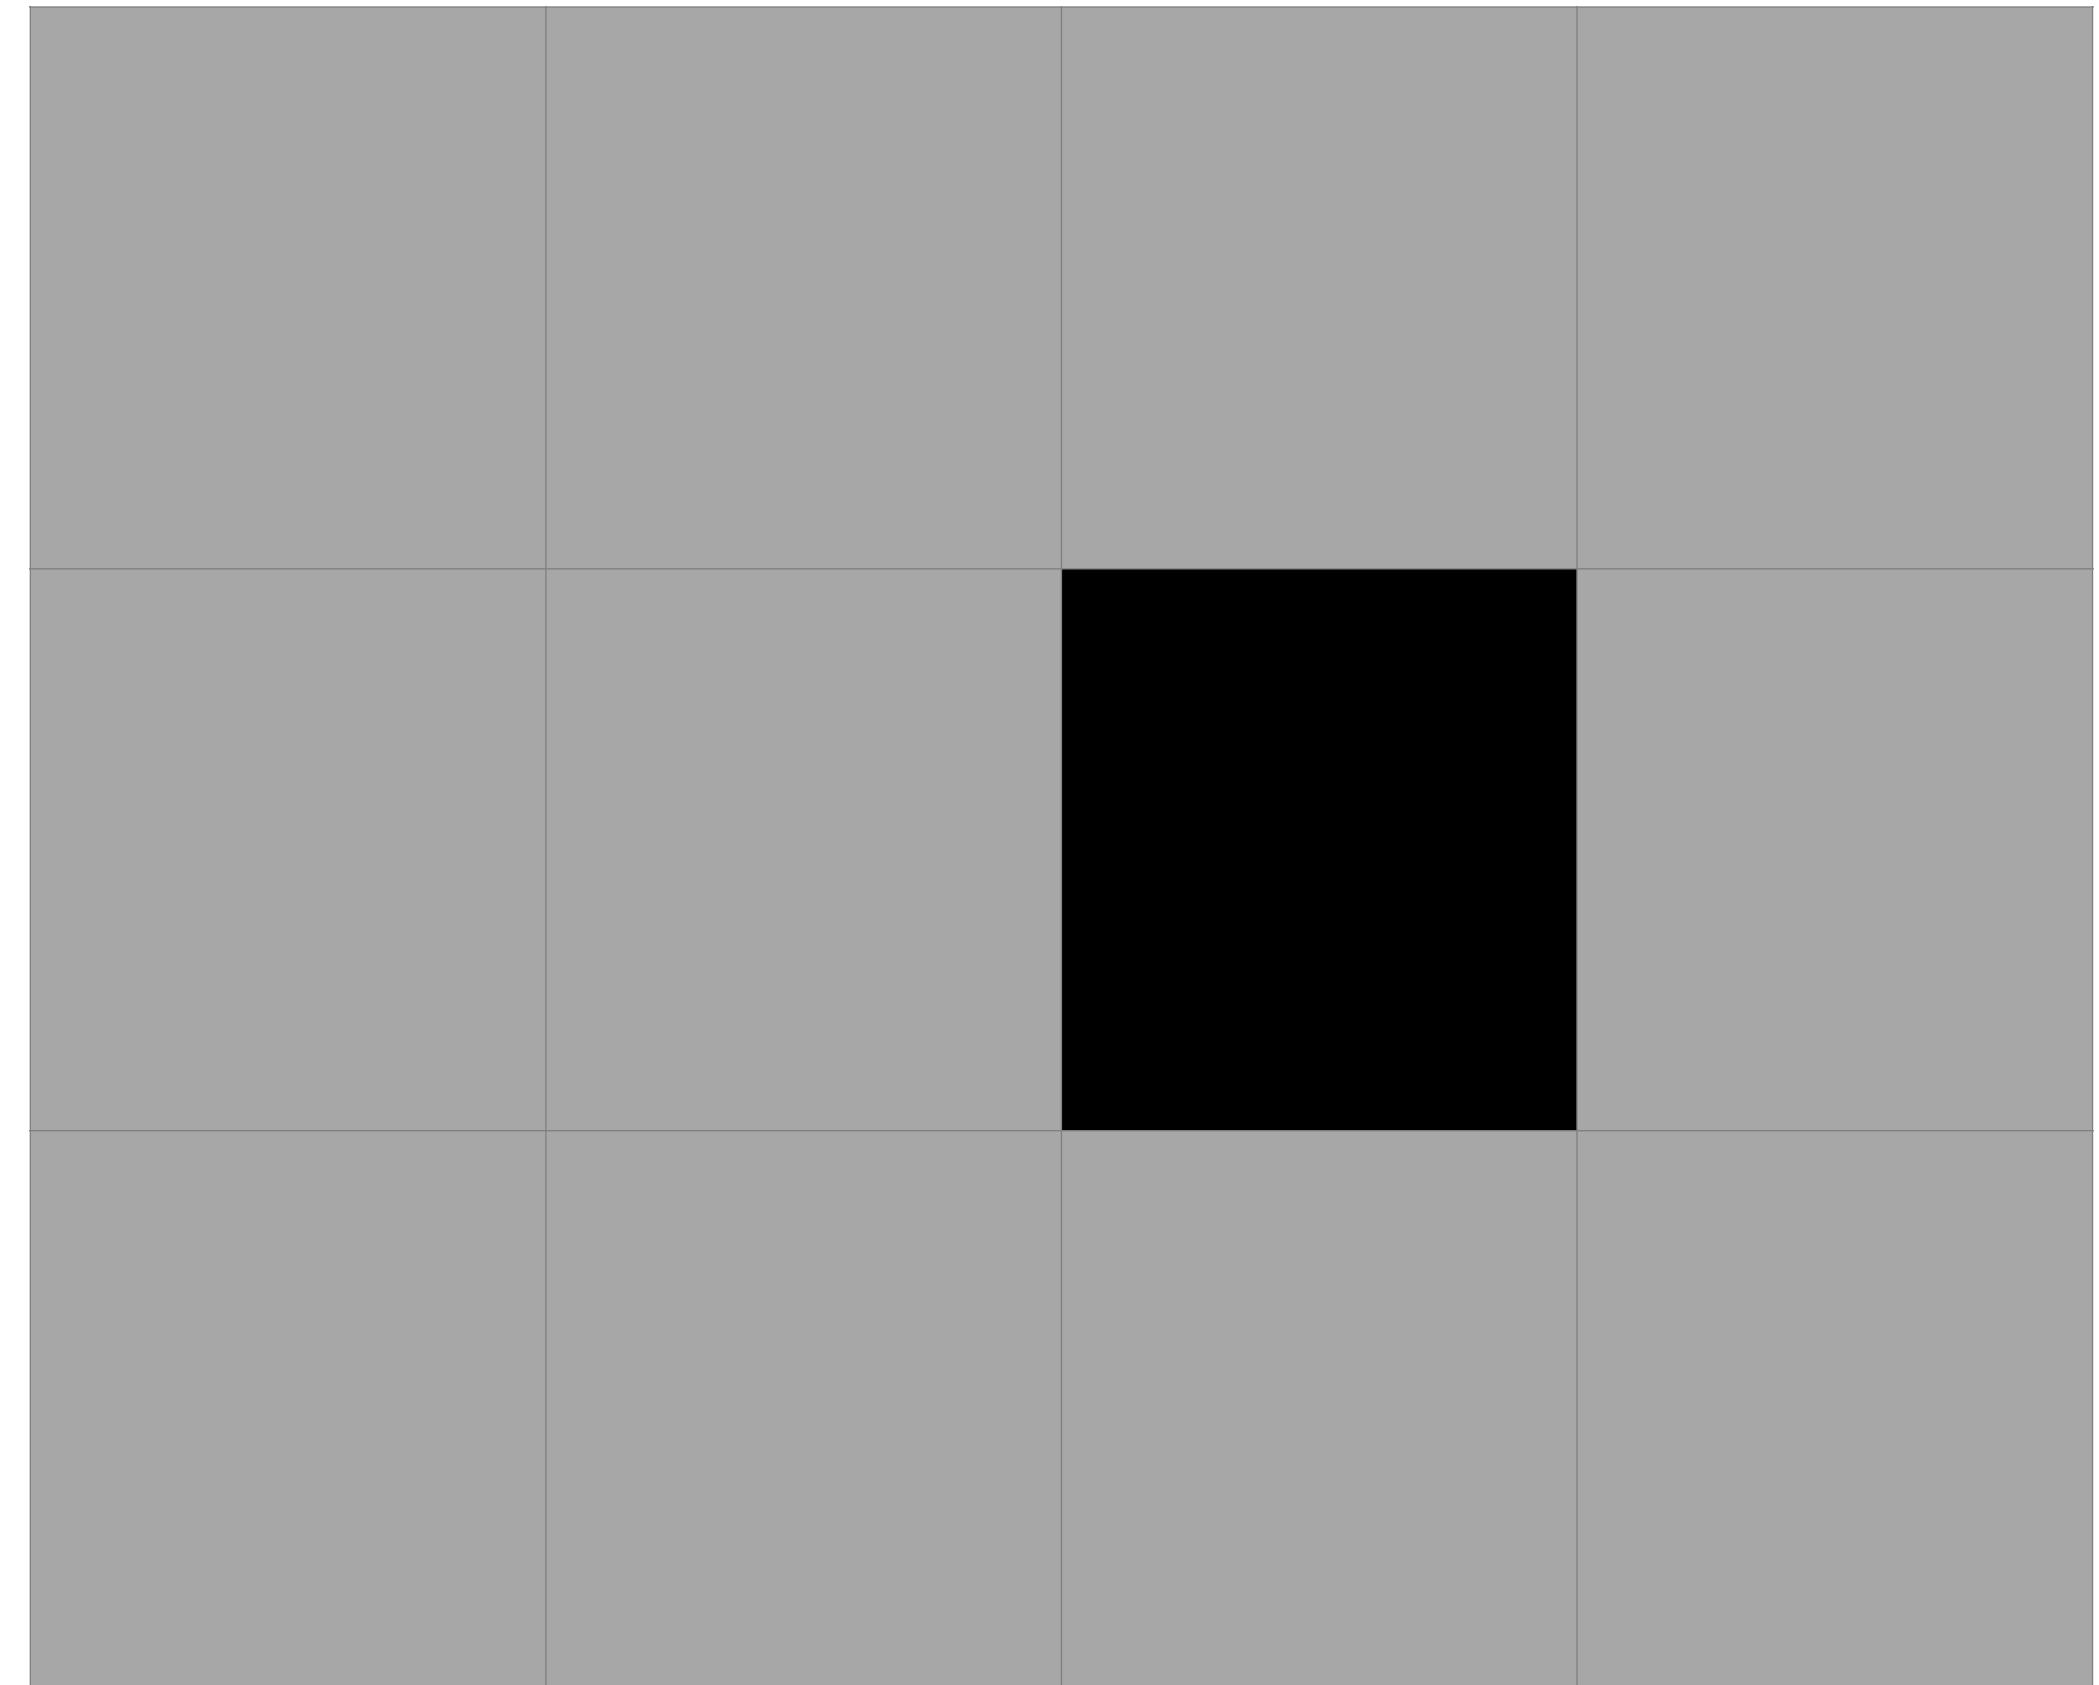
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 1D array
image1D = create_array(width * height)
```

```
# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        index = (y * width) + x
        image1D[index] = 125
```

```
# Access pixel at (row=2, col=3)
index = (row-1 * width) + col
image1D[index] = 0
```



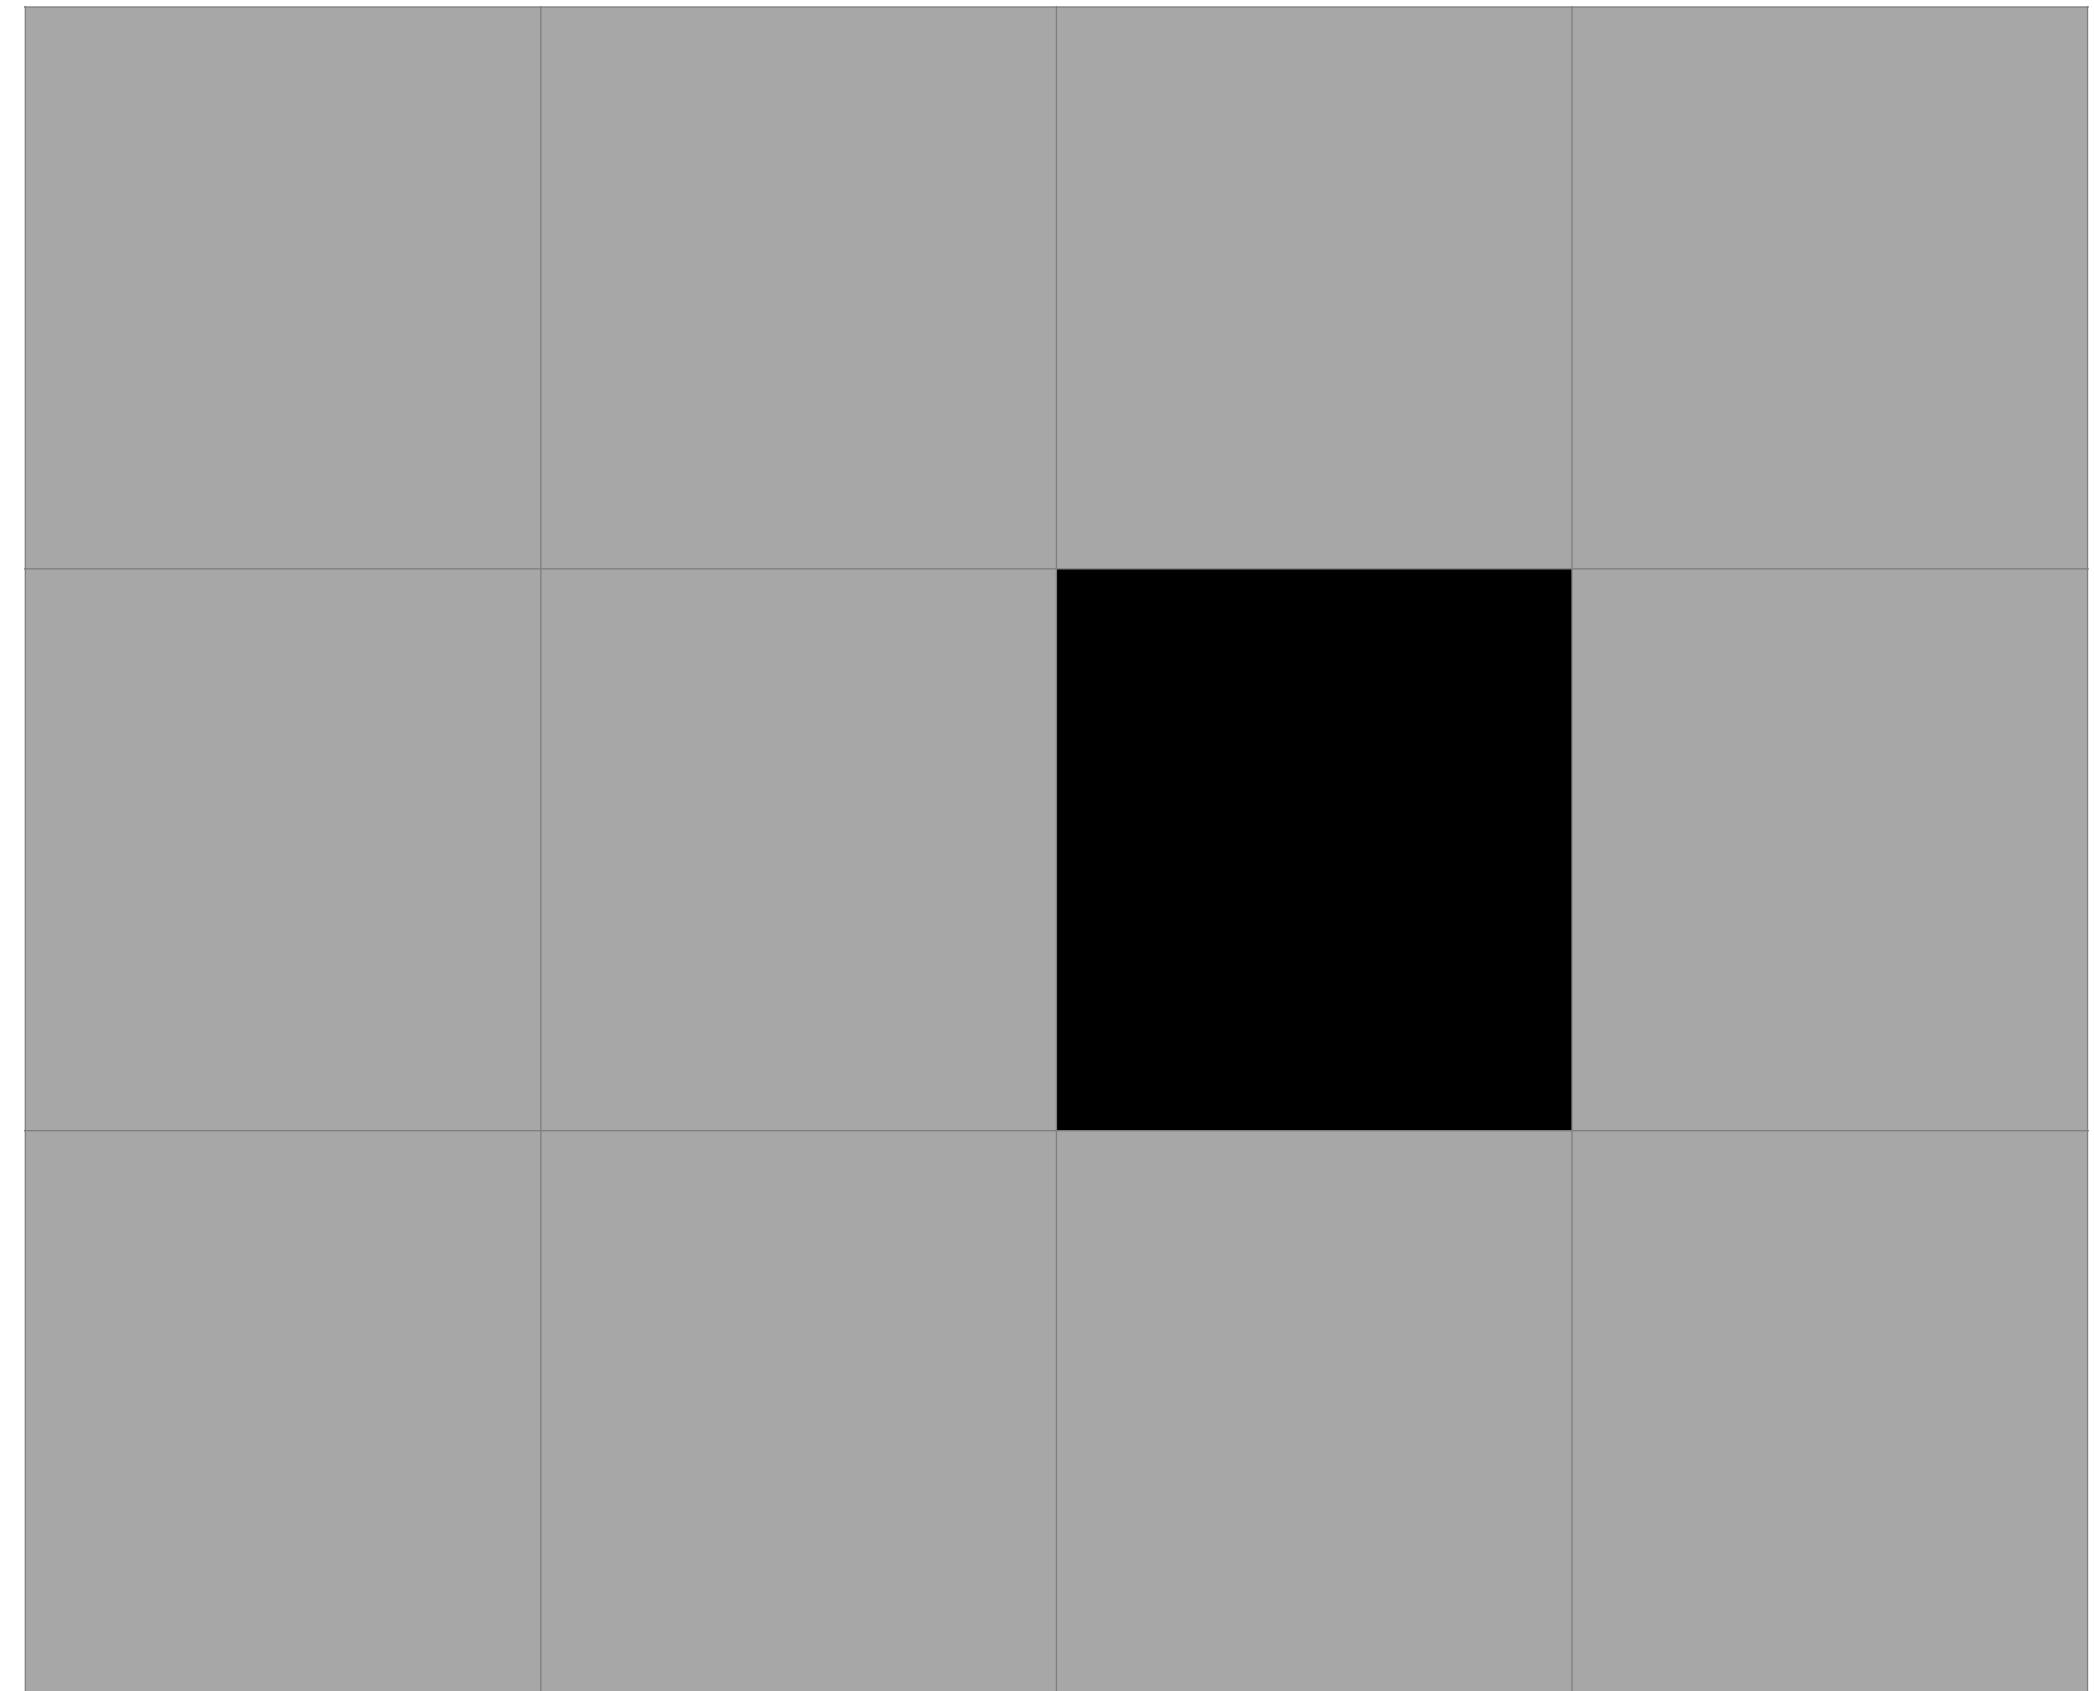
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

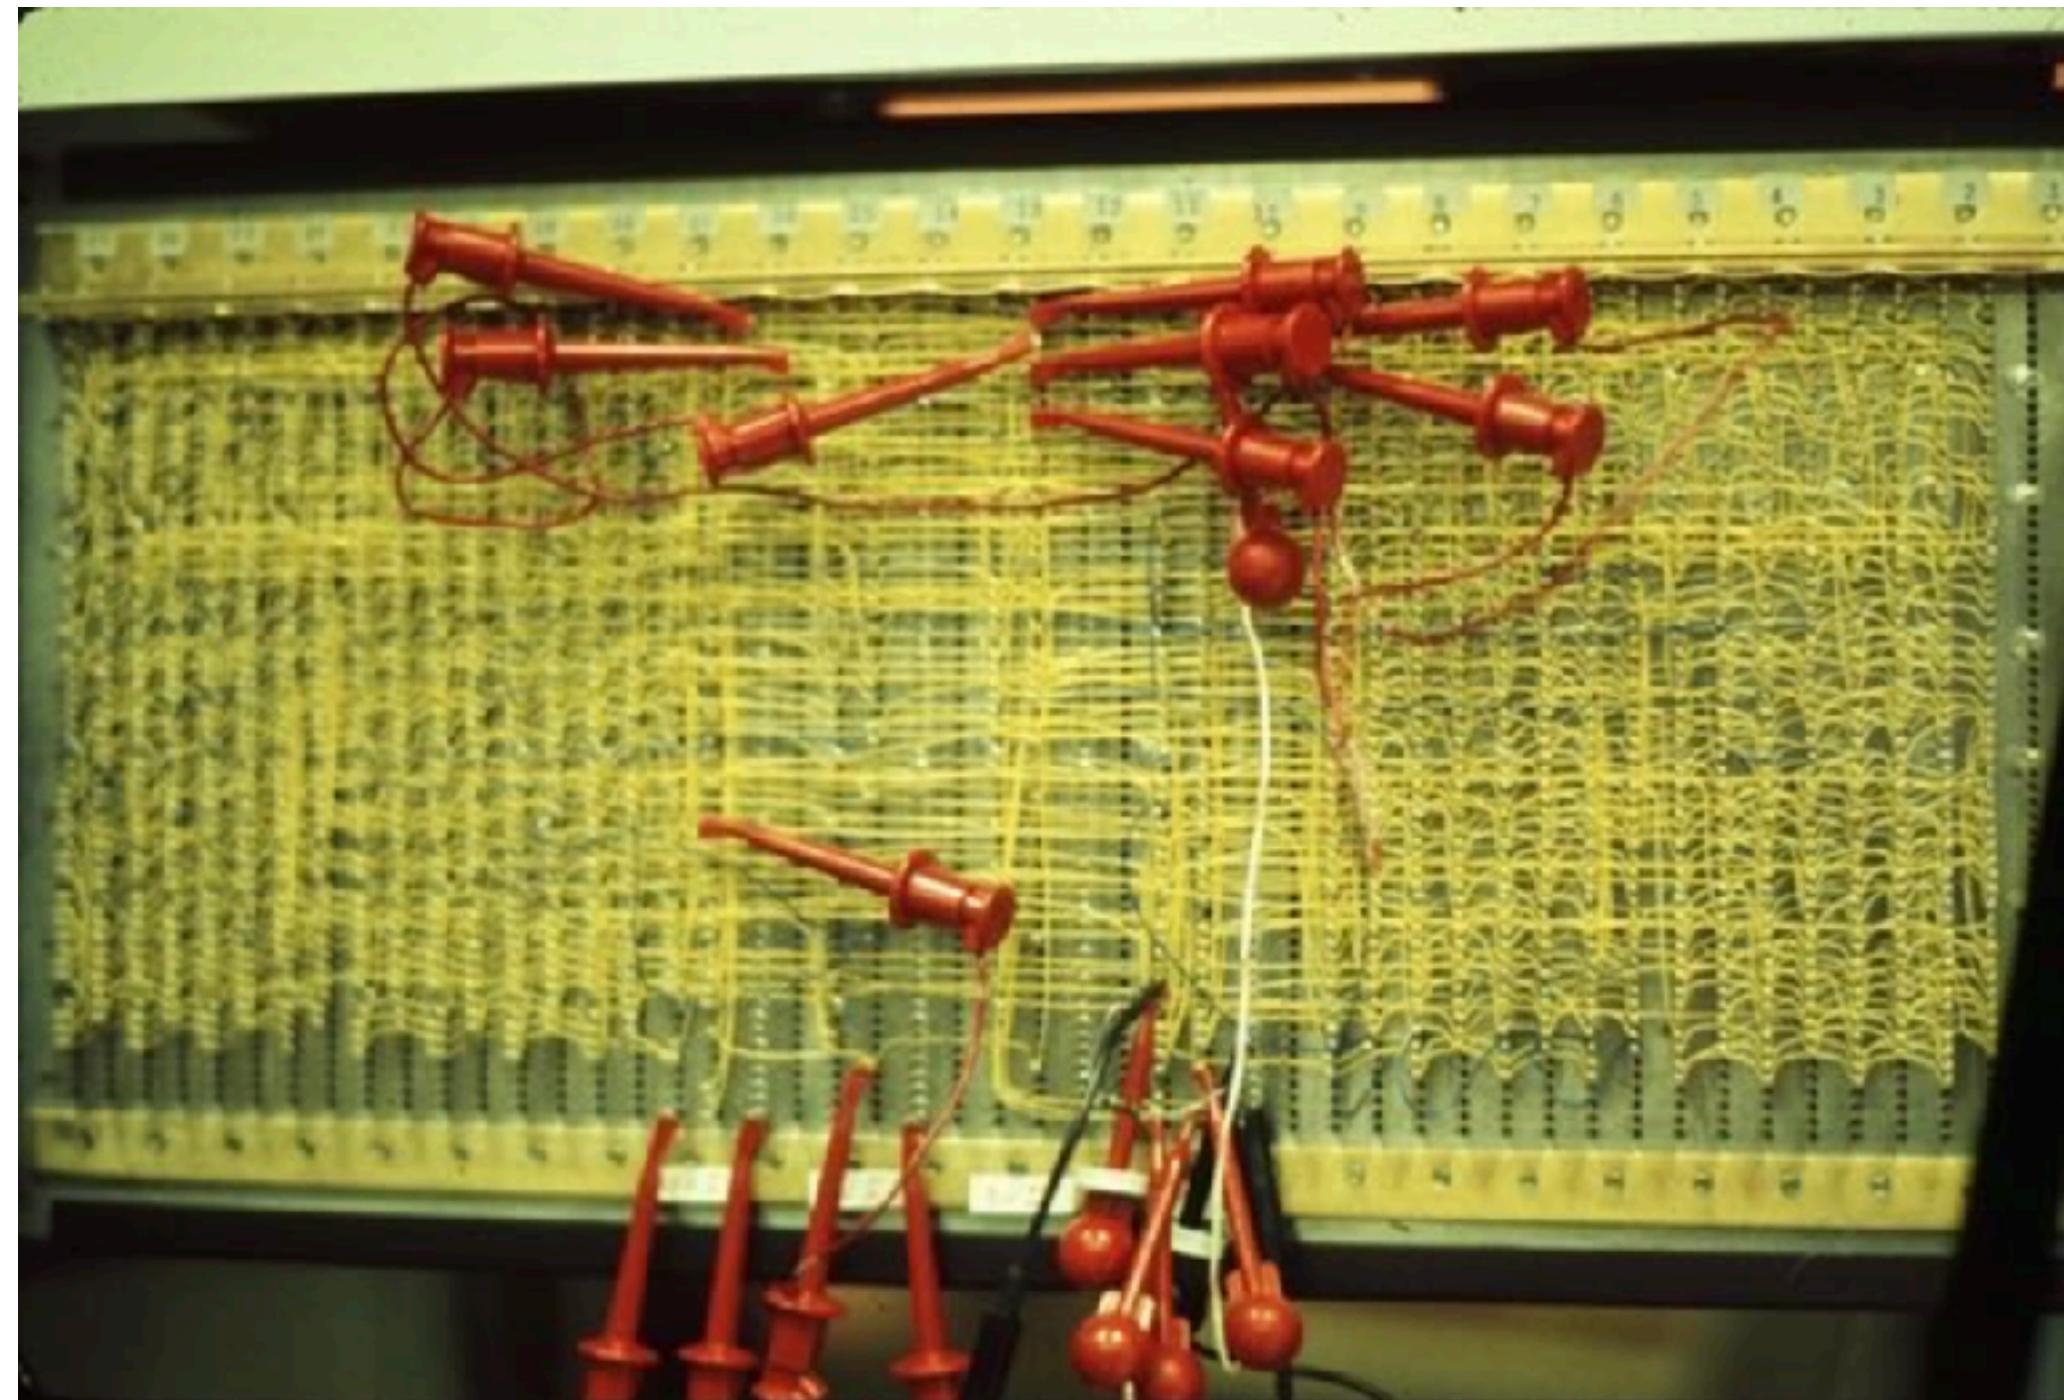
```
# Create a 1D array
image1D = create_array(width * height)
```

```
# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        index = (y * width) + x
        image1D[index] = 125
```

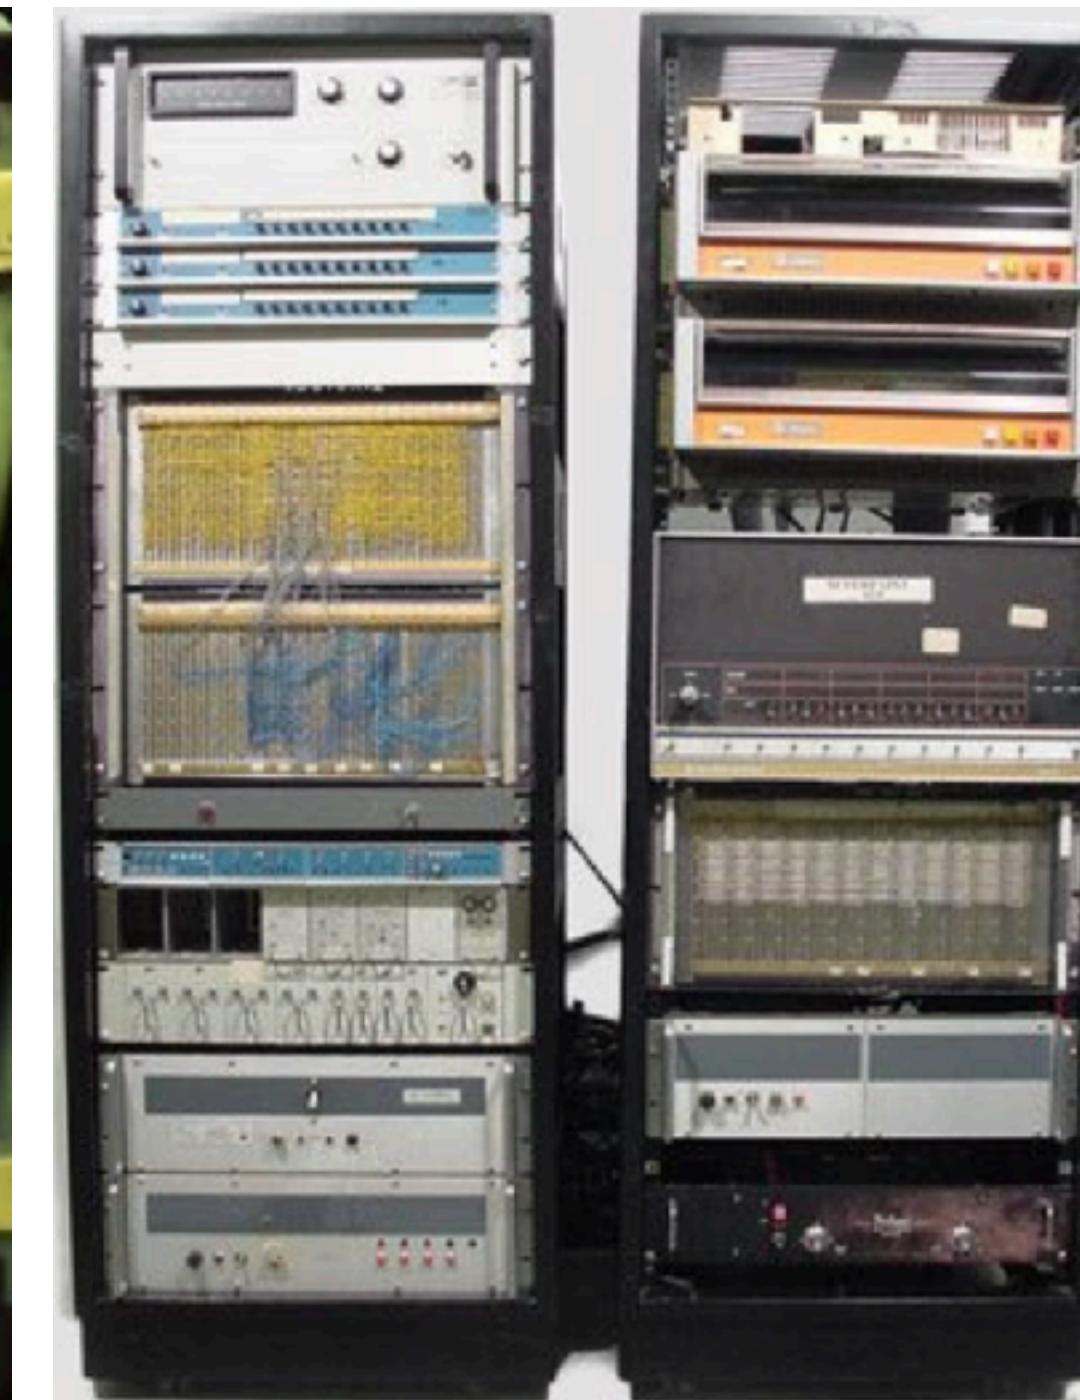
```
image1D = [ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11 ]
           ↑   row0      ↑   row1      ↑   row2
```



# Digital Images and 2D Image Representation

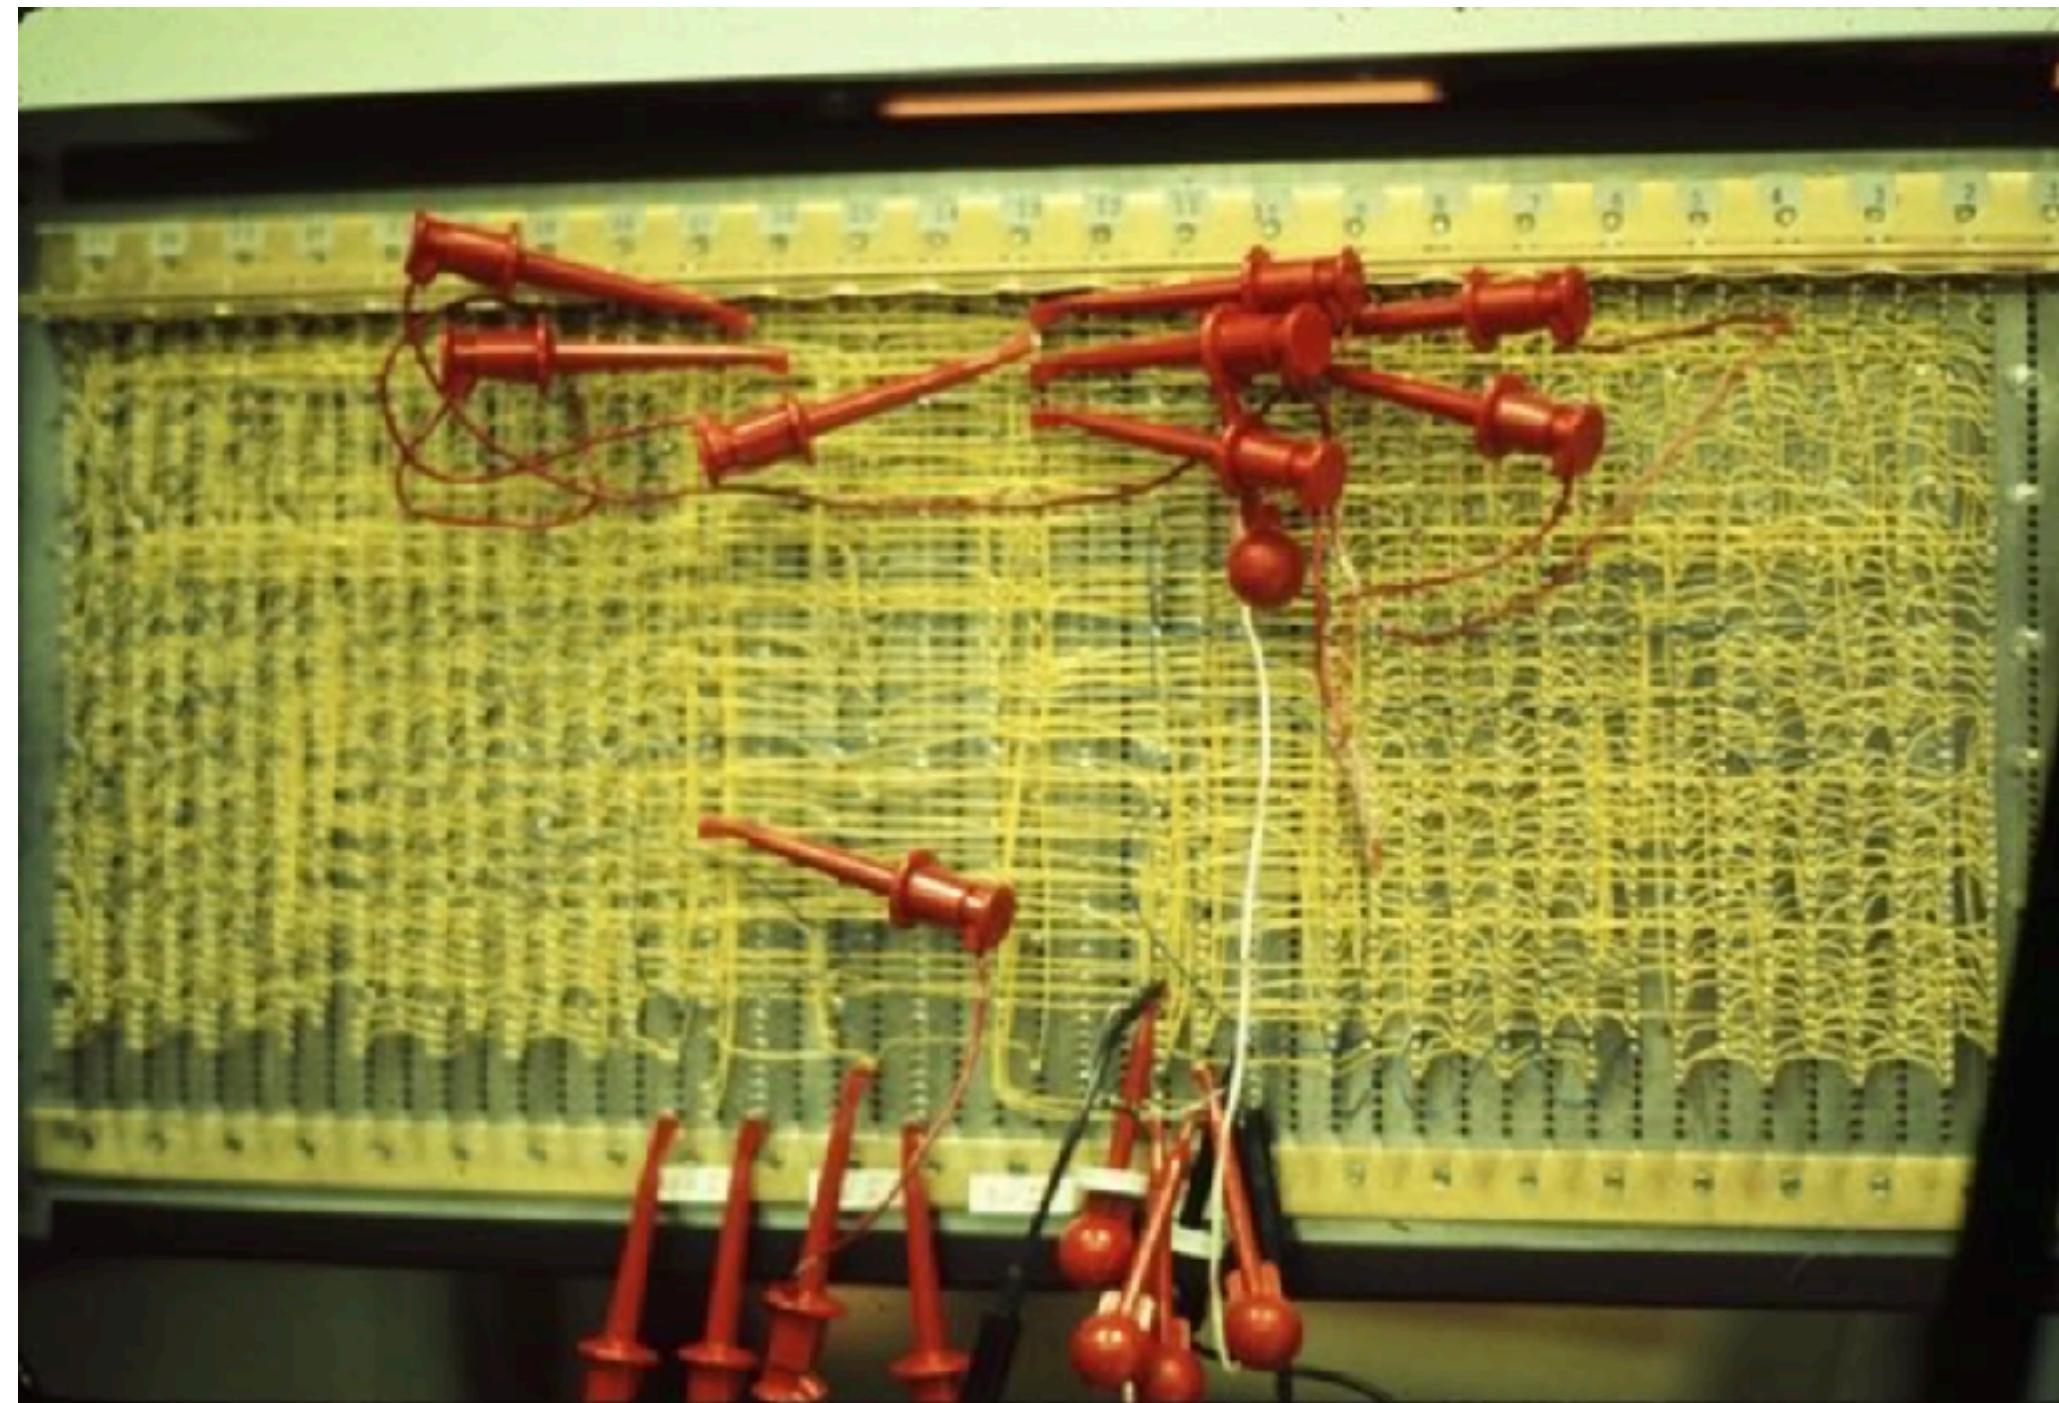


Xerox Parc Super Frame, First Frame  
Buffer



Xerox Parc Super Frame, First Image  
Stored in Frame Buffer

# Digital Images and 2D Image Representation



Xerox Parc Super Frame,  
First Frame  
Buffer



Xerox Parc Super Frame,  
First Image  
Stored in Frame Buffer

- Framebuffer: memory for a raster display

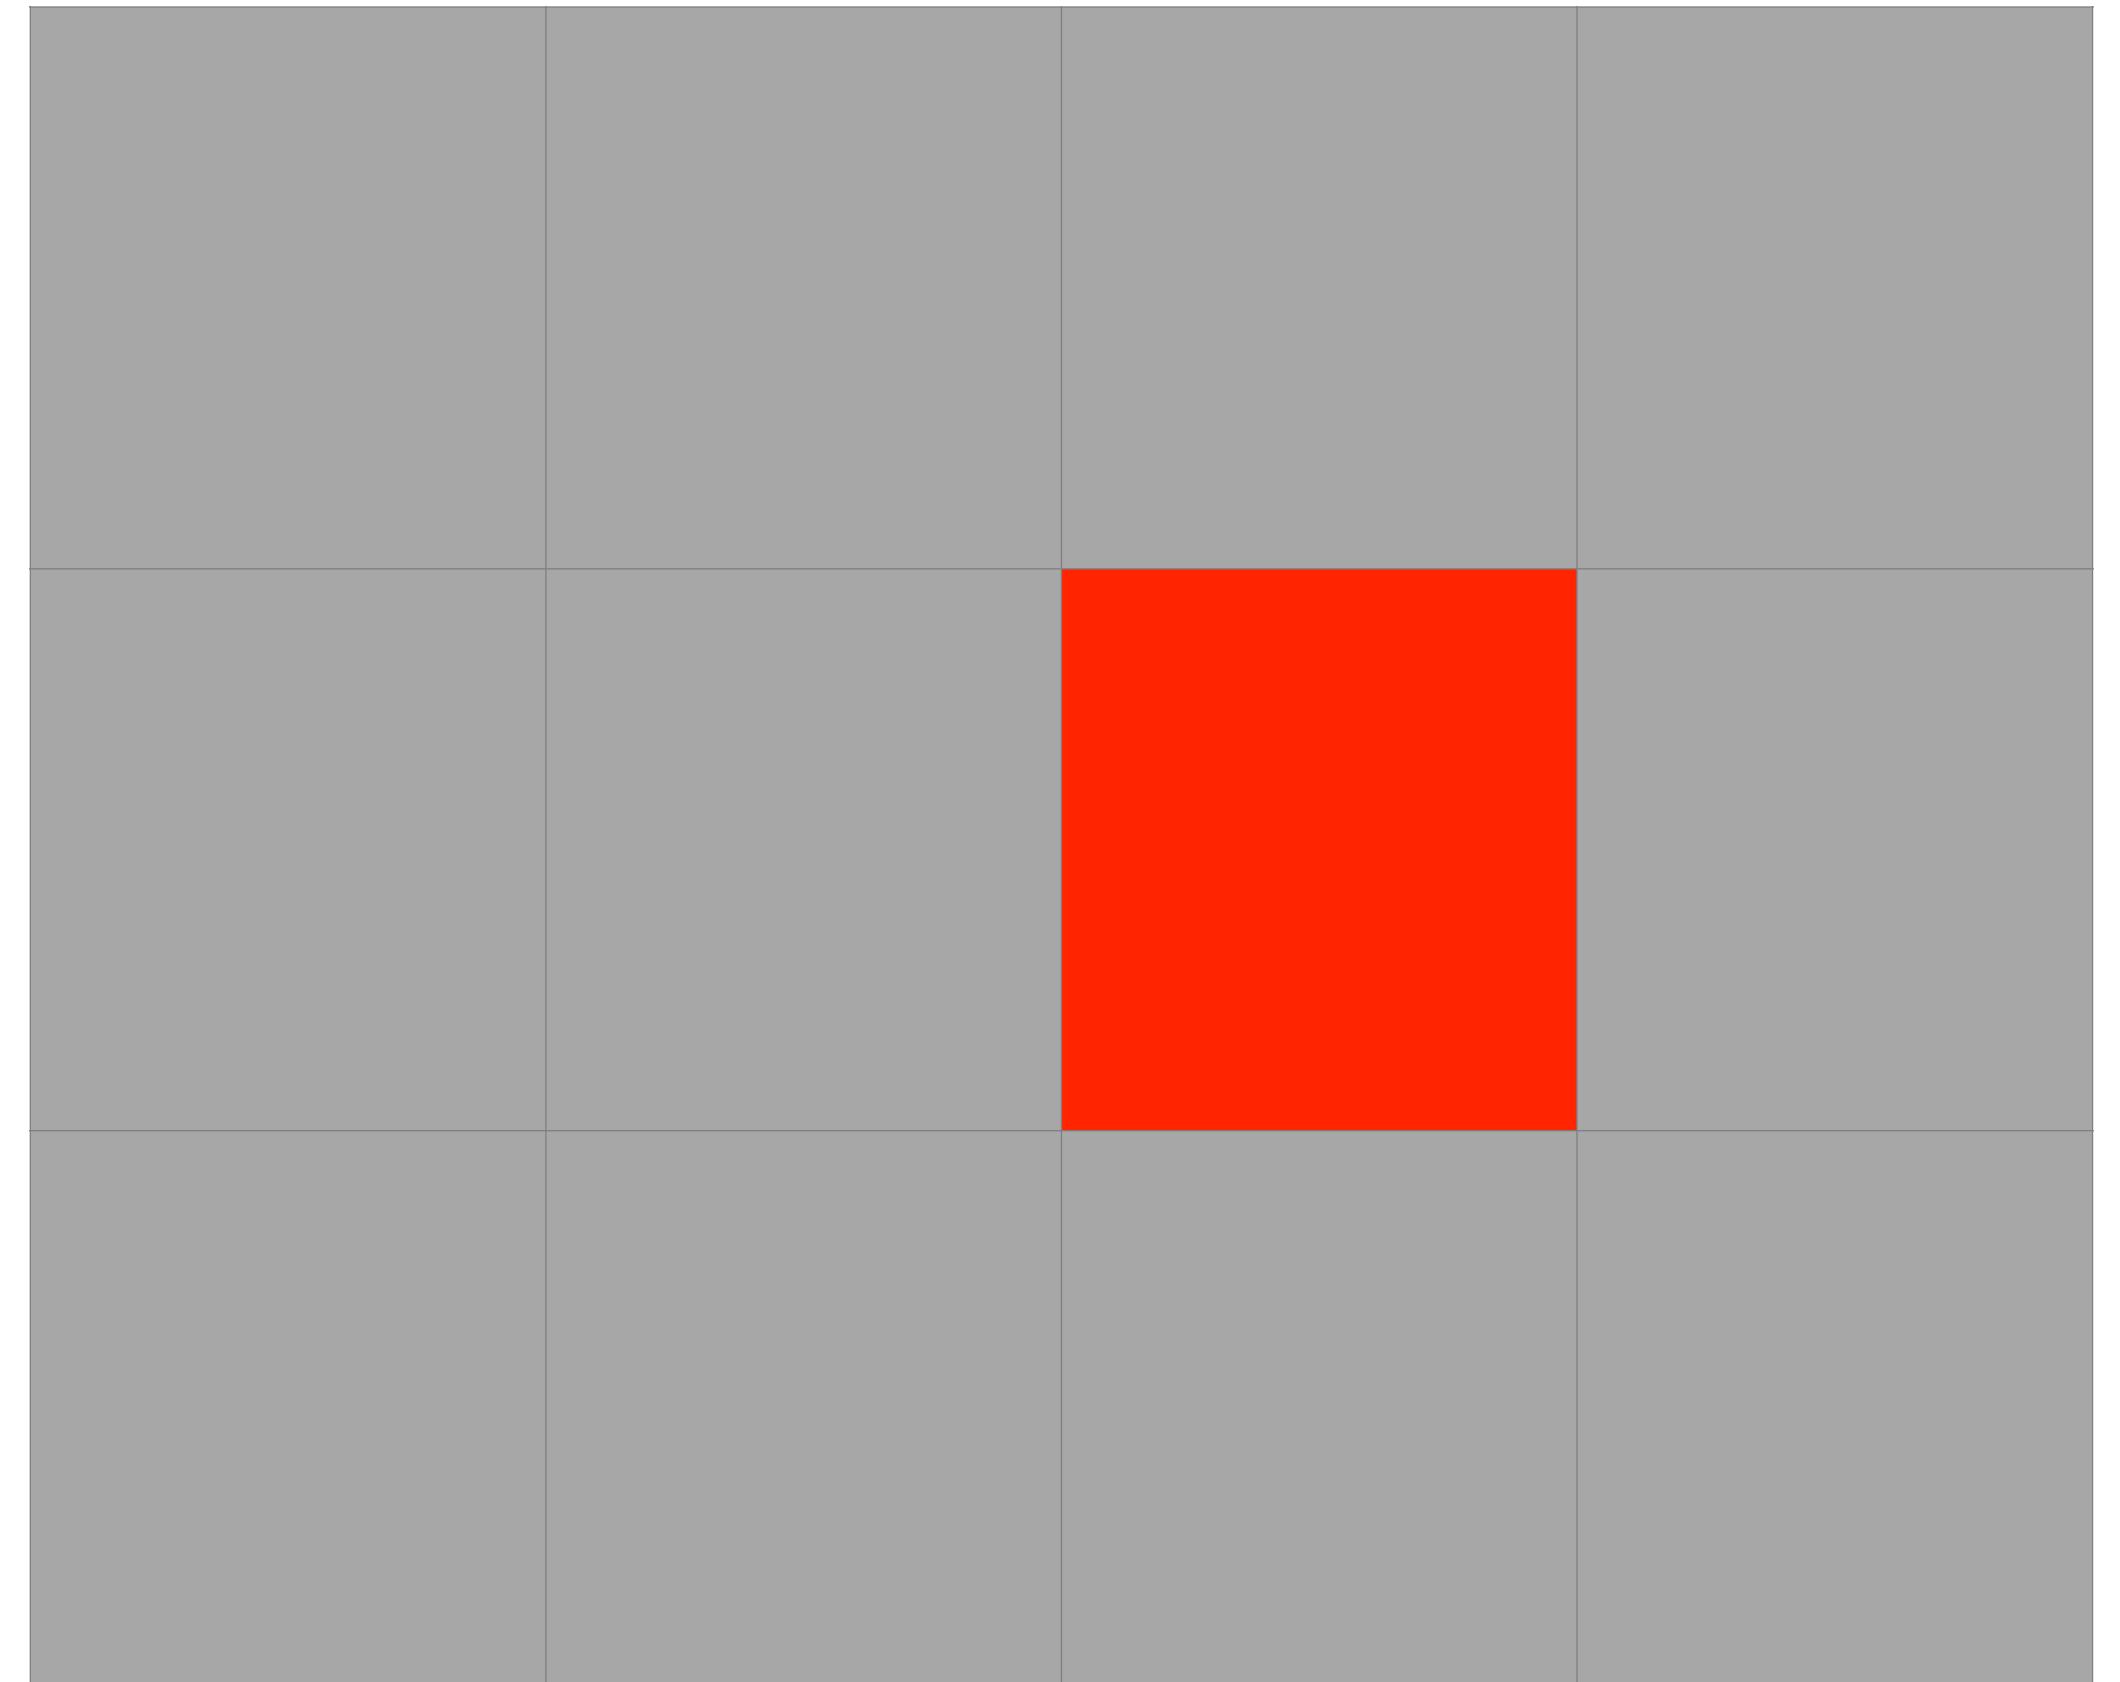
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125

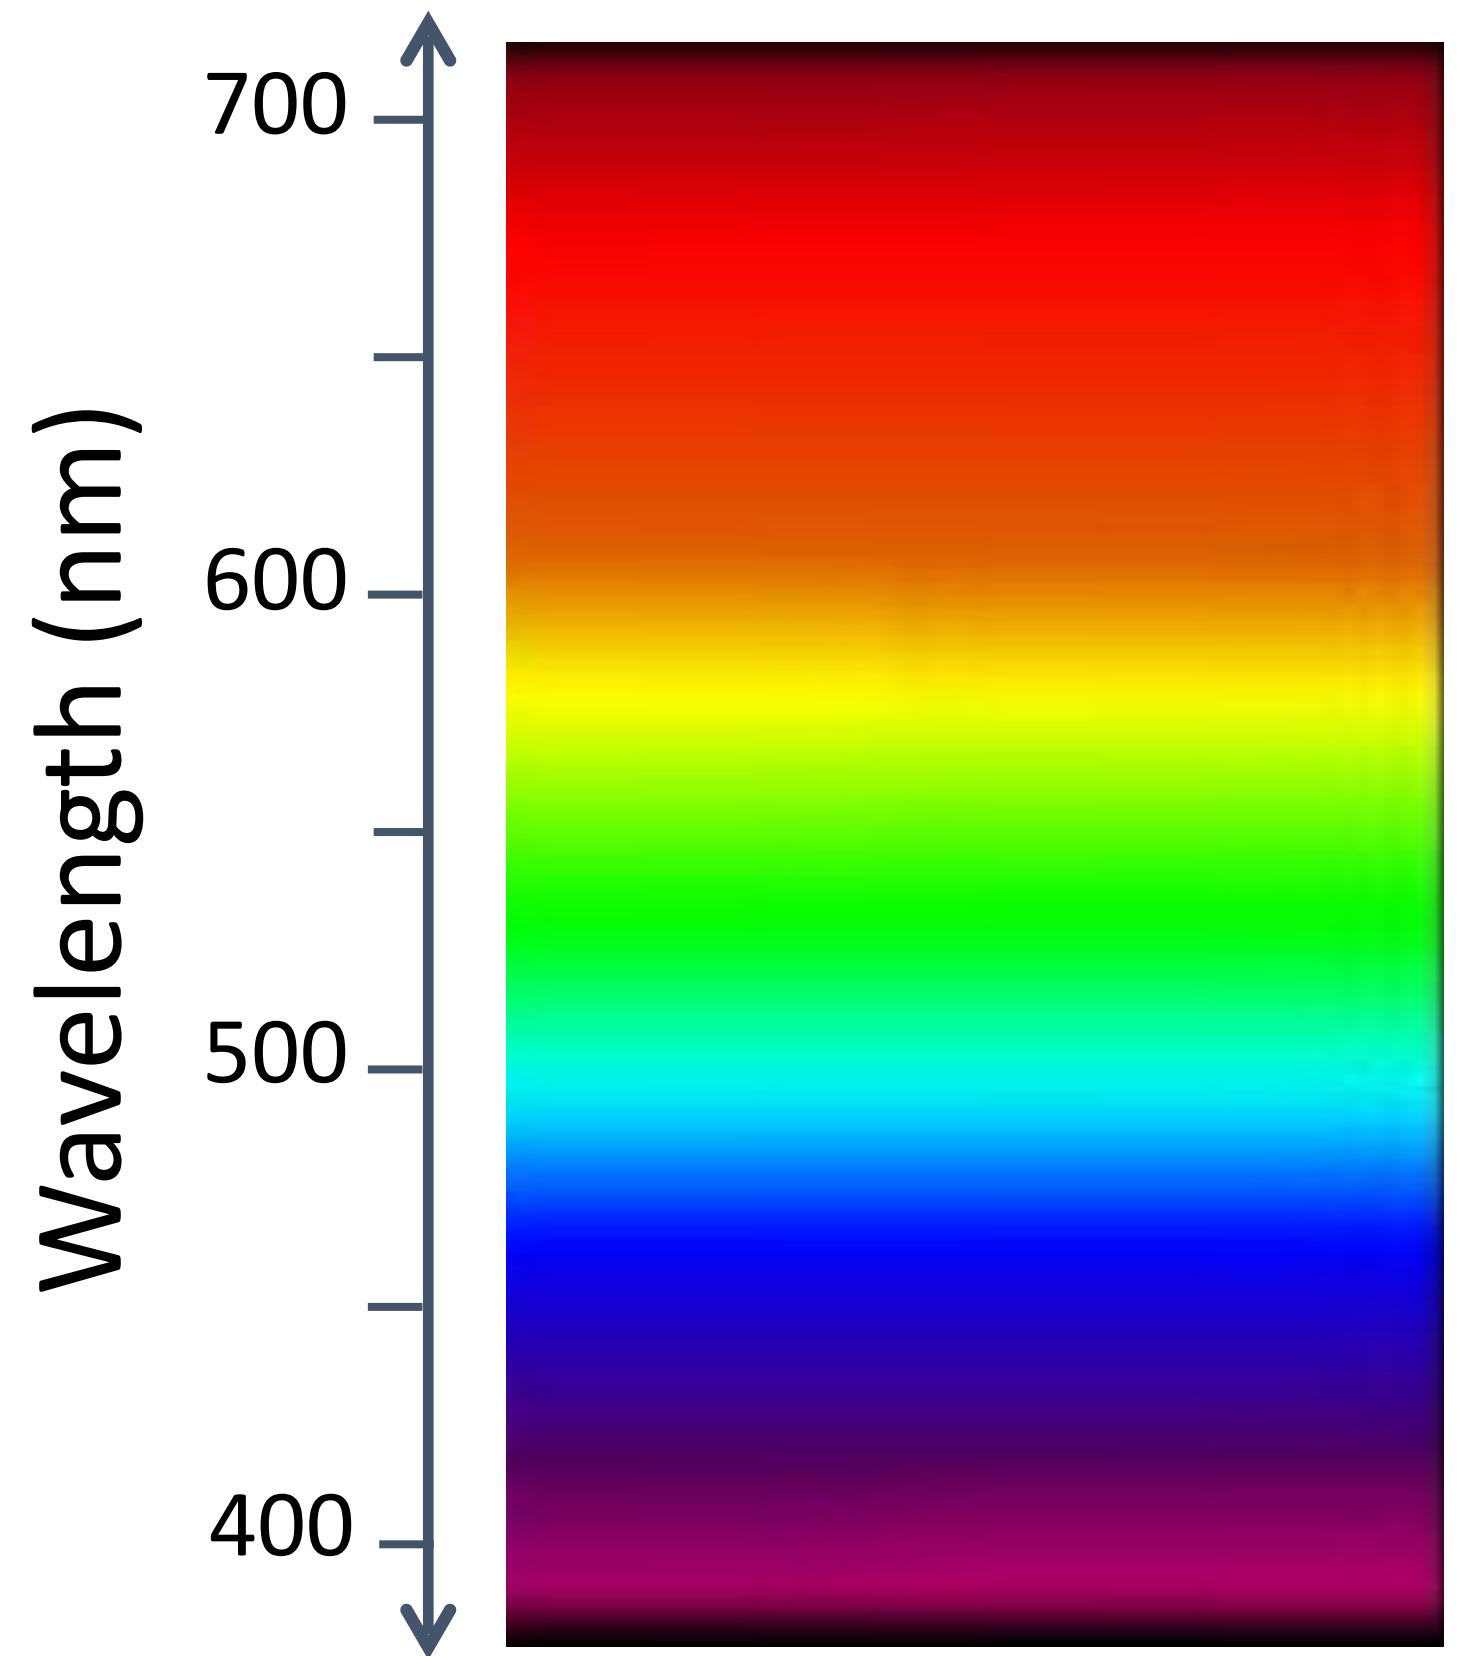
# Access pixel at (row=2, col=3)
image2D[1][2]= ?
```



# Colour

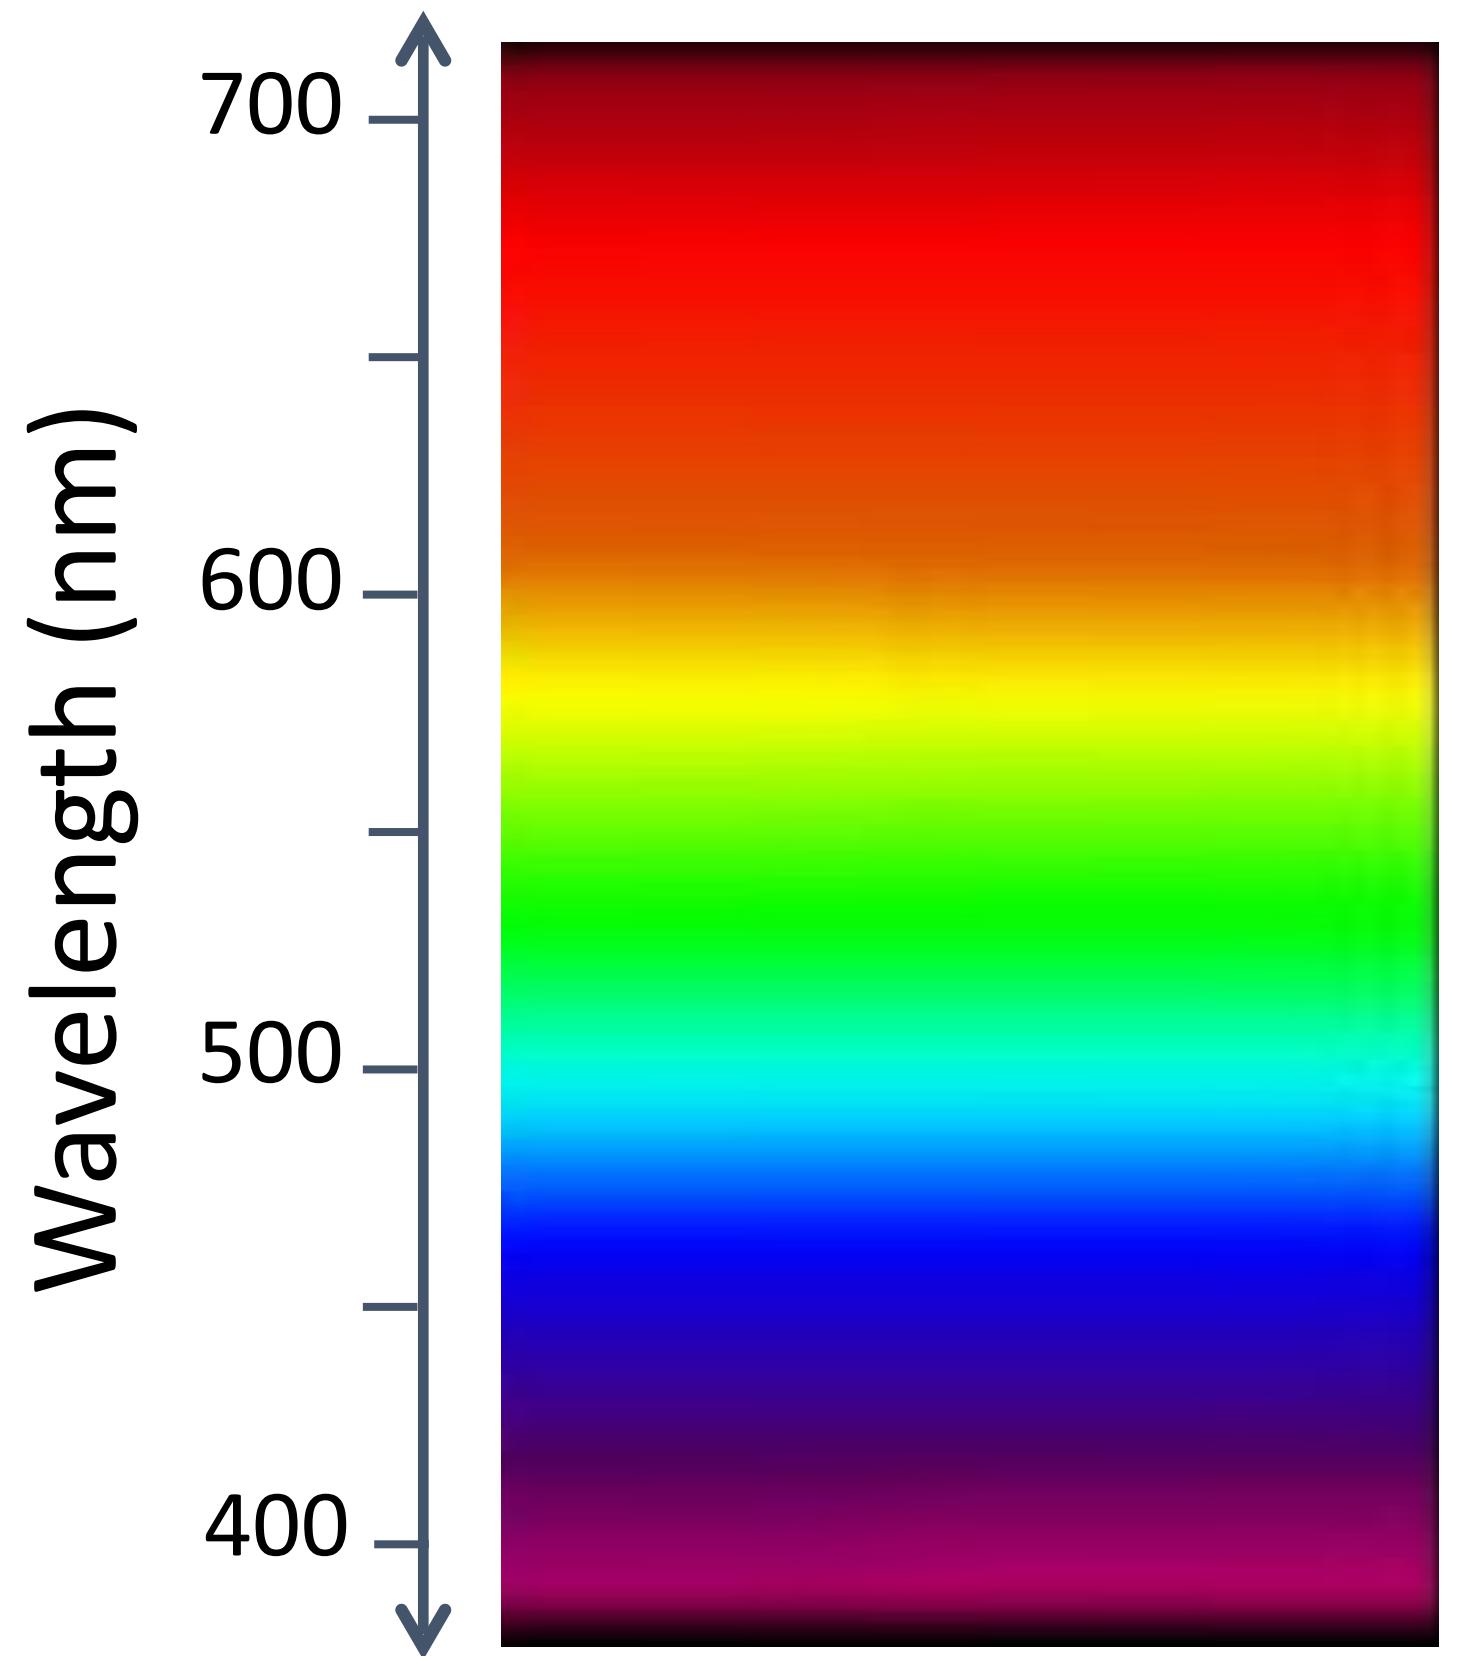
# Colour as a Spectrum

- A physical interpretation of light
  - Light has dual wave/particle nature
  - Colour relates to the frequency or wavelength of photons
  - ‘White’ is a mix of all wavelengths
- This explains many phenomena
  - Rainbows appear because refractive index varies by wavelength
  - Some processes emit photons of specific colour – e.g. sodium lights, LED’s etc



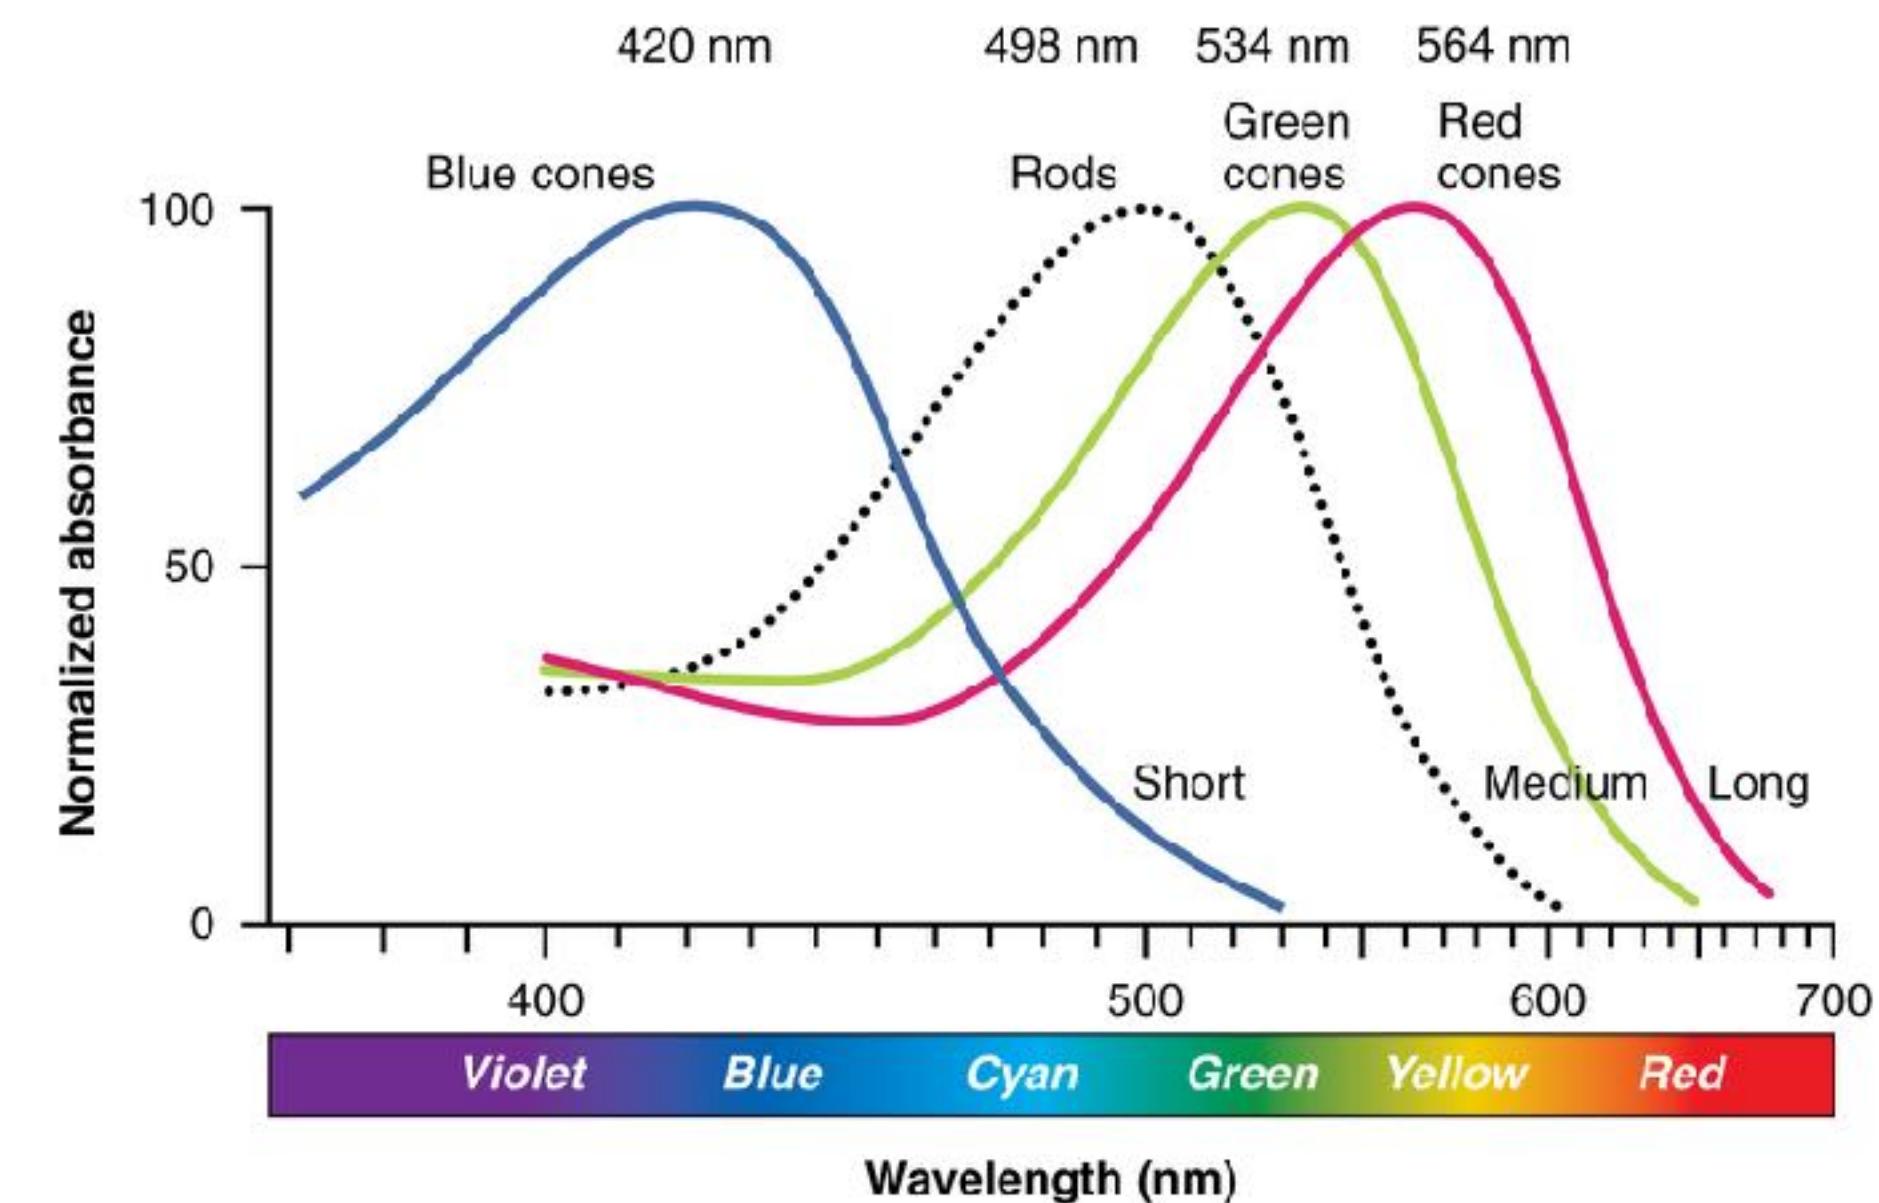
# Colour as a Spectrum?

- This does not explain everything
  - Blue + Green light → Cyan
  - Red + Green light → Yellow
  - Red + Blue light → Magenta
  - Blue + Yellow **paint** → Green
  - Blue + Yellow light → White



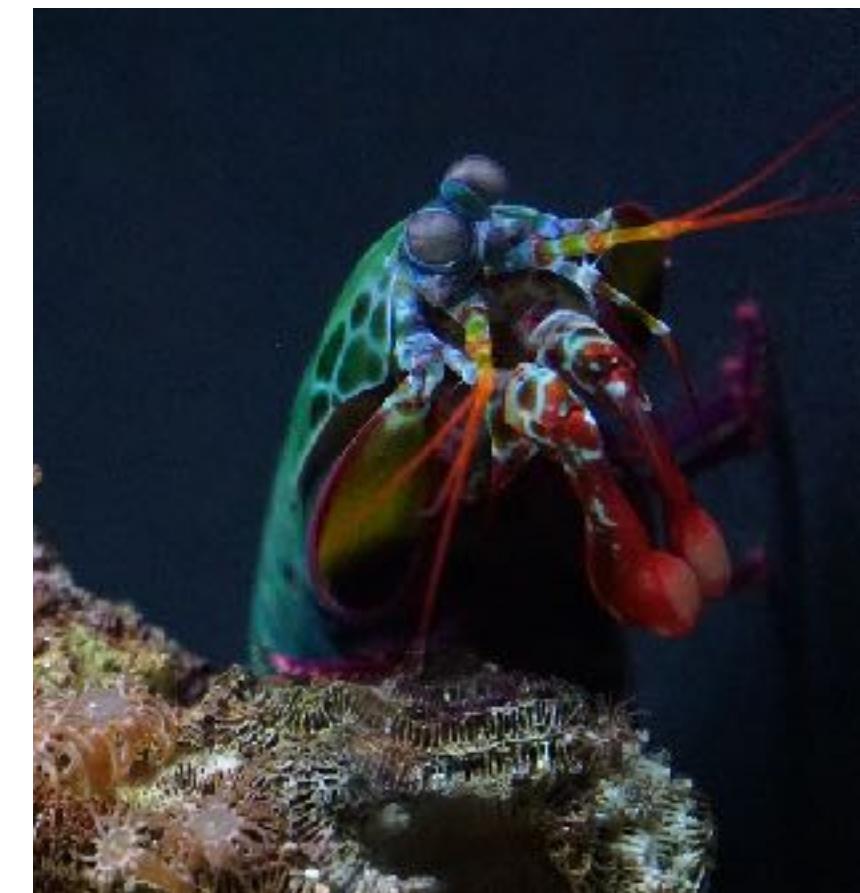
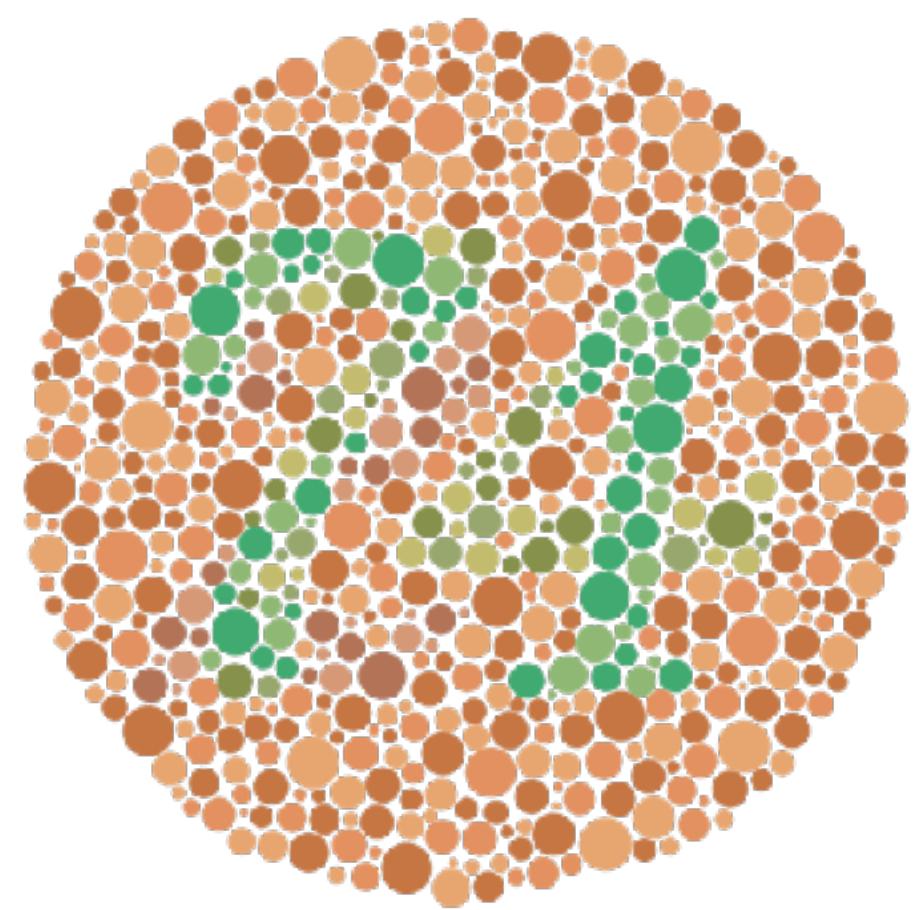
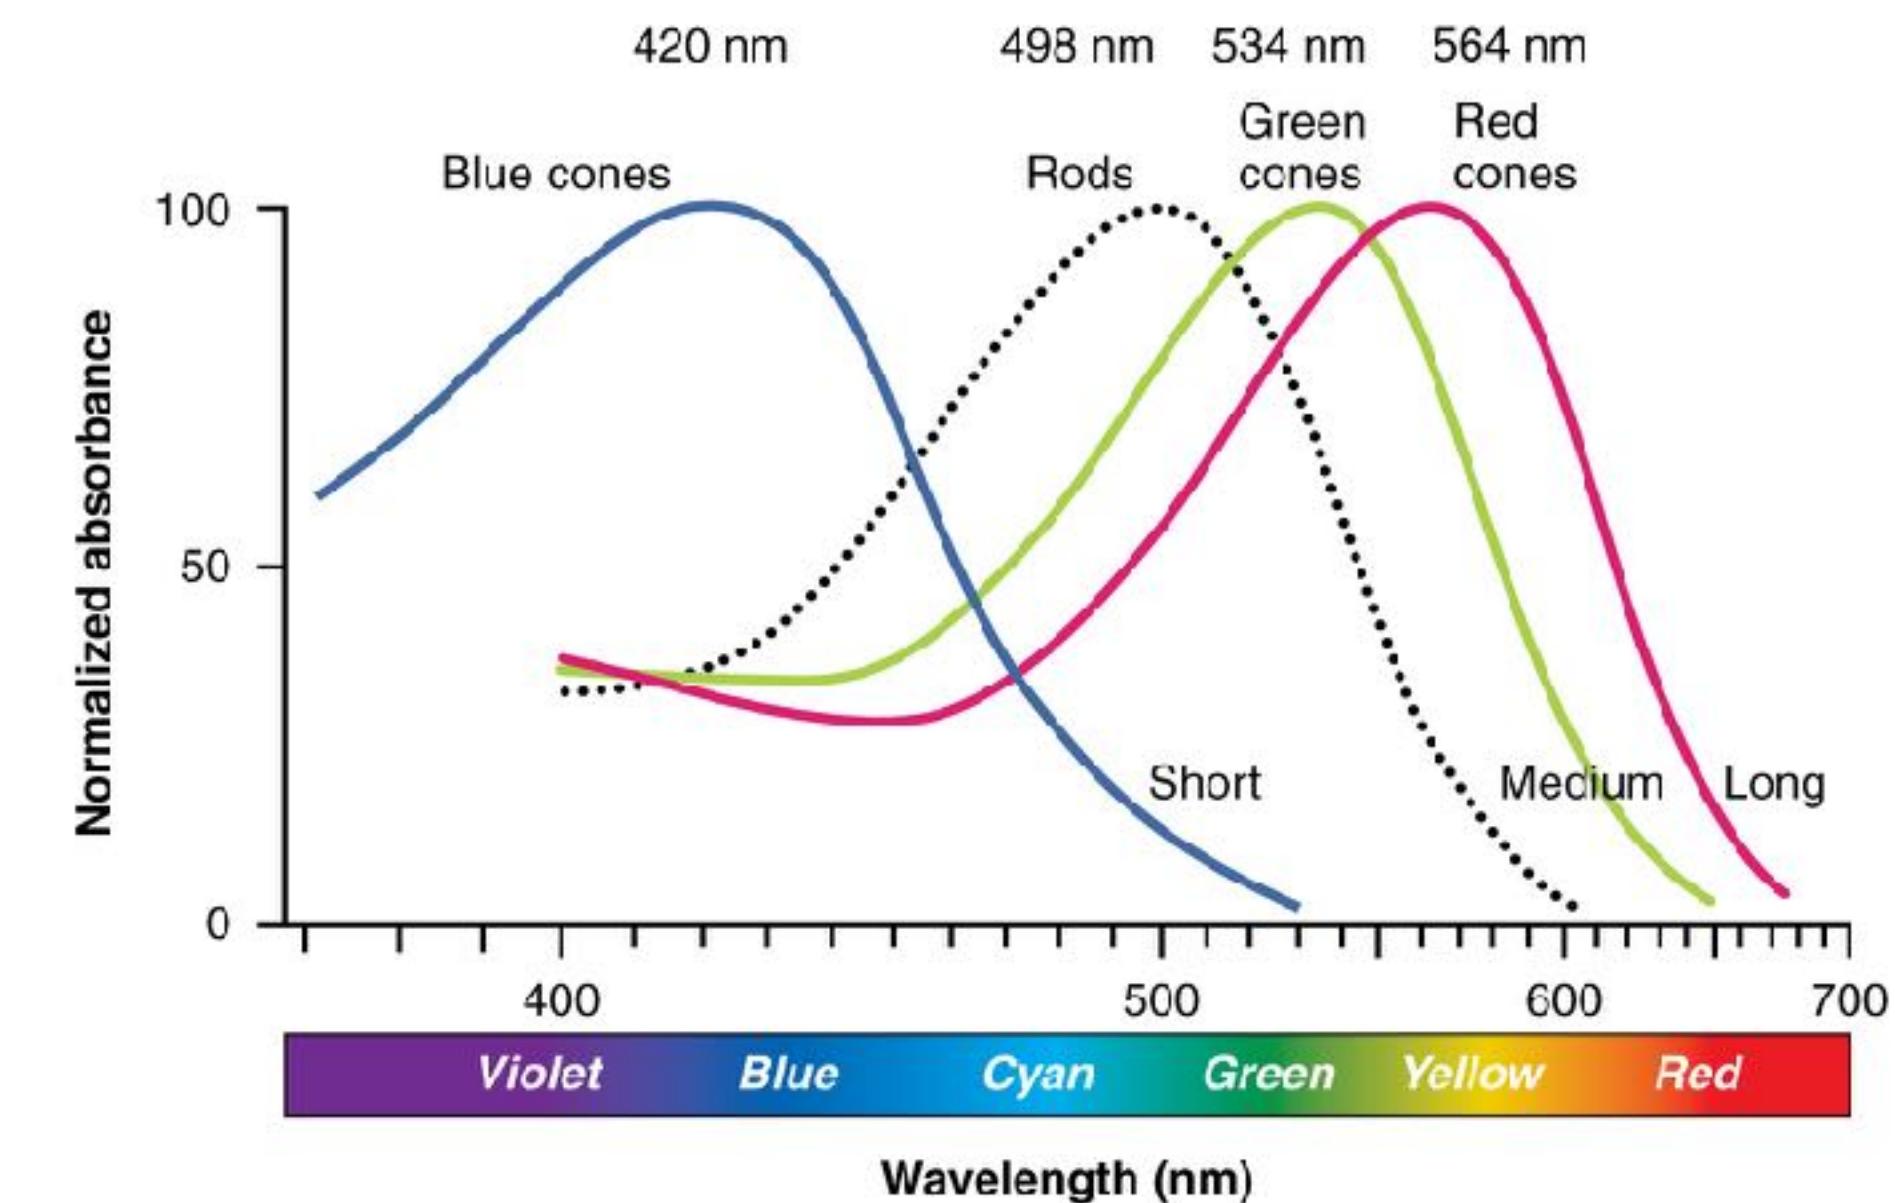
# Colour Perception

- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones



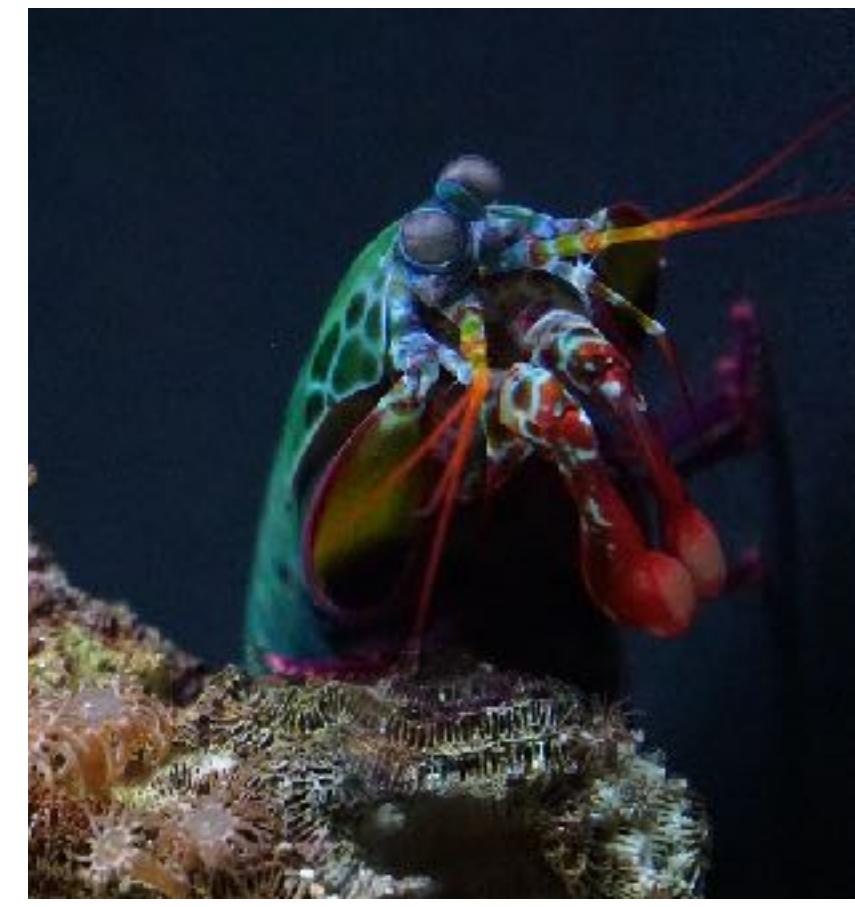
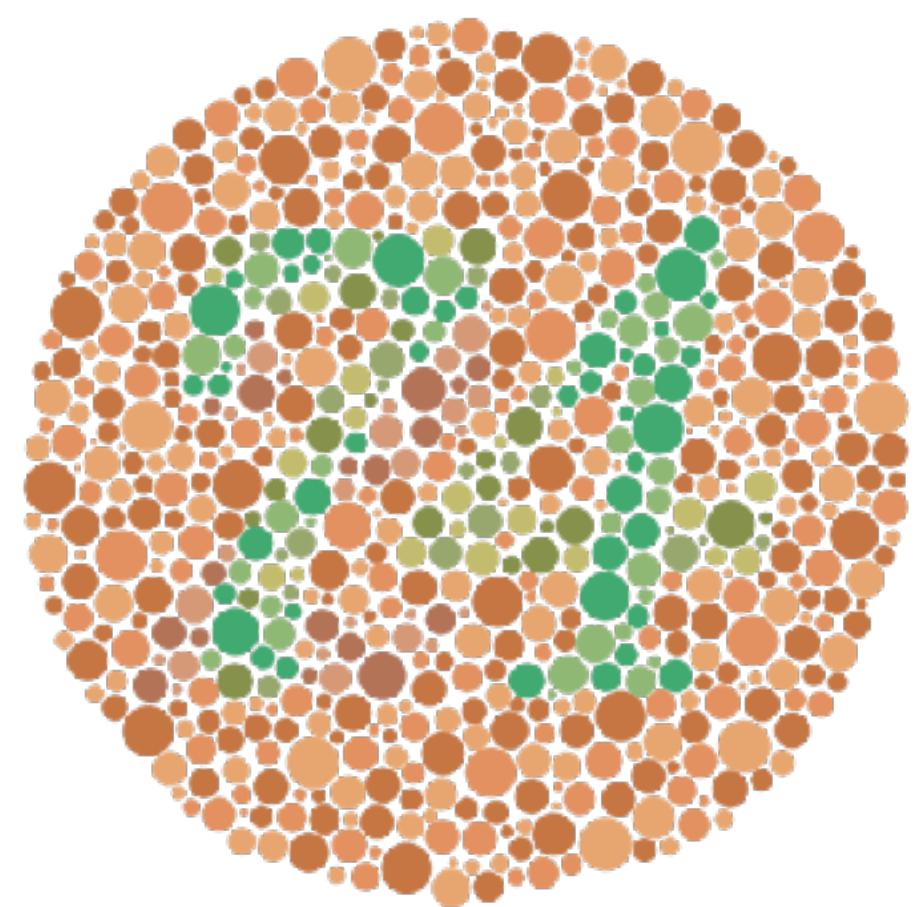
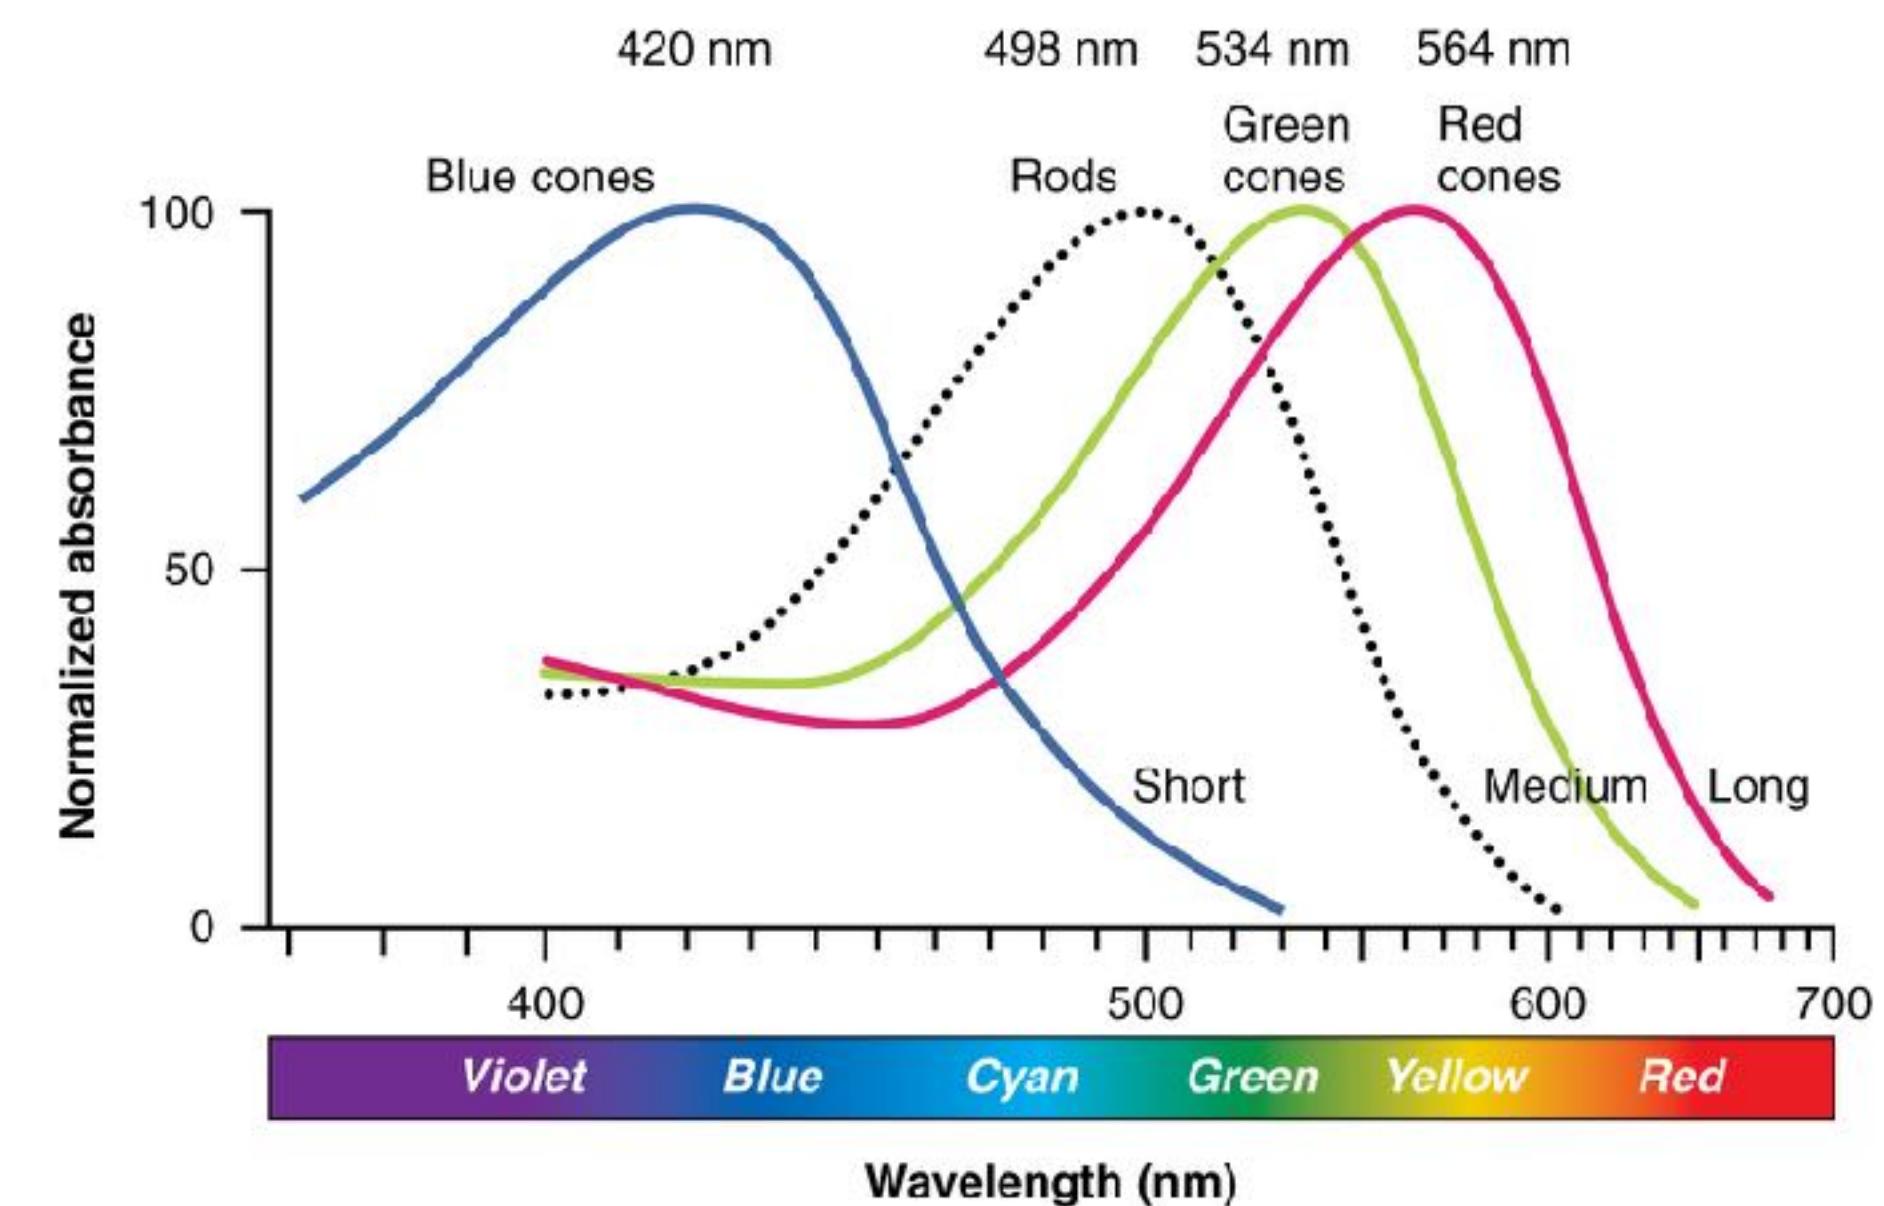
# Colour Perception

- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones



# Colour Perception

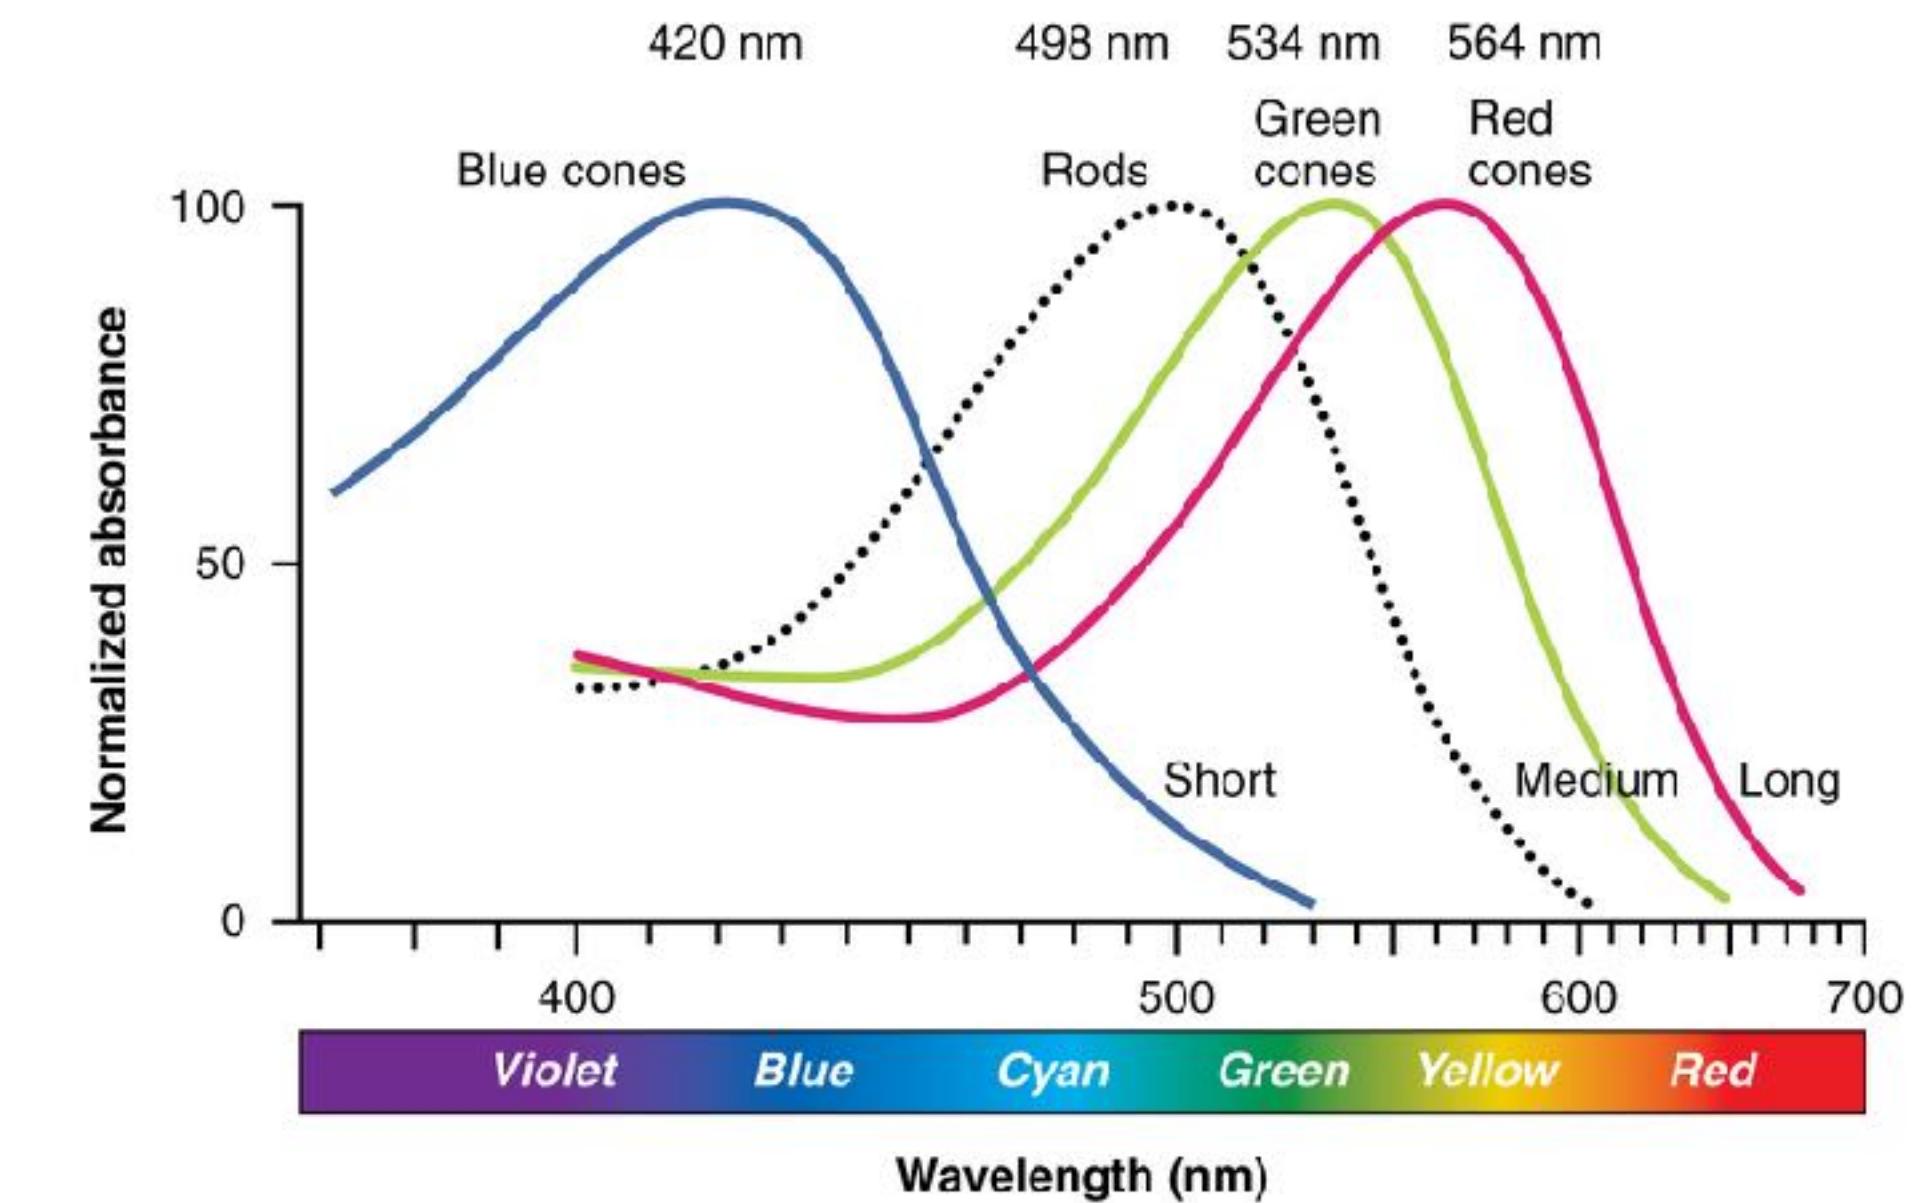
- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones



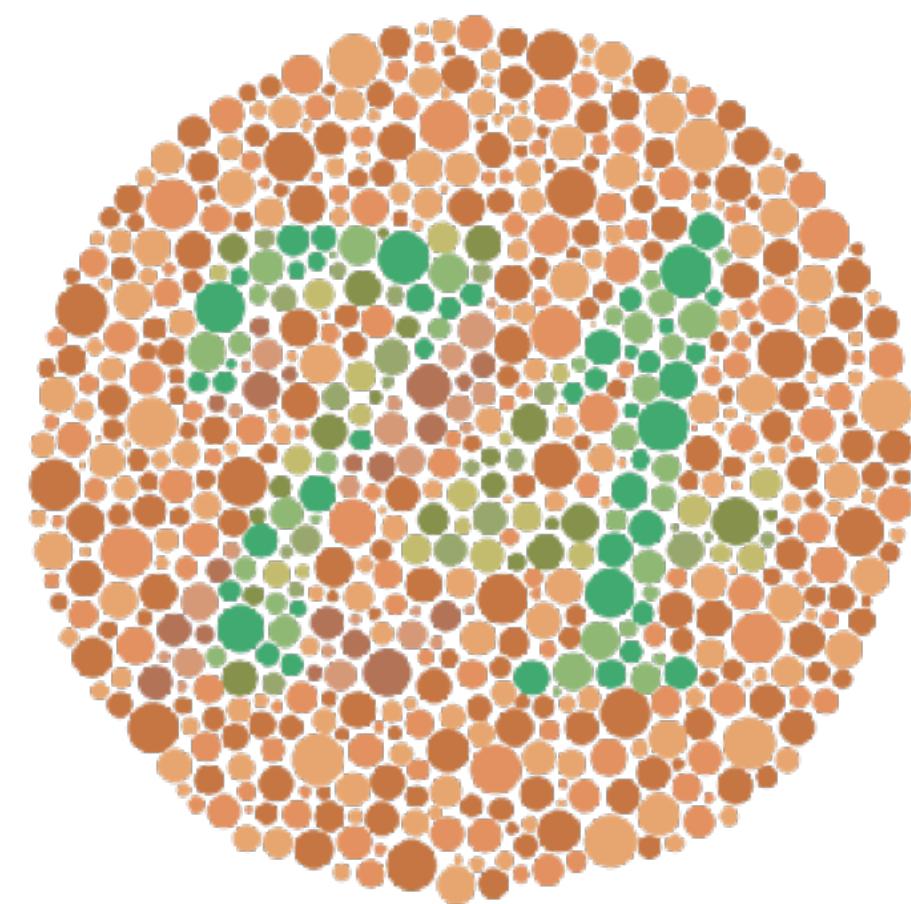
Trichromacy

# Colour Perception

- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones



Monochromacy



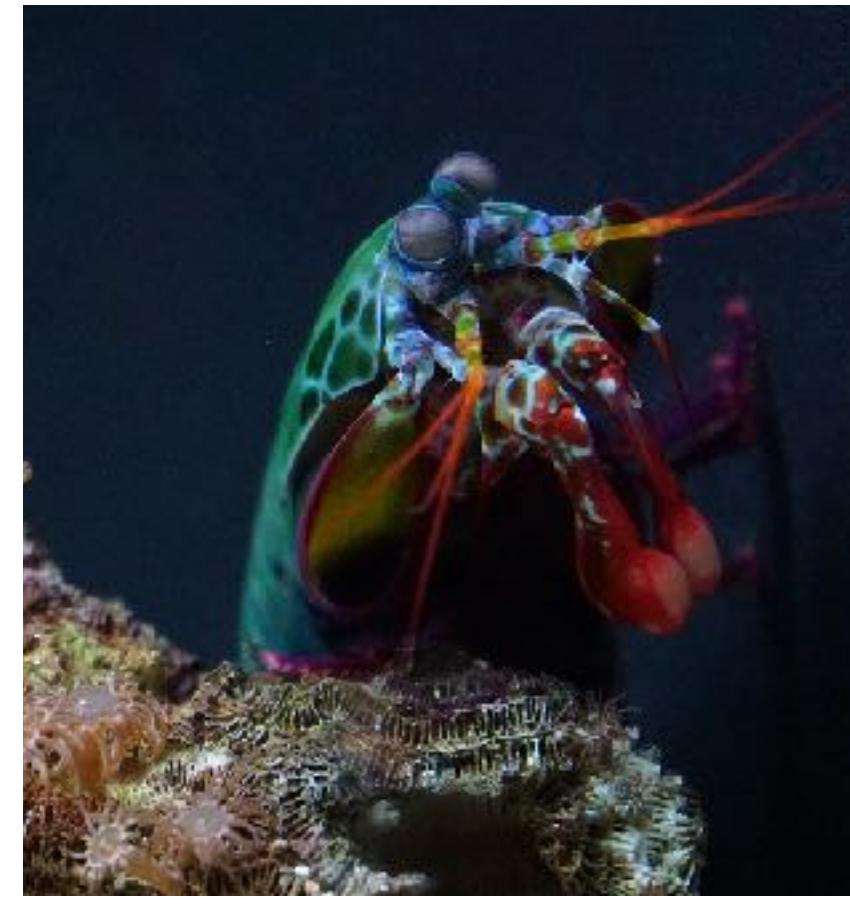
Dichromacy



Trichromacy



Tetrachromacy

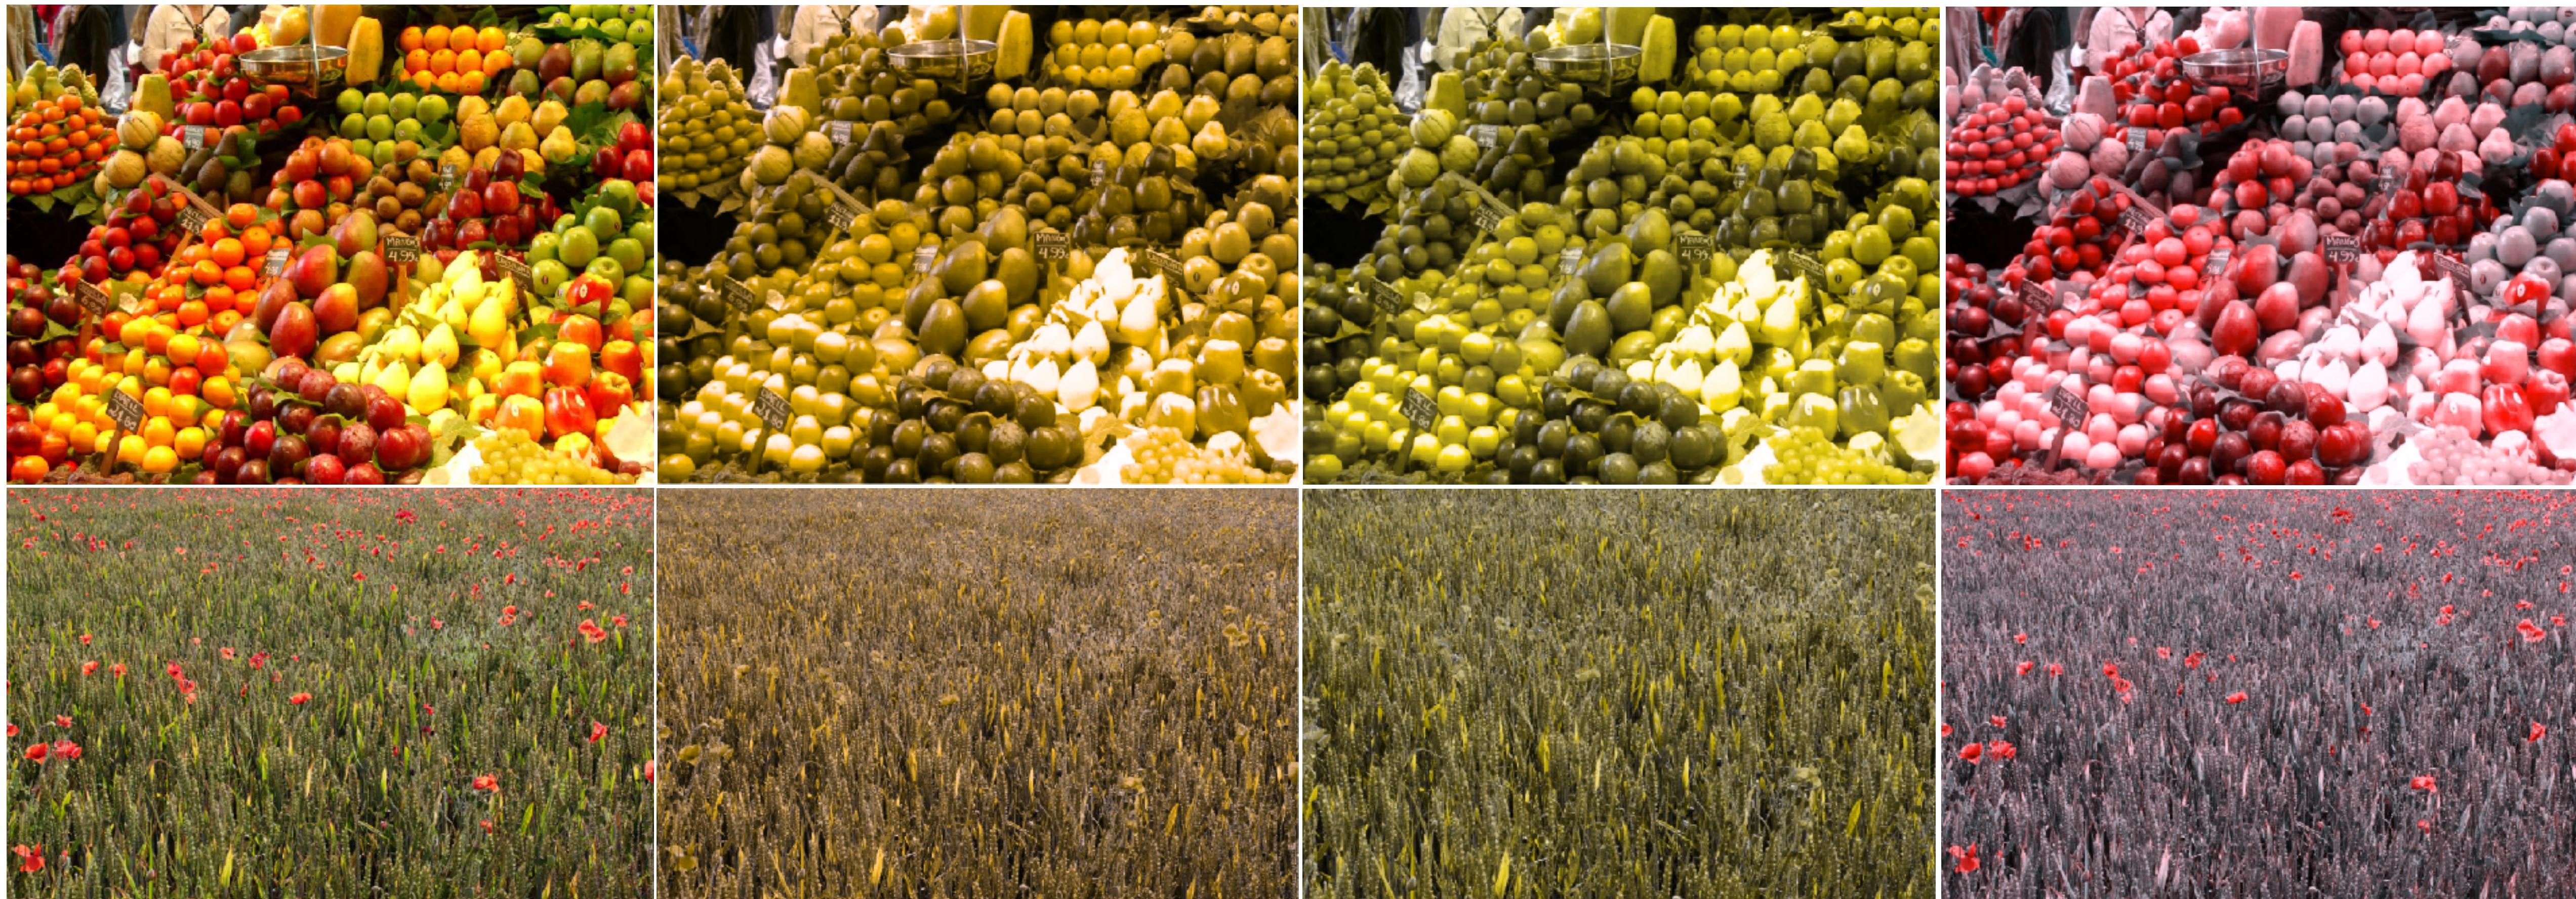


Dodecachromacy  
to Hexadecachromacy

**Sidetrack:**

**Some info on Dichromacy and Colour  
Vision Deficiency**

# Colour Vision Deficiency (CVD)



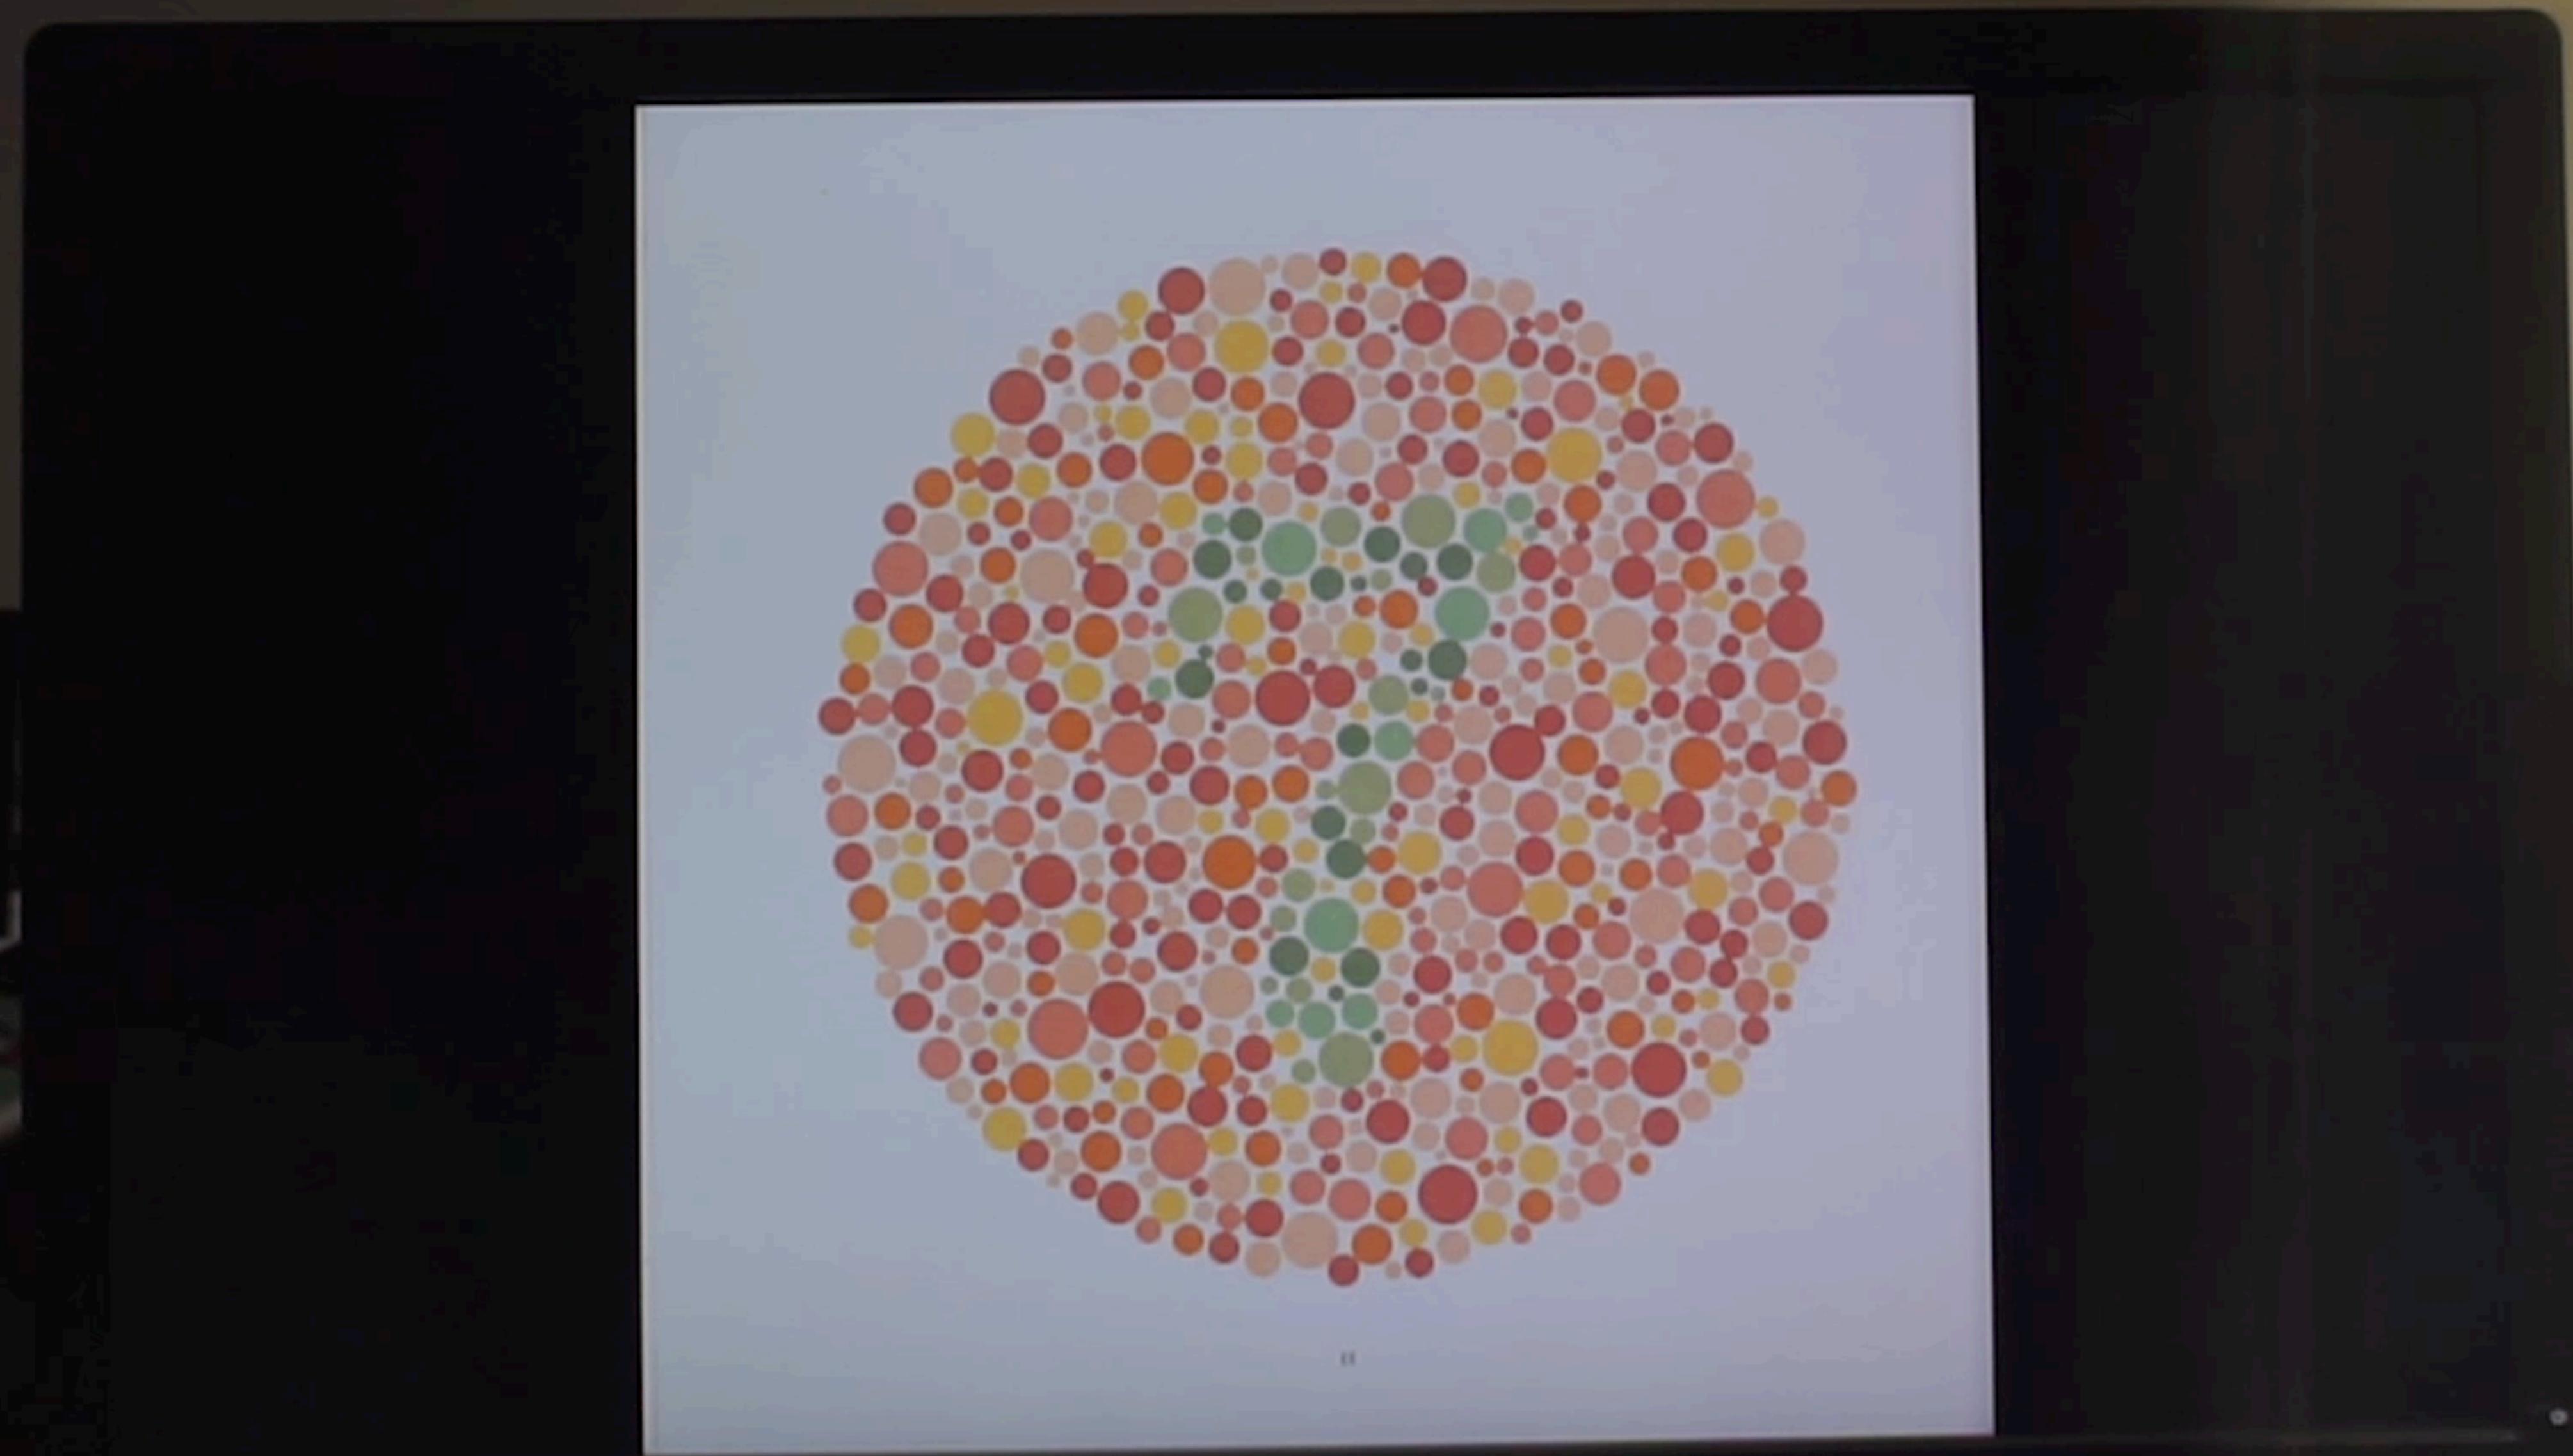
Original

Deutanopia  
(green blind)

Protanopia  
(red blind)

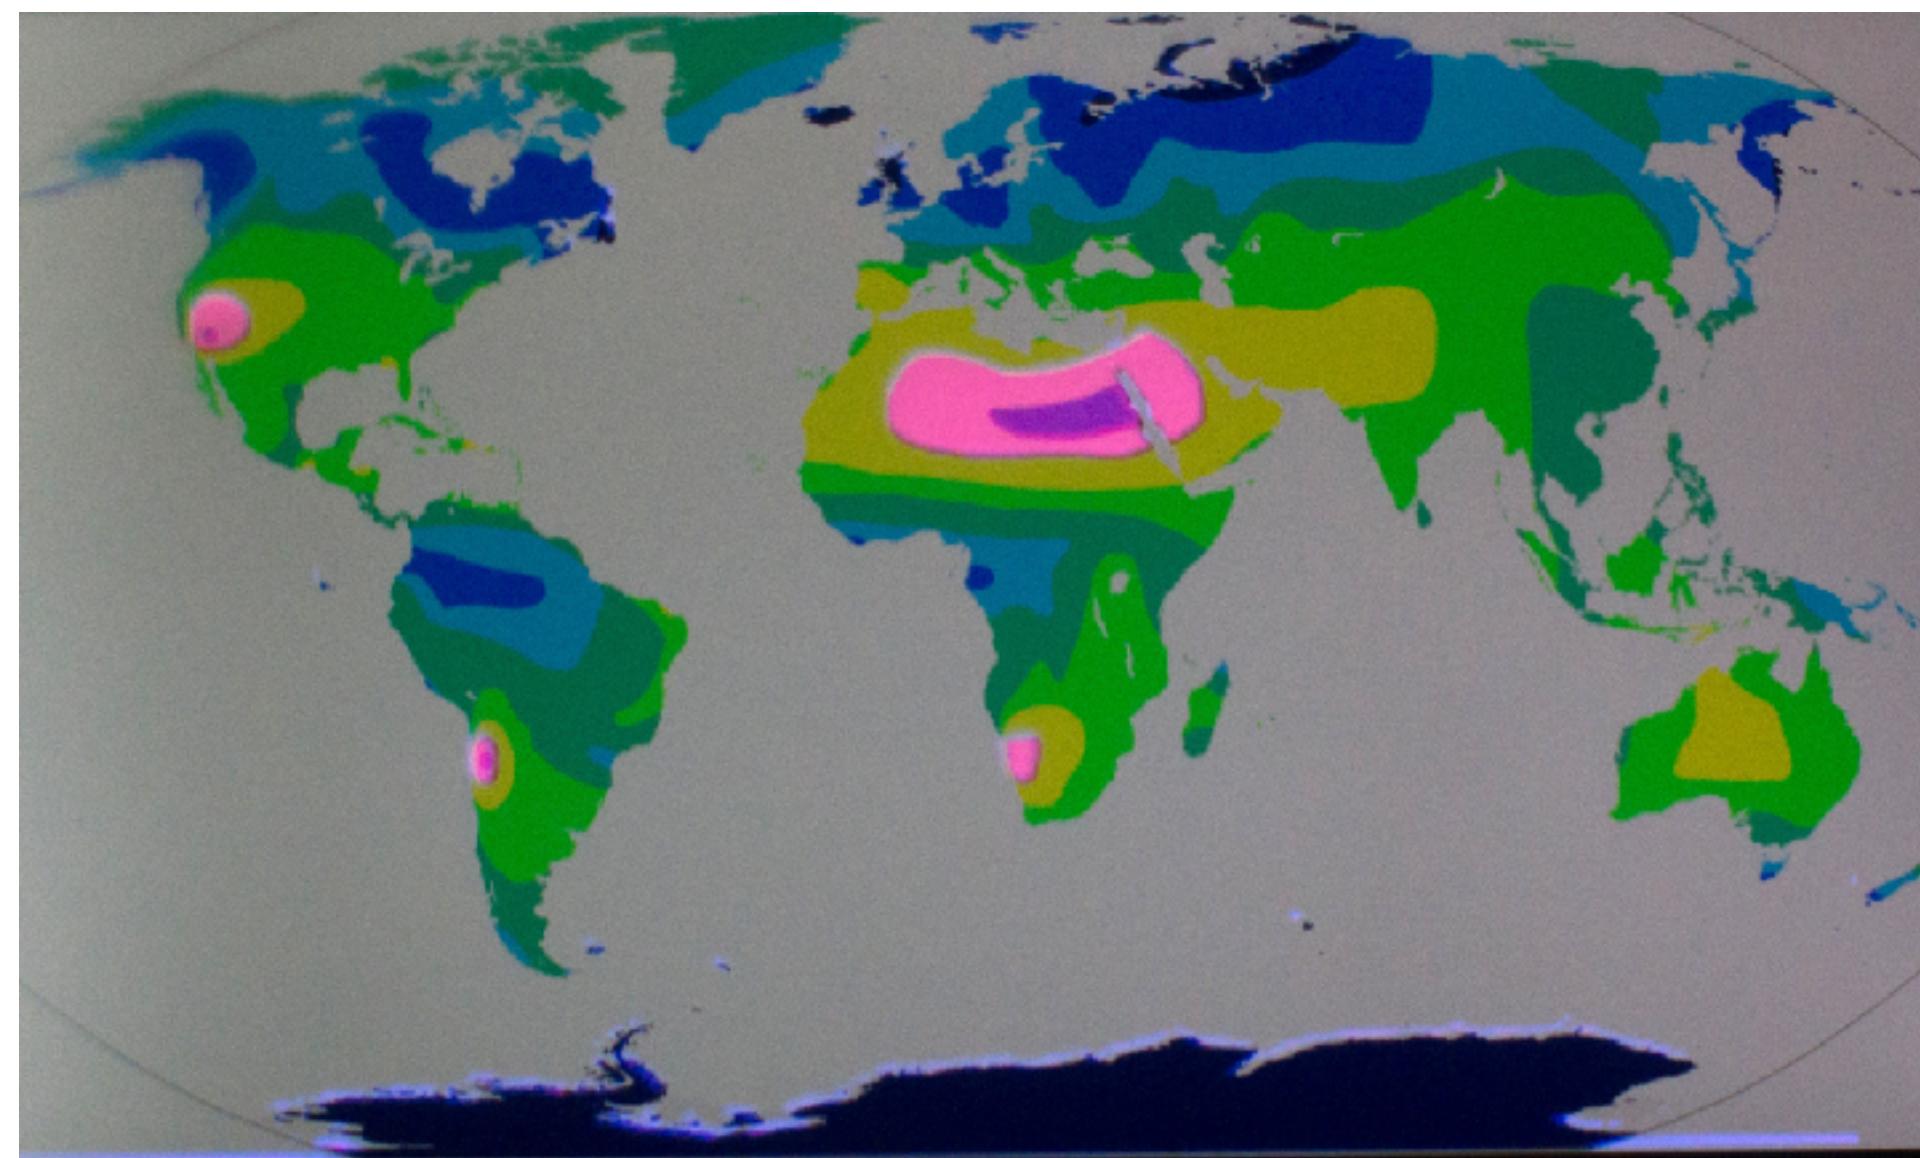
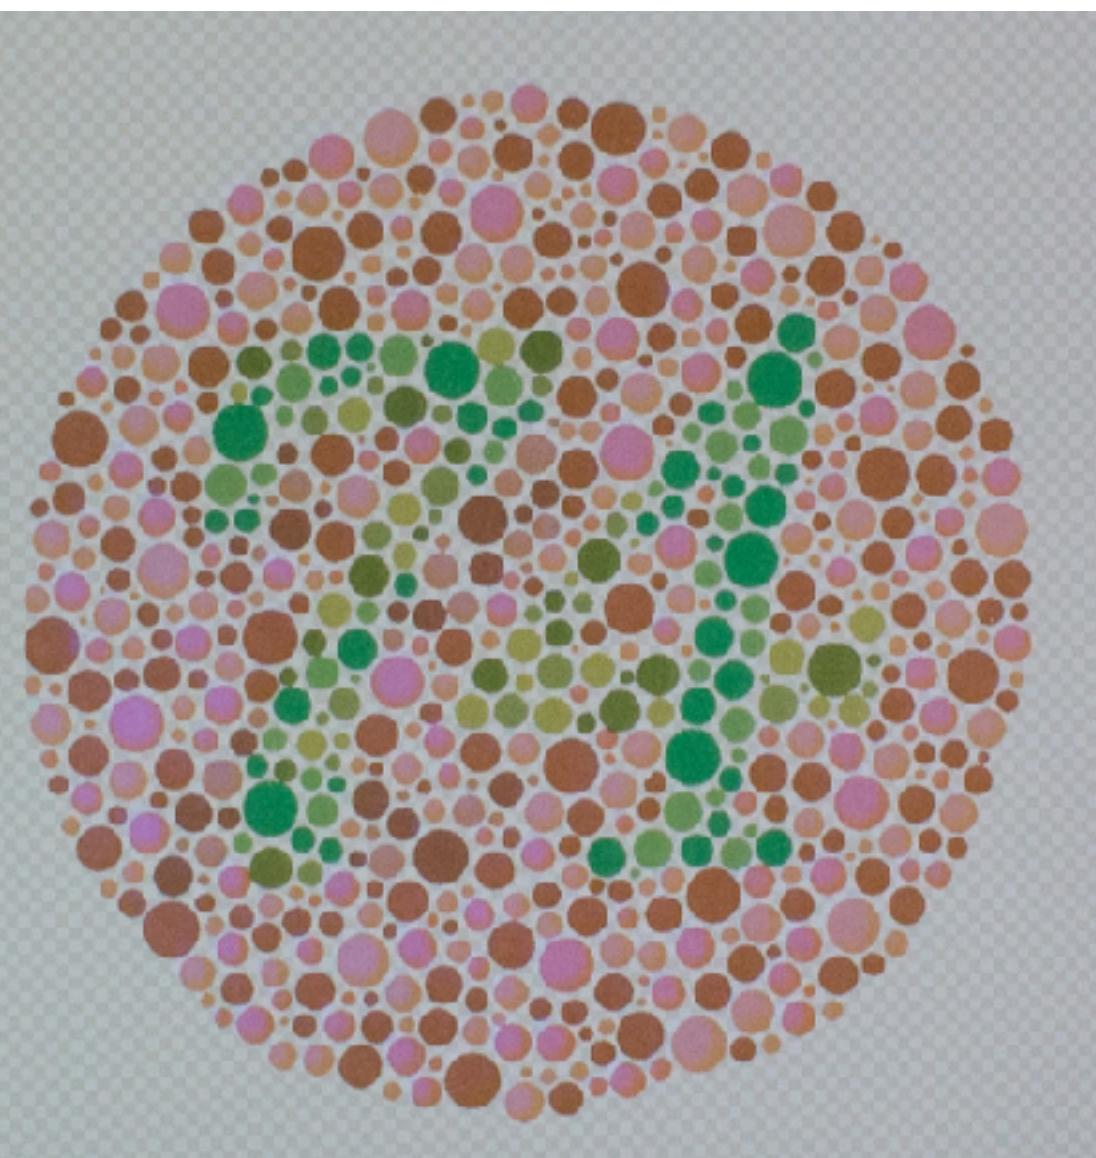
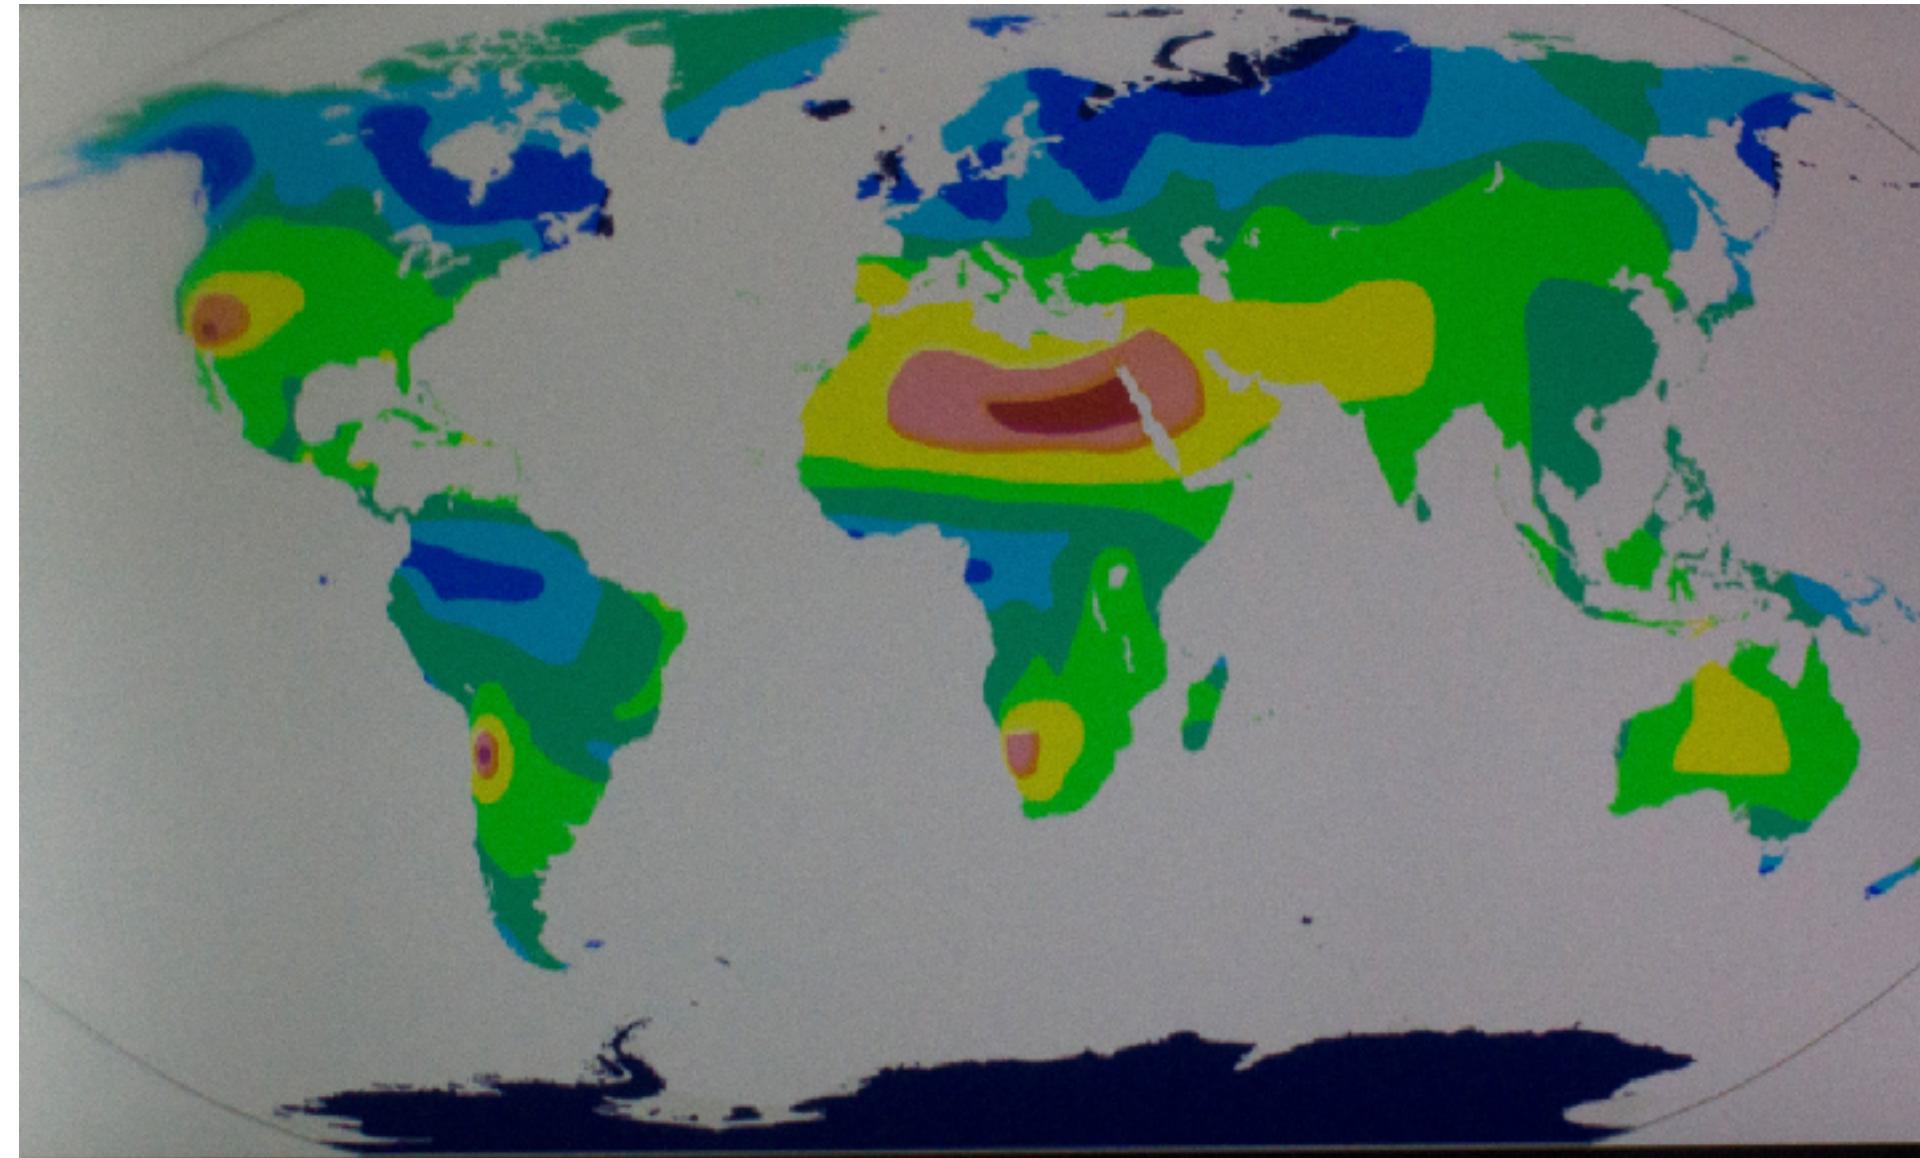
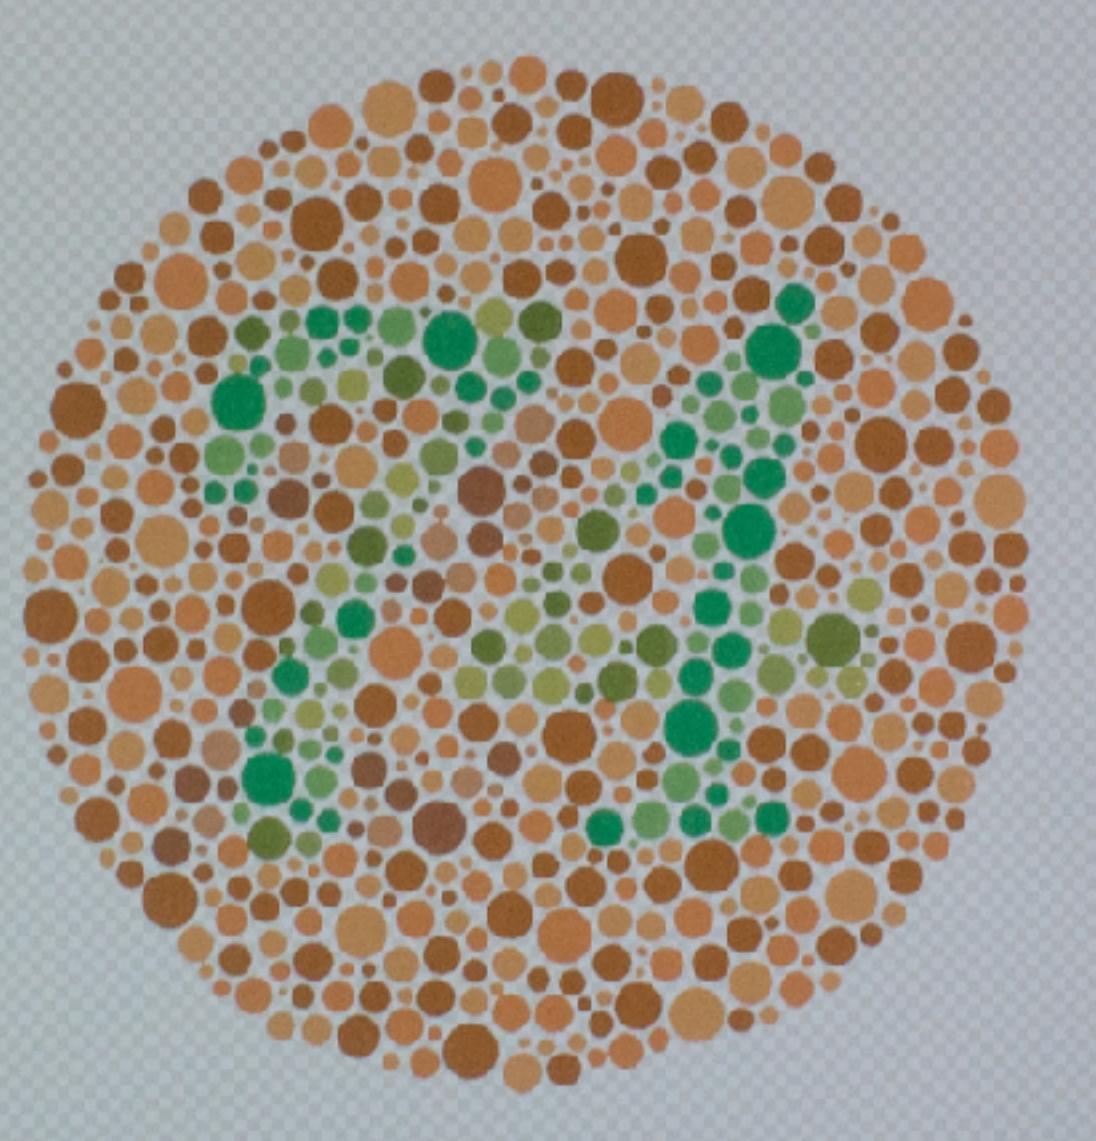
Tritanopia  
(blue blind)





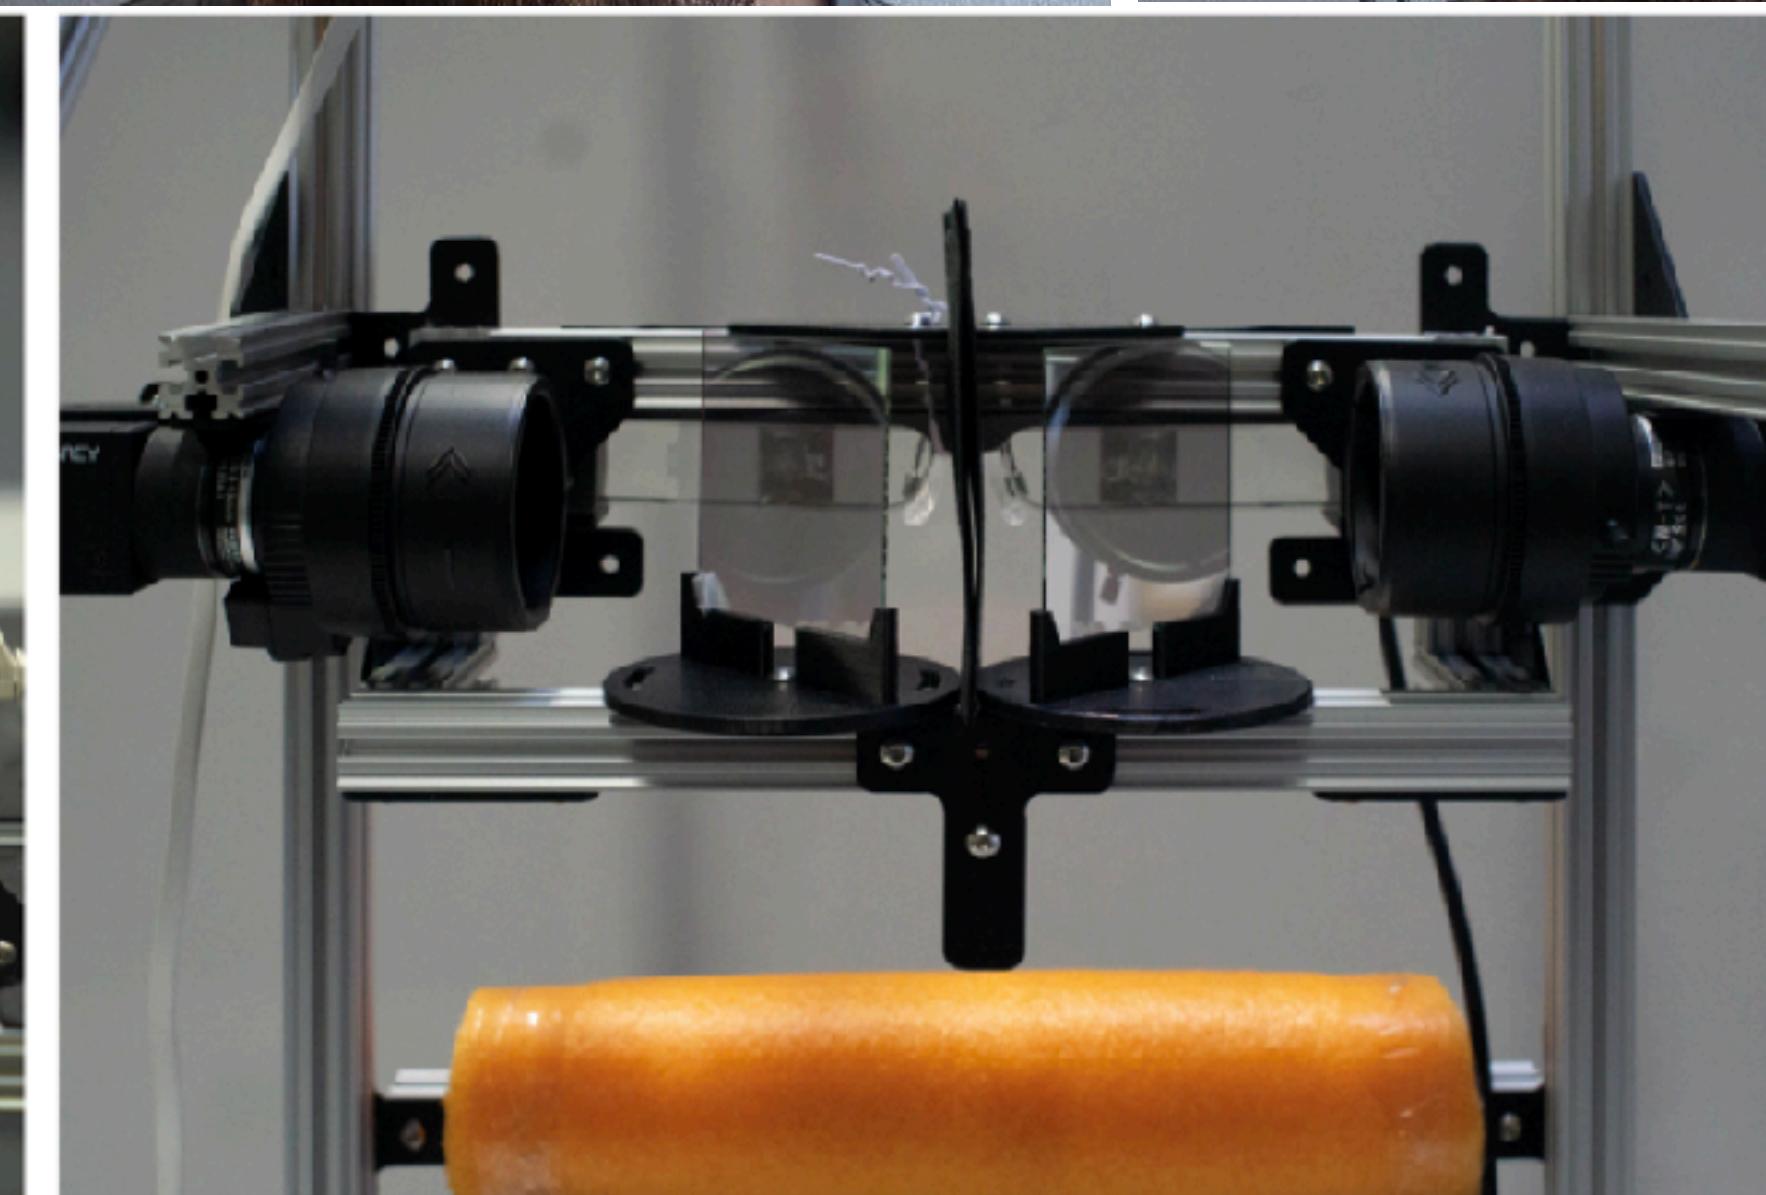
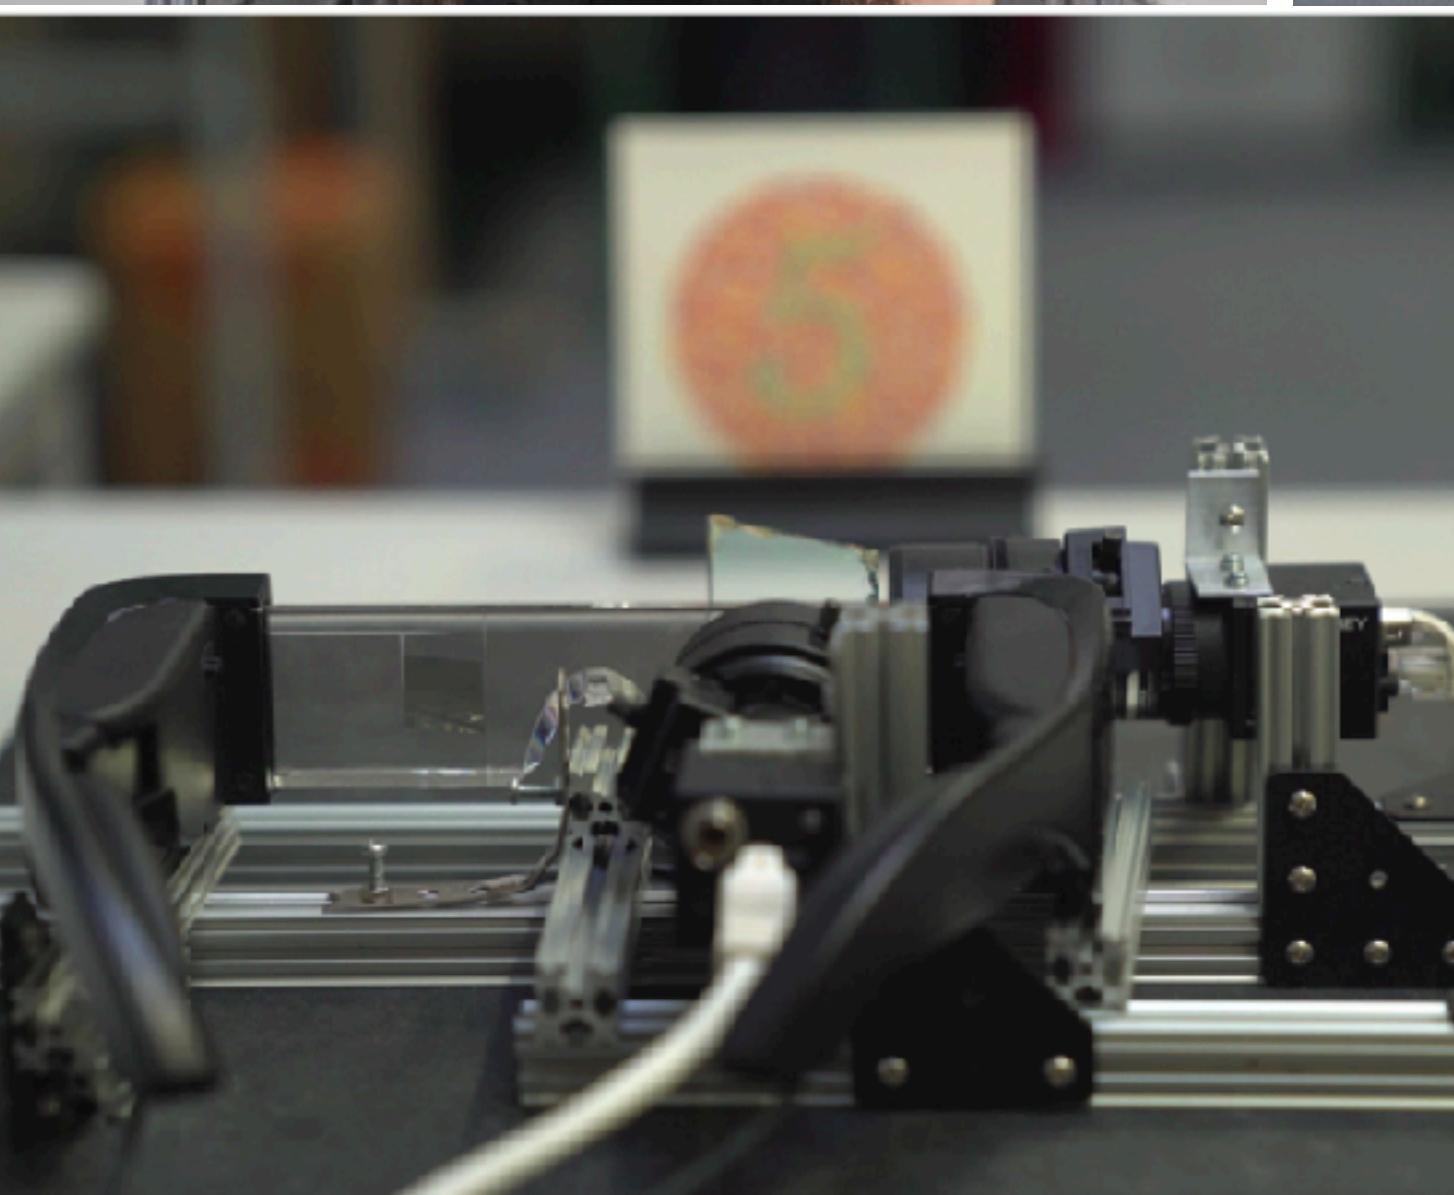
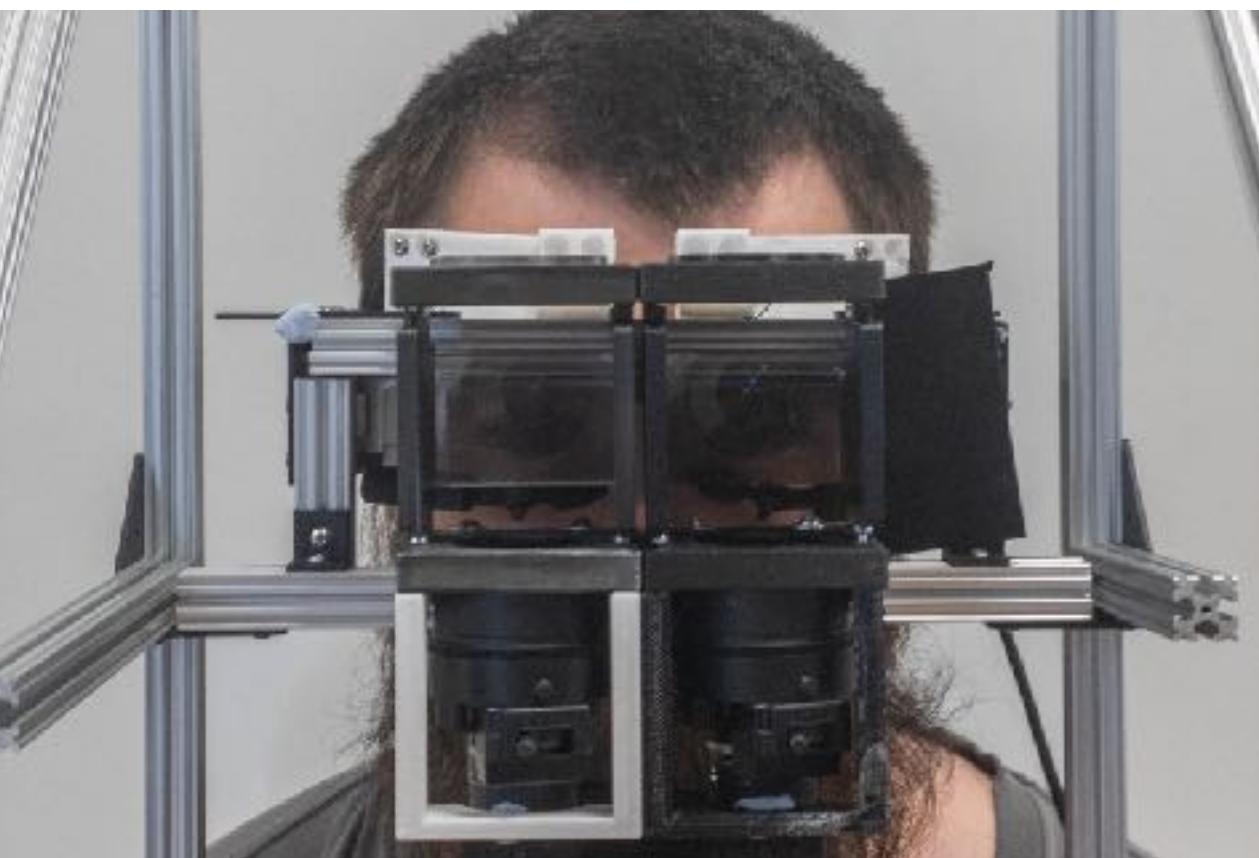
DELL

# Results



Jonathan Sutton, Tobias Langlotz, Alexander Plopski (2022)  
Seeing Colours: Addressing Colour Vision Deficiency with Vision Augmentations using Computational Glasses,  
ACM Transactions on Human-Computer Interaction (ACM TOCHI), 2022

# Prototypes



Tobias Langlotz, Jonathan Sutton, Stefanie Zollmann, Yuta Itoh, and Holger Regenbrecht (2018) ChromaGlasses: Computational Glasses for Compensating Colour Blindness ACM CHI'18 Conference on Human Factors in Computing Systems (ACM CHI 2018), Montreal, 2018,

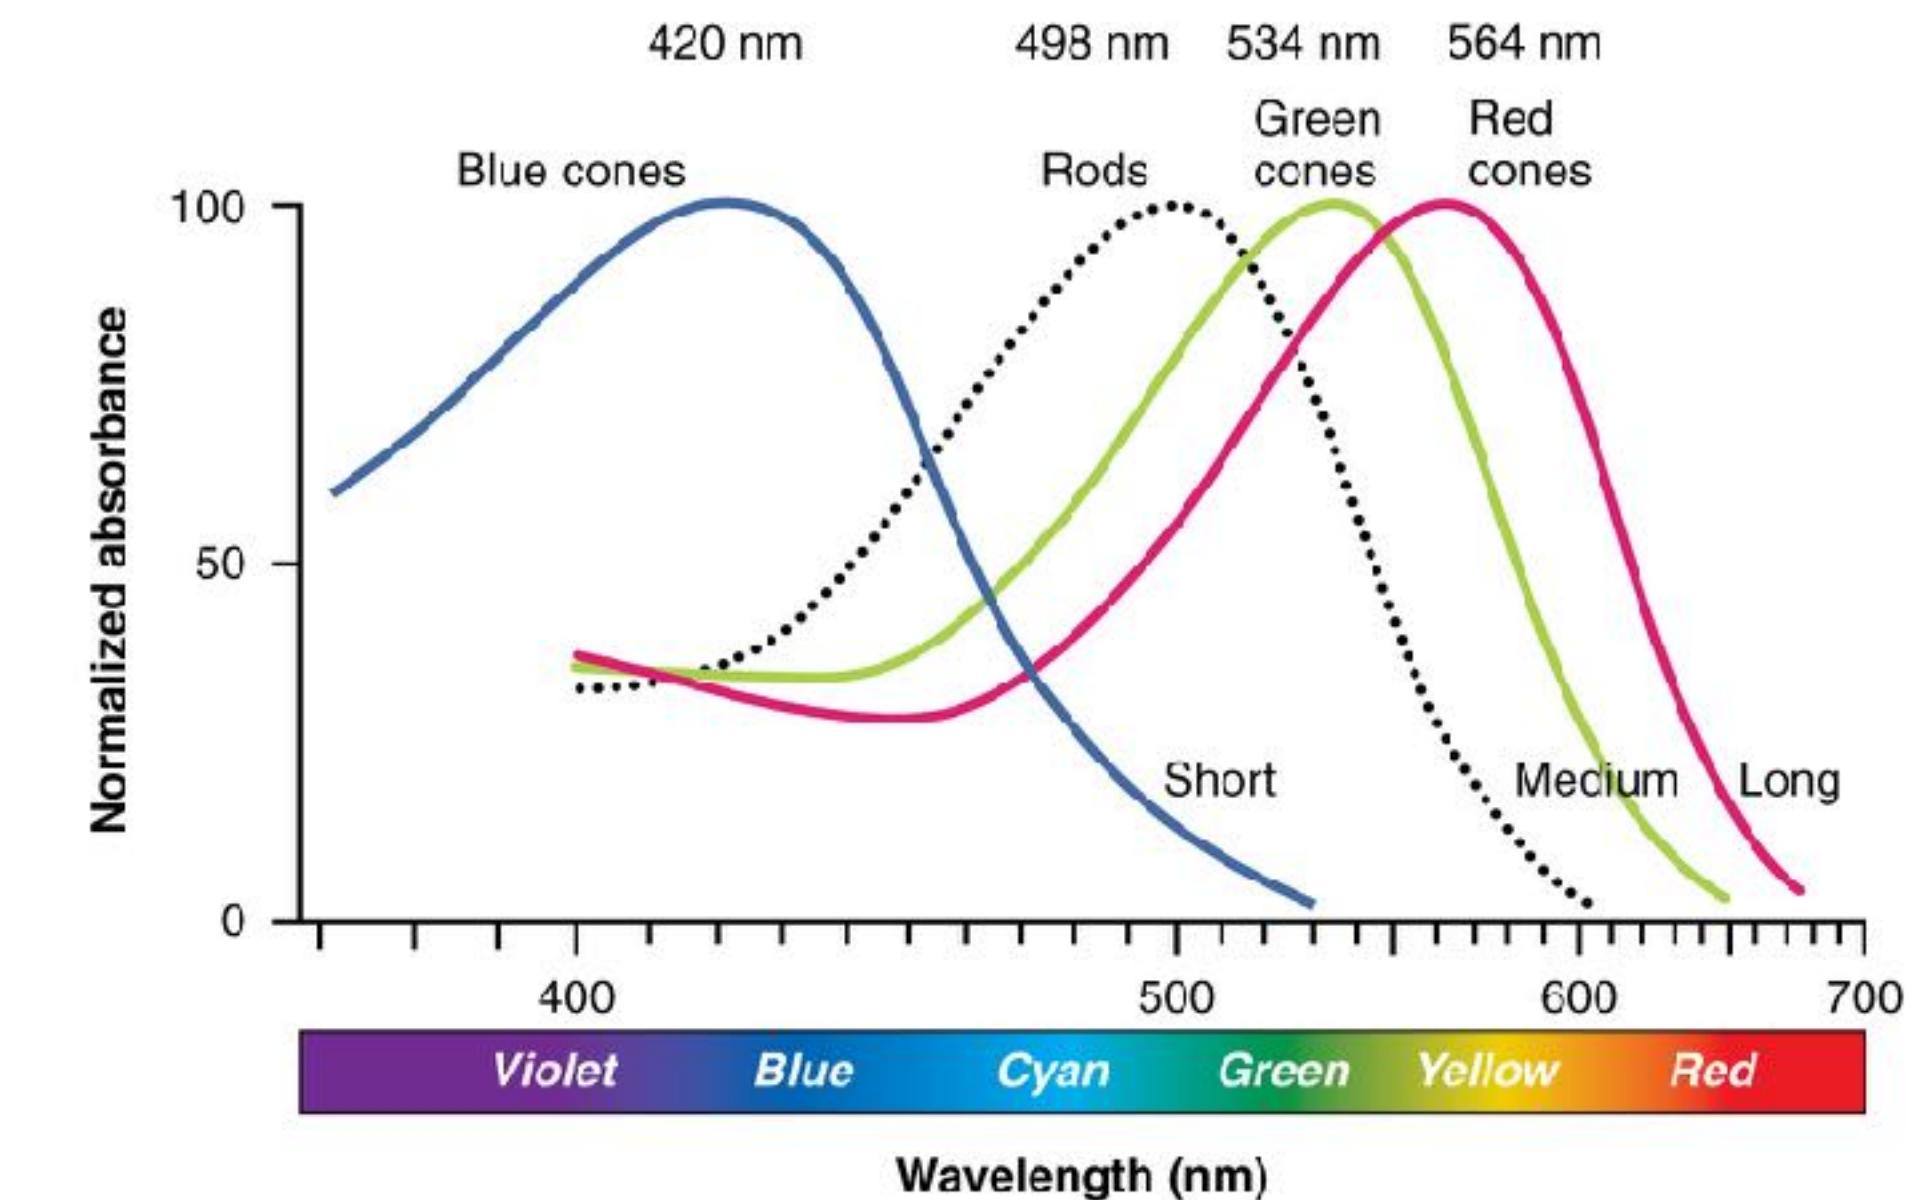
**Honourable Mention Award** 🌟

Jonathan Sutton, Tobias Langlotz, Alexander Plopski (2022)  
Seeing Colours: Addressing Colour Vision Deficiency with Vision Augmentations using Computational Glasses,  
ACM Transactions on Human-Computer Interaction (ACM TOCHI), 2022

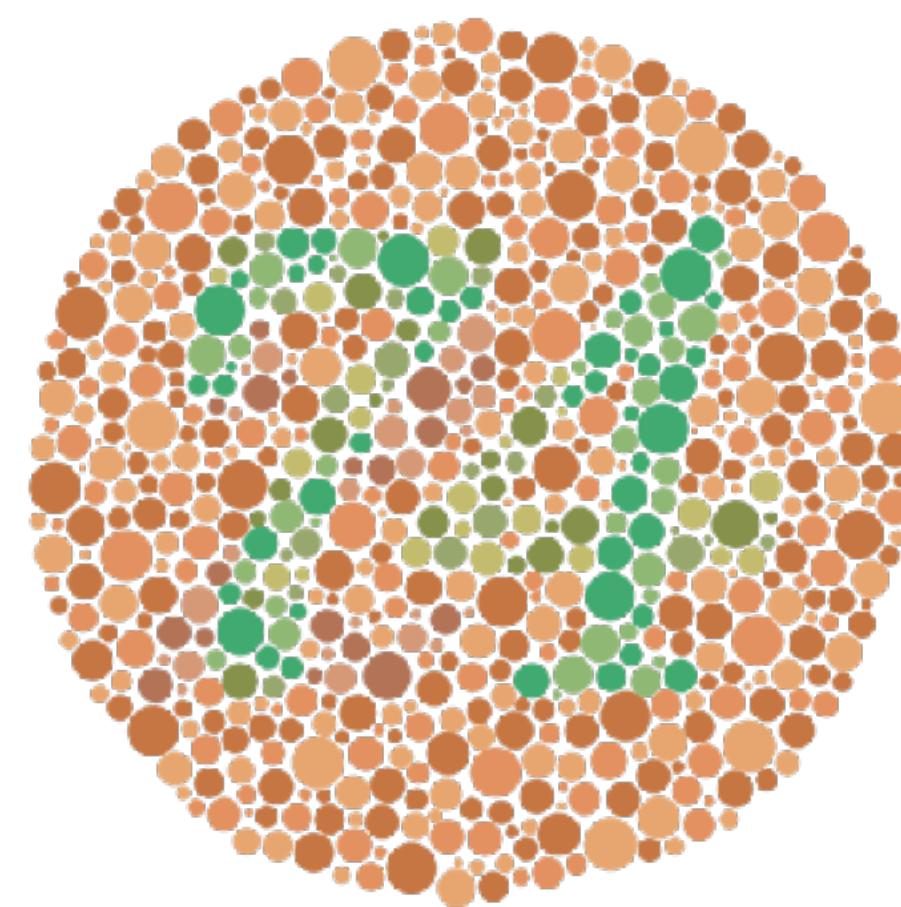
**Back to topic....**

# Colour Perception

- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones



Monochromacy



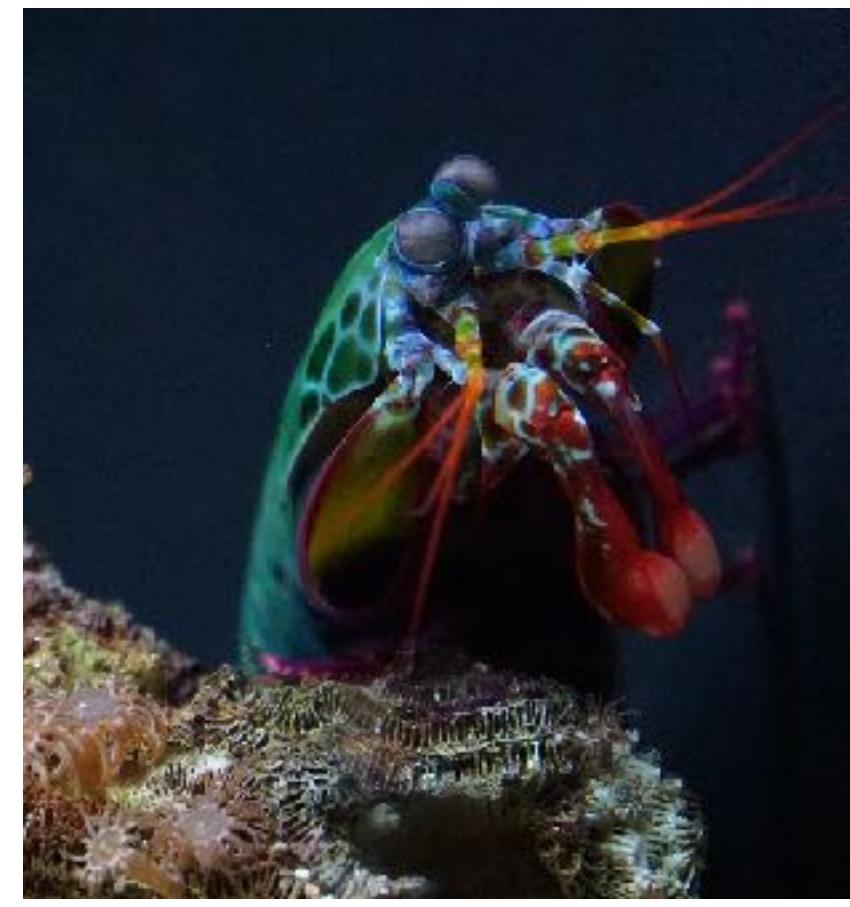
Dichromacy



Trichromacy



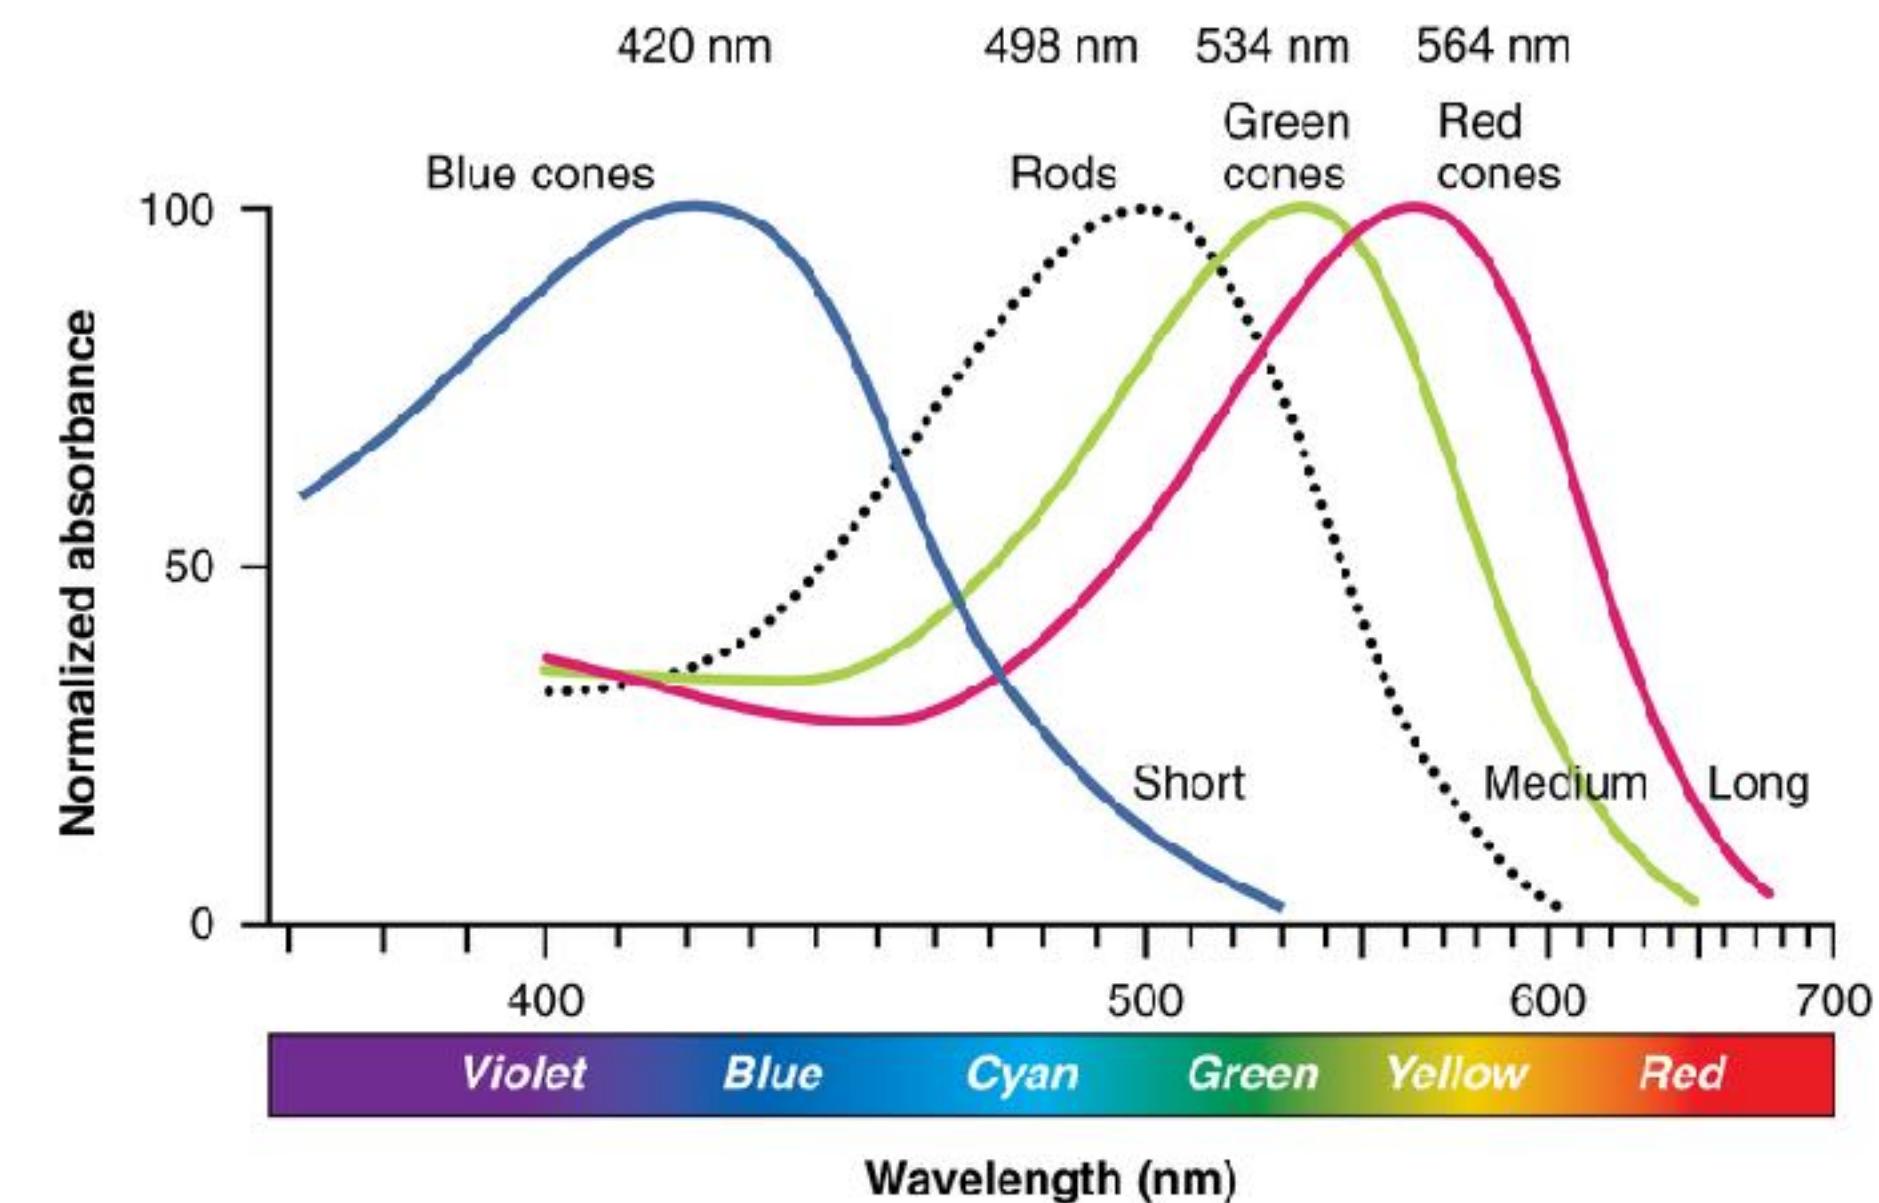
Tetrachromacy



Dodecachromacy  
to Hexadecachromacy

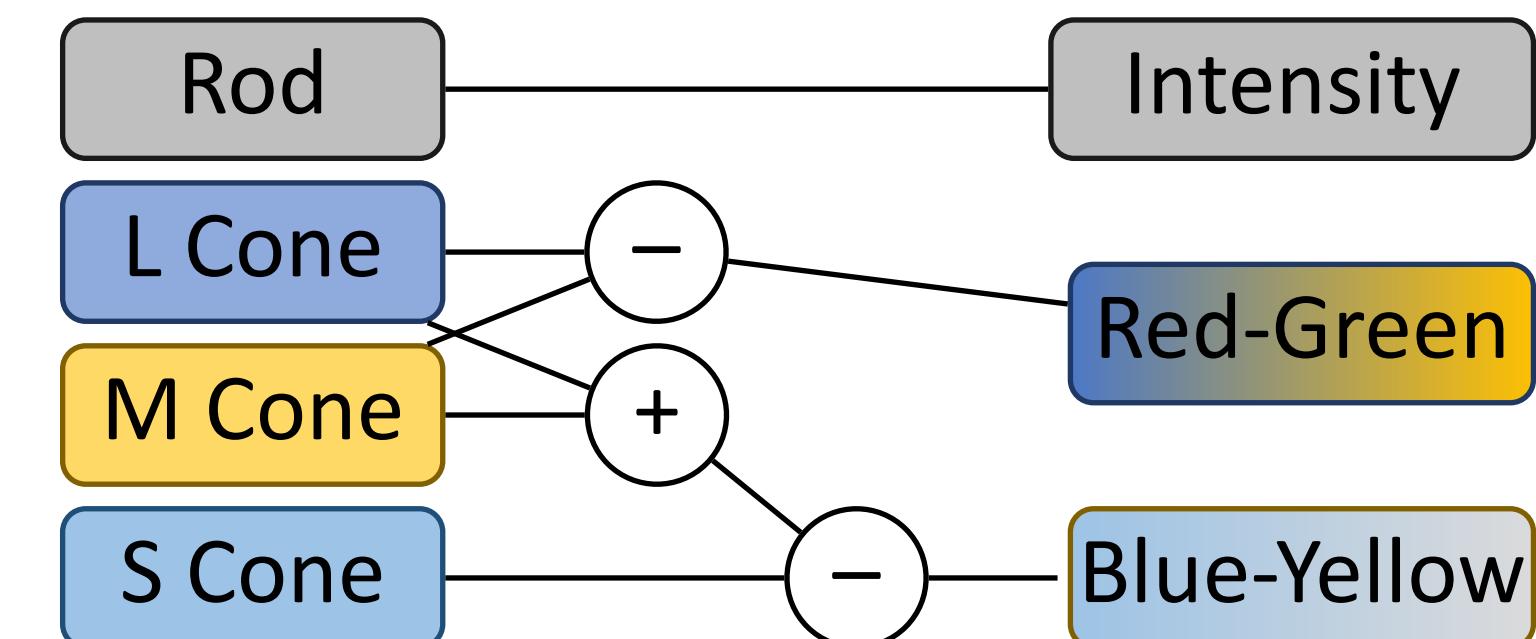
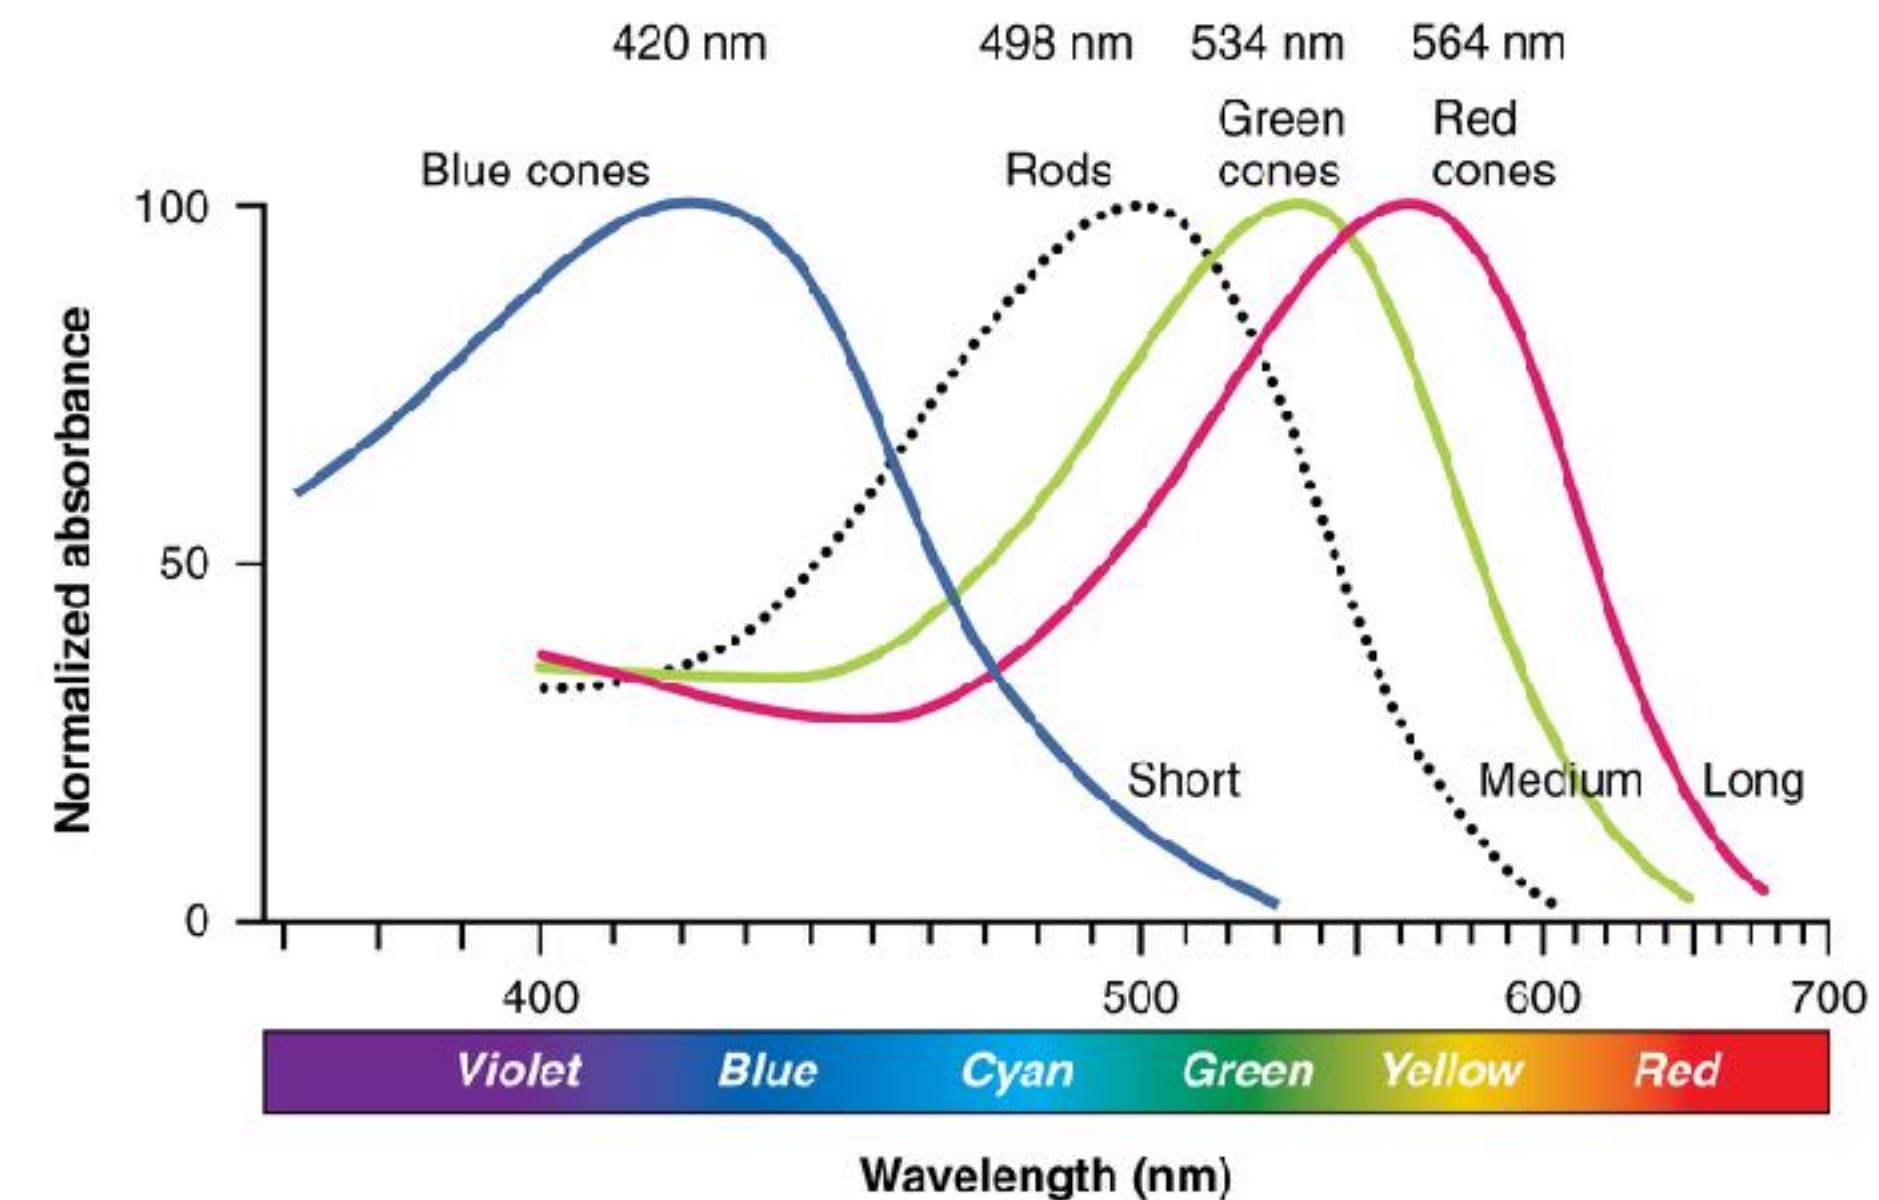
# Colour Perception

- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones



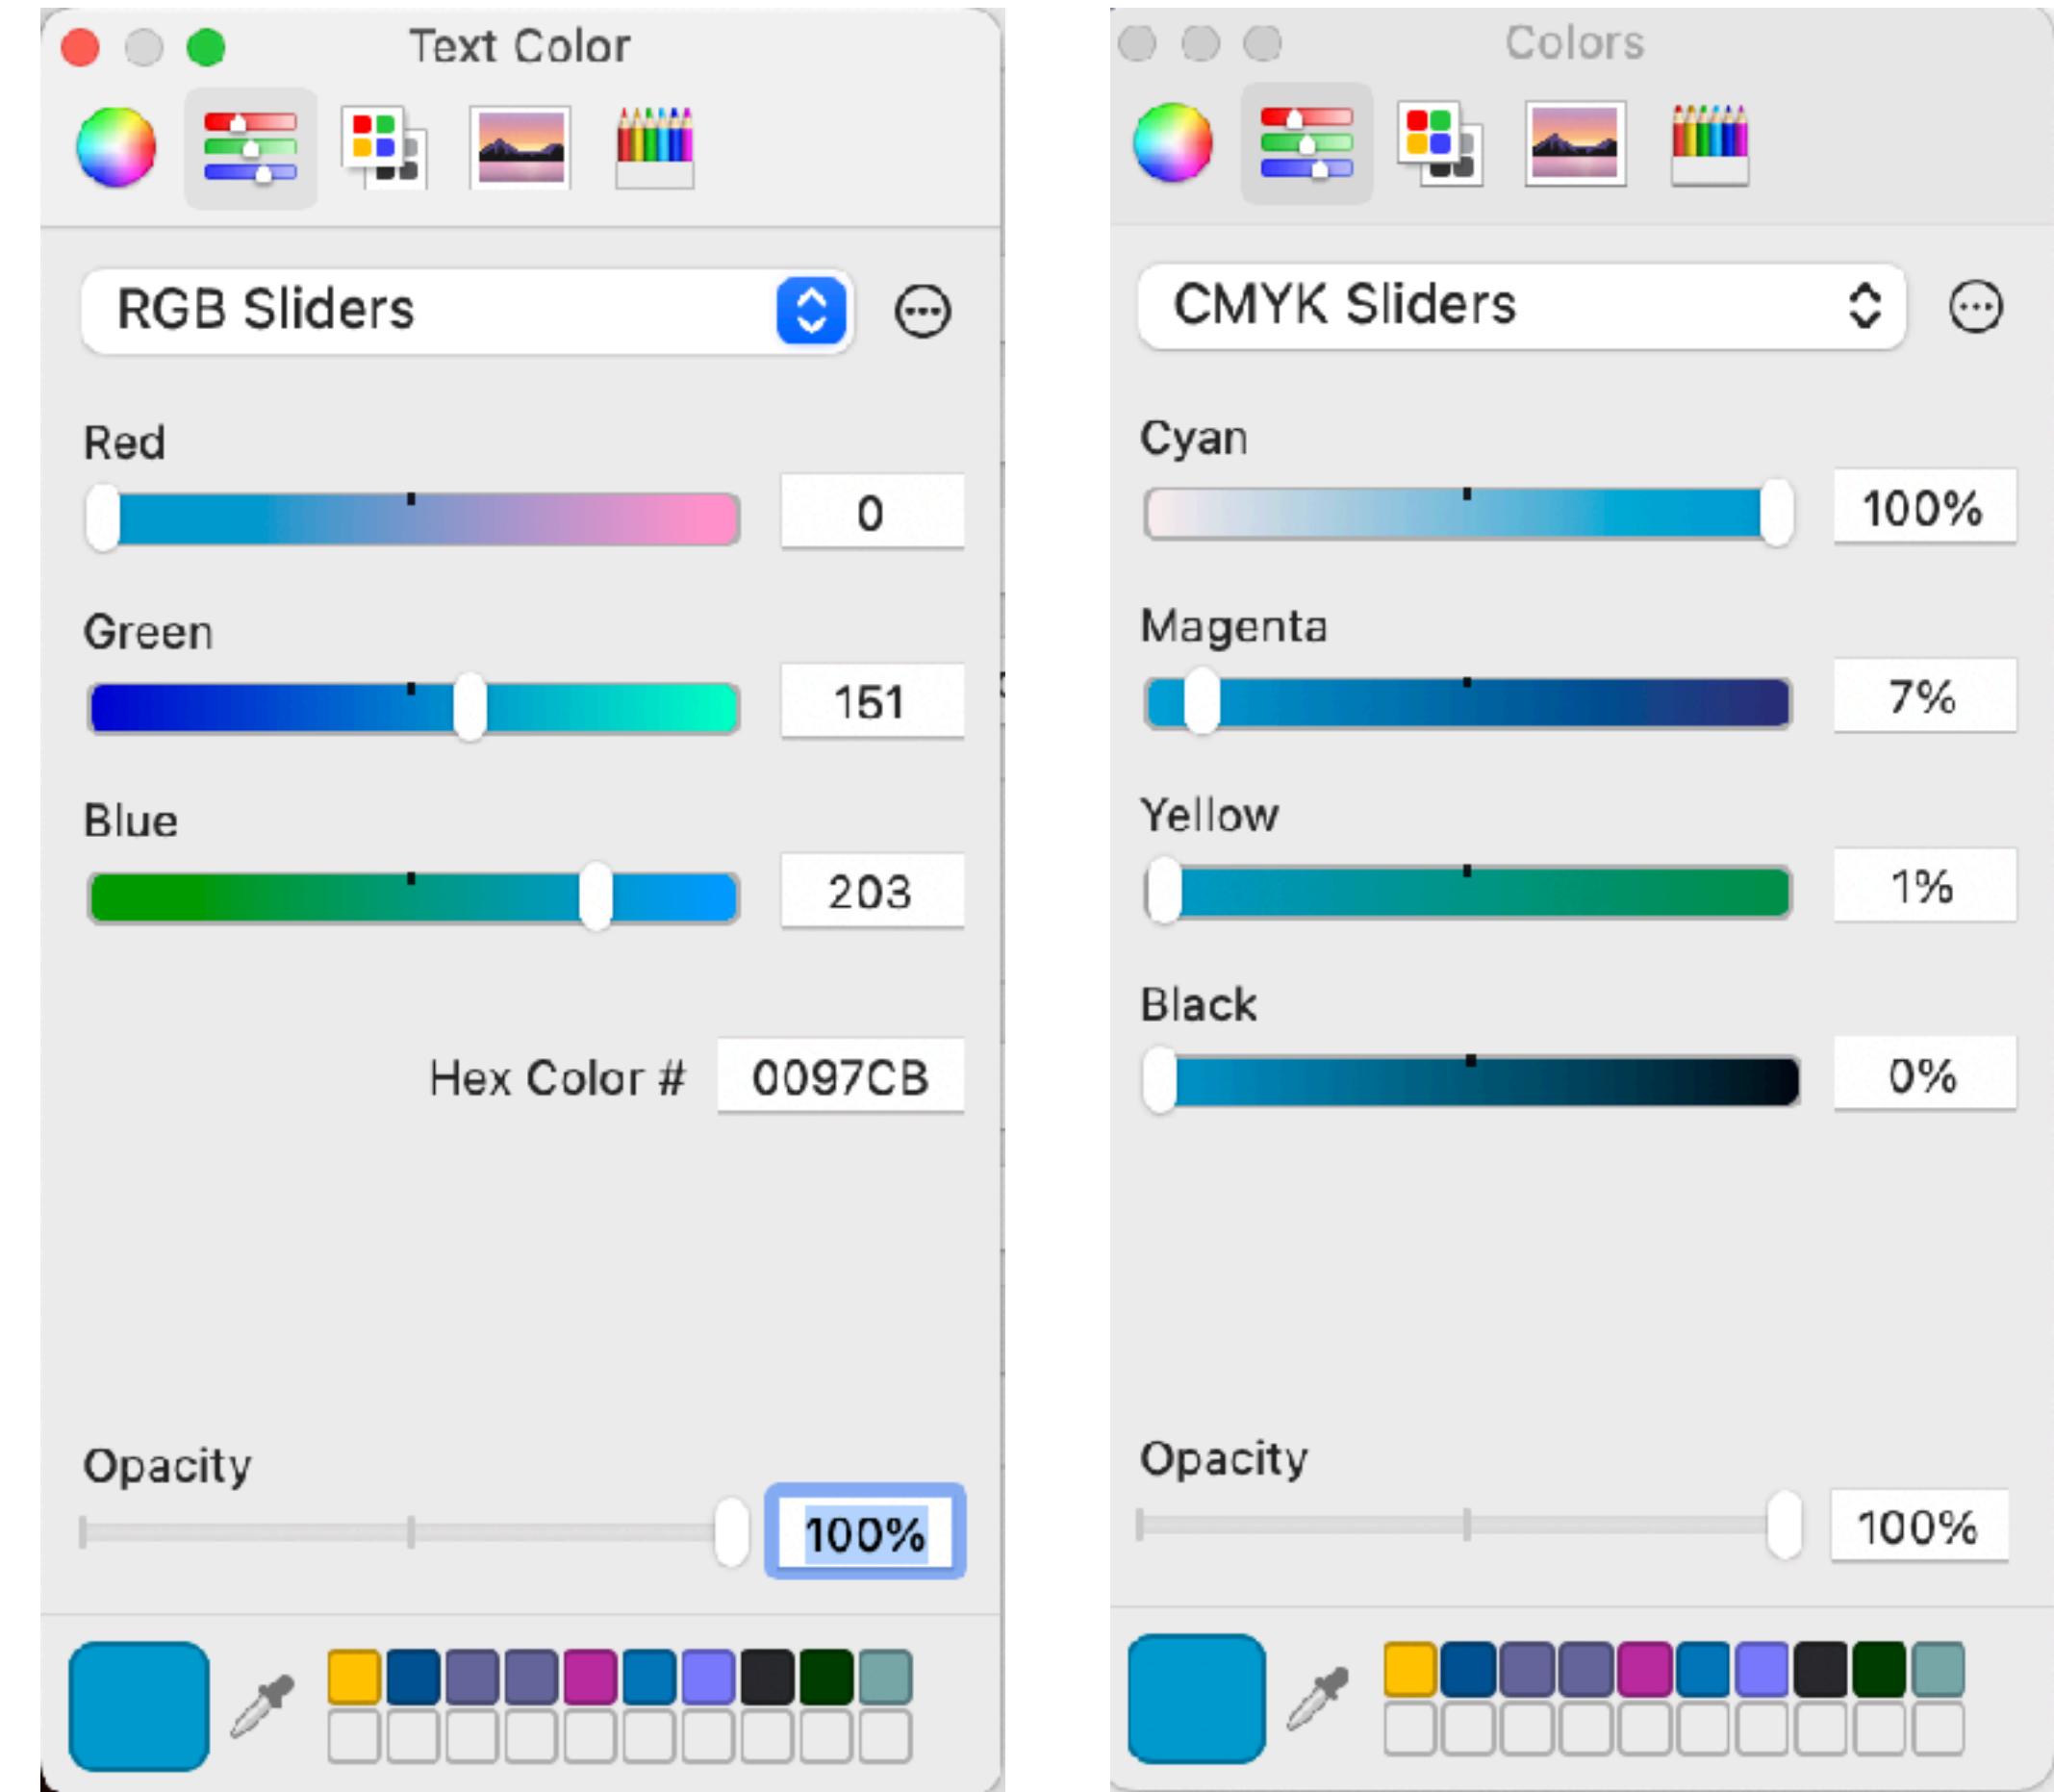
# Colour Perception

- We perceive light on our retina
  - Four types of cells (usually)
  - Rods – just sensitive to brightness
  - 3 types of Cones – sensitive to short/medium/long wavelengths (trichromatic theory of color vision)
  - Often called ‘blue’, ‘green’ or ‘red’ cones
- Brain does not receive raw signal
  - Rather, differences in cone responses
  - One axis is  $L - M$ , or red–green
  - The other is  $S - (L + M)$ , or blue–yellow



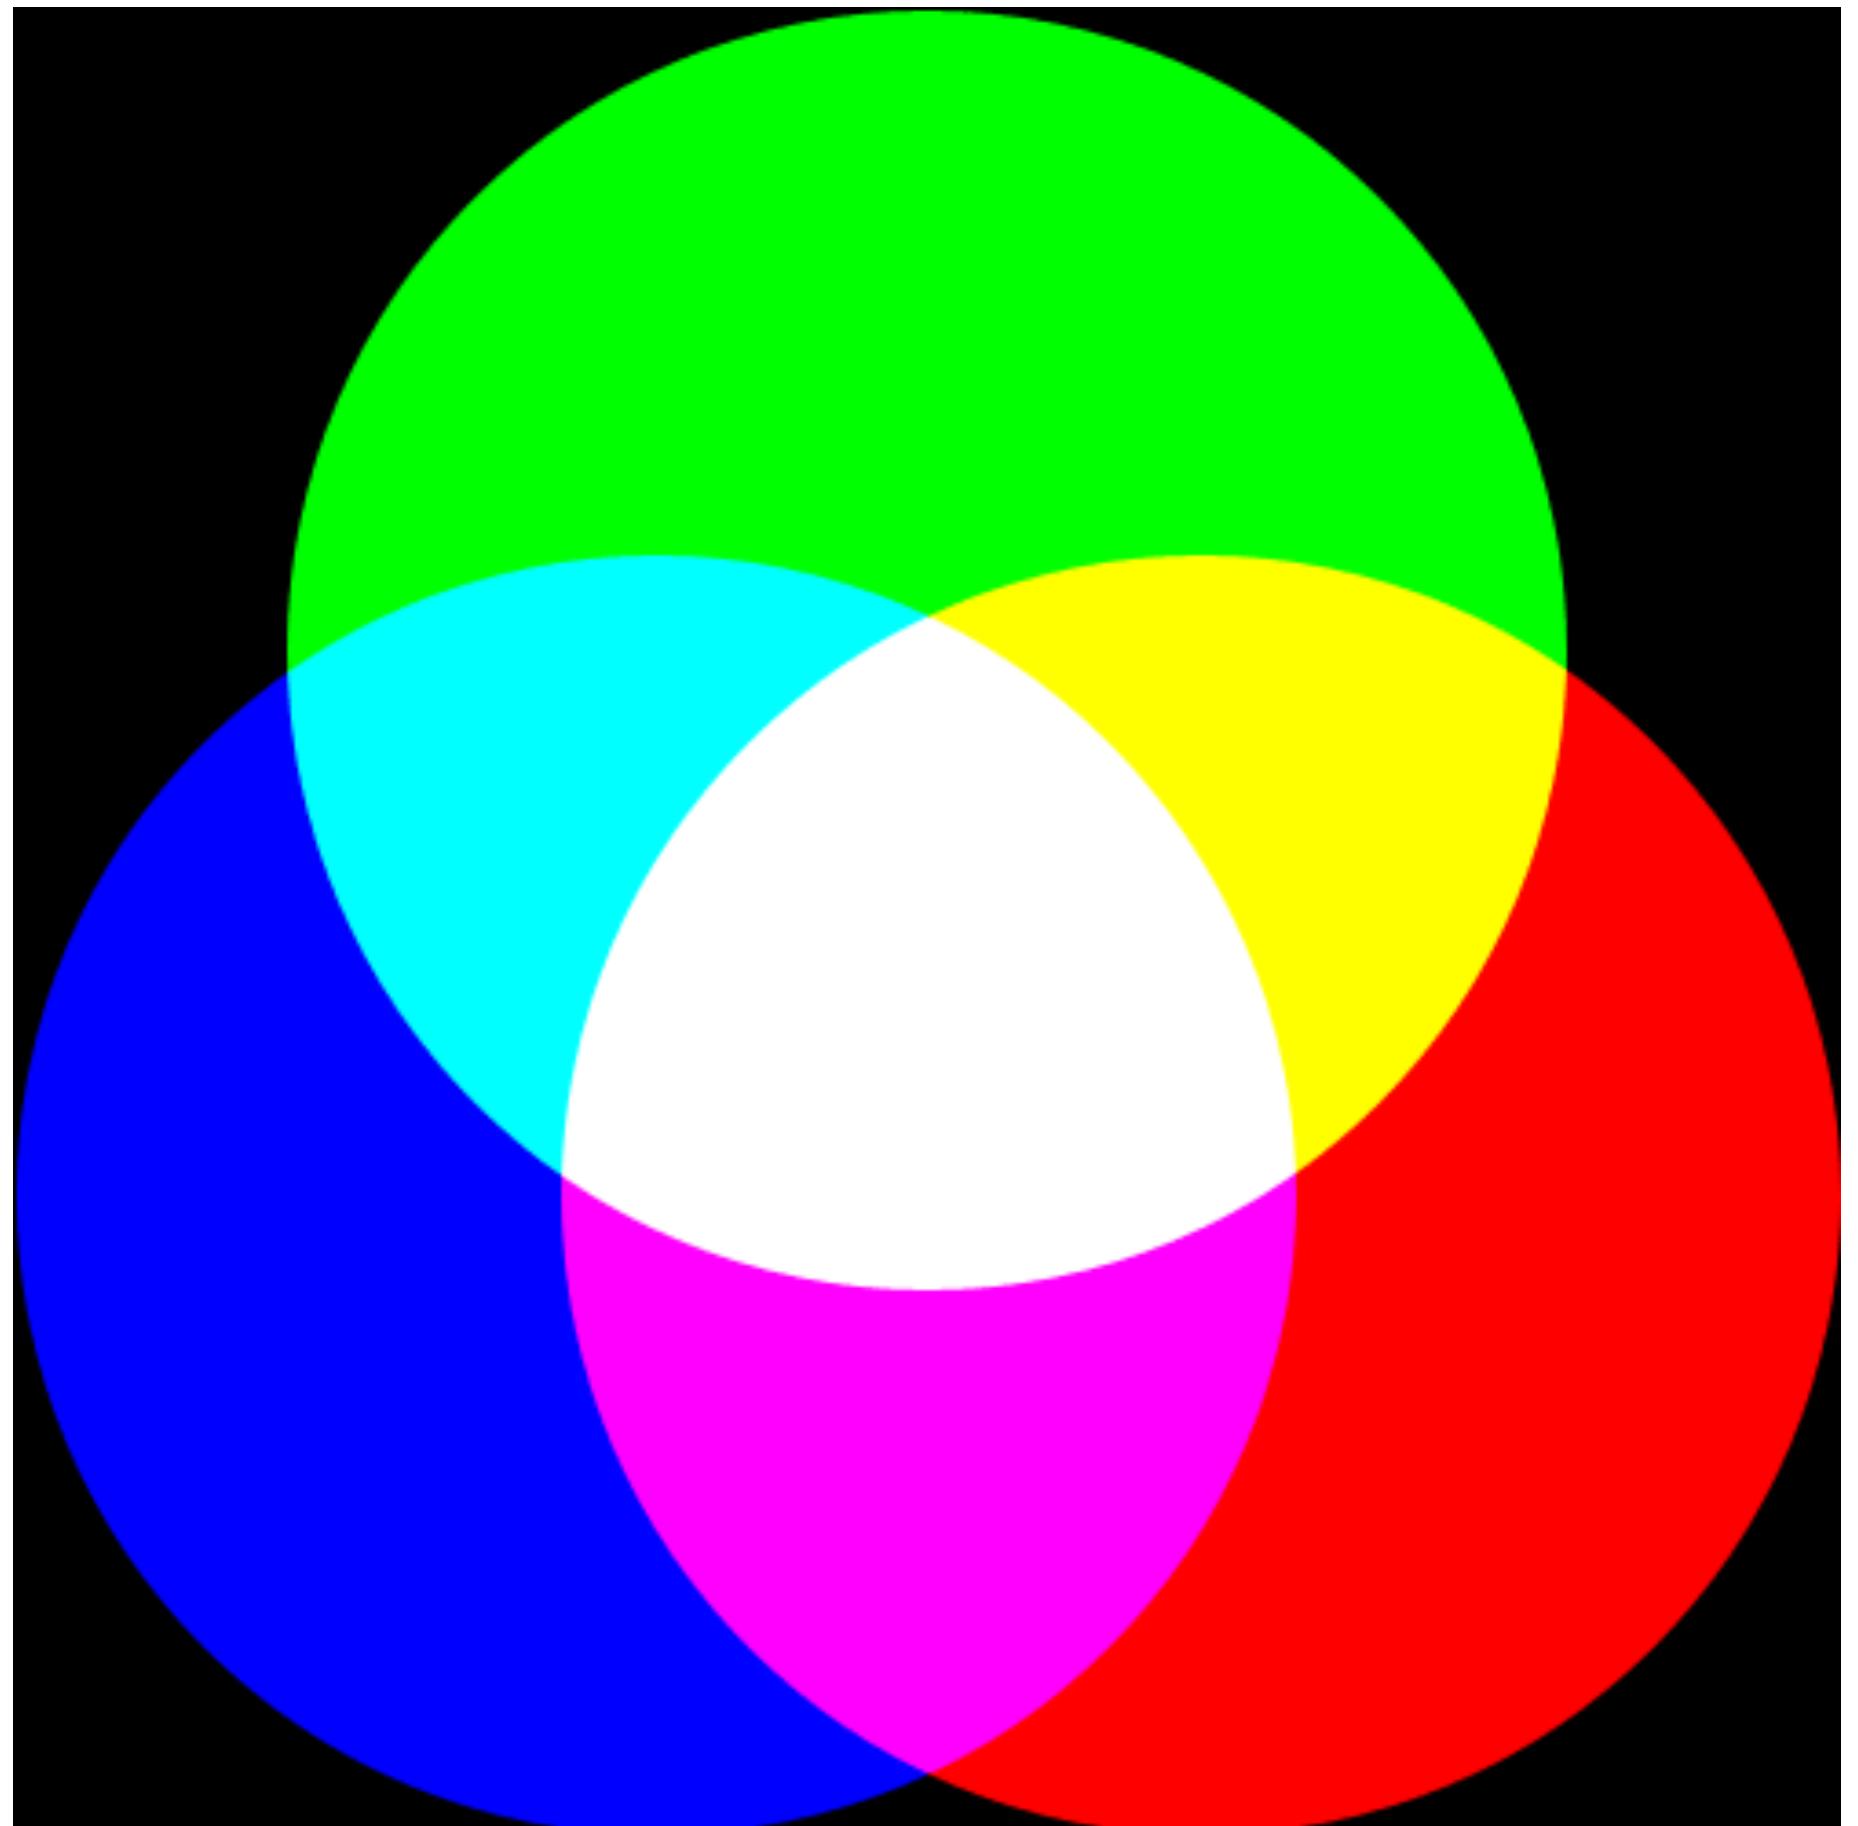
# Representing Colour

- We perceive 3 colour dimensions
  - Short, medium, long cones
  - Intensity, Red-Green, Blue-Yellow
- Colour representations are usually 3 dimensional
  - Red-Green-Blue -> RGB
  - Cyan-Magenta-Yellow -> CMY (+K)
  - Hue-Saturation-Value ->HSV or HSL
  - And others (e.g. YUV, CIE)



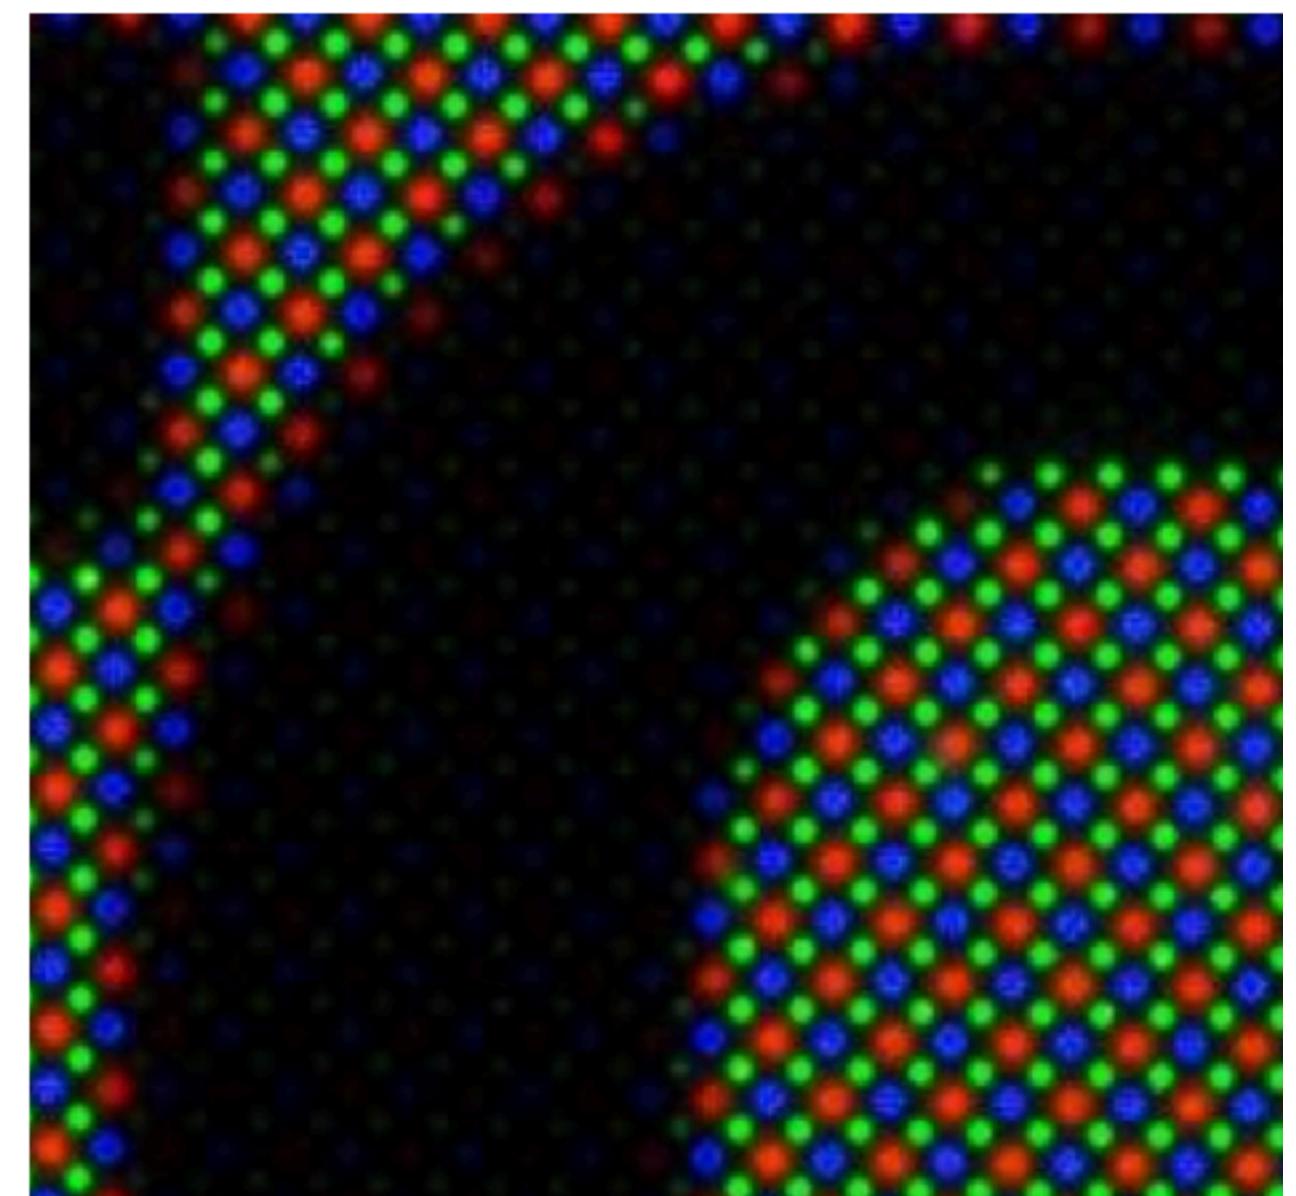
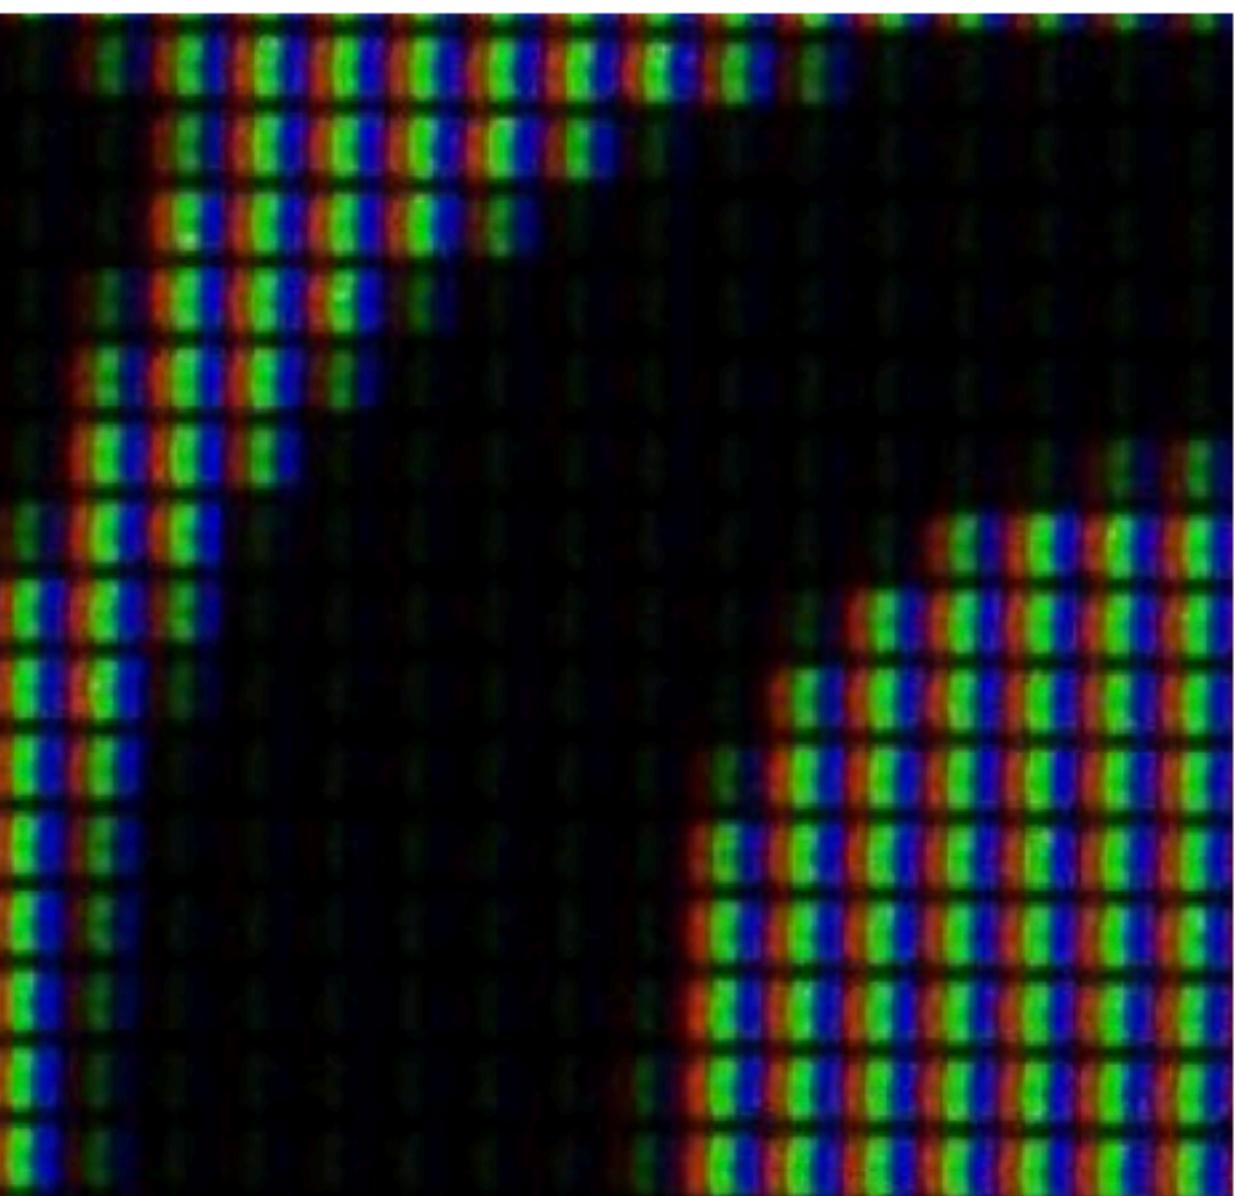
# RGB Model

- Most common in computing
  - Most monitors and projectors
  - Additive colour model
  - Not perceptually uniform
- Record/transmit RGB channels
  - Aligns to the types of cone in the eye
  - Mixes three primary light colours
  - Can represent most colours (not all?)
  - Typically within the range [0,255] or [0,1]
    - E.g. Red as (255,0,0) or (1,0,0)
  - Can be abbreviated as hex values:
    - (183, 24, 92) -> #B7185C



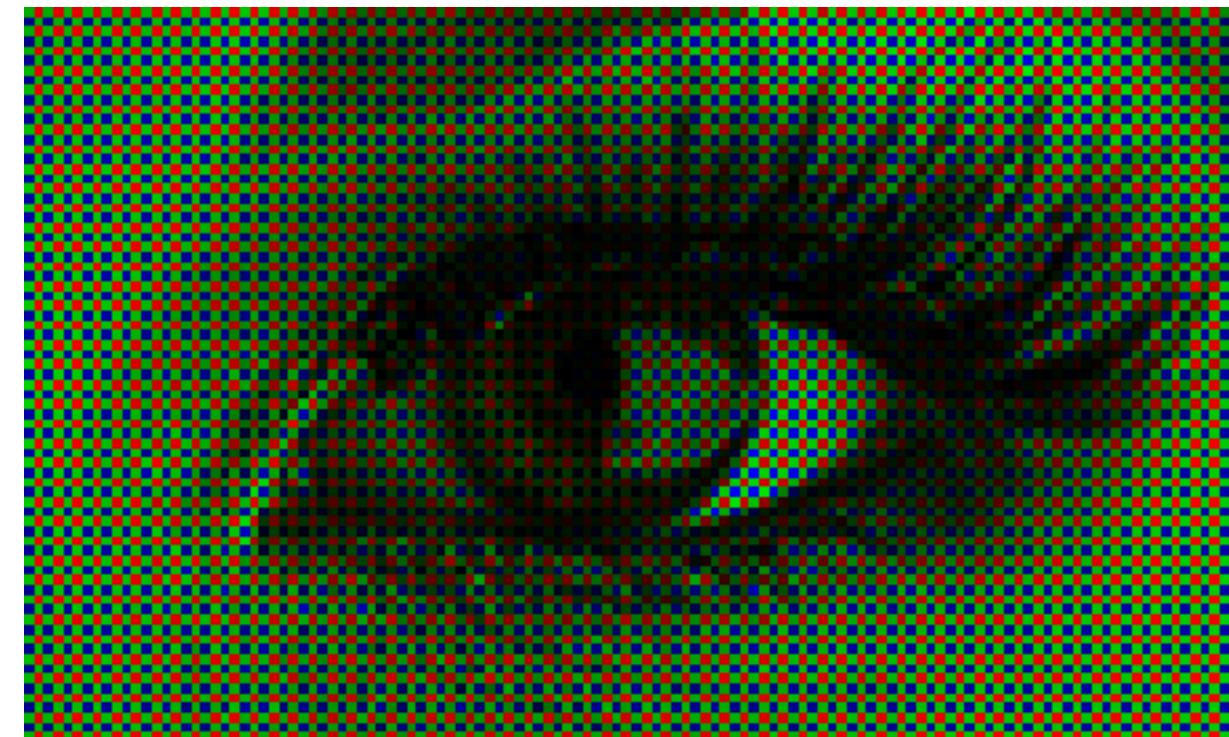
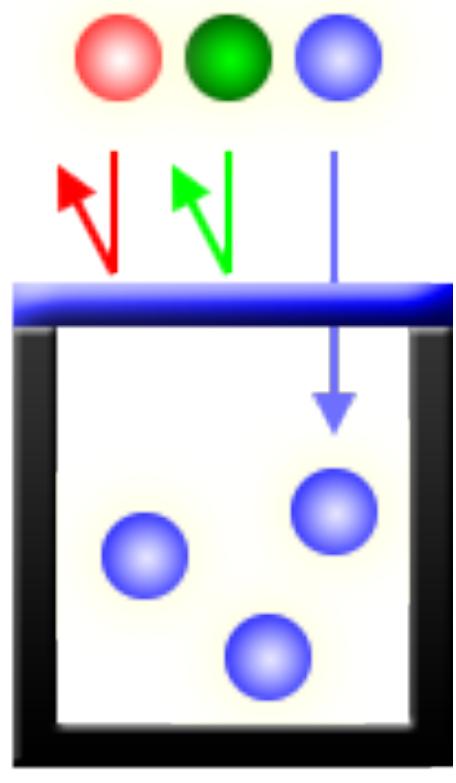
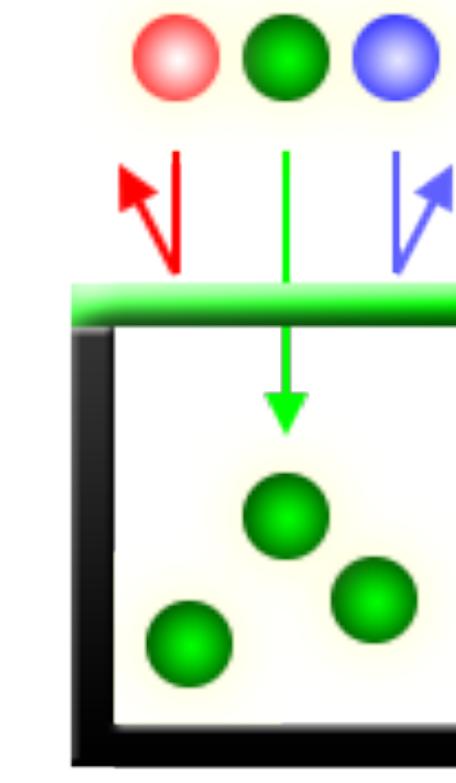
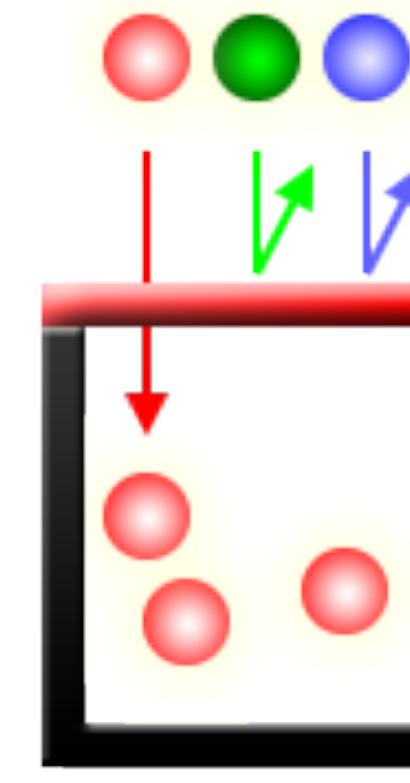
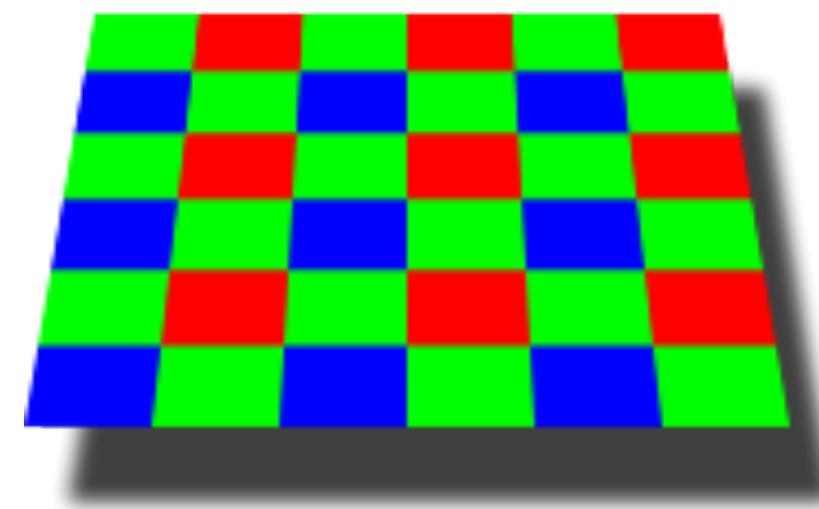
# RGB Model

- Most common in computing
  - Most monitors and projectors
  - Additive colour model
  - Not perceptually uniform
- Record/transmit RGB channels
  - Aligns to the types of cone in the eye
  - Mixes three primary light colours
  - Can represent most colours (not all?)
  - Typically within the range [0,255] or [0,1]
    - E.g. Red as (255,0,0) or (1,0,0)
  - Can be abbreviated as hex values:
    - (183, 24, 92) -> #B7185C

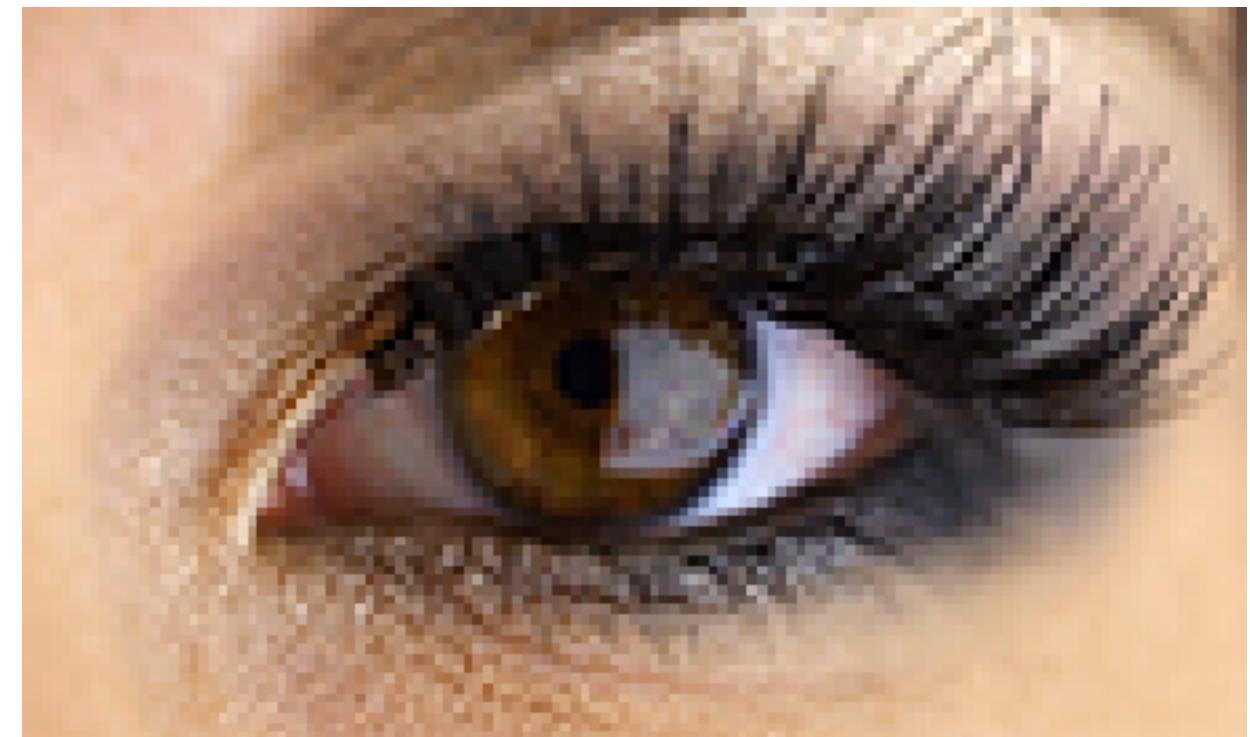


# RGB Model

- Most common in computing
  - Most monitors and projectors
  - Additive colour model
  - Not perceptually uniform
- Record/transmit RGB channels
  - Aligns to the types of cone in the eye
  - Mixes three primary light colours
  - Can represent most colours (not all?)
  - Typically within the range [0,255] or [0,1]
    - E.g. Red as (255,0,0) or (1,0,0)
  - Can be abbreviated as hex values:
    - (183, 24, 92) -> #B7185C



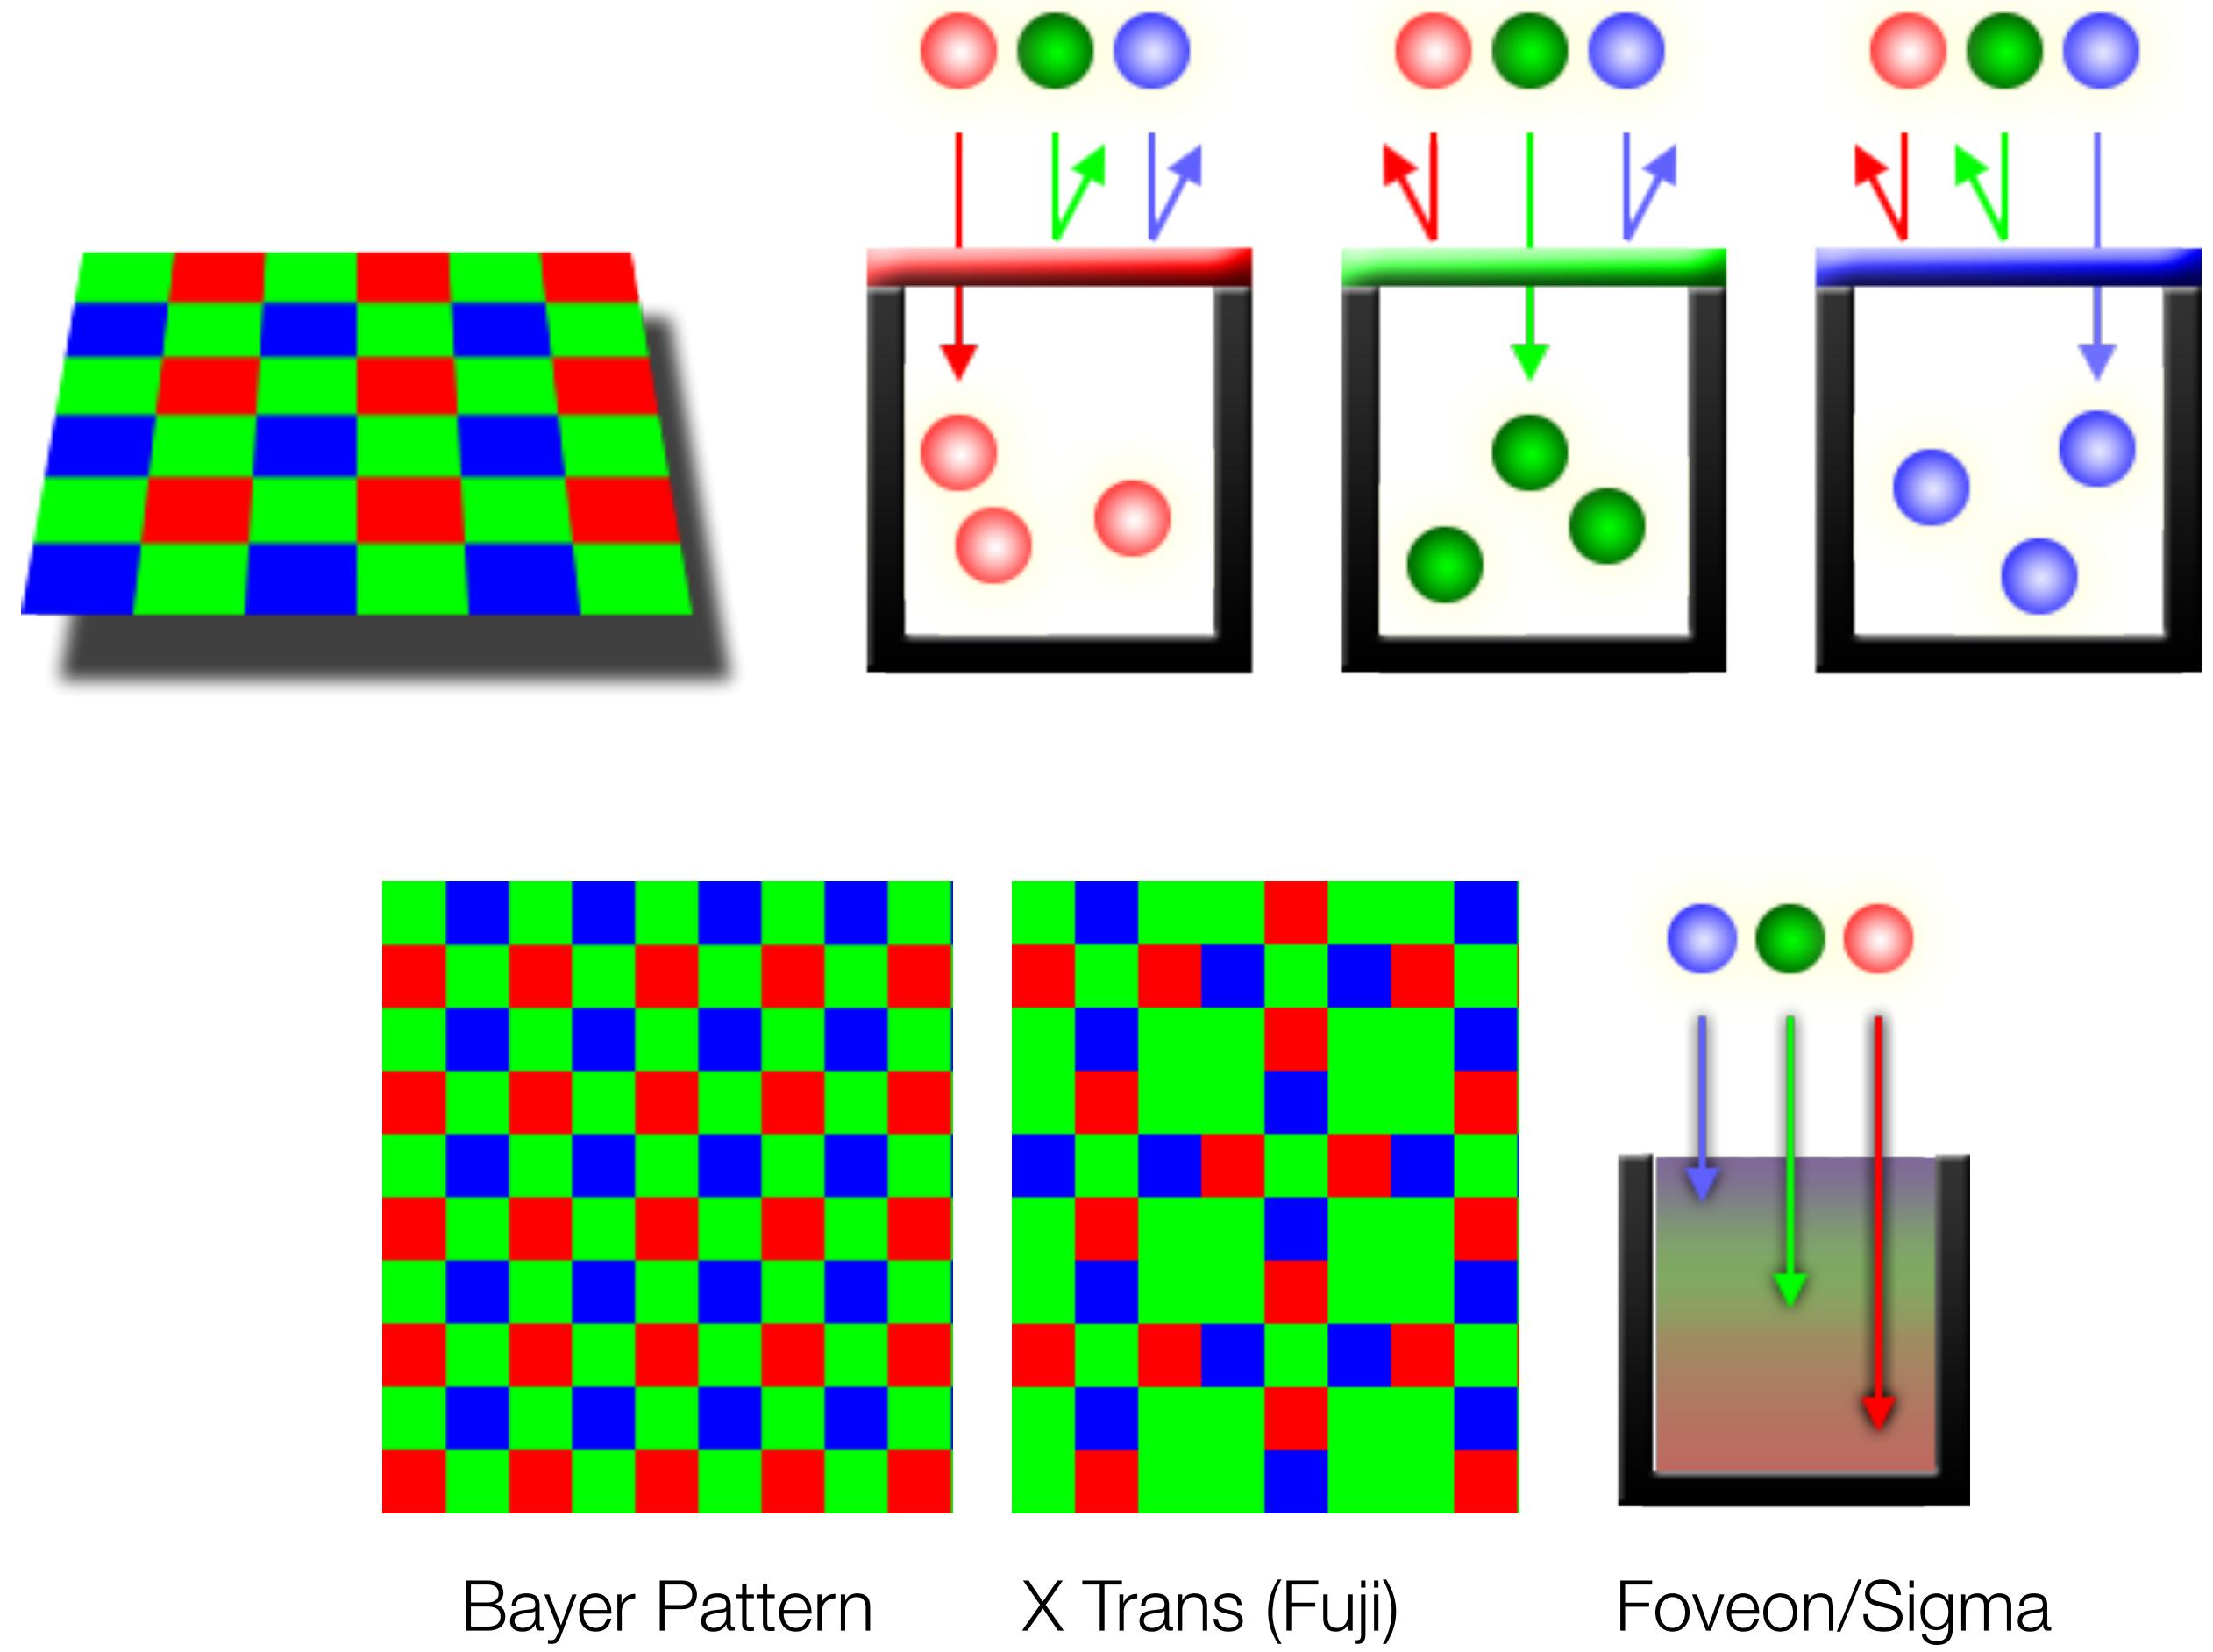
Bayer Image



After Demosaicing

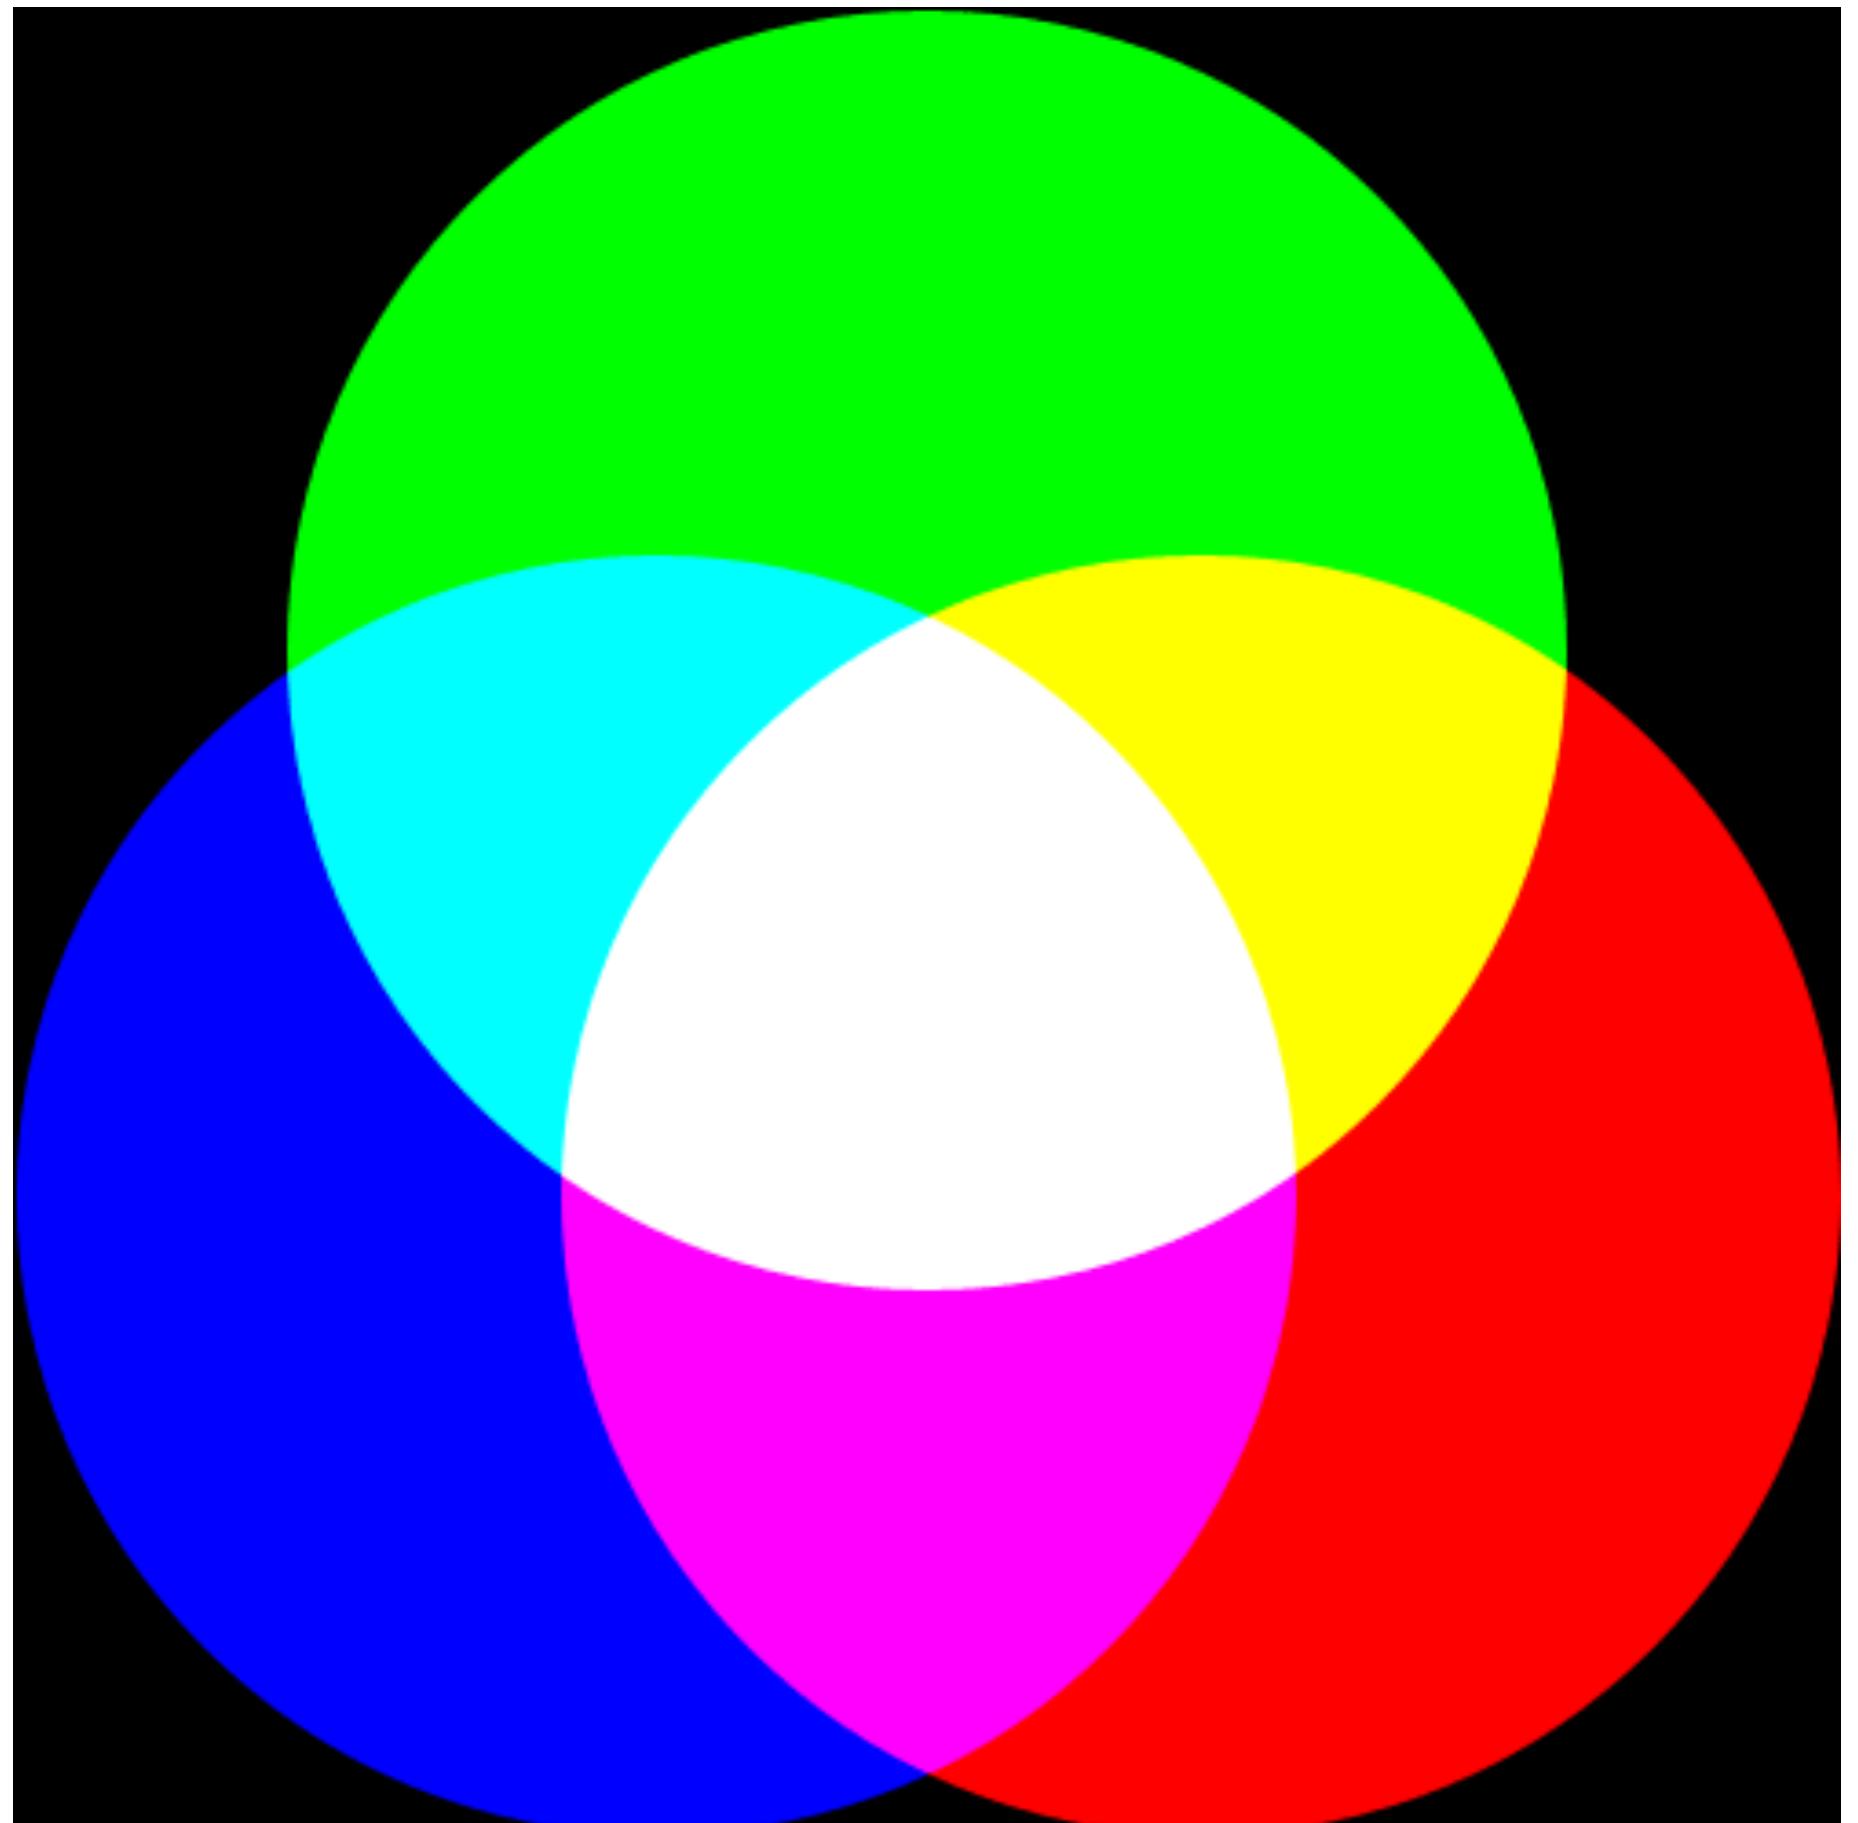
# RGB Model

- Most common in computing
  - Most monitors and projectors
  - Additive colour model
  - Not perceptually uniform
- Record/transmit RGB channels
  - Aligns to the types of cone in the eye
  - Mixes three primary light colours
  - Can represent most colours (not all?)
  - Typically within the range [0,255] or [0,1]
    - E.g. Red as (255,0,0) or (1,0,0)
  - Can be abbreviated as hex values:
    - (183, 24, 92) -> #B7185C



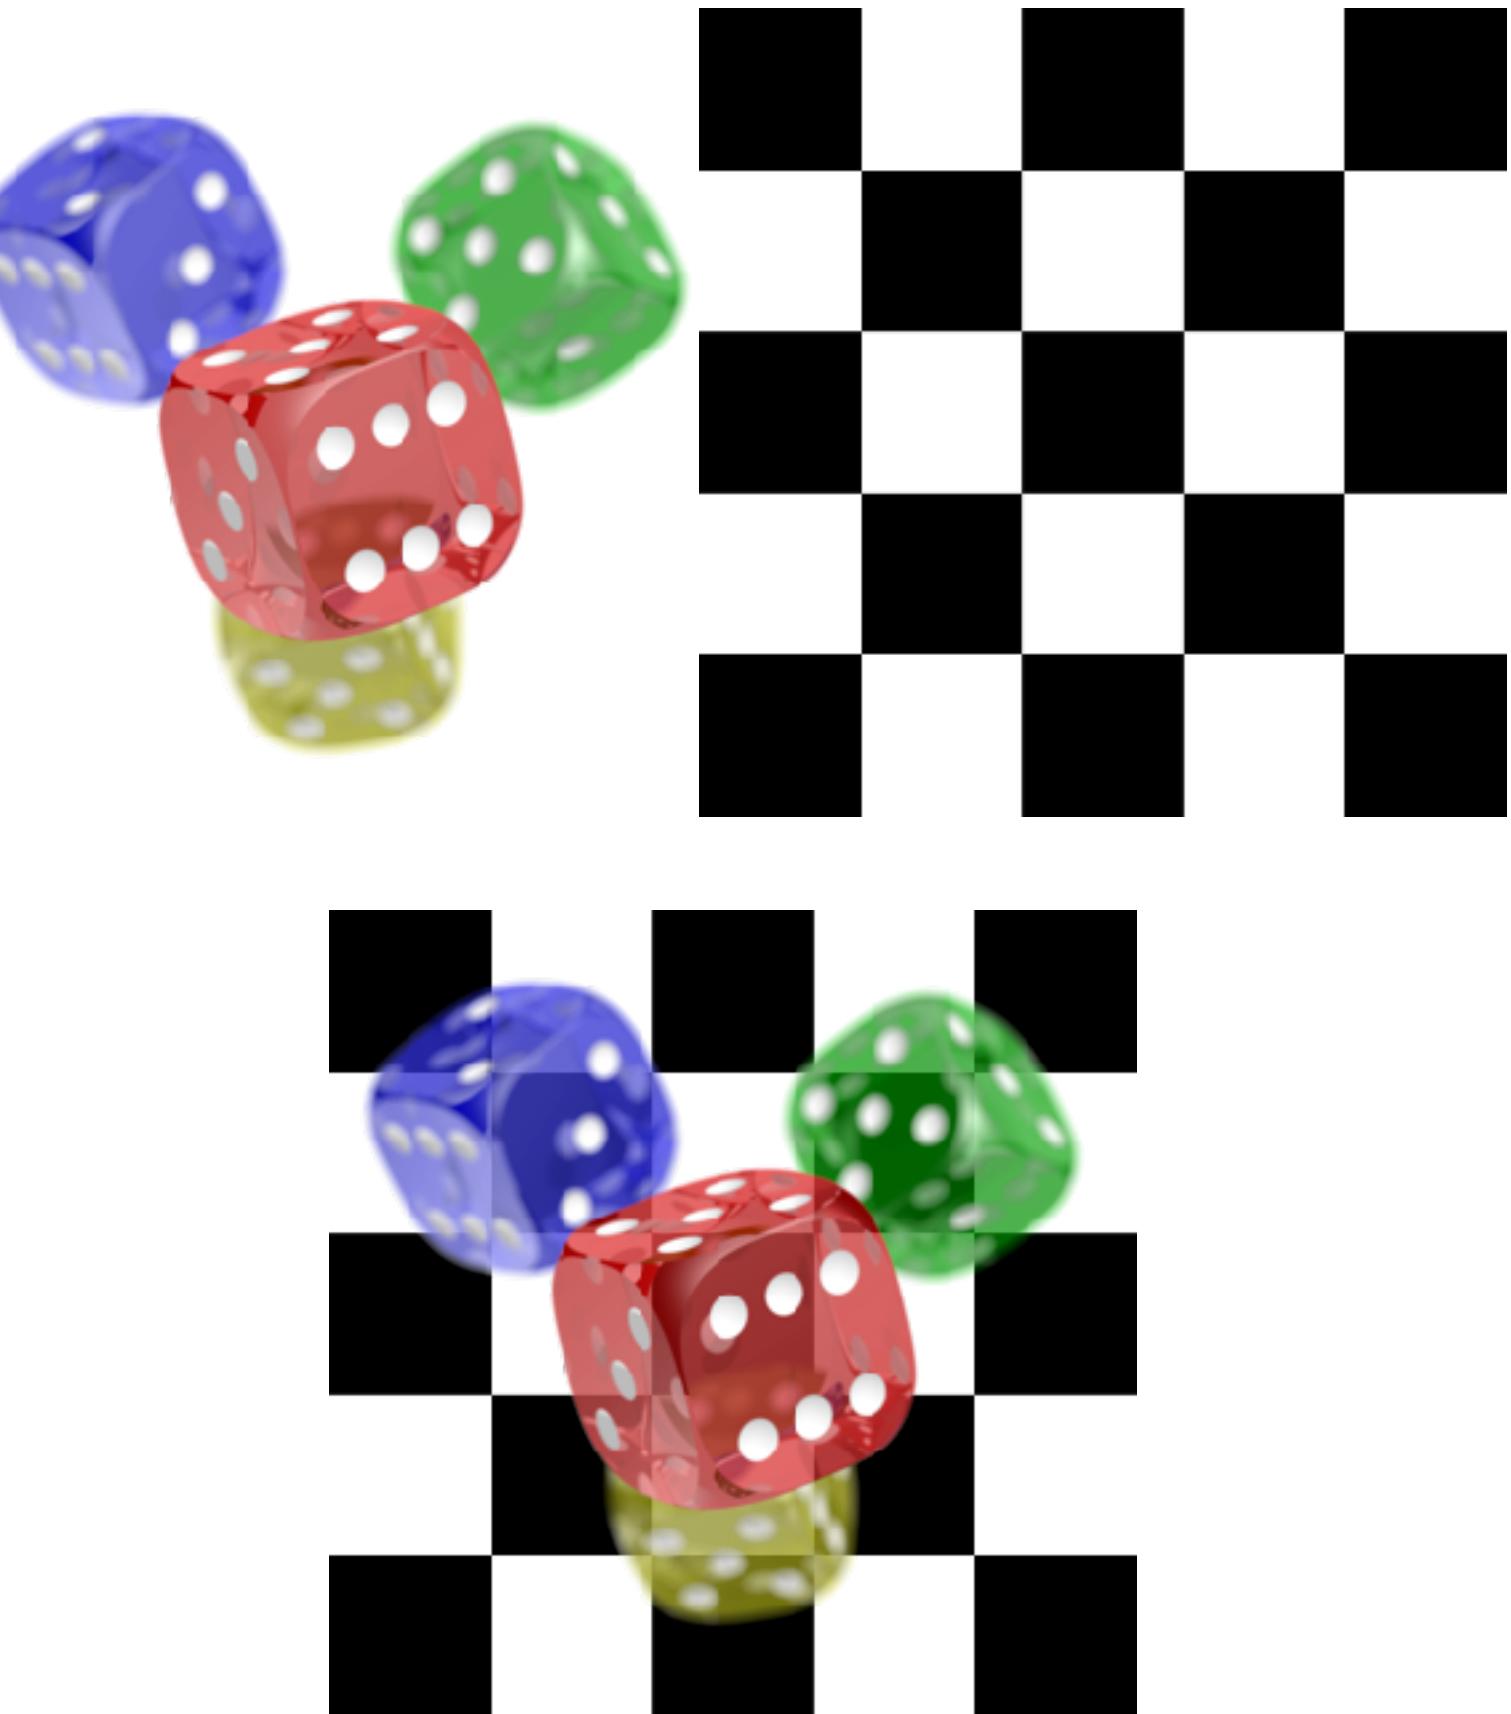
# RGB Model

- Most common in computing
  - Most monitors and projectors
  - Additive colour model
  - Not perceptually uniform
- Record/transmit RGB channels
  - Aligns to the types of cone in the eye
  - Mixes three primary light colours
  - Can represent most colours (not all?)
  - Typically within the range [0,255] or [0,1]
    - E.g. Red as (255,0,0) or (1,0,0)
  - Can be abbreviated as hex values:
    - (183, 24, 92) -> #B7185C



# RGBA - Transparency

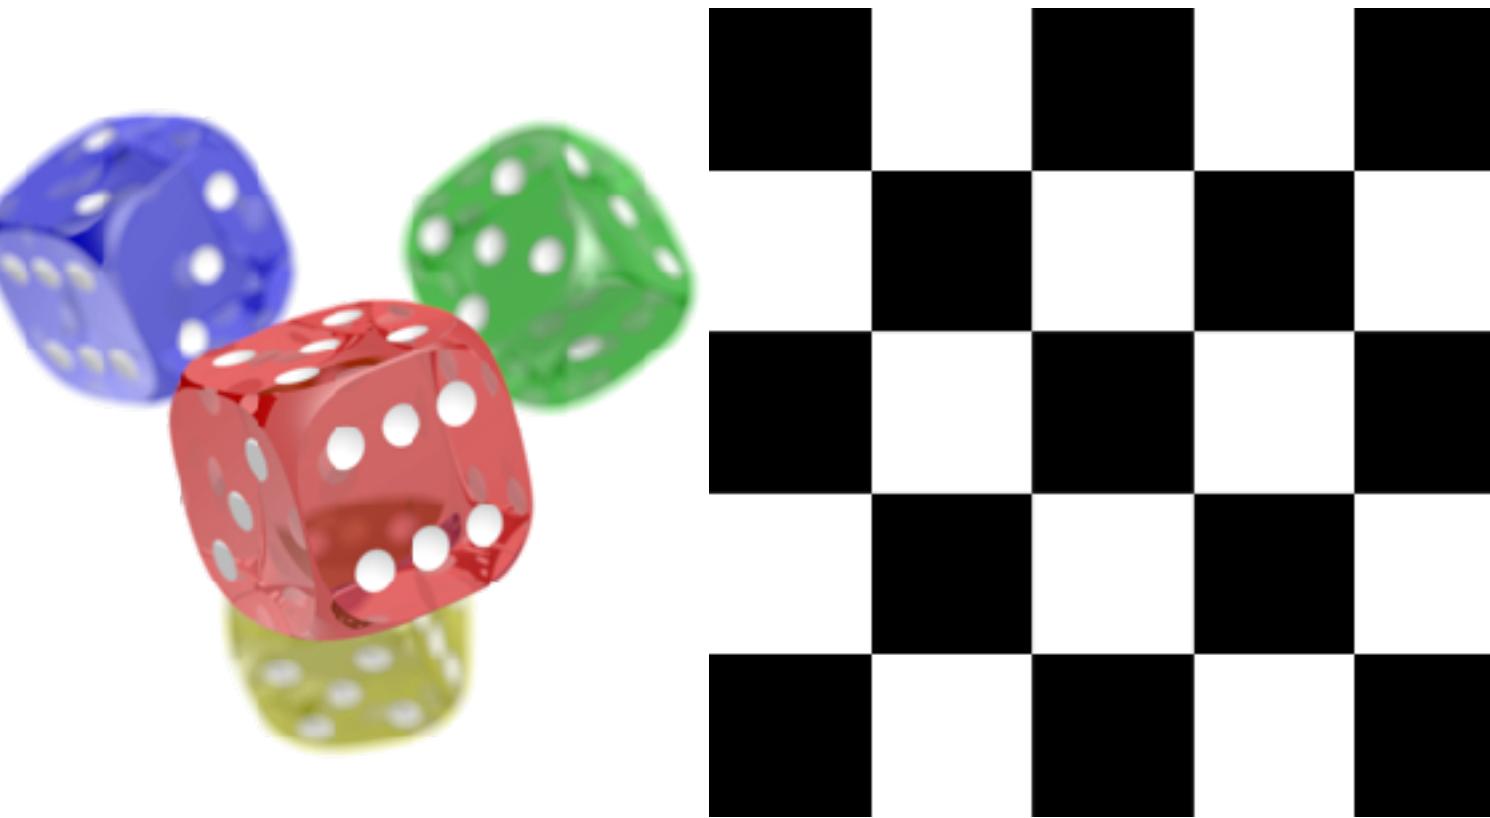
- Adds a fourth ‘colour’ channel
- Typically called the ‘alpha channel’  $\alpha$
- Combining images is alpha-blending
- transparent  $\alpha = 0$ , opaque  $\alpha = 1$
- ARGB or RGBA common byte orders



# RGBA - Transparency

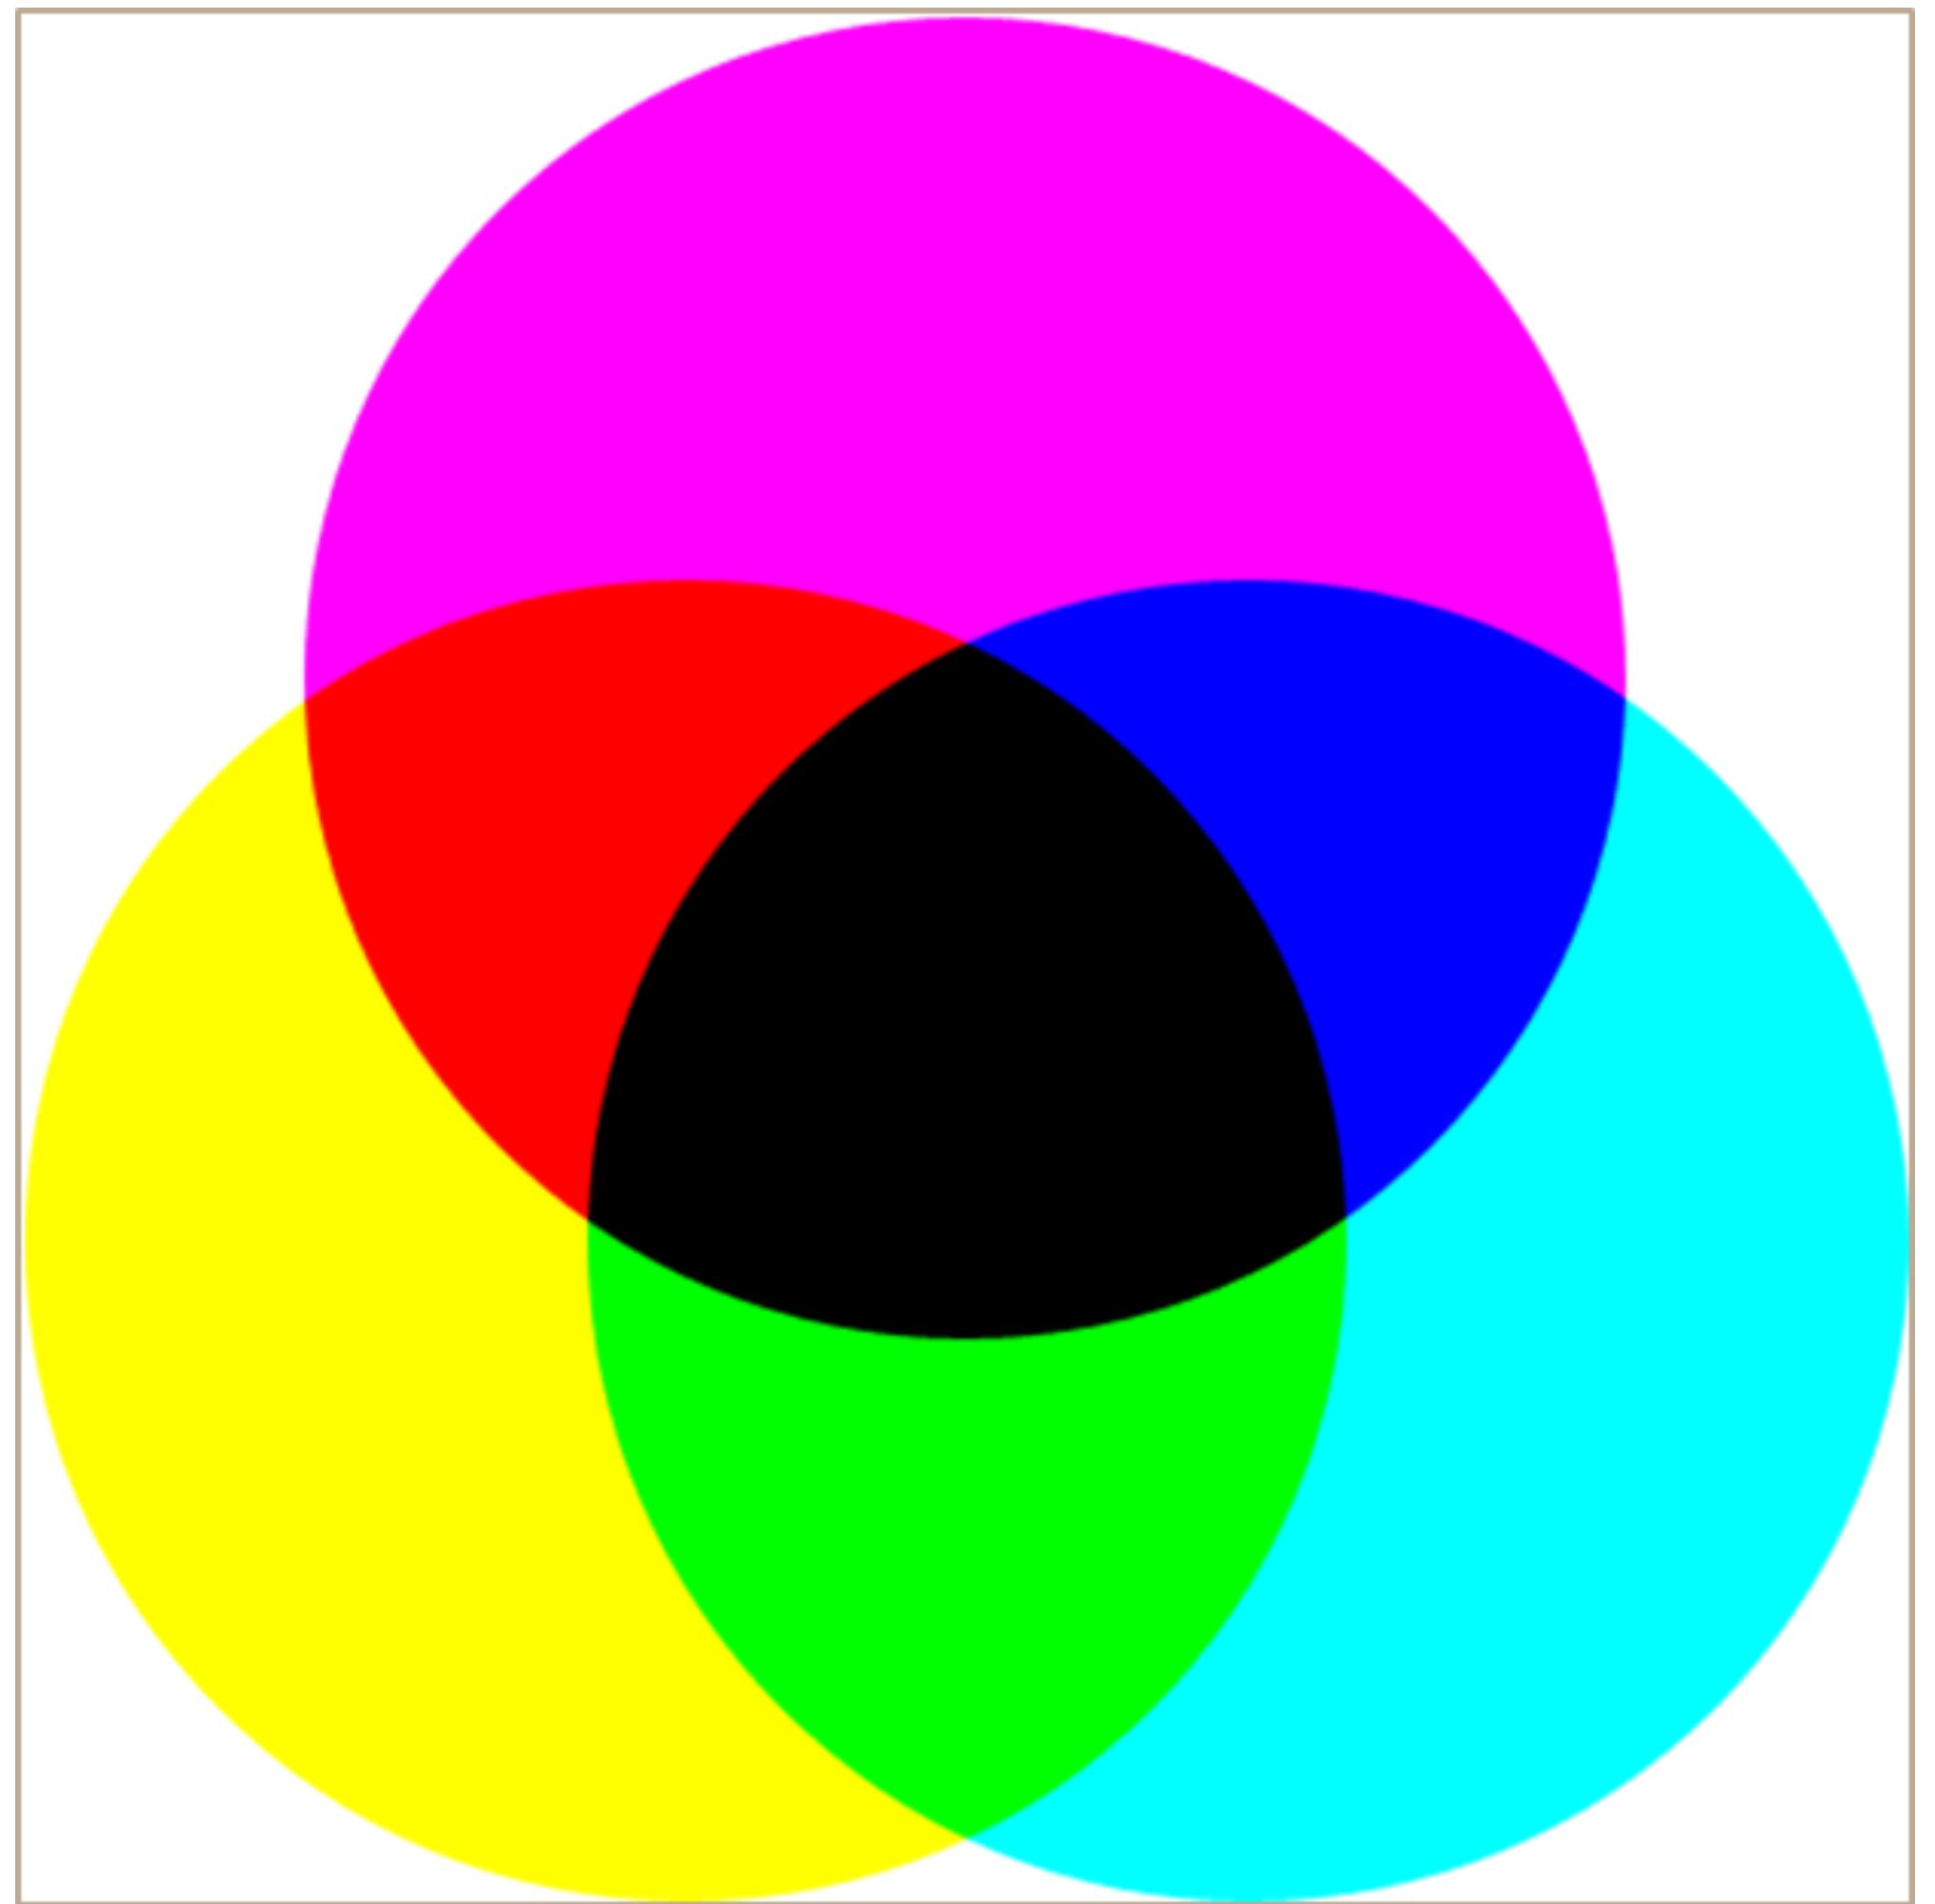
- How do we combine colours?
  - Foreground/front:  $(\alpha_f, r_f, g_f, b_f)$
  - Background/back:  $(\alpha_b, r_b, g_b, b_b)$
  - Output colour is:

$$r_o = \left( \alpha_f r_f + (1 - \alpha_f) \alpha_b r_b \right) / \alpha_o$$
$$g_o = \left( \alpha_f g_f + (1 - \alpha_f) \alpha_b g_b \right) / \alpha_o$$
$$b_o = \left( \alpha_f b_f + (1 - \alpha_f) \alpha_b b_b \right) / \alpha_o$$
$$\alpha_o = \alpha_f + (1 - \alpha_f) \alpha_b$$



# CMY Colour Space

- Most common for printing
  - Subtractive Colour Model
  - Not perceptually uniform
- CMY is the inverse of RGB
  - Cyan Magenta Yellow (Black)
  - Cyan pigment absorbs red light
  - Magenta pigment absorbs green
  - Yellow pigment absorbs blue
- In theory  $C+M+Y = \text{Black}$ 
  - In practice a ‘true Black’ ink is used
  - Black is represented by K (Key)



# CMY Colour Space

- Most common for printing
  - Subtractive Colour Model
  - Not perceptually uniform
- CMY is the inverse of RGB
  - Cyan Magenta Yellow (Black)
  - Cyan pigment absorbs red light
  - Magenta pigment absorbs green
  - Yellow pigment absorbs blue
- In theory  $C+M+Y = \text{Black}$ 
  - In practice a 'true Black' ink is used
  - Black is represented by K (Key)



CYAN

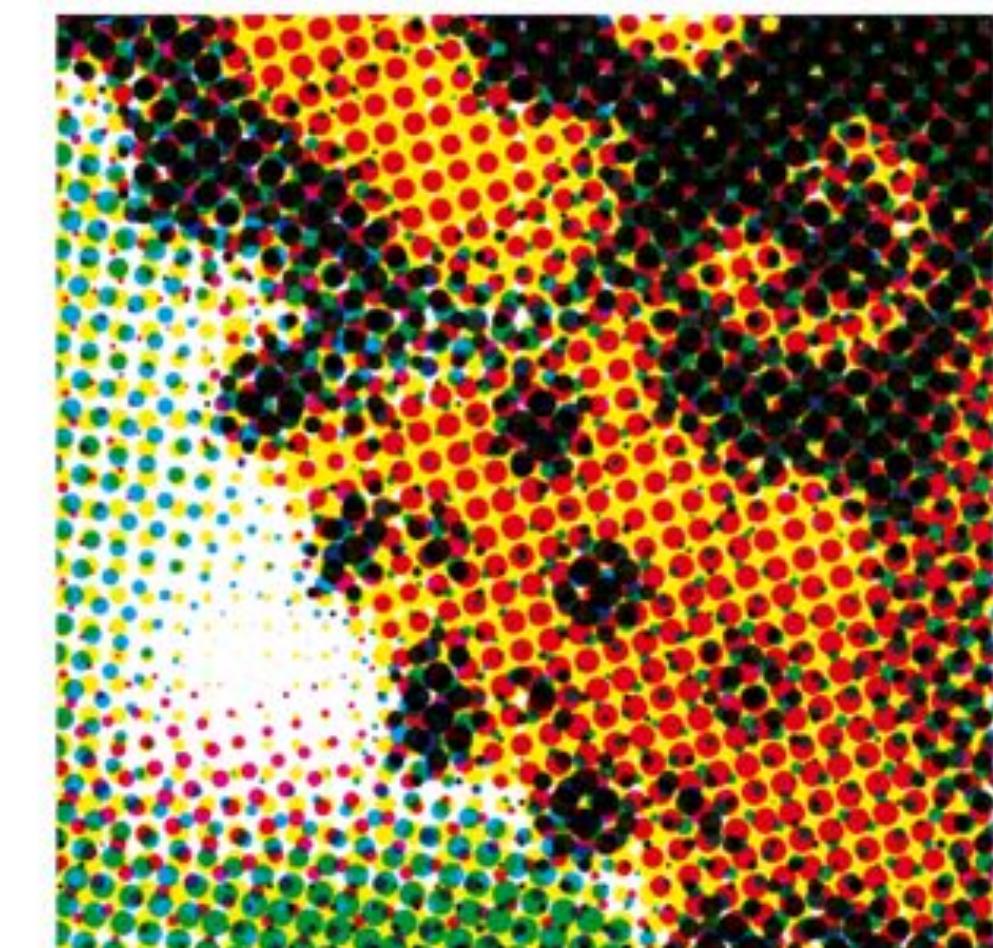
MAGENTA

YELLOW

BLACK



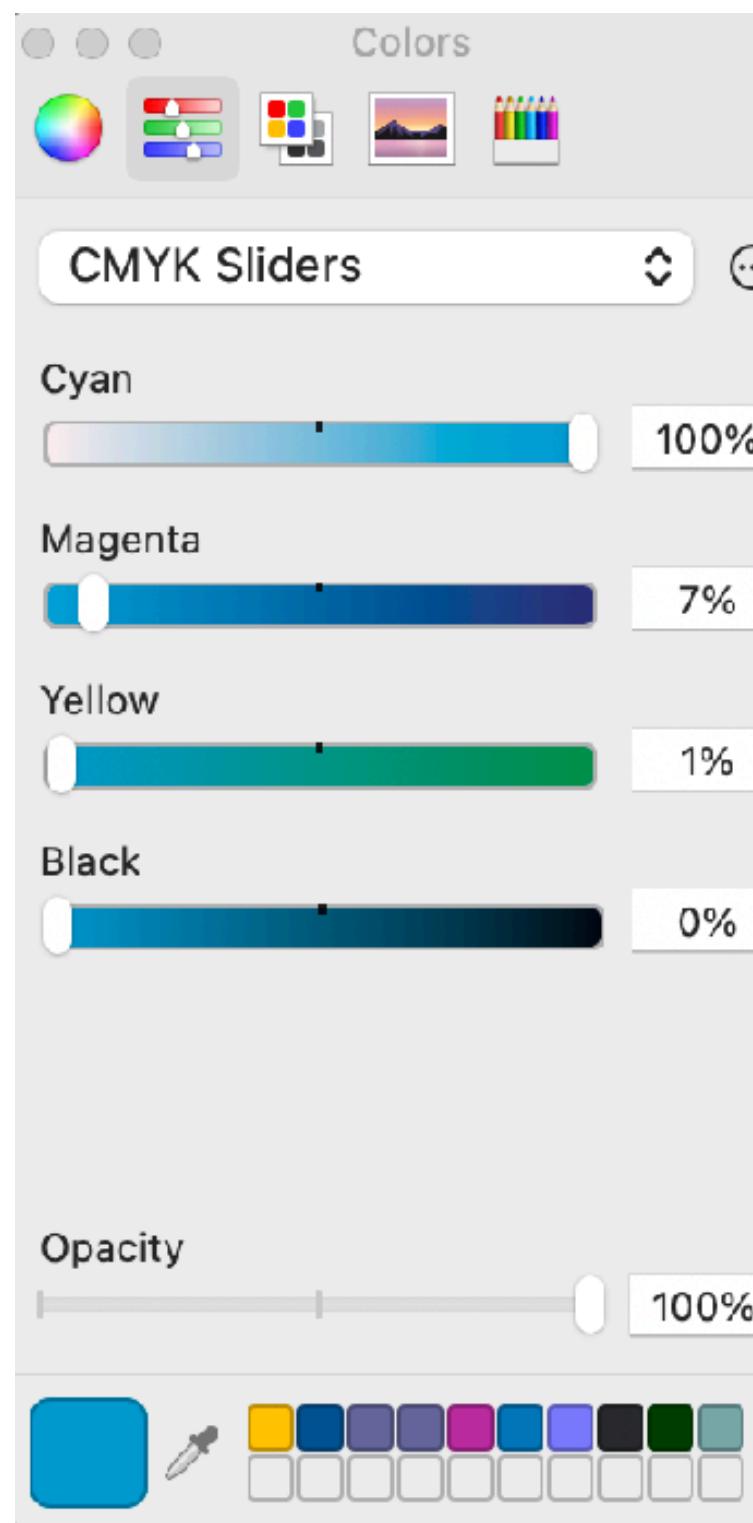
FINAL CMYK



DETAIL VIEW

# CMY Colour Space

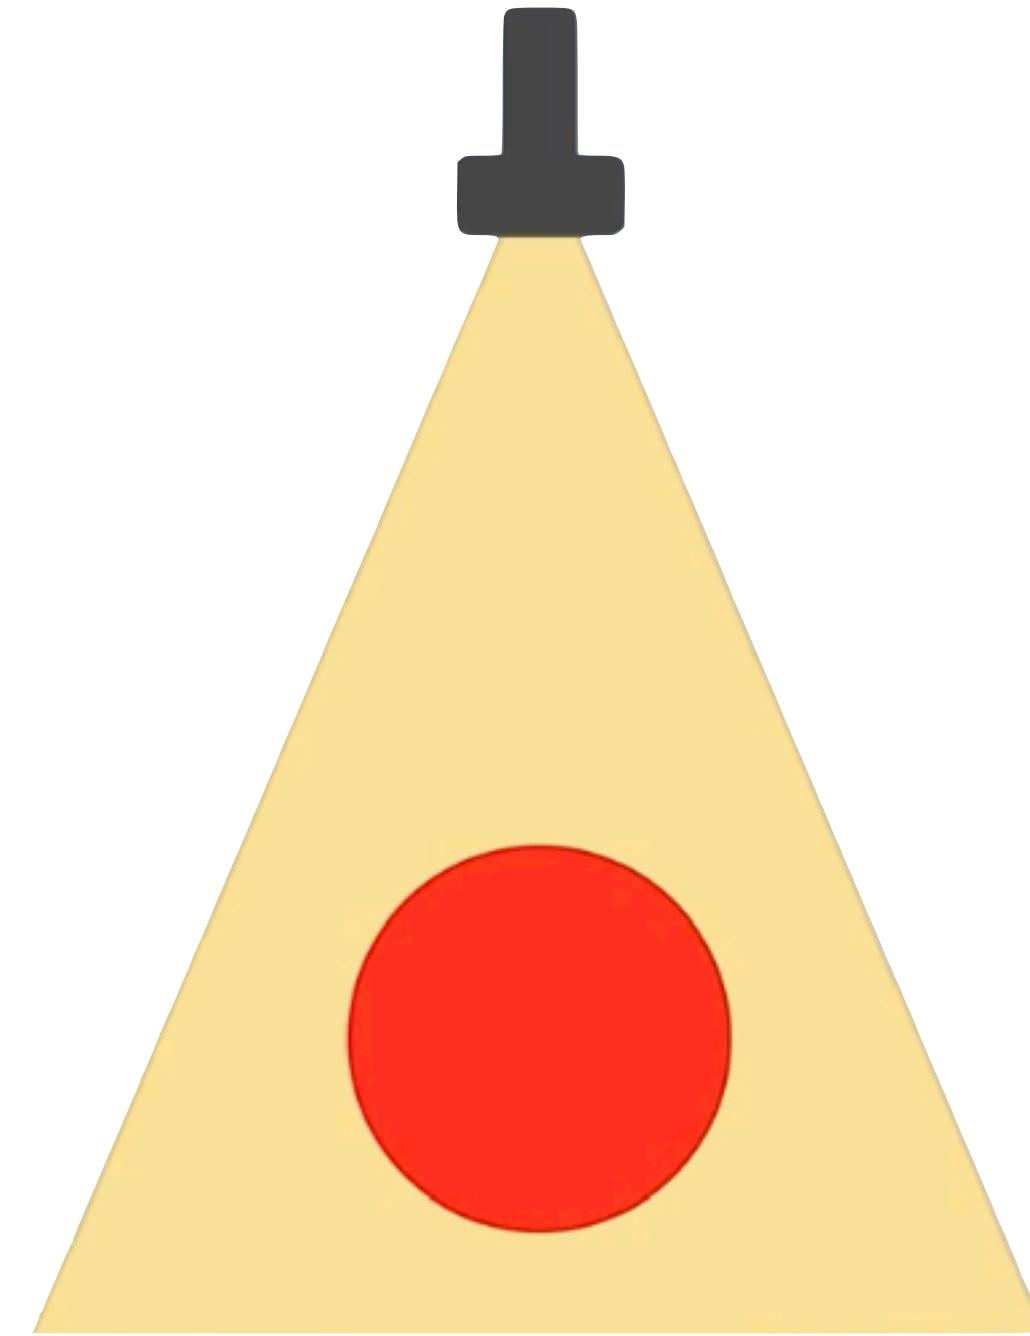
- Most common for printing
  - Subtractive Colour Model
  - Not perceptually uniform
- CMY is the inverse of RGB
  - Cyan Magenta Yellow (Black)
  - Cyan pigment absorbs red light
  - Magenta pigment absorbs green
  - Yellow pigment absorbs blue
- In theory  $C+M+Y = \text{Black}$ 
  - In practice a ‘true Black’ ink is used
  - Black is represented by K (Key)



Color	Color name	(R,G,B)	Hex	(C,M,Y,K)
Black	Black	(0,0,0)	#000000	(0,0,0,1)
White	White	(255,255,255)	#FFFFFF	(0,0,0,0)
Red	Red	(255,0,0)	#FF0000	(0,1,1,0)
Green	Green	(0,255,0)	#00FF00	(1,0,1,0)
Blue	Blue	(0,0,255)	#0000FF	(1,1,0,0)
Yellow	Yellow	(255,255,0)	#FFFF00	(0,0,1,0)
Cyan	Cyan	(0,255,255)	#00FFFF	(1,0,0,0)
Magenta	Magenta	(255,0,255)	#FF00FF	(0,1,0,0)

# Coloured Light on Coloured Object

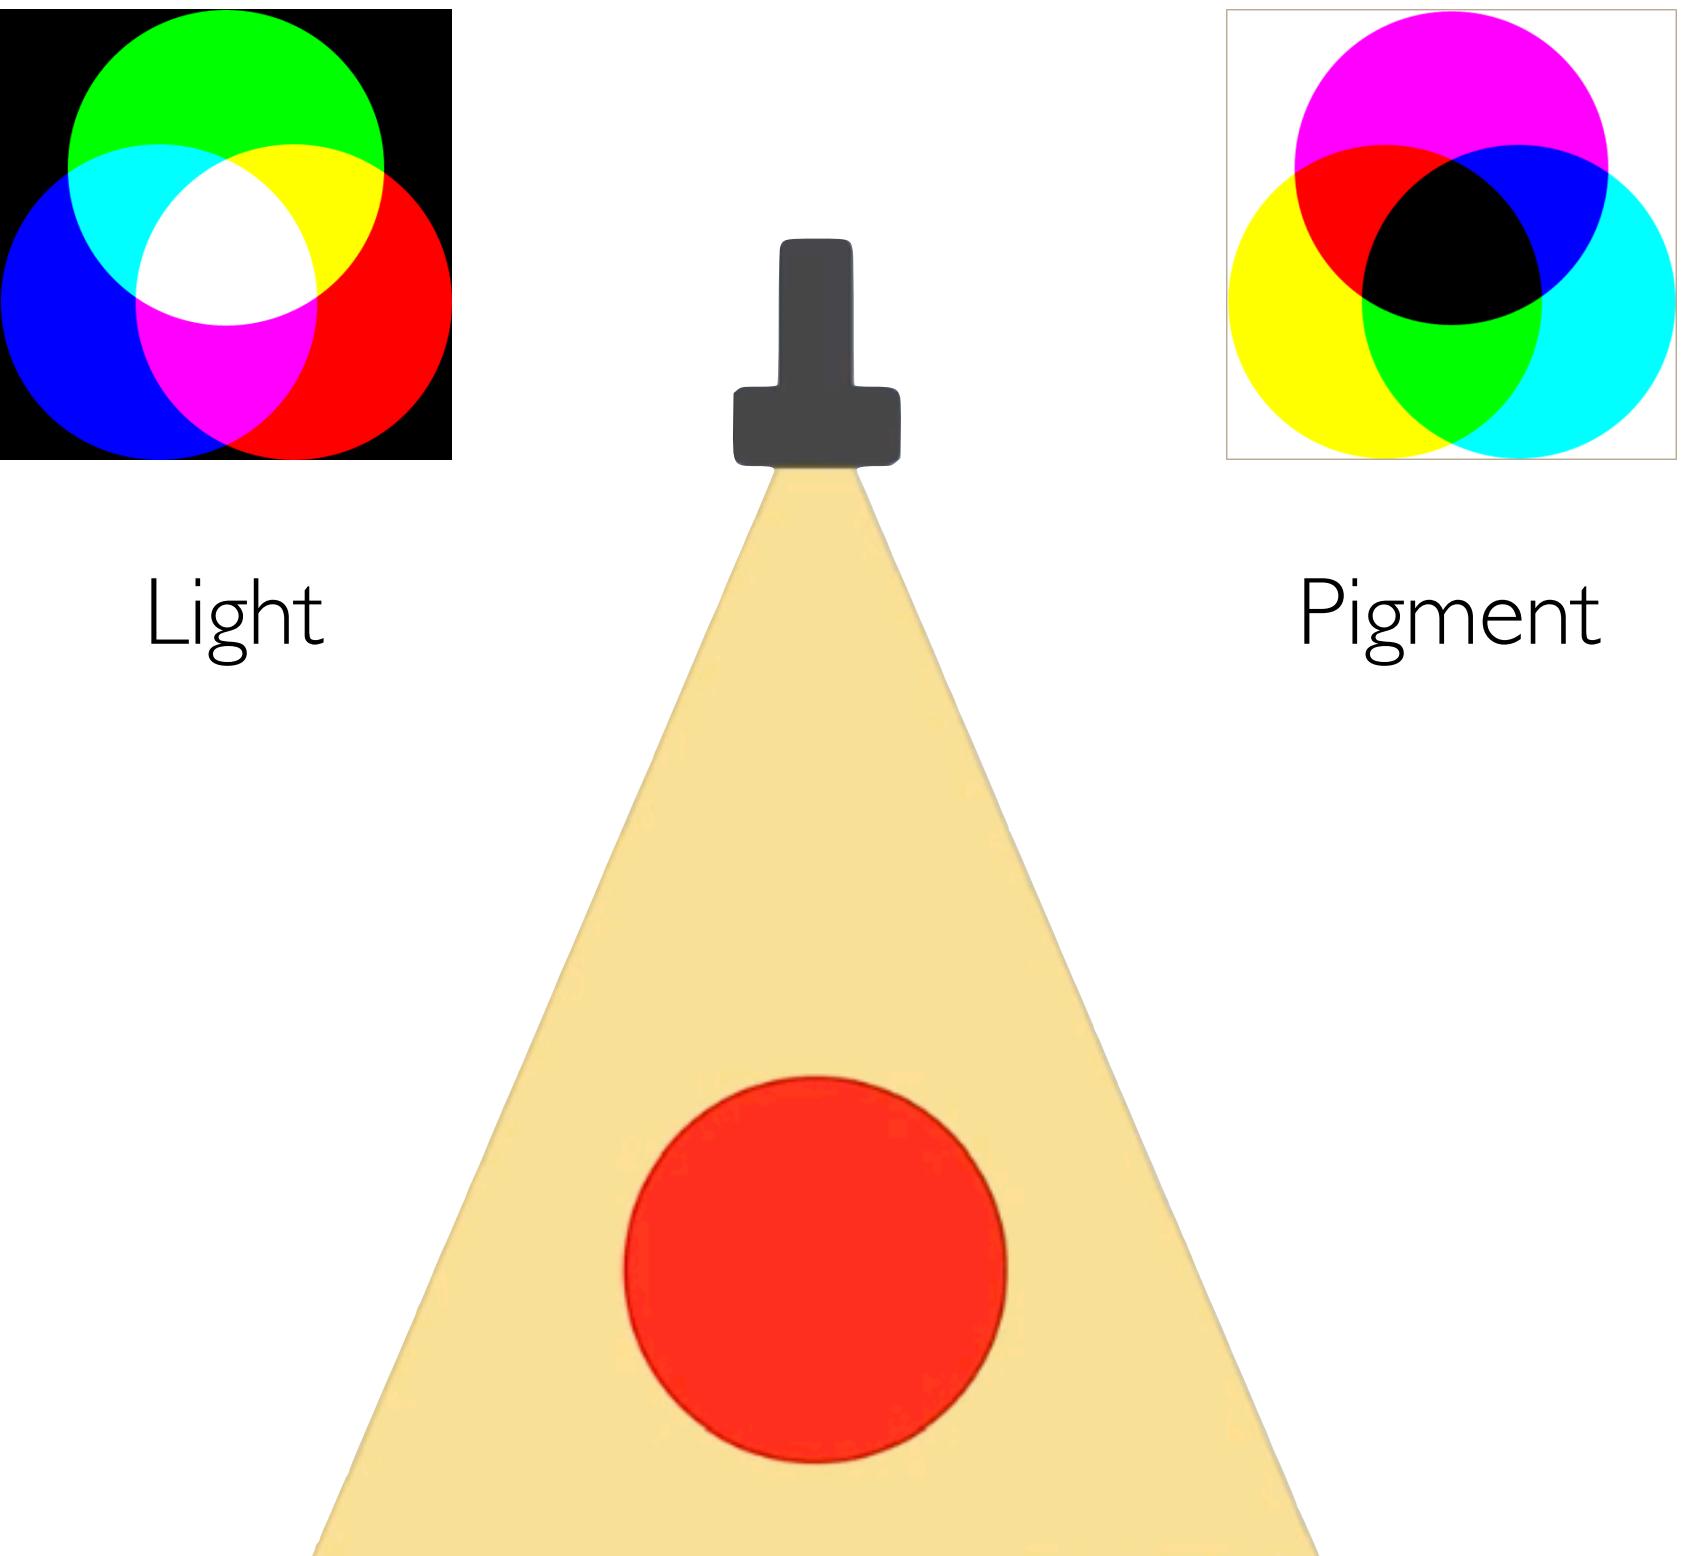
- What do I see if I shine a yellow light on a red ball?



# Coloured Light on Coloured Object

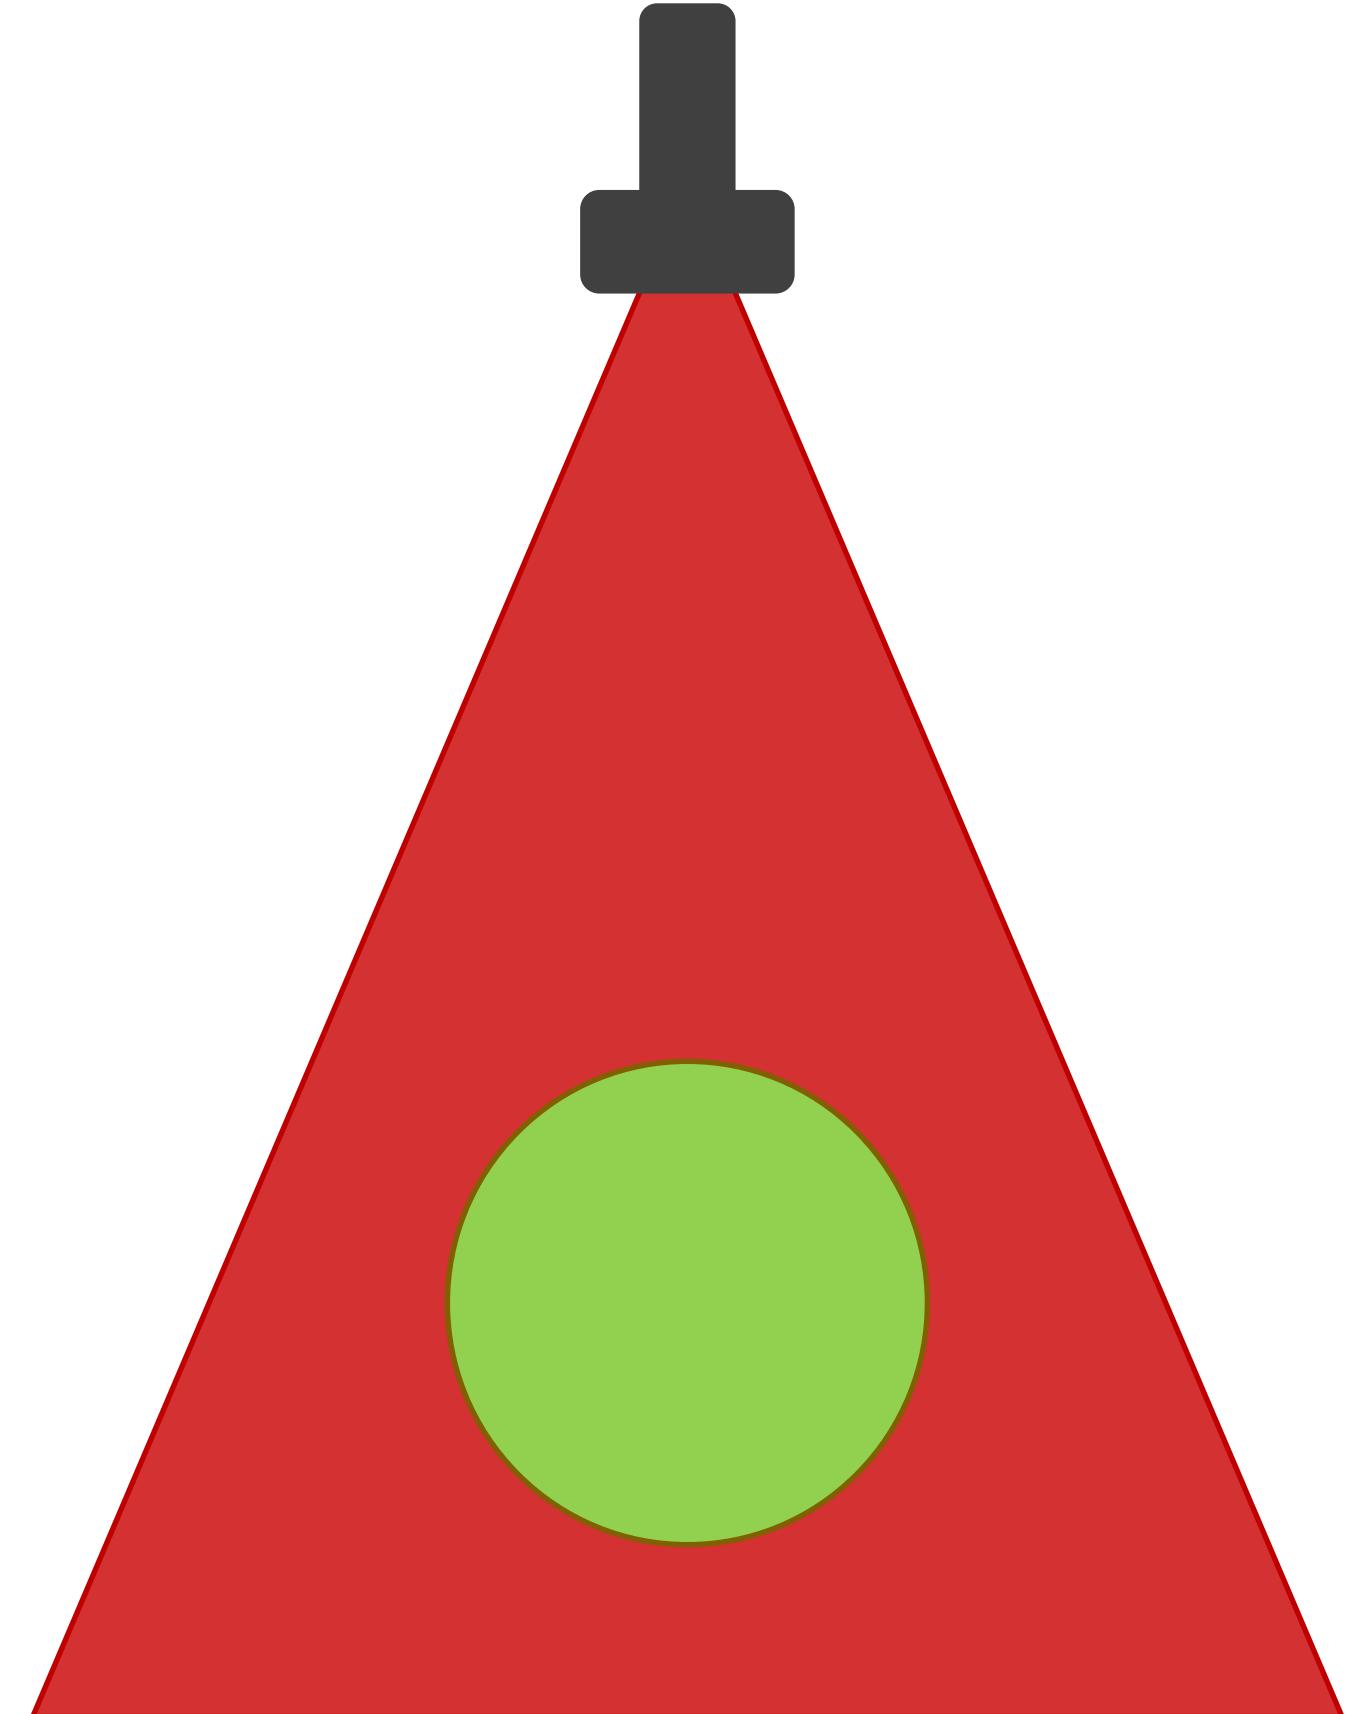
- What do I see if I shine a yellow light on a red ball?

- Yellow light = Red + Green
- Red ball = Magenta + Yellow pigment
- There is no Blue light ->  $B = 0$
- Green light is absorbed by Magenta
- Red light is not absorbed -> reflected
- So we see Red



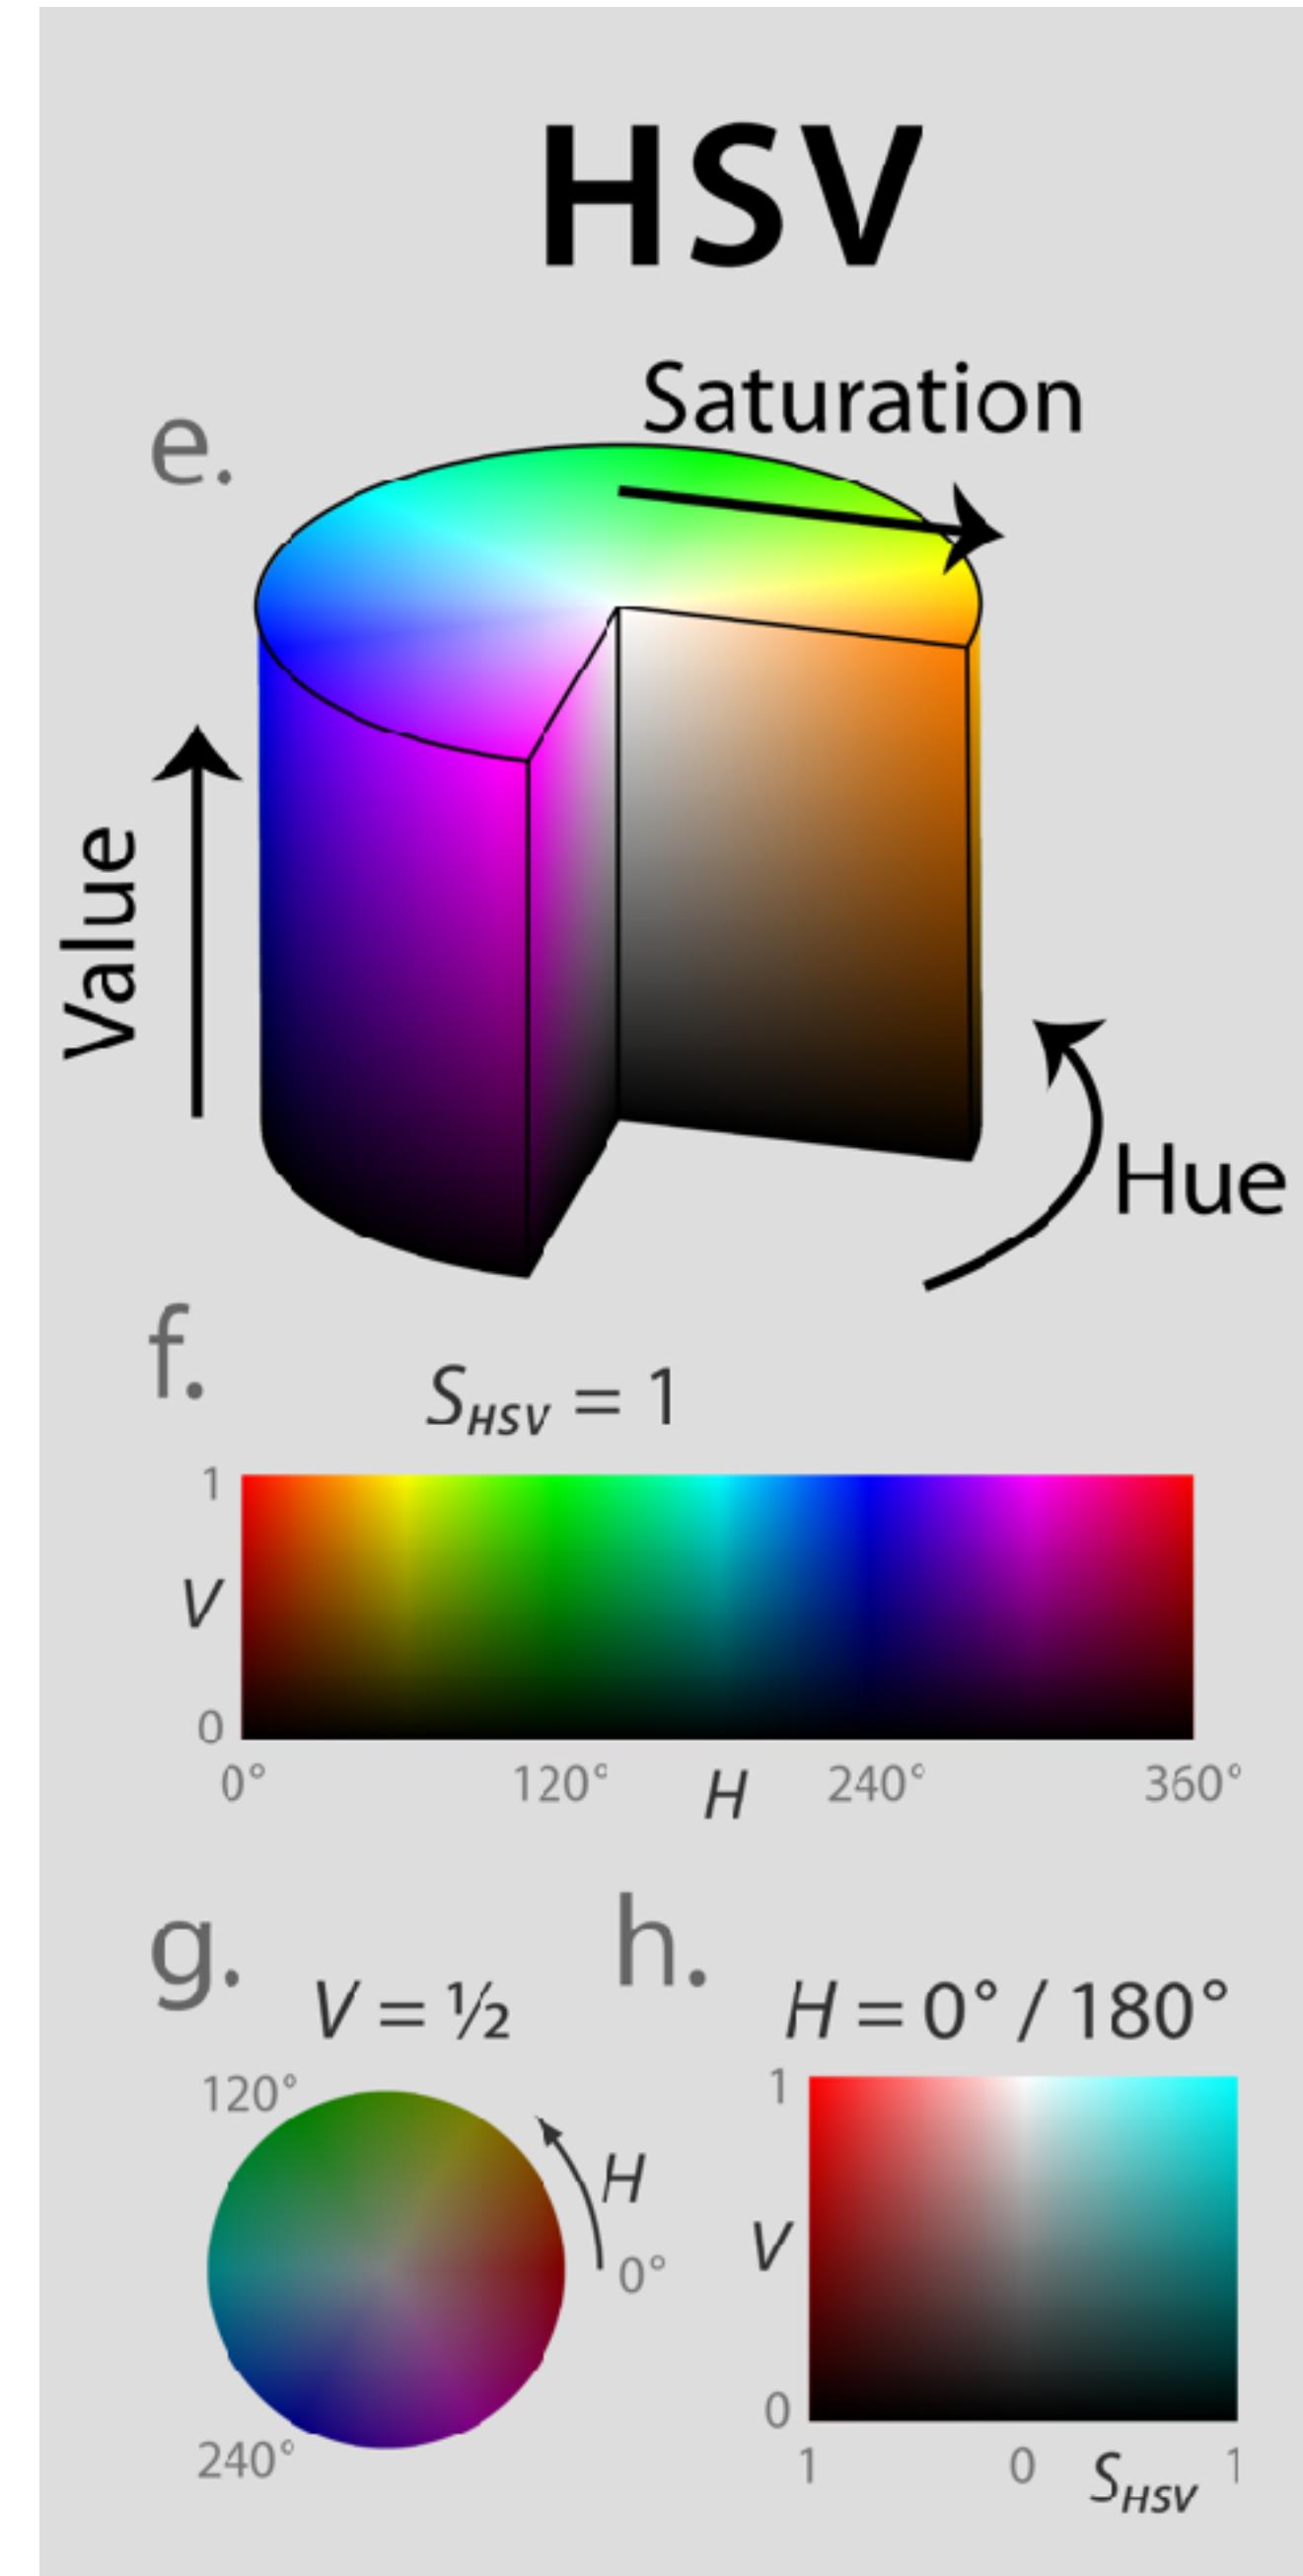
# Coloured Light on Coloured Object

- Think about
  - What colour would you see for:
  - Shining a red light on a green object?
  - A cyan light on a yellow object?
  - Does the order matter when alpha blending colours?



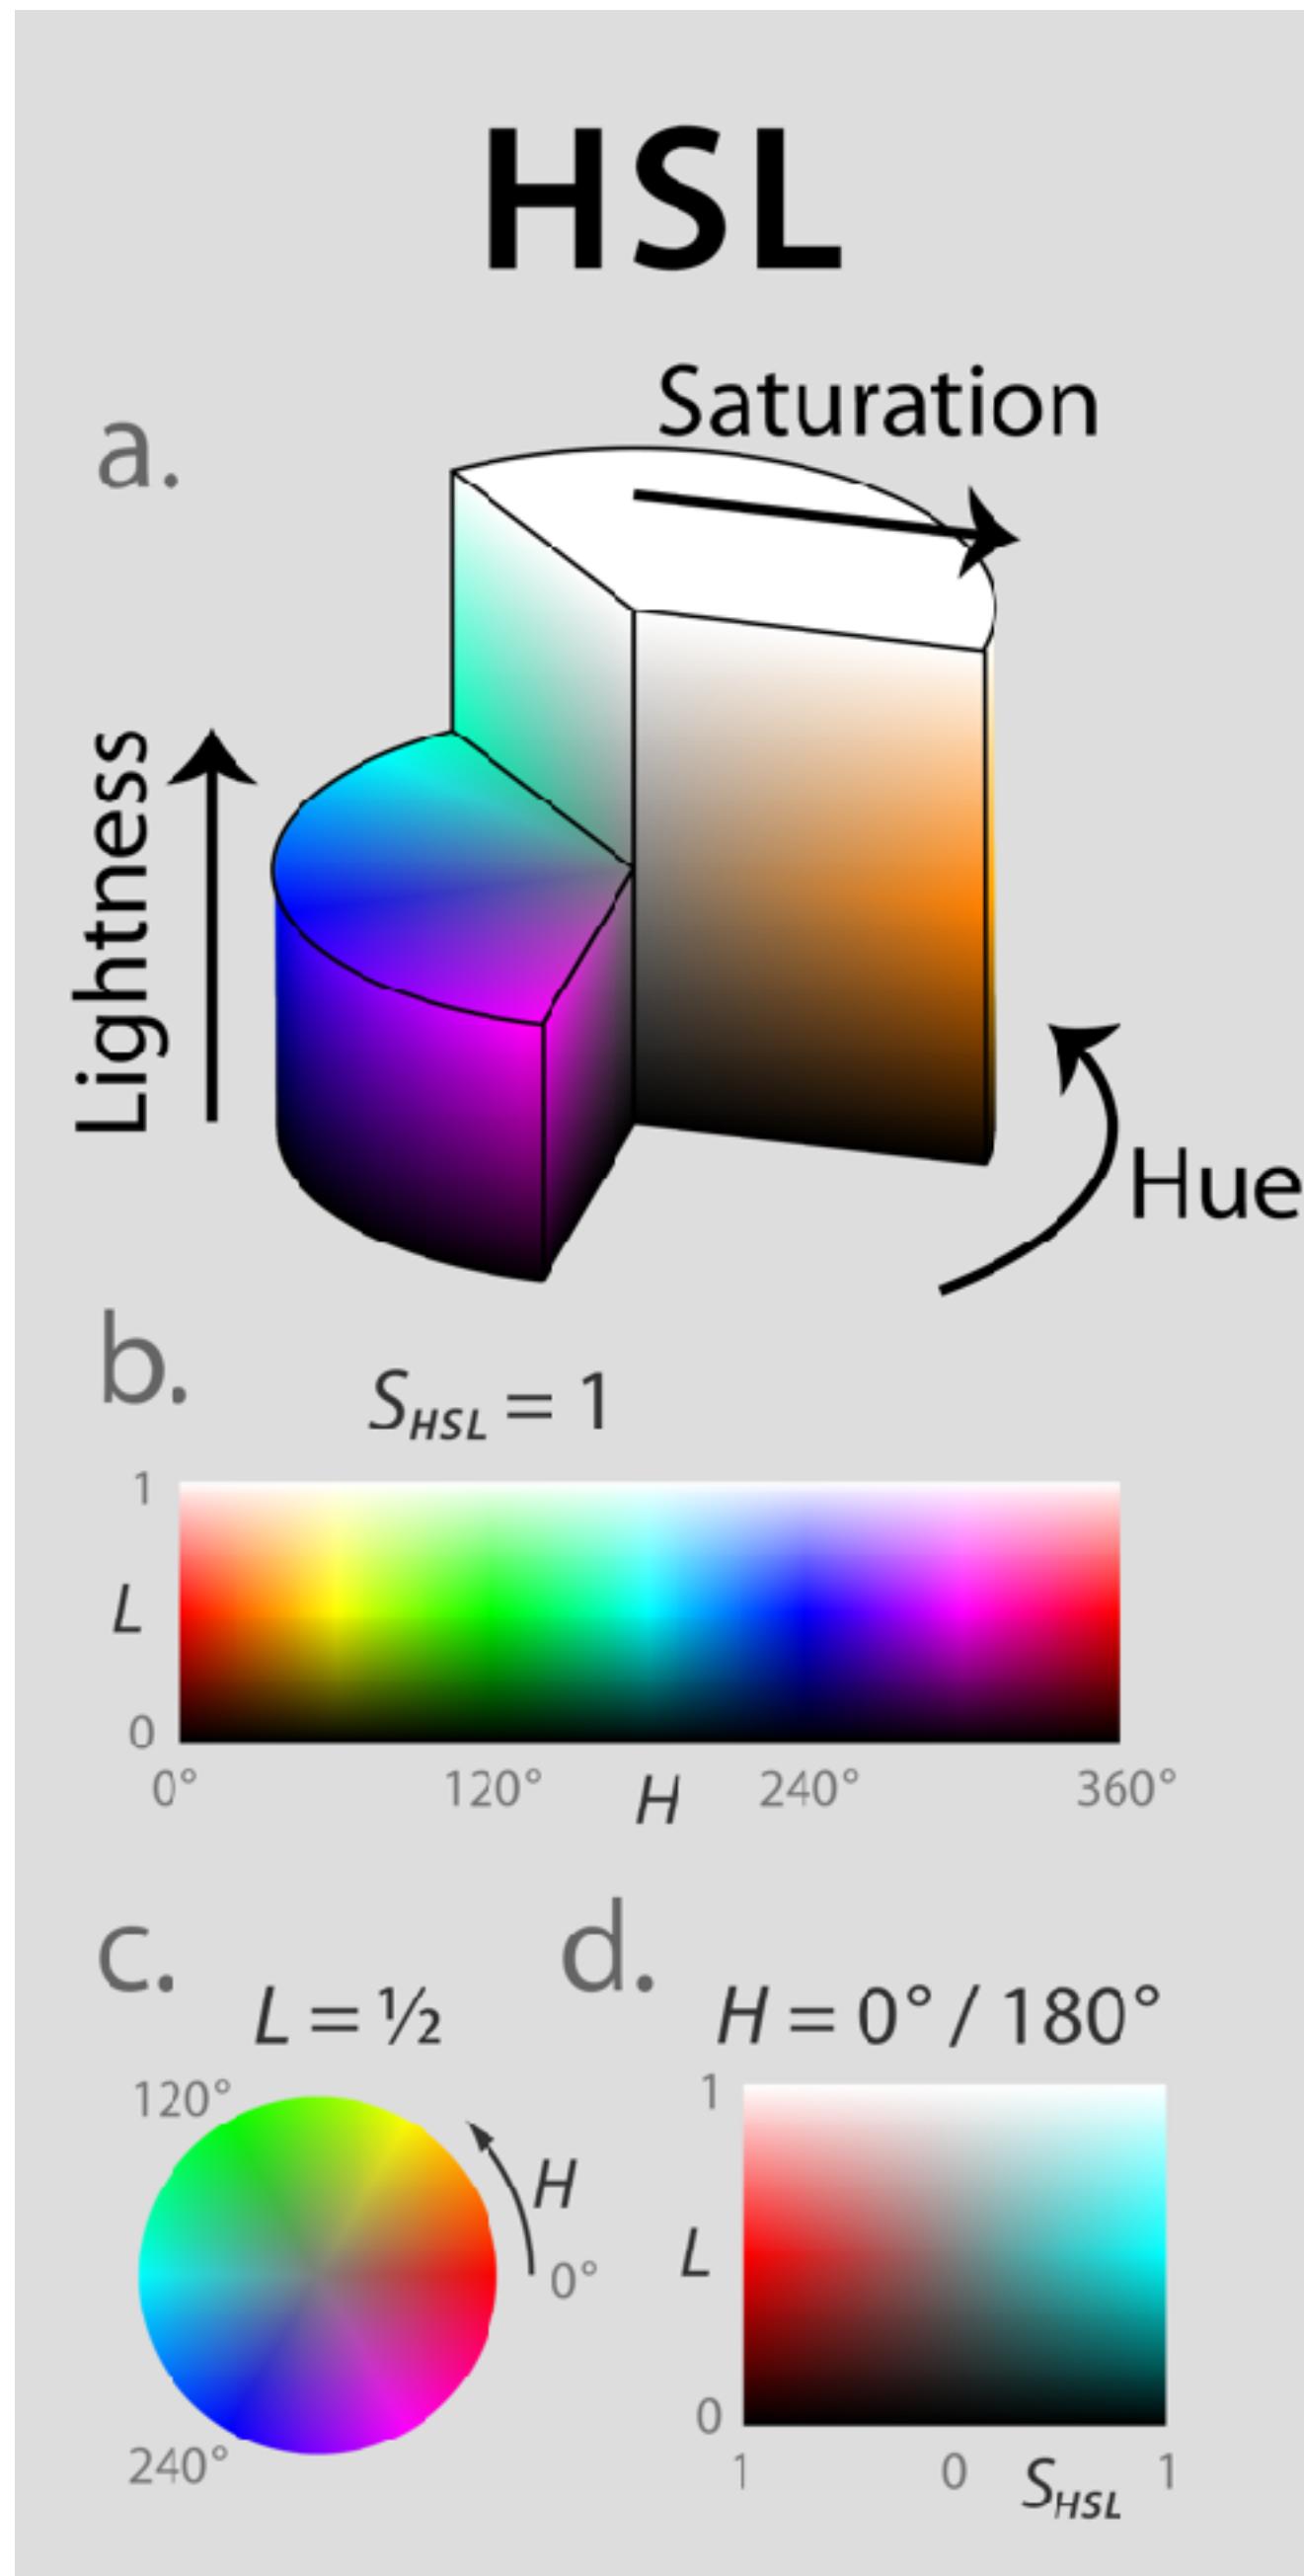
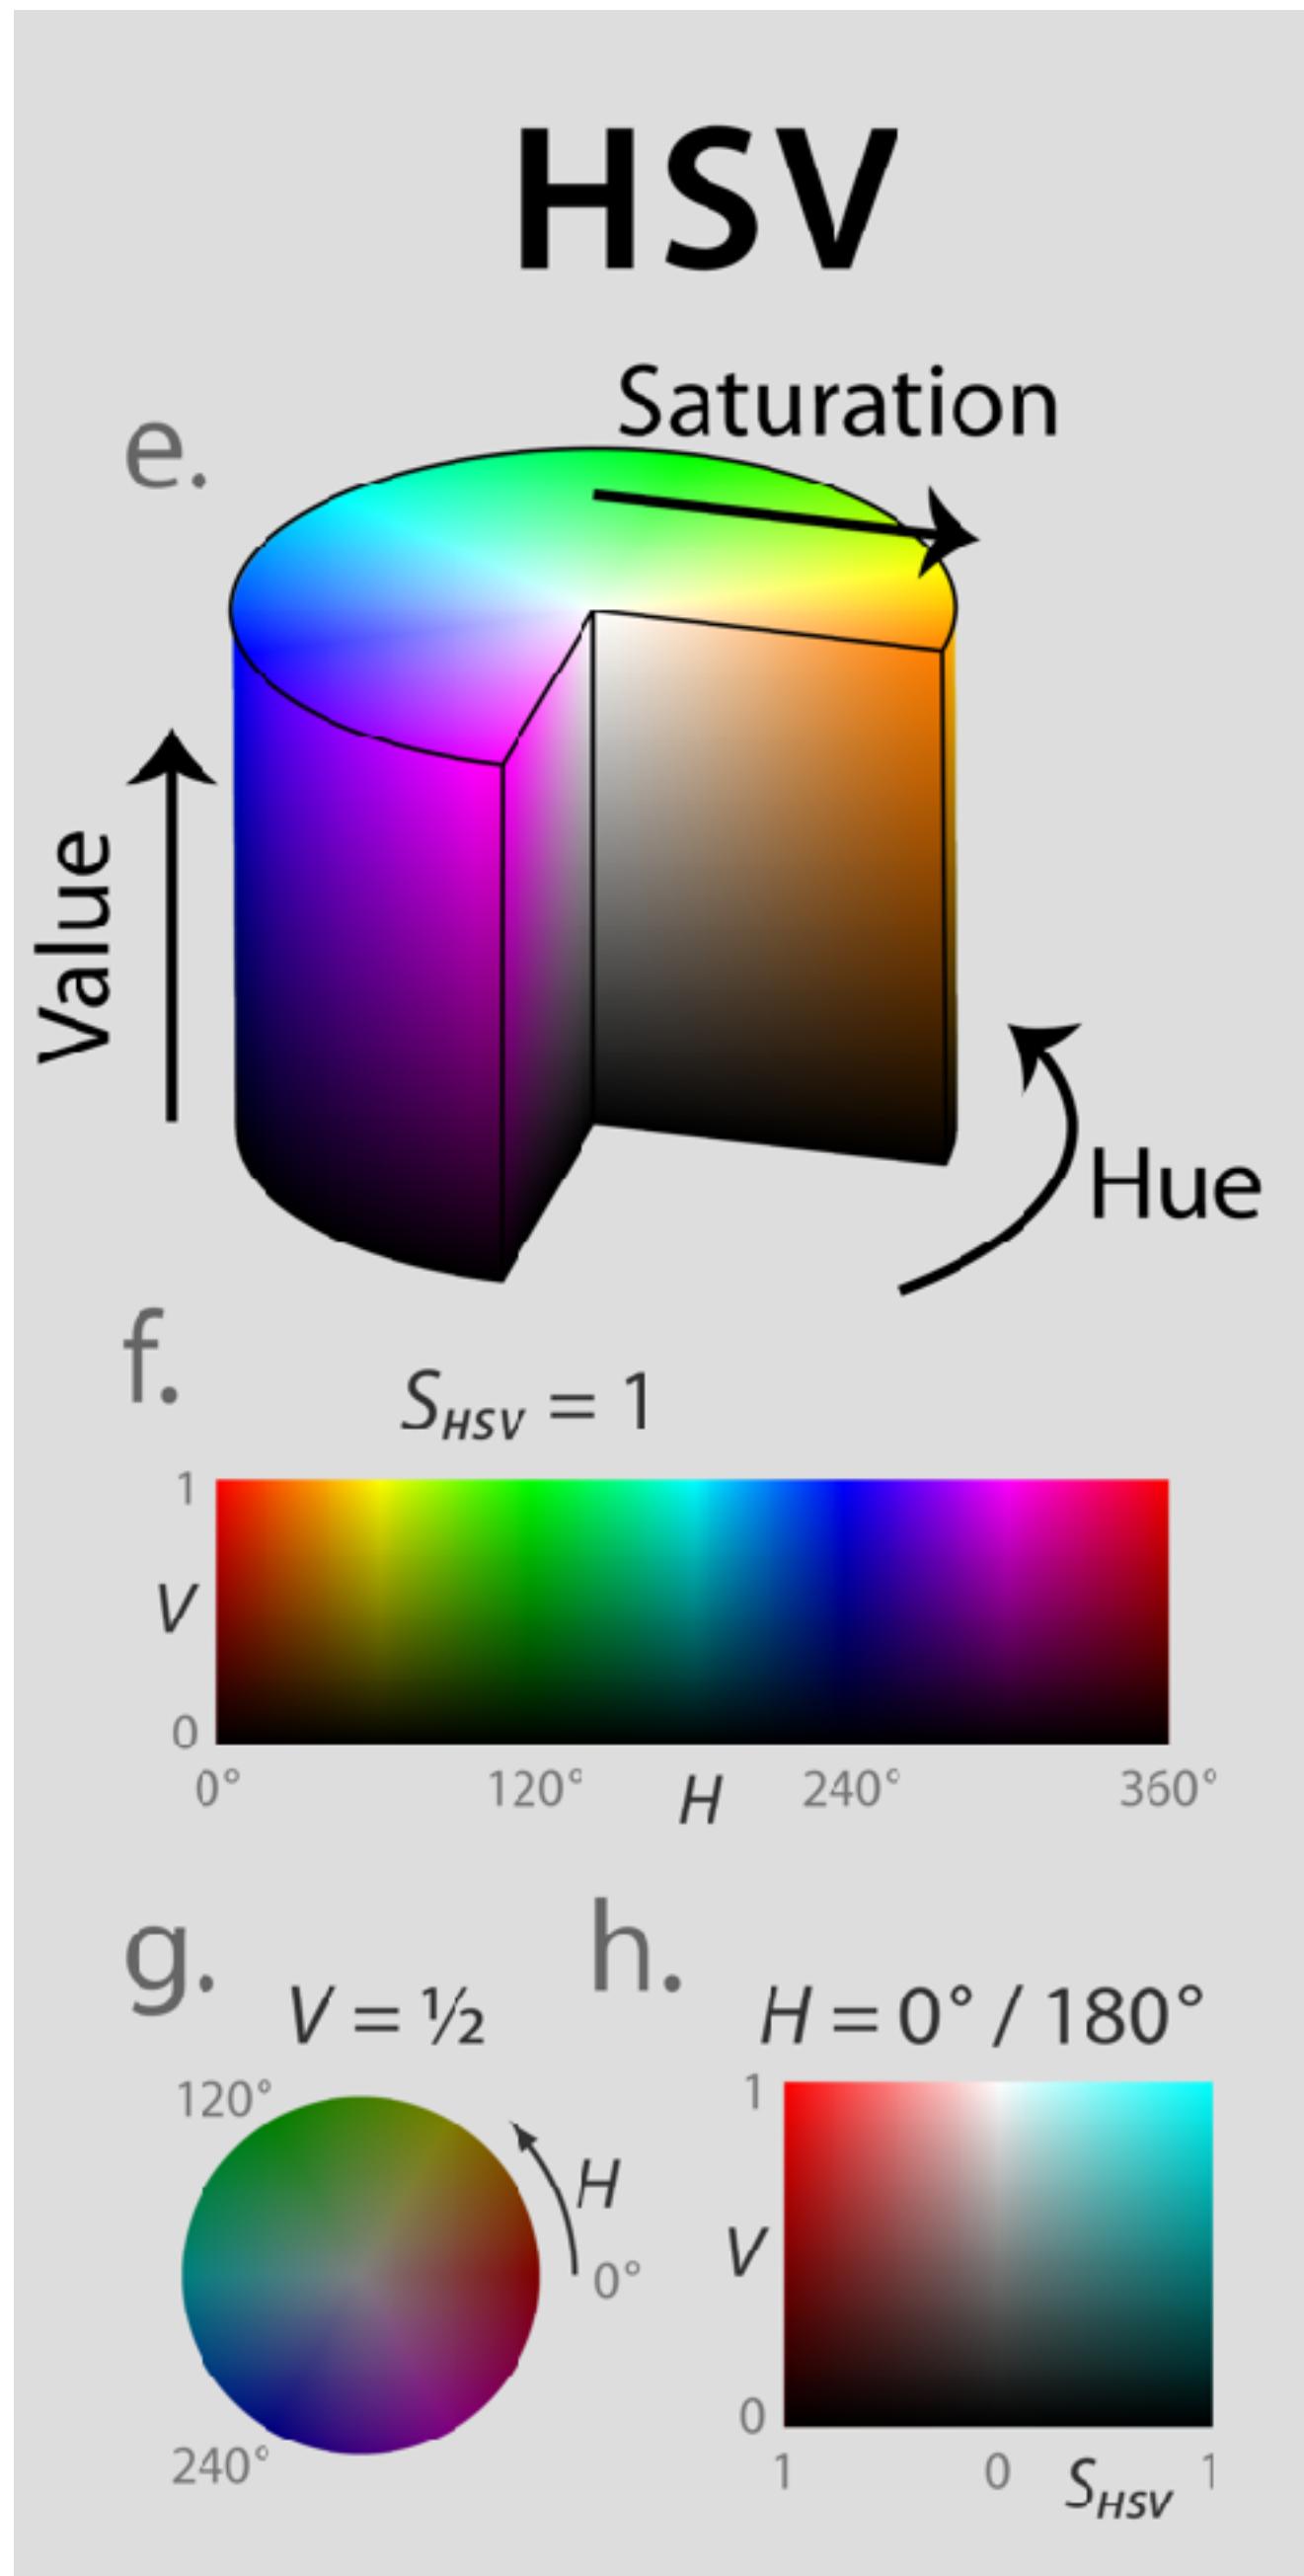
# HSV Colour Space

- How do we describe colour?
  - The sky is a bright blue
  - The trees are a dark green
  - That shirt is a bold colour
- HSV is closer to this:
  - Hue – what general colour is it
  - Saturation – how strong the colour is (distance from grey)
  - Value – how light or dark it is
  - Not perceptually uniform
- Also HSL, L = lightness/luminance
  - from dark to light



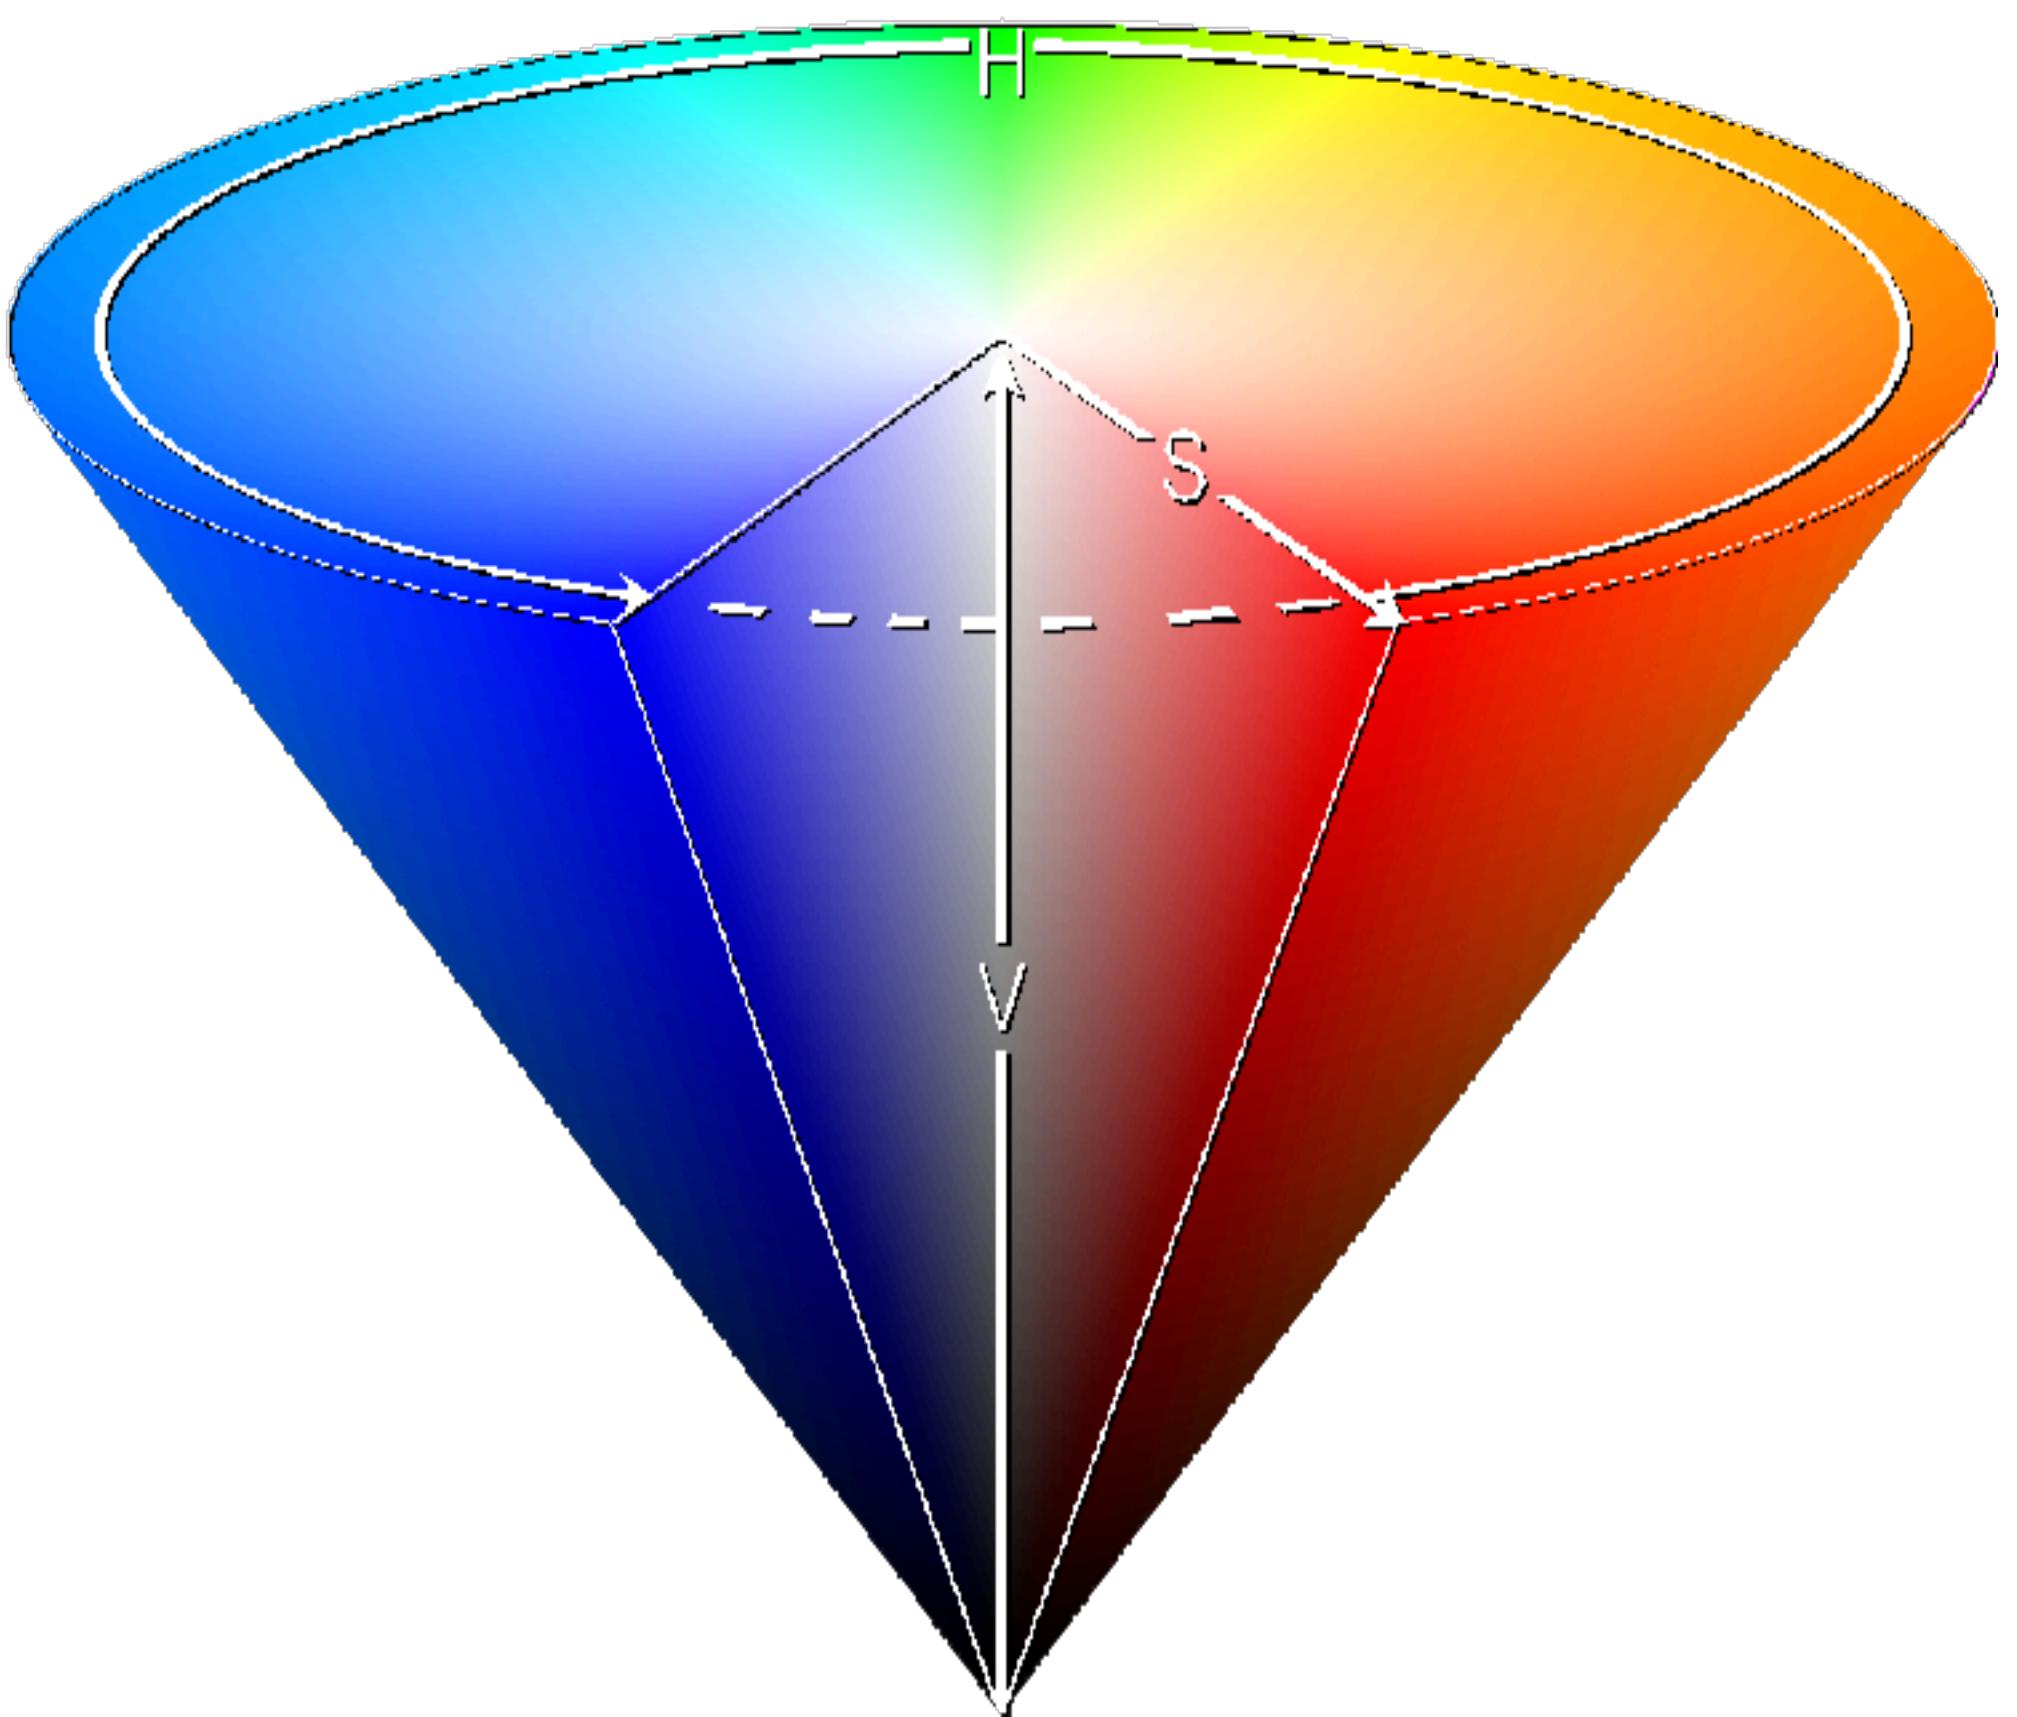
# HSV Colour Space

- How do we describe colour?
  - The sky is a bright blue
  - The trees are a dark green
  - That shirt is a bold colour
- HSV is closer to this:
  - Hue – what general colour is it
  - Saturation – how strong the colour is (distance from grey)
  - Value – how light or dark it is
  - Not perceptually uniform
- Also HSL, L = lightness/luminance
  - from dark to light



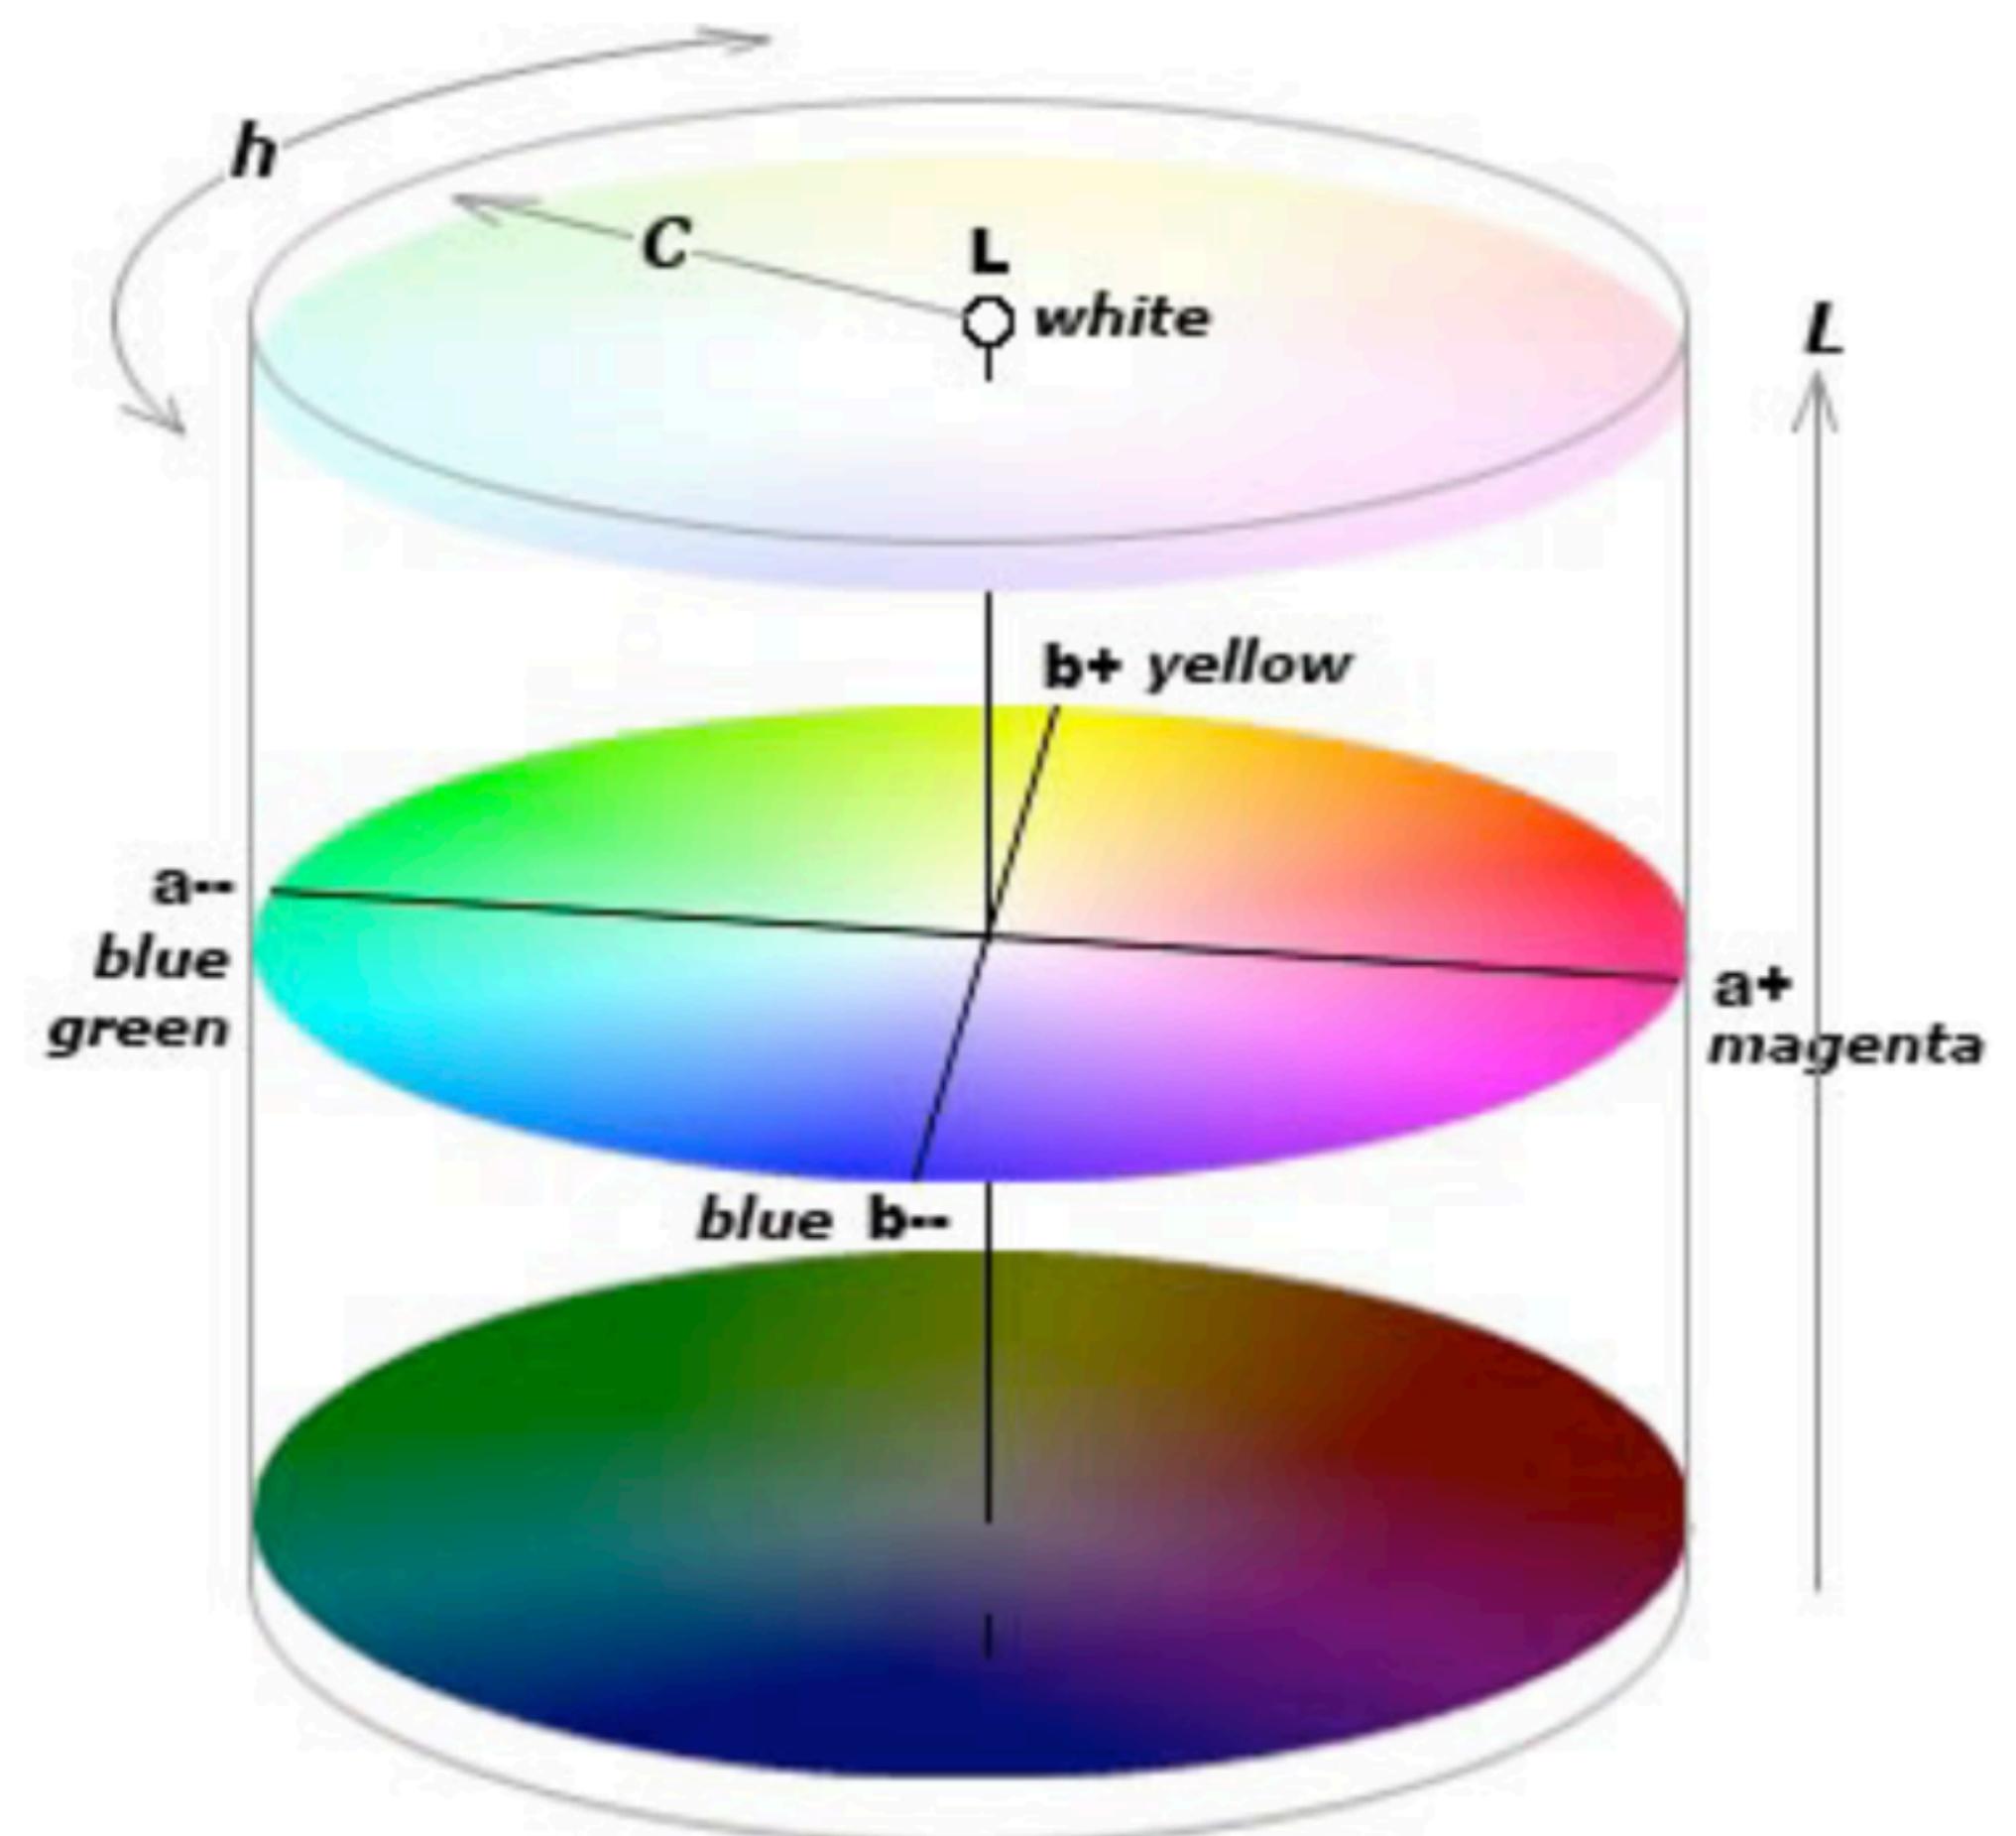
# HSV Colour Space

- How do we describe colour?
  - The sky is a bright blue
  - The trees are a dark green
  - That shirt is a bold colour
- HSV is closer to this:
  - Hue – what general colour is it
  - Saturation – how strong the colour is (distance from grey)
  - Value – how light or dark it is
  - Not perceptually uniform
- Also HSL, L = lightness/luminance
  - from dark to light



# CIELAB Colour Space

- Also LAB or  $L^*a^*b^*$
- Defined by International Commission on Illumination (CIE)
- Perceptually uniform space:
  - Numerical change corresponds to a similar perceived change in colour
- $L^*$  (lightness) coordinate normally ranges from [0,100]
- $a^*$  (red-green) and  $b^*$  (yellow blue) coordinates commonly clamped to [-128,127] or [0,255]



**Back to  
2D Image Representation**

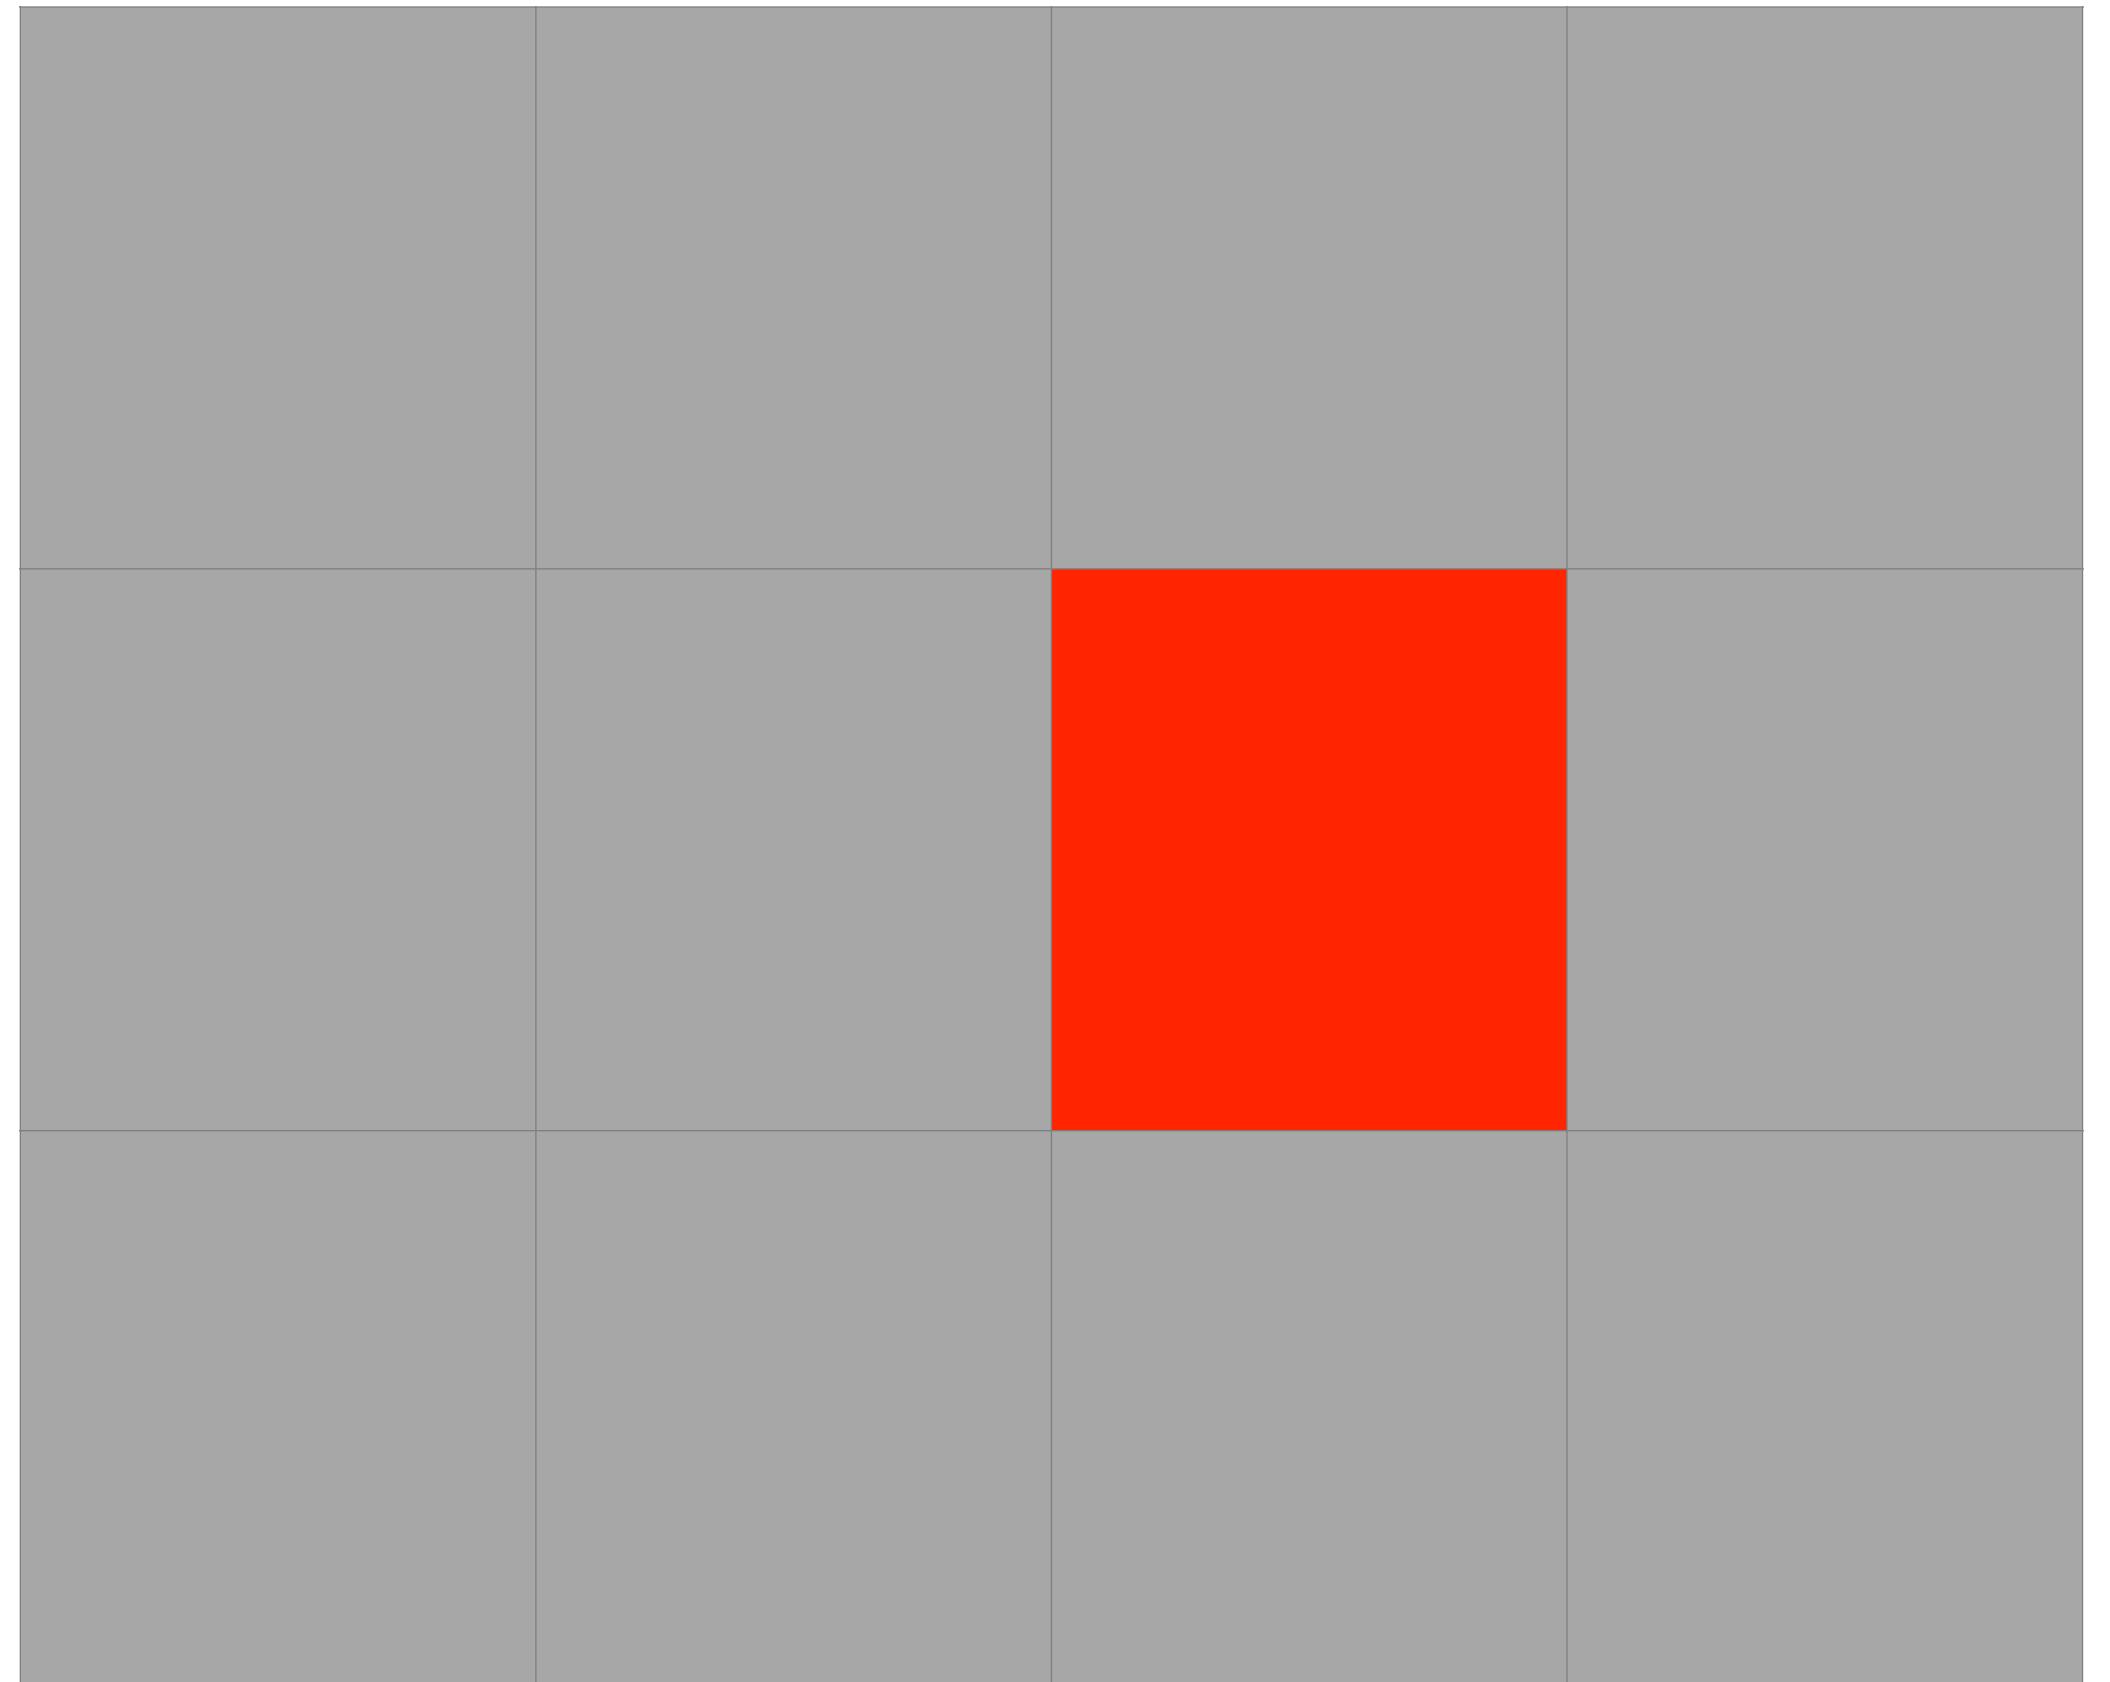
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125

# Access pixel at (row=1, col=2)
image2D[1][2] = ?
```



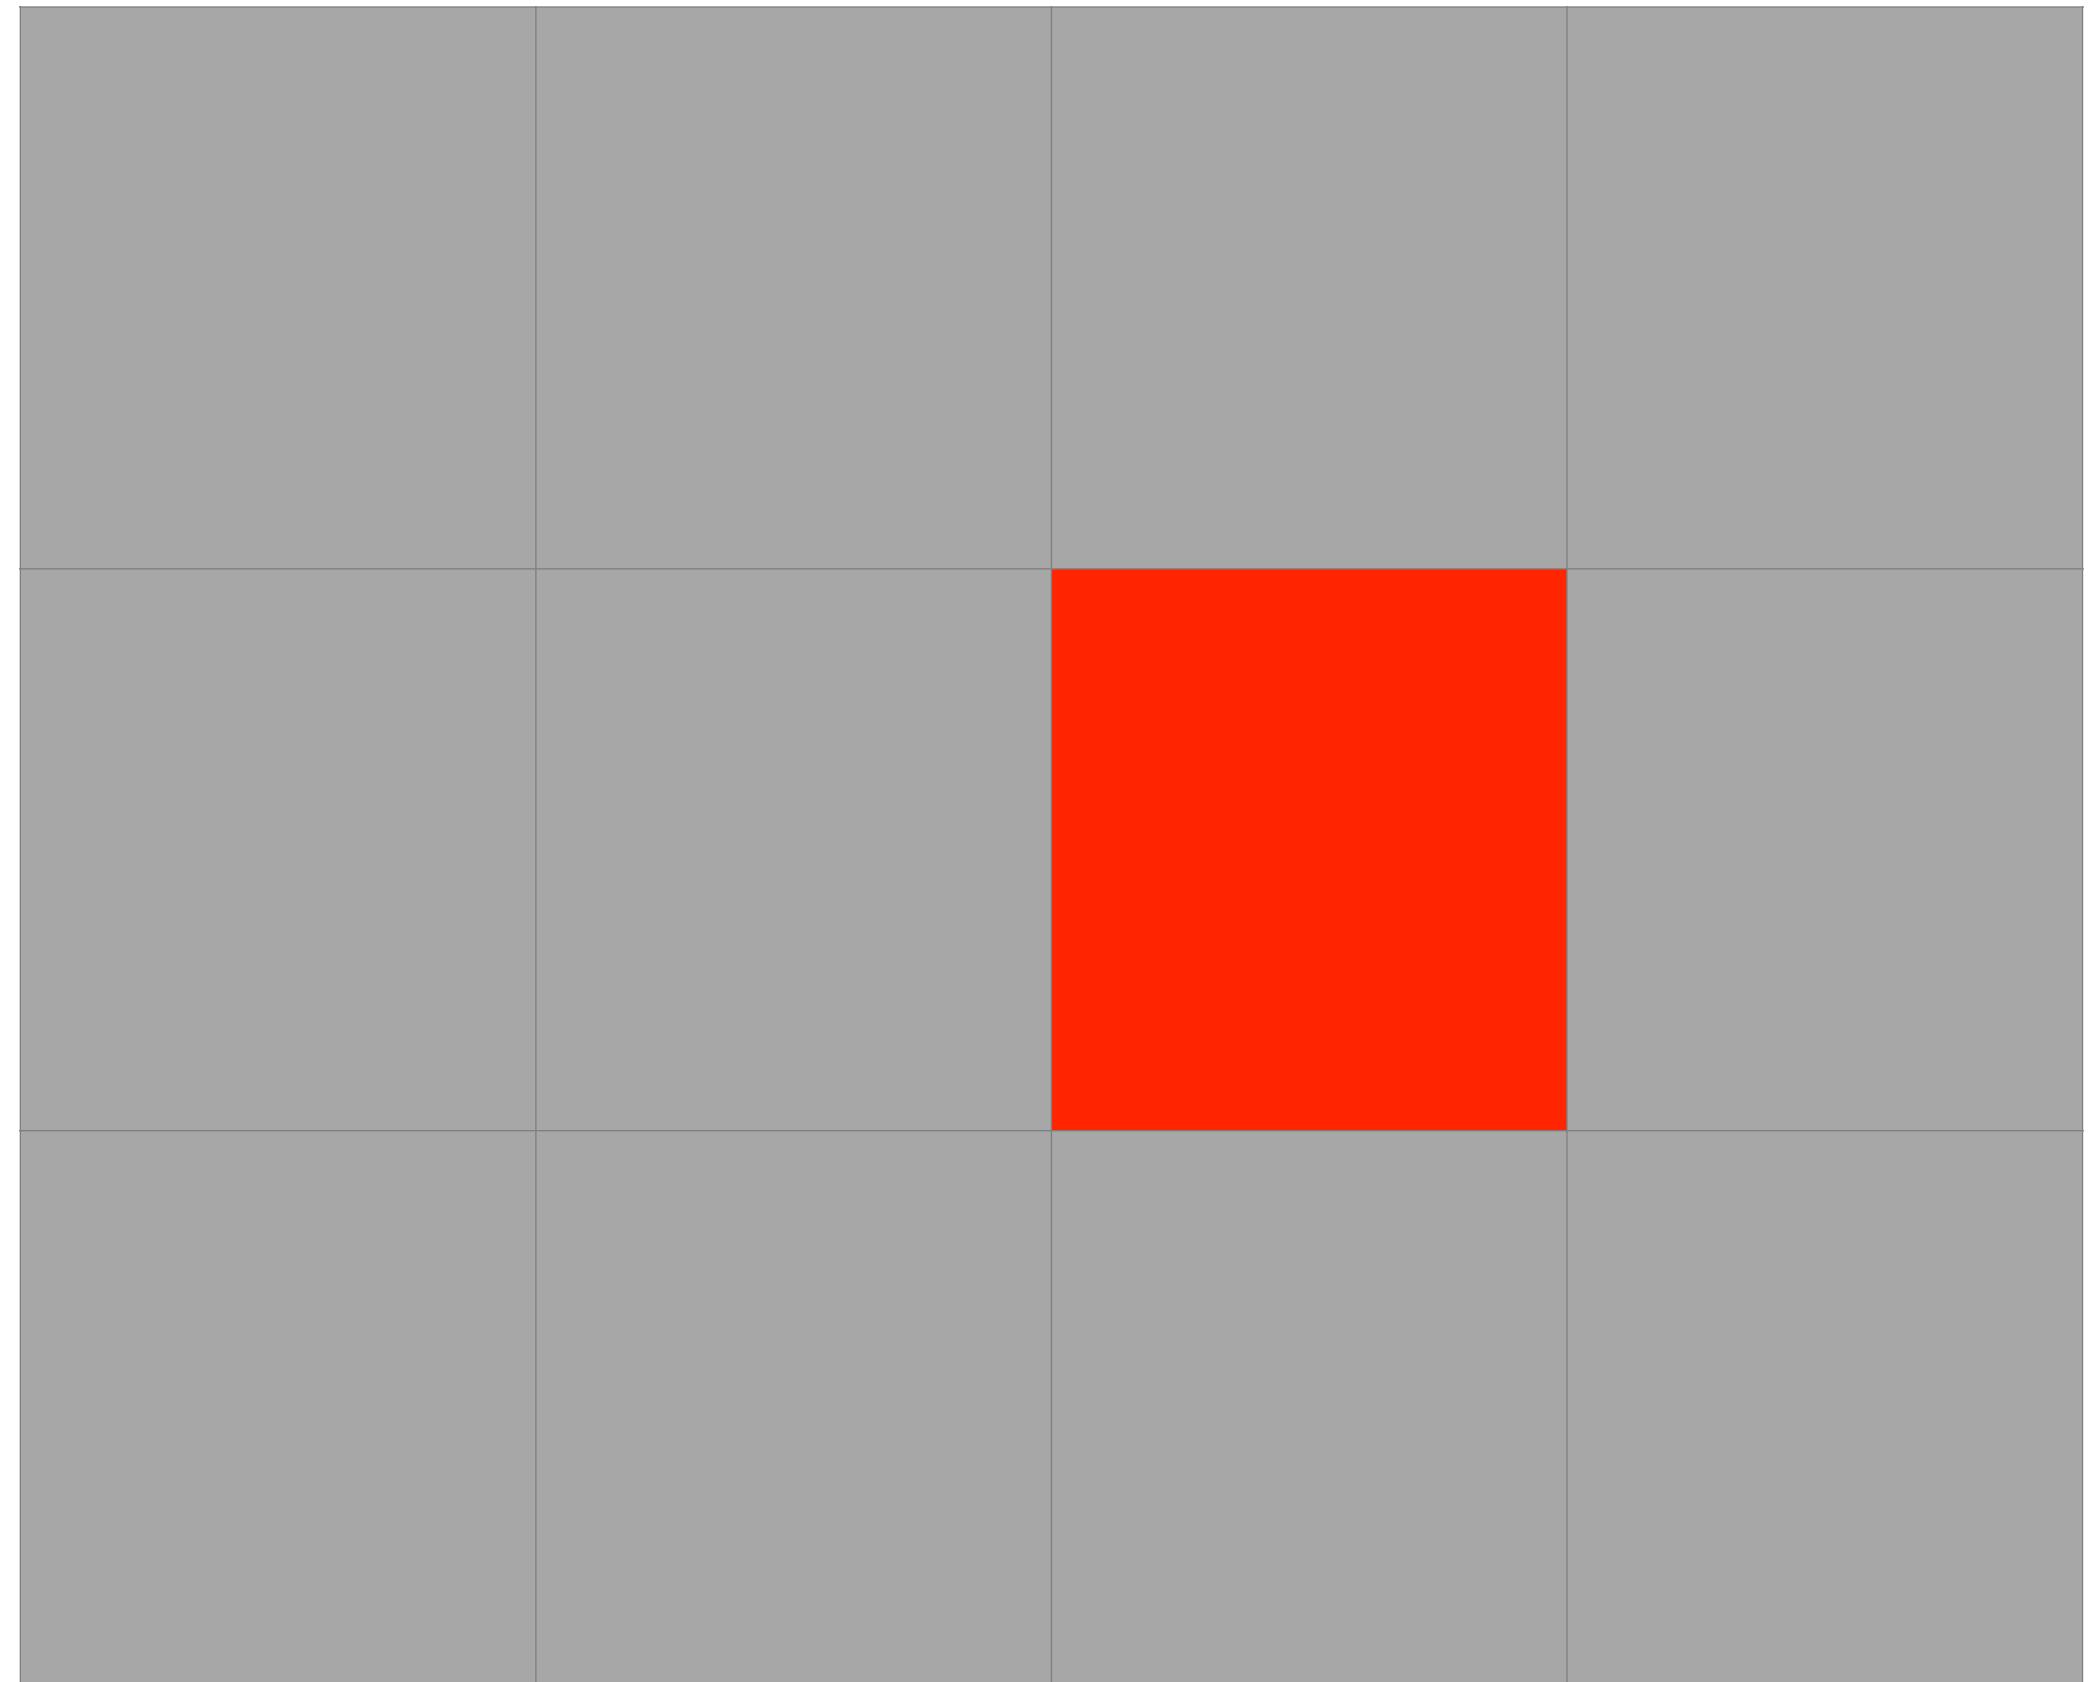
# Digital Images and 2D Image Representation

- 2D array of pixels – a raster
- Pixels have a value, or list of values
  - value: usually brightness
  - values: brightness, colour, transparency

```
# Create a 2D array
image2D = create_array(height, width)

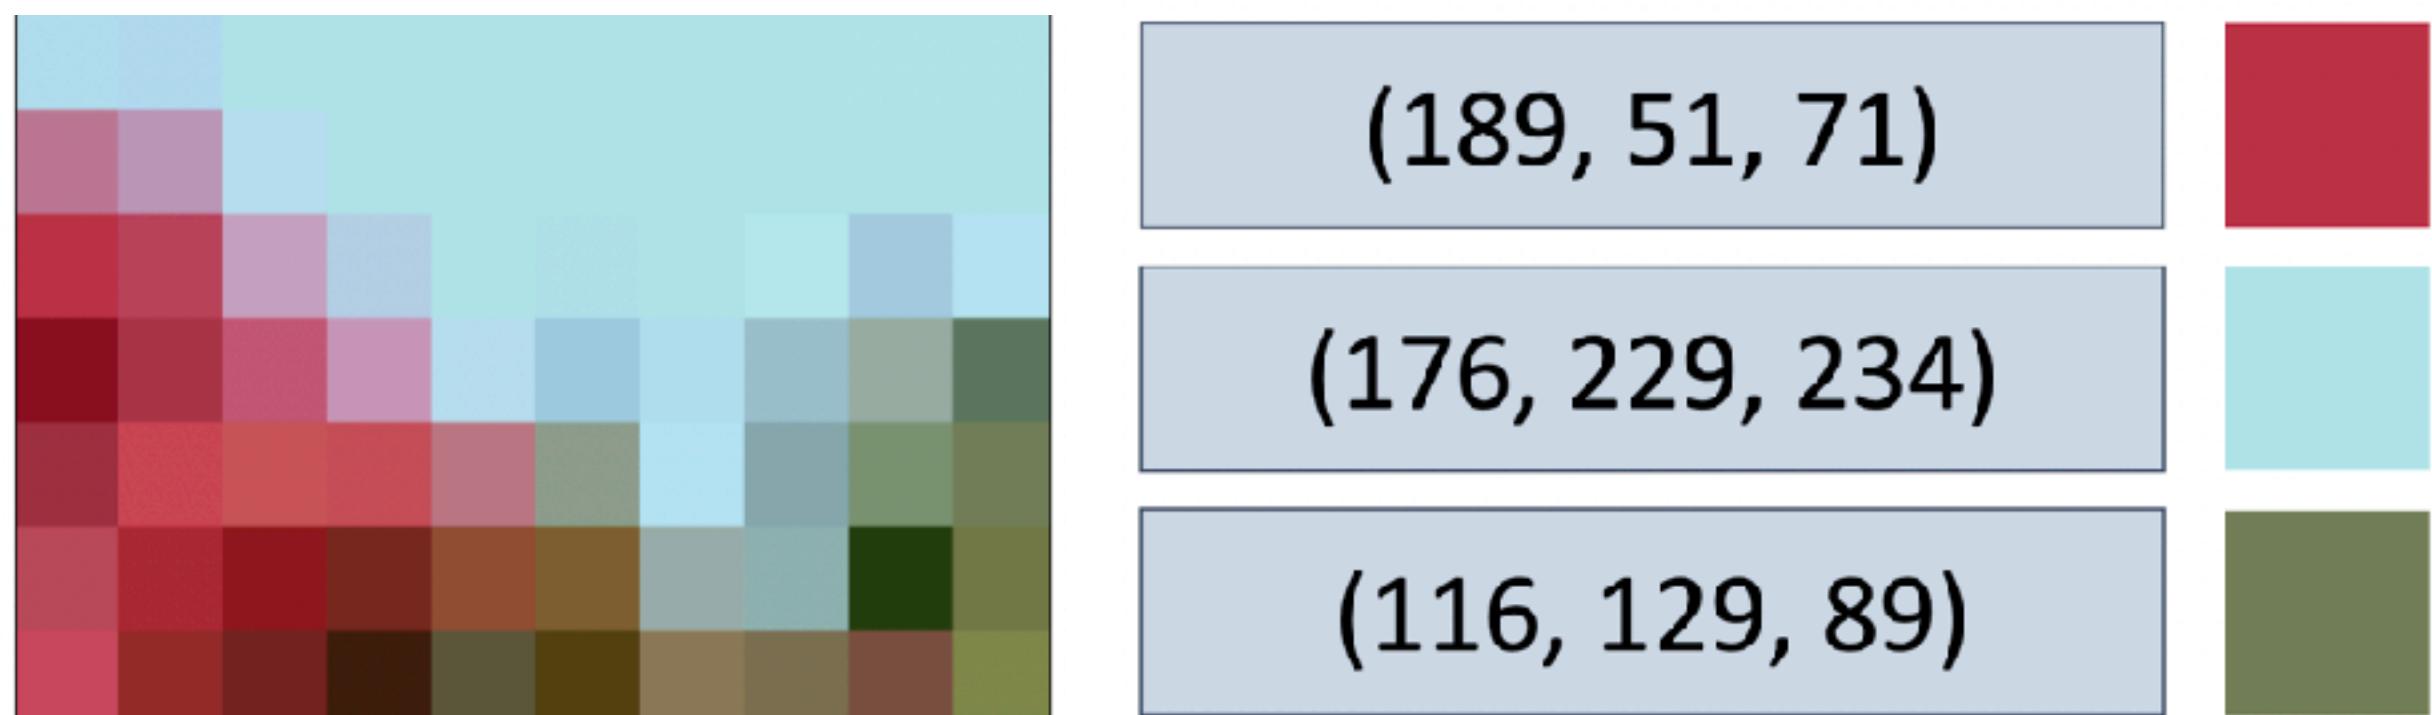
# Fill with example grayscale values
for y from 0 to height-1:
    for x from 0 to width-1:
        image2D[y][x] = 125

# Access pixel at (row=1, col=2)
image2D[1][2] = (255, 0, 0)
```



# Digital Images and 2D Image Representation

- 2D array pixels – a raster
- Pixels have a value, or list of values
- One value – greyscale image
- Colour images – typically 3 values (RGB)
- Values may be integers or floats
  - Integers in the range [0,255] (8bit / 1byte)
  - Floats in the range [0,1]
- High Dynamic Range:
  - Higher bit depth of 10,12,14bit or more per colour channel



# Digital Images and 2D Image Representation

- Images can be very large
- ‘Raw’ images are ~3 bytes per pixel
- 4K image ( $3840 \times 2160$ ) is ~24 MB
- Video is worse – 1 hour HD @ 25 fps:  
 $1920 \times 1080 \times 3 \times 25 \times 60 \times 60 \approx 560 \text{ GB}$   
 $3840 \times 2160 \times 3 \times 25 \times 60 \times 60 \approx 2.3 \text{ TB}$



- High Dynamic Range?
- 4K image ( $3840 \times 2160$ ) is ~36MB (at 12bit)



# Digital Images and 2D Image Representation

- Images can be very large
- ‘Raw’ images are ~3 bytes per pixel
- 4K image ( $3840 \times 2160$ ) is ~24 MB
- Video is worse – 1 hour HD @ 25 fps:  
 $1920 \times 1080 \times 3 \times 25 \times 60 \times 60 \approx 560 \text{ GB}$   
 $3840 \times 2160 \times 3 \times 25 \times 60 \times 60 \approx 2.3 \text{ TB}$



- Compression reduces file sizes
- Lossless compression
- Lossy compression
- Lossless video compression is rare

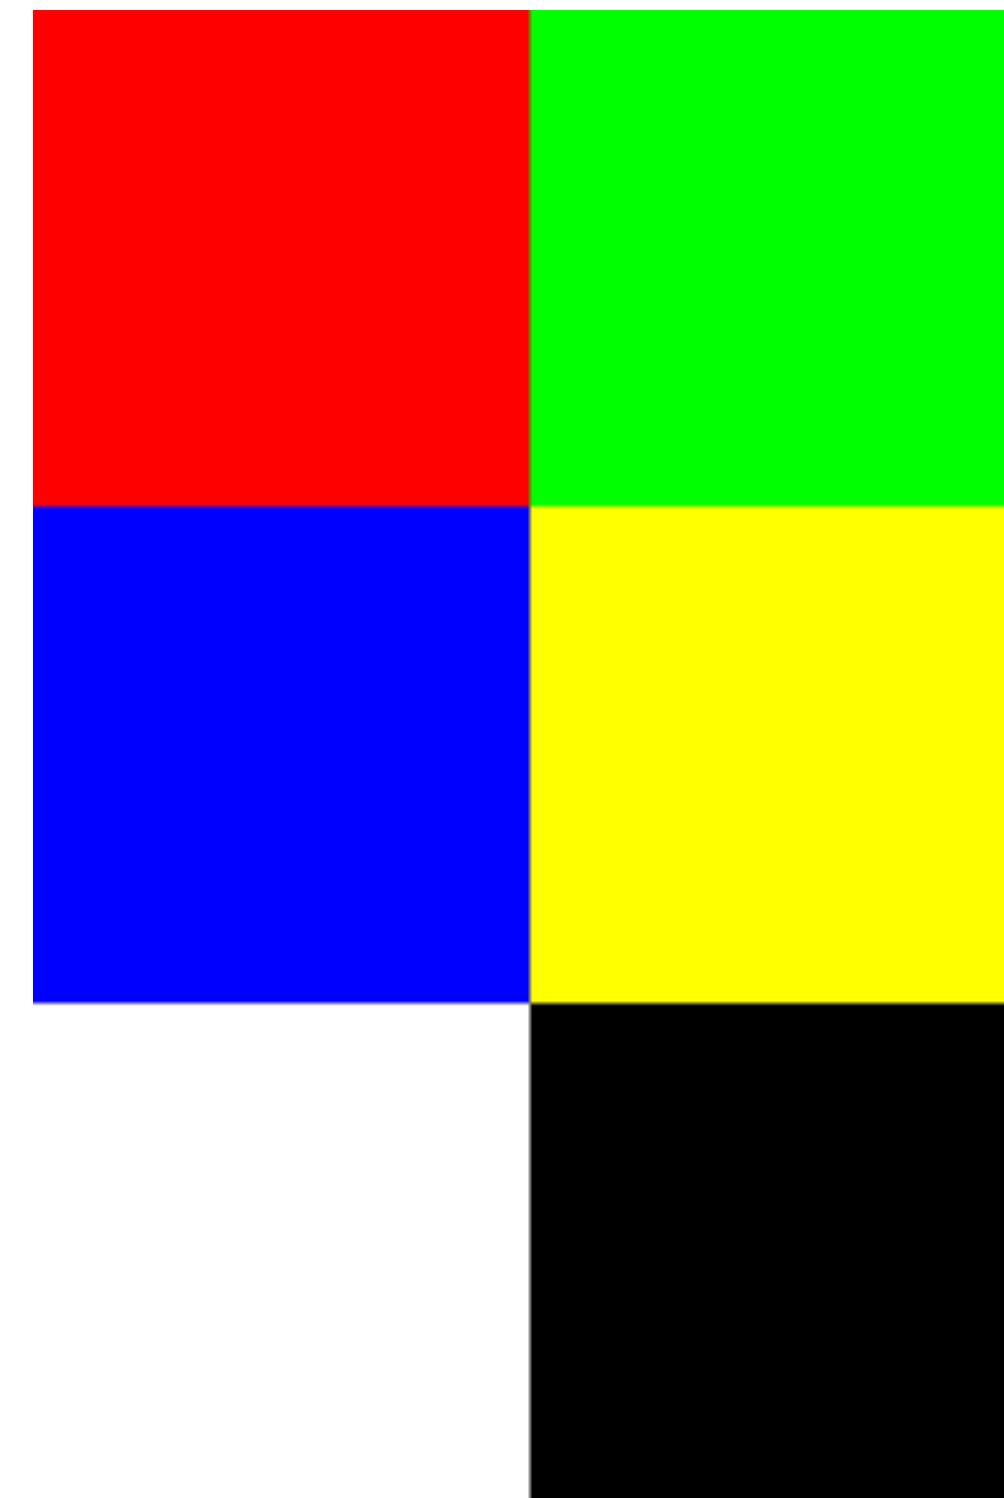


# Image File Formats

<b>Format</b>	<b>Compression</b>	<b>Notes</b>
<b>JPEG / JPG</b>	Lossy	Widely used for photos; high compression with quality trade-offs; supports EXIF metadata
<b>WebP</b>	Lossy or Lossless	Google-developed; better compression than JPEG/PNG; supports transparency and animation
<b>AVIF</b>	Lossy or Lossless	Based on AV1 codec; excellent compression and quality; supports HDR and transparency; growing browser support
<b>PNG</b>	Lossless	Ideal for graphics with sharp edges and transparency; universally supported; open format
<b>GIF</b>	Lossless (palette-based)	Supports animation; limited to 256 colors; largely replaced by APNG or video formats for animation
<b>BMP</b>	Lossless	Simple, uncompressed or RLE; mostly obsolete except for niche or legacy uses
<b>PNM (PBM/PGM/PPM)</b>	None	Very simple, raw or ASCII; mainly for research or tool pipelines
<b>SVG</b>	None (vector)	Scalable vector format for icons, illustrations, and UI; can embed raster images; supports CSS and scripting
<b>High Efficiency Video Coding (HEVC)</b>	Lossy	Used by High Efficiency Video Coding (H.265) but also in Apple HEIF or HEVC image and video files (e.g. heic)

# Sample PPM File

```
P3
# ASCII colour file
# 2x3 image, max value 255
2 3 255
# Content as RGB triples
255 0 0 0 255 0
0 0 255 255 255 0
255 255 255 0 0 0
```



# Sample PPM File

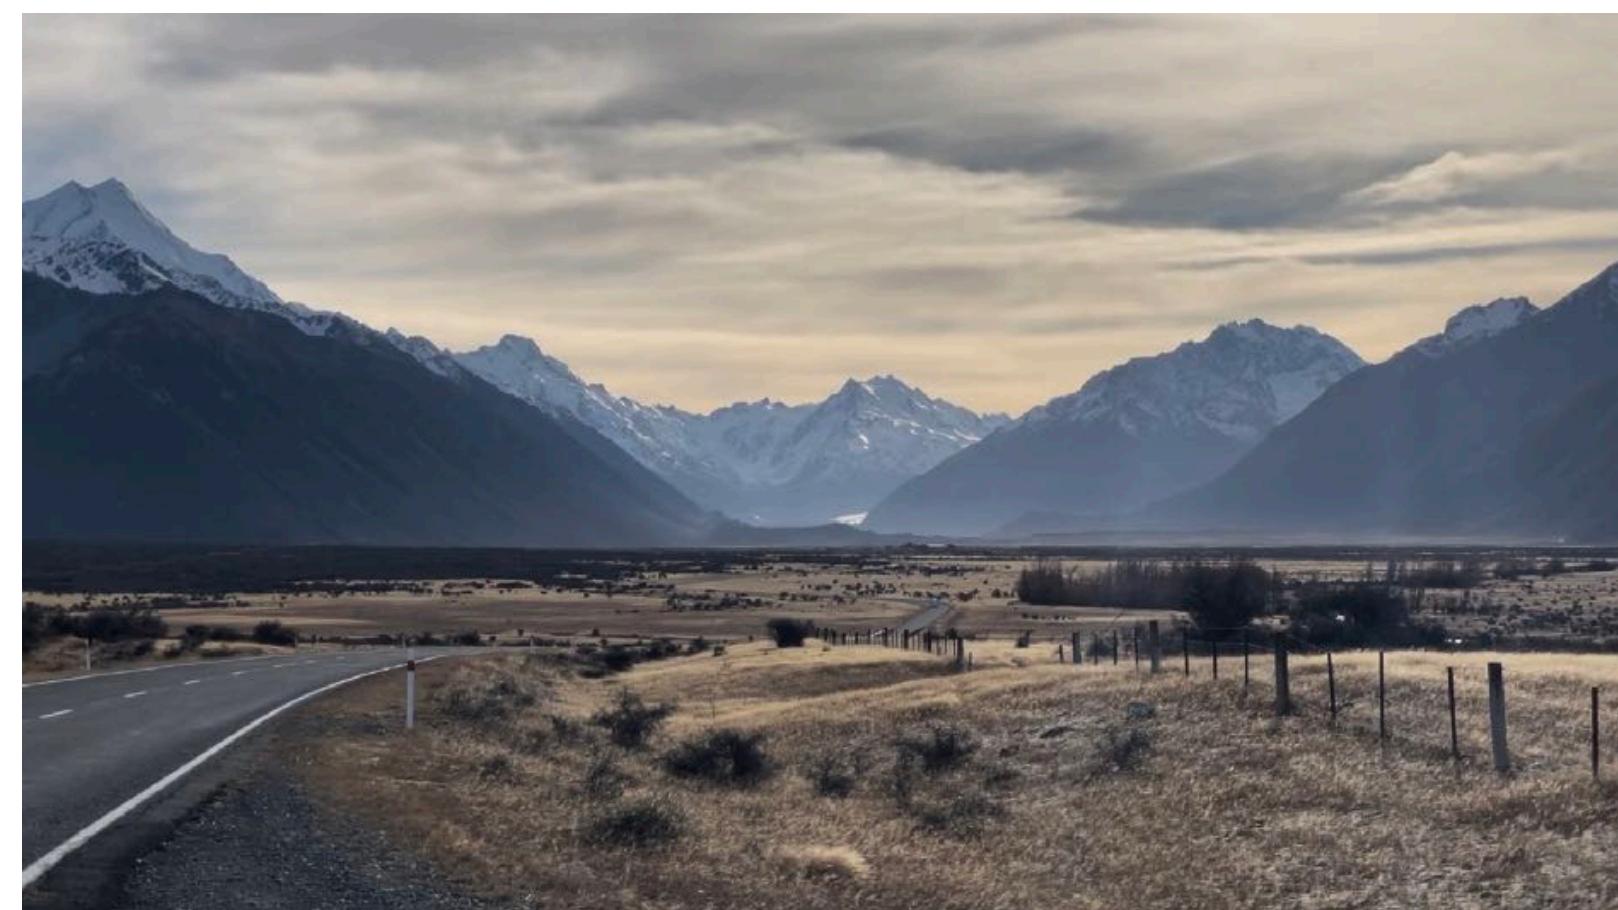
P3  
3433 1902

255

195 194 192 195 194 192 195 194 192 195 194 192 194 193 191 194 193 191 194 192 193 194 192 193 194 192 193 194 192 193 194  
190 190 190 190 190 190 190 190 190 190 190 190 190 188 189 190 188 189 190 188 189 191 189 190 191 189 190 191 189 190 191 189 190  
192 191 189 192 191 189 192 191 189 193 192 190 193 192 190 194 193 191 194 193 191 194 193 191 194 193 191 194 193 191 194  
191 196 192 191 196 192 191 196 192 191 195 191 190 196 192 191 196 192 191 196 192 191 195 191 190 195 191 190 195 191 190  
192 191 189 192 191 189 192 191 189 192 191 189 192 191 189 192 191 189 193 192 190 193 192 190 193 192 190 193 192 190 193  
190 195 194 190 194 193 189 194 193 189 194 193 189 194 193 189 195 194 190 195 194 190 194 193 189 194 193 189 193 192 188 193 192 188  
186 185 183 186 185 183 186 185 183 186 185 183 187 186 184 187 186 184 187 186 184 186 185 183 186 185 183 186 185 183 185 184 182 185  
184 182 182 184 182 182 184 181 181 183 181 181 183 181 181 183 181 181 183 181 181 183 180 180 182 180 180 182 180 180 182  
179 179 181 178 178 180 178 178 180 178 178 180 178 178 180 178 178 180 178 178 180 178 178 180 178 178 180 177 177 179 177 177 179 177  
178 176 176 178 176 176 178 175 176 178 175 176 178 175 176 178 175 176 178 176 176 178 175 175 177 175 175 177 175 175 175  
174 174 176 174 174 176 175 175 177 175 175 177 175 175 177 175 175 177 174 174 176 174 174 176 174 174 176 174 174 176 173  
166 164 164 166 164 164 166 164 164 166 164 164 166 163 163 165 163 163 165 163 163 165 162 162 164 162 162 164 162 162 164  
158 159 163 158 159 163 158 159 163 158 159 163 157 158 162 157 158 162 157 158 162 157 158 162 157 158 162 157 158 162 157  
160 155 156 160 155 156 160 155 156 160 155 156 160 154 155 159 154 155 159 154 155 159 154 155 159 154 155 159 154 155 159  
152 153 157 152 153 157 152 153 157 152 153 157 152 153 157 152 153 157 152 153 157 152 153 157 152 153 157 152 153 157 152 153 157  
154 151 152 154 151 152 154 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156  
151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156  
156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156 151 152 156  
154 153 158 | 154 153 158 154 153 158 154 153 158 154 153 158 154 153 158 154 153 158 154 153 158 154 153 158 154 153 158 154 153 158  
154 152 152 154 152 152 154 152 152 154 152 152 154 152 151 156 152 151 156 152 151 156 152 151 156 152 151 156 152 151 156  
151 150 155 151 150 155 151 150 155 151 150 155 151 150 155 151 150 155 151 150 155 151 150 155 151 150 155 151 150 155 152 151 156  
152 148 147 152 148 147 152 148 147 152 148 147 152 148 147 152 148 147 152 148 147 152 148 147 152 148 147 152 148 147 152  
144 145 149 144 145 149 144 145 149 144 145 149 144 145 149 144 145 149 144 145 149 144 145 149 144 145 149 144 145 149 144 145 149  
148 143 144 148 143 144 148 143 144 148 143 144 148 143 144 148 143 144 148 143 144 148 143 144 148 143 144 148 143 144 148  
144 143 148 144 143 148 144 143 148 143 144 146 143 144 146 143 144 146 143 144 146 143 144 146 143 144 146 143 144 146  
147 145 145 147 145 145 147 145 145 147 145 145 147 145 145 147 145 145 147 145 145 147 145 145 147 145 145 147 145 145 147  
149 149 151 149 149 151 149 149 151 149 149 151 149 149 151 149 149 151 149 149 151 149 149 151 149 149 151 149 149 151 149 149 151  
153 151 151 153 151 151 153 151 151 153 151 152 152 154 152 152 154 152 152 154 152 152 154 152 152 154 152 152 154 152 152 154  
153 153



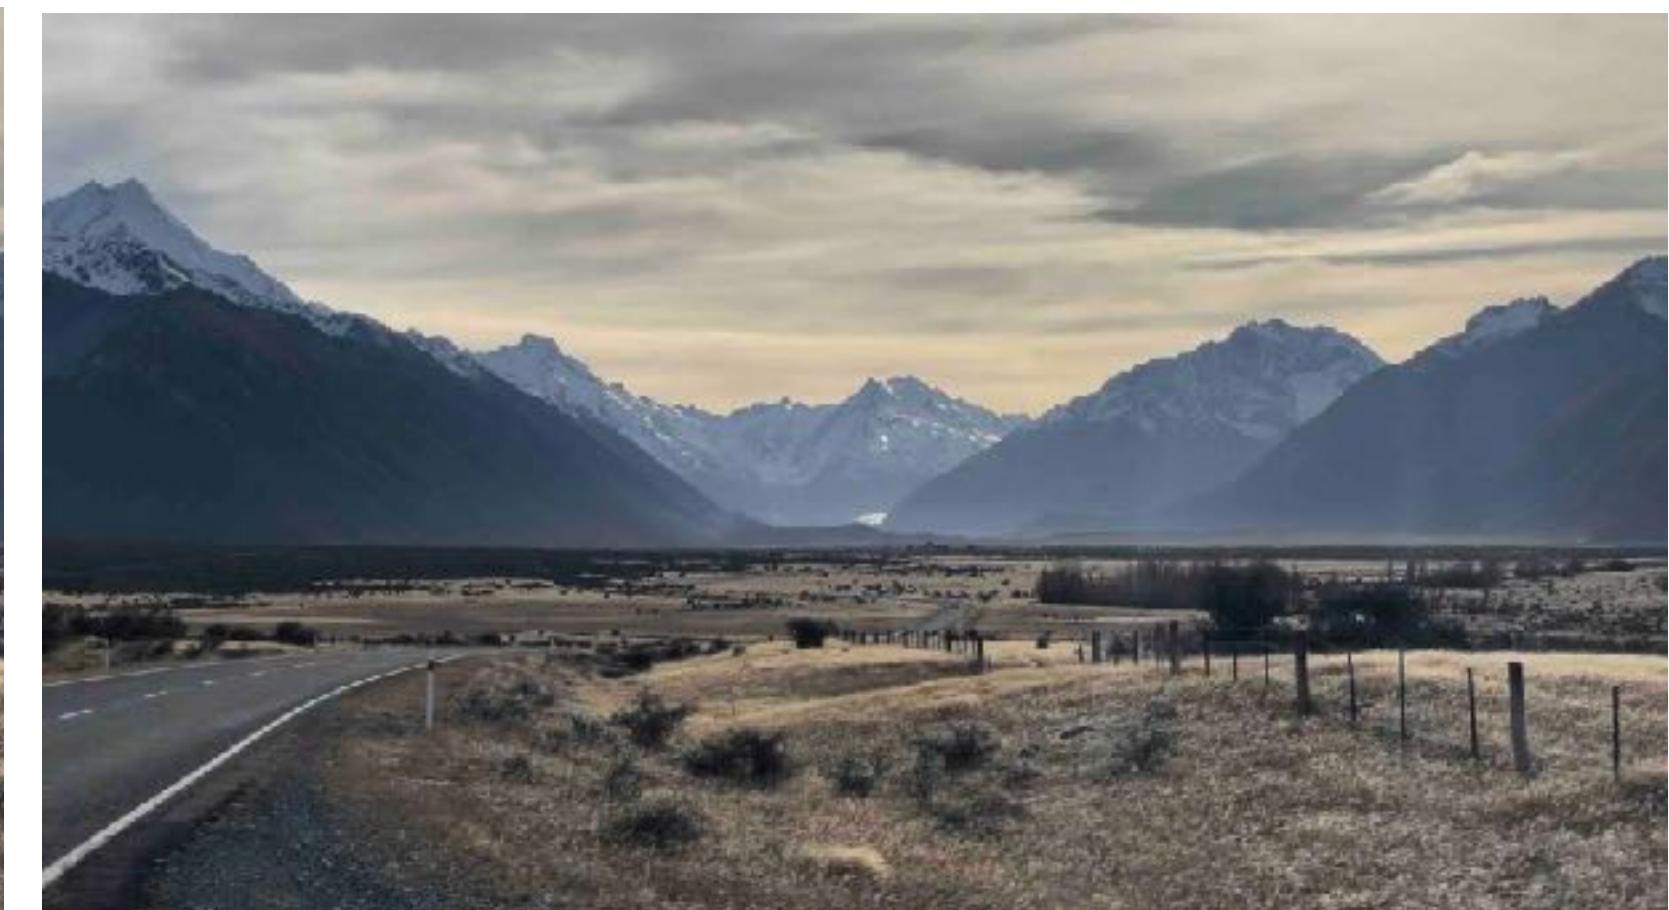
# JPEG Compression Artifacts



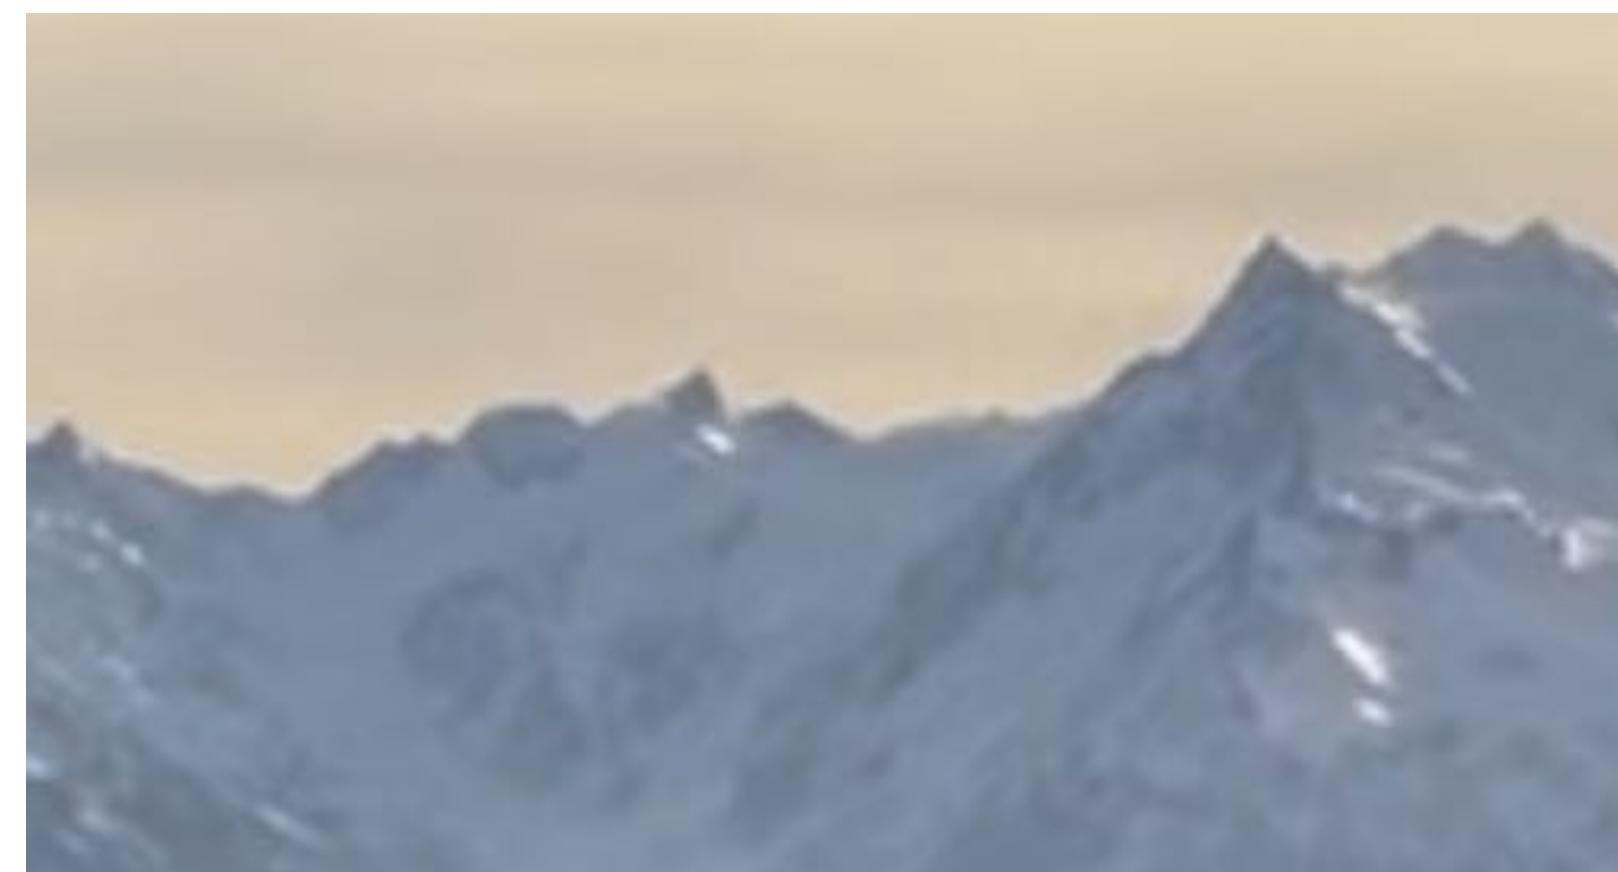
Quality factor 100



Quality factor 50



Quality factor 10



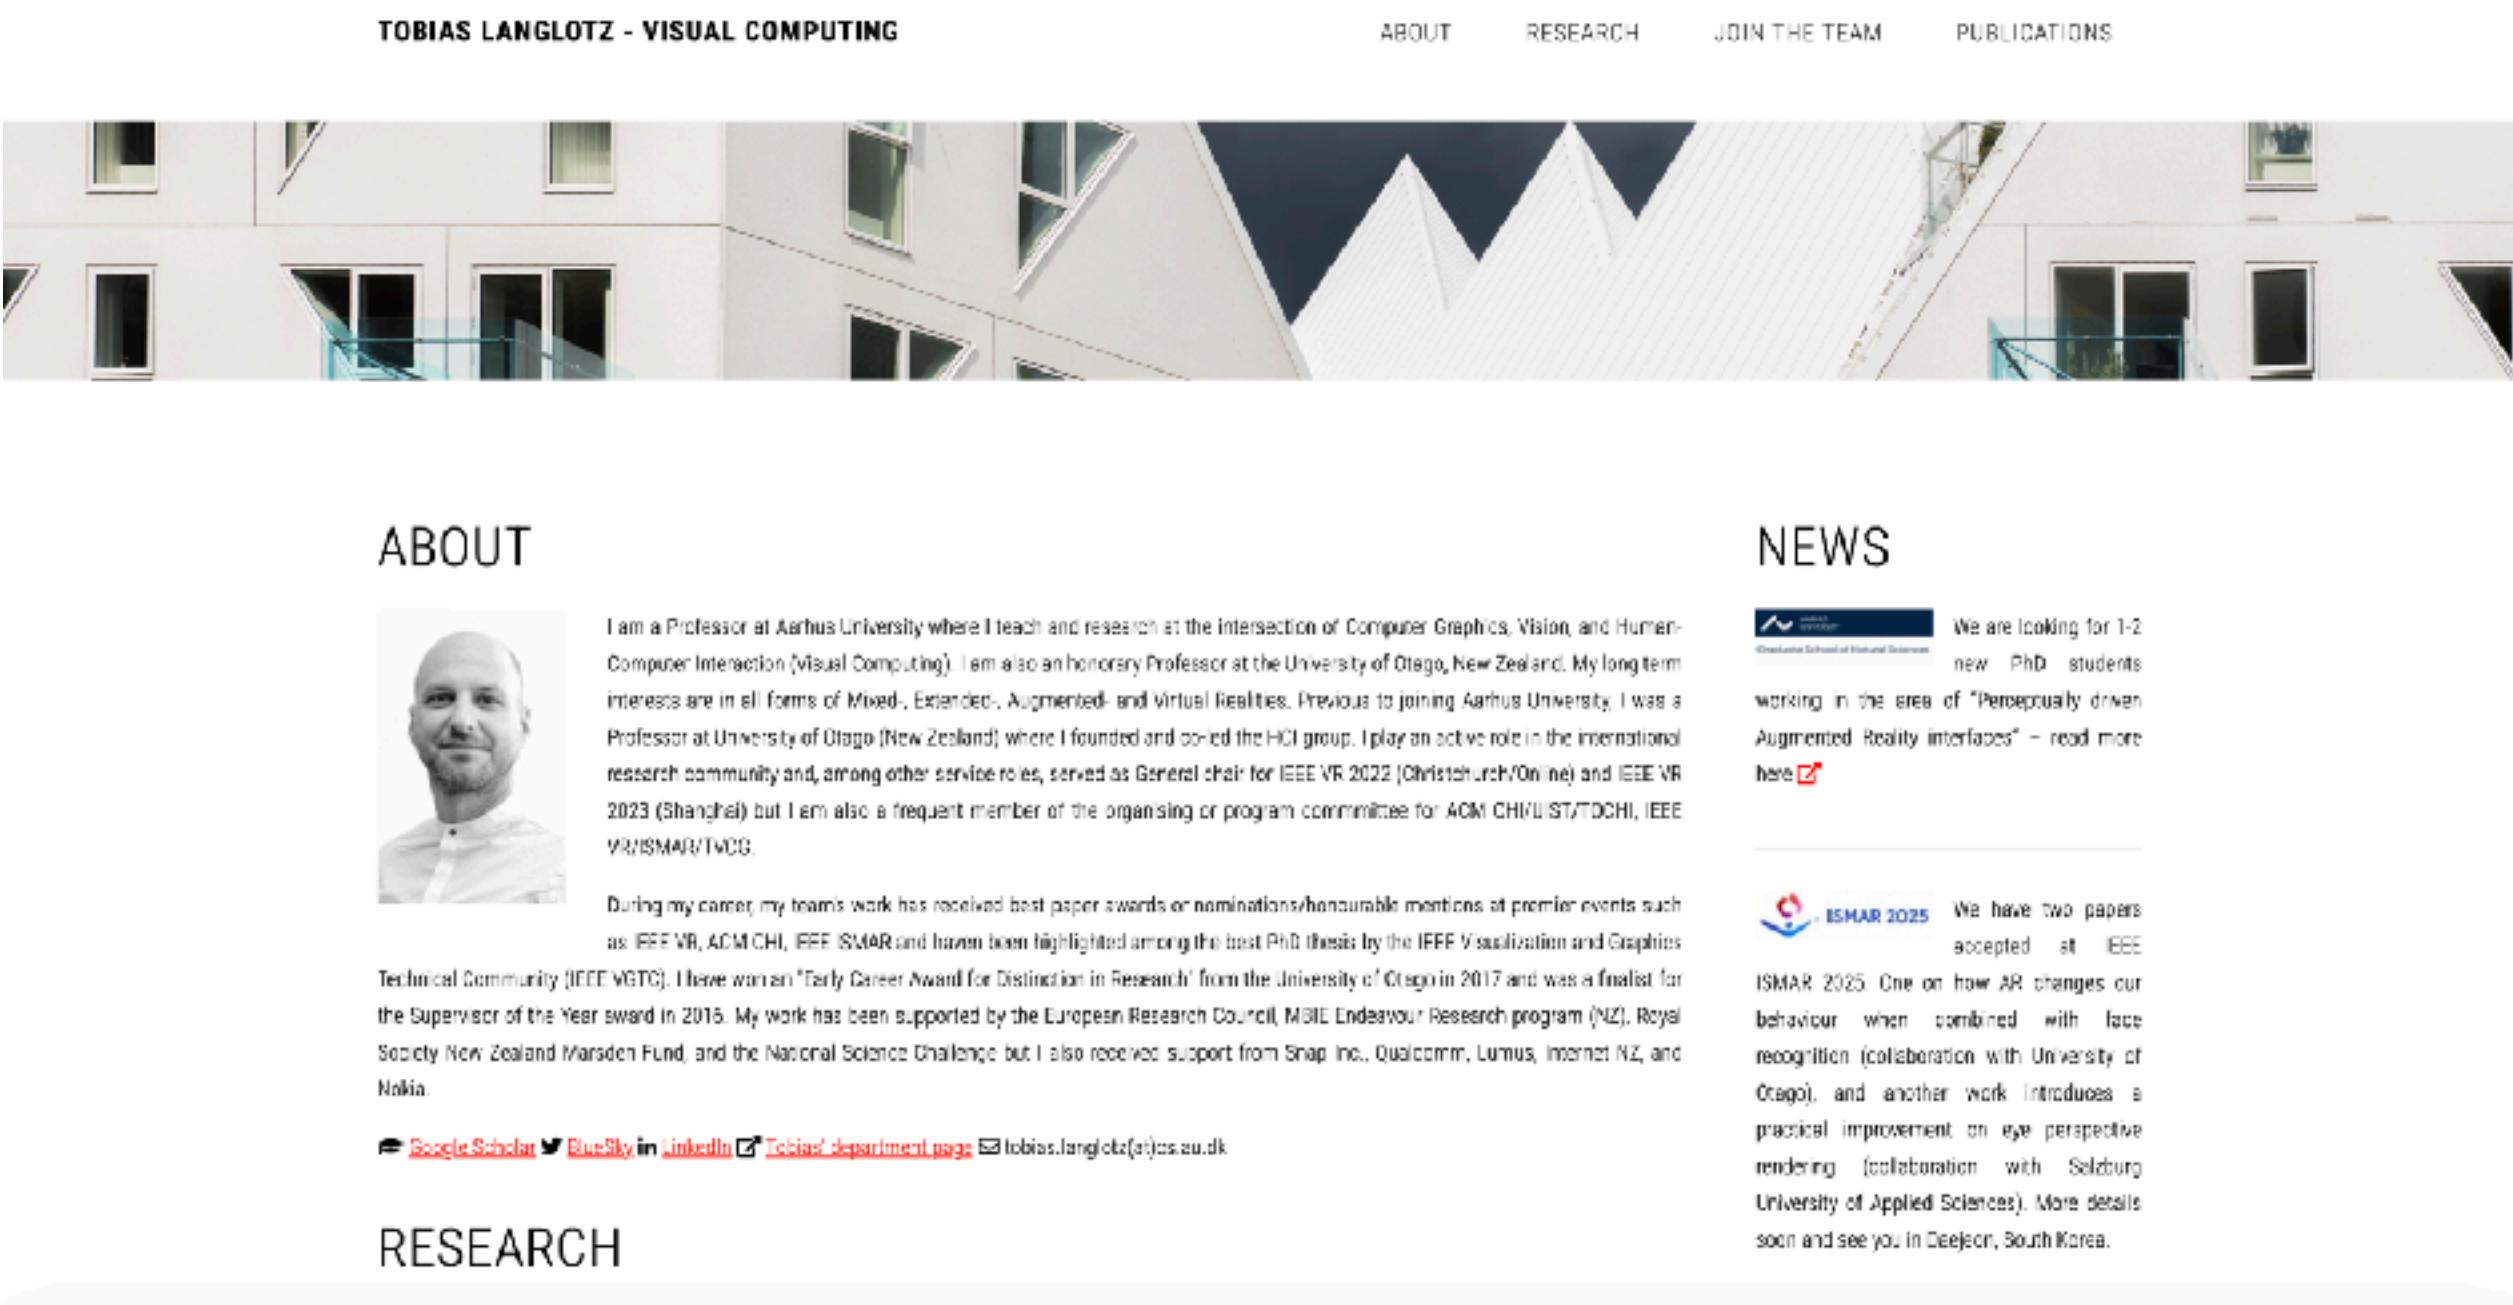
# Format Size Comparison

<i>Format</i>	<i>Size</i>
<b>JPEG (QF 90)</b>	<b>995 KB</b>
<b>JPEG (QF 50)</b>	<b>379 KB</b>
<b>JPEG (QF 10)</b>	<b>205 KB</b>
<b>GIF</b>	<b>2.3 MB</b>
<b>PNG</b>	<b>4.9 MB</b>
<b>BMP</b>	<b>19.5 MB</b>
<b>PPM (Binary)</b>	<b>19.7 MB</b>
<b>PPM (ASCII)</b>	<b>91.4 MB</b>



# Task

- Think about
  - What file format(s) would be best
    - For photographs on a website?
    - For logos?
    - For professional photo editing?
  - For scientific / medical imaging?
  - What about videos?
    - Is a series of JPEGs good enough?
    - What could help us do better?



The screenshot shows a modern building facade with large windows and a geometric patterned roof. The website header reads "TOBIAS LANGLOTZ - VISUAL COMPUTING" with links for "ABOUT", "RESEARCH", "JOIN THE TEAM", and "PUBLICATIONS". Below the header is a large image of the building. The "ABOUT" section features a portrait of Tobias Langlotz and text about his academic background and research interests. The "NEWS" section includes a link to a job posting for PhD students at the University of Aarhus.

**ABOUT**

I am a Professor at Aarhus University where I teach and research at the intersection of Computer Graphics, Vision, and Human-Computer Interaction (Visual Computing). I am also an honorary Professor at the University of Otago, New Zealand. My long term interests are in all forms of Mixed-, Extended-, Augmented- and Virtual Realities. Previous to joining Aarhus University, I was a Professor at University of Otago (New Zealand) where I founded and co-led the HCI group. I play an active role in the international research community and, among other service roles, served as General chair for IEEE VR 2022 (Christchurch/NZ) and IEEE VR 2023 (Shanghai) but I am also a frequent member of the organising or program committee for ACM CHI/IUI/STOC/CHI, IEEE VR/ISMAR/VGTC.

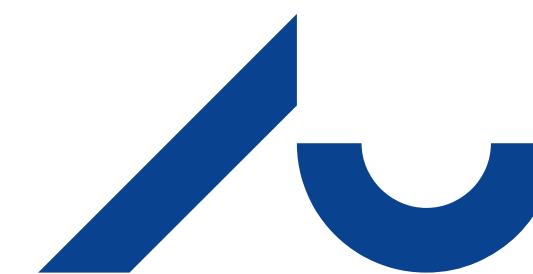
During my career, my team's work has received best paper awards or nominations/honorable mentions at premier events such as IEEE VR, ACM CHI, IEEE ISMAR and have been highlighted among the best PhD thesis by the IEEE Visualization and Graphics Technical Community (IEEE VGTC). I have won an "Early Career Award for Distinction in Research" from the University of Otago in 2017 and was a finalist for the Supervisor of the Year award in 2016. My work has been supported by the European Research Council, MLE Endeavour Research program (NZ), Royal Society New Zealand Marsden Fund, and the National Science Challenge but I also received support from Snap Inc., Qualcomm, Lumus, Internet NZ, and Nokia.

[Google Scholar](#) [DuoSky](#) [LinkedIn](#) [Tobias' department page](#) [tobias.langlotz@cs.au.dk](mailto:tobias.langlotz@cs.au.dk)

**NEWS**

We are looking for 1-2 new PhD students working in the area of "Perceptually driven Augmented Reality interfaces" – read more [here](#).

**ISMAR 2025** We have two papers accepted at IEEE ISMAR 2025. One on how AR changes our behaviour when combined with face recognition (collaboration with University of Otago), and another work introduces a practical improvement on eye perspective rendering (collaboration with Salzburg University of Applied Sciences). More details soon and see you in Daegu, South Korea.

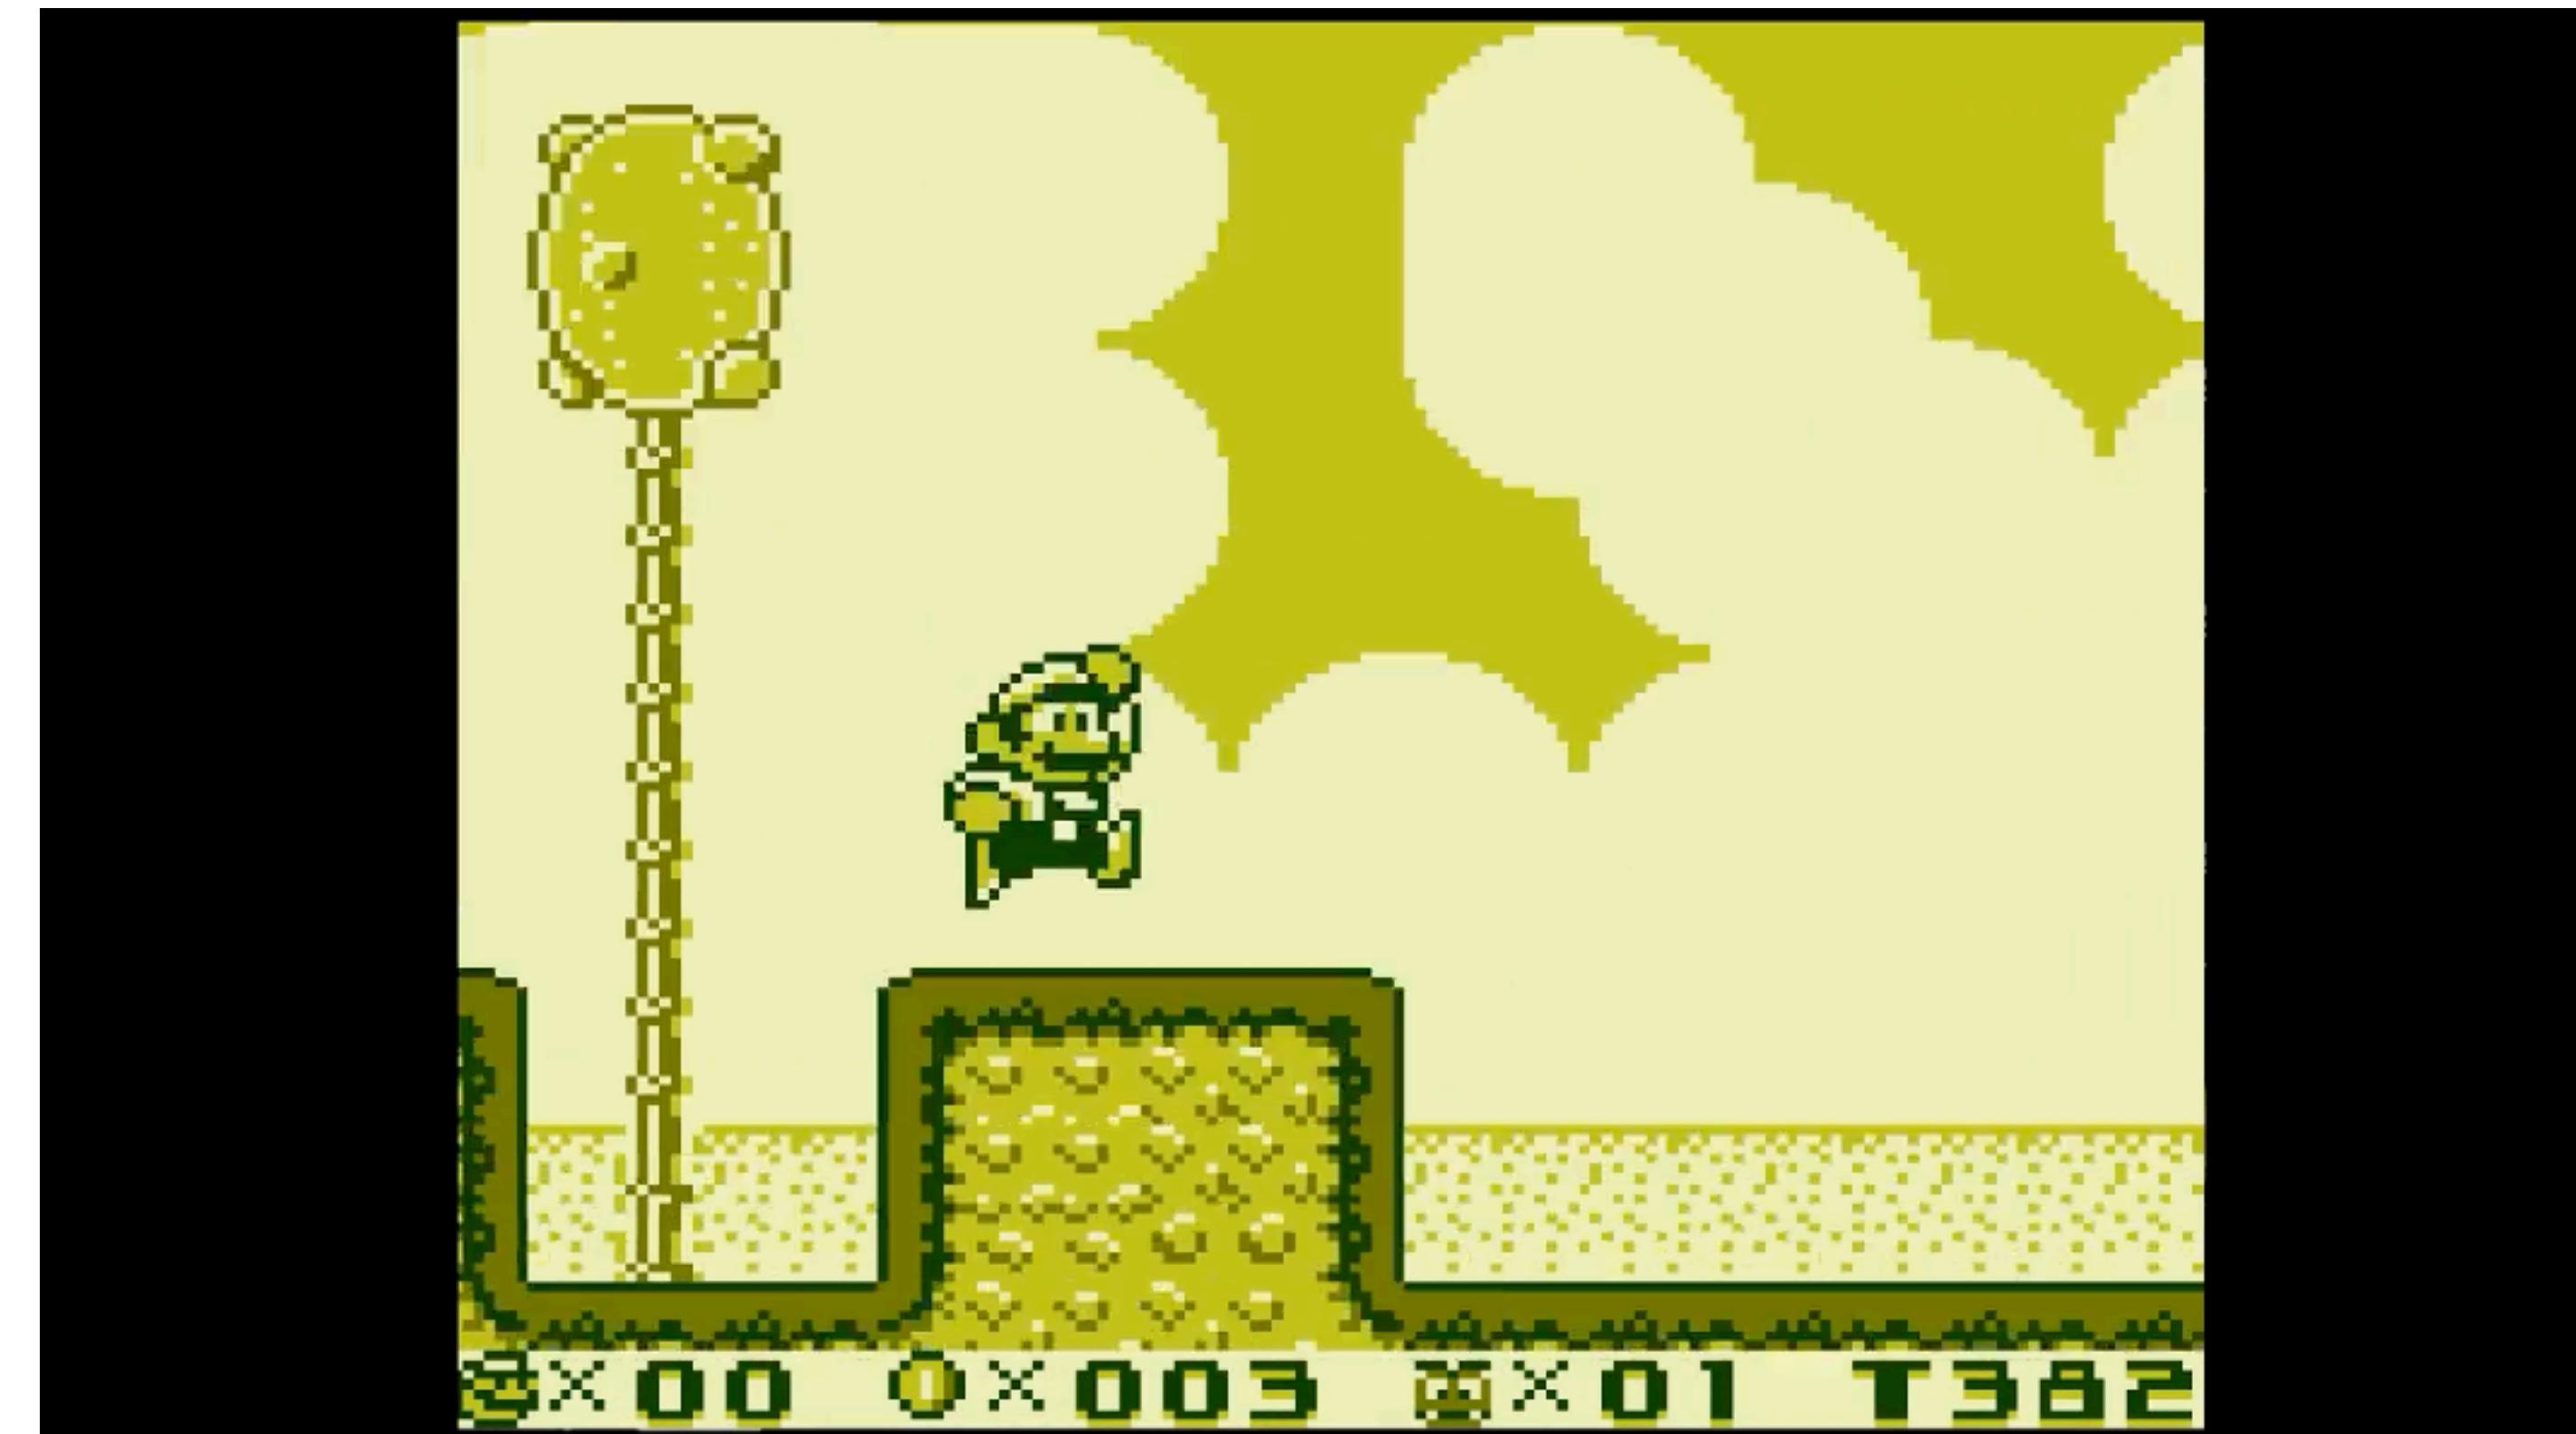


AARHUS  
UNIVERSITY  
DEPARTMENT OF COMPUTER SCIENCE

**Sidetrack:  
Can you do something cool with it?**

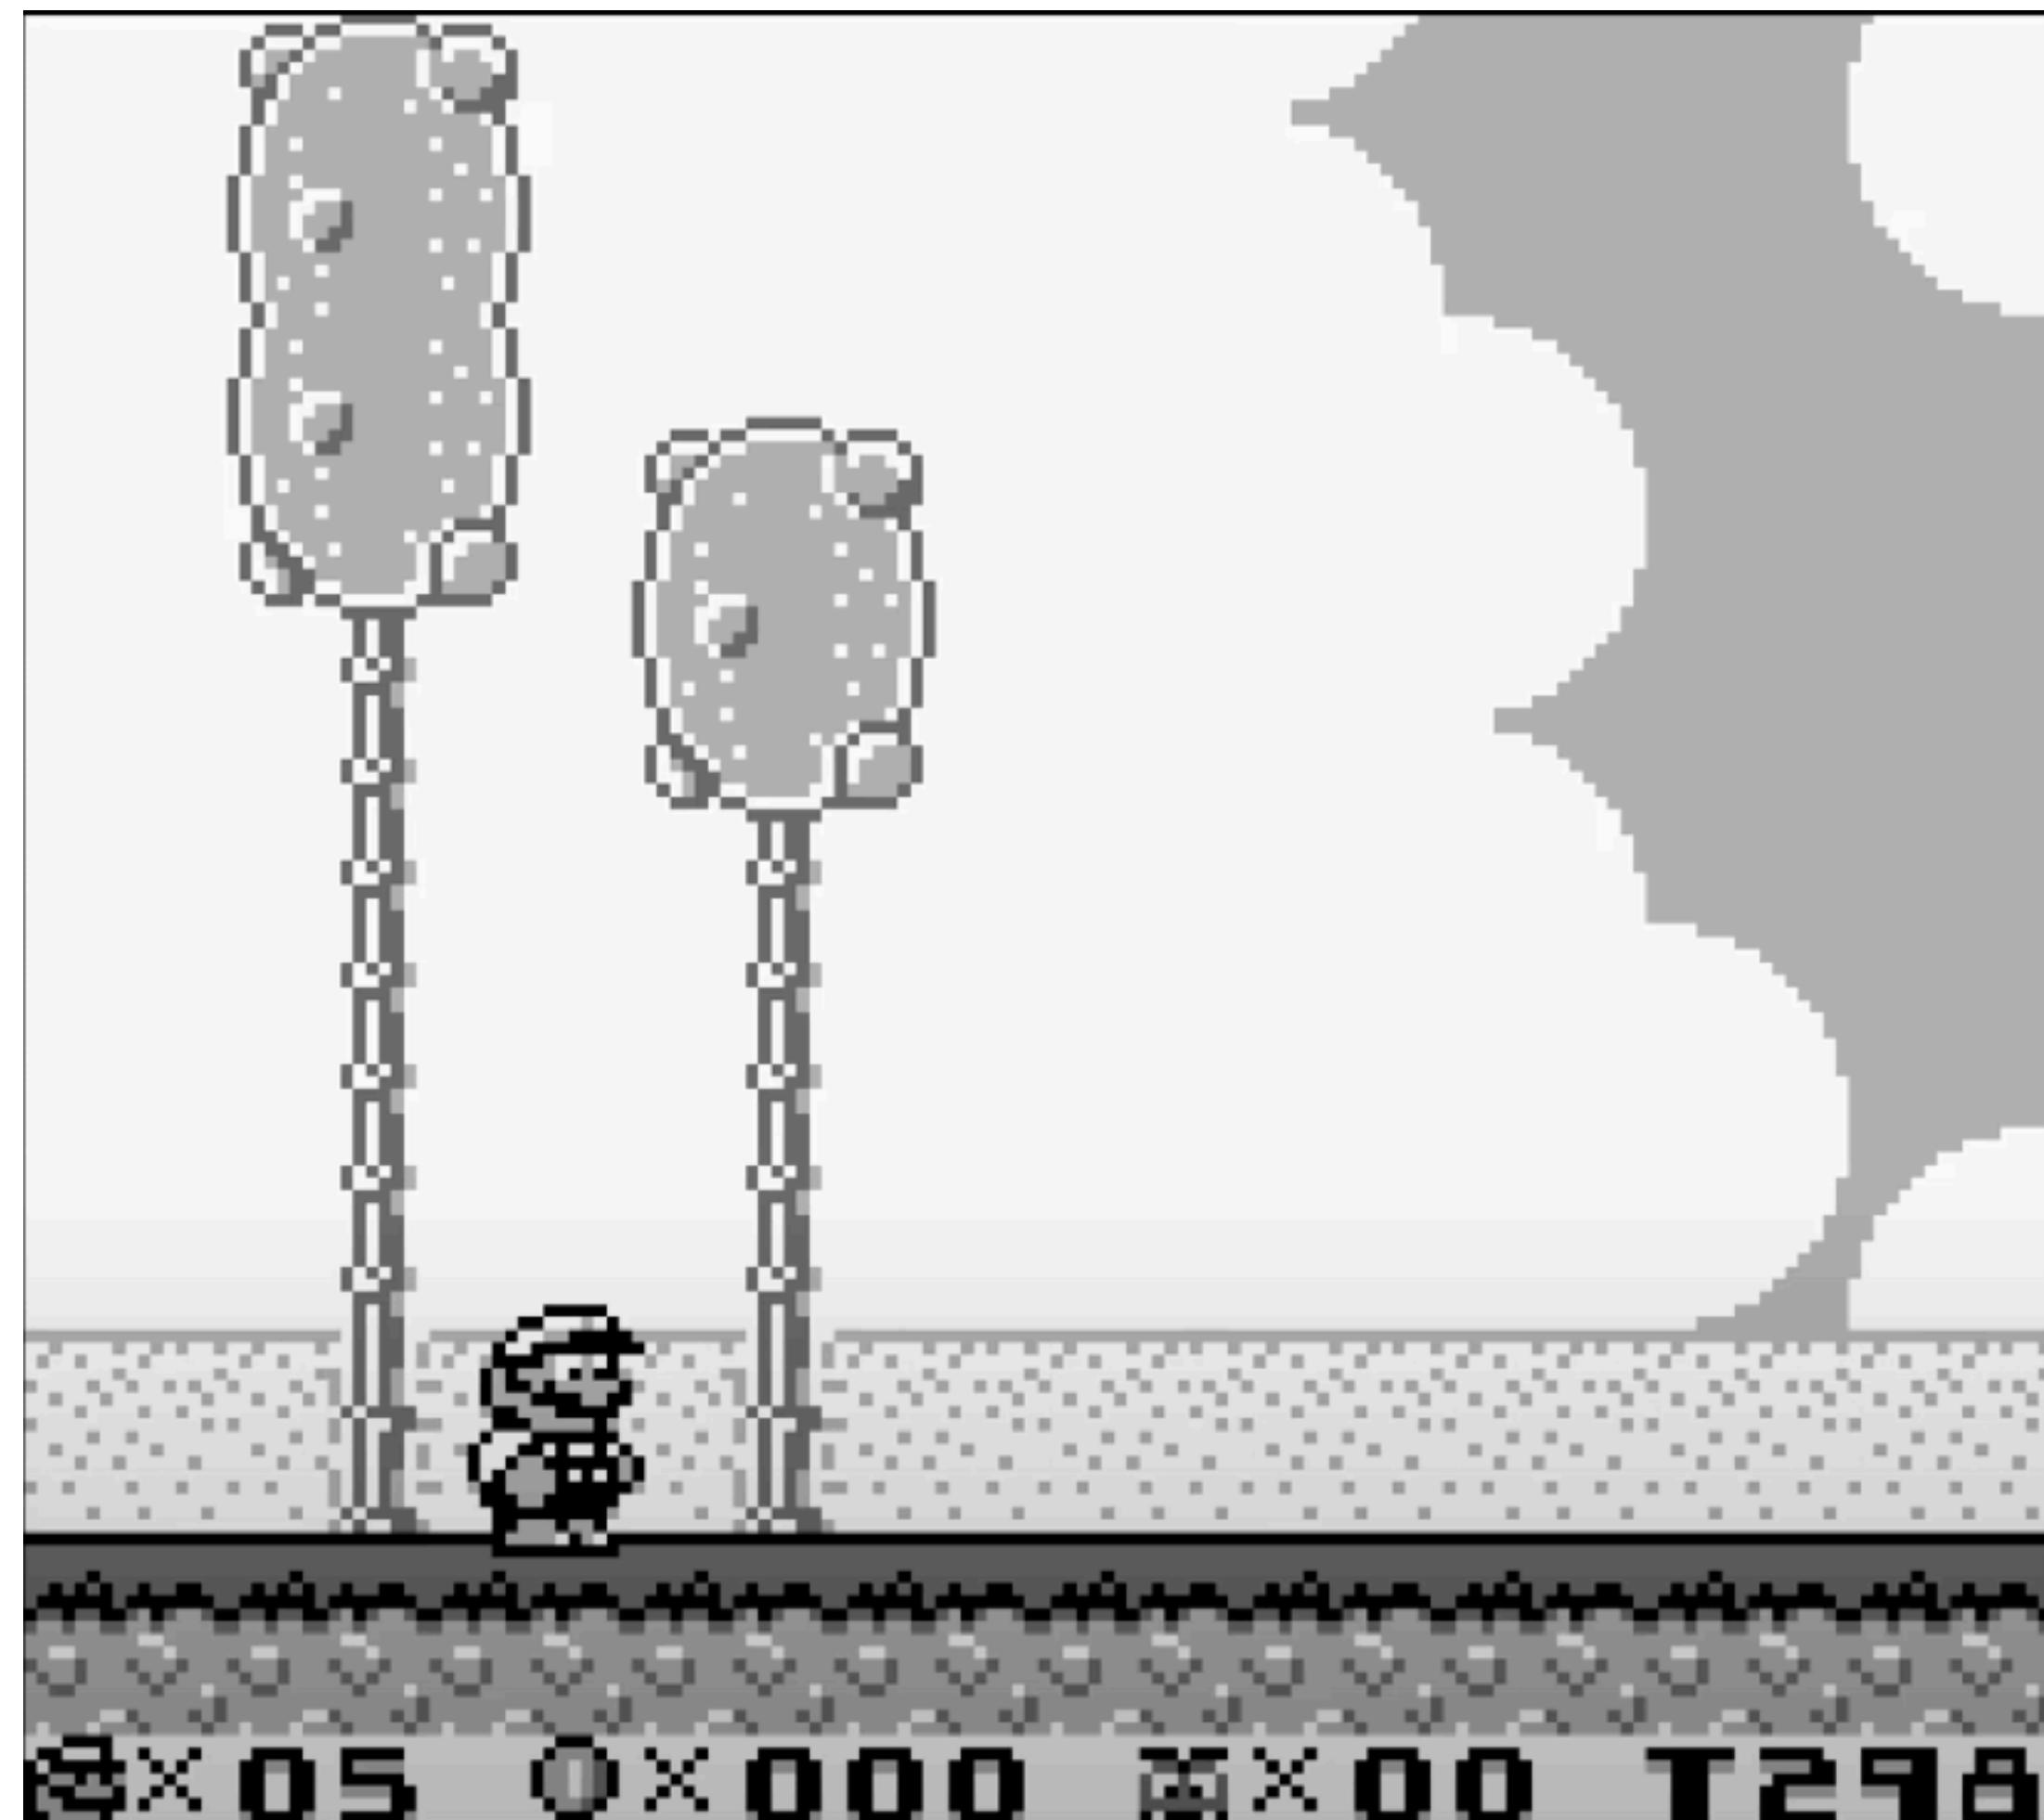
# Case study: Retro Gaming / Gameboy

- Context:
  - Implement a 2D game using only 2D graphics (images)



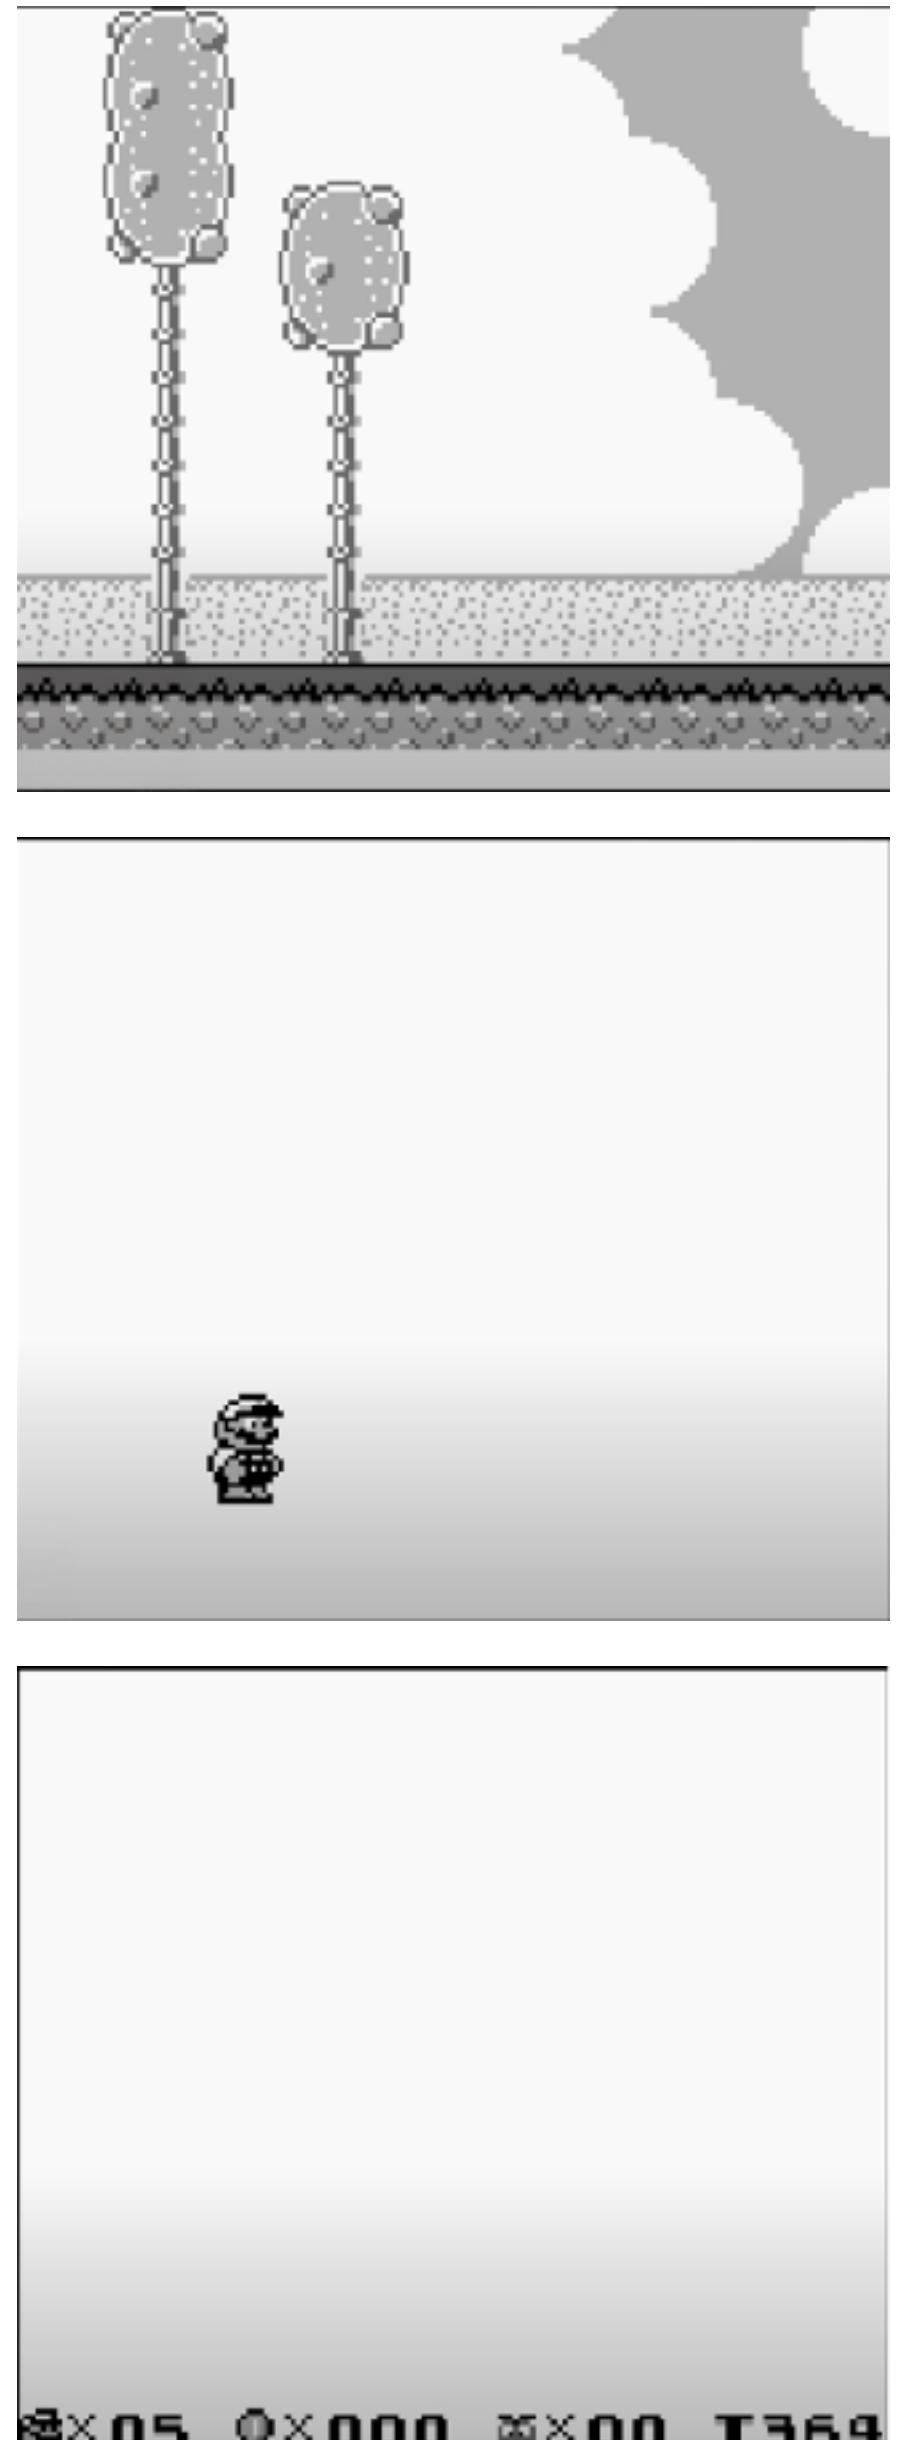
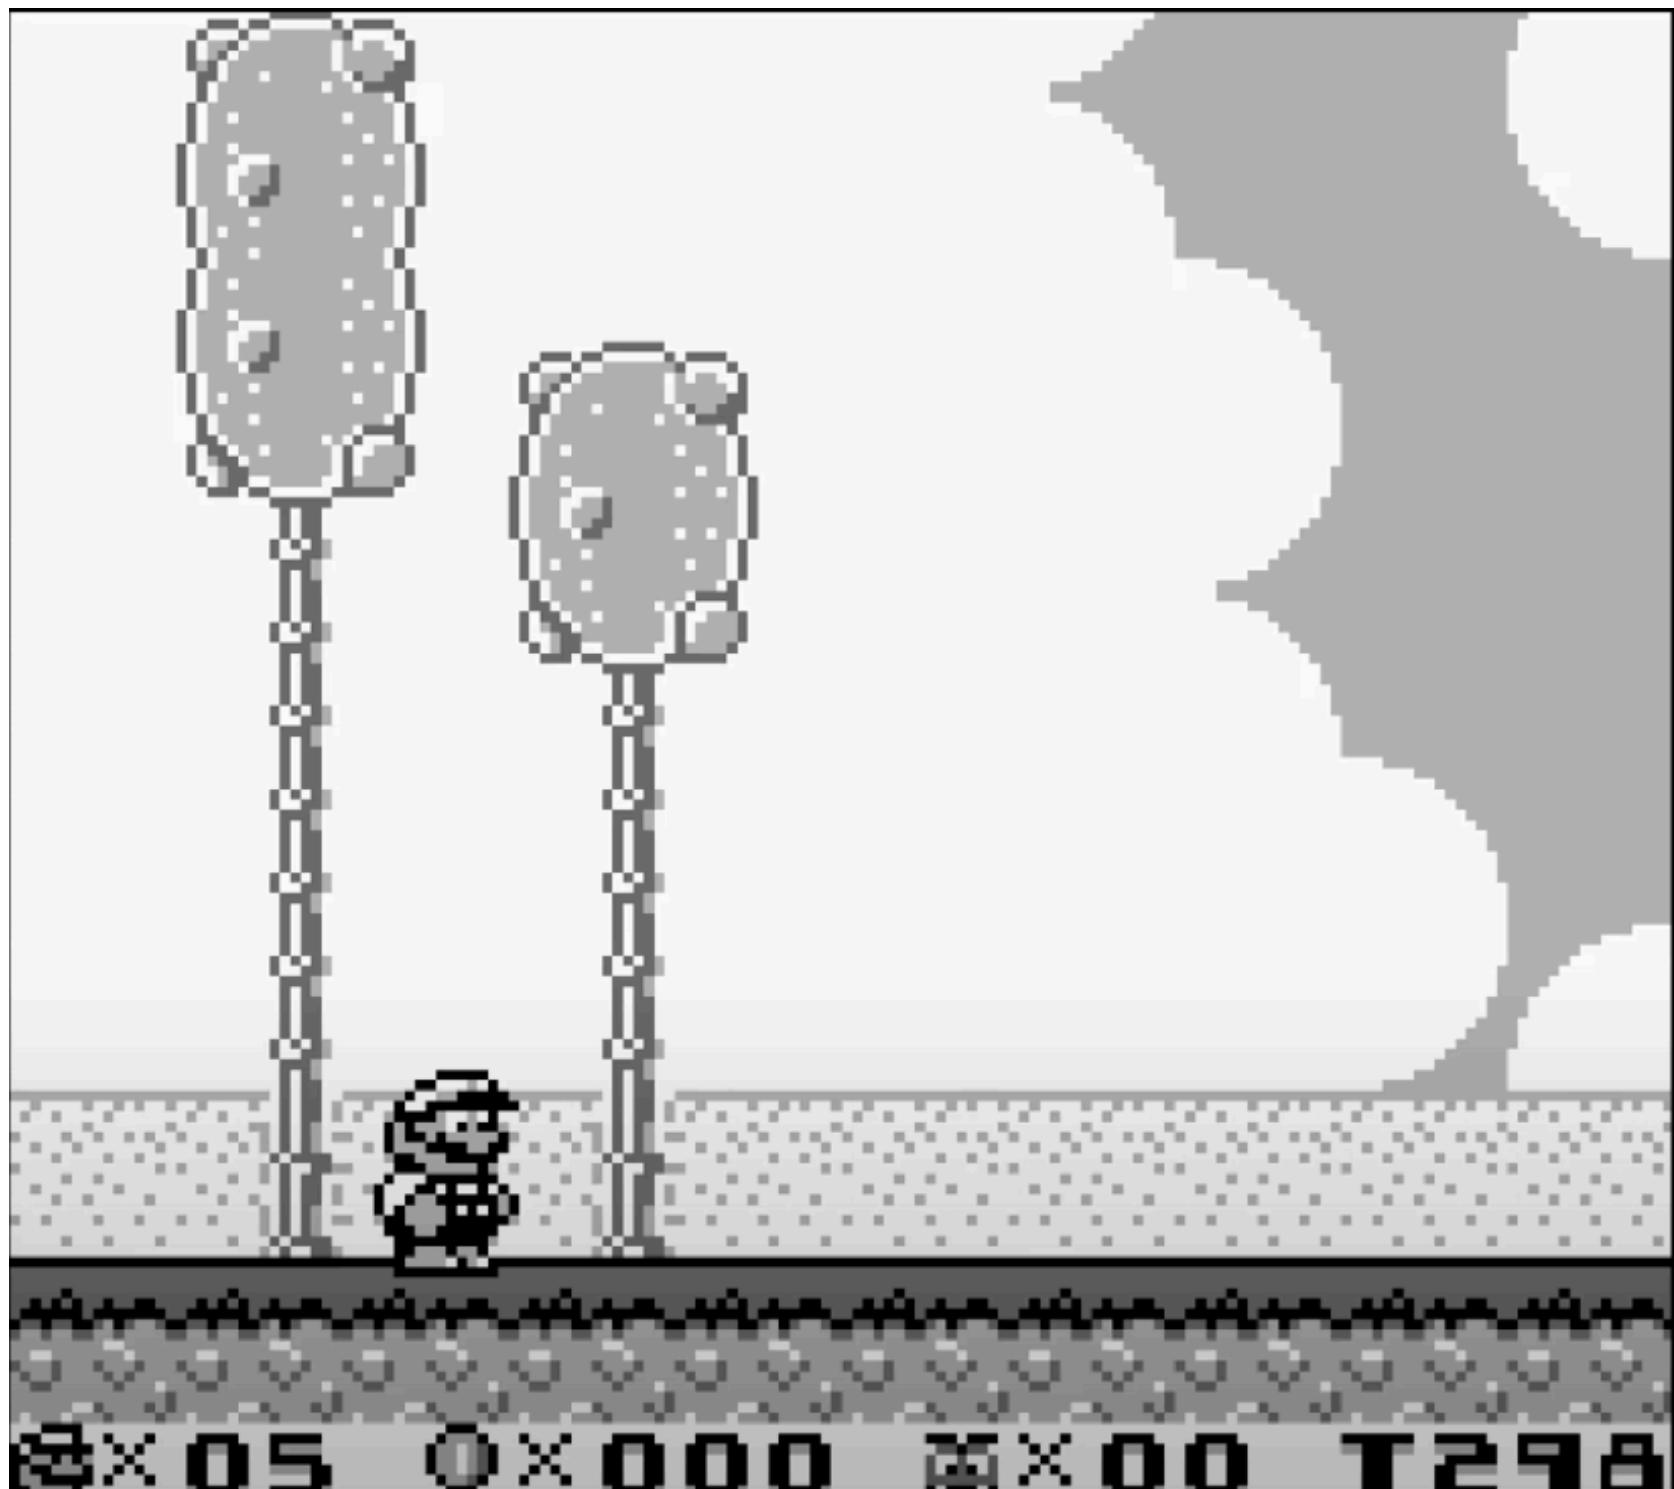
# Case study: Retro Gaming / Gameboy

- Context:
  - No 3D graphics
  - Slow hardware
  - Grey scale (2 bits, 4 shades)
  - 160x144 pixel LCD screen
  - Tile-based rendering (images are constructed from small 8x8 pixel tiles, total 20x18 tiles)
  - Three layers: background, window, and objects (sprites)



# Case study: Retro Gaming / Gameboy

- Context:
  - No 3D graphics
  - Slow hardware
  - Grey scale (2 bits, 4 shades)
  - 160x144 pixel LCD screen
  - Tile-based rendering (images are constructed from small 8x8 pixel tiles, total 20x18 tiles)
  - Three layers: background, window, and objects (sprites)



# Case study: Retro Gaming / Gameboy

- Context:
  - No 3D graphics
  - Slow hardware
  - Grey scale (2 bits, 4 shades)
  - 160x144 pixel LCD screen
  - Tile-based rendering (images are constructed from small 8x8 pixel tiles, total 20x18 tiles)
  - Three layers: background, window, and objects (sprites)



# Case study: Retro Gaming / Gameboy

- Context:
  - No 3D graphics
  - Slow hardware
  - Grey scale (2 bits, 4 shades)
  - 160x144 pixel LCD screen
  - Tile-based rendering (images are constructed from small 8x8 pixel tiles, total 20x18 tiles)
  - Three layers: background, window, and objects (sprites)



# **What do we miss? 2D Transformations / Animations**

**Next time :)**

**The end!**