

Visual Computing I:

Interactive Computer Graphics and Vision



Digital Image Processing (Manipulation and Analysis)

Stefanie Zollmann and Tobias Langlotz

Last time..

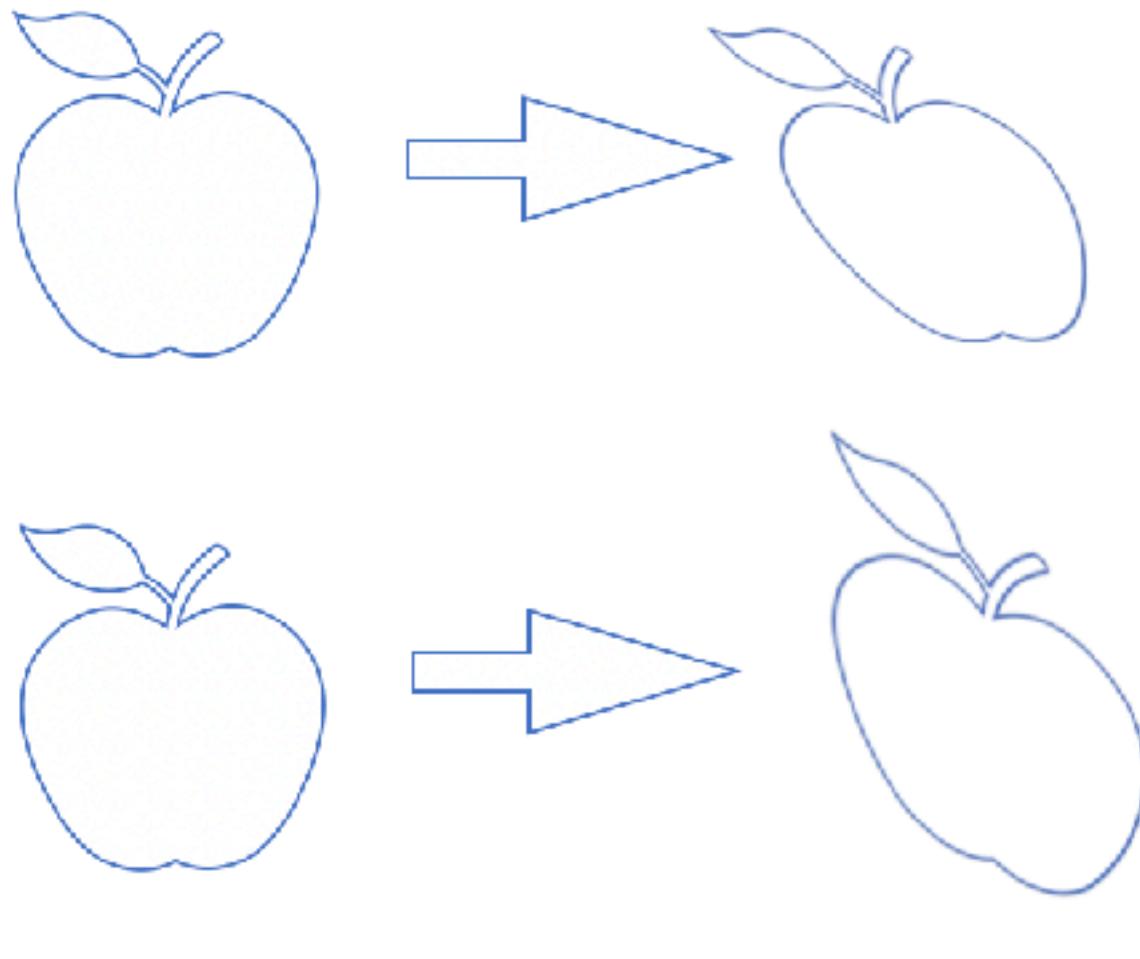
Homogenous Transforms - Shearing

- Distorts the shape by shifting one axis relative to the other
- Example: Turning a rectangle into a parallelogram
- Shearing by a factor of sh_x in x direction:

$$\begin{bmatrix} 1 & sh_x & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- Shearing by a factor of sh_y in your direction:

$$\begin{bmatrix} 1 & 0 & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



2D Animations

- Animation = applying transformations over time
- Common techniques:
 - Translation over time -> moving an object
 - Scaling over time -> growing/shrinking
 - Rotation over time -> spinning
- Achieved by updating transformation parameters frame by frame



2D Image Manipulation

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Image Luminance

Luminance

- Measures perceived “gray-level” of pixel
 - Not so good:
 - $(R+G+B) / 3$
 - Our example: $L = 120$



Luminance

- Measures perceived “gray-level” of pixel
 - Not so good:
 - $(R+G+B) / 3$
 - Our example: $L = 120$
 - Better:
 - $L = 0.30*R + 0.59*G + 0.11*B$
 - Luminosity method (human perception)
 - Our example: $L = 106$



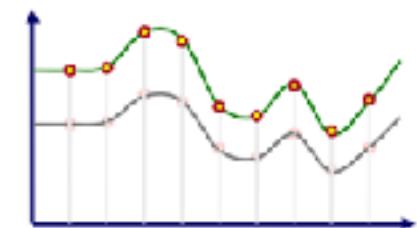
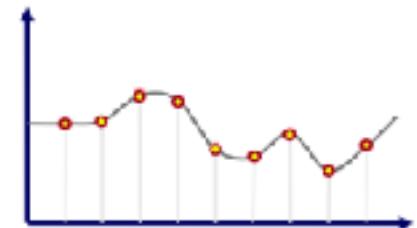
Luminance

- Measures perceived “gray-level” of pixel
 - Not so good:
 - $(R+G+B) / 3$
 - Our example: $L = 120$
 - Better:
 - $L = 0.30*R + 0.59*G + 0.11*B$
 - Luminosity method (human perception)
 - Our example: $L = 106$
- What must be done to the RGB values to make this image brighter?



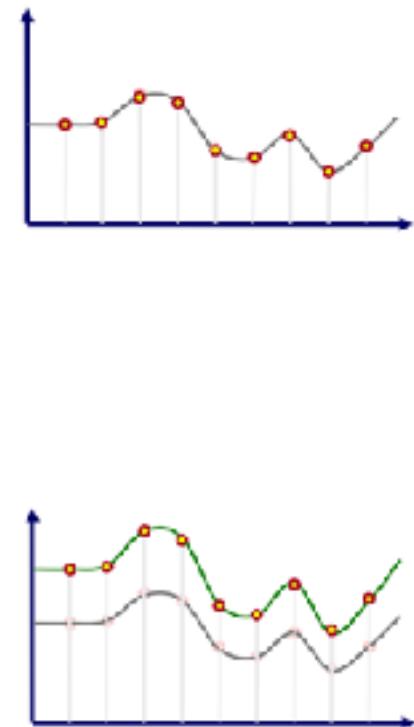
Luminance

- Change luminance:
- Method 1: Convert to HSL, scale L, convert back
 - (more on this shortly...)
- Method 2: Scale R, G, and B directly
 - Multiply each of red, green, and blue by a factor
 - Must clamp to [0..1] ... always
 - ([0..1] in floating point but often [0,255] for fixed point)



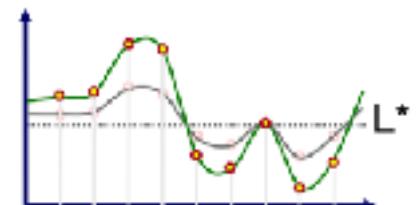
Luminance

- Change luminance:
- Method 1: Convert to HSL, scale L, convert back
 - (more on this shortly...)
- Method 2: Scale R, G, and B directly
 - Multiply each of red, green, and blue by a factor
 - Must clamp to [0..1] ... always
 - ([0..1] in floating point but often [0,255] for fixed point)



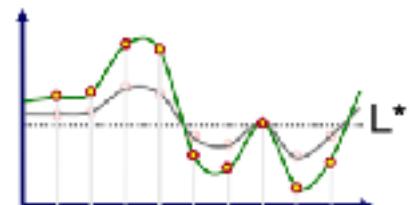
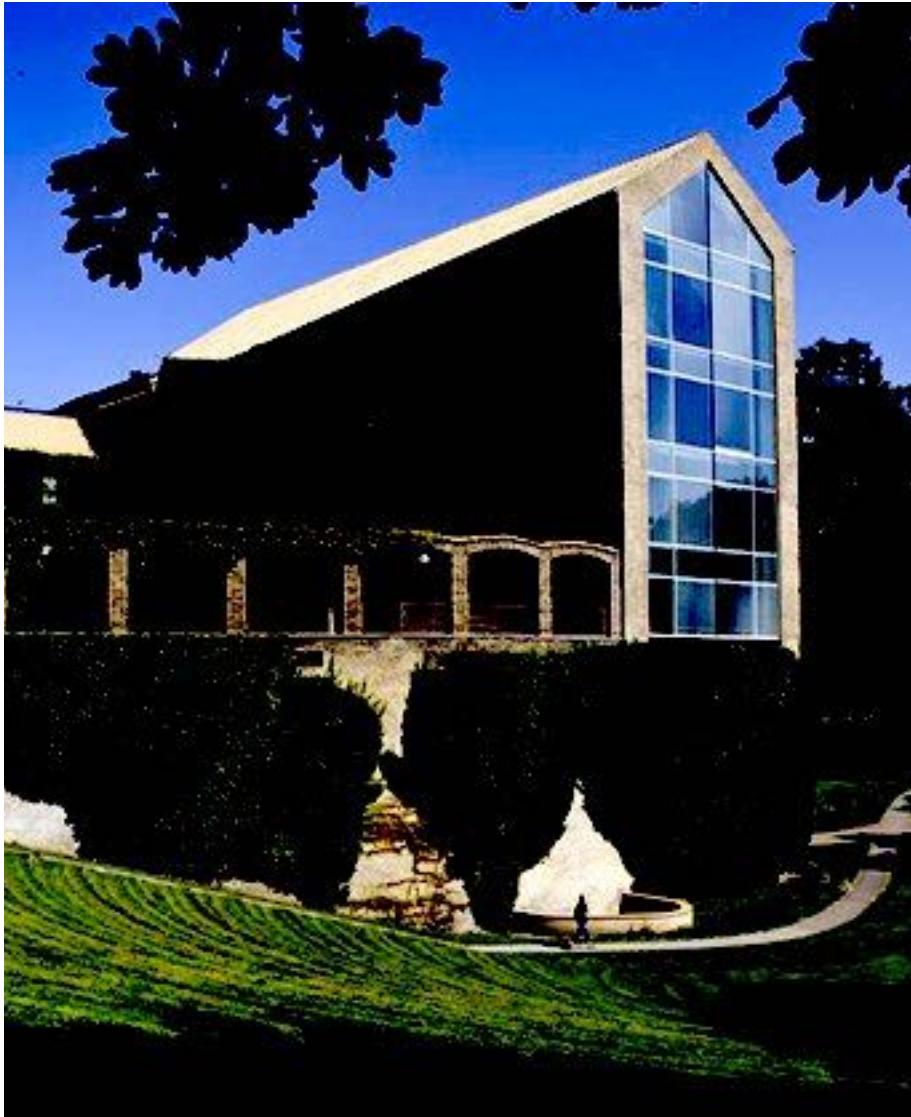
Contrast

- Contrast is the difference in luminance or colour that makes an object (or its representation in an image or display) stand out from its background
- Simpler: Contrast is the difference in luminance or colour for each pixel from the overall mean
 - Compute mean luminance L^* over whole image
 - Scale deviation from L^* for each pixel



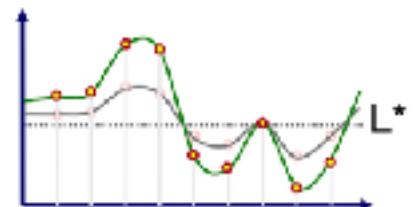
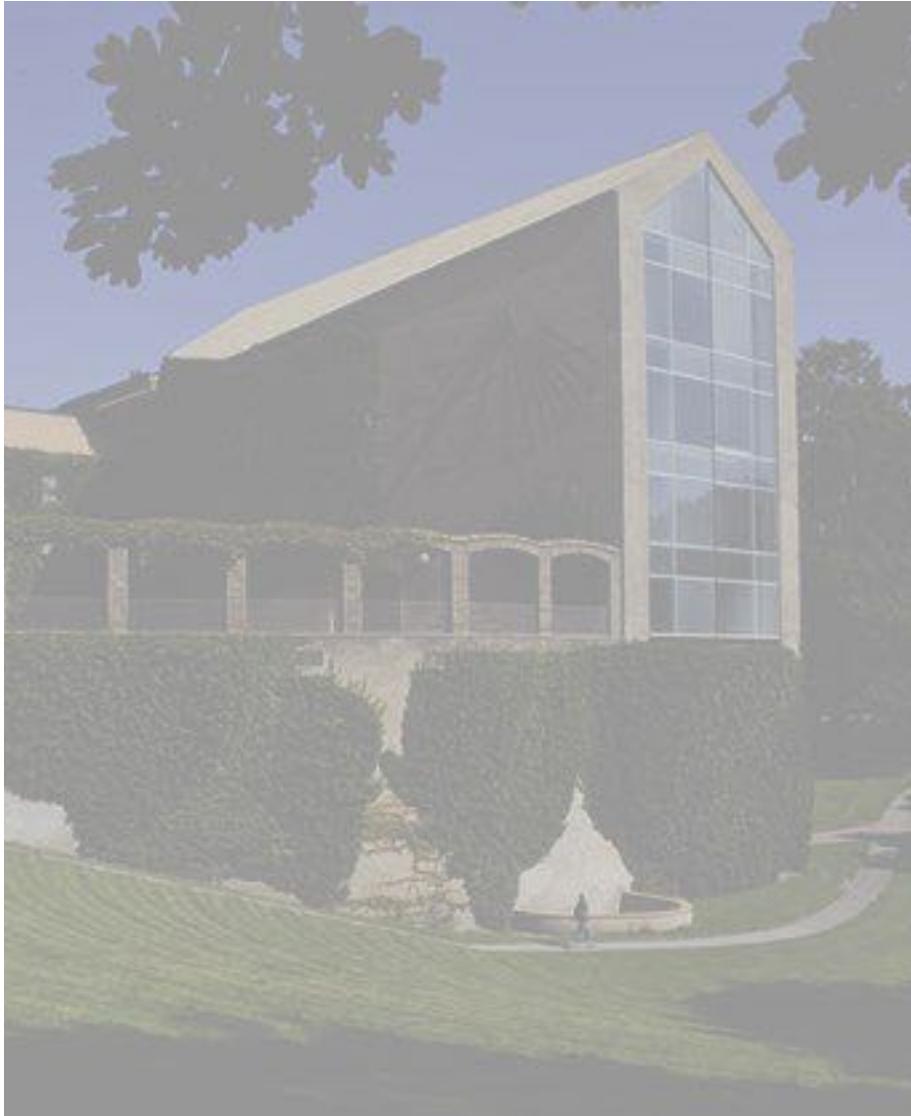
Contrast

- Contrast is the difference in luminance or colour that makes an object (or its representation in an image or display)
- Simpler: Contrast is the difference in luminance or colour for each pixel from the overall mean
 - Compute mean luminance L^* over whole image
 - Scale deviation from L^* for each pixel



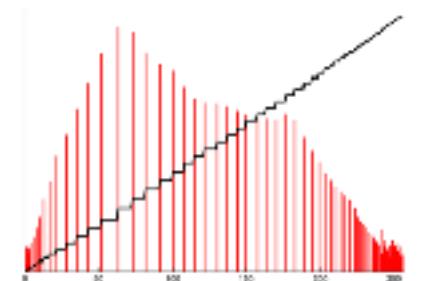
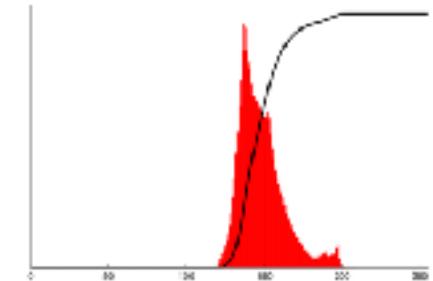
Contrast

- Contrast is the difference in luminance or colour that makes an object (or its representation in an image or display)
- Simpler: Contrast is the difference in luminance or colour for each pixel from the overall mean
 - Compute mean luminance L^* over whole image
 - Scale deviation from L^* for each pixel



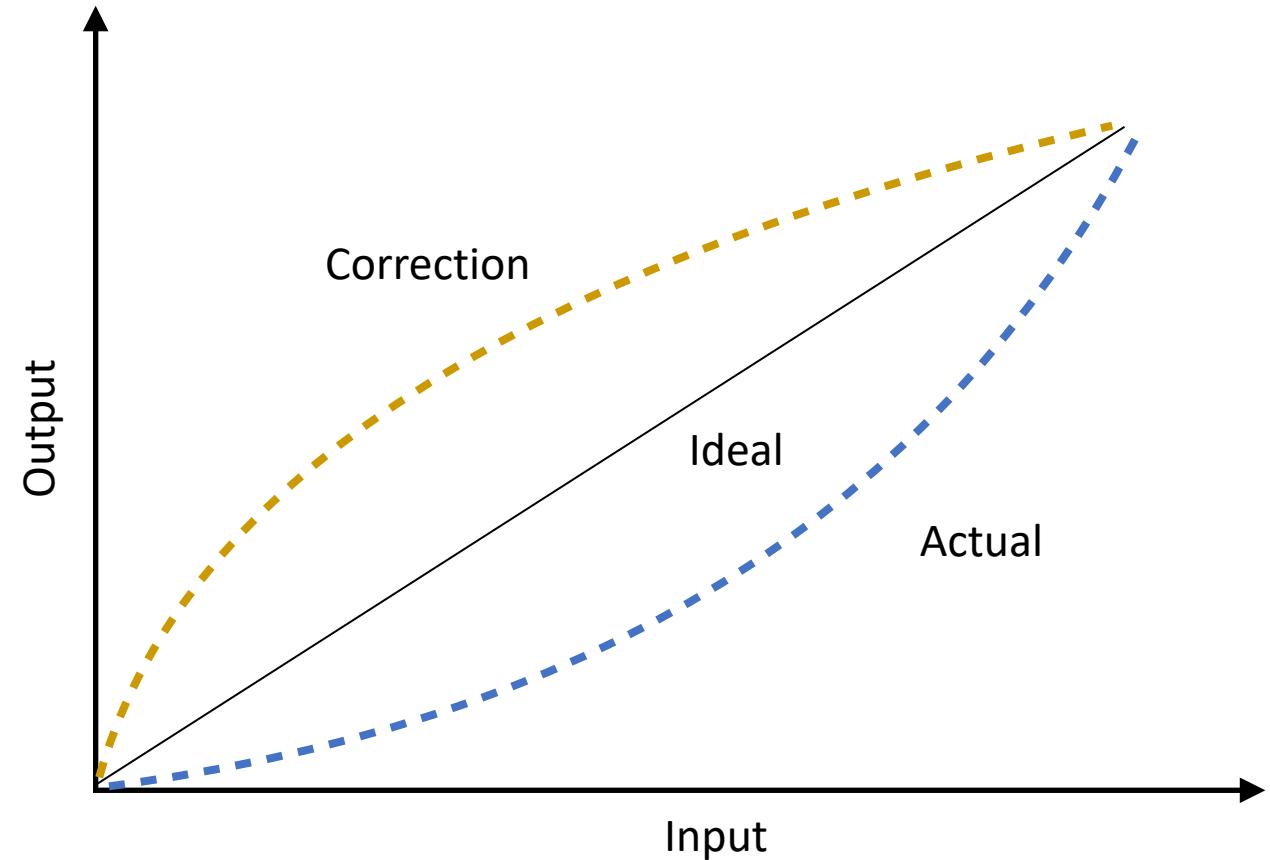
Histogram Equalization / Normalisation

- Different form of Contrast
- Change distribution of luminance values to cover full range [0-1] / [0,255]



Gamma Correction

- Apply non-linear function to account for the difference between brightness and perceived brightness of displays
- $I_{out} = I_{in}^\gamma$
- γ depends on camera and monitor



2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

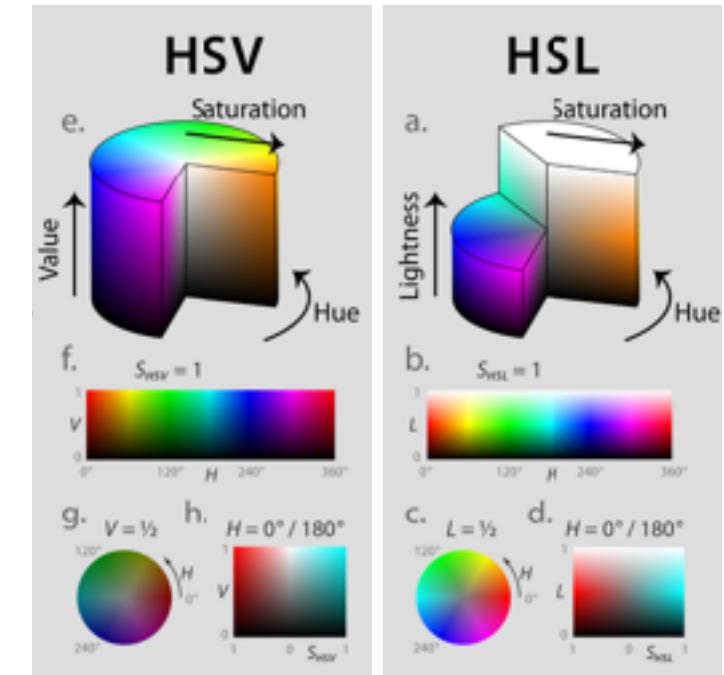
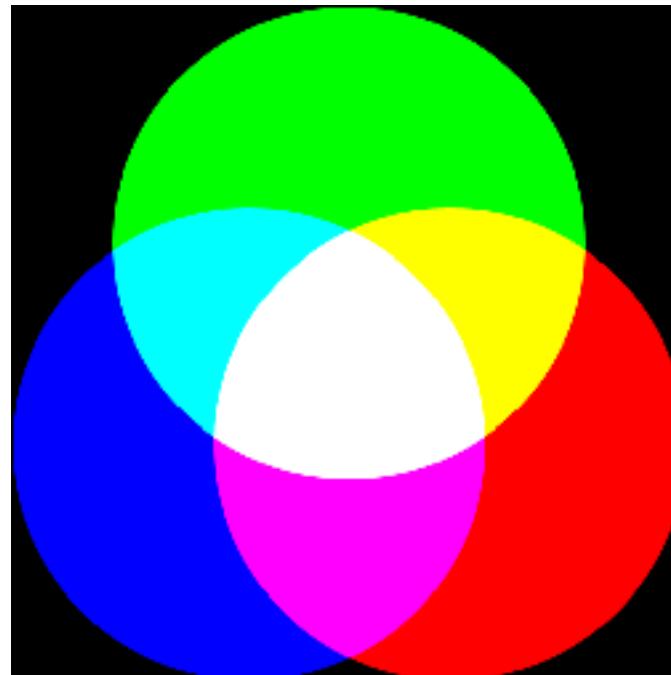
2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Color

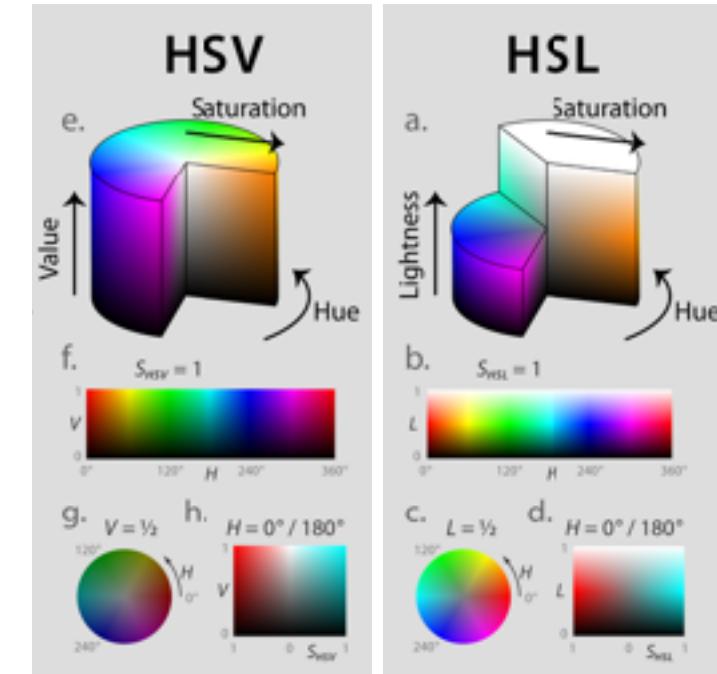
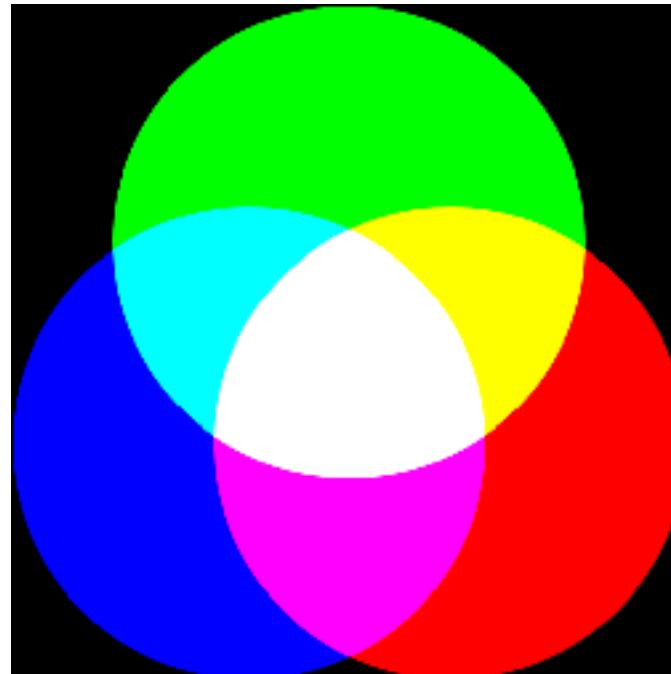
Colour Processing

- Make use of different colour models
- Colour models (previous lec.)
 - RGB
 - CMY(K)
 - HSV
 - HSL
 - La^*b^*
 - Etc. (XYZ, ..)



Colour Processing

- Make use of different colour models
- Colour models (previous lec.)
 - RGB
 - CMY(K)
 - HSV
 - HSL
 - La^*b^*
 - Etc. (XYZ, LMS..)



- Which one is good for conversion to grayscale? Saturation?

Greyscale

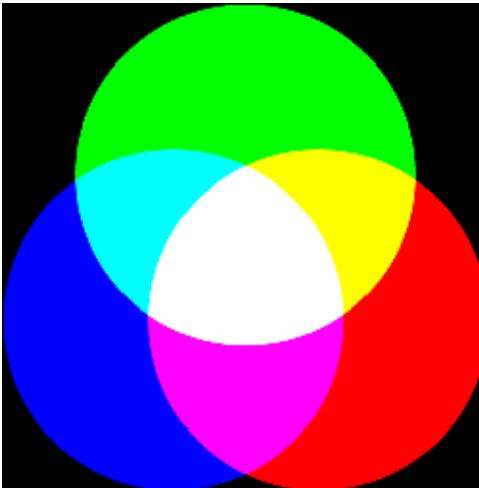


Original



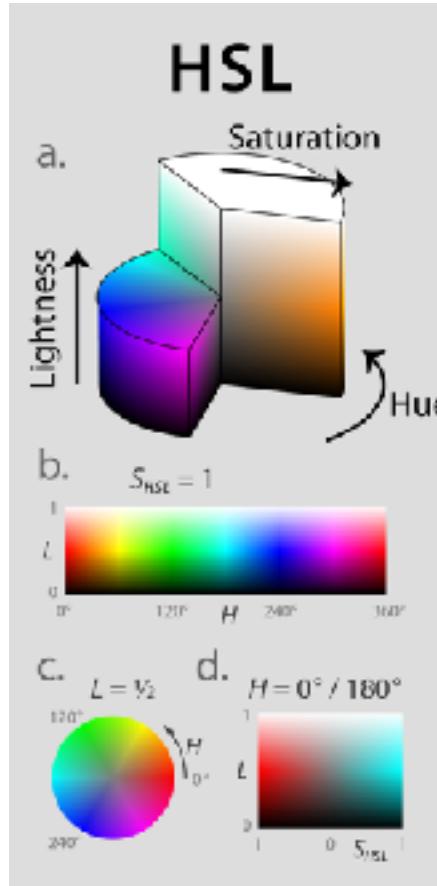
Greyscale
(“black&white” photo)

Greyscale



- Convert from colour to grey-levels
 - Method 1: (not so good): $R+G+B/3$

Greyscale



- Convert from colour to grey-levels
 - Method 2: Convert to HSL, set S=0, convert back to RGB
 - Method 3: Set RGB of every pixel to (L,L,L)

Saturation



Original



Adjusted Saturation



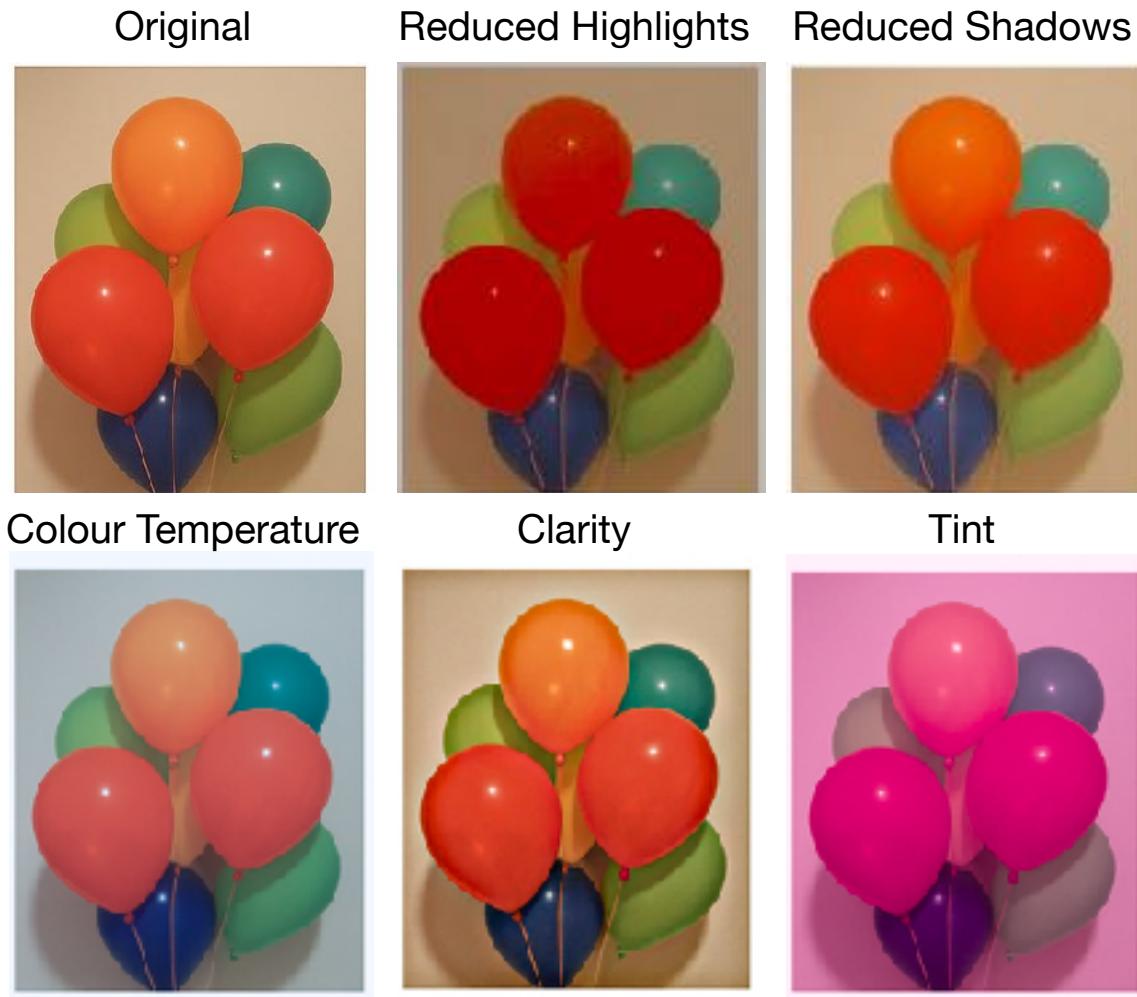
Saturation



- Increase/decrease color saturation of every pixel
 - Method 1: Convert to HSL, scale S, convert back
 - Method 2: $R' = L + \text{scale} * (R-L)$... same for G&B

Other

- Many other colour processing methods:
 - White balancing
 - Colour Temperature
 - Tint
 - Texture
 - Clarity
 - Dehazing
 - Highlights/Shadows
 - Levels/Histogram
- General rule: Think about the best colour models to operate in



2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

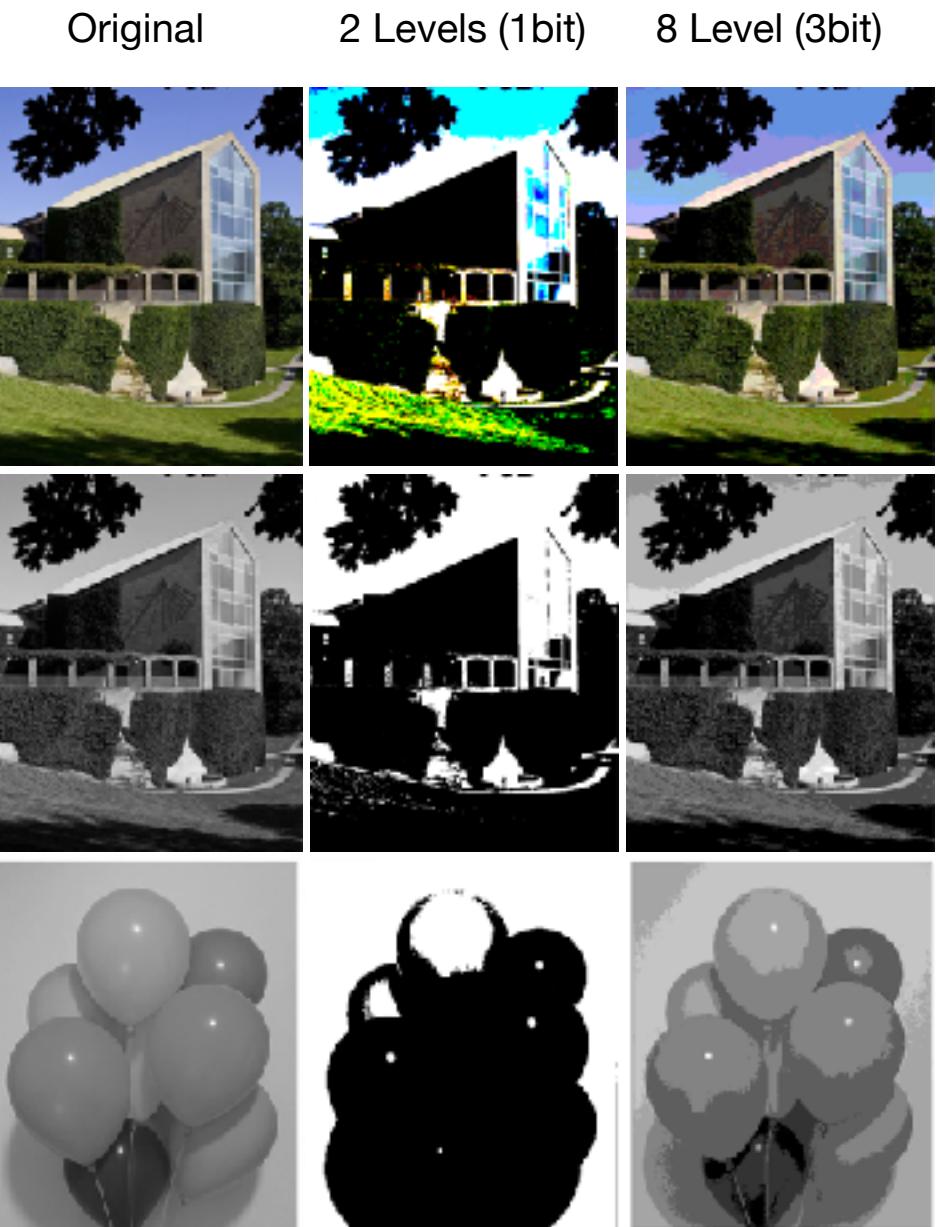
2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Dithering

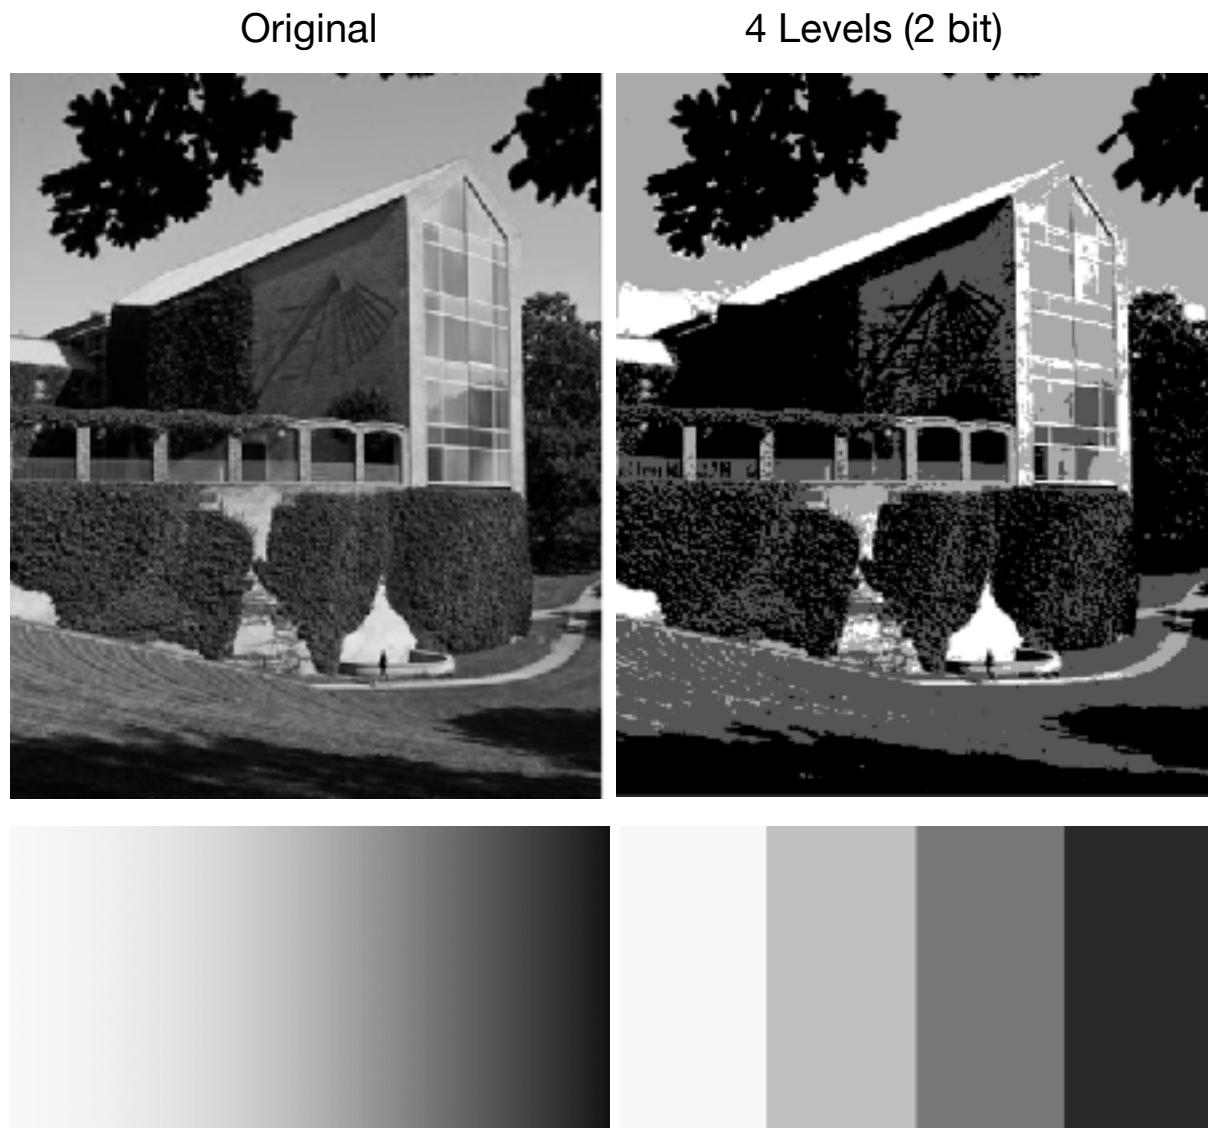
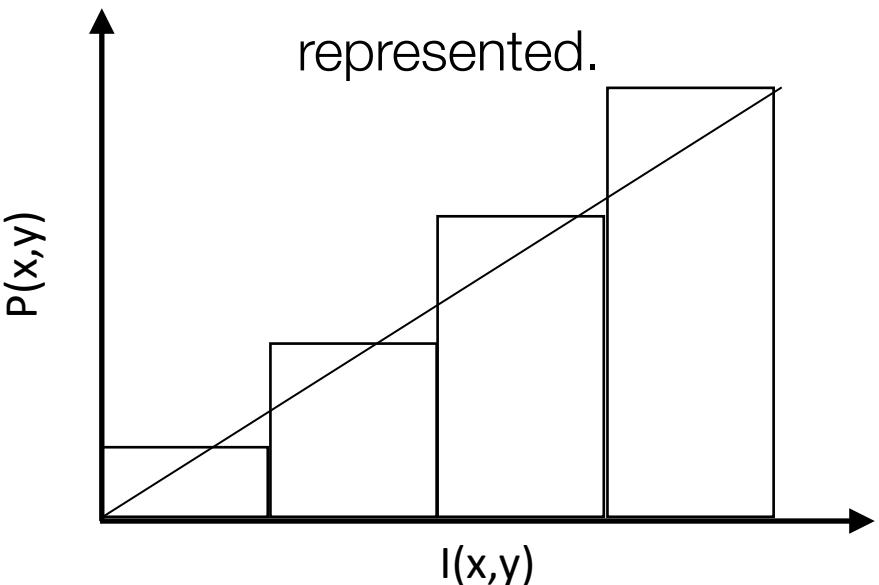
Quantization and Dithering

- Quantization:
 - Reduce intensity/colour resolution
 - Greyscale 8bit to 1bit
 - RGB 12 bit to 8 bit
 - RGB 8bit to 5/6 bits (RGB565)
 - Frame buffers have a limited number of bits per pixel
 - Physical devices have a limited dynamic range
 - Reduce image size



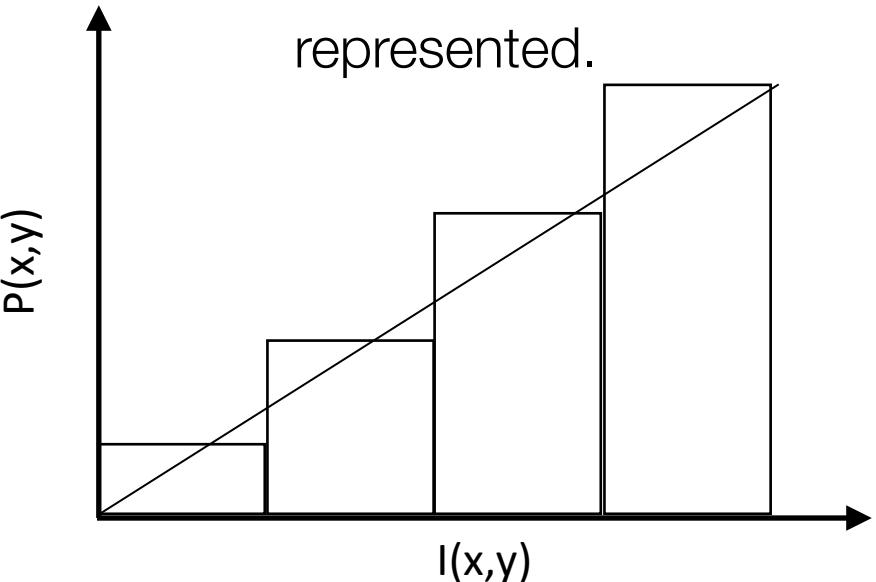
Quantization and Dithering

- Quantization:
 - $P(x, y) = \text{round}(I(x, y))$
 - Where $\text{round}()$ chooses nearest value that can be represented.

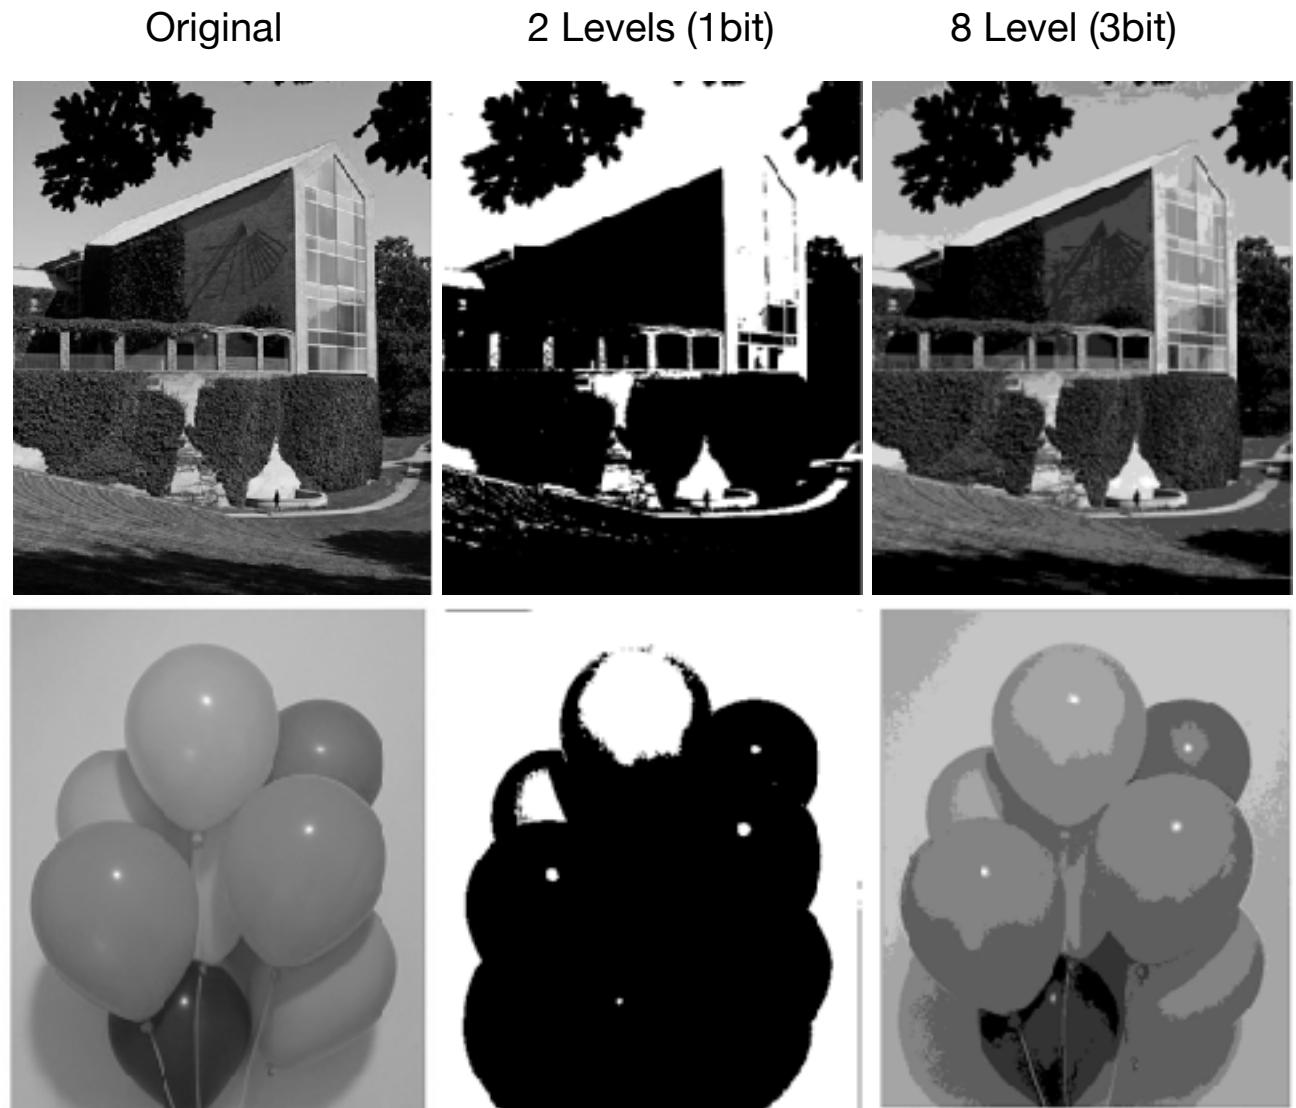


Quantization and Dithering

- Quantization:
 - $P(x, y) = \text{round}(I(x, y))$
 - Where $\text{round}()$ chooses nearest value that can be represented.



- Contouring or Banding



Quantization and Dithering

- Dither use noise or patterns to randomise or reduce quantisation error
- Reducing large-scale patterns (contouring or banding)

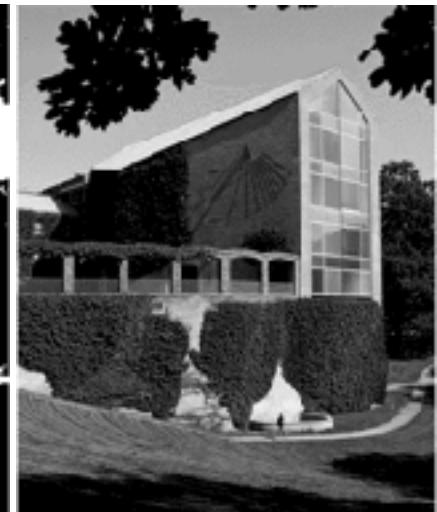
Original



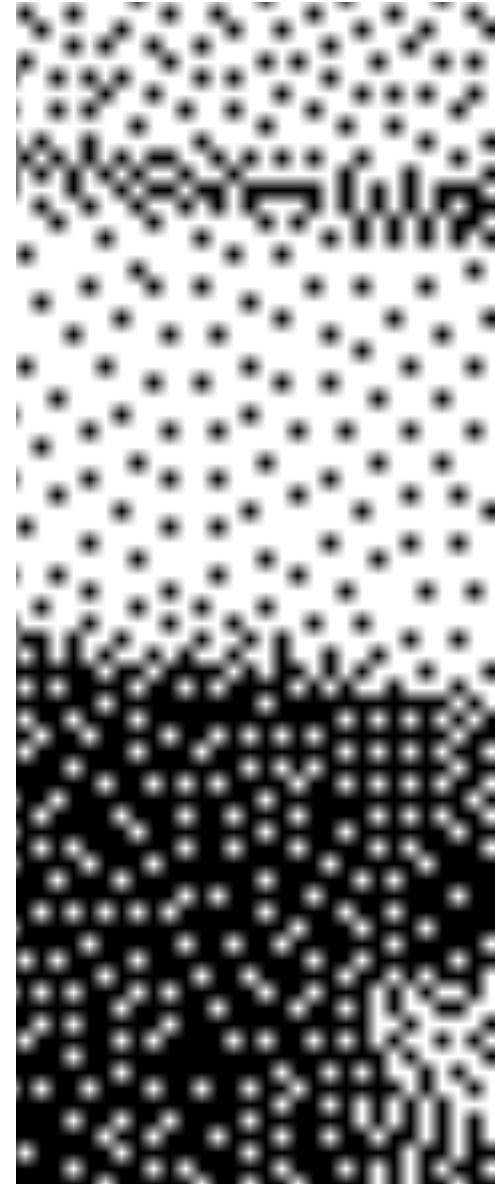
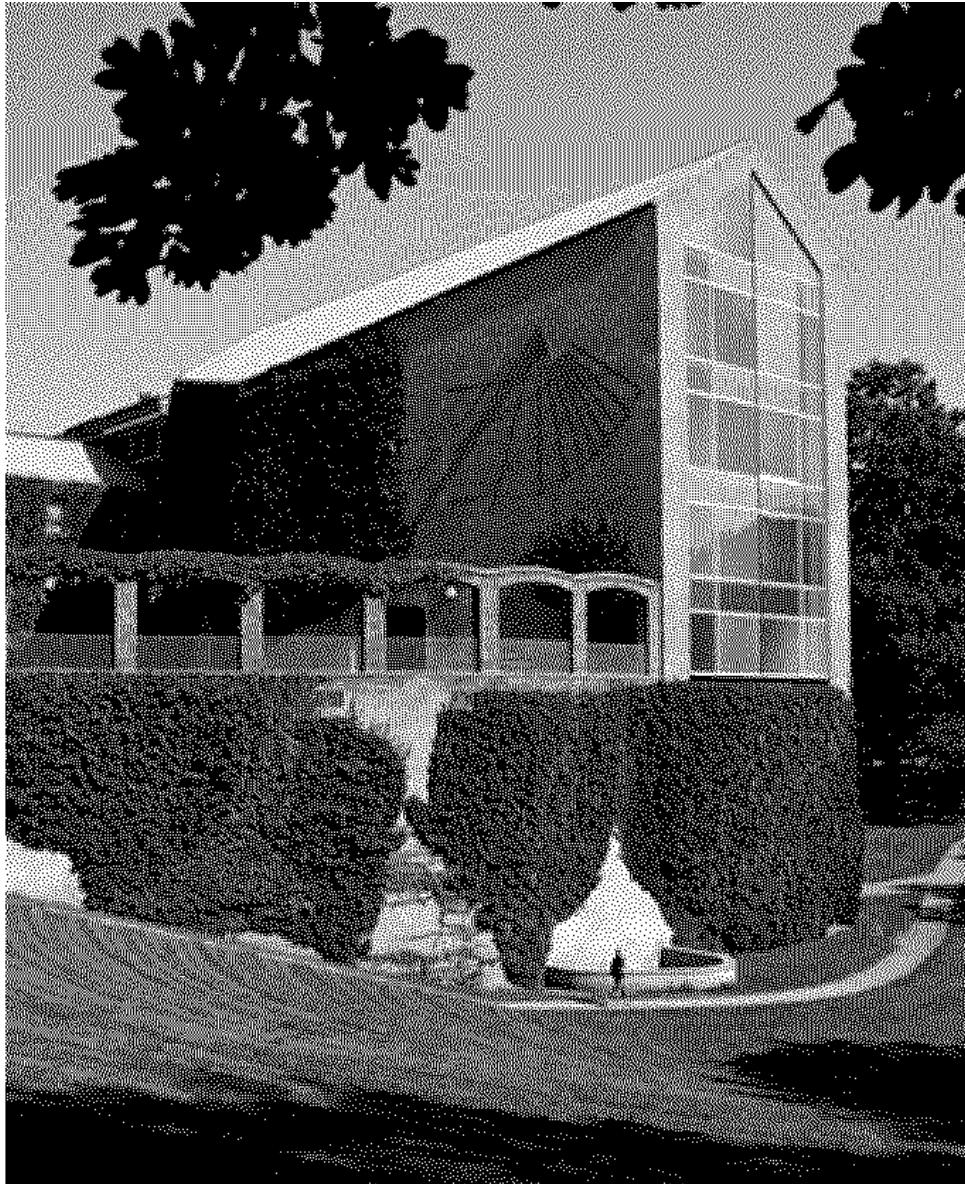
2 Levels (1bit)



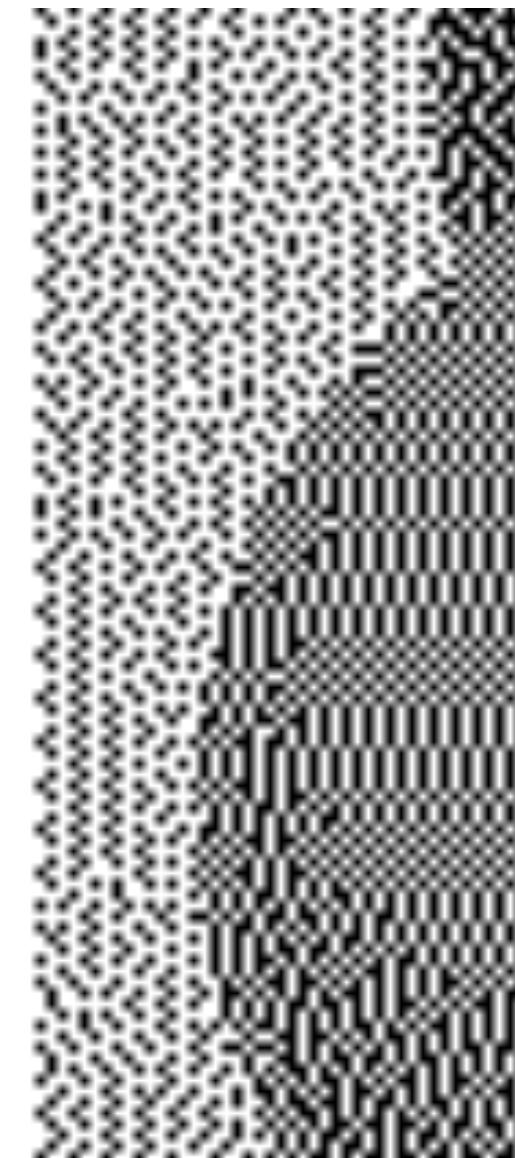
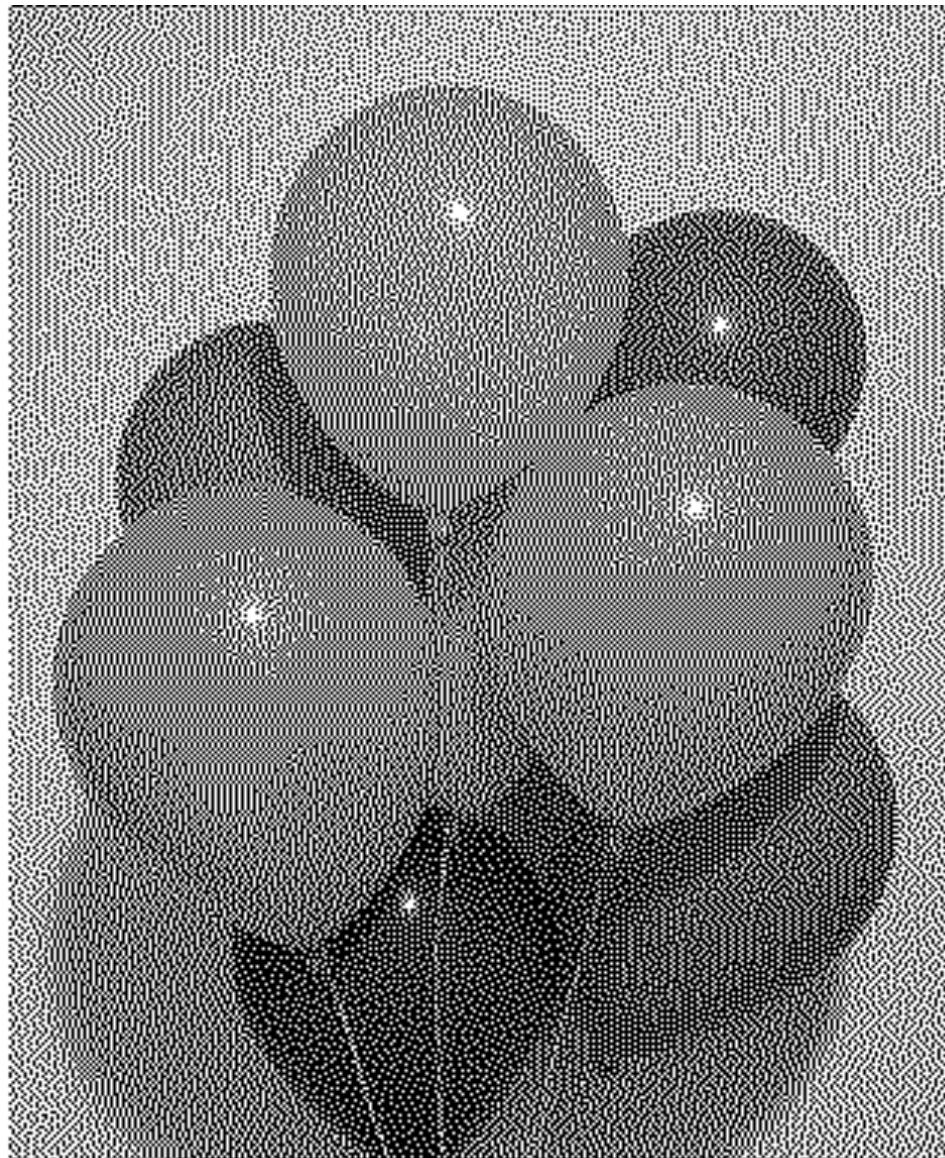
2 Levels (1bit with Dither)



Quantization and Dithering

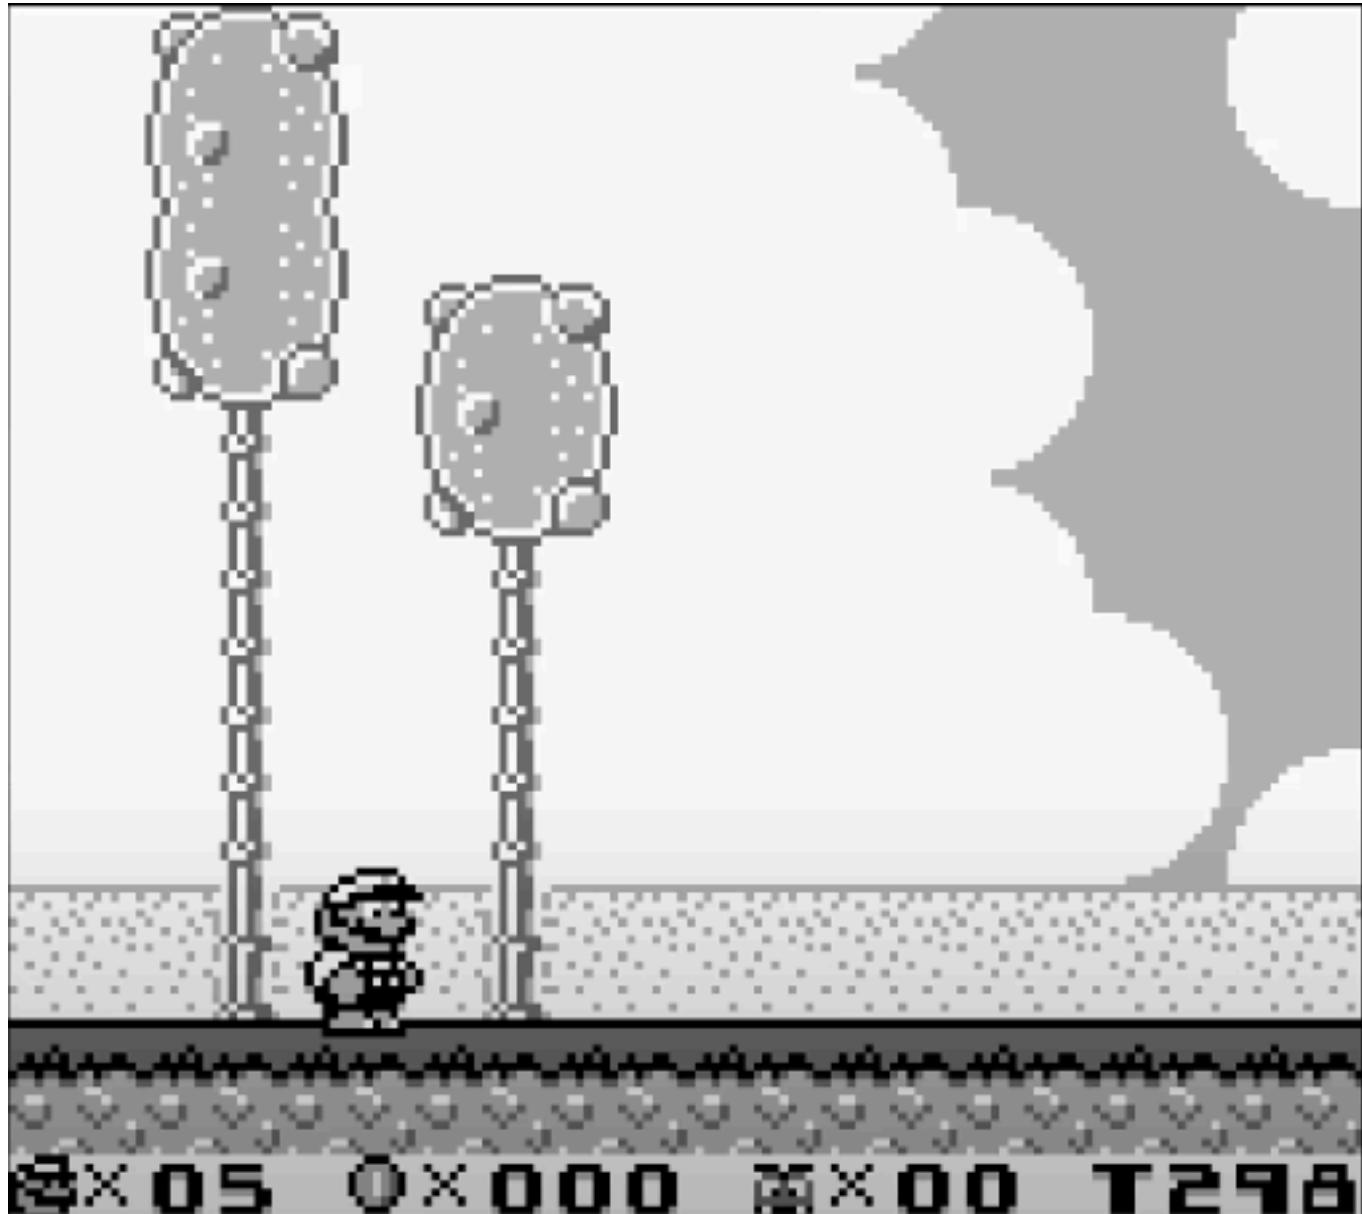
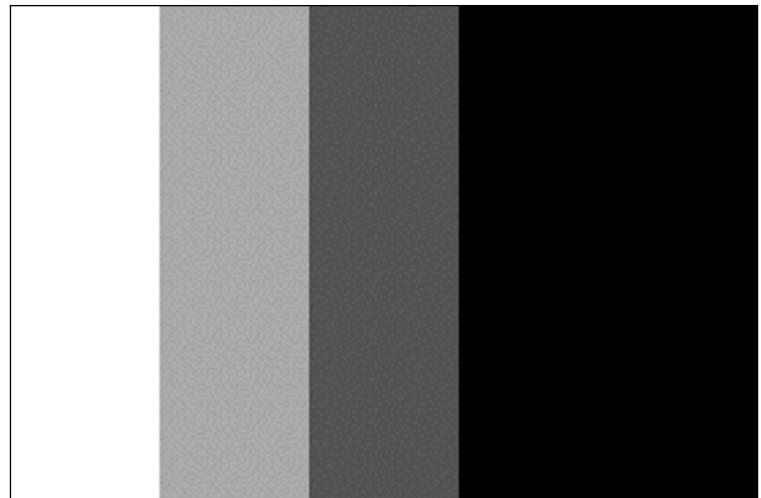


Quantization and Dithering



Quantization and Dithering

- From the 2nd Lecture
 - Super Mario World (Gamboy)
 - No 3D graphics
 - Grey scale (2 bits, 4 shades)



Quantization and Dithering

- Dither uses noise or patterns to randomise or reduce quantisation error
- Reducing large-scale patterns (contouring or banding)
- Distribute errors among pixels
 - Exploit perceptual spatial integration
 - Display greater range of perceptible intensities



Quantization and Dithering

- Random Dither:

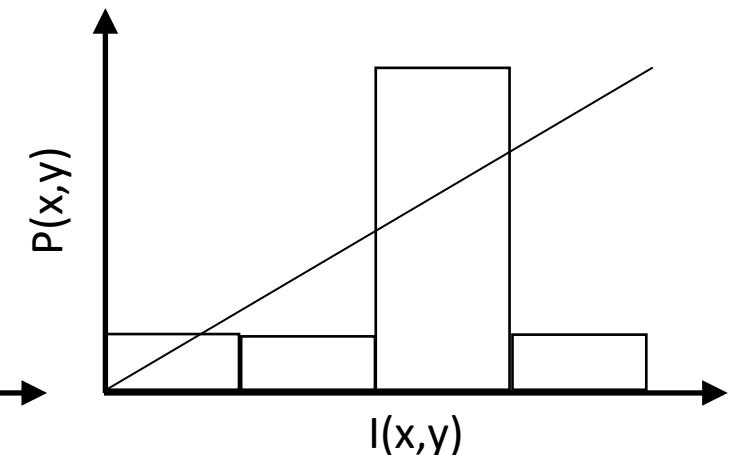
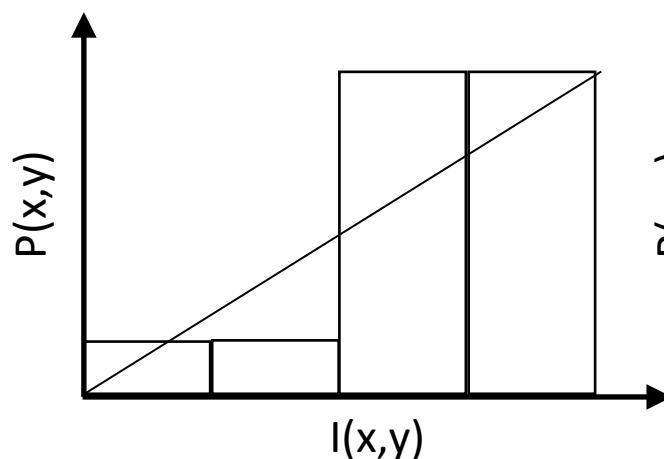
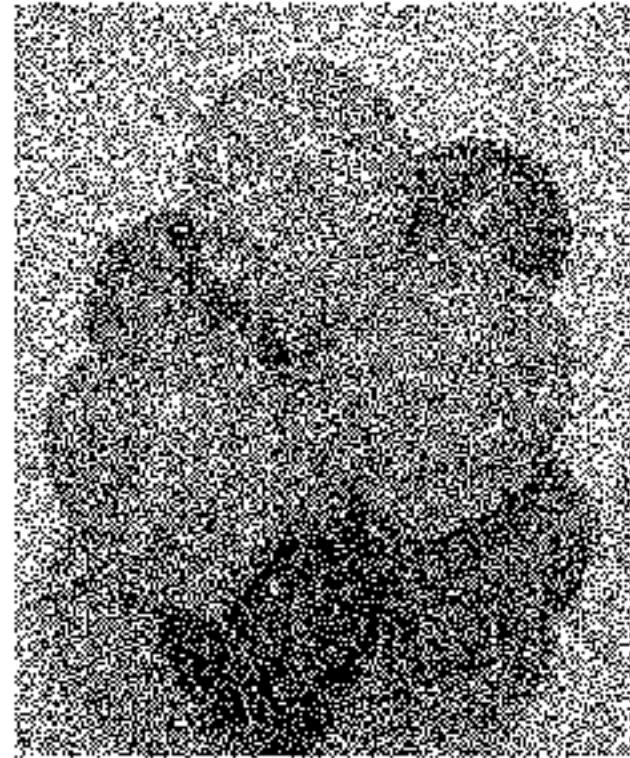
- $P(x, y) = \text{round}(I(x, y) + \text{noise}(x, y))$

- Where $\text{round}()$ chooses nearest value that can be represented.

Original



2 Levels (1bit with Random Dither)



Quantization and Dithering

- Random Dither:

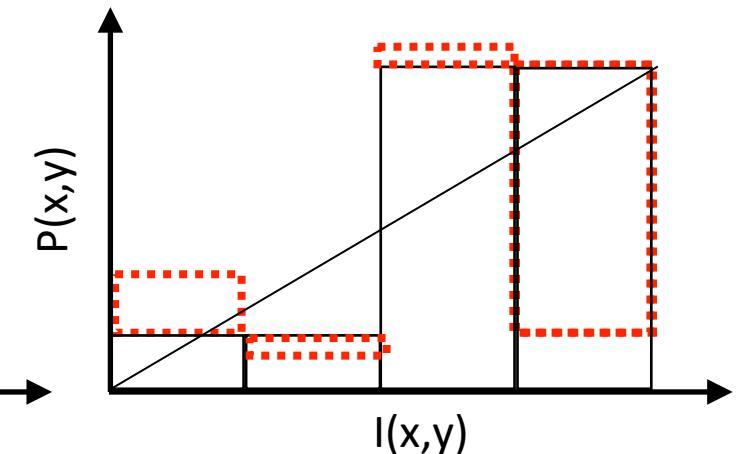
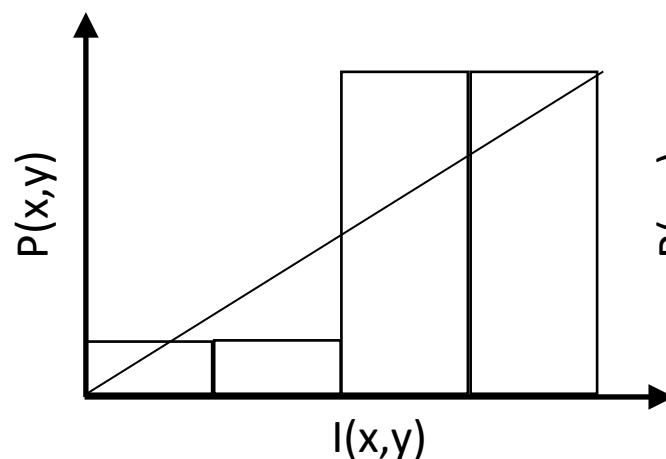
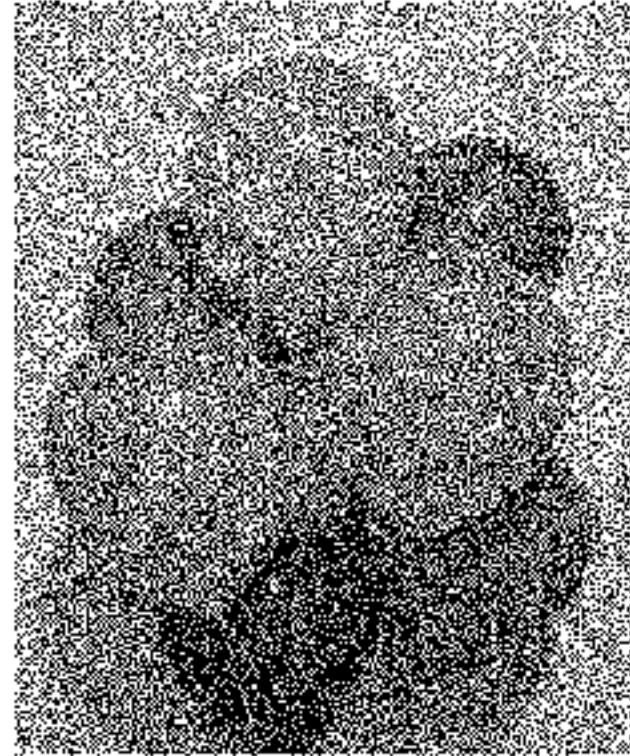
- $P(x, y) = \text{round}(I(x, y) + \text{noise}(x, y))$

- Where $\text{round}()$ chooses nearest value that can be represented.

Original



2 Levels (1bit with Random Dither)



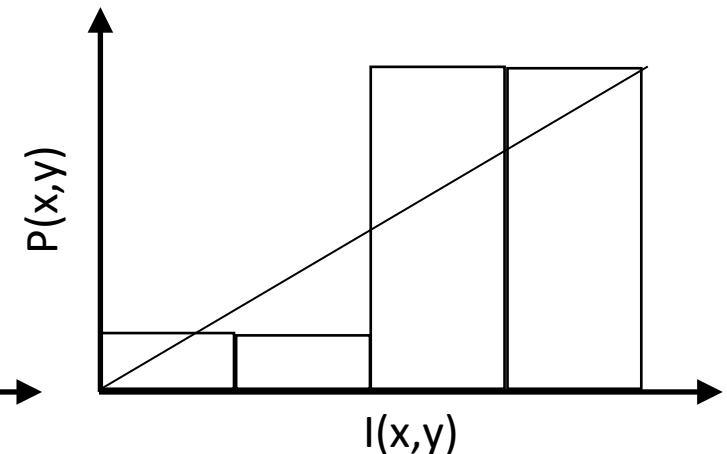
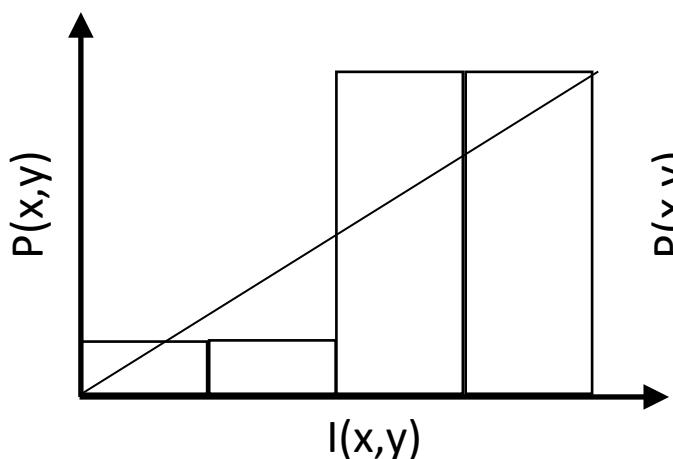
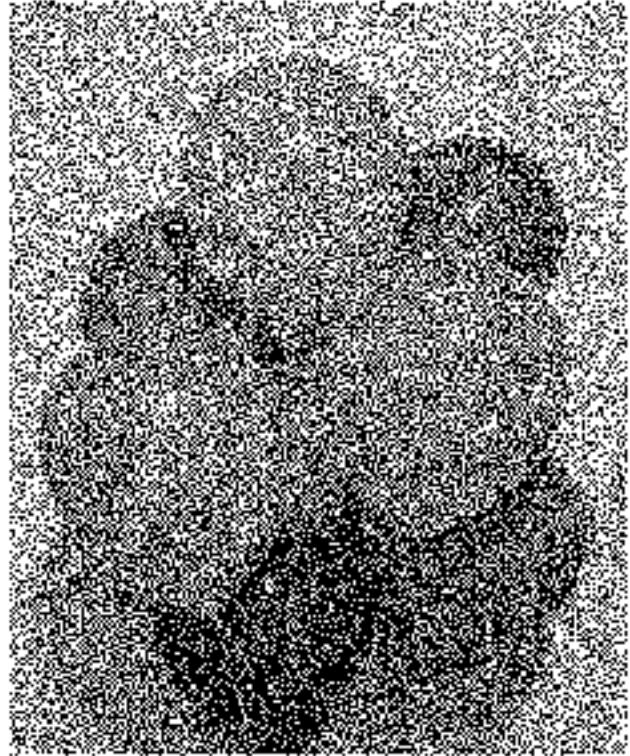
Quantization and Dithering

- Random Dither:
 - $P(x, y) = \text{round}(I(x, y) + \text{noise}(x, y))$
 - Where $\text{round}()$ chooses nearest value that can be represented.

Original



2 Levels (1bit with Random Dither)



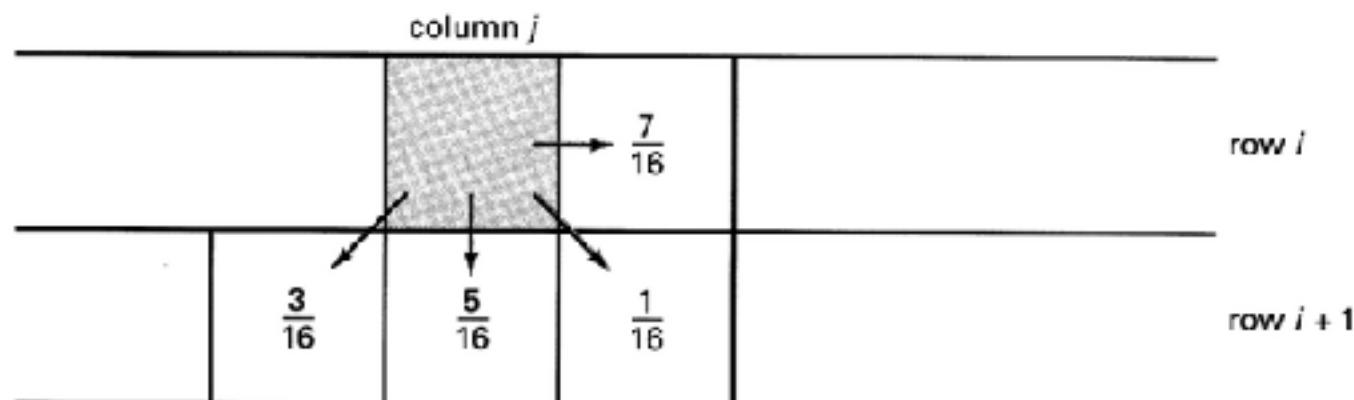
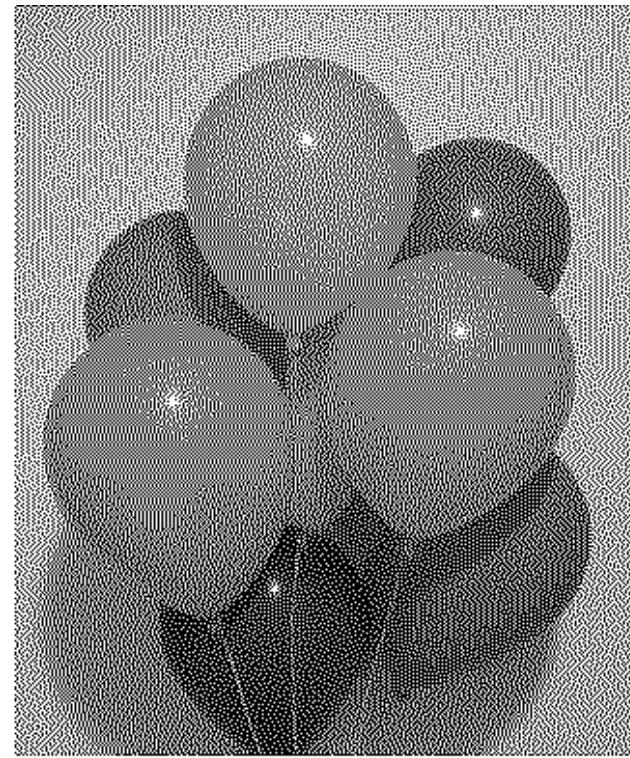
Quantization and Dithering

- Diffusion Dither:
 - E.g. Floyd–Steinberg dithering
 - Spread quantisation error over neighbouring pixels
 - Error dispersed to pixels right and below
 - Floyd-Steinberg weights:
 - $3/16 + 5/16 + 1/16 + 7/16 = 1.0$

2 Levels (1bit with Random Dither)

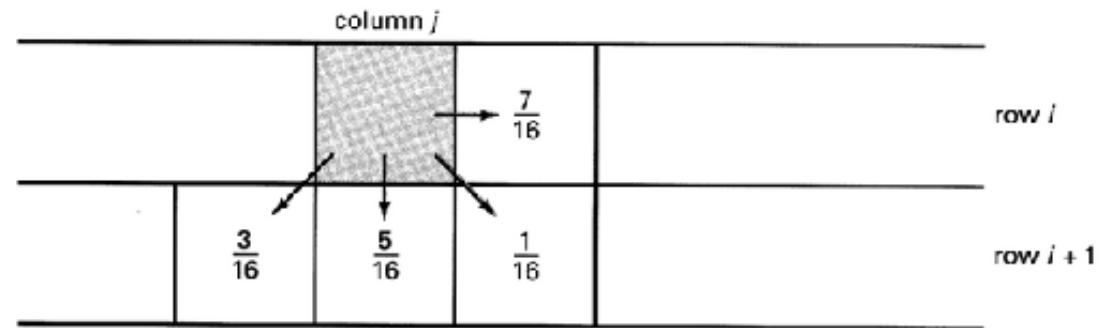


2 Levels (1bit with Diffusion Dither)



Quantization and Dithering

- Diffusion Dither:
 - E.g. Floyd–Steinberg dithering

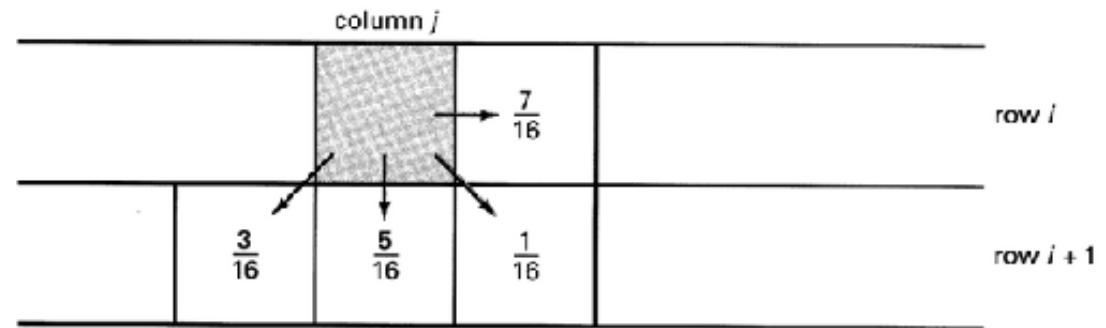


```
for each y from top to bottom do
    for each x from left to right do
        oldpixel = pixels[x][y]
        newpixel = find_closest_palette_color(oldpixel)
        pixels[x][y] = newpixel
        quant_error = oldpixel - newpixel
        pixels[x + 1][y      ] = pixels[x + 1][y      ] + quant_error × 7 / 16
        pixels[x - 1][y + 1] = pixels[x - 1][y + 1] + quant_error × 3 / 16
        pixels[x      ][y + 1] = pixels[x      ][y + 1] + quant_error × 5 / 16
        pixels[x + 1][y + 1] = pixels[x + 1][y + 1] + quant_error × 1 / 16
```

Quantisation

Quantization and Dithering

- Diffusion Dither:
 - E.g. Floyd–Steinberg dithering



```
for each y from top to bottom do
    for each x from left to right do
        oldpixel = pixels[x][y]
        newpixel = find_closest_palette_color(oldpixel)
        pixels[x][y] = newpixel
        quant_error = oldpixel - newpixel
        pixels[x + 1][y      ] = pixels[x + 1][y      ] + quant_error × 7 / 16
        pixels[x - 1][y + 1] = pixels[x - 1][y + 1] + quant_error × 3 / 16
        pixels[x      ][y + 1] = pixels[x      ][y + 1] + quant_error × 5 / 16
        pixels[x + 1][y + 1] = pixels[x + 1][y + 1] + quant_error × 1 / 16
```

Diffusion dithering by distributing quantisation error

Quantization and Dithering

8 colour quantisation



8 colour quantisation with Diffusion Dither



2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Image Filters -

Linear Filter

Linear Filters



Linear Filters and Image Convolutions

- Image filtering with linear filters
 - Construct array \mathbf{R} of same size as original image \mathbf{F}
 - Fill each location in \mathbf{R} with a weighted sum of the pixel values of corresponding location's neighbourhood in \mathbf{F}
 - Use the same weights each time
 - Different sets of weights represent different filters
 - **Convolution** with a (filter) **kernel**
- **Shift-invariant:** filtered result depends on image neighbourhood but not on position of neighbourhood
- **Linear:** filtering the sum of two images is equivalent as summing the two filtered images
- **Shift-Invariant Linear Systems (SILS)**



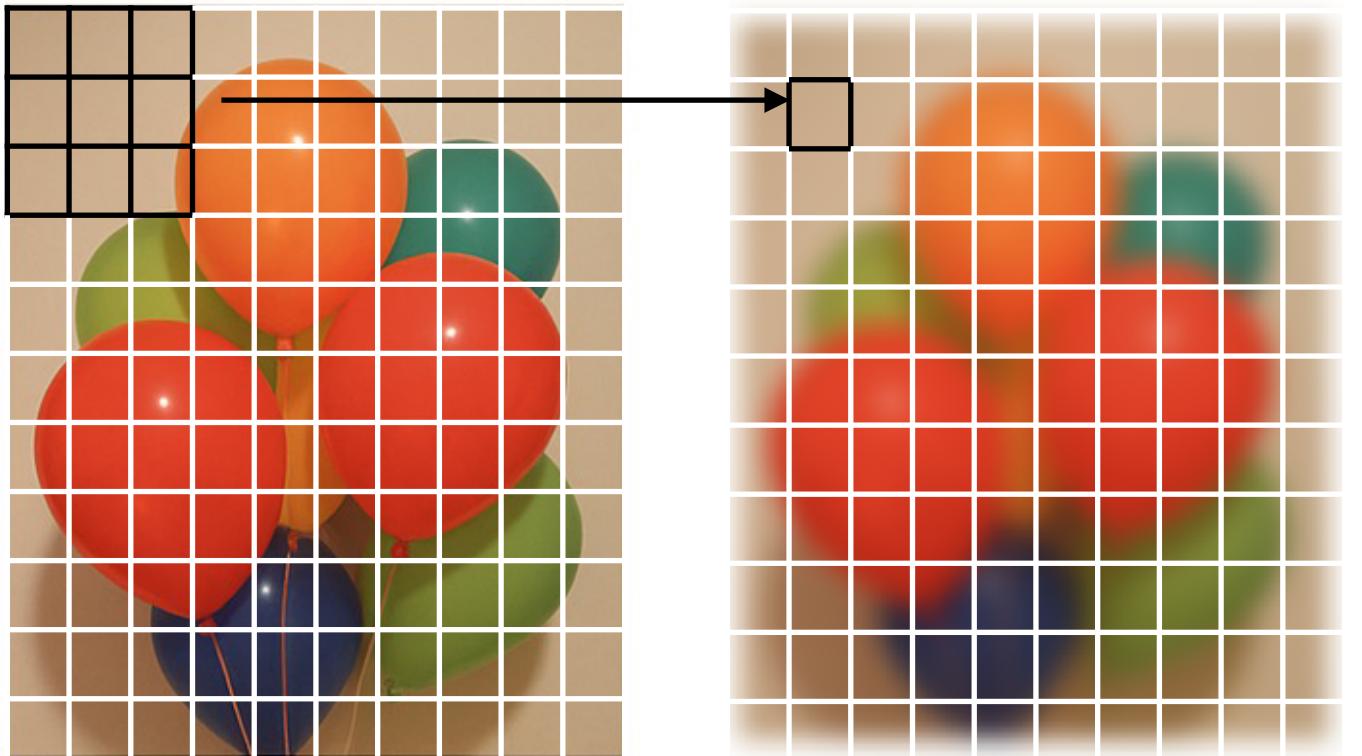
\mathbf{F}



\mathbf{R}

Linear Filters and Image Convolutions

- Image filtering with linear filters
 - Construct array \mathbf{R} of same size as original image \mathbf{F}
 - Fill each location in \mathbf{R} with a weighted sum of the pixel values of corresponding location's neighbourhood in \mathbf{F}
 - Use the same weights each time
 - Different sets of weights represent different filters
 - ***Convolution*** with a (filter) ***kernel***
- **Shift-invariant:** filtered result depends on image neighbourhood but not on position of neighbourhood
- **Linear:** filtering the sum of two images is equivalent as summing the two filtered images
- **Shift-Invariant Linear Systems (SILS)**

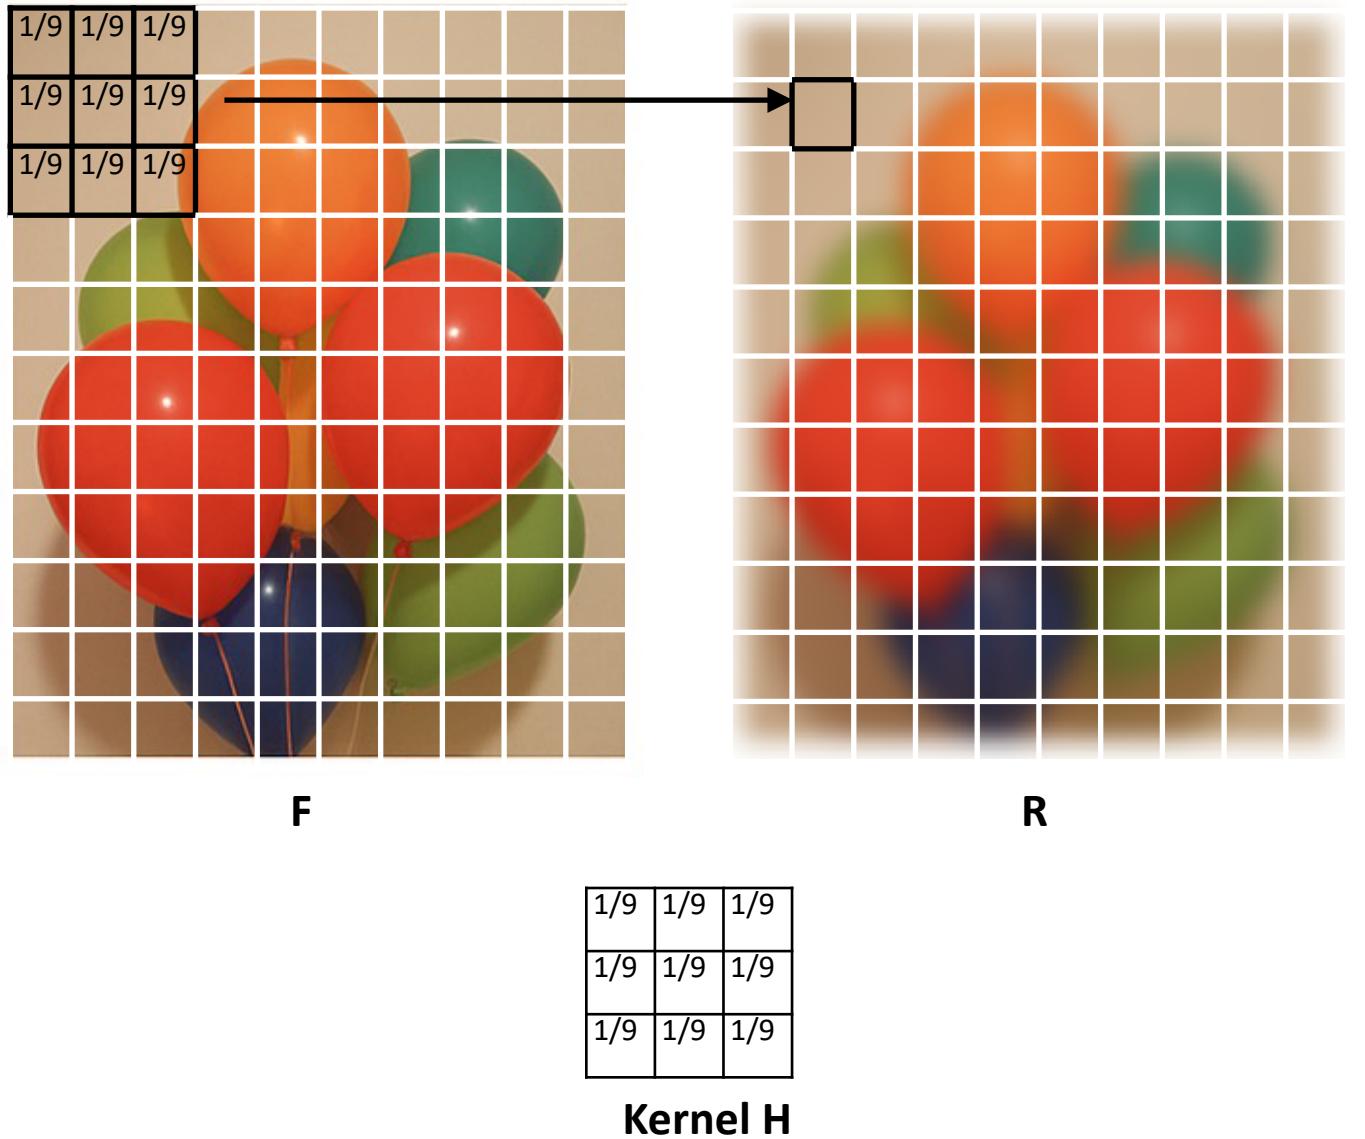


F

R

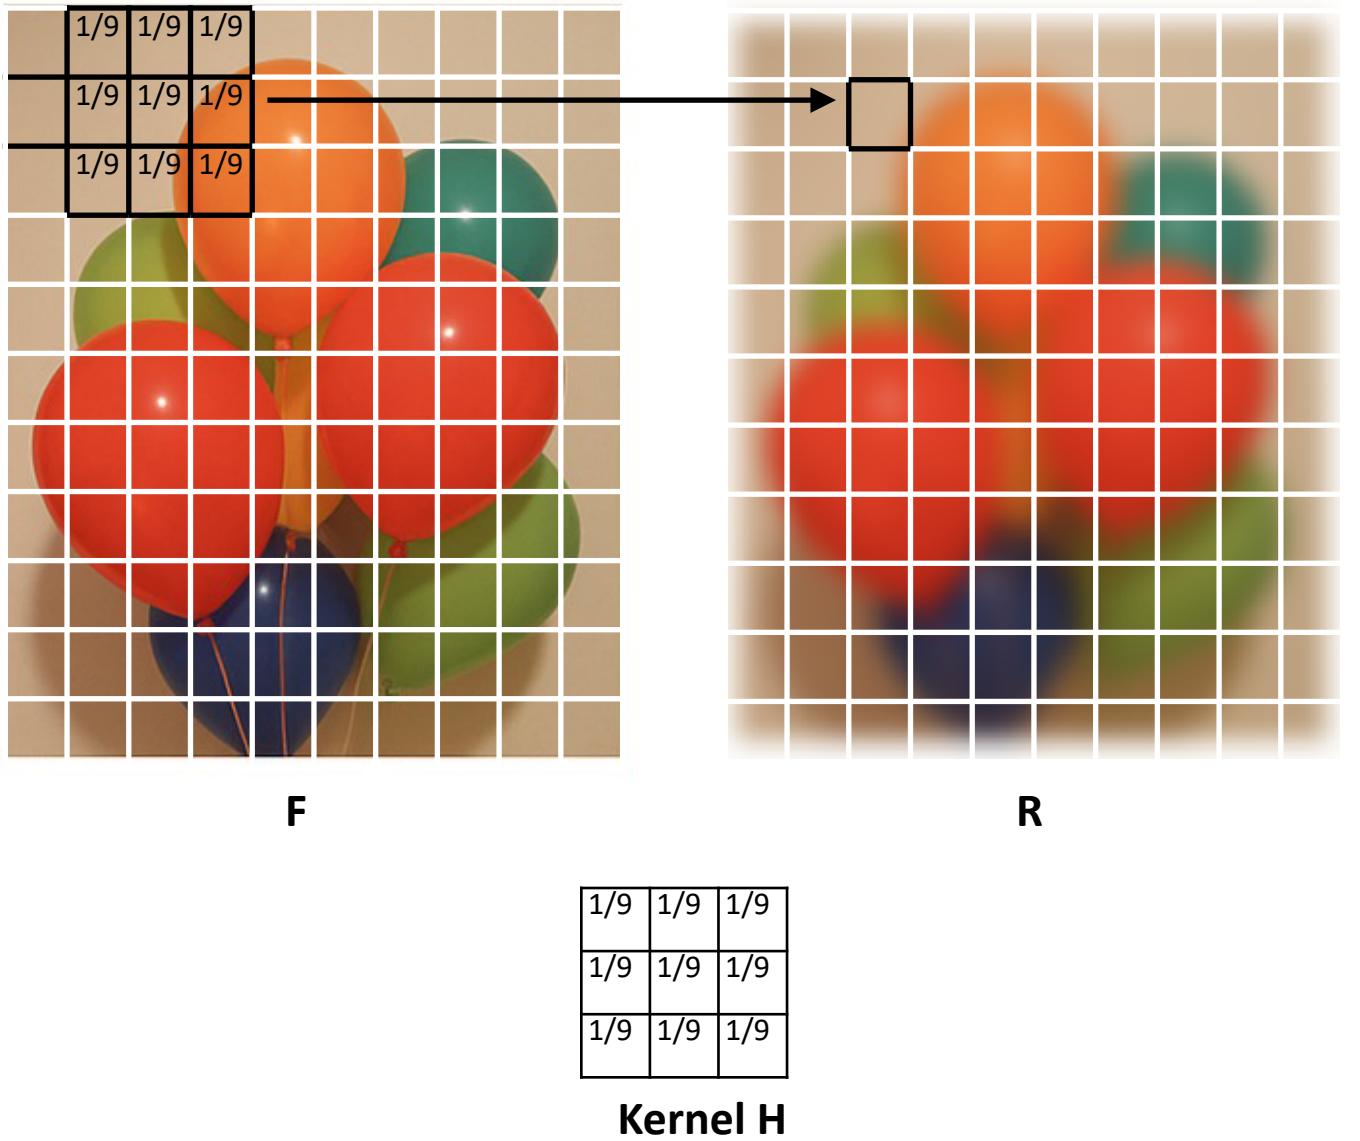
Linear Filters and Image Convolutions

- Image filtering with linear filters
 - Construct array **R** of same size as original image **F**
 - Fill each location in **R** with a weighted sum of the pixel values of corresponding location's neighbourhood in **F**
 - Use the same weights each time
 - Different sets of weights represent different filters
 - **Convolution** with a (filter) **kernel**
- **Shift-invariant:** filtered result depends on image neighbourhood but not on position of neighbourhood
- **Linear:** filtering the sum of two images is equivalent as summing the two filtered images
- **Shift-Invariant Linear Systems (SILS)**



Linear Filters and Image Convolutions

- Image filtering with linear filters
 - Construct array \mathbf{R} of same size as original image \mathbf{F}
 - Fill each location in \mathbf{R} with a weighted sum of the pixel values of corresponding location's neighbourhood in \mathbf{F}
 - Use the same weights each time
 - Different sets of weights represent different filters
 - **Convolution** with a (filter) **kernel**
- **Shift-invariant:** filtered result depends on image neighbourhood but not on position of neighbourhood
- **Linear:** filtering the sum of two images is equivalent as summing the two filtered images
- **Shift-Invariant Linear Systems (SILS)**



Linear Filters and Image Convolutions

- Convolution
 - The patterns of weights is usually referred to as **kernel H**
 - The process of applying H to an image F is known as **convolution**
 - Say: “*The convolution of H with F results in R*”
 - The kernel H can be defined as a continuous function or a finite matrix of discrete weights
 - What is the kernel for our local average example?
 - How large is it?
 - $(2k)^2$ weights

$$R_{i,j} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{i+k} \sum_{v=j-k}^{j+k} F_{u,v}$$

$$R_{i,j} = \sum_{u=i-k}^{i+k} \sum_{v=j-k}^{j+k} H_{i-u, j-v} \cdot F_{u,v}$$

$$R = H \otimes F$$

$$H_{i,j} = \frac{1}{(2k+1)^2}$$

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

$$H$$

Linear Filters and Image Convolutions

- Kernels used for filters
- Array of weighting values
- Typically 3×3 , 5×5 , etc.
- At each pixel in the image:
 - The kernel is overlaid so that the middle number is over the current pixel
 - Each weight is multiplied by the corresponding pixel value
 - The result of the filter is the sum of these products

9	6	9	8	1	5
8	7	6	2	2	3
9	7	4	6	7	4
8	9	3	7	5	7
6	2	4	5	2	1
4	1	3	5	3	4

$$H = \frac{1}{16} \begin{vmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{vmatrix}$$

Linear Filters and Image Convolutions

- Kernels used for filters
- Array of weighting values
- Typically 3×3 , 5×5 , etc.
- Example

$$\begin{aligned}[F * H](3, 4) &= \frac{1}{16}(7 \times 1 + 4 \times 2 + 6 \times 1 + \\ &\quad 9 \times 2 + 3 \times 4 + 7 \times 2 + \\ &\quad 2 \times 1 + 4 \times 2 + 5 \times 1) \\ &= \frac{1}{16}(7 + 8 + 6 + 18 + 12 + 14 + 2 + 8 + 5) \\ &= 5\end{aligned}$$

9	6	9	8	1	5
8	7	6	2	2	3
9	7	4	6	7	4
8	9	3	7	5	7
6	2	4	5	2	1
4	1	3	5	3	4

$$H = \frac{1}{16} \begin{vmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{vmatrix}$$

Linear Filters and Image Convolutions

Original



$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 0 & -\frac{1}{2} & 0 \\ -\frac{1}{2} & 3 & -\frac{1}{2} \\ 0 & -\frac{1}{2} & 0 \end{bmatrix}$$

Linear Filters and Image Convolutions

Original



$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Box Blur



$$\begin{bmatrix} \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \\ \frac{1}{9} & \frac{1}{9} & \frac{1}{9} \end{bmatrix}$$

Gaussian Blur



$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Sharpen



$$\begin{bmatrix} 0 & -\frac{1}{2} & 0 \\ -\frac{1}{2} & 3 & -\frac{1}{2} \\ 0 & -\frac{1}{2} & 0 \end{bmatrix}$$

Linear Filters and Image Convolutions

Original



Edge horizontal



Edge vertical



Edge (Sobel Operator)



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$

$$\begin{matrix} -\frac{1}{2} & 0 & \frac{1}{2} \\ -1 & 0 & 1 \\ -\frac{1}{2} & 0 & \frac{1}{2} \end{matrix}$$

$$\begin{matrix} -\frac{1}{2} & -1 & -\frac{1}{2} \\ -0 & 0 & 0 \\ \frac{1}{2} & 1 & \frac{1}{2} \end{matrix}$$

Linear Filters and Image Convolutions

Original



Edge horizontal



Edge vertical



Edge (Sobel Operator)



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix}$$

$$\begin{matrix} -\frac{1}{2} & 0 & \frac{1}{2} \\ -1 & 0 & 1 \\ -\frac{1}{2} & 0 & \frac{1}{2} \end{matrix}$$

$$\begin{matrix} -\frac{1}{2} & -1 & -\frac{1}{2} \\ -0 & 0 & 0 \\ \frac{1}{2} & 1 & \frac{1}{2} \end{matrix}$$

Linear Filters and Image Convolutions

- Edge filter (Sobel Operator):

9	6	9	8	1	5
8	7	6	2	2	3
9	9	9	9	7	4
8	1	1	1	5	7
6	1	1	1	2	1
4	1	3	5	3	4

$$[F * K]_{2,3} = (-0.5 \times 9 + -1 \times 9 + -0.5 \times 9 + \dots)$$

$$= -16$$

$$H = \begin{bmatrix} -\frac{1}{2} & -1 & -\frac{1}{2} \\ 0 & 0 & 0 \\ \frac{1}{2} & 1 & \frac{1}{2} \end{bmatrix}$$

Linear Filters and Image Convolutions

- Implementation:

```
// Image is width x height, filtered into Result
// Kernel has 'radius' r, so is (2r+1) square
for (y = 0; y < height; ++y) {
    for (x = 0; x < width; ++x) {
        Result[x][y] = 0;
        for (dy = -r; dy <= r; ++dy) {
            for (dx = -r; dx <= r; ++dx) {
                Result[x][y] += Kernel[r+dx][r+dy] * Image[x+dx][y+dy];
            }
        }
    }
}
```

Linear Filters and Image Convolutions

- Problematic cases for Image Convolutions
- Edges of images:
 - Out of bounds errors
 - Can stay away from edges
 - Or use nearest valid pixel
- Results out of range
 - What if the sum is less than 0 or greater than 255?
 - Some filters give results in different ranges (e.g. edge detection)

9	6	9	8	1	5	
8	7	6	2	2	3	
9	7	4	6	7	4	?
8	9	3	7	5	7	?
6	2	4	5	2	1	?
4	1	3	5	3	4	

Linear Filters and Image Convolutions

- Convolution Rules for SILS:

- **Superimposition:** the sum of two filtered images is equivalent to the filtered image with the sum of the two kernels (component-wise addition)

$$H_1 \otimes R + H_2 \otimes R = (H_1 + H_2) \otimes R$$

- **Scaling:** scaling of a filtered image is equivalent to filtering the image with a scaled kernel

$$(kH) \otimes R = k(H \otimes R)$$

- **Symmetric:** the convolution of two filter to an image is symmetric

$$H_1 \otimes (H_2 \otimes R) = H_2 \otimes (H_1 \otimes R)$$

- **Associative:** convolution is associative, which means that we can find a single kernel that behaves like the composition of two kernels

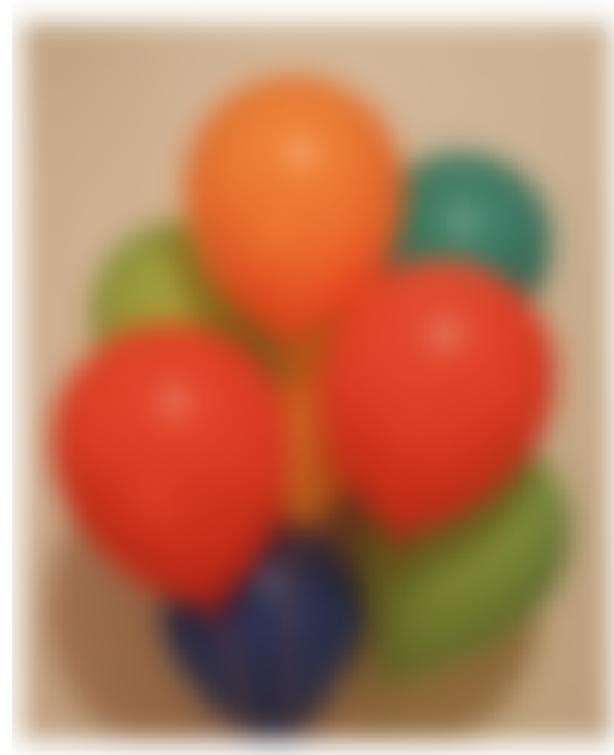
$$H_1 \otimes (H_2 \otimes R) = (H_1 \otimes H_2) \otimes R$$

Linear Filters and Image Convolutions

- Something to discuss:
 - Many convolutions have kernels where the values add up to 1. Why is this the case?
 - What are the largest and smallest values you might get from the different filters?
 - How can you apply these filters to colour images?



F



R

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Kernel H

Image Pyramids

Convolutions

- Image Pyramids (here **Gaussian Pyramid**):

- An image pyramid is a collection of representations of the same image at different scales
- Typically, each layer is half the width and half the height of the previous layer
- At each layer, a Gaussian pyramid stores a smoothed (Gaussian kernel) and downsampled version of the previous layer



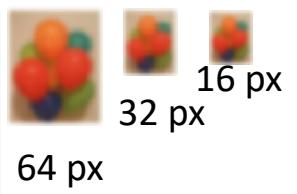
512px



256px



128px



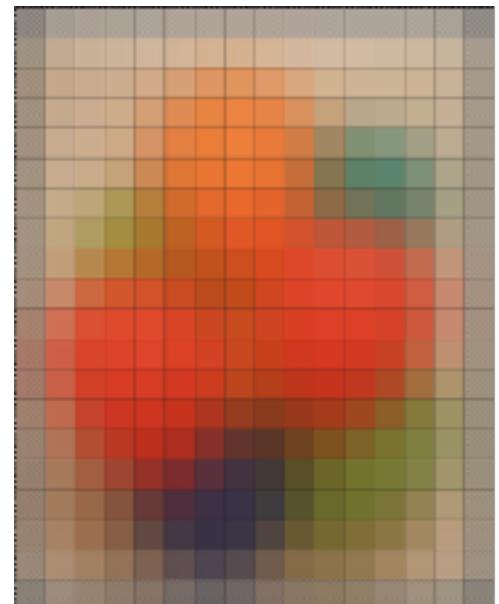
64 px



32 px



16 px



Convolutions

- What are possible applications for Gaussian pyramids?
 - **Search over scale:** scale invariant pattern matching, search with same pattern over multiple scales
 - **Spatial search:** search matches in coarse layers of two and track downwards, thereby search regions are decreased
 - **Feature tracking:** accept features found in fine-scale images only if they can be tracked across different coarser scales to suppress noise
 - More about that next time!



512px



256px



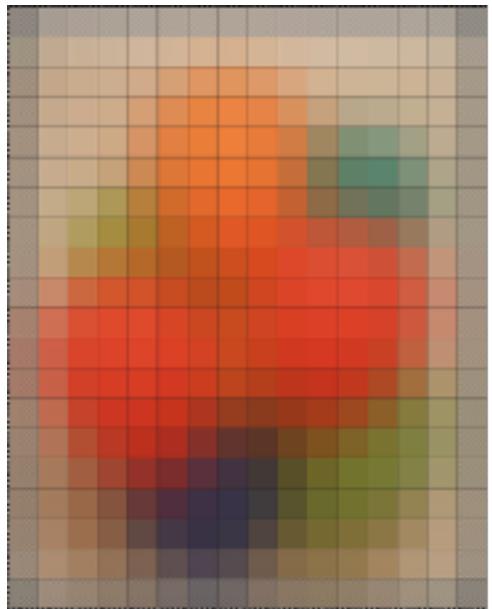
128px



64 px

32 px

16 px



Convolutions

- A Gauss filter is a low-pass filter
 - Gaussian pyramid successively removes high frequencies
 - What happens if two consecutive layers are subtracted?
 - The resulting image contains the removed (via Gauss filter) spatial frequencies
 - What happens if this is done for each layer of the Gaussian pyramid?
 - The result is another image pyramid, called **Laplacian** pyramid, that can be thought of as the response of a band-pass filter

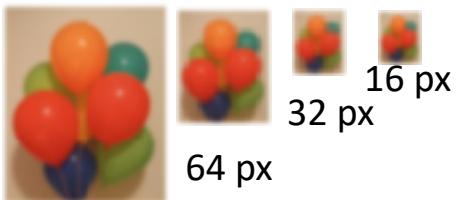
Gaussian Pyramid



512px



256px



64 px

32 px

16 px



Laplacian Pyramid



2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

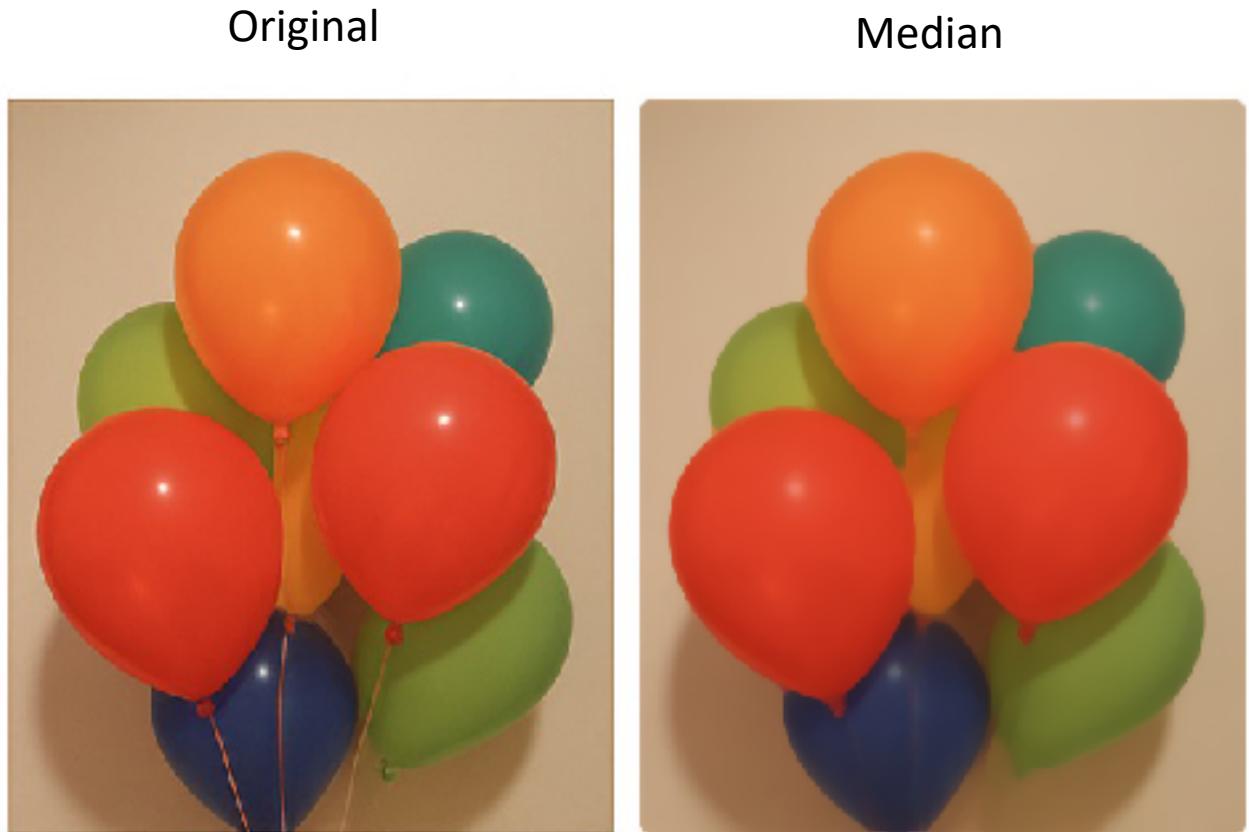
2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Image Filters - Non-Linear Filters

Non-Linear Filters

- Not all filters are implemented as convolution operations
- E.g. Median filter
 - Take a local neighbourhood (e.g. a 3×3 or 5×5 window)
 - Sort the pixel values
 - Use the middle value as the result
 - Has a smoothing effect, but preserves sharp edges



2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

2D Image Processing

- Luminance
 - Brightness
 - Contrast
 - Histogram equalization
 - Gamma
- Color
 - Grayscale
 - Saturation
- Dithering
 - Quantization
 - Random dither
 - Floyd-Steinberg
- Linear filtering
 - Convolution
 - Blur & Sharpen
 - Edge detection
 - Image Pyramids
- Non-linear filtering
 - Median
- Image Feature Detection
 - Edges (e.g Canny Edge)
 - Corners (e.g. Shi-Tomasi, Harris Corner)
 - Blobs (e.g. Difference of Gaussians)

Image Features & Feature Detection

Image Features & Feature Detection

- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views



What are good image features?

Image Features & Feature Detection

- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views
- Feature descriptors
 - Used to tell features apart
 - Ideally robust to changes in lighting, viewpoint, contrast, etc.
- Three common types of features
 - Edges – gradient-based
 - Corners – gradient-based
 - Blobs – bright/dark centre/surround



What are good image features?

Image Features & Feature Detection

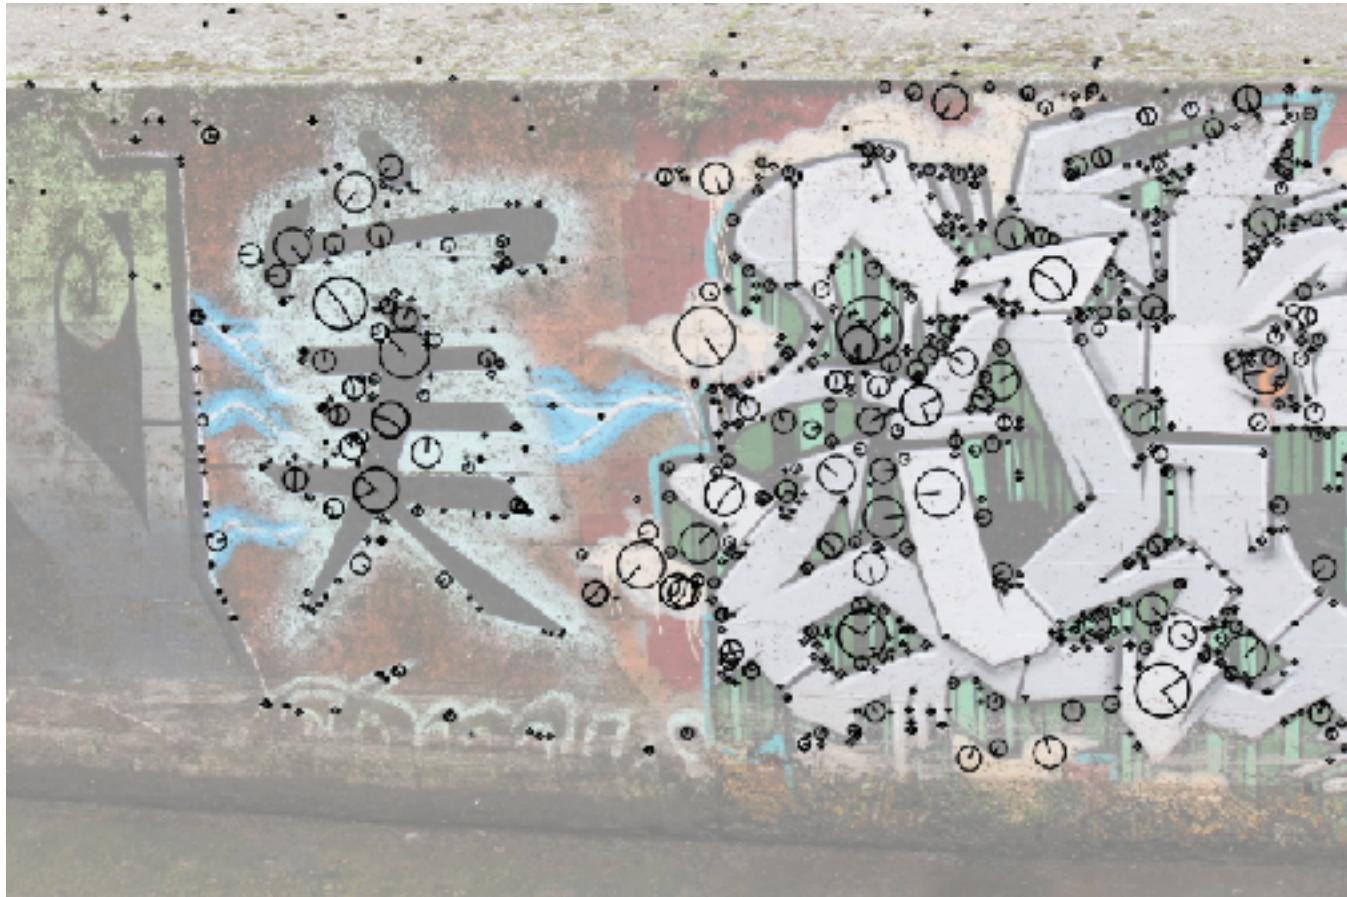
- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views
- Feature descriptors
 - Used to tell features apart
 - Ideally robust to changes in lighting, viewpoint, contrast, etc.
- Three common types of features
 - Edges – gradient-based
 - Corners – gradient-based
 - Blobs – bright/dark centre/surround



Example: Corners

Image Features & Feature Detection

- Often required:
 - Features in an image that can be reliably detected in different images of the same scene
- A ‘feature’ or ‘key point’ is
 - A 2D location in an image
 - Accurately/precisely located
 - Can be found in different views
- Feature descriptors
 - Used to tell features apart
 - Ideally robust to changes in lighting, viewpoint, contrast, etc.
- Three common types of features
 - Edges – gradient-based
 - Corners – gradient-based
 - Blobs – bright/dark centre/surround



Example: Blobs

Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise ($>$ Gaussian filter G_σ)

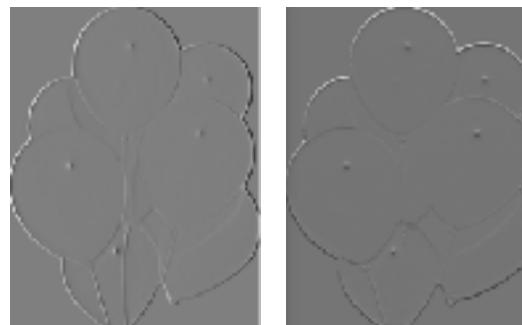
Original



Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)
 - Step 2: Compute gradients in x and y direction (> Sobel operators H_u, H_v), gradient is a vector
$$\mathbf{g} = [H_u \quad H_v]^T$$

Original

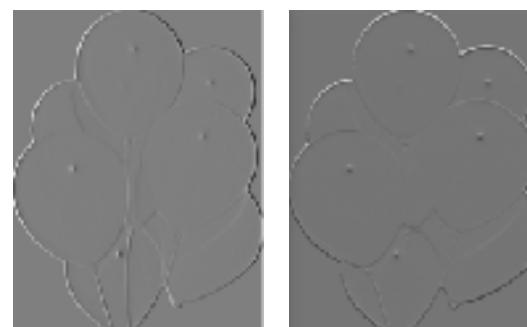


$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)
 - Step 2: Compute gradients in x and y direction (> Sobel operators H_u, H_v), gradient is a vector
$$\mathbf{g} = [H_u \quad H_v]^T$$
 - Step 3: compute gradient magnitudes (edge strength) and gradient directions (edge directions)

Original



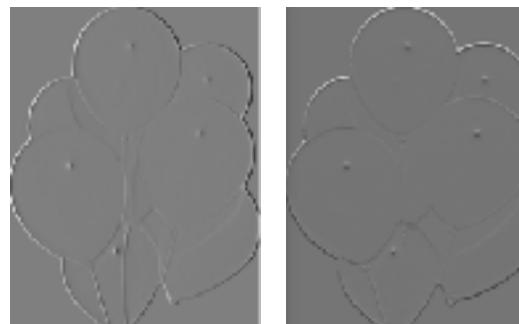
$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



Magnitude

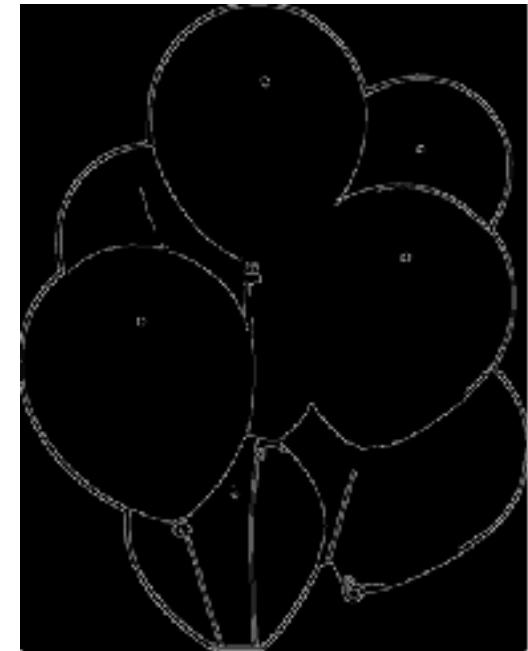
Edge Detection - Canny Edge Detector

- Canny Edge Detector
 - Step 1: Reduce noise (> Gaussian filter G_σ)
 - Step 2: Compute gradients in x and y direction (> Sobel operators H_u, H_v), gradient is a vector
$$\mathbf{g} = [H_u \quad H_v]^T$$
 - Step 3: compute gradient magnitudes (edge strength) and gradient directions (edge directions)
 - Step 4: select points with maximum magnitudes (above threshold T1)
 - Step 5: non-maximum suppression with hysteresis:
 - For all points with magnitudes above high-threshold T1
 - Trace along the edge (in edge direction) suppress all pixels that are not on traced line (pixel by pixel) and mark as visited
 - Stop tracing when point's magnitude is below low-threshold T2 or already visited



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Edges using Canny Edge Detector



Magnitude

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

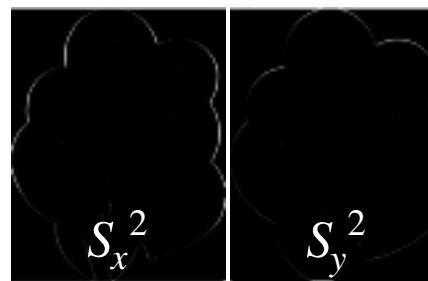


Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

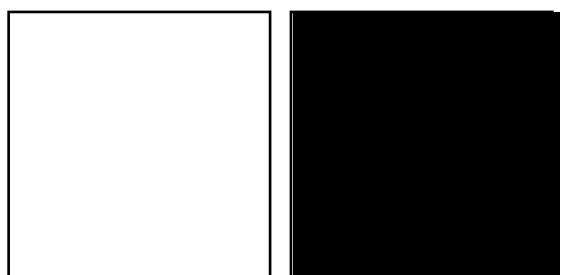
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



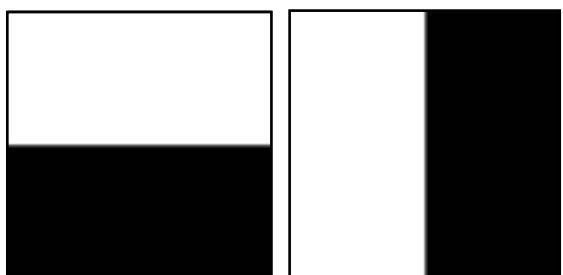
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

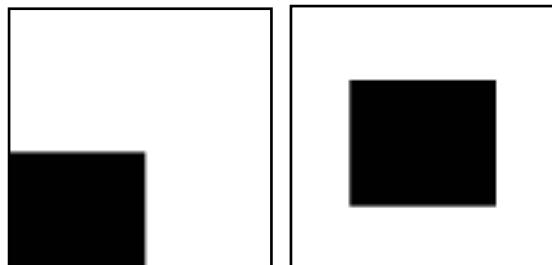


$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

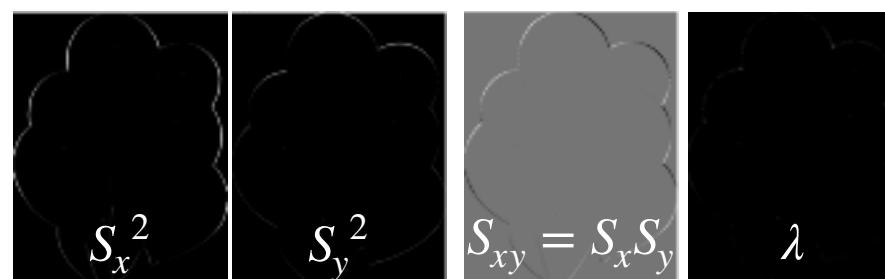
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
 - Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$

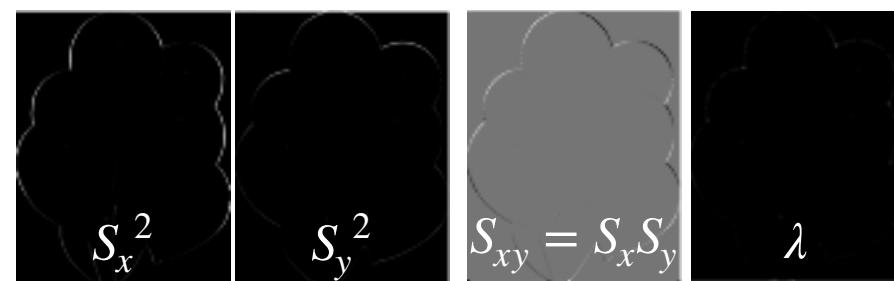
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
 - Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge
 - Step 4 (Corner Detection): A point is identified if the eigenvalues or their approximation are large (meaning the local image patch changes significantly in all directions)
 - 6. Post-processing:
Non-maximal suppression to keep only the strongest corners



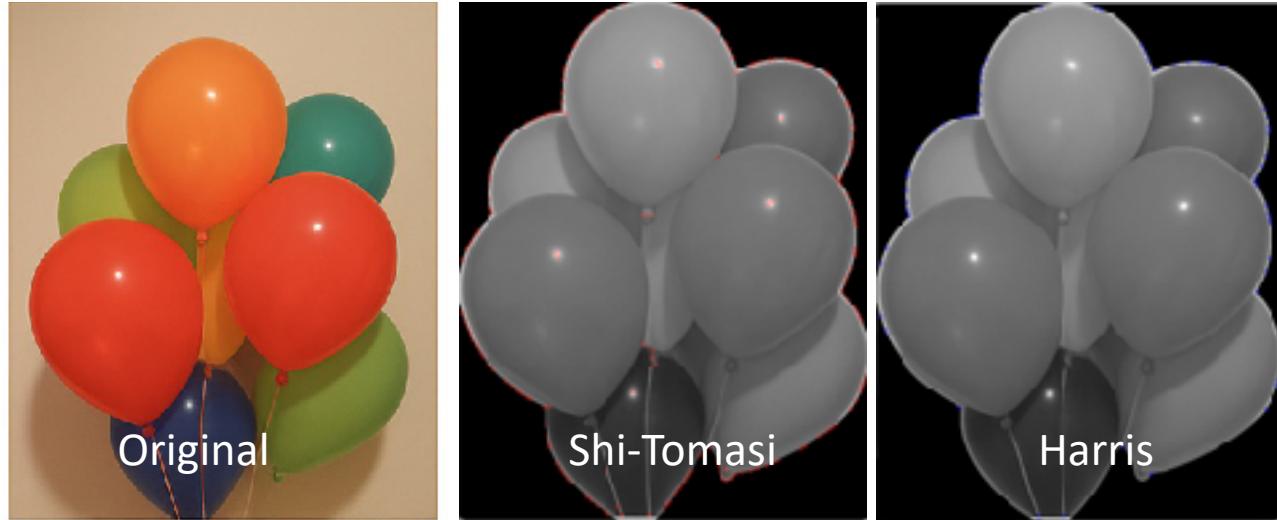
$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



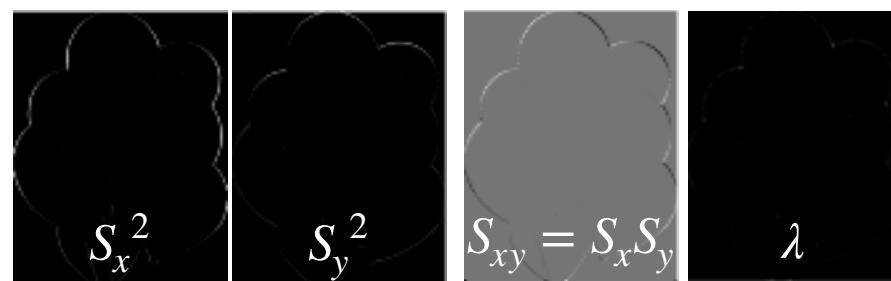
$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection
 - Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
 - Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
 - Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge
 - Step 4 (Corner Detection): A point is identified if the eigenvalues or their approximation are large (meaning the local image patch changes significantly in all directions)
 - 6. Post-processing:
Non-maximal suppression to keep only the strongest corners



$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_x S_y \\ S_x S_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_x S_y$$

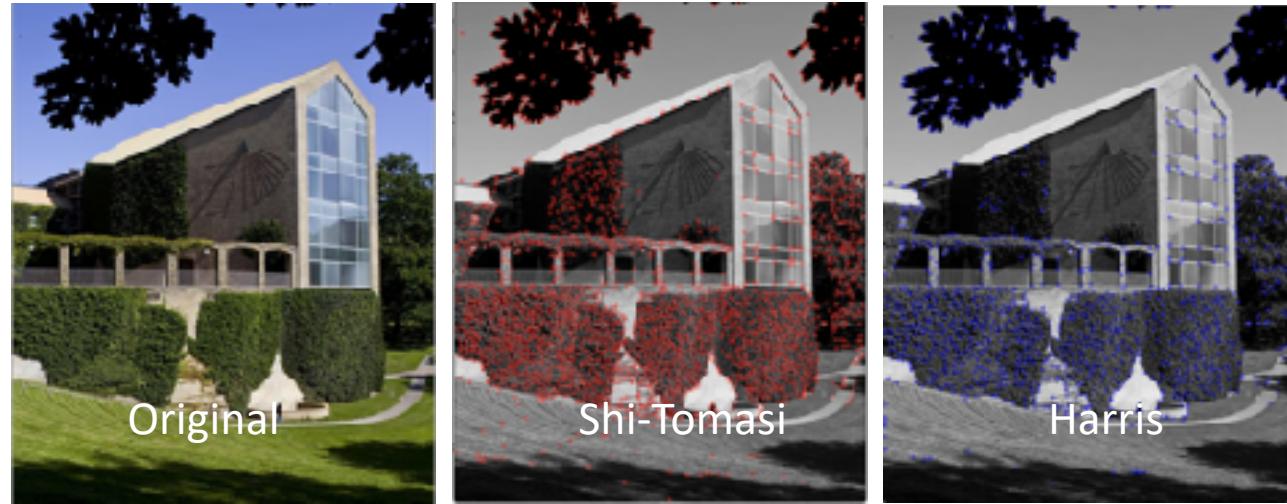
Corner Detection - Shi-Tomasi / Harris detector



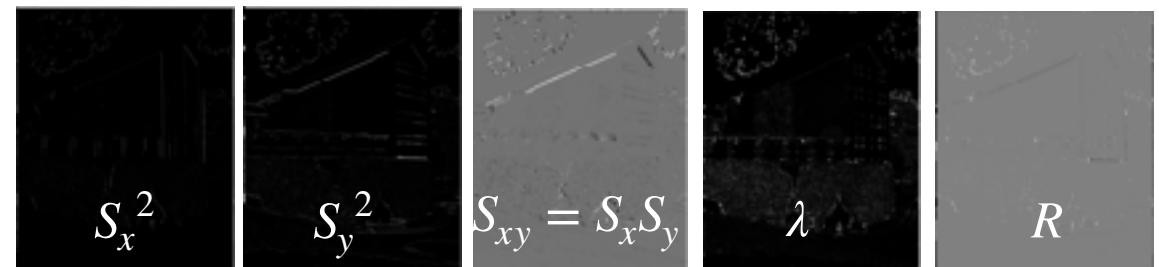
Corner Detection - Shi-Tomasi / Harris detector

- Corner Detection

- Step 1 (Gradients): For each pixel, the algorithm calculates the image's x and y derivatives S_x, S_y (gradients) using filters like Sobel.
- Step 2 (Second Moment Matrix): The derivatives are used to compute the second moment matrix (M) for a small window around the pixel:
 - Properties of M :
 - Constant windows: A and B are low, C is low
 - Hori./vert. edge: A or B is high, C is low
 - Diag. edge: A,B and C are high (removed later)
 - Corner: A,B and C are high
- Step 3 (Evaluation of M): Computing of Eigenvalues (λ) for M (Shi-Tomasi) or approximation using determinant and trace of M (Harris), can be used to correctly remove diagonal edge
- Step 4 (Corner Detection): A point is identified if the eigenvalues or their approximation are large (meaning the local image patch changes significantly in all directions)
- 6. Post-processing:
Non-maximal suppression to keep only the strongest corners

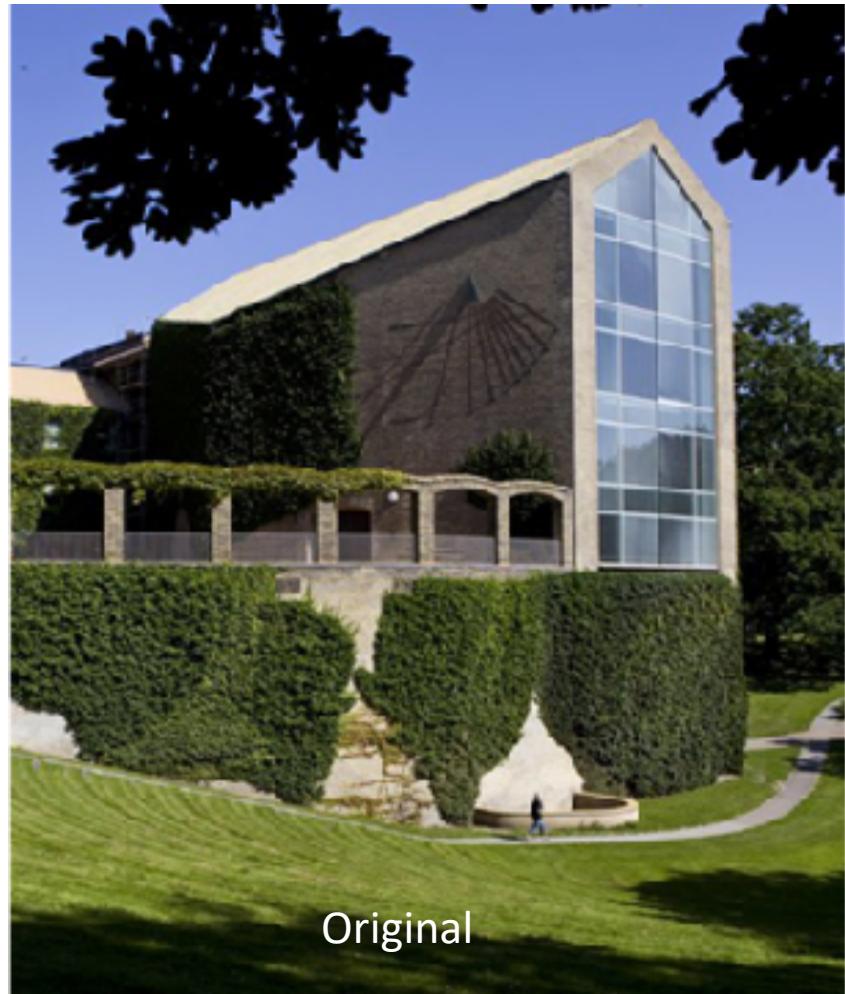


$$H_u = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad H_v = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$M = \begin{bmatrix} A & C \\ C & B \end{bmatrix} = \begin{bmatrix} (S_x)^2 & S_xS_y \\ S_xS_y & (S_y)^2 \end{bmatrix}$$
$$A = S_x^2, B = S_y^2, C = S_xS_y$$

Corner Detection - Shi-Tomasi / Harris detector



**Tell me more about blobs!
And what can I do with all that stuff?**

Next time :)

The end!