

# Visual Computing I:

Interactive Computer Graphics and Vision



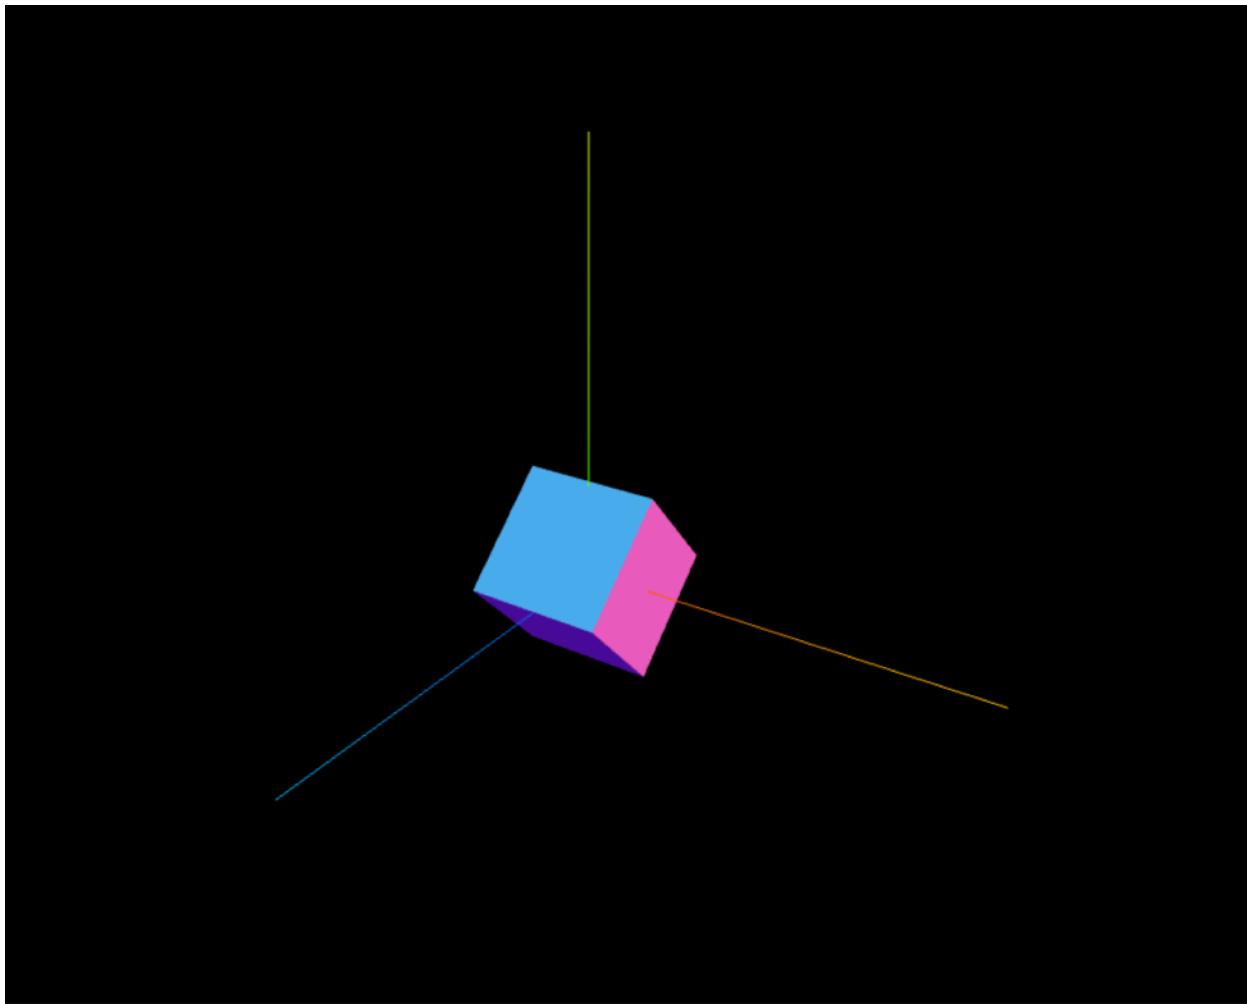
Camera Calibration and Stereo Vision

Stefanie Zollmann and Tobias Langlotz

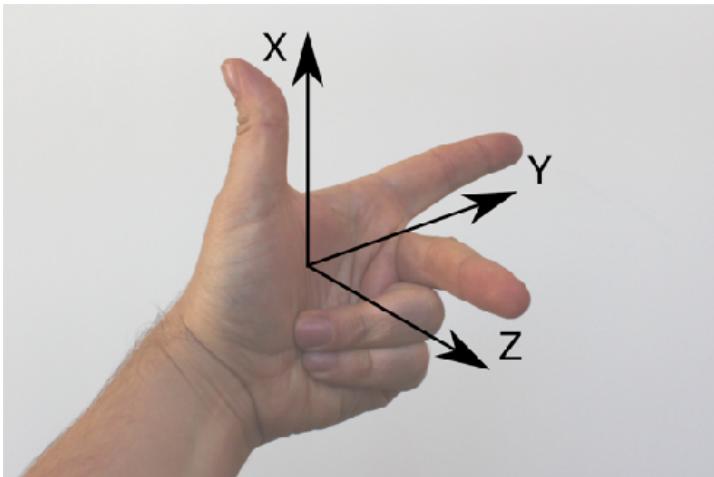
**Last time..**

# Three Dimensions

- We moved from 2D to 3D
- Many ideas the same:
  - Homogeneous co-ordinates
  - Scaling and translation
- Some things get tricky
  - Choice of left- or right-handed coordinates
  - Rotations get complicated
- Projection from 3D to 2D

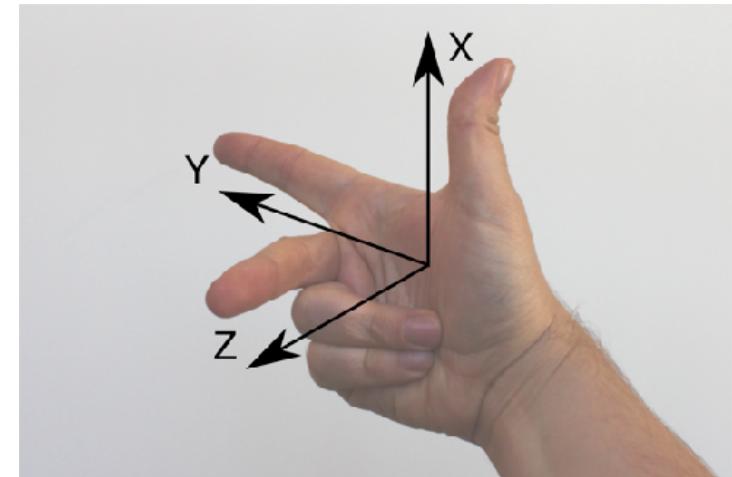


# Left & Right-Handed Coordinates



Left handed

Thumb is axis  
Forefinger is axis  
Middle finger is axis



Right-Handed



## Transform 3D

### Input

```
[0, -1, 0, 0],  
[1, 0, 0, 0],  
[0, 0, 1, 0],  
[0, 0, 0, 1]
```

**Apply Matrix!**

**Reset Matrix!**

### Output

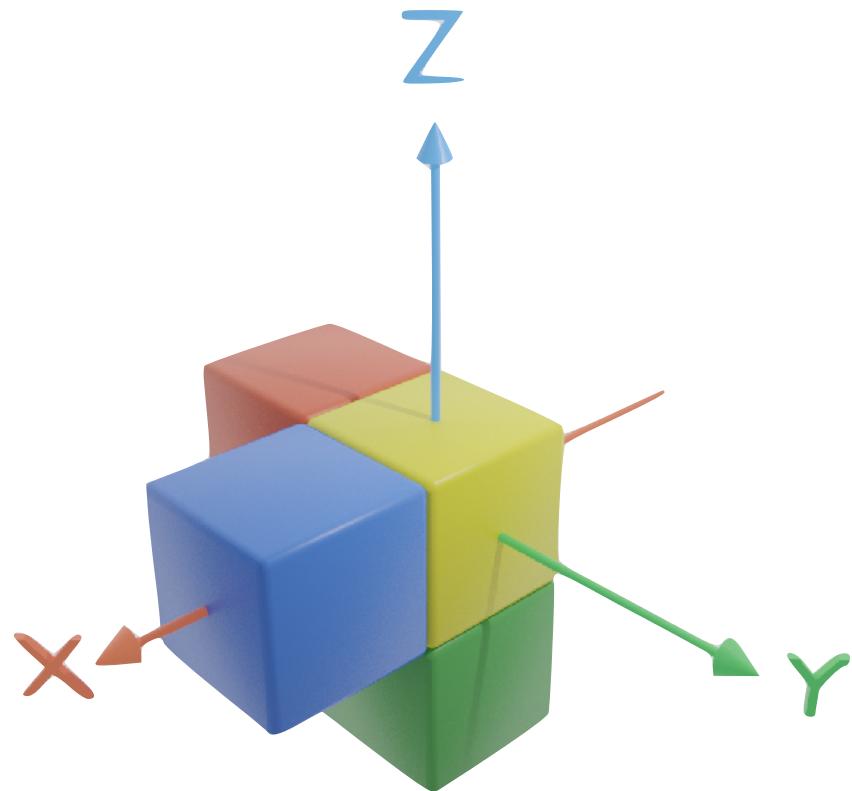
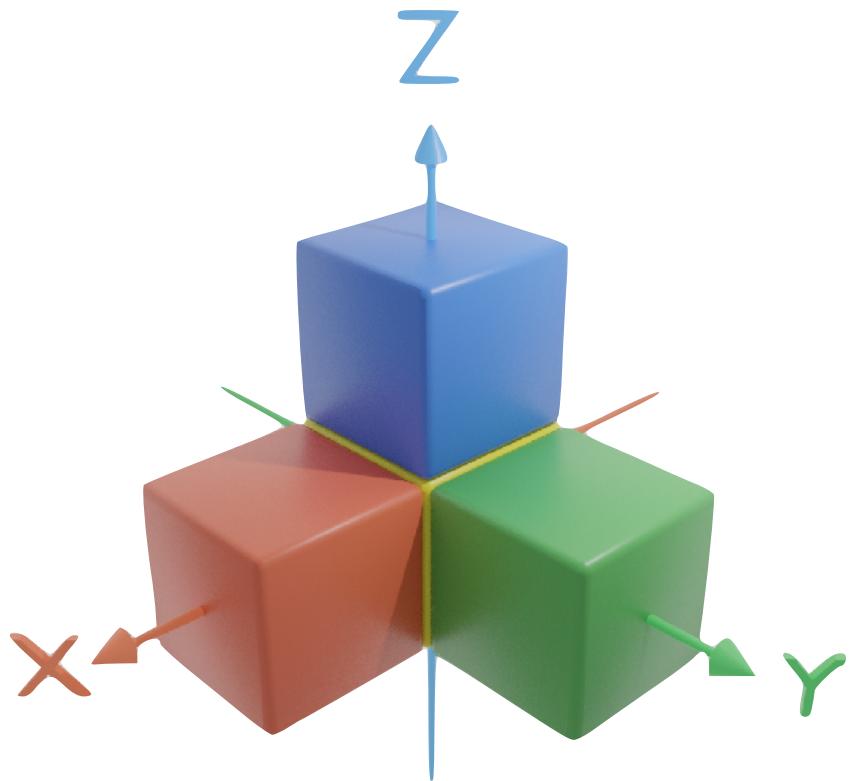
```
0.0000, -1.0000, 0.0000, 0.0000,  
1.0000, 0.0000, 0.0000, 0.0000,  
0.0000, 0.0000, 1.0000, 0.0000,  
0.0000, 0.0000, 0.0000, 1.0000,
```



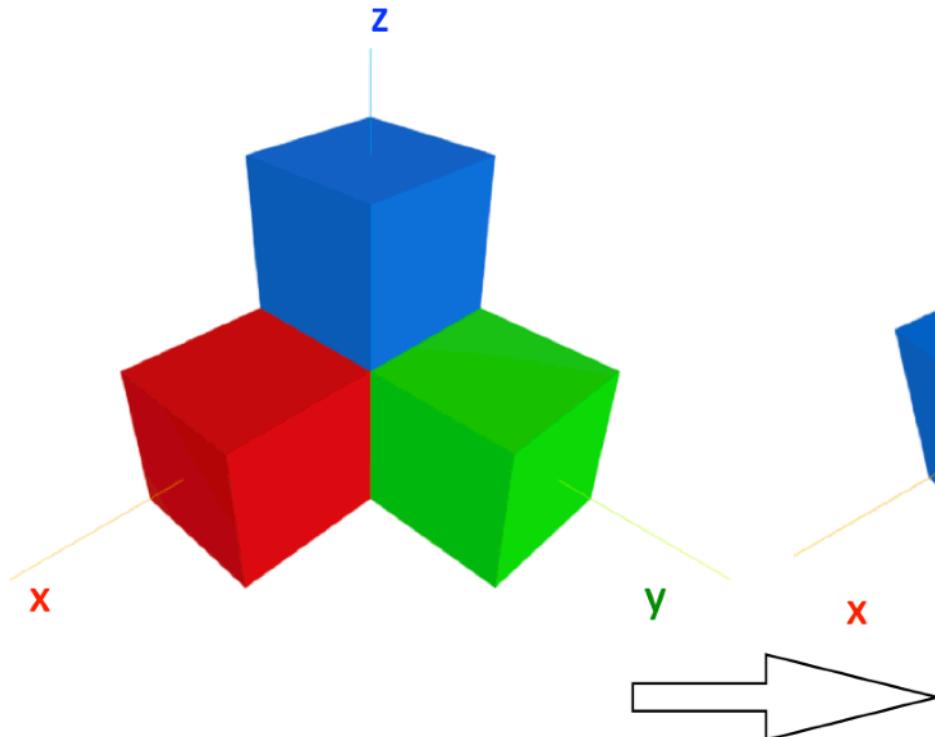
$$R_Z = \begin{bmatrix} \cos(\theta) & -\sin(\theta) & 0 & 0 \\ \sin(\theta) & \cos(\theta) & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$



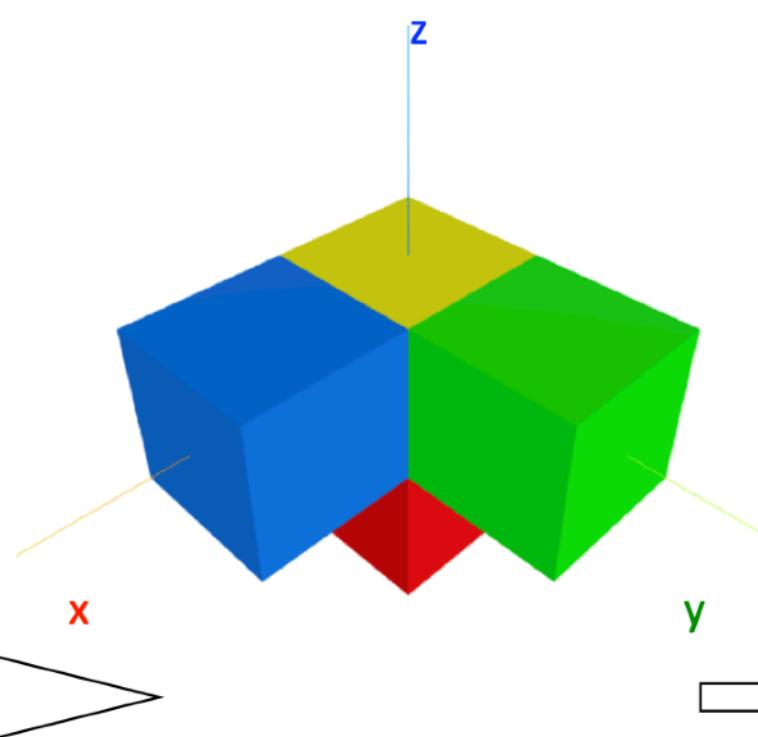
# Quiz: Rotation in 3D



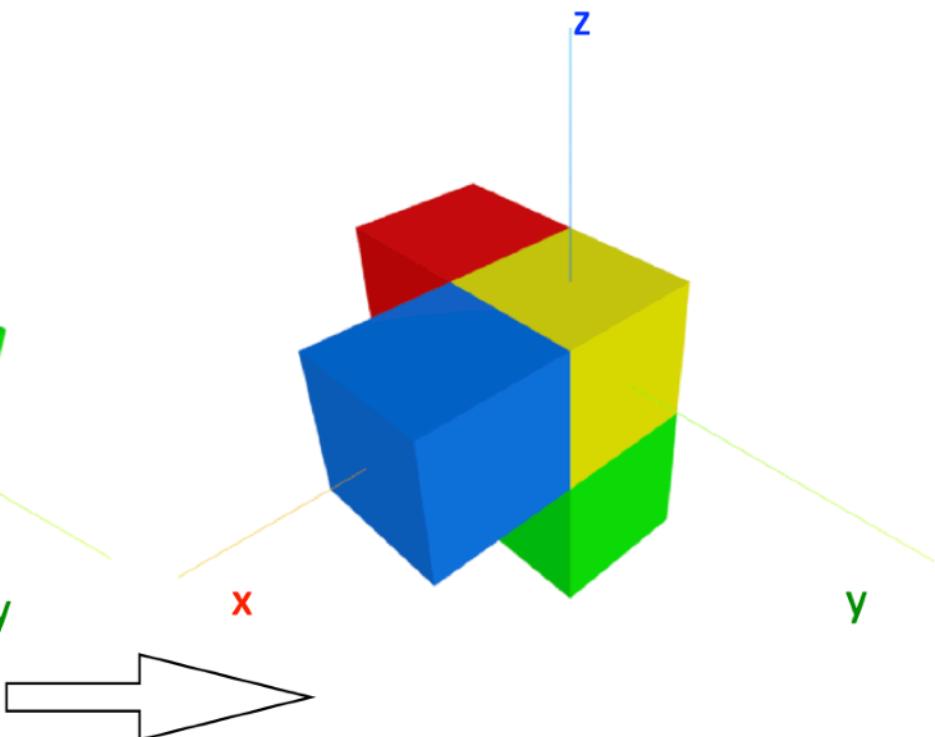
# Answer



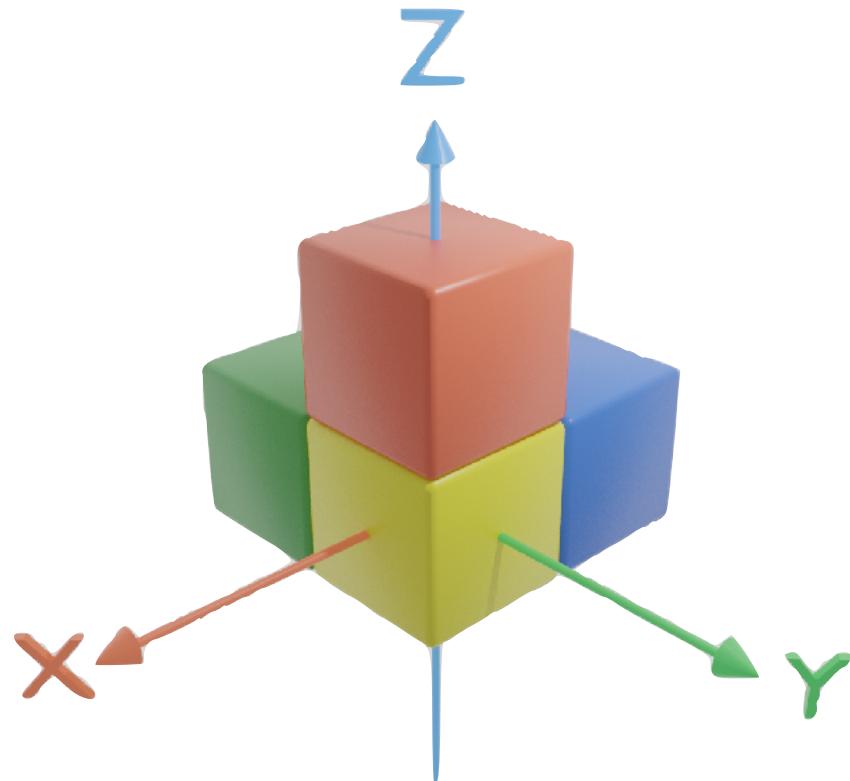
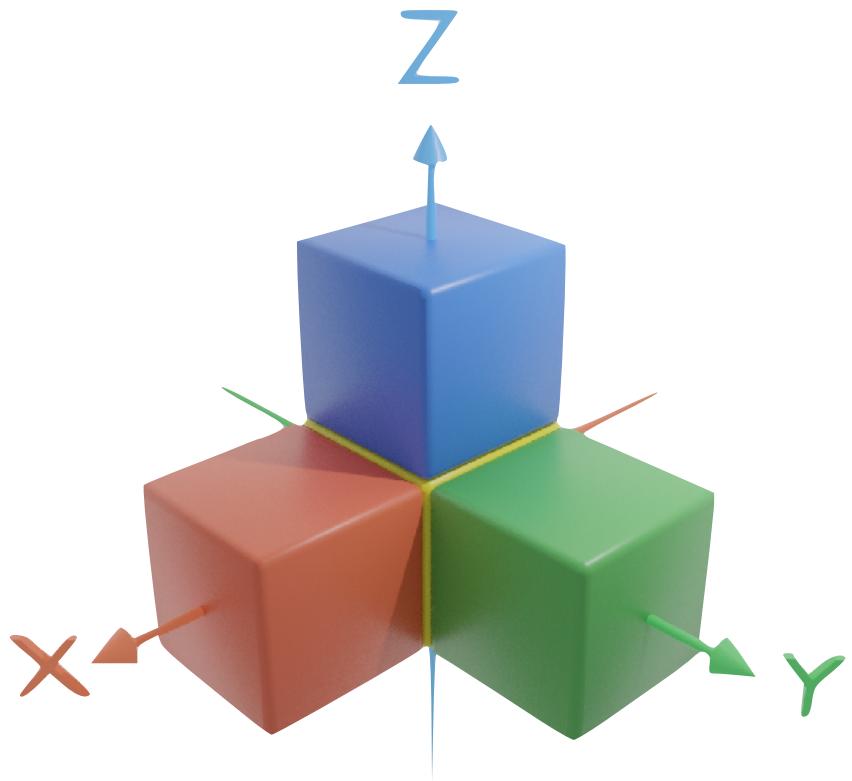
1. 90 Degree around y- Axis



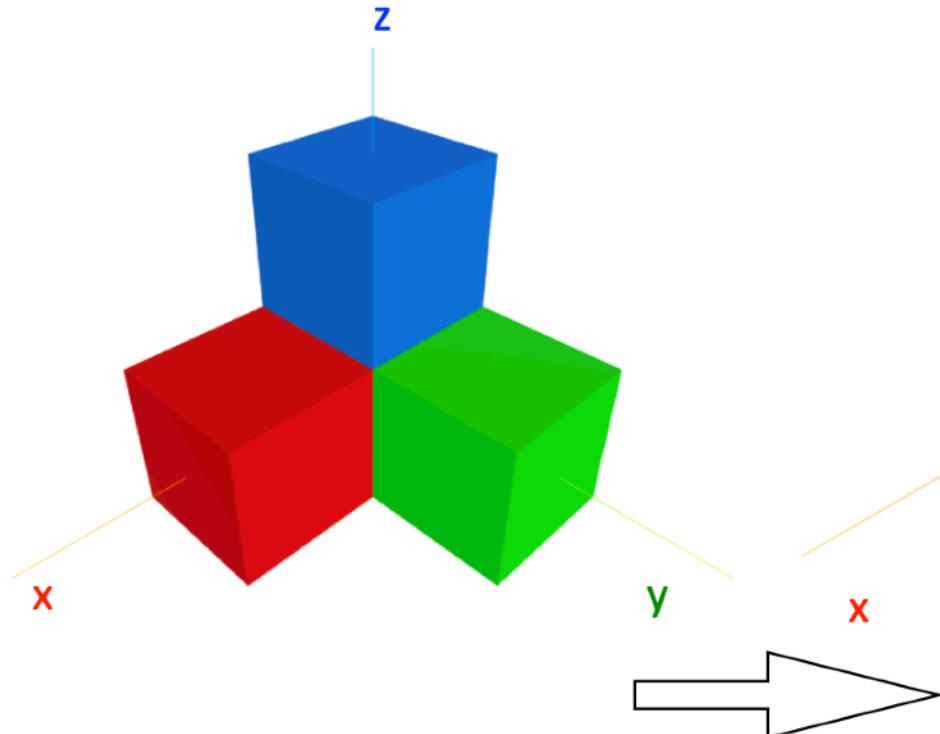
2. 270 Degree around y-Axis



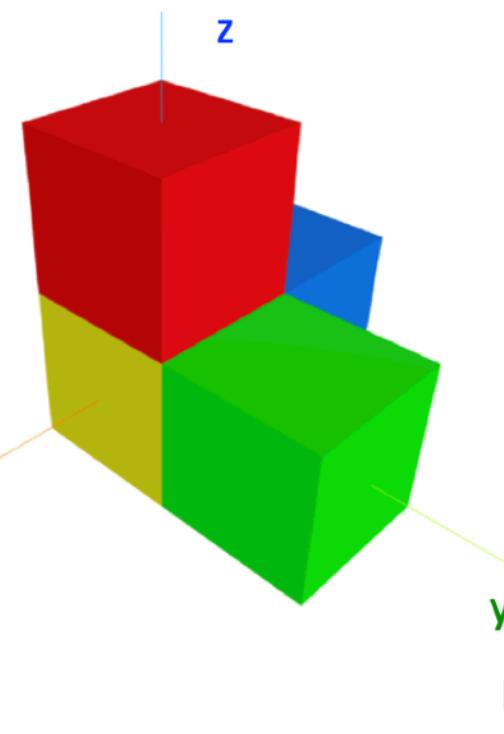
# Quiz: Rotation in 3D



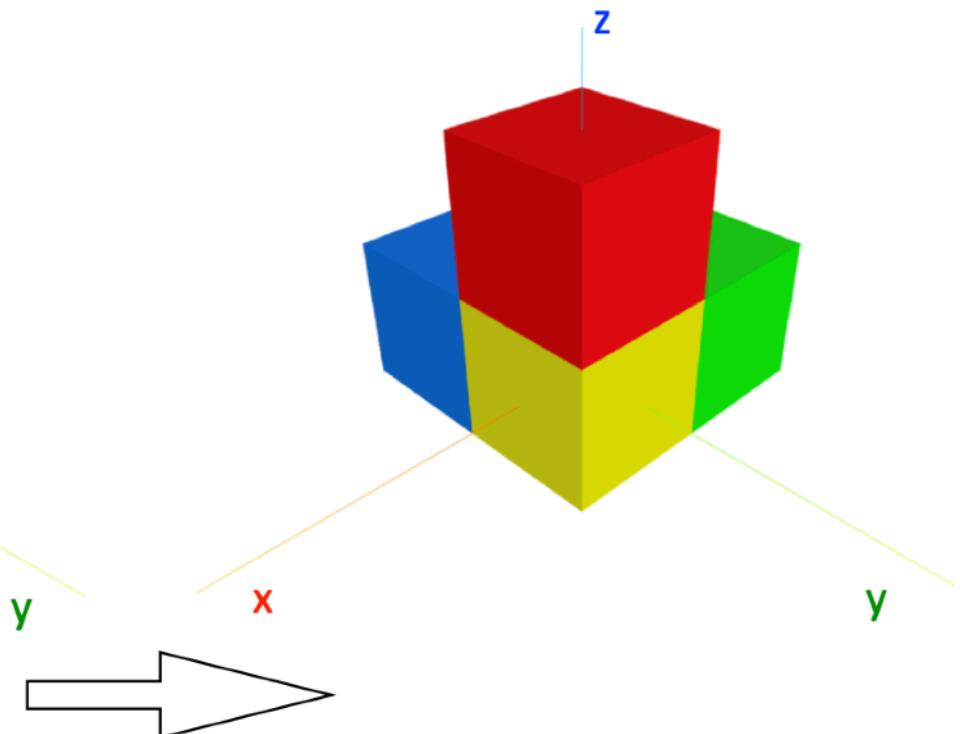
# Answer: Not Possible



1. 270 Degree around y- Axis



2. 90 Degree around z-Axis



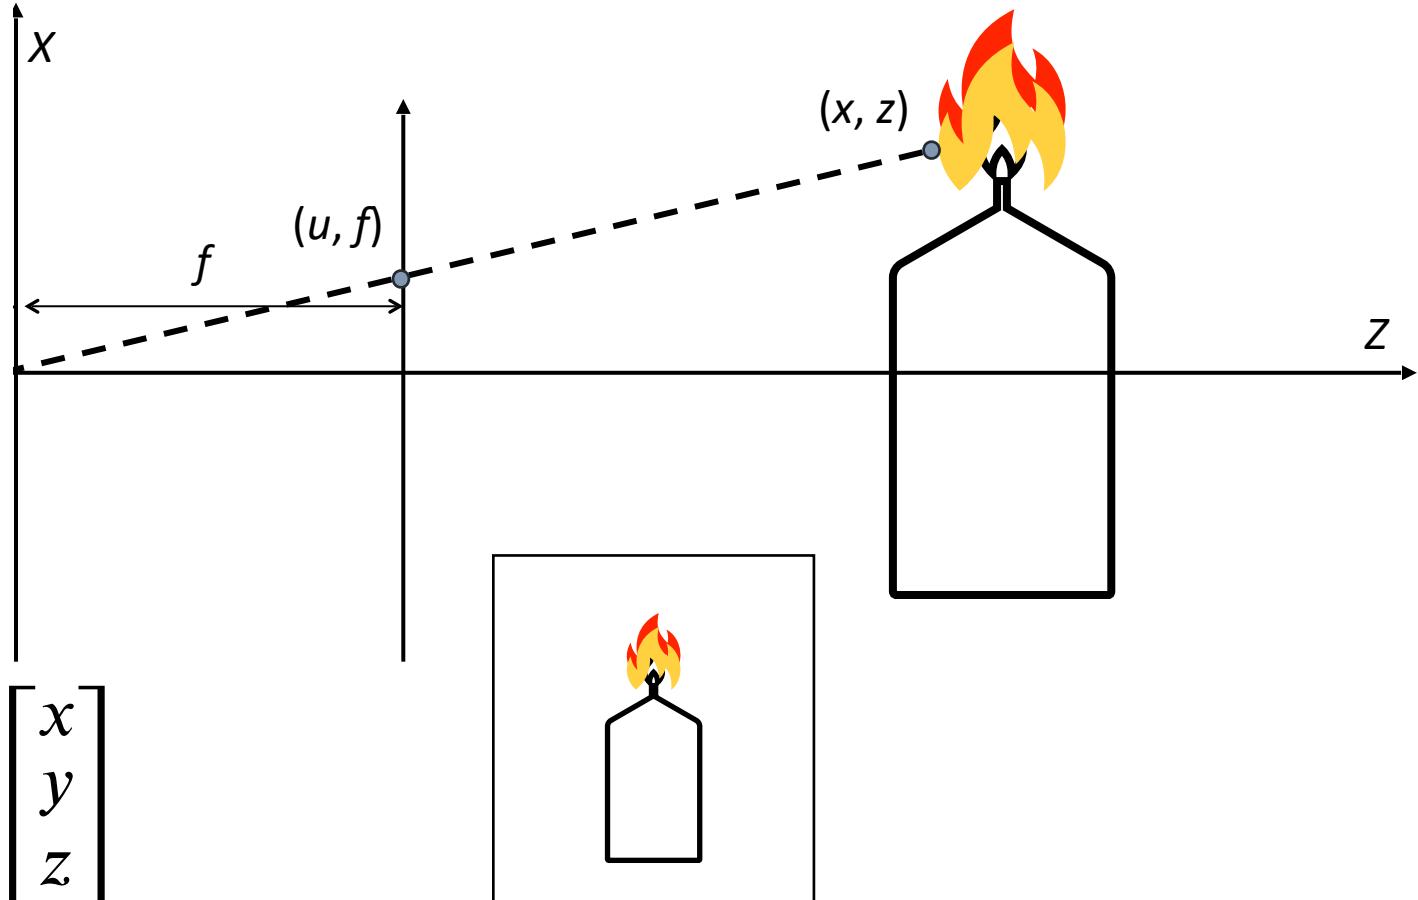
Flip needed

# The Pinhole Camera Model and Projective Geometry

- Removing the sign change

- We can put the image plane in front of the pinhole
  - Removes the sign change
  - Not practical for real cameras
  - The maths works out just fine

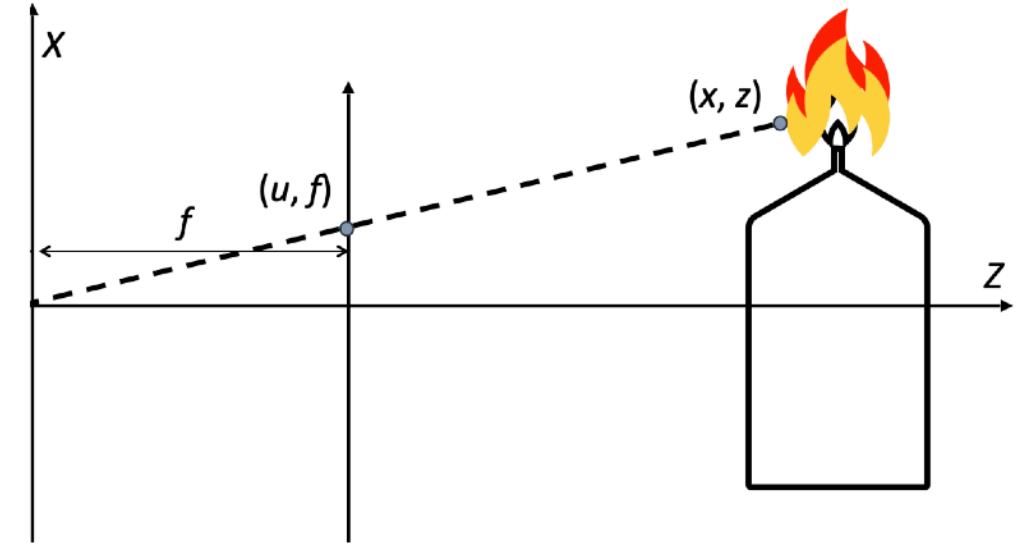
$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$



# **Camera Parameters and Calibration**

# Camera Parameters - Intrinsic and Extrinsic

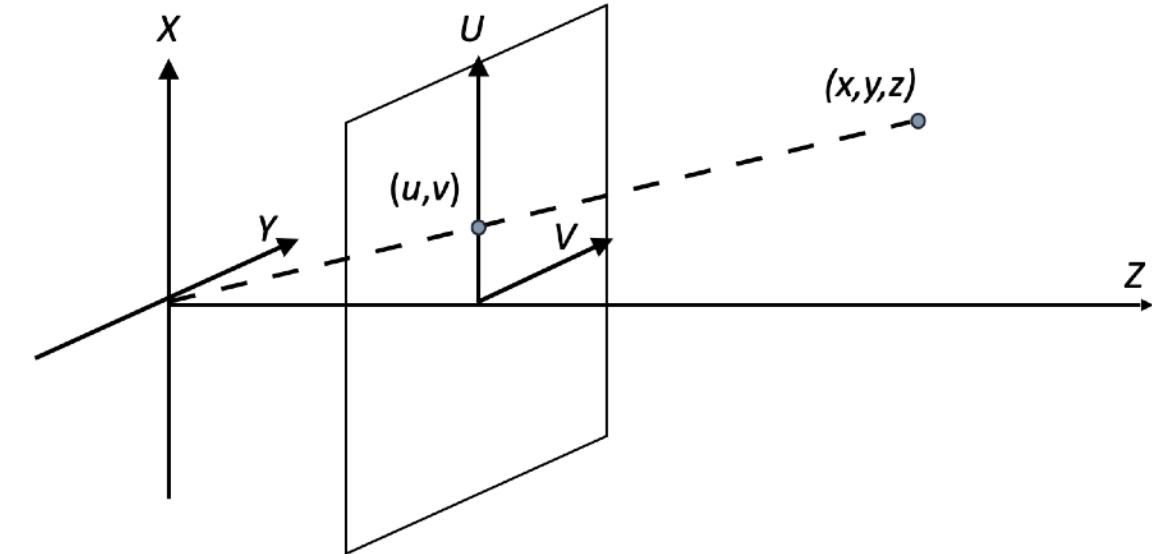
- Two main types of parameters:
  - Intrinsic parameters relate the camera's coordinate system to idealized coordinate system described earlier
  - Extrinsic parameters relate the camera's coordinate system to a world coordinate system
- Let's ignore lens issues for a moment
- Estimating intrinsic and extrinsic camera parameters is known as geometric camera calibration



# Camera Parameters - Intrinsics and Extrinsics

- The pinhole camera model
  - Projects from 2D to 3D
  - Models perspective projection well
  - Basic model:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \equiv \begin{bmatrix} f & 0 & 0 & 0 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

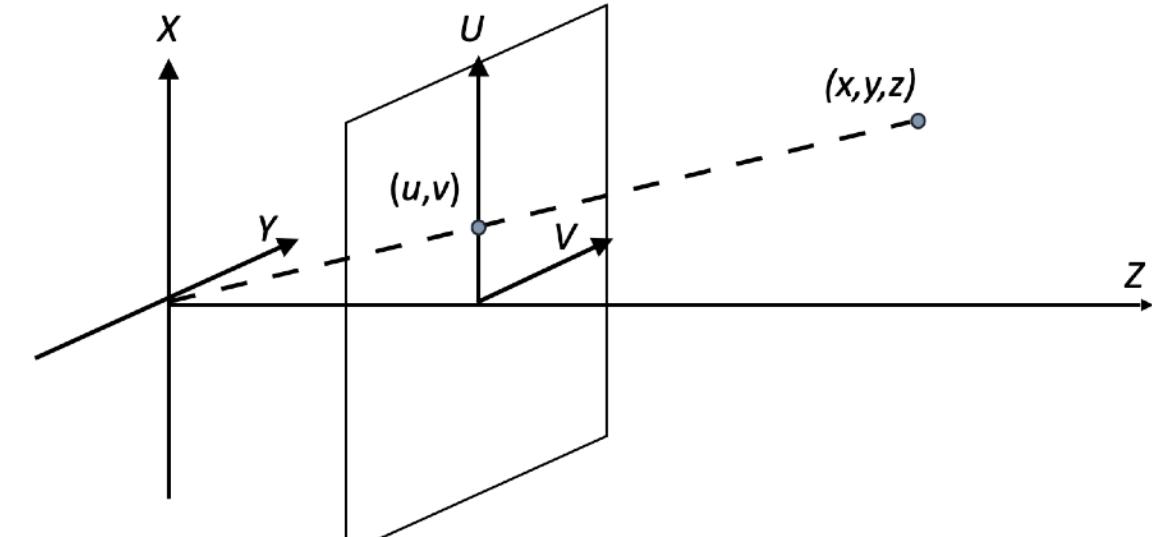


# Camera Parameters - Intrinsic and Extrinsic

- Often break this down into

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \equiv \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Intrinsic                              Extrinsic



- Most simple case:  $\mathbf{u} = \mathbf{K}[\mathbf{I} \mid \mathbf{0}]\mathbf{x}$
- $\mathbf{K}$ : camera calibration or Intrinsic parameters
- $[\mathbf{I} \mid \mathbf{0}]$ : camera pose or Extrinsic parameters

# Camera Parameters - Intrinsic and Extrinsic

- Model has image origin in the centre

- Remember in CV:

- We usually put this at the top left corner (0,0)

- Can fix this with a translation

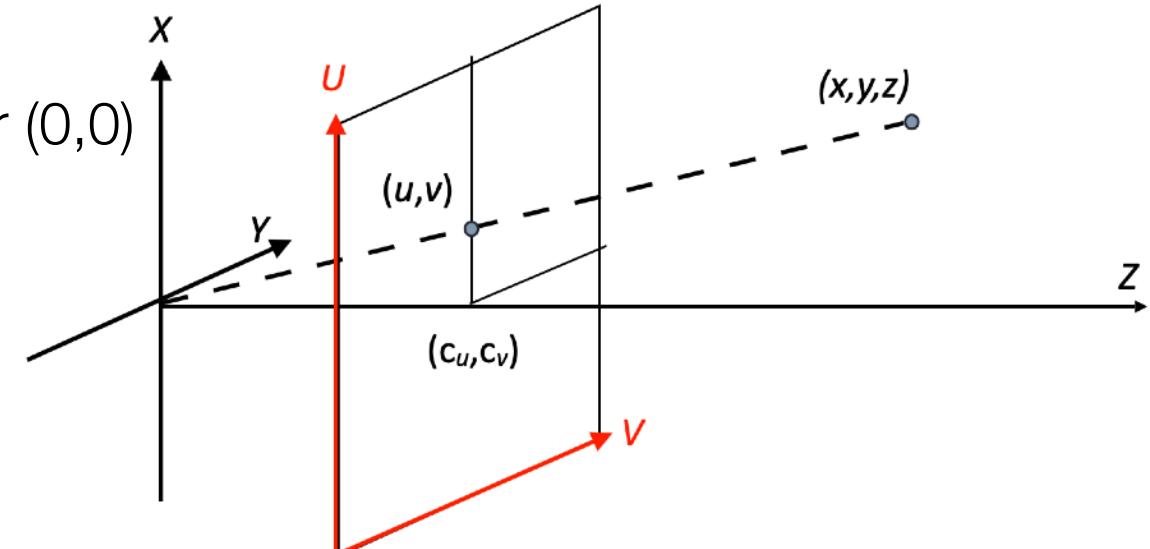
- If the centre is at  $(c_u, c_v)$  we get:

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \equiv \begin{bmatrix} f & 0 & c_u \\ 0 & f & c_v \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Intrinsic                      Extrinsic

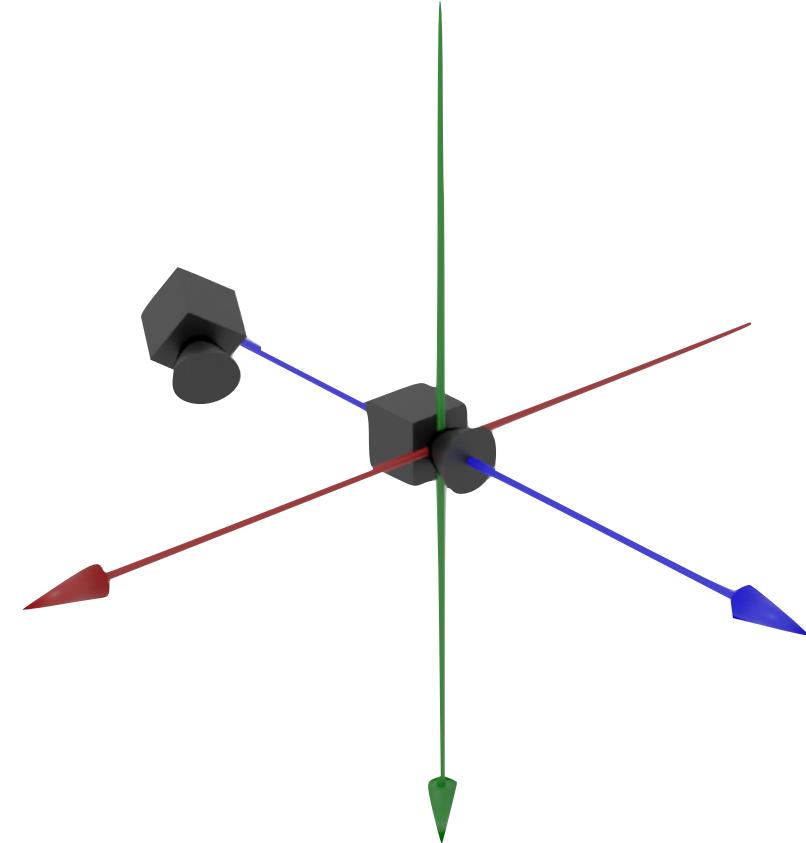
- $(c_u, c_v)$  is called principal point

- True geometric center of the image (For perfectly manufactured 1000x800 pixel sensor (500, 400) - Reality?)



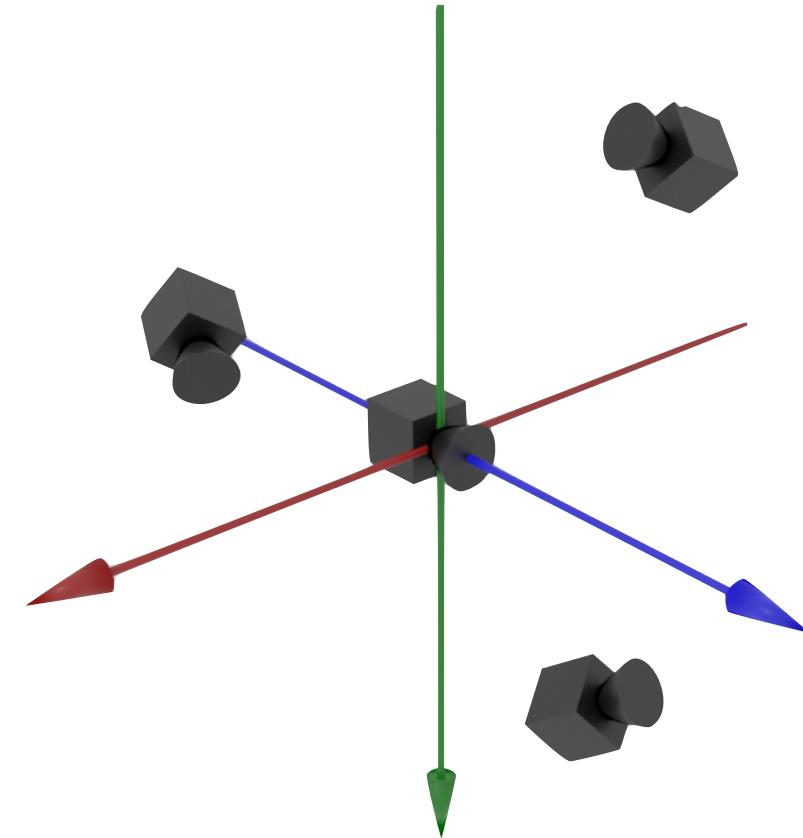
# Camera Parameters - Intrinsic and Extrinsic

- We have assumed
  - A camera at the origin
  - Pointing along the positive  $Z$  axis
- Let's consider the case where camera coordinate system C does not match the world coordinate system W
- So far: Most simple case:  $\mathbf{u} = K [I \mid \mathbf{0}] \mathbf{x}$
- We will need the general case:  $\mathbf{u} = K [R \mid \mathbf{t}] \mathbf{x}$ 
  - Move the camera to any location (Translation  $\mathbf{t}$ )
  - Point the camera in any direction (Rotation  $R$ )
- The **extrinsic parameters** consist of a  $3 \times 3$  rotation matrix R and a 3-dimensional translation vector t



# Camera Parameters - Intrinsic and Extrinsic

- To transform a camera by  $T$ 
  - Apply inverse,  $T^{-1}$ , to points
  - To move the camera left 3 units, move the world right 3 units
  - To rotate the camera  $45^\circ$  about  $Z$ , rotate the world  $-45^\circ$  about  $Z$
  - The relative motion of the camera and the world is the same



# Camera Parameters - Intrinsic and Extrinsic

- Want to rotate a camera by  $\mathbf{R}$ , then shift by  $\mathbf{t}$ 
  - Equivalently shift the points by  $-\mathbf{t}$ , then rotate them by  $\mathbf{R}^{-1} = \mathbf{R}^T$ :

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} f & 0 & c_u \\ 0 & f & c_v \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} r_{11} & r_{21} & r_{31} & 0 \\ r_{12} & r_{22} & r_{32} & 0 \\ r_{13} & r_{23} & r_{33} & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & -t_x \\ 0 & 1 & 0 & -t_y \\ 0 & 0 & 1 & -t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

- Or, in more compact form:

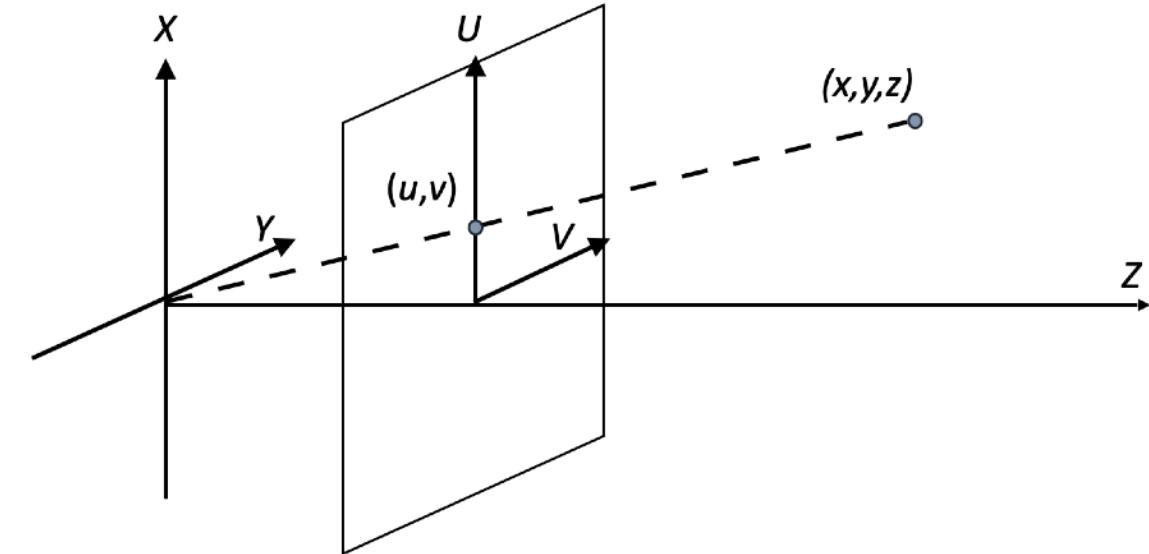
$$\mathbf{u} \equiv \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \begin{bmatrix} \mathbf{R}^T & \mathbf{0} \\ \mathbf{0}^T & 1 \end{bmatrix} \begin{bmatrix} \mathbf{I} & -\mathbf{t} \\ \mathbf{0}^T & 1 \end{bmatrix} \mathbf{x} = \mathbf{K} [\mathbf{R}^T \quad -\mathbf{R}^T \mathbf{t}] \mathbf{x}$$

# **Practical Camera Calibration**

# Pinhole Camera Review

$$\mathbf{u} \equiv \mathbf{K}[\mathbf{R} \quad \mathbf{t}]\mathbf{x}$$

- $\mathbf{u} = [u \quad v \quad 1]^T$  is a 2D image point
- $\mathbf{K}$  is a  $3 \times 3$  calibration matrix
- $\mathbf{R}$  is a  $3 \times 3$  rotation matrix
- $\mathbf{t}$  is a 3D translation vector
- $\mathbf{x} = [x \quad y \quad z \quad 1]^T$  is a 3D point
- Camera calibration: determine  $\mathbf{K}$

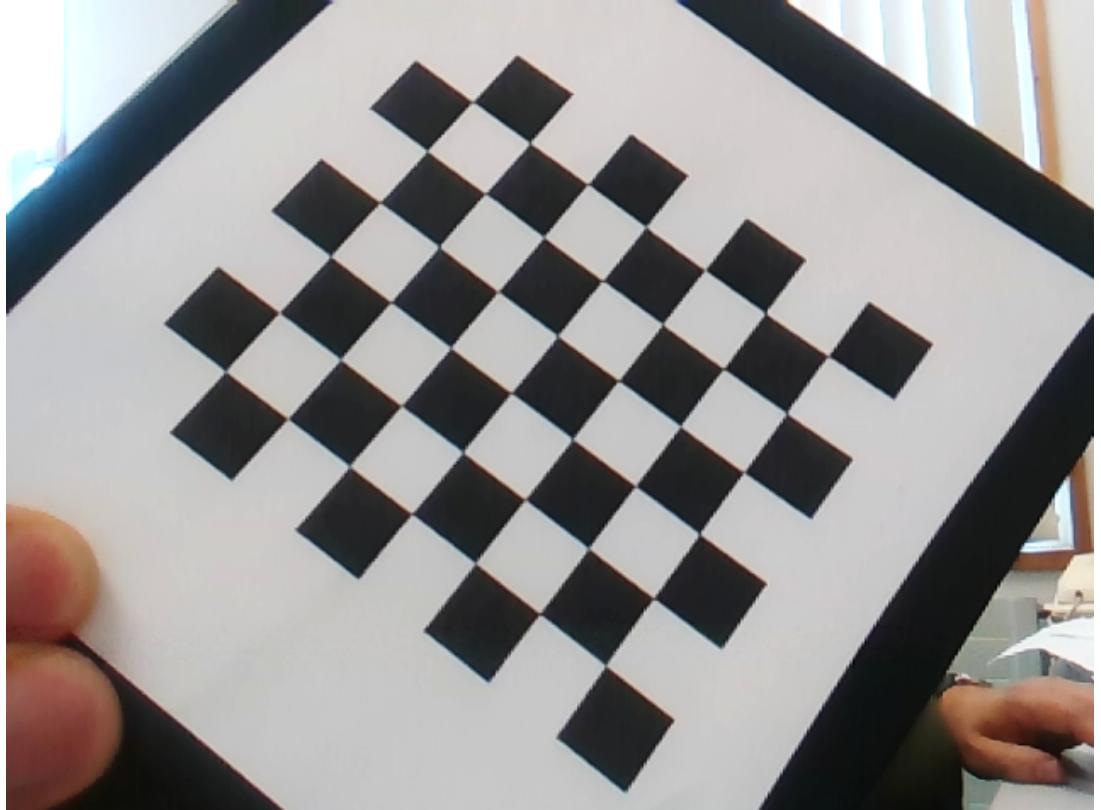


# Calibration in OpenCV

```
double cv::calibrateCamera(  
    // Input parameters  
    //std::vector<std::vector<cv::Point3f>> objectPoints,  
    //std::vector<std::vector<cv::Point2f>> imagePoints,  
    //cv::Size imageSize,  
    // Output parameters  
    //cv::Mat cameraMatrix,  
    //std::vector<double> distCoeffs,  
    //std::vector<cv::Mat> rotationVectors,  
    //std::vector<cv::Mat> translationVectors);
```

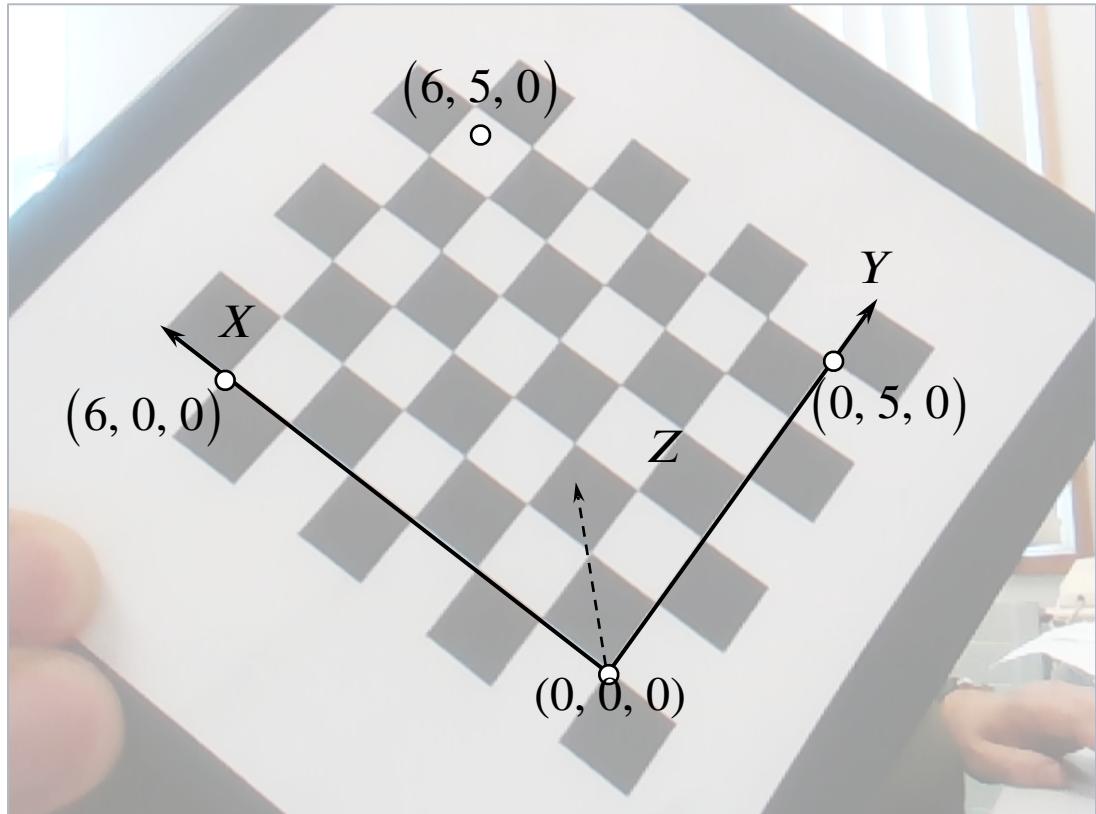
# Calibration in OpenCV

- Calibration targets:
  - Input is 3D-2D matches
  - Want easy-to find 2D points
  - Need known 3D co-ordinates
- Planar targets common
  - Easy to make with a printer
  - Chess/Checkerboards
  - Grids of dots or lines
- Is a 2D pattern enough?



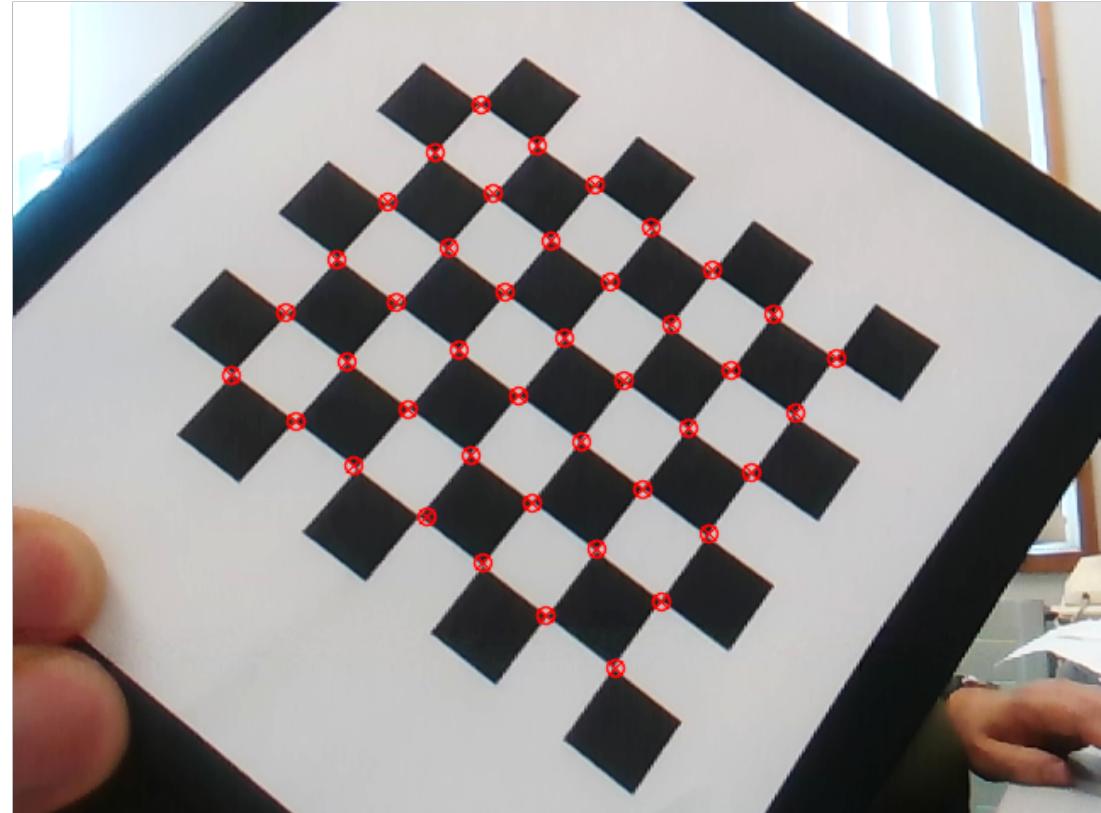
# Calibration in OpenCV

- 3D Point Locations
- Choice of co-ordinates
  - $X$ - $Y$  plane is the target
  - Origin at one internal corner
  - $Z$  runs into the target plane
- Need to decide on units
  - Best to use real-world units
  - Here 1 square = 1cm
  - Can use 1 square = 1 unit



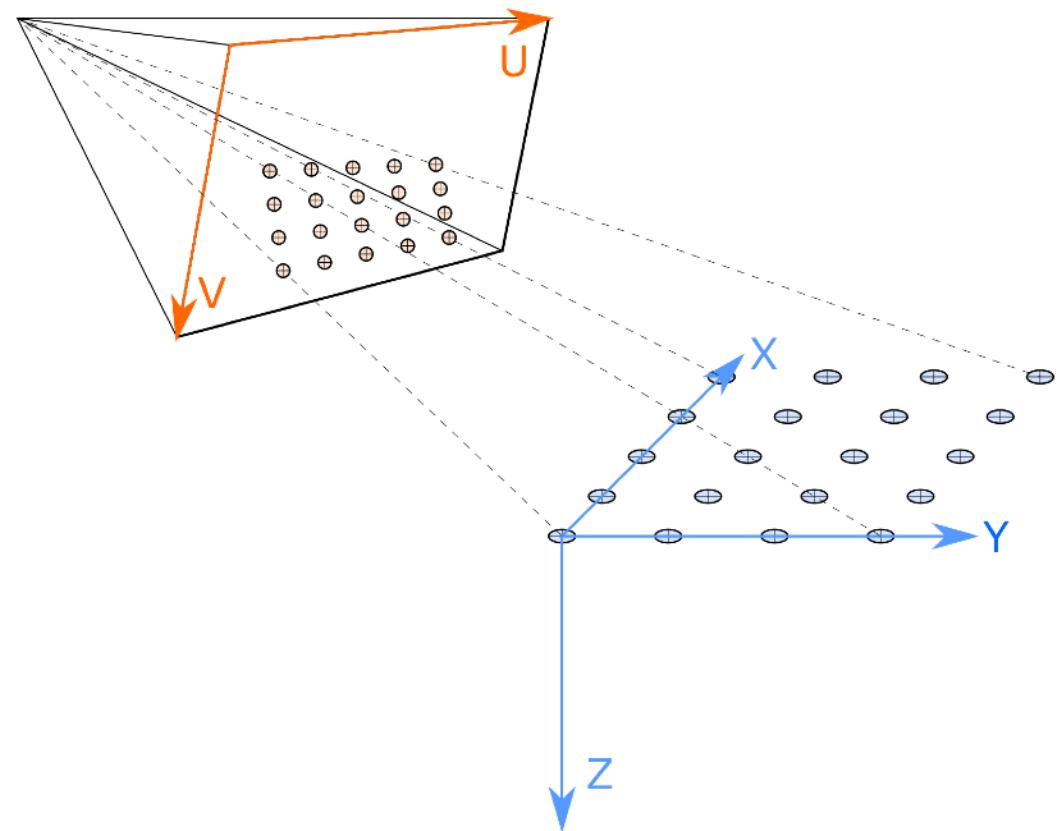
# Calibration in OpenCV

- Finding Checkerboard Corners
- OpenCV method (in brief)
  - Threshold image to black and white
  - Look for quads
  - Link quads → checkerboard
  - Sub-pixel refinement



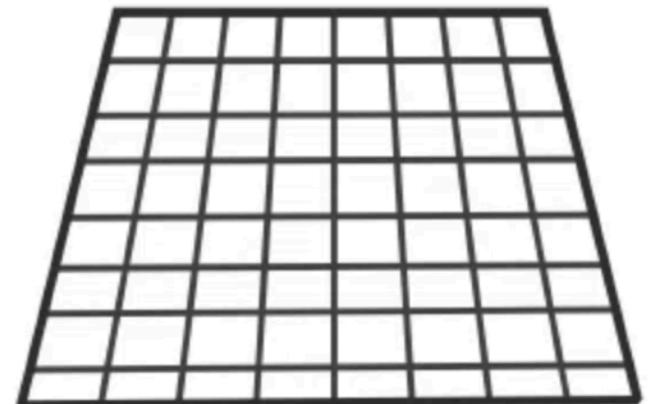
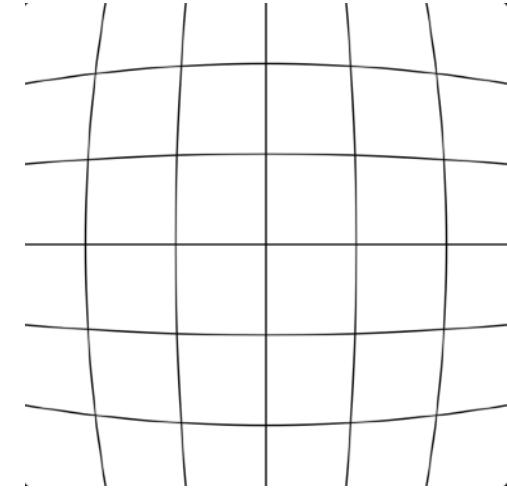
# Camera Calibration Problem

- Have  $k$  images,  $i = 1 \dots k$ 
  - Each has a set of 2D points
$$\mathbf{u}_{i,j} = [u_{i,j} \ v_{i,j} \ 1]^T$$
  - These match 3D points
$$\mathbf{x}_{i,j} = [x_{i,j} \ y_{i,j} \ z_{i,j} \ 1]^T$$
  - Related by
$$\mathbf{u}_{i,j} \equiv \mathbf{K}[\mathbf{R}_i \ \mathbf{t}_i] \mathbf{x}_{i,j}$$
- Find  $\mathbf{K}$ , get  $\mathbf{R}_i$ s and  $\mathbf{t}_i$ s too



# Modelling Lens Distortion

- Lens Distortion: Real lenses are imperfect and cause lens distortion, making straight lines in the world appear curved in the image
- Two main types:
  - Radial distortion (barrel or pincushion effect, where lines curve towards/away from the image center)
  - Tangential distortion (when the lens is not perfectly parallel to the sensor).
- The function calculates a set of distortion coefficients ( $k_1, k_2, p_1, p_2, k_3\dots$ ) that mathematically describe how to correct for these curves



# Modelling Lens Distortion

- Barrel/pincushion distortion

$$\mathbf{u}' = \mathbf{u}(1 + k_1 r^2 + k_2 r^4 + \dots)$$

- $(u, v)$  measured from image centre

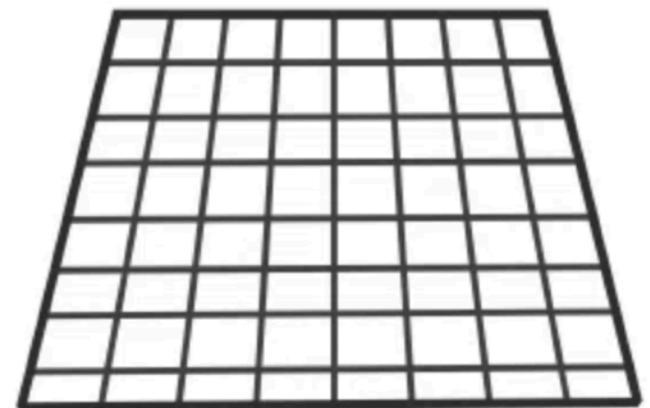
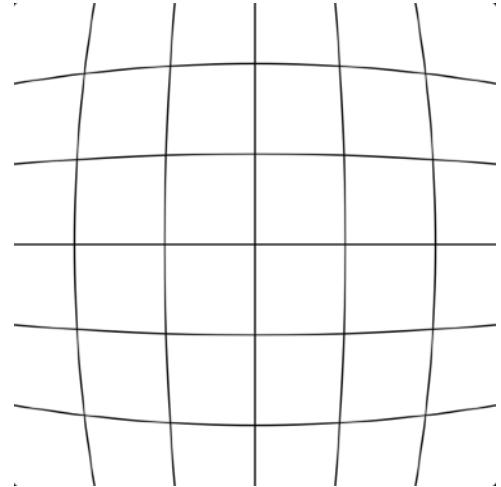
- $r = \sqrt{u^2 + v^2}$  : distance from centre

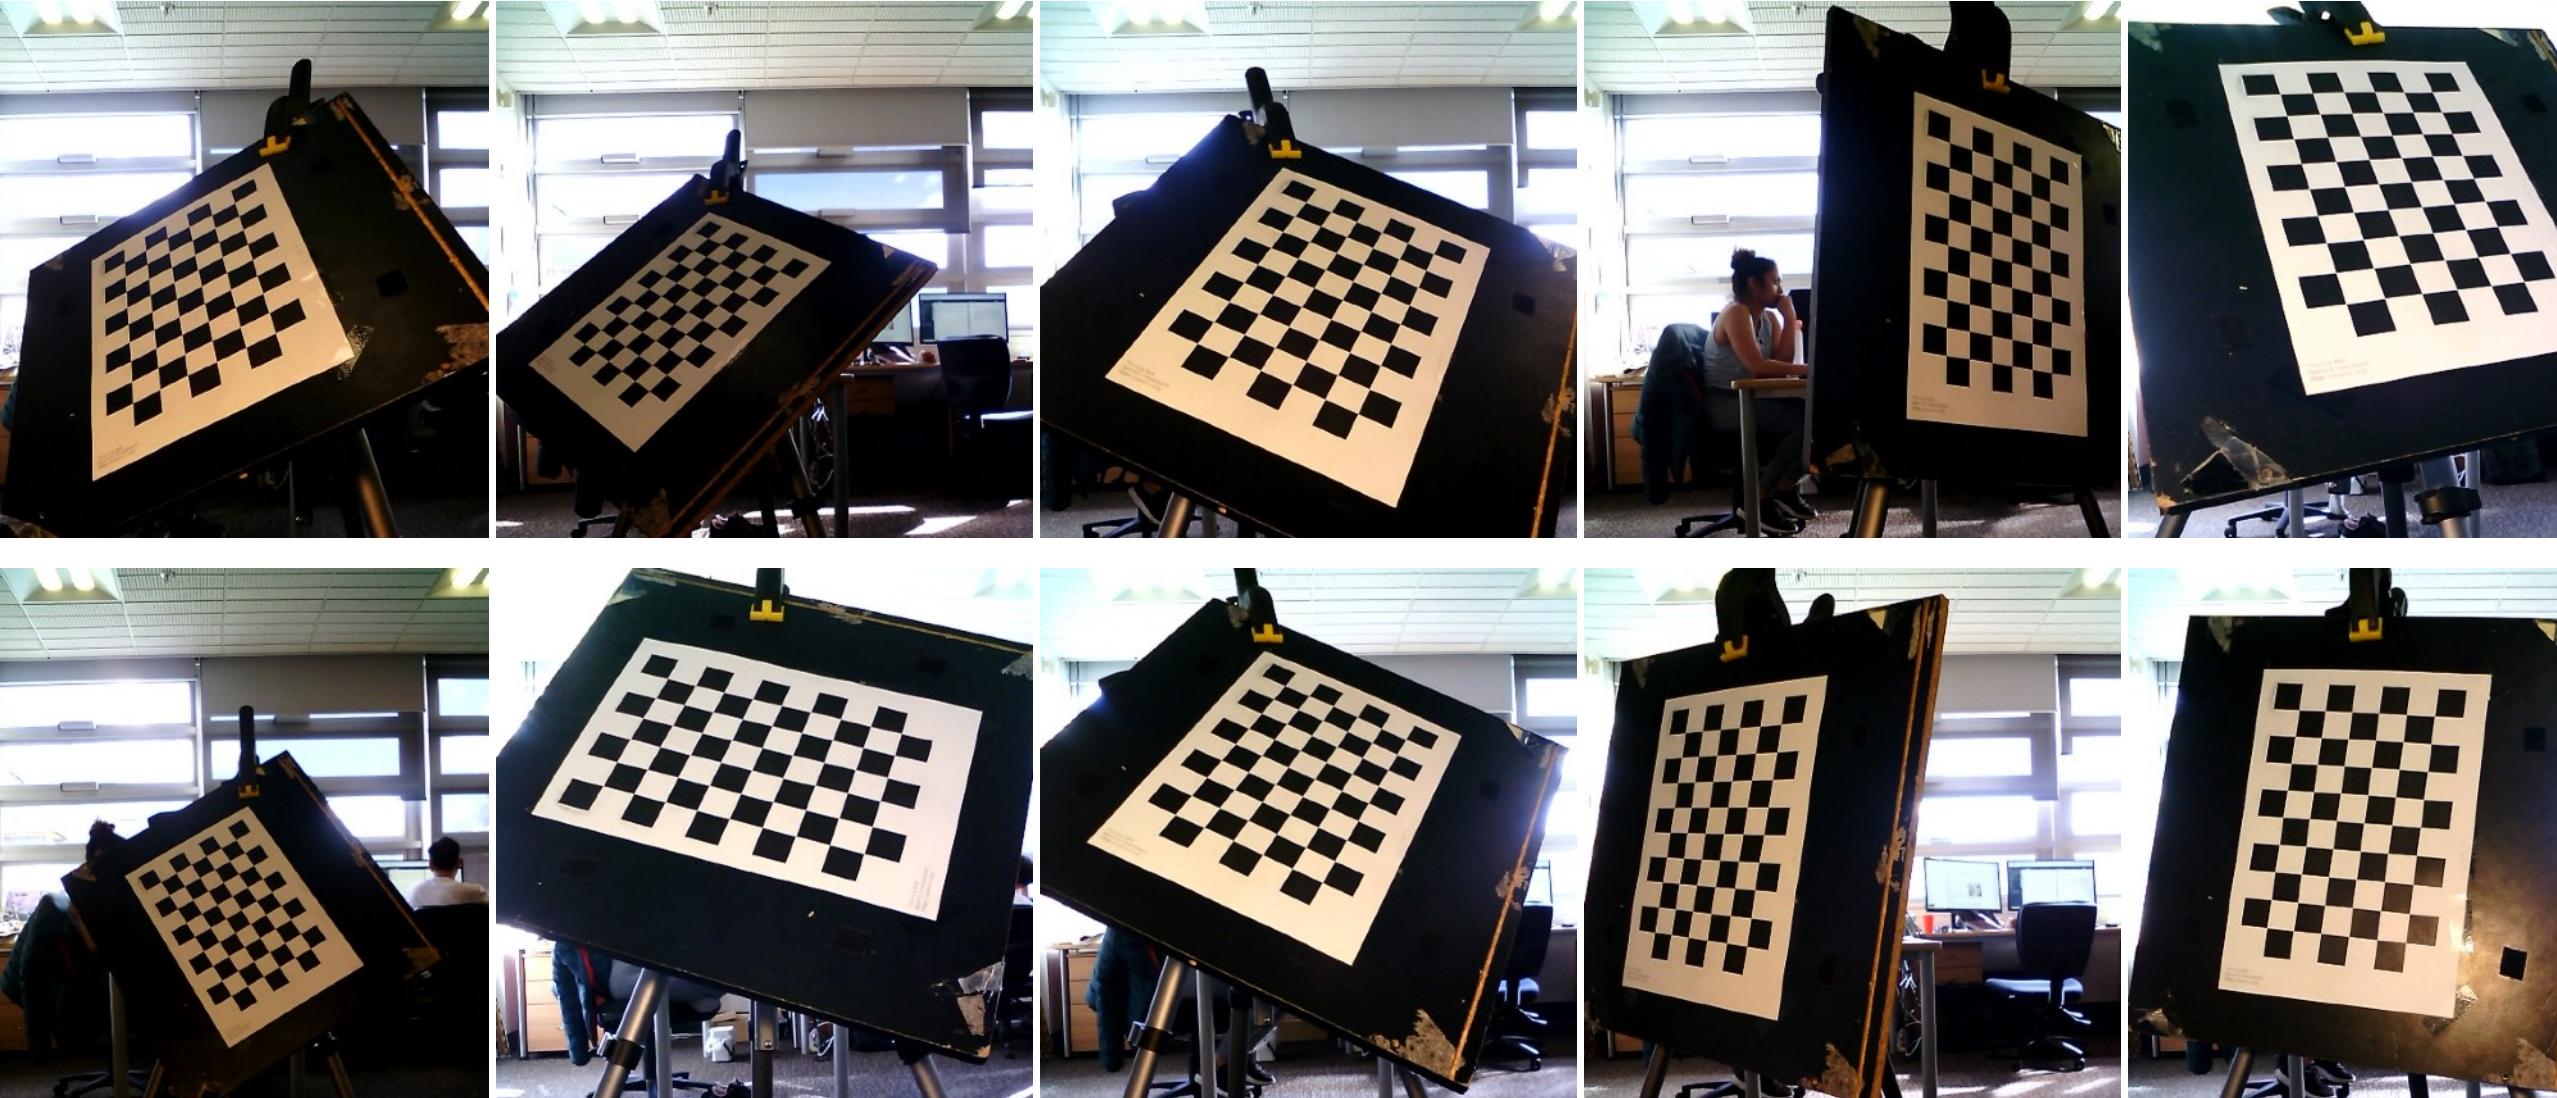
- Tangential distortion

- Lens not parallel to image plane

$$u' = u + (2p_1 uv + p_2(r^2 + 2u^2))$$

$$v' = v + (2p_2 uv + p_1(r^2 + 2u^2))$$





# **The Geometry of Multiple Views and Scene Reconstruction**

# The Geometry of Multiple Views and Scene Reconstruction

- Scene Reconstruction
- Epipolar Geometry
  - Epipolar Constraint
  - Baseline, Epipolar Plane, Epipoles, Epipolar Lines
- The Uncalibrated Case
  - Fundamental Matrix
  - Eight-Point Algorithm
  - Seven-Point Algorithm
- The Calibrated Case
  - Essential Matrix
- More cameras -> Multi-View Geometry

# Scene Reconstruction

- So far, we have estimated the parameters of a single camera
- The depth of a scene point is not directly accessible in a single image
- What can we do?
  - With at least two pictures, depth can be measured through triangulation
  - Parameters of all cameras have to be known
  - Search for point correspondence!



# **Stereo Computation**

# A Simple Stereo System

- Assume a set of 2 pinhole cameras

$$\mathbf{u} = \mathbf{K}[\mathbf{R} \quad \mathbf{t}]\mathbf{x}$$

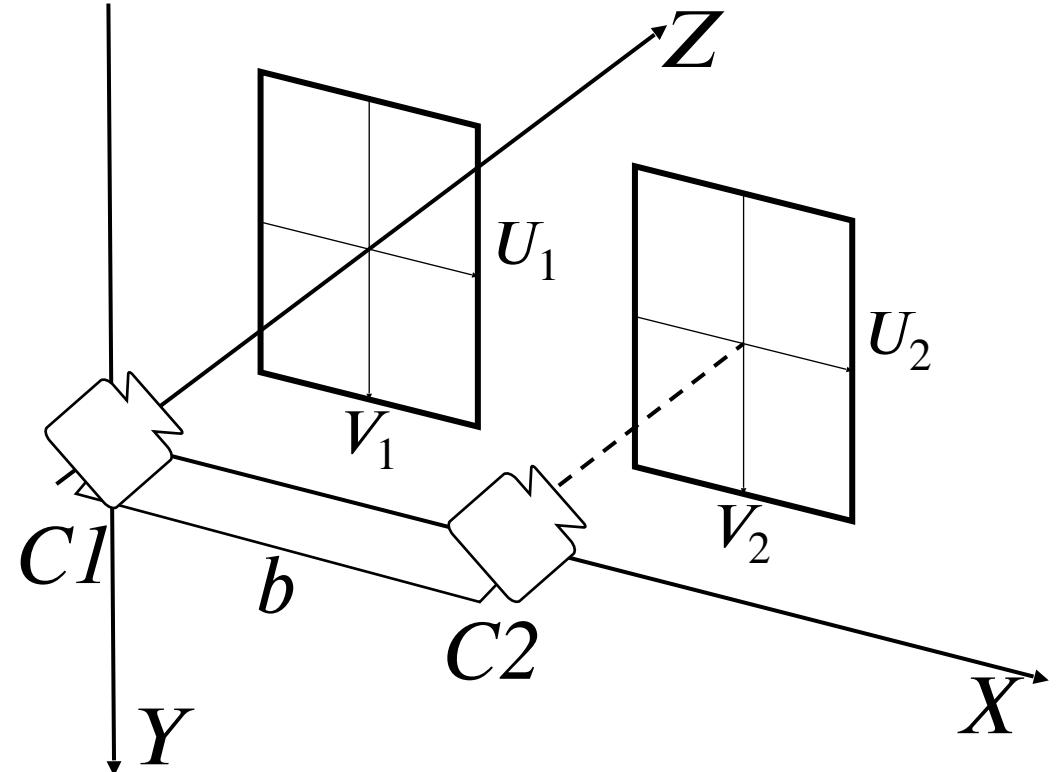
- Simplify for now:  $\mathbf{K} = \mathbf{I}$

- Camera 1 (C1):  $\mathbf{R} = \mathbf{I}, \mathbf{t} = \mathbf{0}$

- No rotation, looks along  $Z$

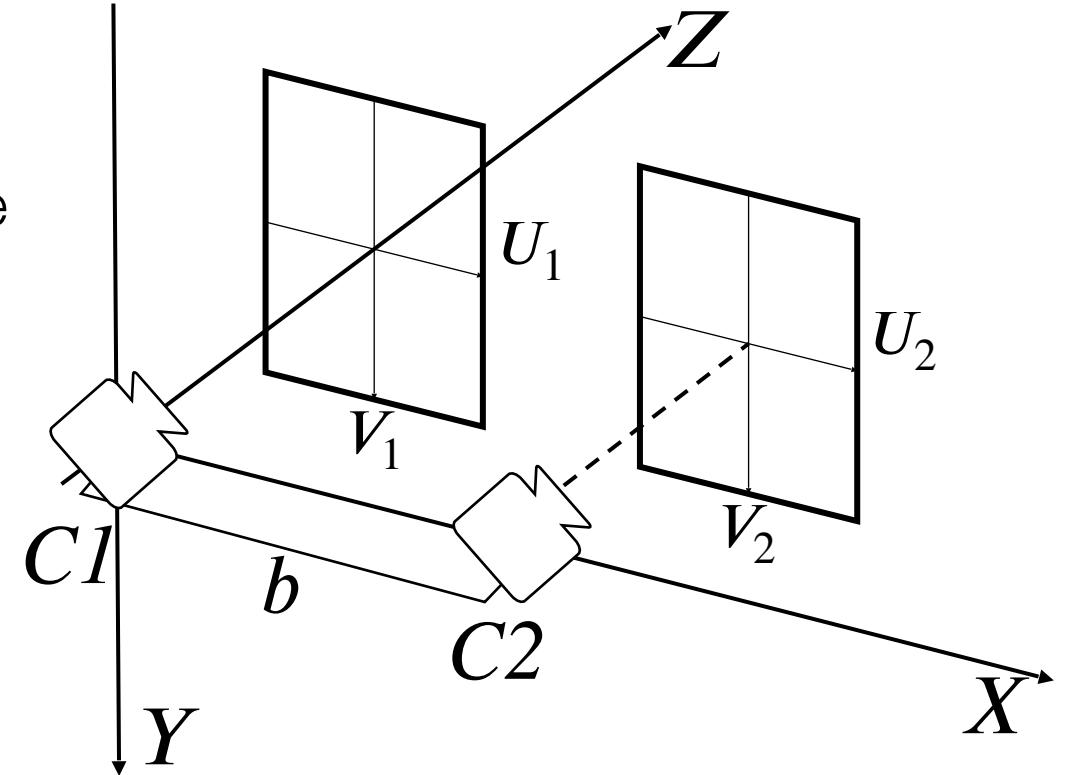
- Camera 2 (C2): shifted along  $X$

$$\mathbf{R} = \mathbf{I}, \mathbf{t} = \begin{bmatrix} b \\ 0 \\ 0 \end{bmatrix}$$



# A Simple Stereo System

- Perfectly Parallel: The principal axes (the lines pointing straight out from the center of the lenses) of both cameras are parallel to each other and perpendicular to the baseline
- Perfectly Aligned: Image planes of the two cameras lie on the same plane. This means there is no vertical shift between the cameras; their y-axes are aligned
- Known Baseline: The distance between the centres of the two cameras, called the baseline ( $b$ ), is precisely known



# Project a Point into the Cameras

- Camera 1:

$$\begin{bmatrix} u_1 \\ v_1 \\ 1 \end{bmatrix} \equiv \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

$$= \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} \equiv \begin{bmatrix} x/z \\ y/z \\ 1 \end{bmatrix}$$

- Camera 2:

$$\begin{bmatrix} u_2 \\ v_2 \\ 1 \end{bmatrix} \equiv \begin{bmatrix} 1 & 0 & 0 & -b \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

$$= \begin{bmatrix} x - b \\ y \\ z \\ 1 \end{bmatrix} \equiv \begin{bmatrix} (x - b)/z \\ y/z \\ 1 \end{bmatrix}$$

# Depth and Disparity

- Difference between the views is the disparity

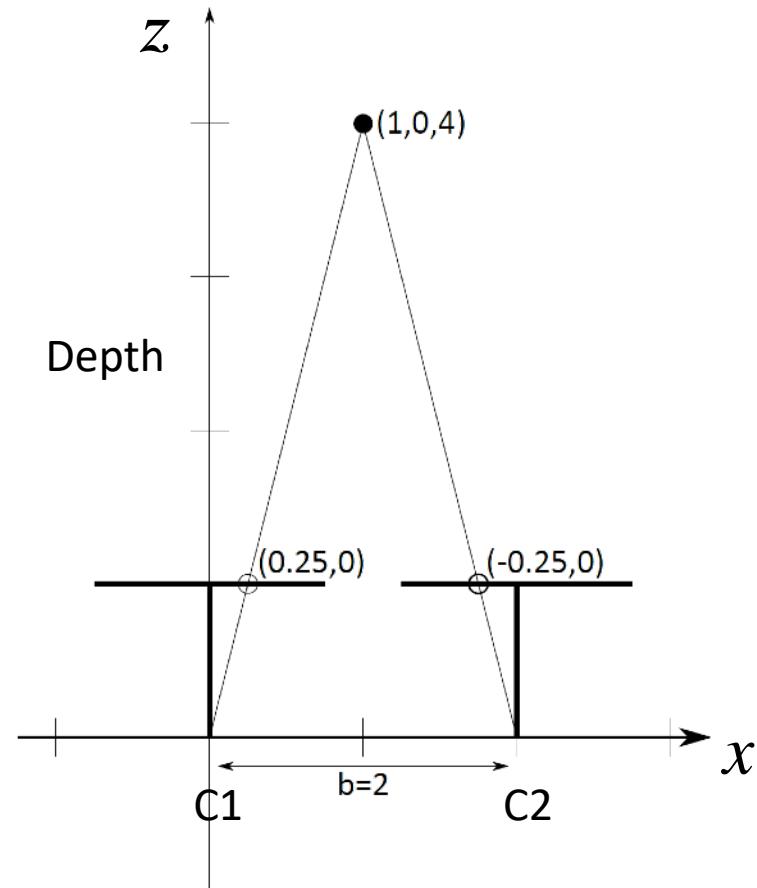
- In this case it is in  $u$ :

$$u_1 - u_2 = \frac{x}{z} - \frac{x-b}{z} = \frac{b}{z}$$

- The  $v$  values are the same
- If we know  $b$ , can find  $z$ :

$$z = \frac{b}{u_1 - u_2}$$

Top Down View



# Depth and Disparity



<https://juniorxsound.github.io/THREE.SixDOF/examples/video.html>

# More General Camera Model

- Camera 1:

$$\begin{bmatrix} u_1 \\ v_1 \\ 1 \end{bmatrix} \equiv \begin{bmatrix} f & 0 & c_u & 0 \\ 0 & f & c_v & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

$$= \begin{bmatrix} fx + c_u z \\ fy + c_v z \\ z \end{bmatrix} \equiv \begin{bmatrix} \frac{fx}{z} + c_u \\ \frac{fy}{z} + c_v \\ 1 \end{bmatrix}$$

- Camera 2:

$$\begin{bmatrix} u_2 \\ v_2 \\ 1 \end{bmatrix} \equiv \begin{bmatrix} f & 0 & c_u & -fb \\ 0 & f & c_v & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

$$\equiv \begin{bmatrix} \frac{f(x-b)}{z} + c_u \\ \frac{fy}{z} + c_v \\ 1 \end{bmatrix}$$

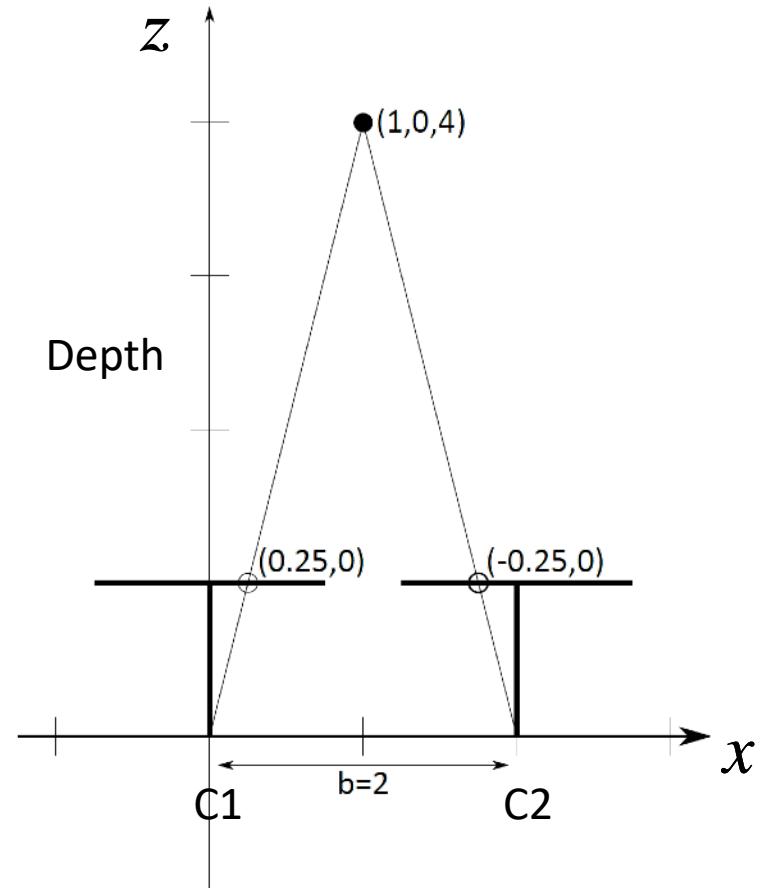
# Depth Estimation

Top Down View

- And we solve for  $z$ :

$$u_1 - u_2 = \frac{fx}{z} - \frac{f(x-b)}{z}$$

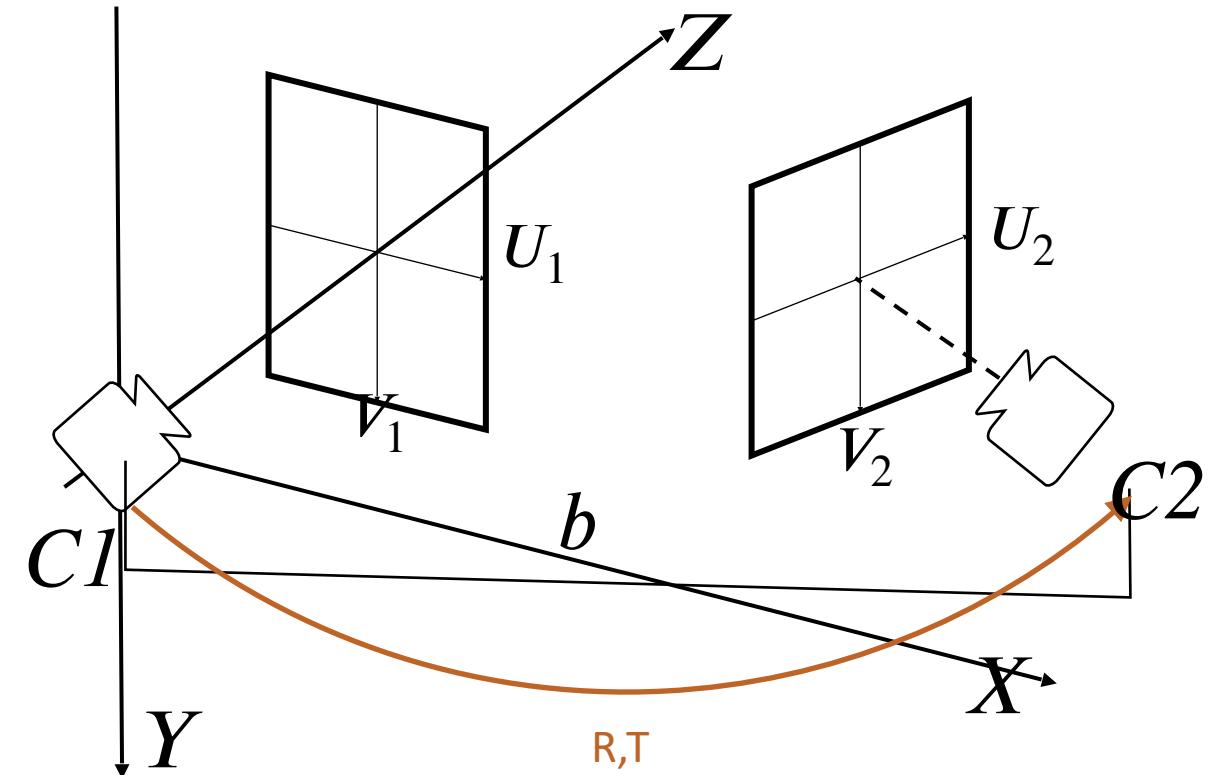
$$z = \frac{fb}{u_1 - u_2}$$



# **More General Stereo**

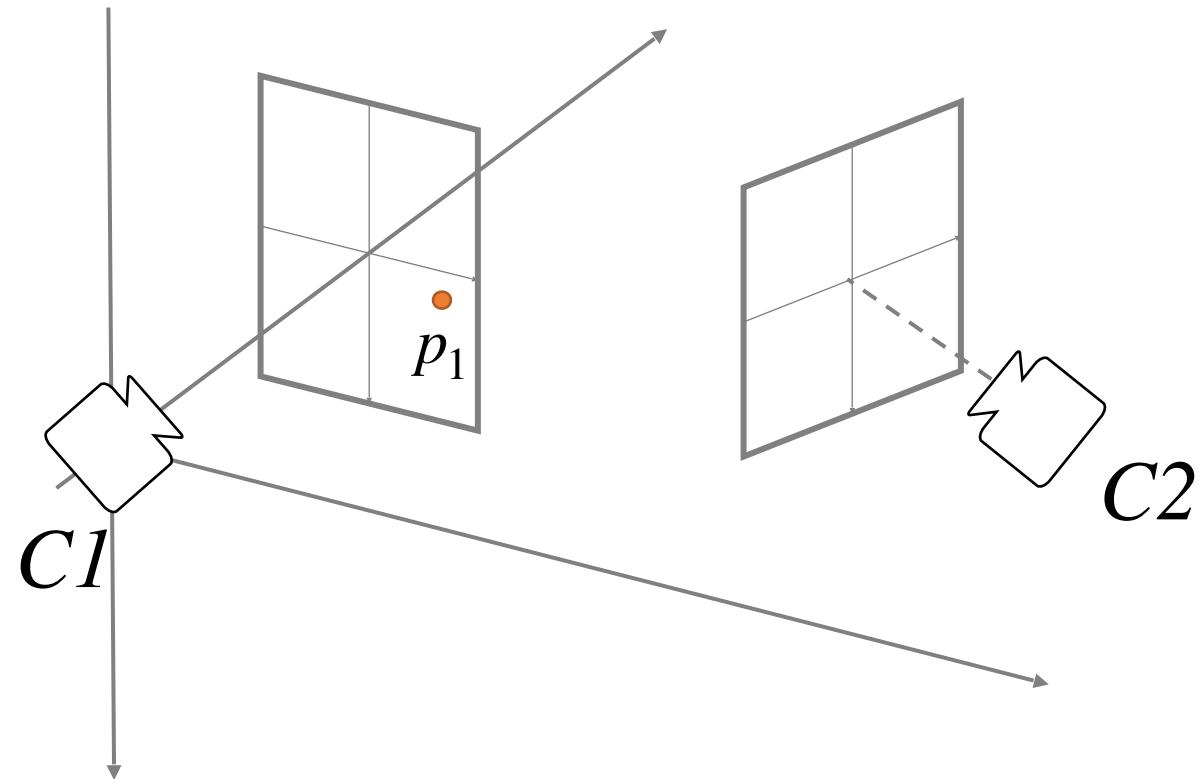
# More General Stereo

- In the more general case:
  - The cameras have different calibration parameters
  - There is a rotation between the cameras
  - There is translation along all three axes
- We can still align the XYZ coordinate system with camera C1



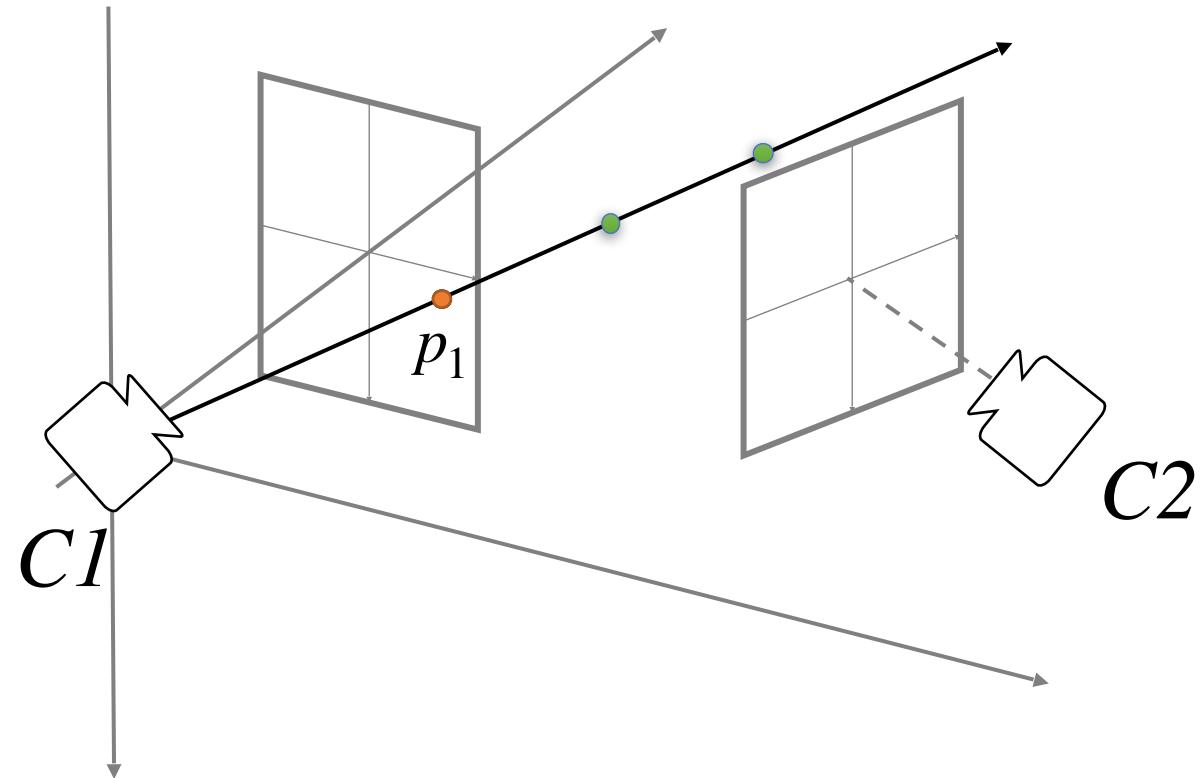
# Epipolar Geometry

- Consider a point  $p_1$  in the first camera



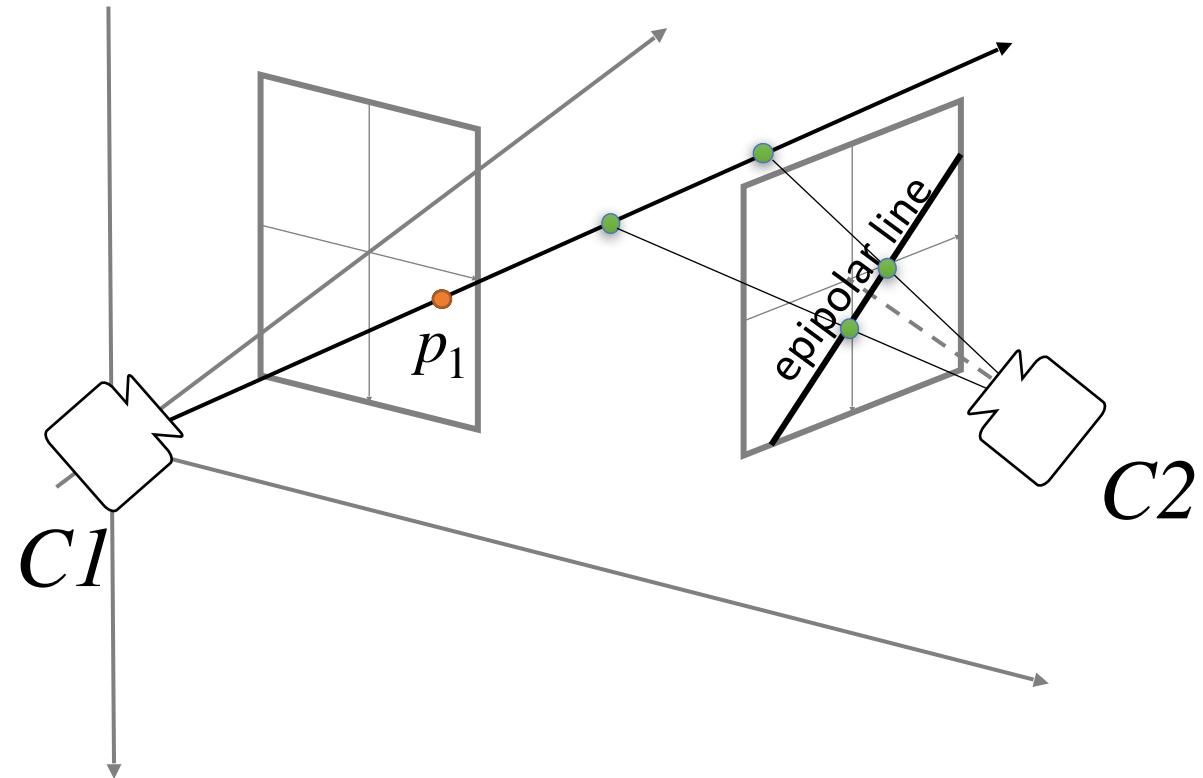
# Epipolar Geometry

- Consider a point  $p_1$  in the first camera
  - This gives a 3D line, along which the point must lie



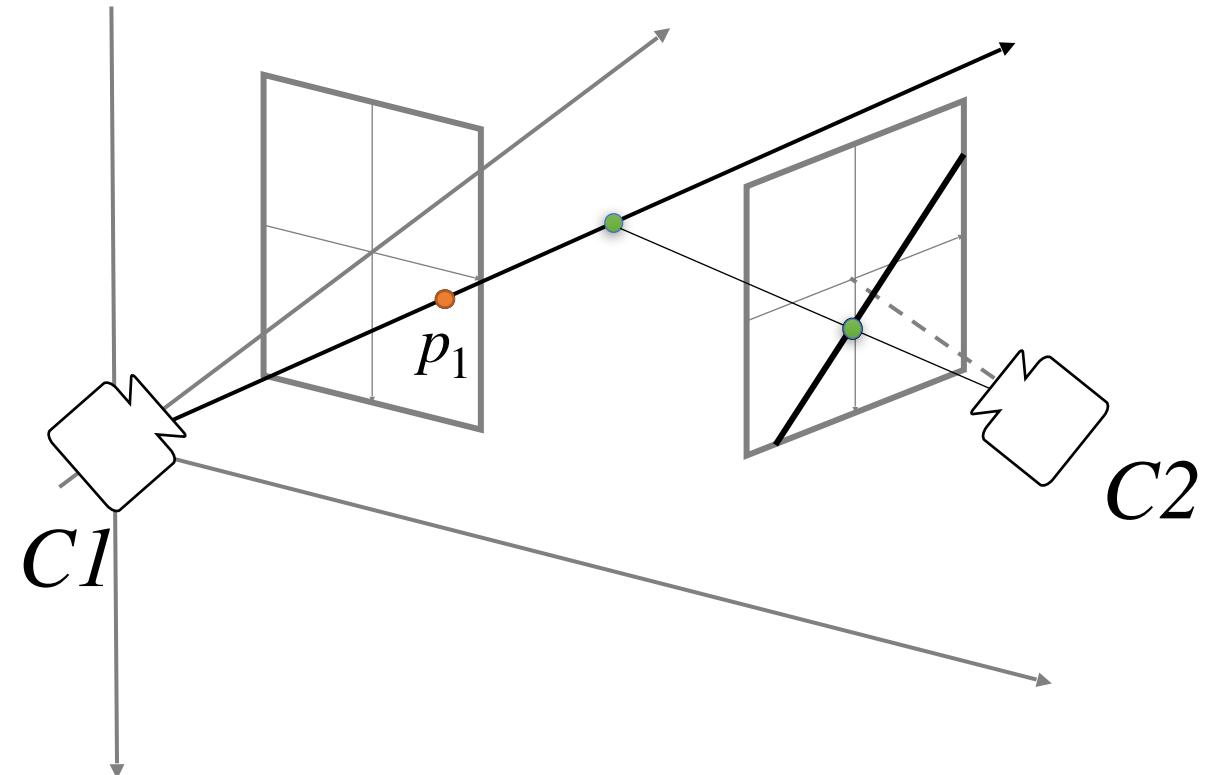
# Epipolar Geometry

- Consider a point  $p_1$  in the first camera
  - This gives a 3D line, along which the point must lie
  - The 3D line projects to a 2D line in the second image – this is called an epipolar line



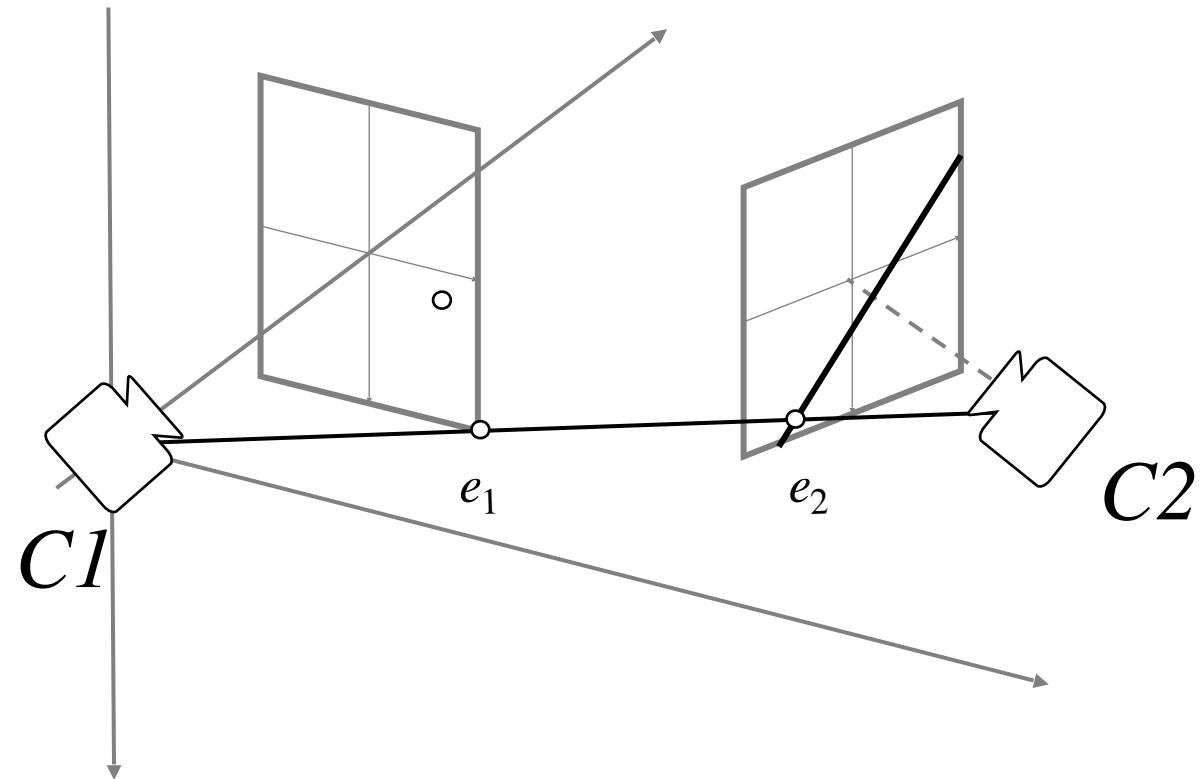
# Epipolar Geometry

- Consider a point  $p_1$  in the first camera
  - This gives a 3D line, along which the point must lie
  - The 3D line projects to a 2D line in the second image – this is called an **epipolar line**
  - The matching point must lie on this line



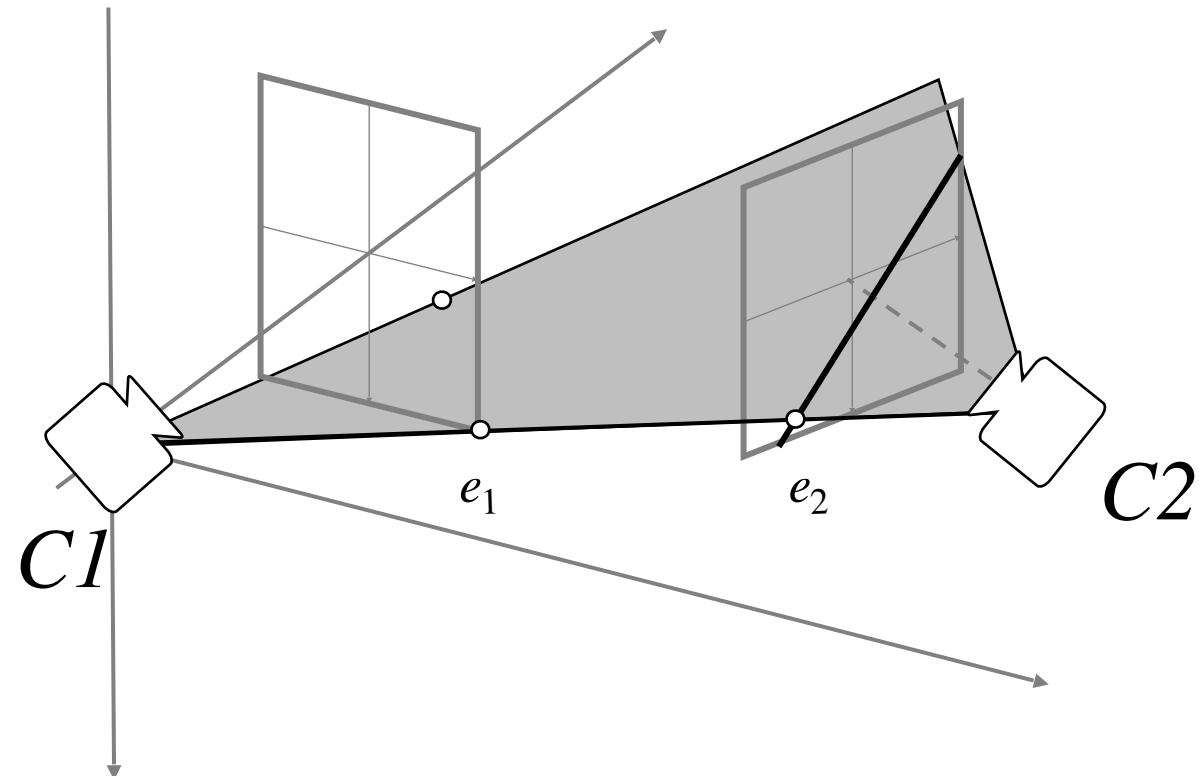
# Epipoles and Epipolar line

- Draw a line between the camera centres
  - This crosses the image planes at the epipoles ( $e_1, e_2$ )



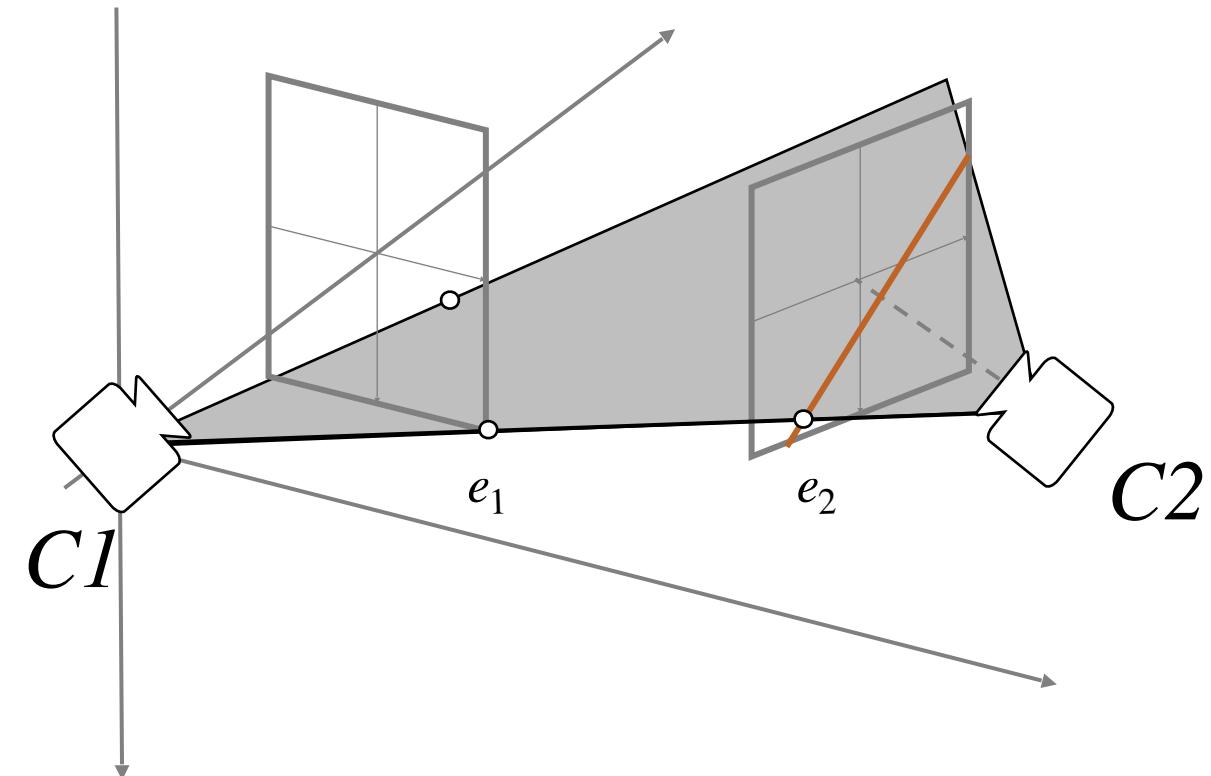
# Epipoles and Epipolar line

- Draw a line between the camera centres
  - This crosses the image planes at the epipoles ( $e_1, e_2$ )
  - Make a plane through a point in one image and the two camera centres: epipolar plane



# Epipoles and Epipolar line

- Draw a line between the camera centres
  - This crosses the image planes at the epipoles ( $e_1, e_2$ )
  - Make a plane through a point in one image and the two camera centres: epipolar plane
  - This intersects the second image at the epipolar line



# Epipoles and Epipolar line

- Each point in one image makes an epipolar line
  - These all meet at the epipole in the other image
- For simple (parallel) case
  - Line through camera centres is the X axis
  - Epipoles are at infinity
  - Epipolar lines are the rows

