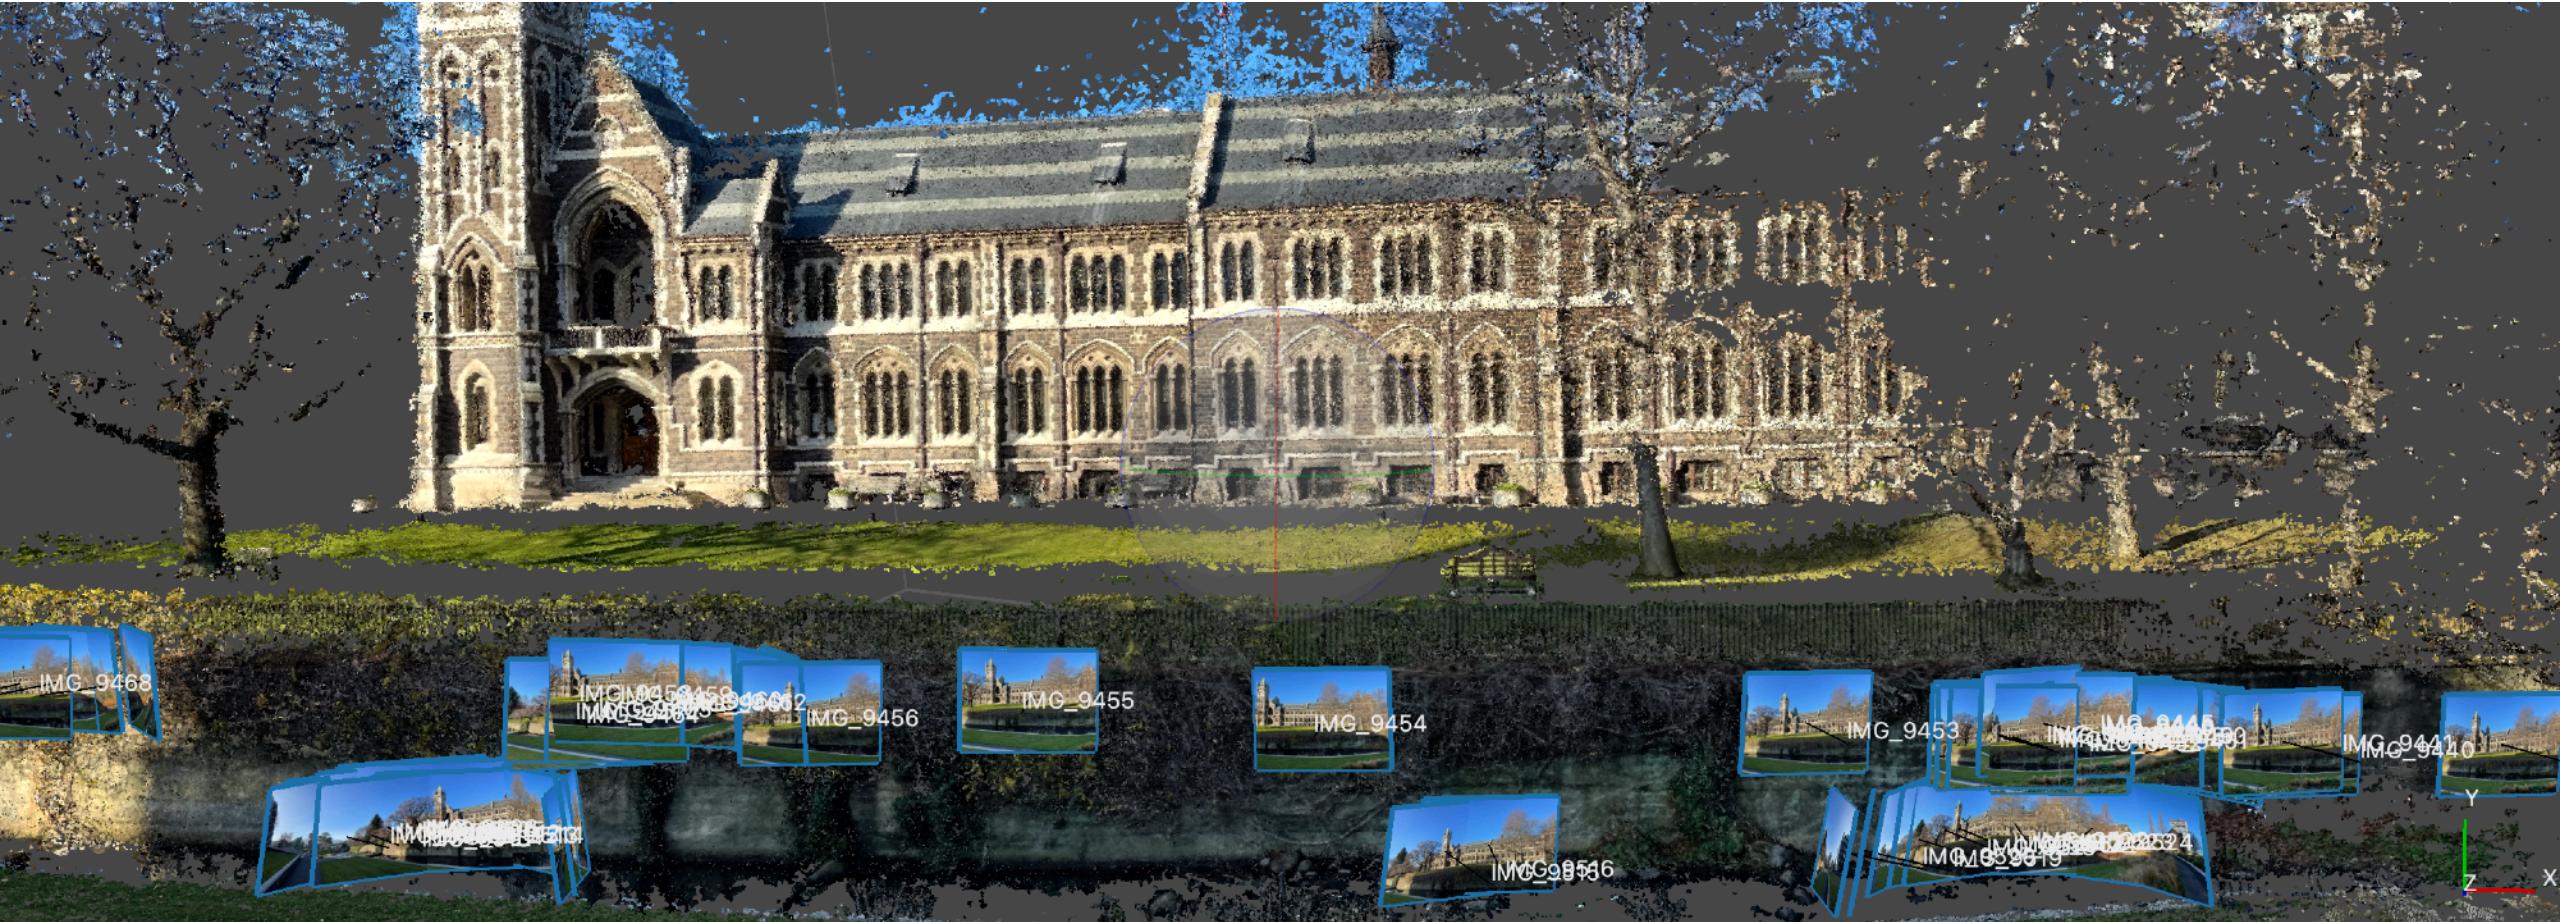


Visual Computing I:

Interactive Computer Graphics and Vision



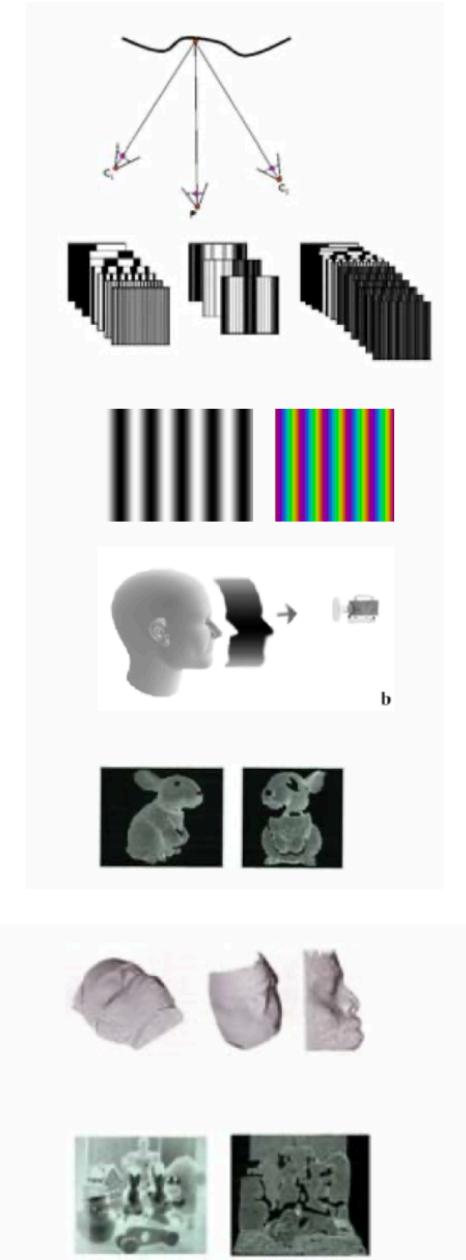
Depth Maps Processing, 3D Reconstruction

Stefanie Zollmann and Tobias Langlotz

Last time..

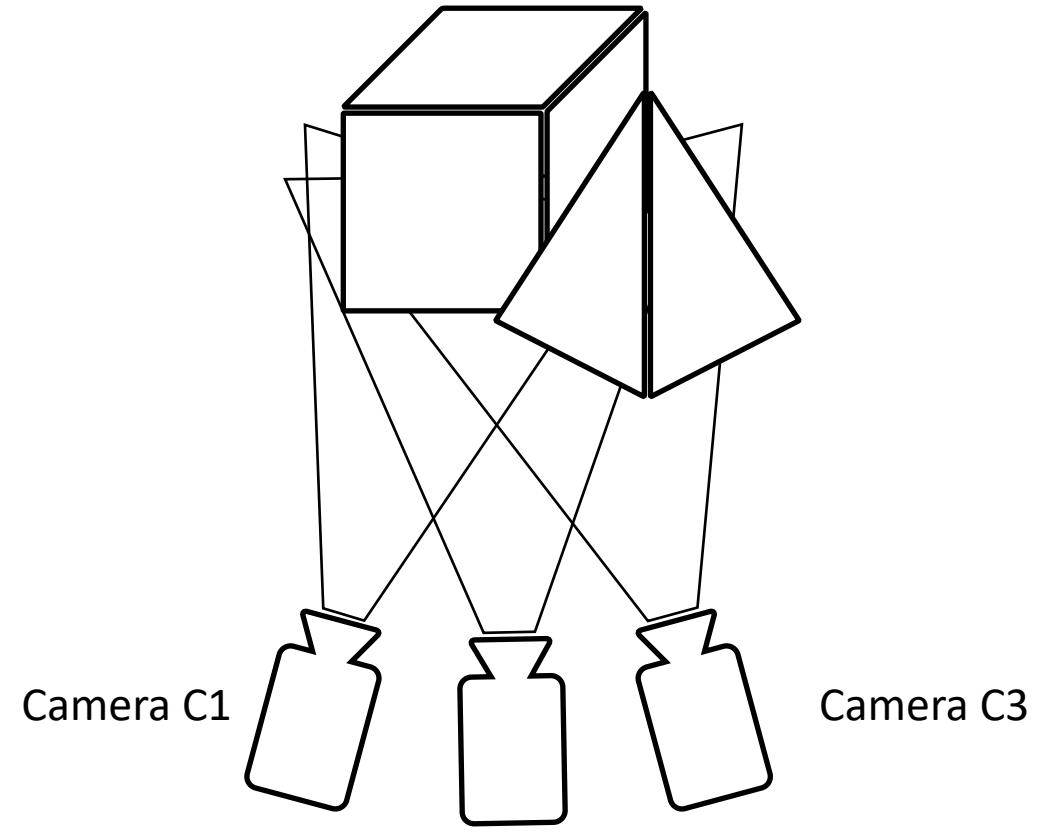
Overview

- Range Scanning
 - Passive:
 - Stereo or Multi-Camera
 - Active:
 - Structured Light Projection
 - Optical Time-of-Flight
 - Direct ToF
 - Indirect ToF
 - Registration of Depth Maps
 - Processing Depth Maps



Passive Range Scanning

- Two or more cameras
 - Stereo camera
 - Two cameras
 - One camera takes several photos from different perspectives
 - Multiple cameras
- Which one is the easiest and why?
- What is the main challenge?



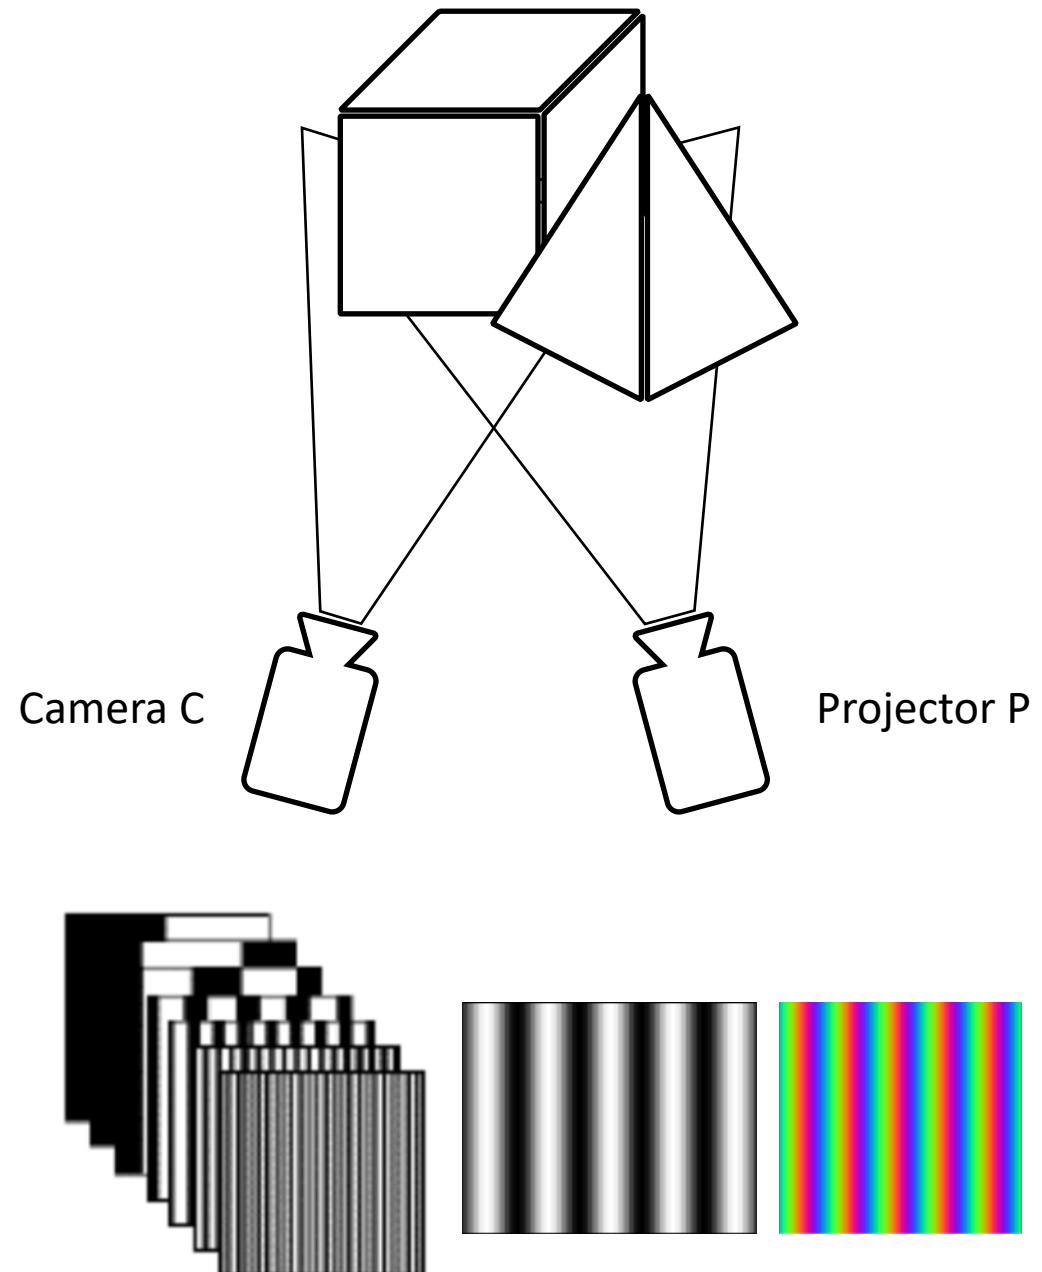
PointGrey Bumble Bee



Zed Mini

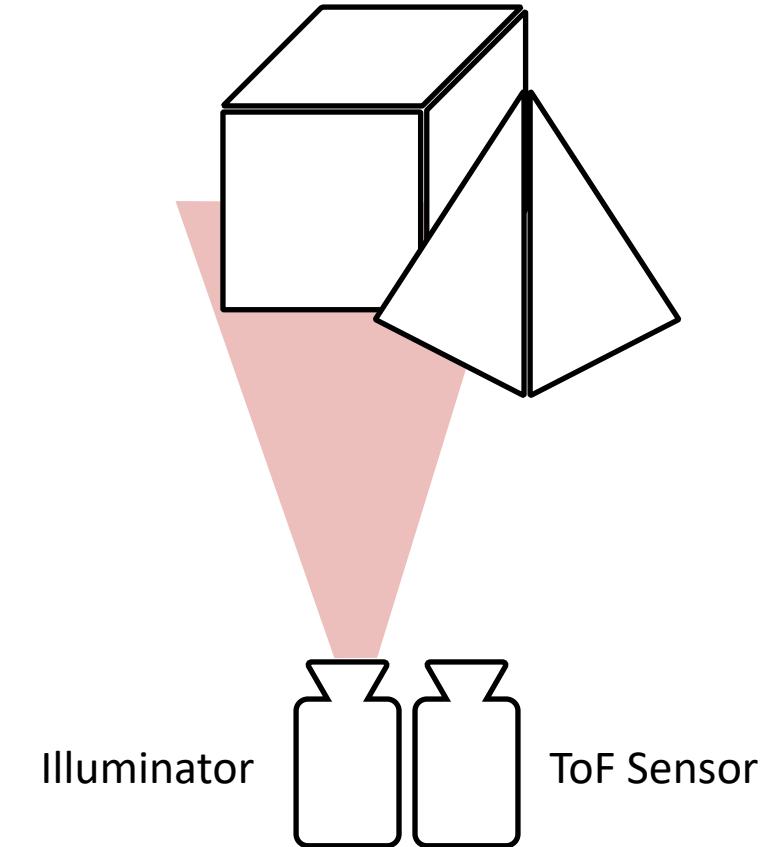
Structured Light Projection

- Many structured light techniques exist that follow the two objectives:
 - Determine projector-camera correspondences with as few images as possible
 - Be robust against surface modulation and noise
- Codes include:
 - Binary codes, such as the Gray code (most common): <100 images, quite robust
 - Intensity codes that apply cosine patterns and phase shifts: 6 images, but are not robust against strong absorptions
 - Color coding: <3 images, but very fragile (does not work on coloured surfaces)



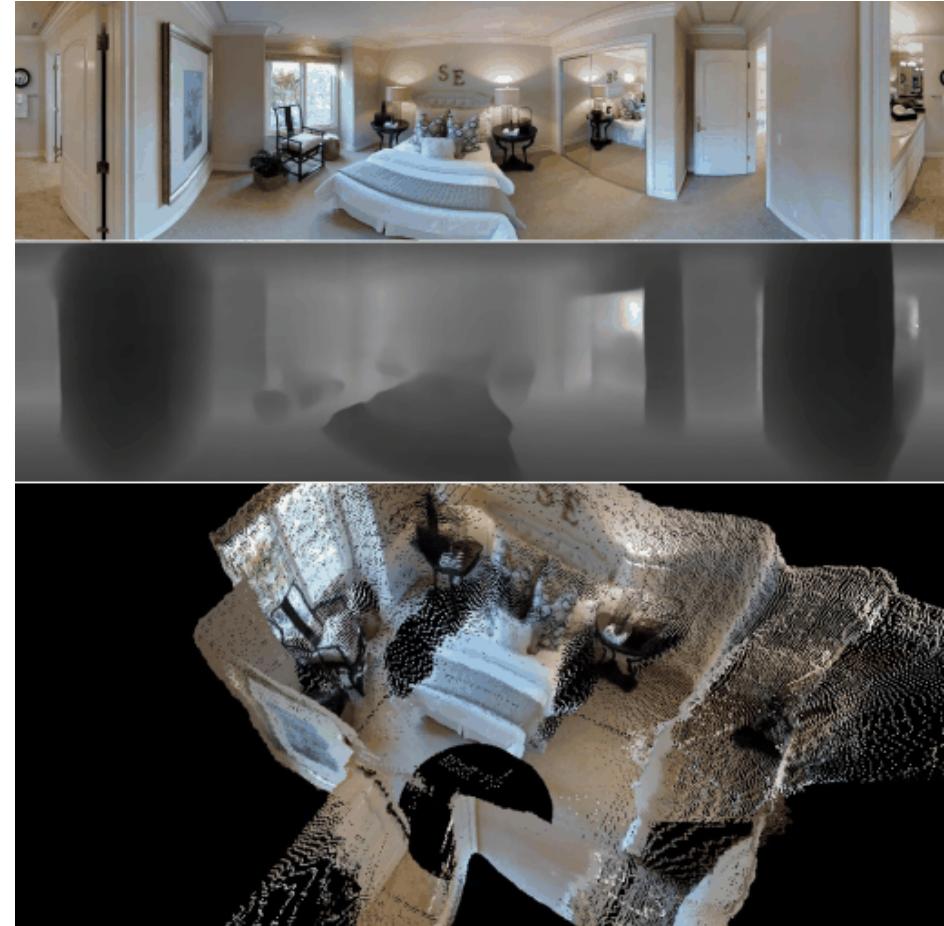
Optical Time-of-Flight (ToF)

- General idea:
 - Compute depth by round-trip estimation of a light wave emitted and its reflection back to the sensor
 - $D = C \cdot \Delta t/2$
- Use low-power IR illumination for pulsed illumination
 - Can be switched on/off with <1ns
 - Generates a “light wall” that is reflected back by objects



Depth from Mono / Depth Estimation using ML

- Convolutional neural networks (CNNs) have been applied to monocular depth estimation
- Common architectures include U-Net, ResNet, and DenseNet
- During training:
 - Network is given pairs of images and their corresponding ground-truth depth maps
 - Minimize the difference between predicted depth and ground truth depth

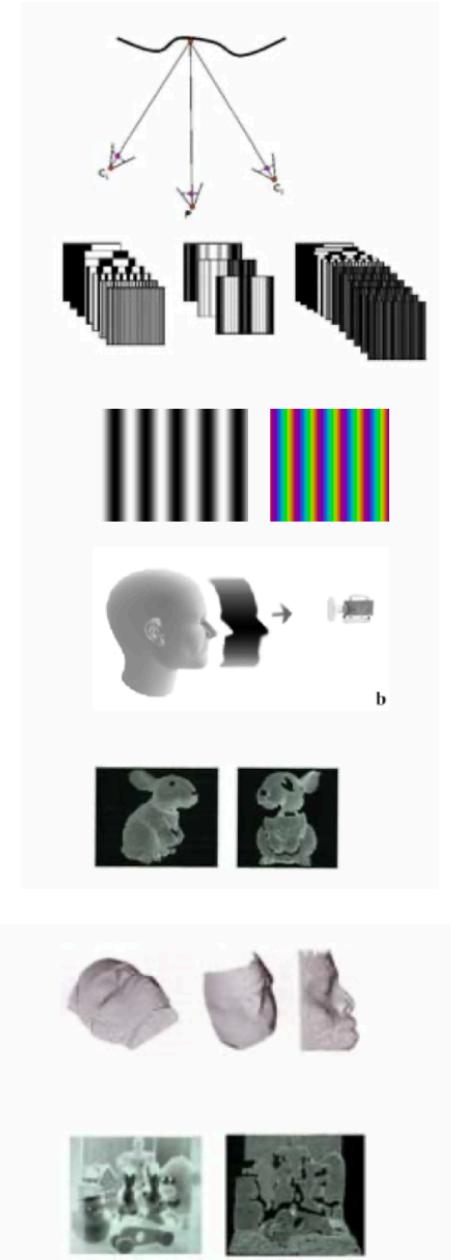


<https://github.com/yuhuanyeh/BiFuse>

Today:
Processing of Depth Maps, Relative Pose
and Multi-view Stereo

Overview

- Range Scanning
 - Passive:
 - Stereo or Multi-Camera
 - Active:
 - Structured Light Projection
 - Optical Time-of-Flight
 - Direct ToF
 - Indirect ToF
 - Registration of Depth Maps
 - Processing Depth Maps



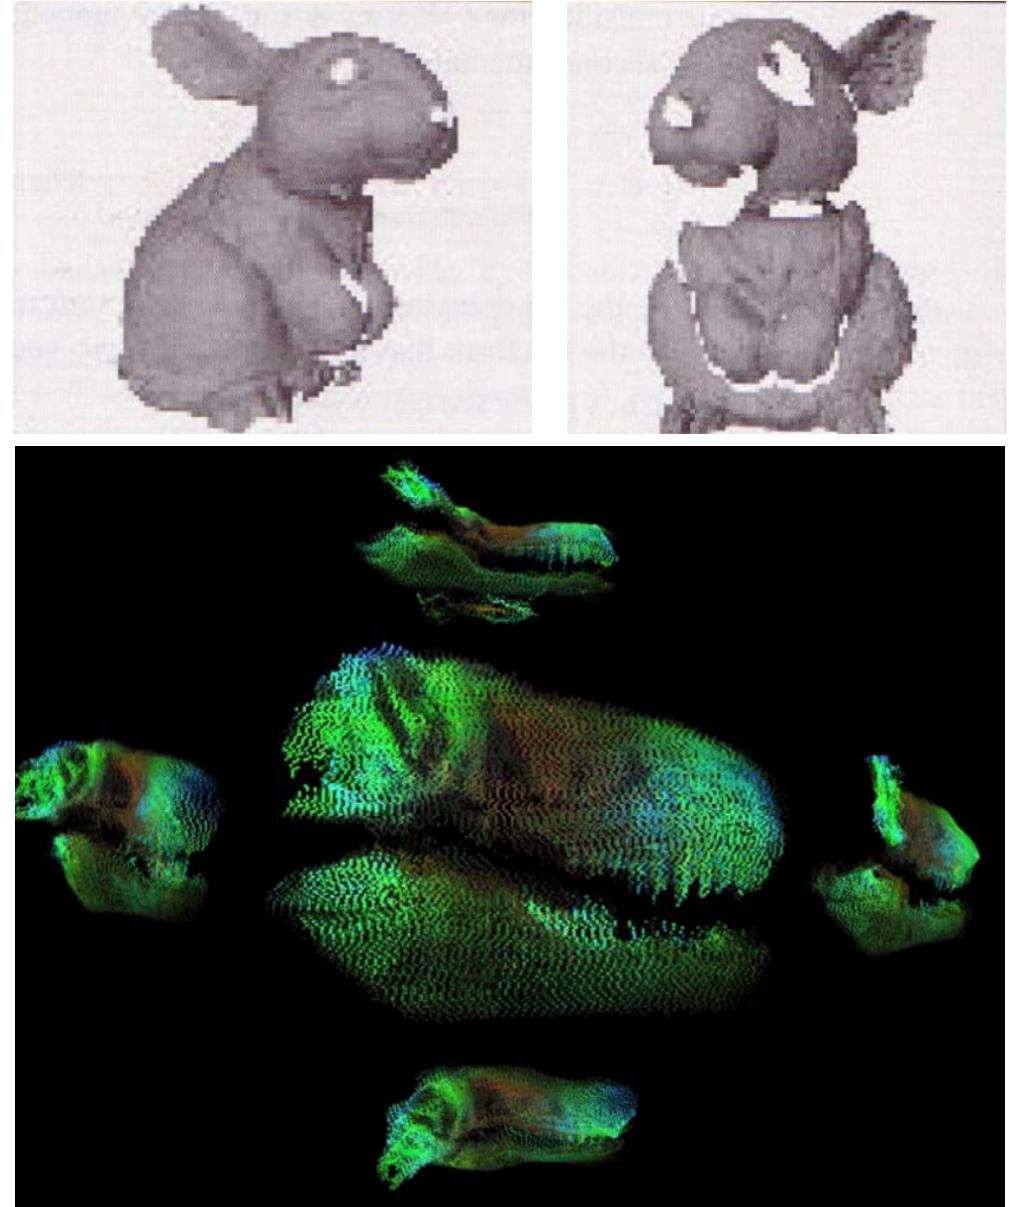
Registration of Depth Maps

Registration of Depth Maps

- Depth map gives us distance for each pixel, can compute 3D points:

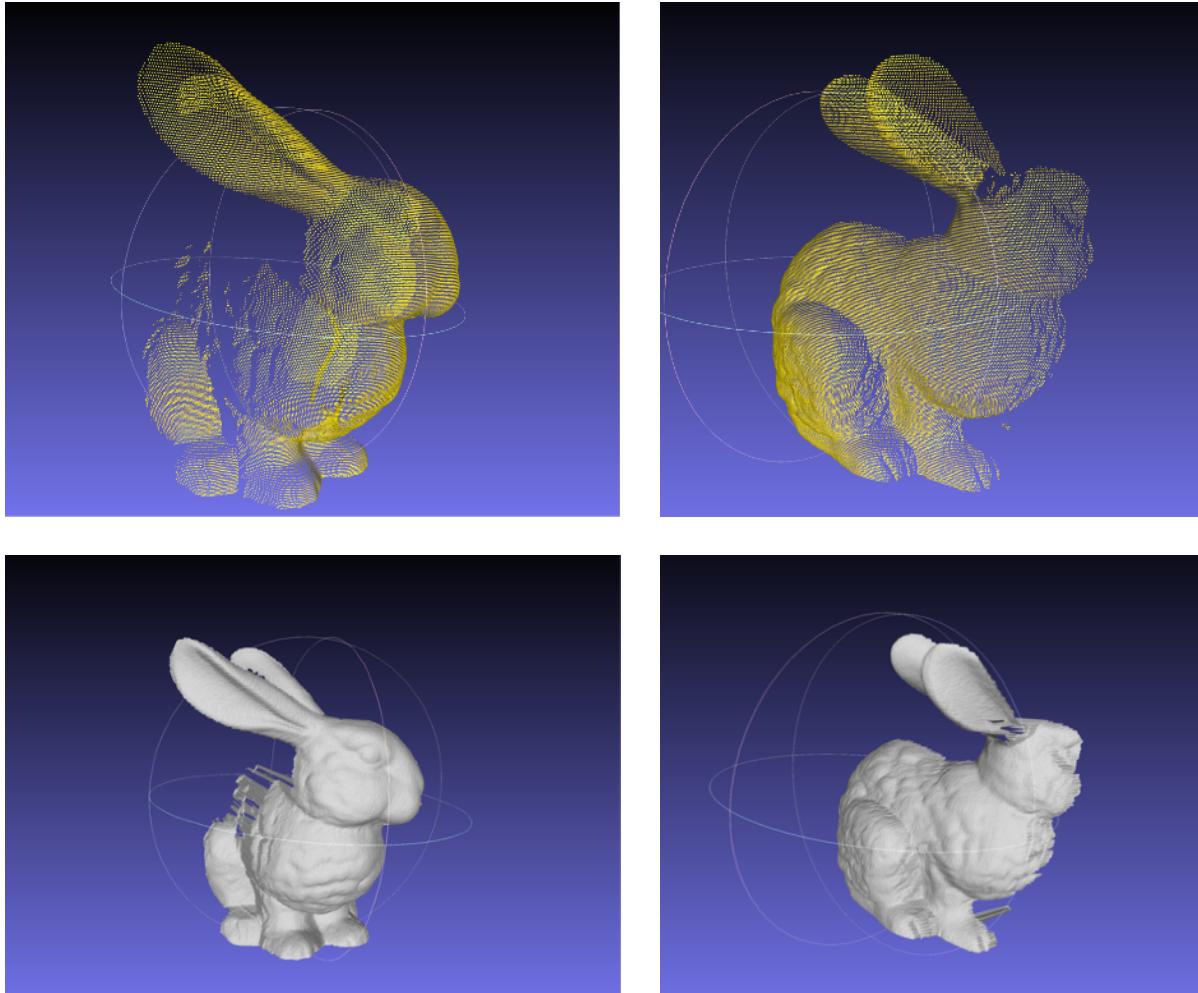
$$\mathbf{X}_{\text{camera}} = D(u, v) \cdot K^{-1} \begin{bmatrix} u \\ v \\ 1 \end{bmatrix}$$

- Having multiple depth maps of the same scene, how can they be merged?
- This process is known as registration
- If camera poses are unknown, we need to align the 3D points purely based on geometry



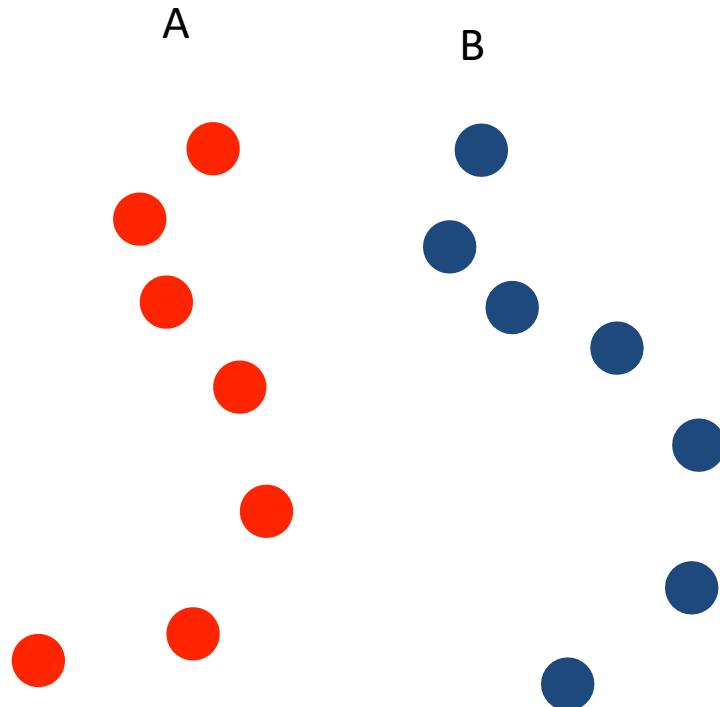
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



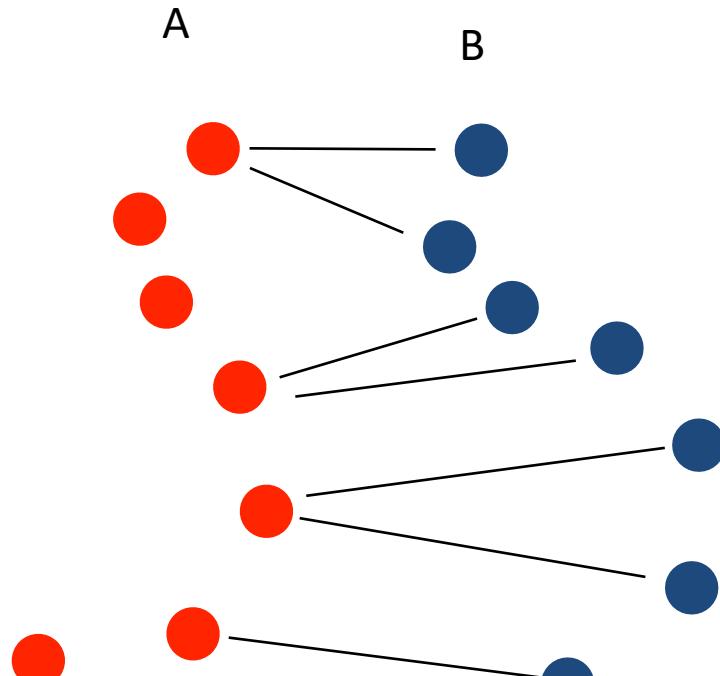
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



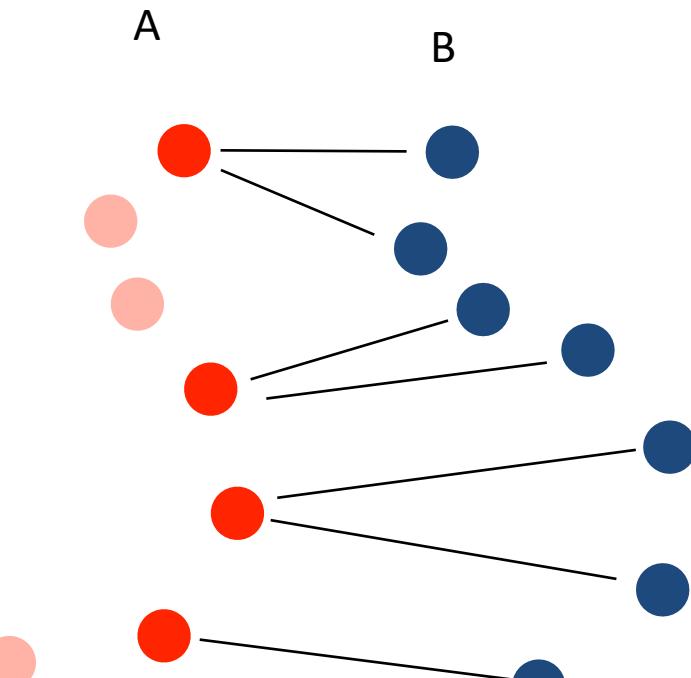
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



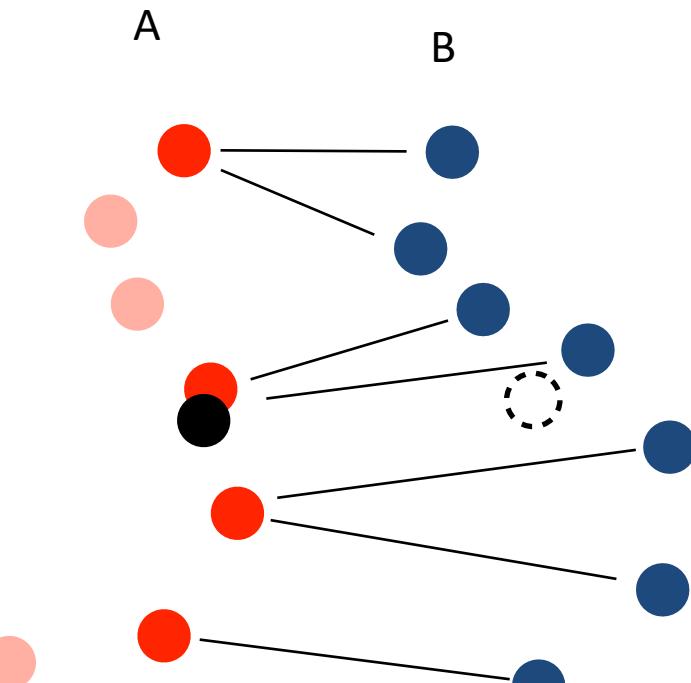
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



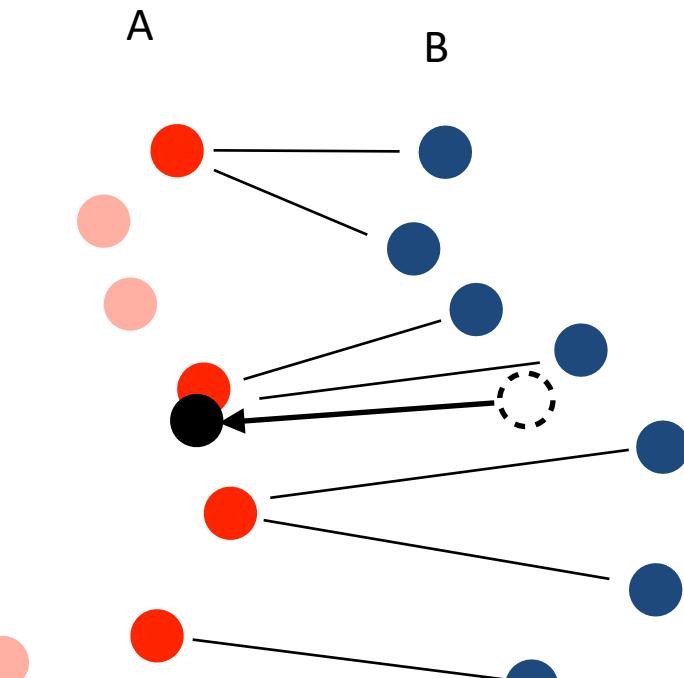
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



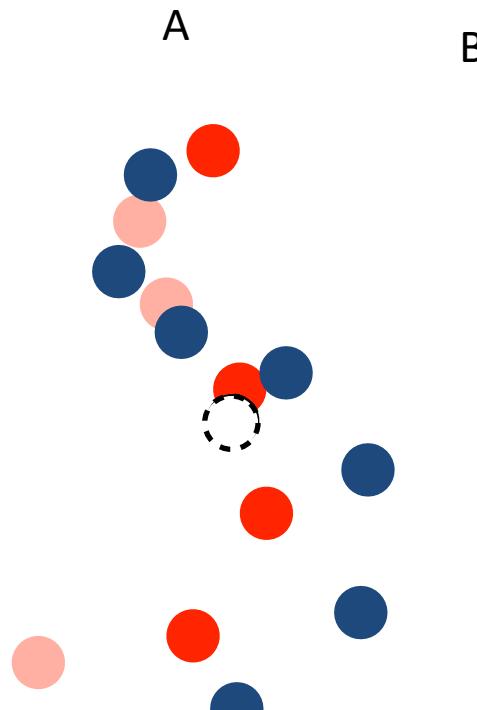
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



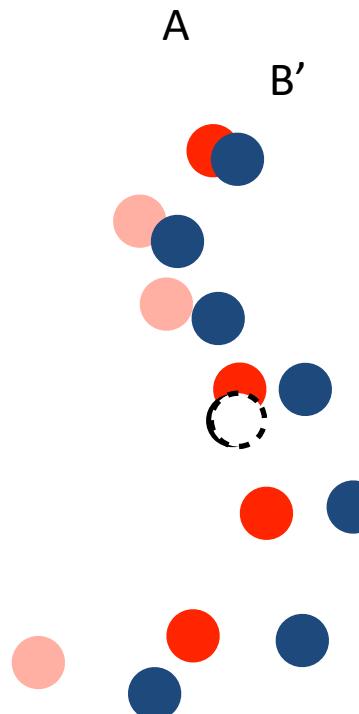
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



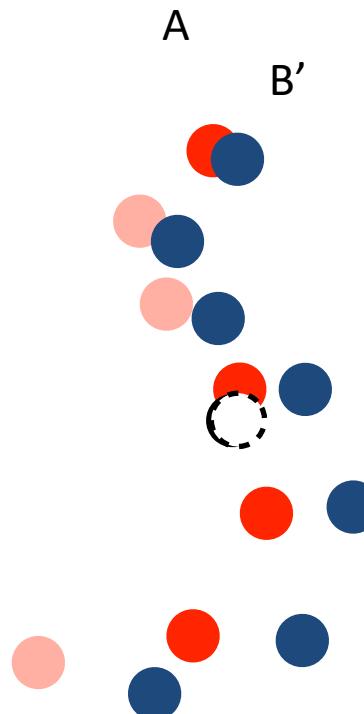
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



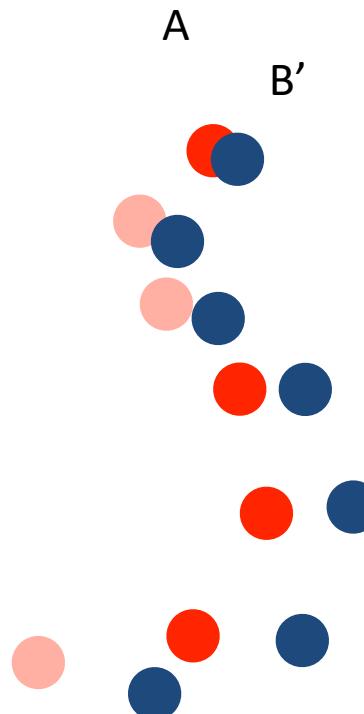
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



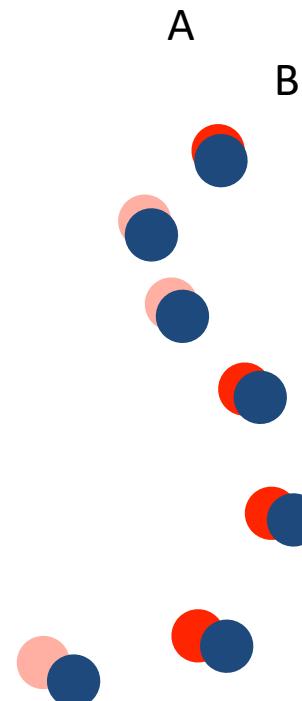
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



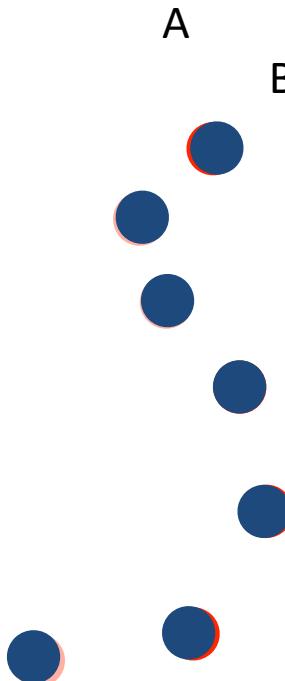
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



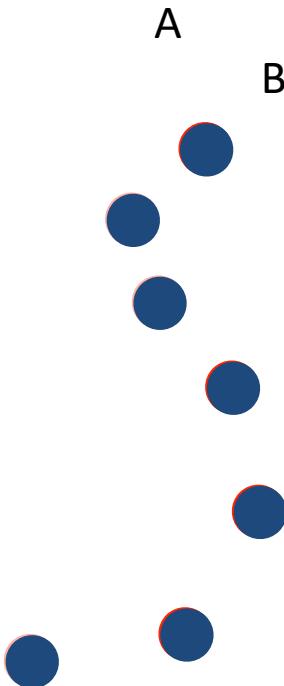
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



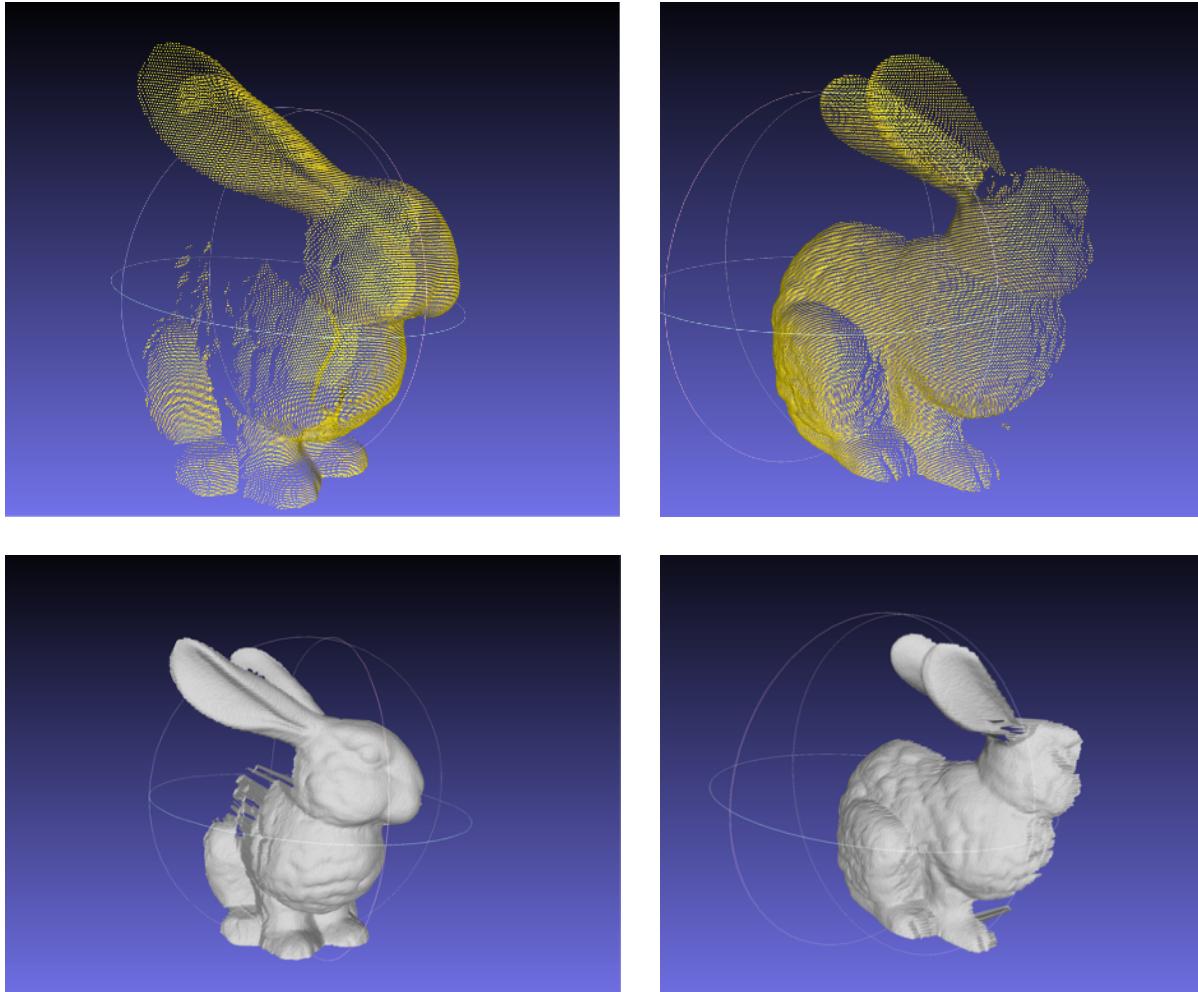
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



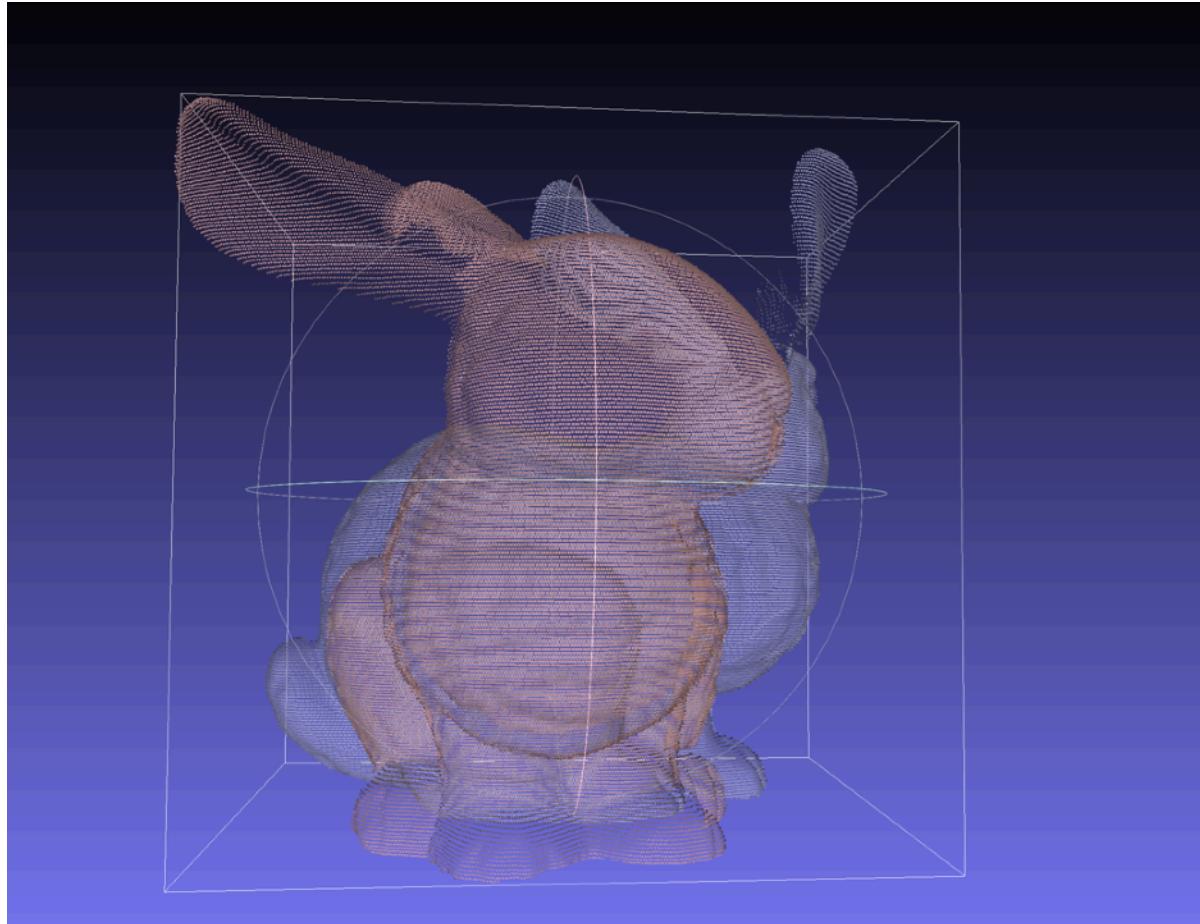
Iterative Closest-Point Method

- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



Iterative Closest-Point Method

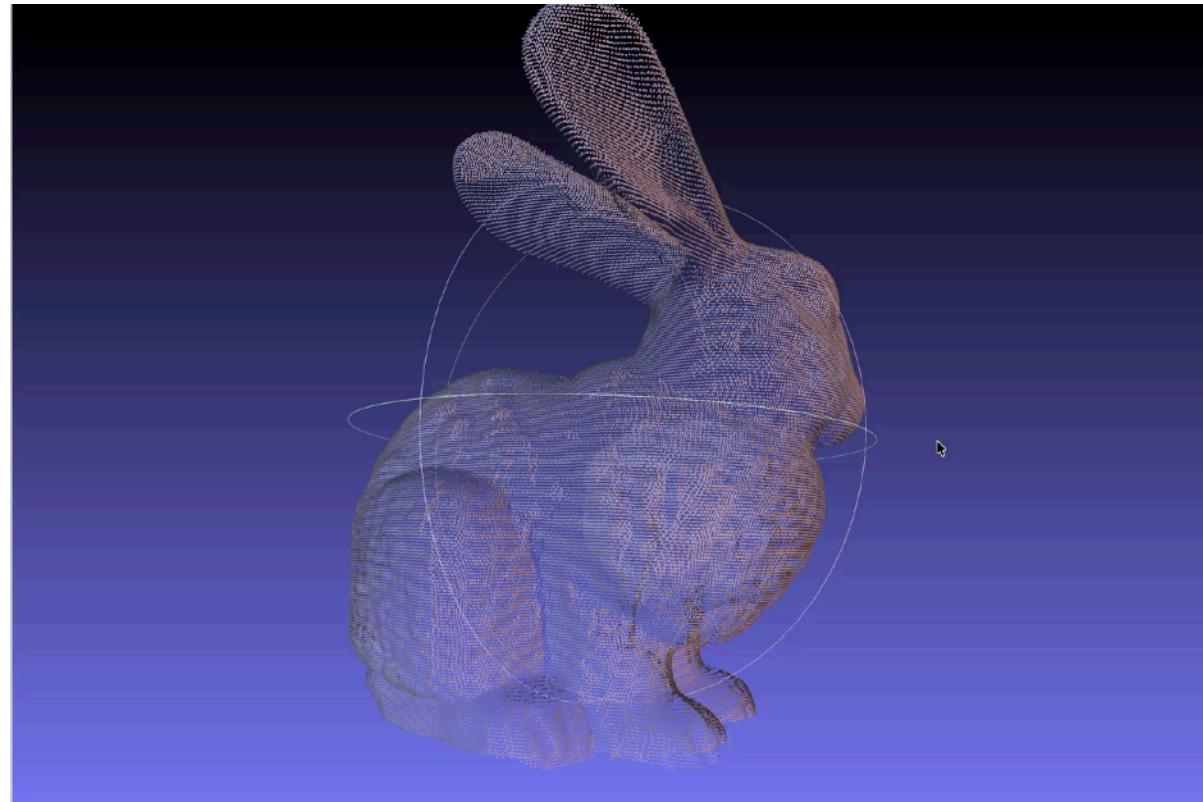
- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



ICP in Meshlab

Iterative Closest-Point Method

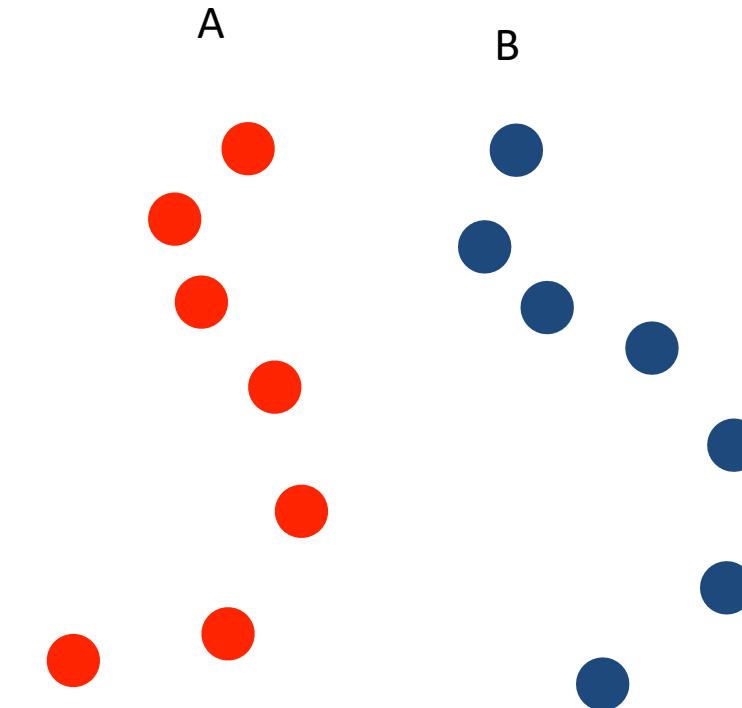
- Registration of two sets of three-dimensional points requires the computation of a rigid transformation that maps the first point set (B) to the second one (A)
- Iteratively minimise the average distance between the two point sets:
 - Find correspondences between the sets by matching every point in B with the point closest to it in A
 - Compute rigid transformation that maps B to A and apply it to B (B')
 - Compute average distance E between A and B'
 - Repeat until E is minimal



ICP in Meshlab

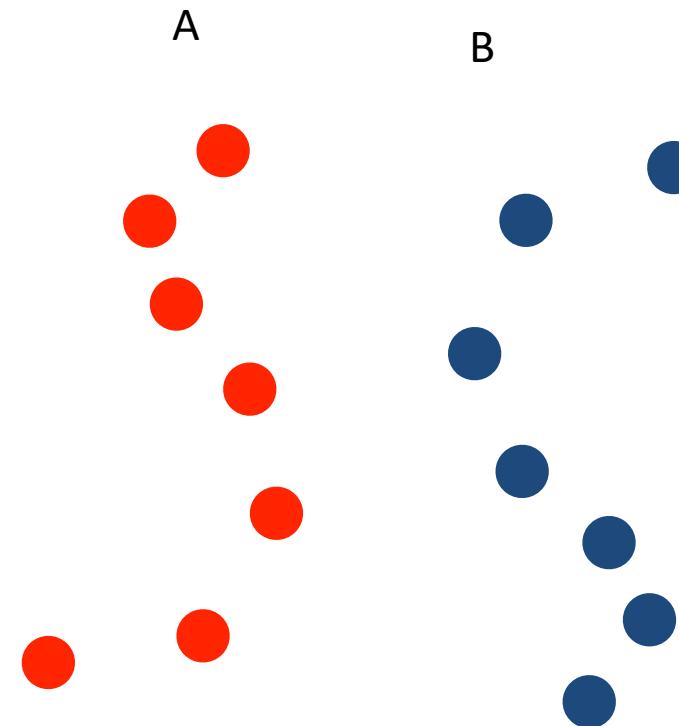
Iterative Closest-Point Method

- How to find a reasonable initial guess?



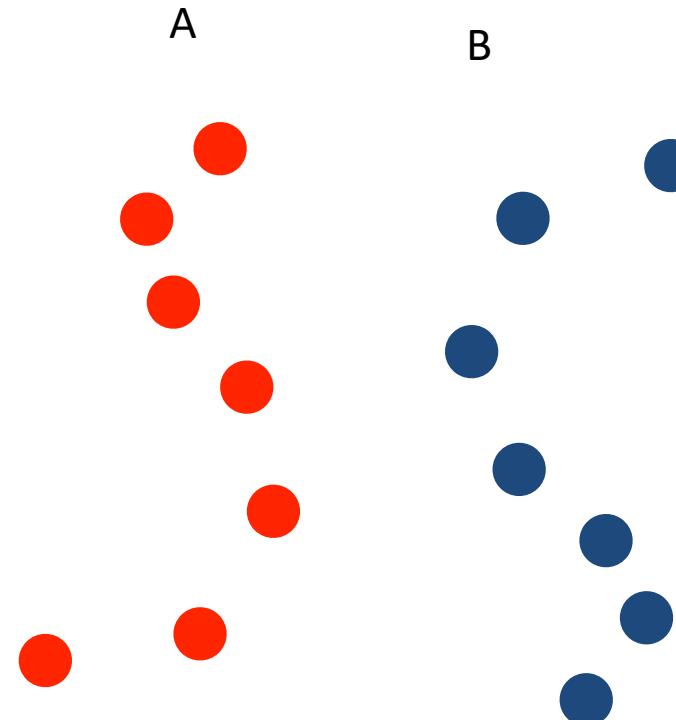
Iterative Closest-Point Method

- How to find a reasonable initial guess?

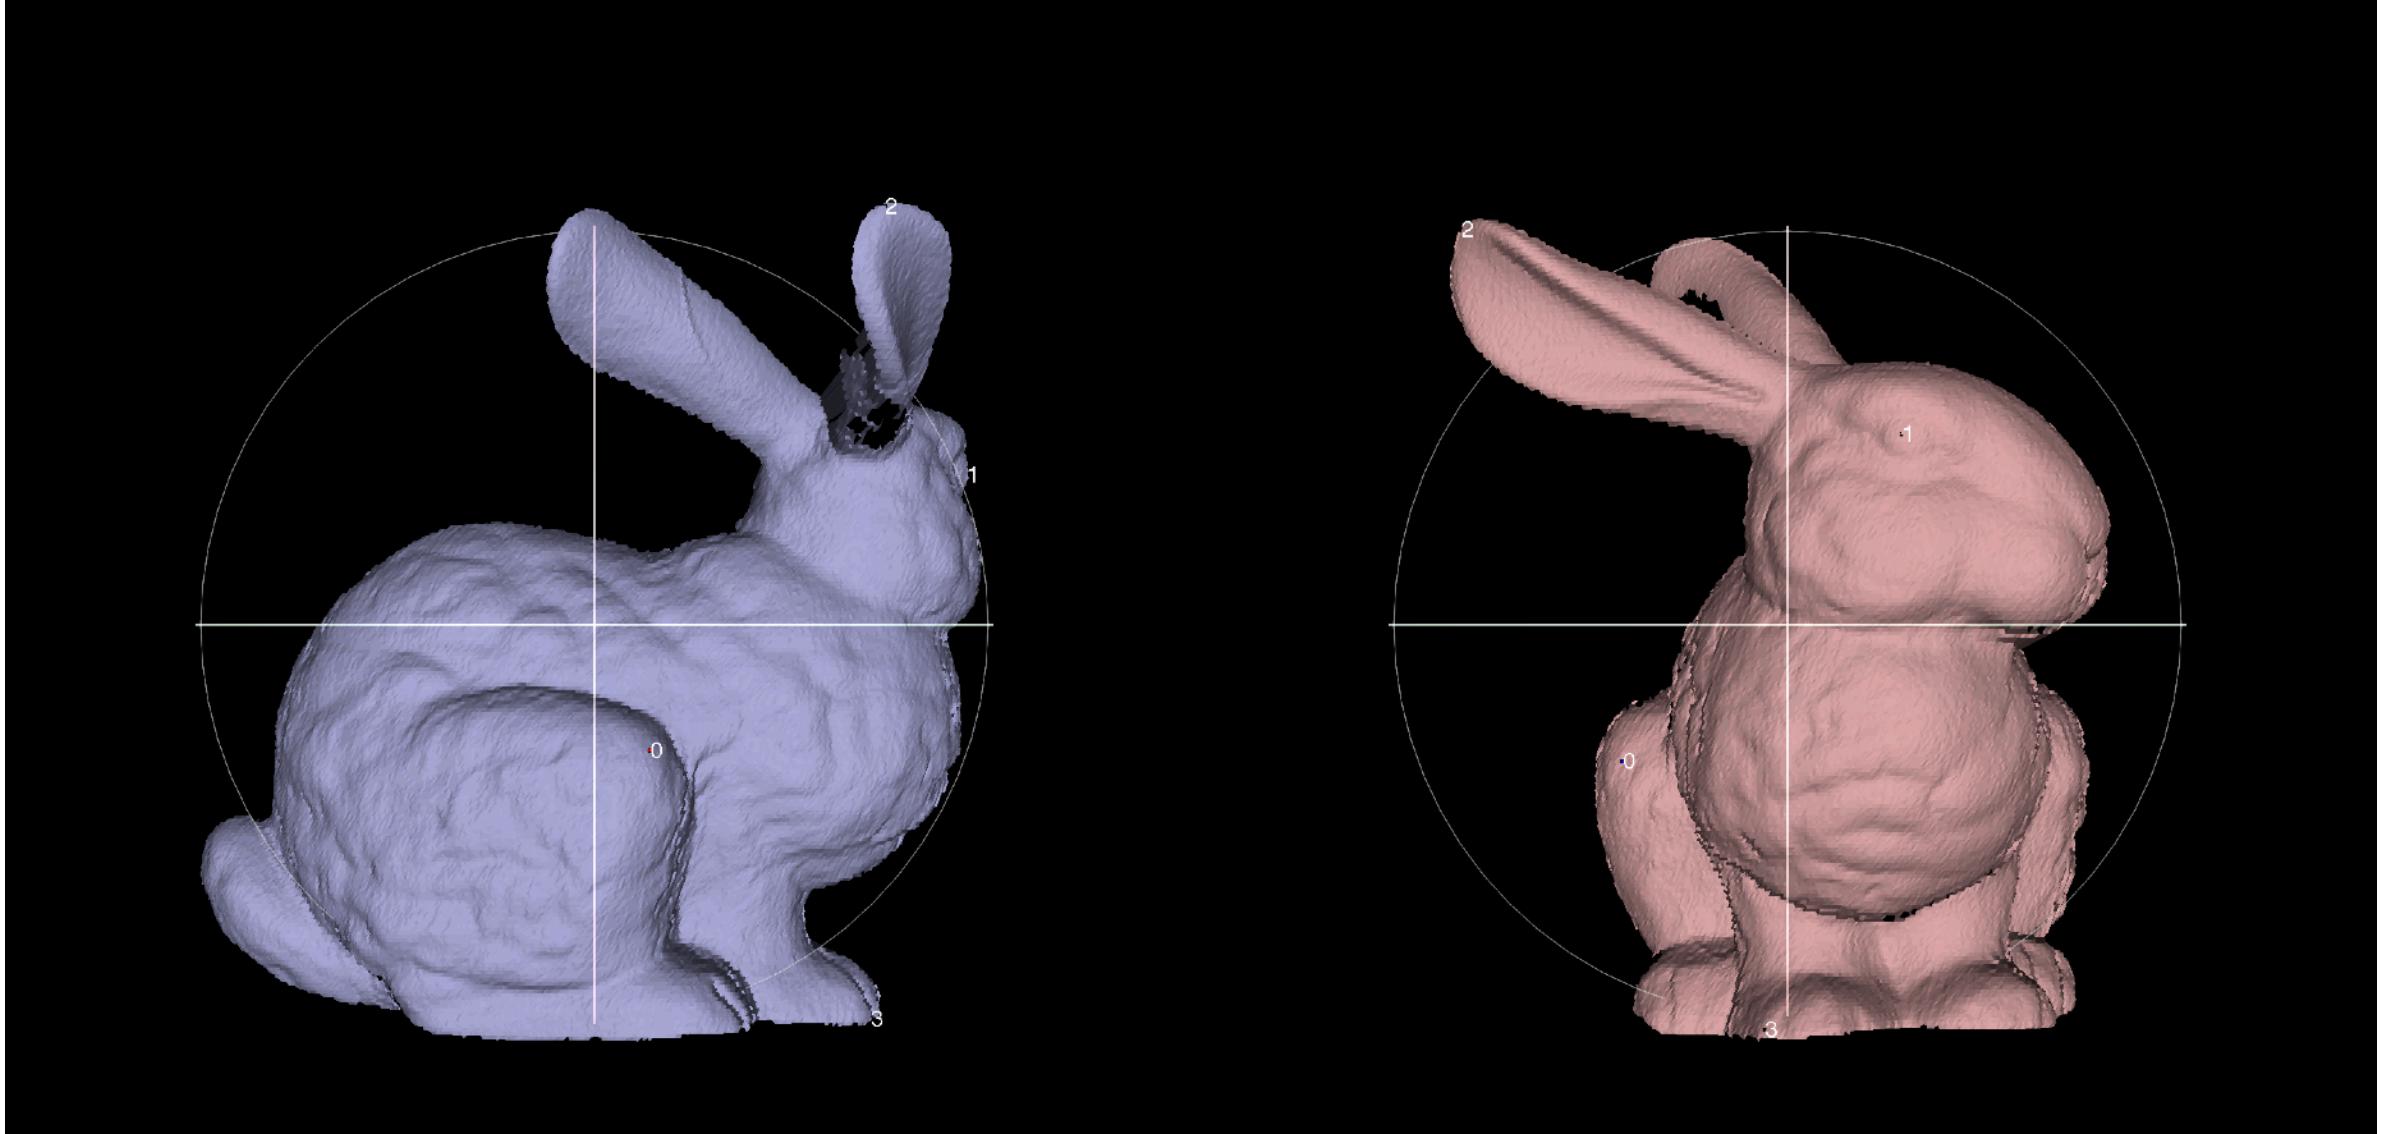


Iterative Closest-Point Method

- How to find a reasonable initial guess?
 - Manually
 - Matching distinctive features between clouds
 - Using sensor data
 - Global registration algorithms
(algorithms not needing initial guess)
- At every iteration finding the closest point M in A to a given point S in B takes (for n points) how long?
 - $O(n^2)$ brute-force
- Can this be done faster?
 - $O(\log n)$ for nearest-neighbour matches with k-d trees



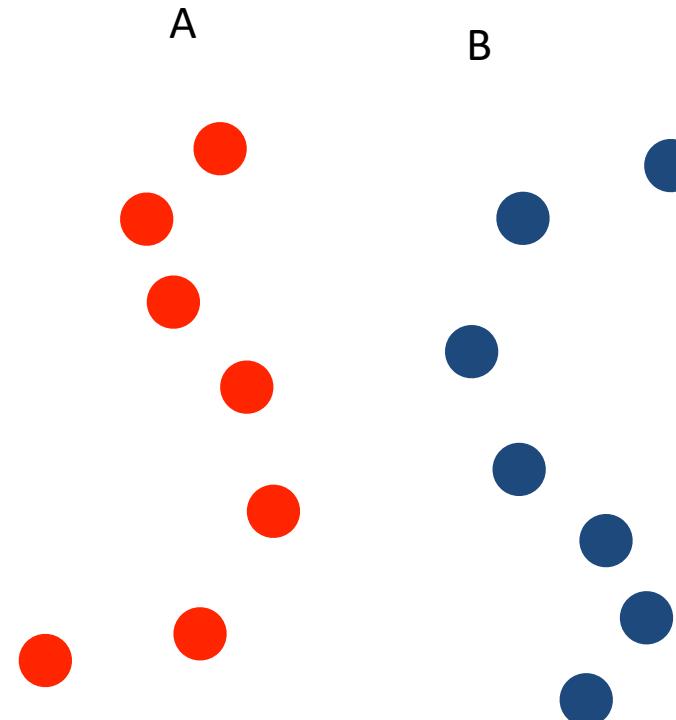
Iterative Closest-Point Method



Manual alignment in Meshlab to initialise ICP

Iterative Closest-Point Method

- How to find a reasonable initial guess?
 - Manually
 - Matching distinctive features between clouds
 - Using sensor data
 - Global registration algorithms
(algorithms not needing initial guess)
- At every iteration finding the closest point M in A to a given point S in B takes (for n points) how long?
 - $O(n^2)$ brute-force
- Can this be done faster?
 - $O(\log n)$ for nearest-neighbour matches with k-d trees



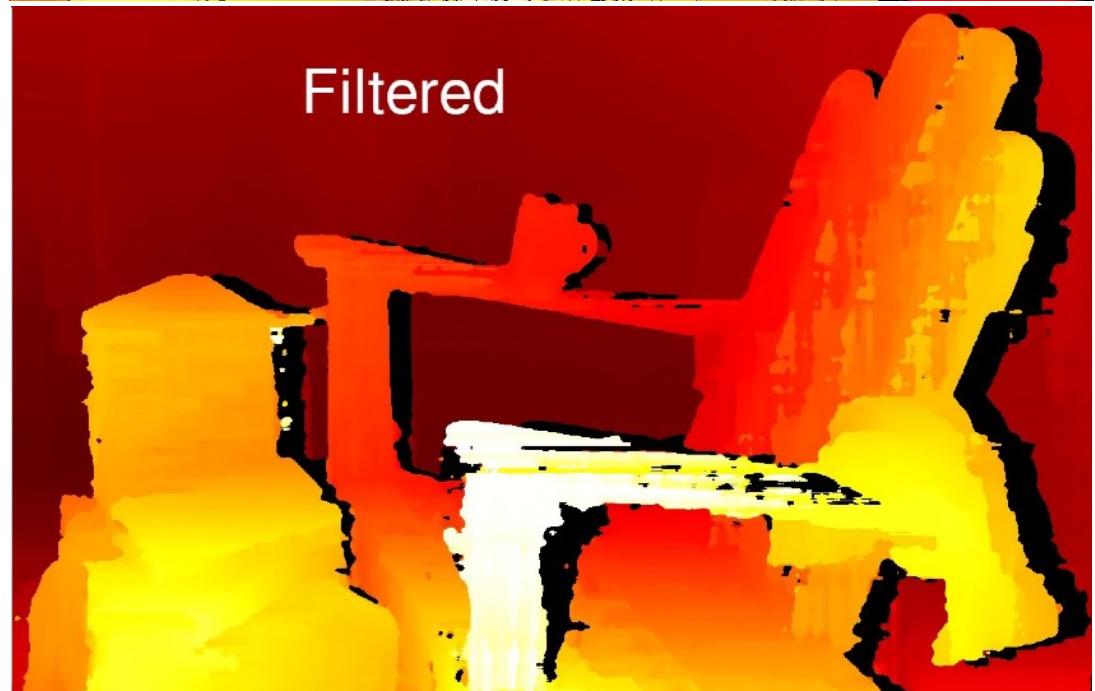
Processing Depth Maps

Processing Depth Maps

- Depth maps can be processed like images
- Image operators, such as filters, can be applied to:
 - Handle discontinuities and overlaps
 - Fill missing portions
 - Smoothen geometric noise
 - Analyse depth images
 - Find features such as edges
 - Segmentation of objects
 - Etc.



Unfiltered



Filtered

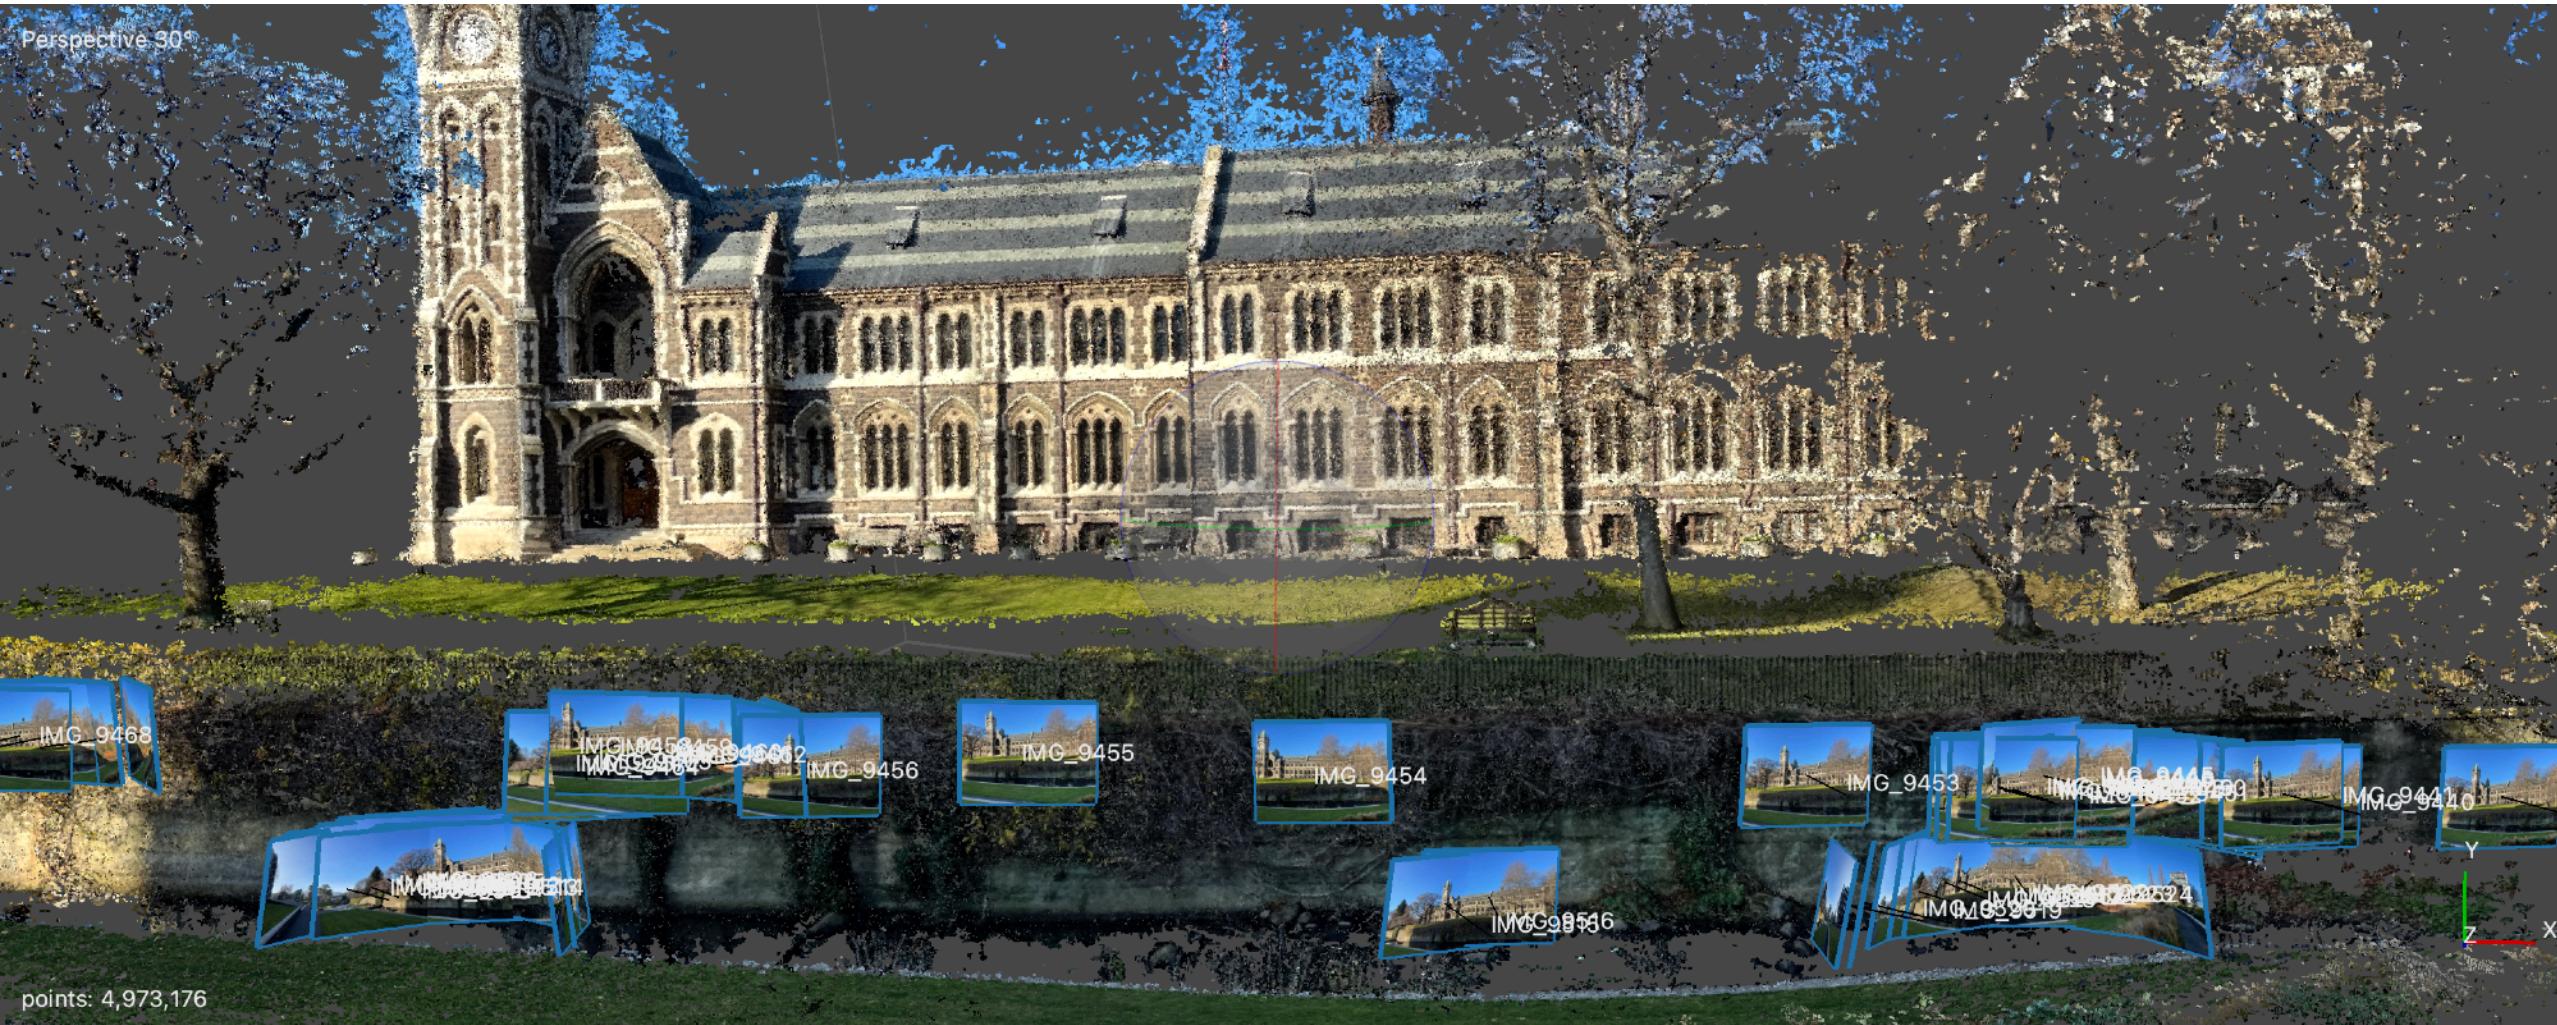
3-Min Discussion: Where do you see challenges with this approach for creating larger 3D models?



03:00

3D Reconstruction

Goal



Workflow

Unstructured Images



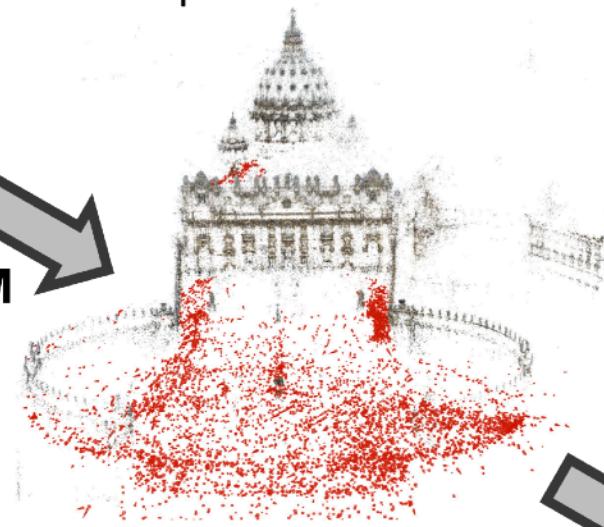
Assoc.

Scene Graph



SFM

Sparse Model



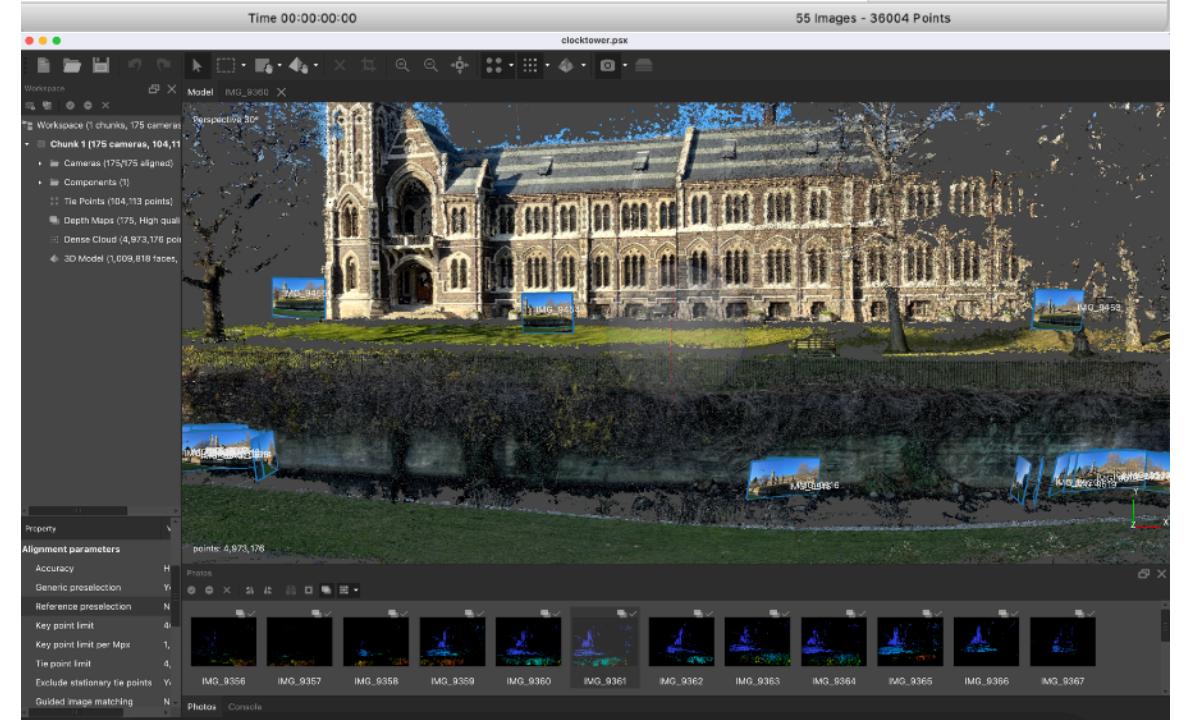
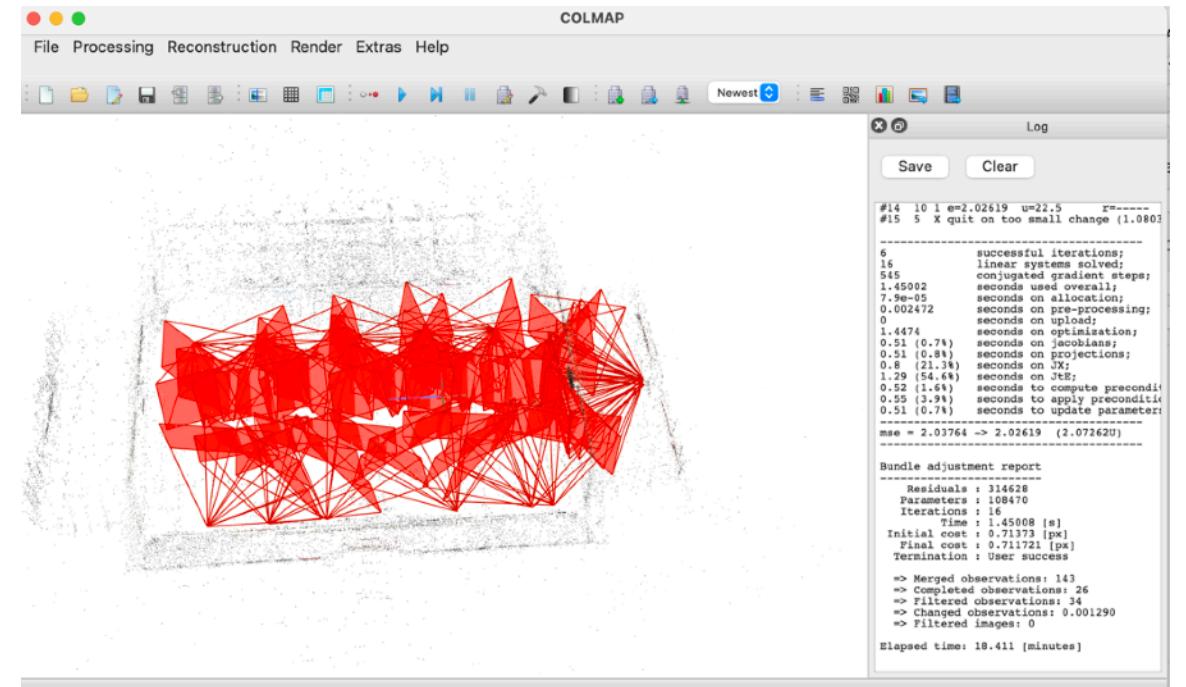
MVS

Dense Model

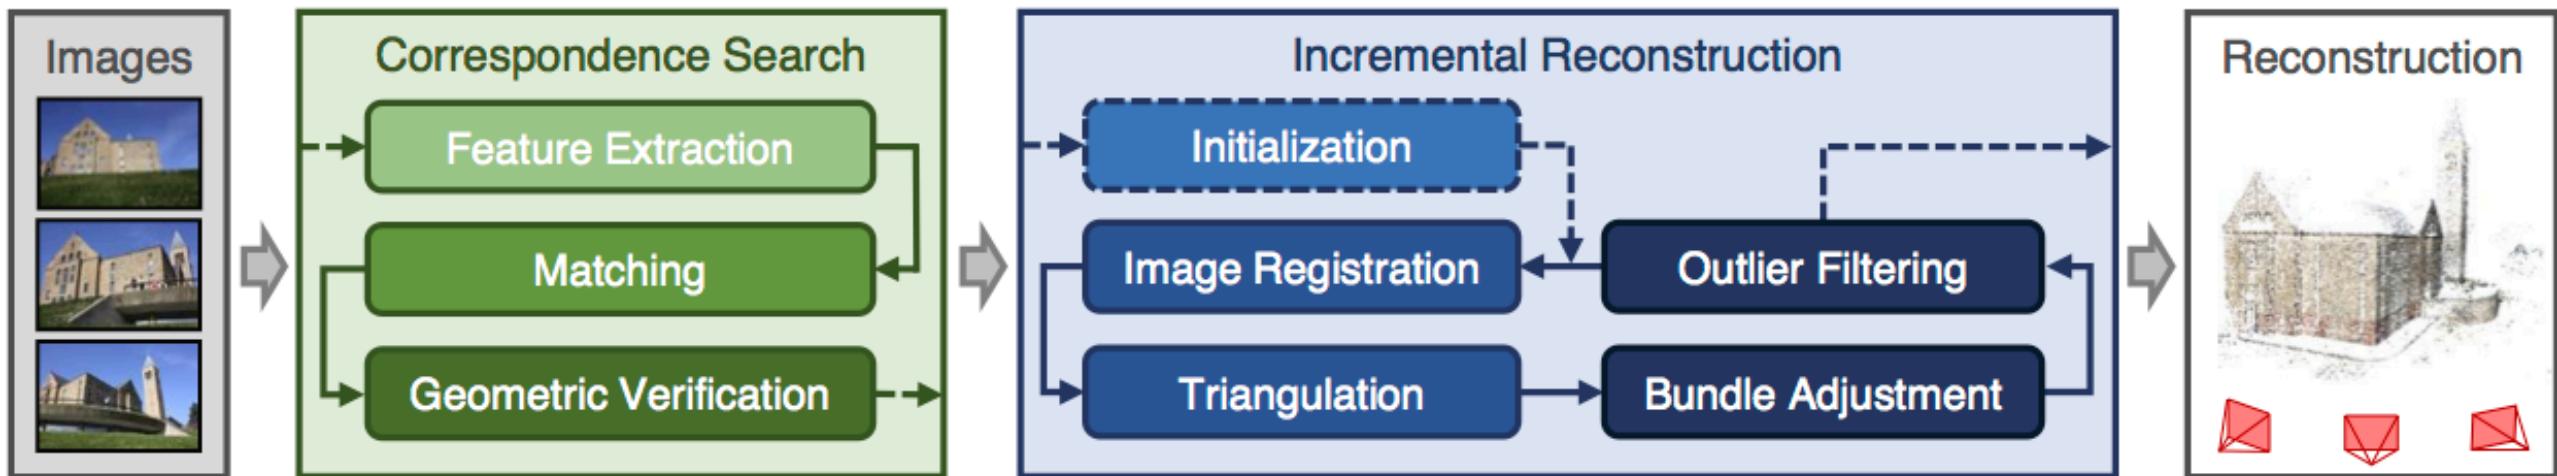


Selection of Tools

- COLMAP <https://colmap.github.io>
- Metashape
- Meshroom
- VisualSFM
- Pix4D

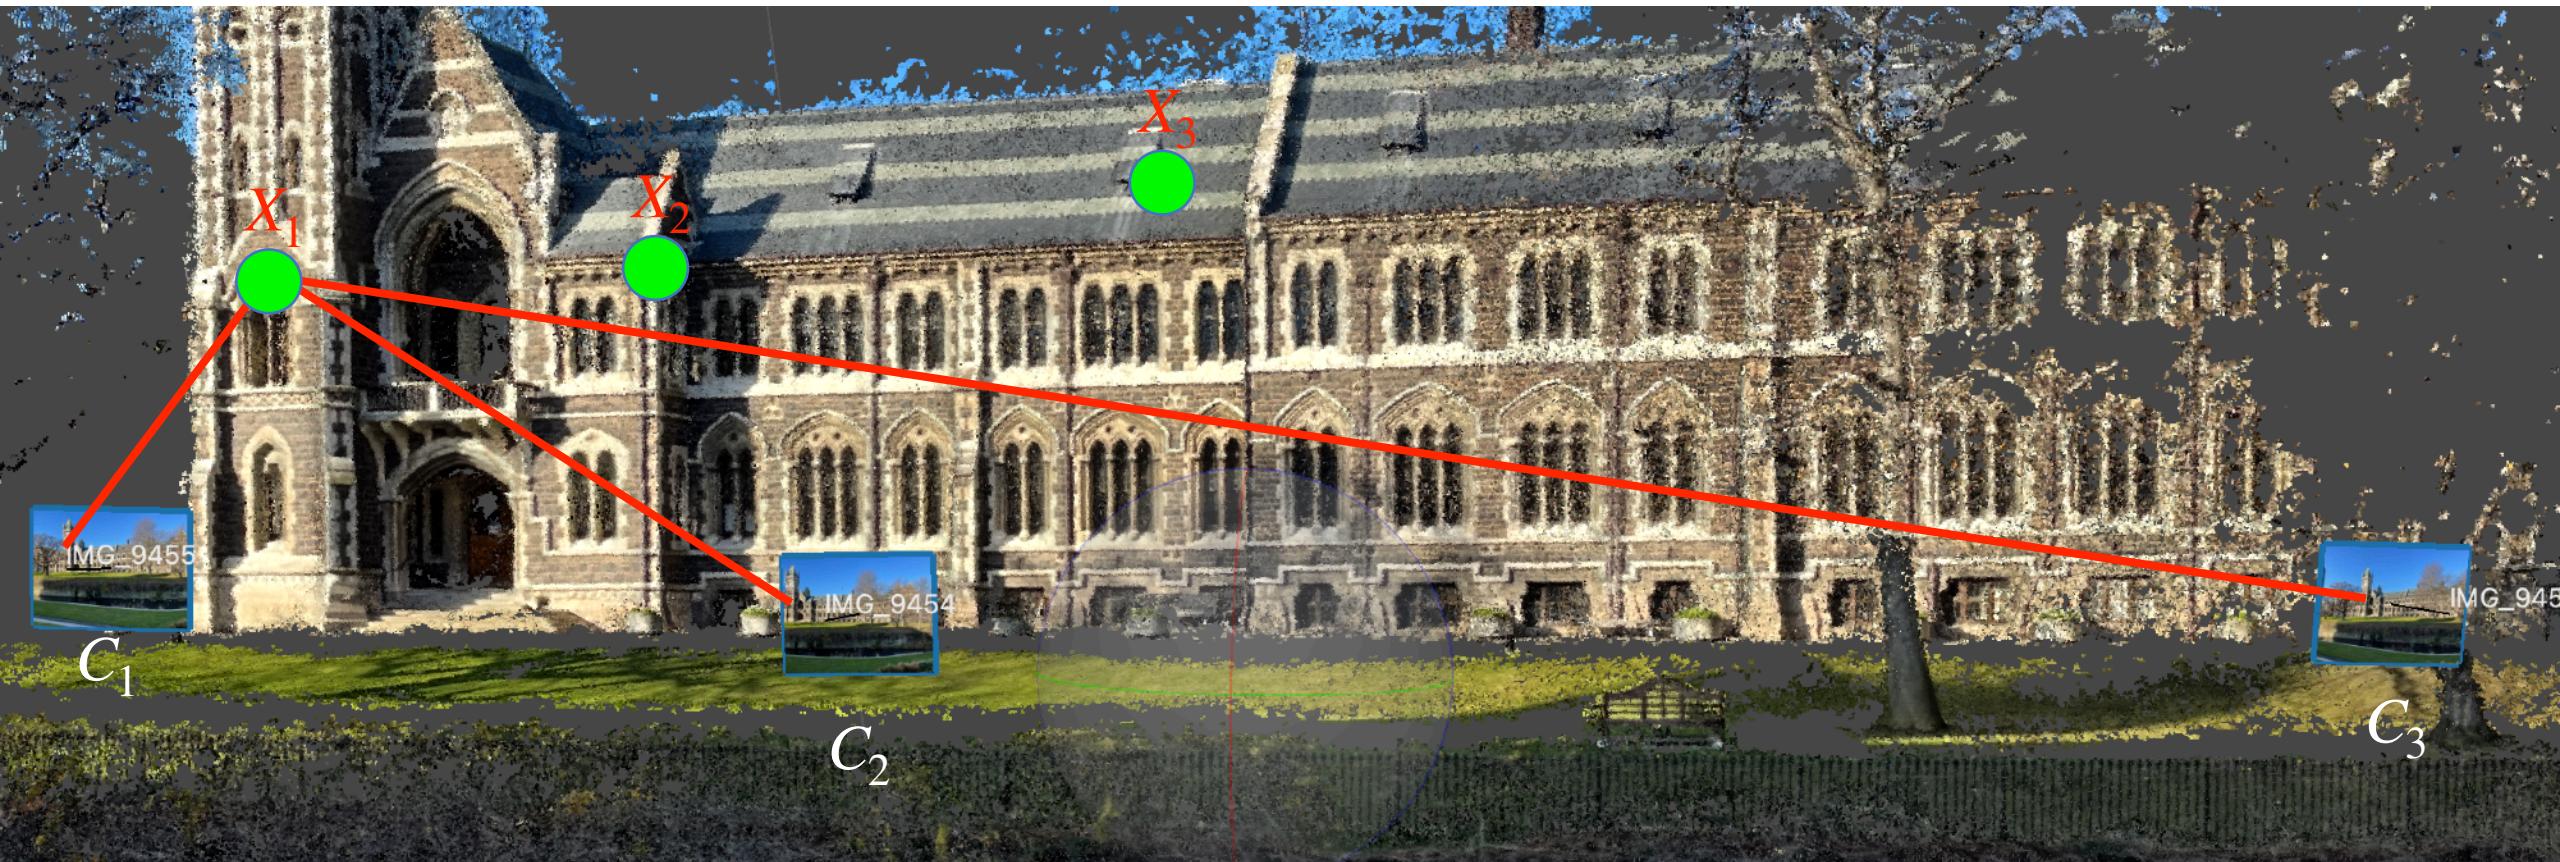


Colmap Workflow

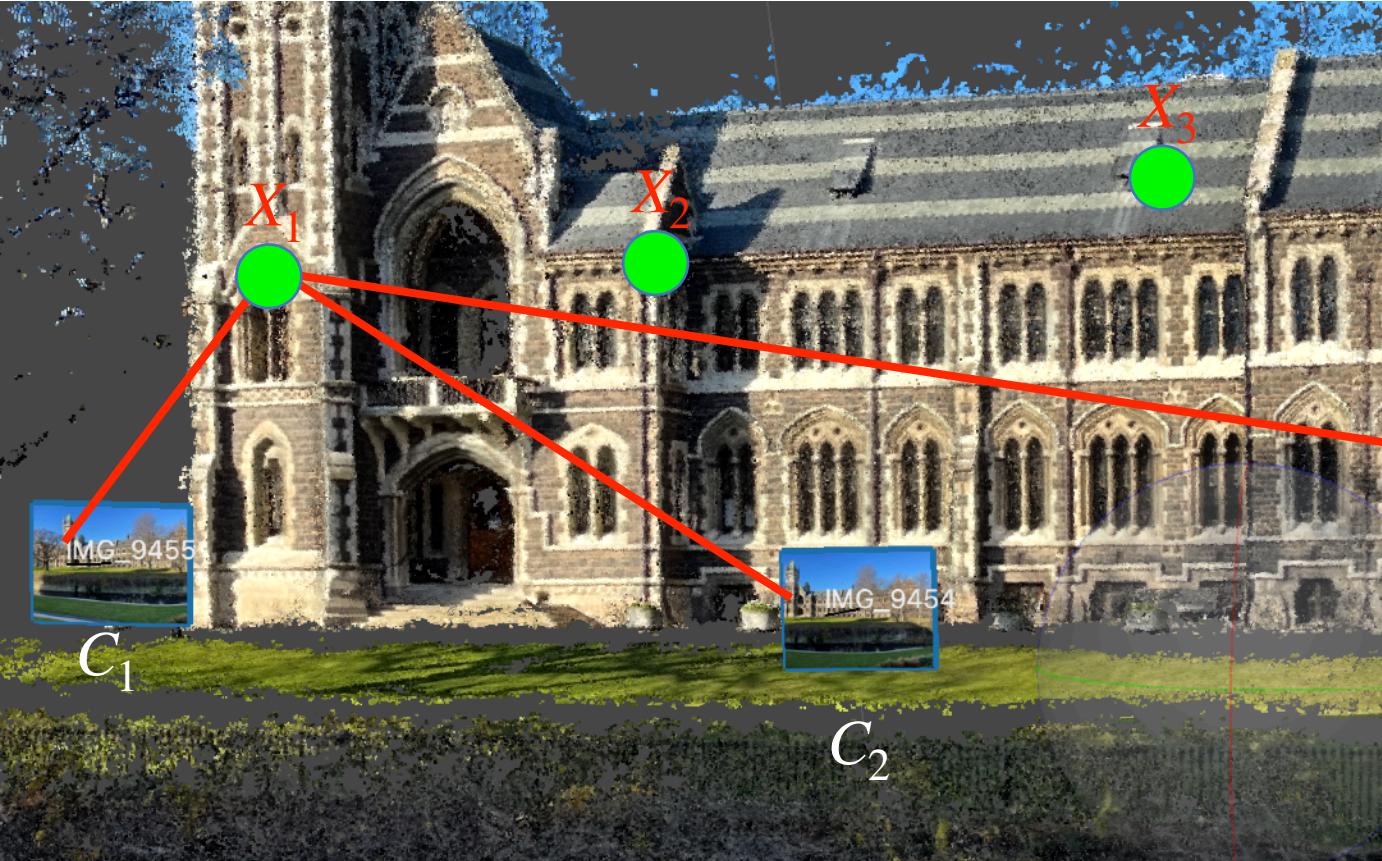


<https://colmap.github.io/tutorial.html>

Structure-from-Motion



Structure-from-Motion



- Estimation of 3D points X_i and camera poses C_j
- Based on motion (camera views from different viewpoints)

Relative Pose Estimation

The Essential Matrix

- Reminder: Fundamental matrix

$$F = K_2^{-T} [t]_x R K_1^{-1}$$

- For calibrated cameras

- We can factor out K_1 and K_2 in

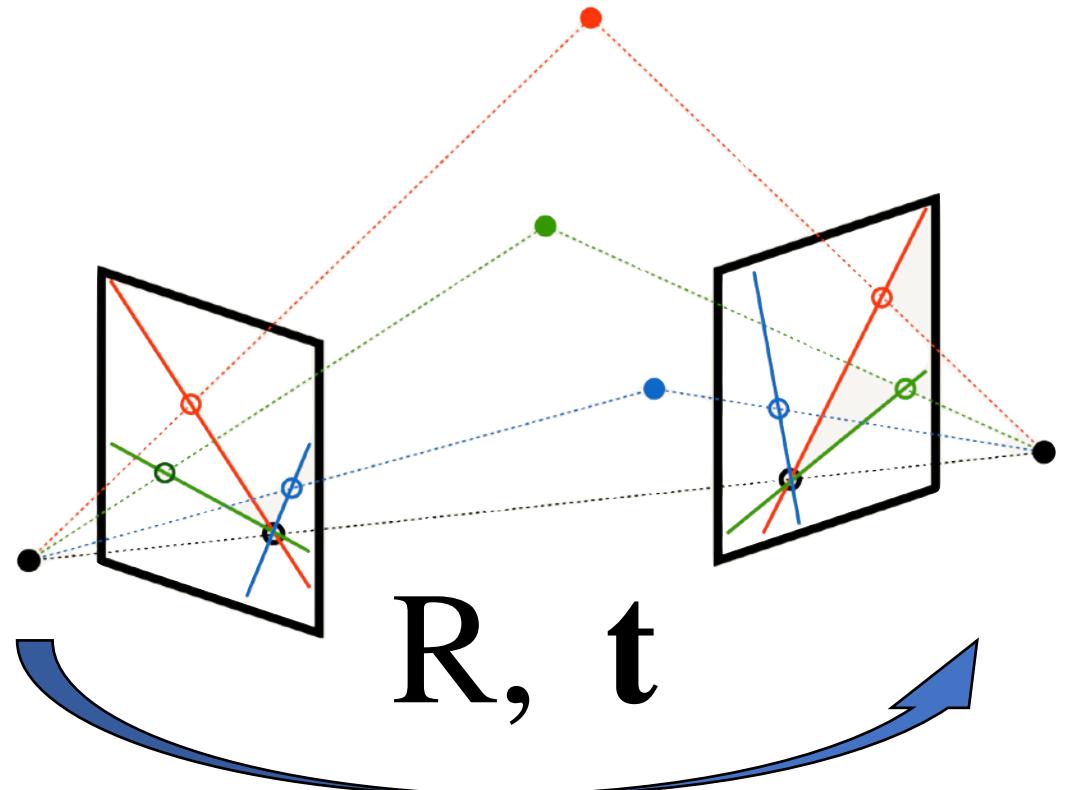
$$F = K_2^{-T} [t]_x R K_1^{-1}$$

- This gives us the essential matrix

$$E = K_2^T F K_1 = [t]_x R$$

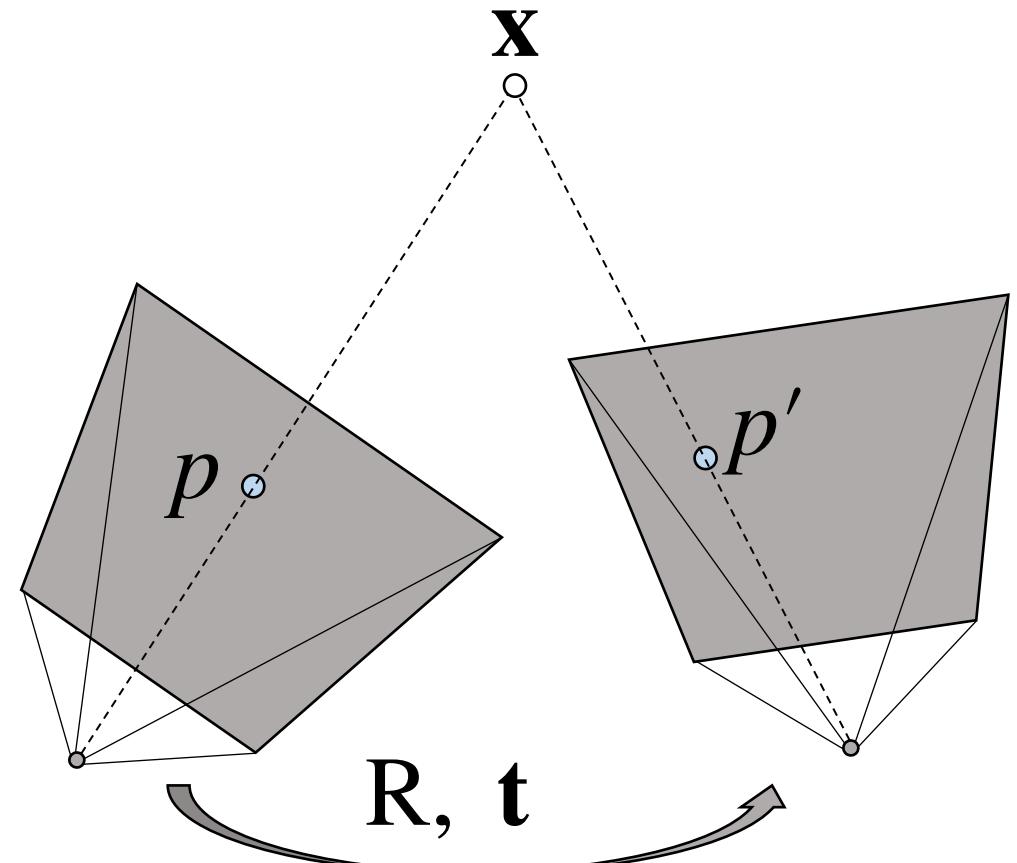
- E captures the relative pose of the two cameras

- A rotation, R , and translation, t



Relative Pose Estimation

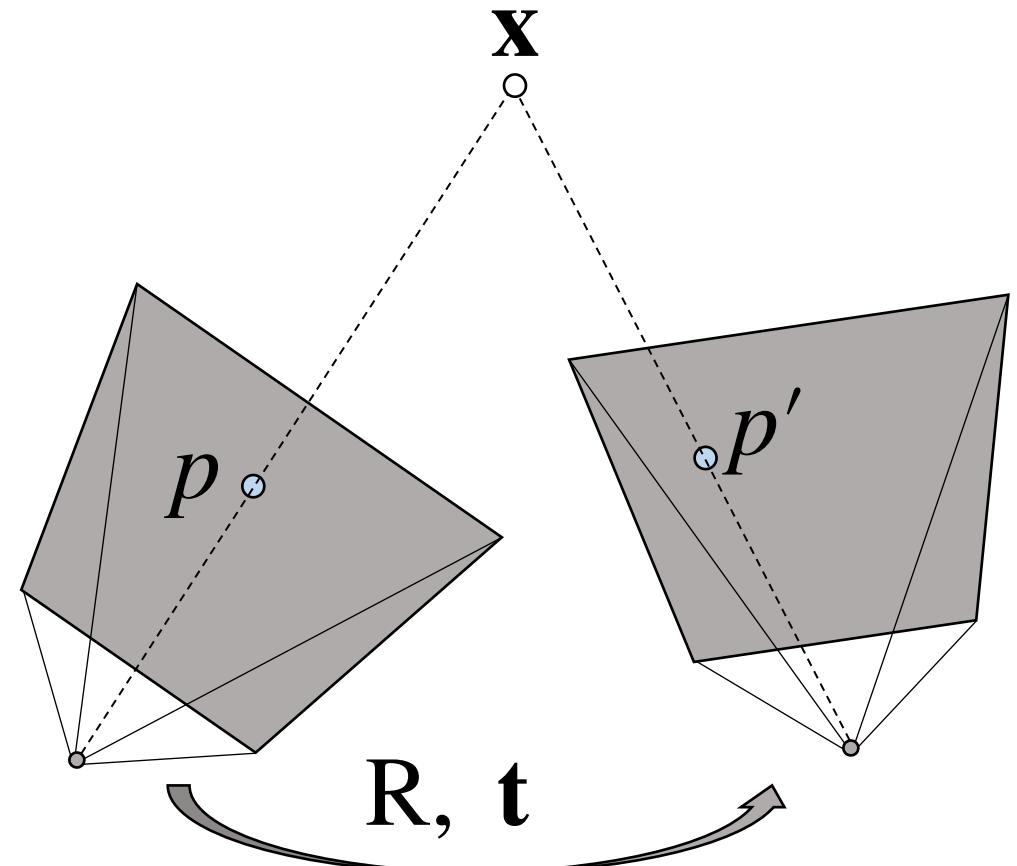
- To recover 3D world point, \mathbf{x}
 - We need to know the rotation, \mathbf{R} , and translation, \mathbf{t} , between the cameras
 - This is known as their relative pose
- Conceptually, we intersect rays
 - From camera centres through the matching image points $p \leftrightarrow p'$
 - Intersection (or nearest approach) is an estimate of \mathbf{x}



Relative Pose Estimation

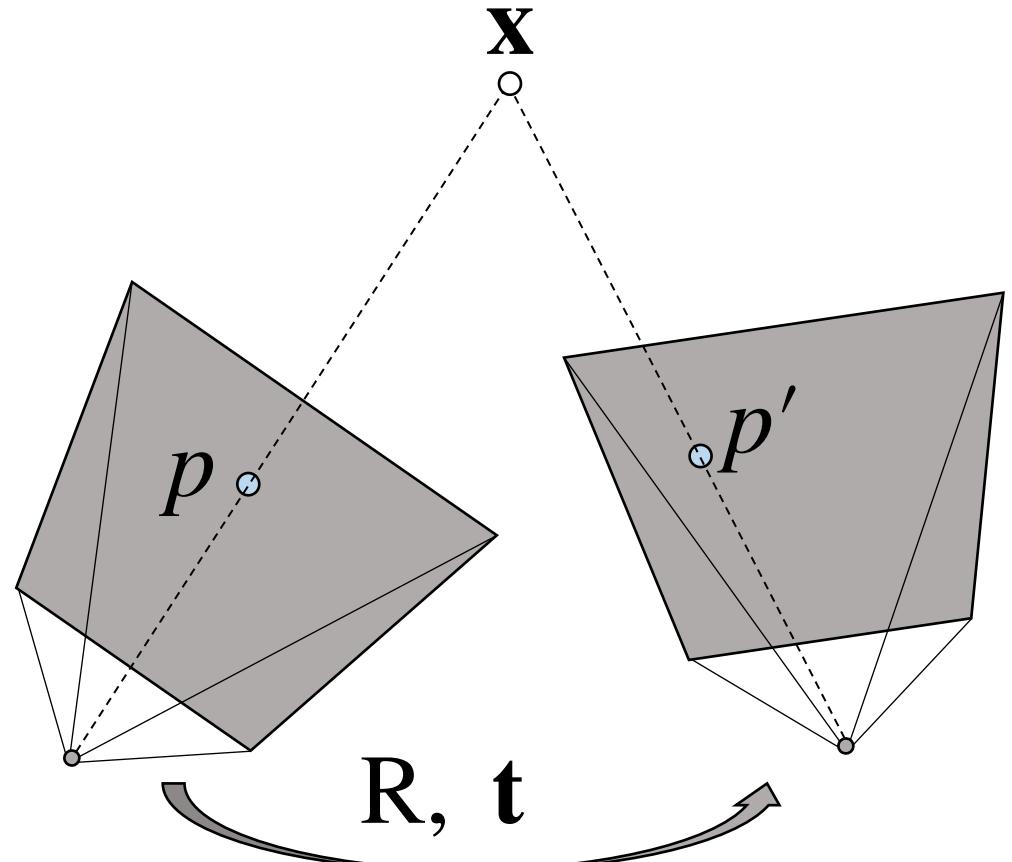
- E encodes R and t : $E = [t]_x R$
 - $[t]_x$ is the matrix that has the effect of giving the cross product with t
- How do we recover R and t from E ?

$$[t]_x = \begin{bmatrix} 0 & -t_z & t_y \\ t_z & 0 & -t_x \\ -t_y & t_x & 0 \end{bmatrix}$$

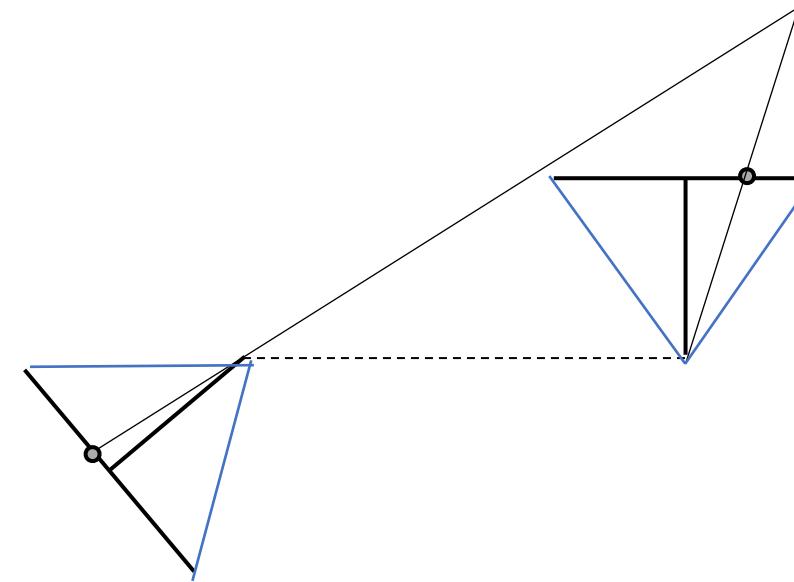
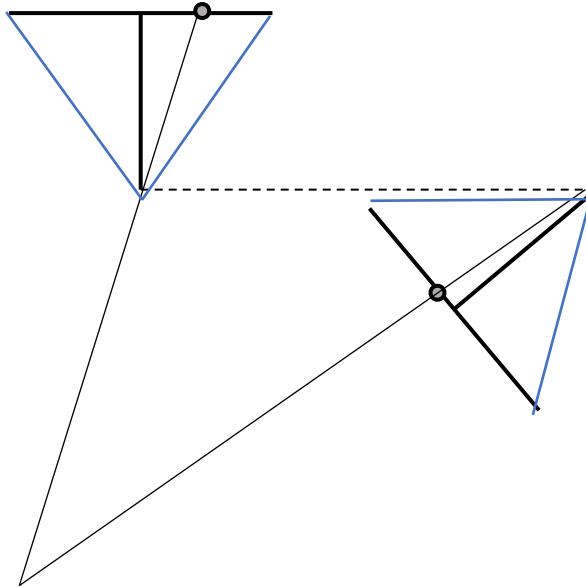
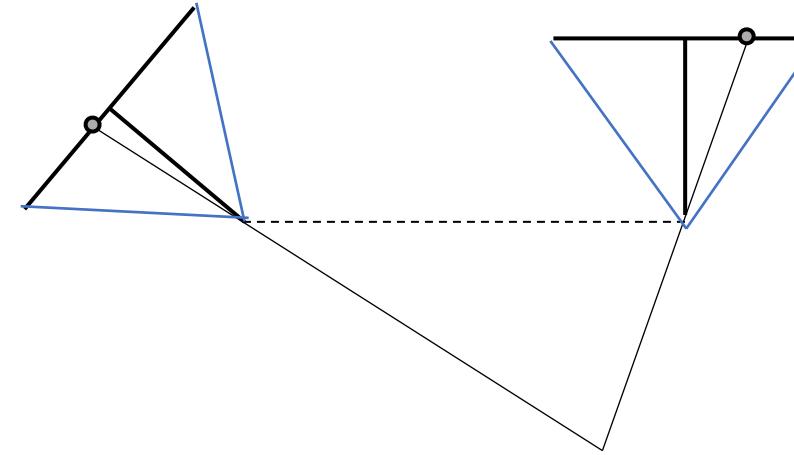
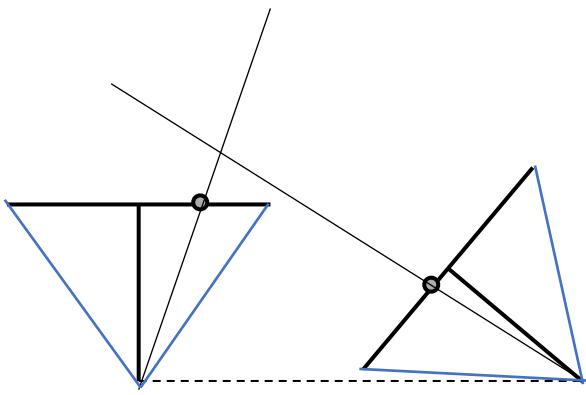


Decomposing the Essential Matrix

- To recover R and t
 - Compute SVD of E
 - SVD: gives us three matrices U S and V
transpose: $E = USV^T$
 - From this decomposition, we get:
 - Two possible rotations (R)
 - Two possible translations (t)
 - That means there are four candidate solutions in total.
 - To find the correct one: check which solution places the 3D points in front of both cameras



Which Solution is Correct?



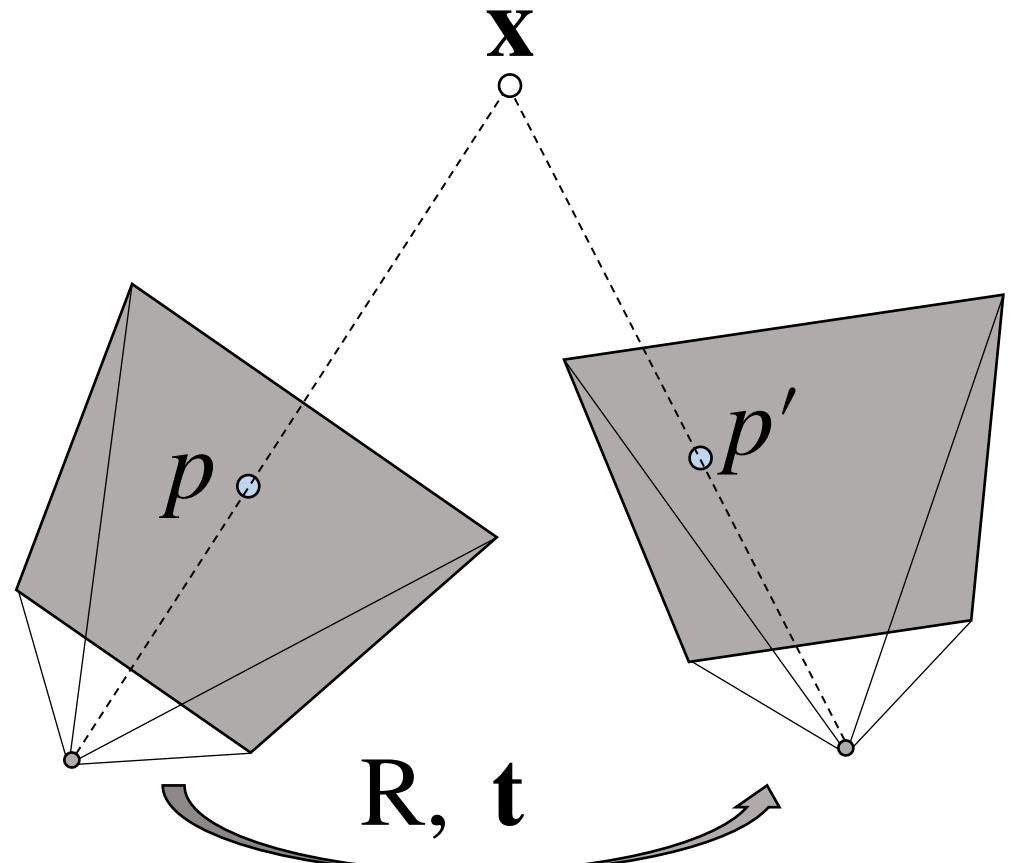
3D Reconstruction

3D Structure Estimation

- We now know
 - Calibration matrices, \mathbf{K} and \mathbf{K}'
 - Relative pose, \mathbf{R} and \mathbf{t}
 - Matching points, $p \leftrightarrow p'$
- These give us a 3D world point, \mathbf{x}

$$p \equiv \mathbf{K} [\mathbf{I} \ \mathbf{0}] \mathbf{x} = \mathbf{P} \mathbf{x}$$

$$p' \equiv \mathbf{K}' [\mathbf{R} \ \mathbf{t}] \mathbf{x} = \mathbf{P}' \mathbf{x}$$



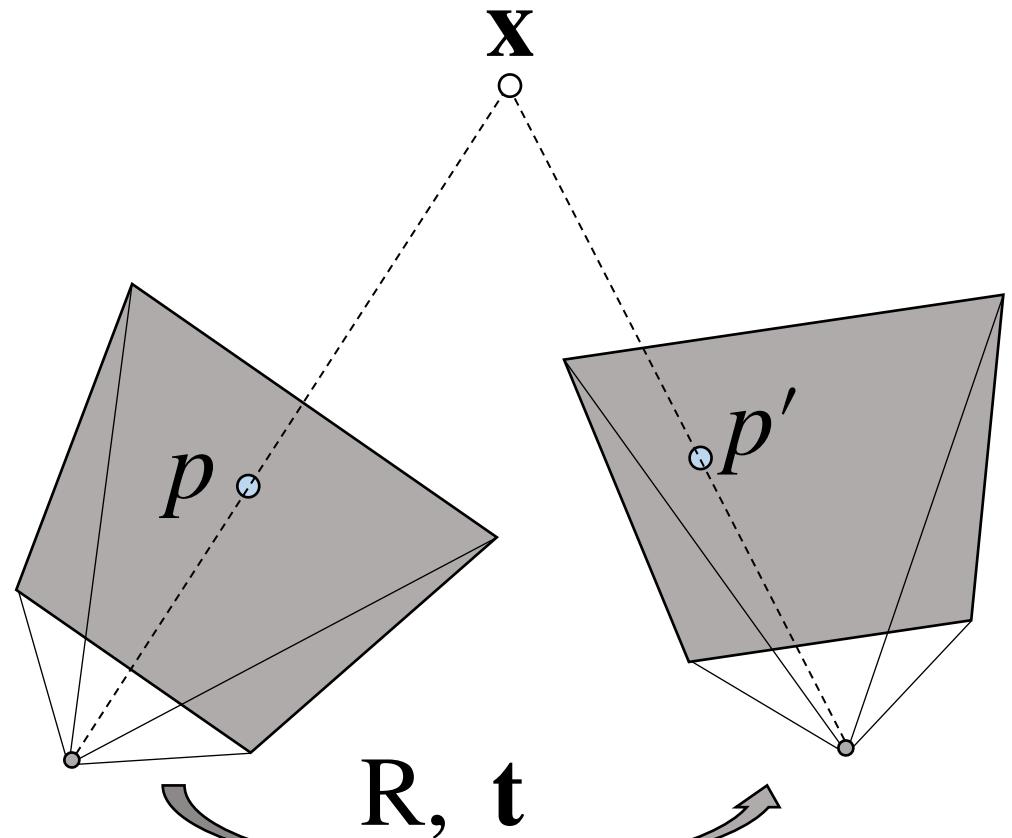
3D Structure Estimation

- If we let \mathbf{r}_i^T be the i th row of \mathbf{P} :

$$p \times \mathbf{Px} = \mathbf{0}$$

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \times \begin{bmatrix} \mathbf{r}_1^T \mathbf{x} \\ \mathbf{r}_2^T \mathbf{x} \\ \mathbf{r}_3^T \mathbf{x} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} (v\mathbf{r}_3^T - \mathbf{r}_2^T)\mathbf{x} \\ (\mathbf{r}_1^T - u\mathbf{r}_3^T)\mathbf{x} \\ (u\mathbf{r}_2^T - v\mathbf{r}_1^T)\mathbf{x} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

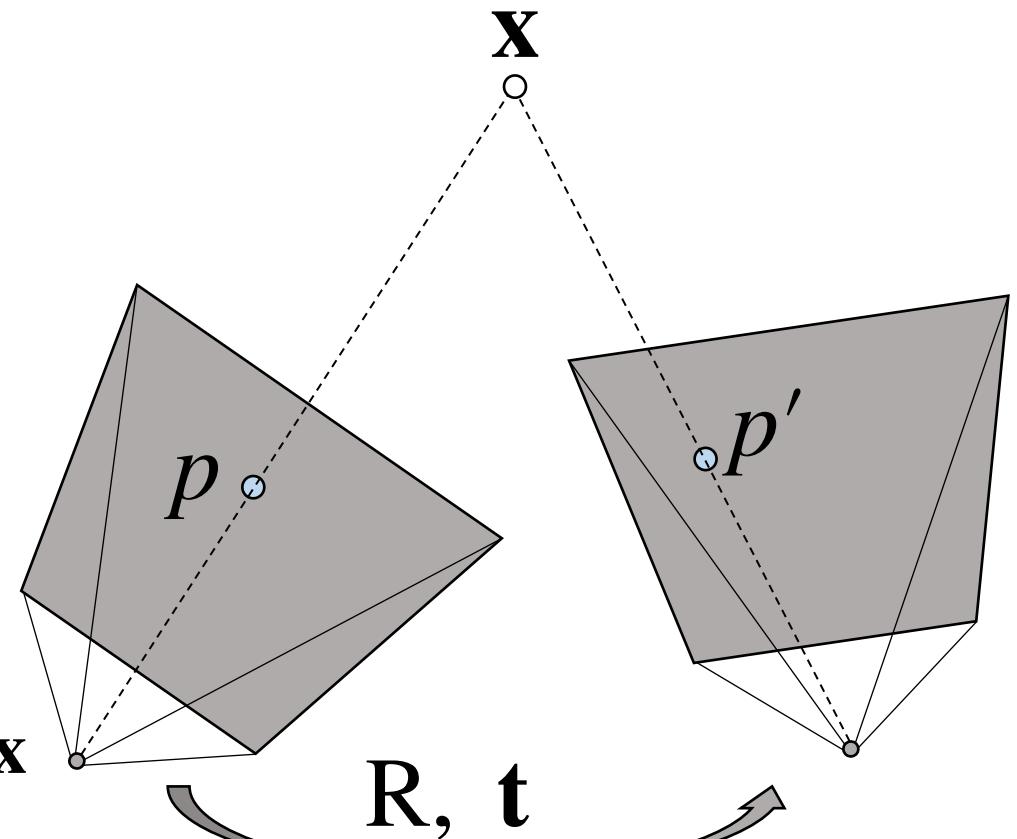


3D Structure Estimation

- This looks like 3 equations in \mathbf{x}
 - They are not linearly independent
 - Can only use two
 - Usually take the first two (symmetry)
 - Get two more from $p' \times P'\mathbf{x} = \mathbf{0}$

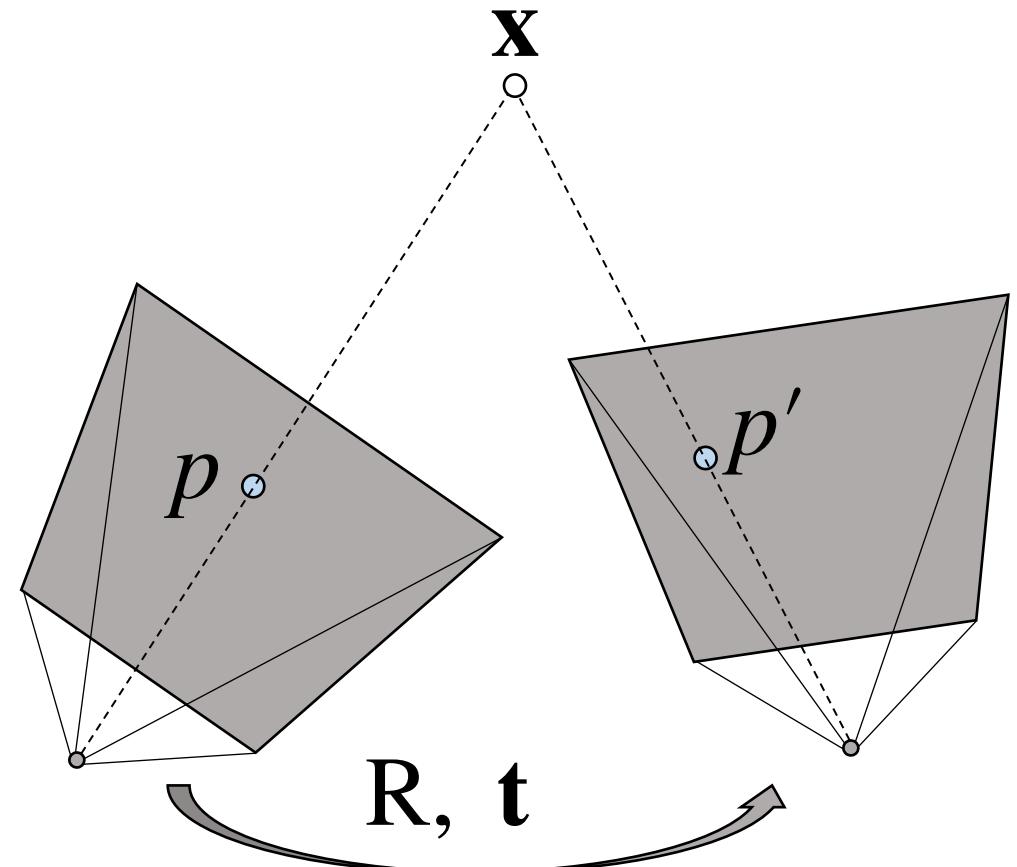
$$\begin{aligned} (v\mathbf{r}_3^T - \mathbf{r}_2^T)\mathbf{x} &= 0 \\ (\mathbf{r}_1^T - u\mathbf{r}_3^T)\mathbf{x} &= 0 \\ (v'\mathbf{r}_3^T - \mathbf{r}_2^T)\mathbf{x} &= 0 \\ (\mathbf{r}_1^T - u'\mathbf{r}_3^T)\mathbf{x} &= 0 \end{aligned}$$

} Four equations in \mathbf{x}
Solve $A\mathbf{x} = \mathbf{0}$



Reprojection Error

- From our measurements
 - Point correspondences $p_i \leftrightarrow p'_i$
 - Camera calibration matrices \mathbf{K} , and \mathbf{K}'
- We have estimated
 - The fundamental/essential matrices
 - The camera's relative pose, \mathbf{R} , \mathbf{t}
 - The 3D world points, \mathbf{x}_i
- Have minimised algebraic error
 - Convenient, but not meaningful

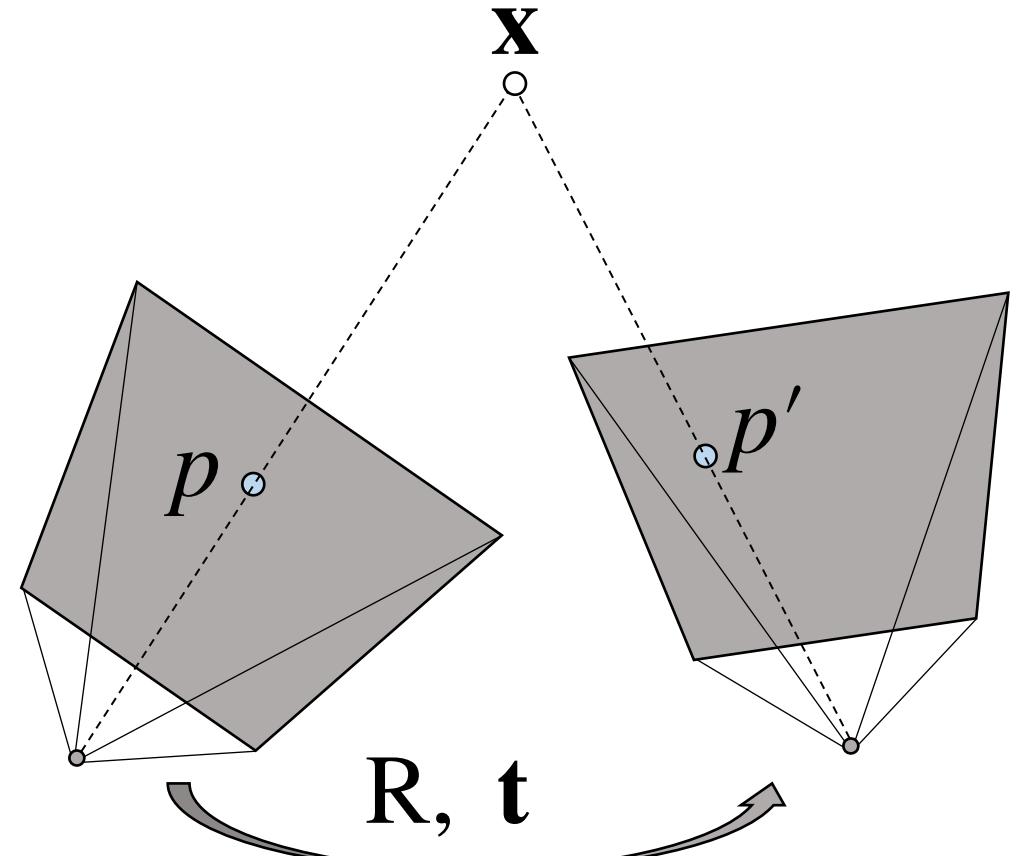


Reprojection Error

- There are true values for
 - The camera calibration matrices $\tilde{\mathbf{K}}$, $\tilde{\mathbf{K}}'$
 - Their relative pose, $\tilde{\mathbf{R}}$, $\tilde{\mathbf{t}}$
 - The 3D world points, \mathbf{x}_i
- These predict our measurements

$$\tilde{\mathbf{p}}_i = \tilde{\mathbf{K}} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \tilde{\mathbf{x}}_i$$

$$\tilde{\mathbf{p}}'_i = \tilde{\mathbf{K}}' \begin{bmatrix} \tilde{\mathbf{R}} & \tilde{\mathbf{t}} \end{bmatrix} \tilde{\mathbf{x}}_i$$



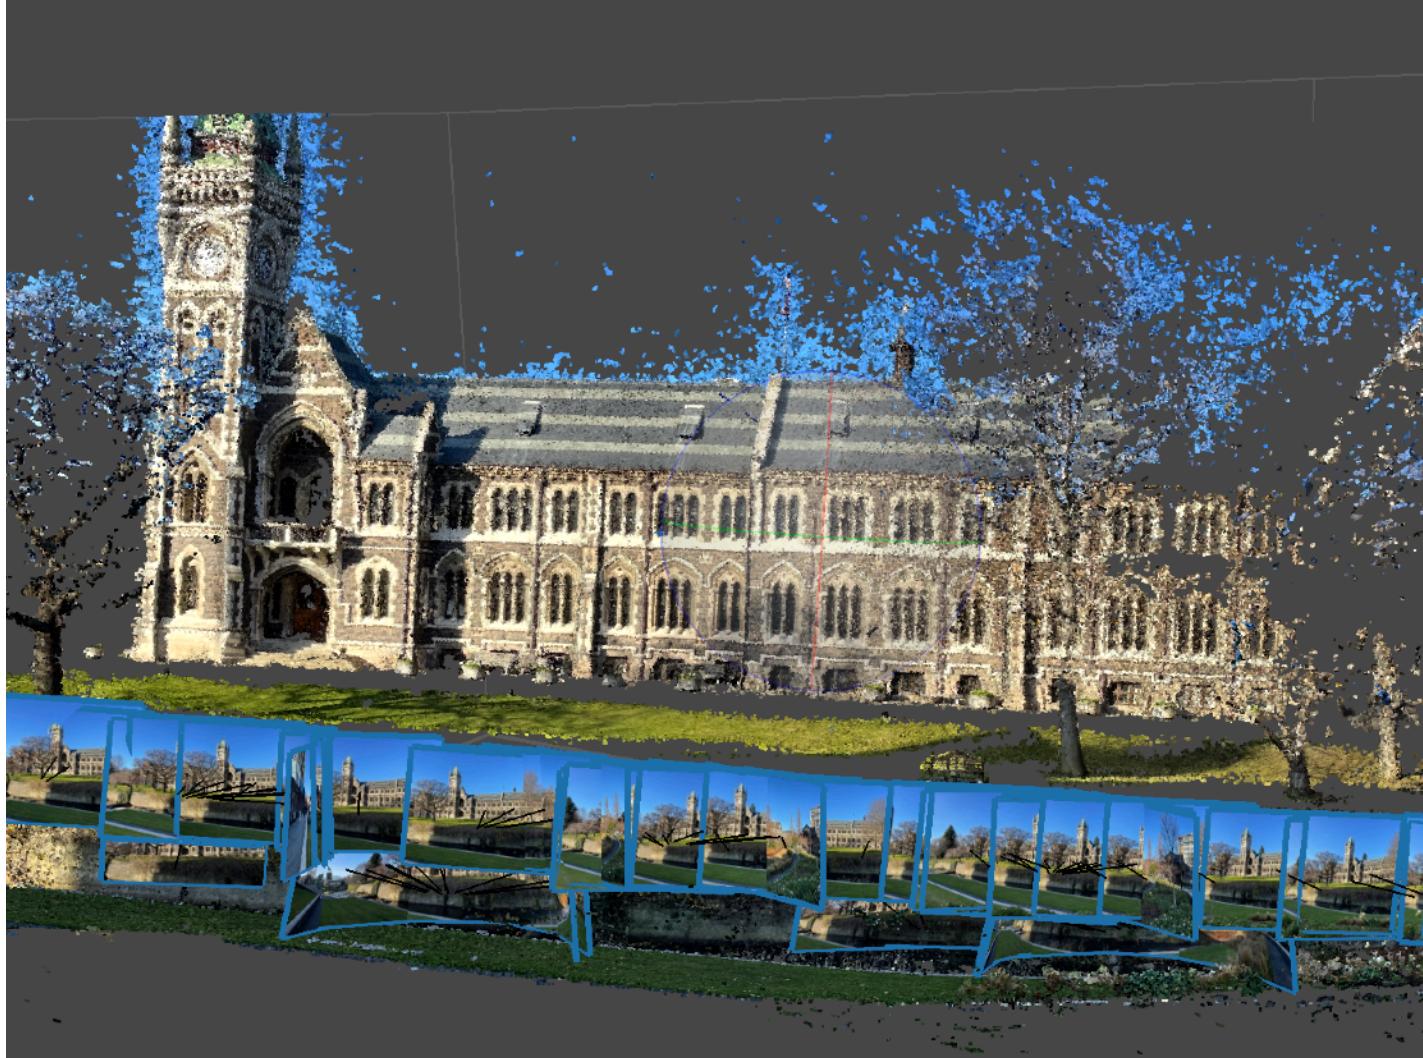
Reprojection Error

$$\sum \left((p_i - \tilde{p}_i)^2 + (p'_i - \tilde{p}'_i)^2 \right) = \sum \left((p_i - \tilde{\mathbf{K}}[\mathbf{I} \ \mathbf{0}] \tilde{\mathbf{x}}_i)^2 + (p'_i - \tilde{\mathbf{K}}'[\tilde{\mathbf{R}} \ \tilde{\mathbf{t}}] \tilde{\mathbf{x}}_i)^2 \right)$$

- This is a non-linear function with parameters $\tilde{\mathbf{K}}$, $\tilde{\mathbf{K}}'$, $\tilde{\mathbf{R}}$, $\tilde{\mathbf{t}}$, $\tilde{\mathbf{x}}_i$
 - Have initial estimates from minimising algebraic error
 - Minimise via Levenberg-Marquardt

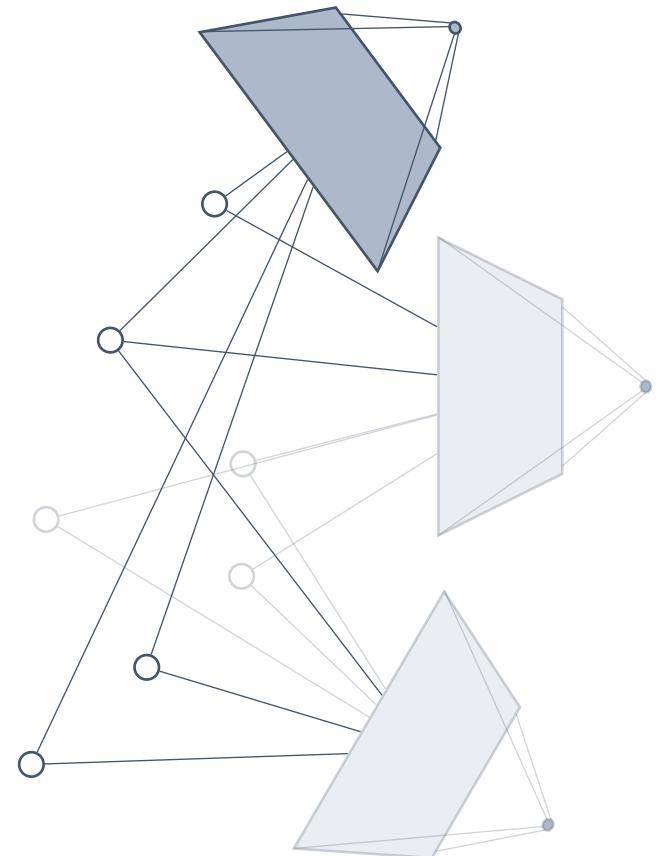
Multi-View Stereo

- More than two images
 - Multiple camera rigs
 - Single moving camera
 - Multiple moving cameras
- More information
 - Reconstruct larger areas
 - Resolve more details
 - More accurate models



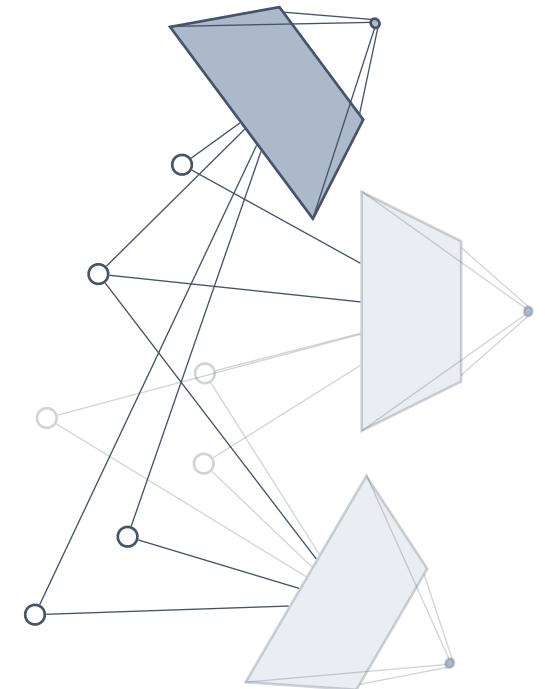
Discussion (1min):

Can we just find relative poses
between each pair?



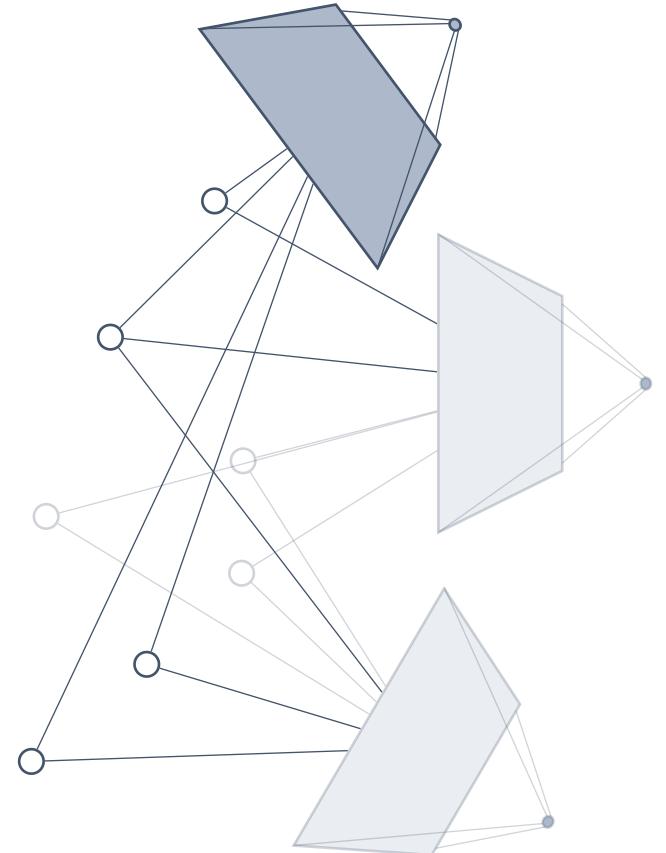
**3-Min Discussion: Is finding relative poses
between each pair enough? What are the
challenges?**

03:00



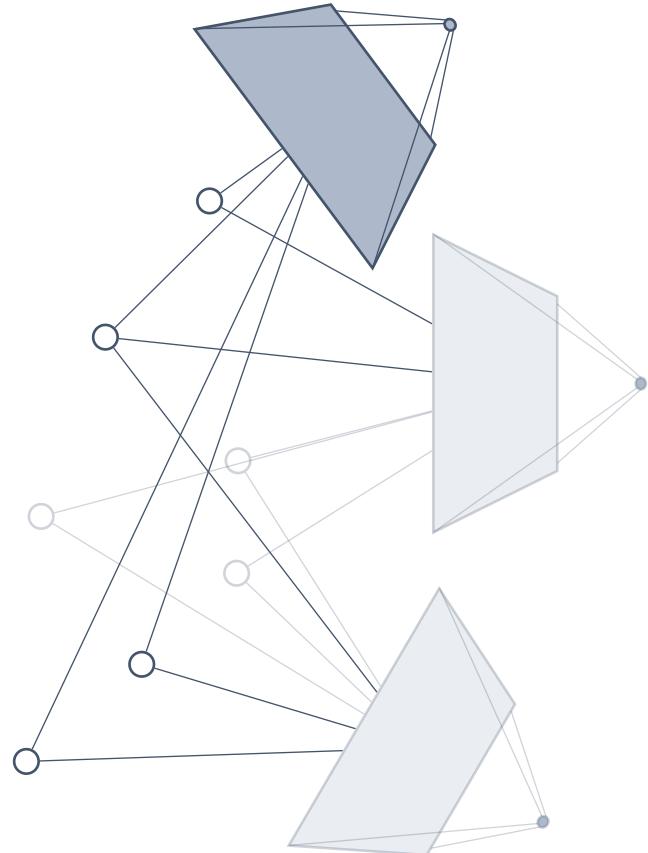
Multi-View Stereo

- Two-view methods
 - Find fundamental matrices between each pair
 - Scales are not independent
 - Non-overlapping views
- Incremental approach
 - Start with two views
 - Recover motion & structure
 - Find next camera's pose
 - Recover more structure
 - Tasks to solve
 - Determine order of images
 - Recover absolute pose



Multi-View Stereo Pipeline

- Identify best matching pair
 - Two-view stereo for relative pose
 - Reconstruct 3D points
 - Minimise reprojection error
- While still images to add
 - Select next best view
 - Determine absolute pose
 - Reconstruct more 3D points
 - Minimise reprojection error



Multi-View Stereo Pipeline

- Identify best matching pair
 - Two-view stereo for relative pose
 - Reconstruct 3D points
 - Minimise reprojection error
- While still images to add
 - Select next best view
 - Determine absolute pose
 - Reconstruct more 3D points
 - Minimise reprojection error

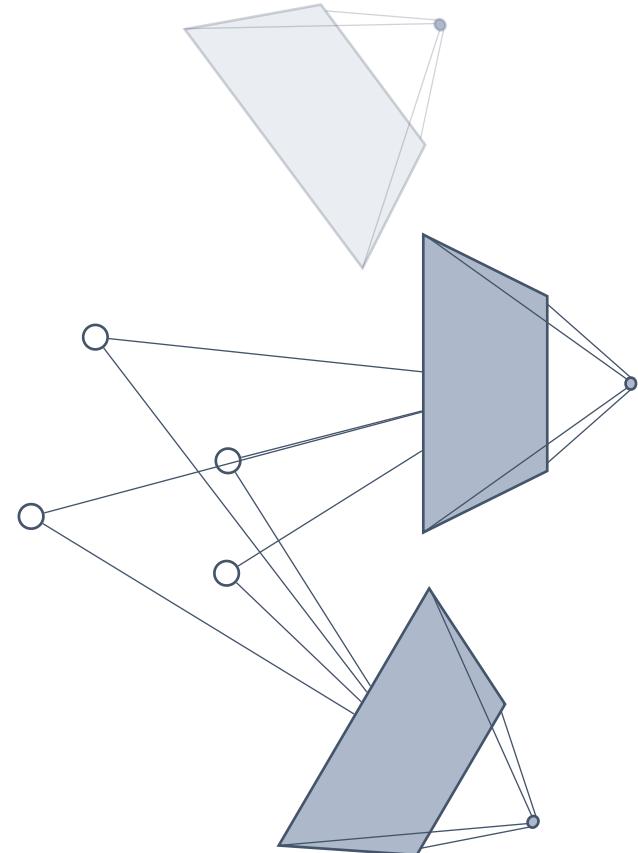


Image Bag-of-Words

Multi-View Stereo Pipeline

- Start by finding image features
- Matching all pairs is expensive
 - For n images, $O(n)^2$ pairs
 - Feature matching between one pair is quite expensive (as we've seen)
- Need a quick check of pairs
 - How similar are two images?
 - How many likely matches are there?

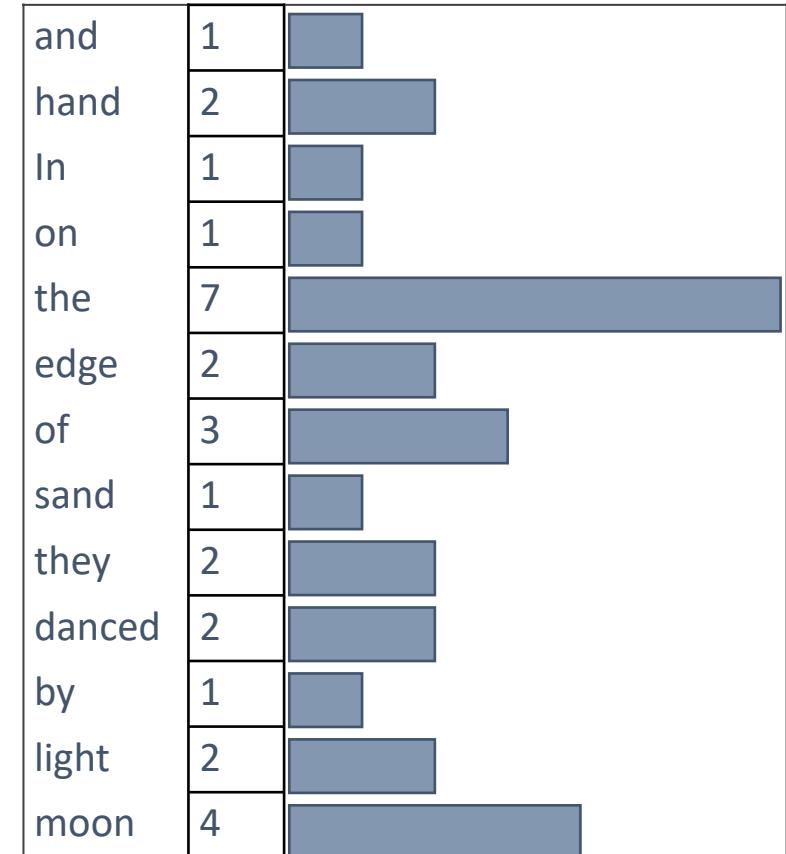


Bags of Words (BoW)

*And hand in hand on the edge of the sand
they danced by the light of the moon,
the moon, the moon
They danced by the light of the moon*

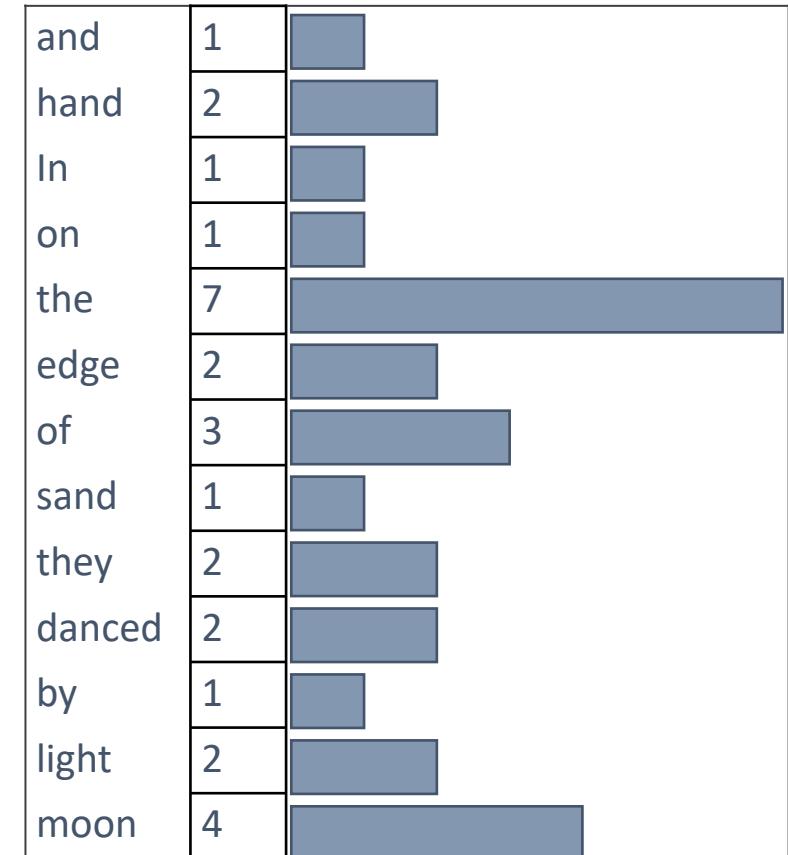
*Edward Lear
from The Owl and the Pussycat*

- From information retrieval
 - Group features into ‘words’
 - Count how many times each ‘word’ appears in each image -> histogram
 - Compare histograms



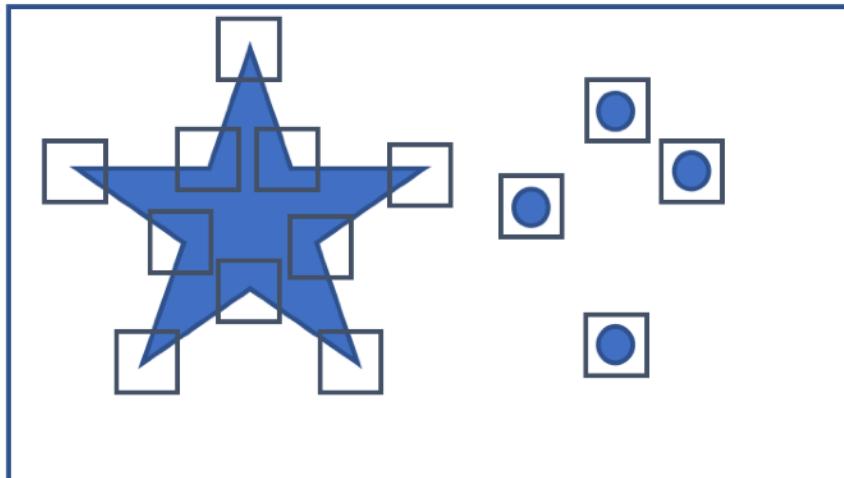
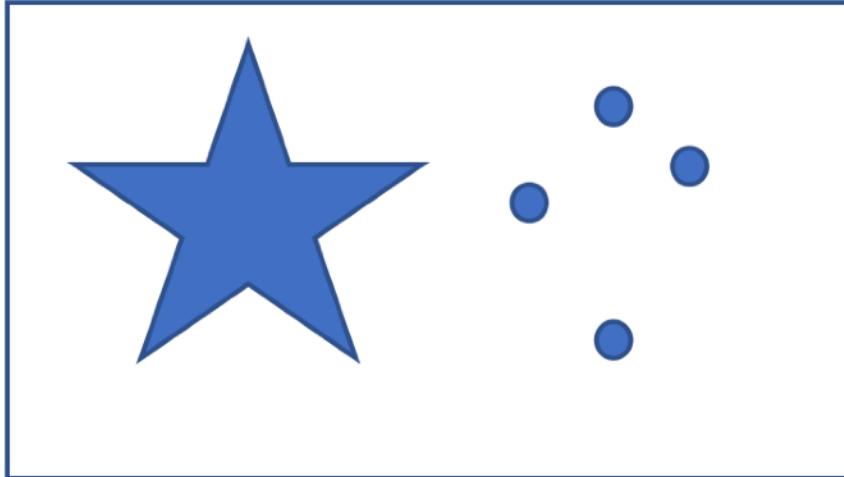
Bags of Words (BoW)

- Bag of Words (BoW) vectors:
 - Ignores word order
 - Can ignore counts – binary there/not
 - Content, not structure
- For text, group similar words
 - Computer, computing, compute...
- How to compare vectors?
- How to apply to images?



Bags of Words (BoW)

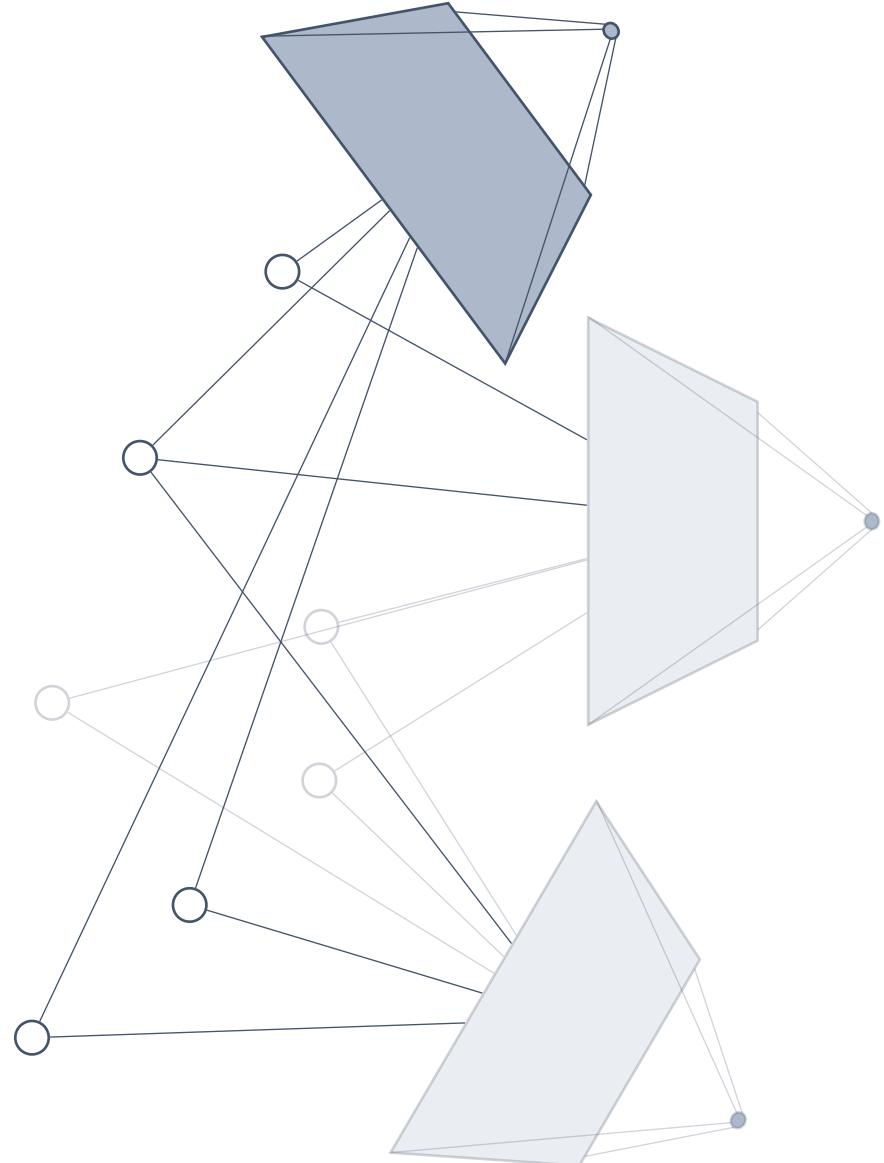
- Need equivalent of words
- Cluster features – can pre-train
 - Choose some number, k , of ‘words’
 - Make k clusters – e.g. k -means
 - Cluster centres become words
- Bag of words for images
 - Detect and describe feature
 - Count features closest to each word



Absolute Pose

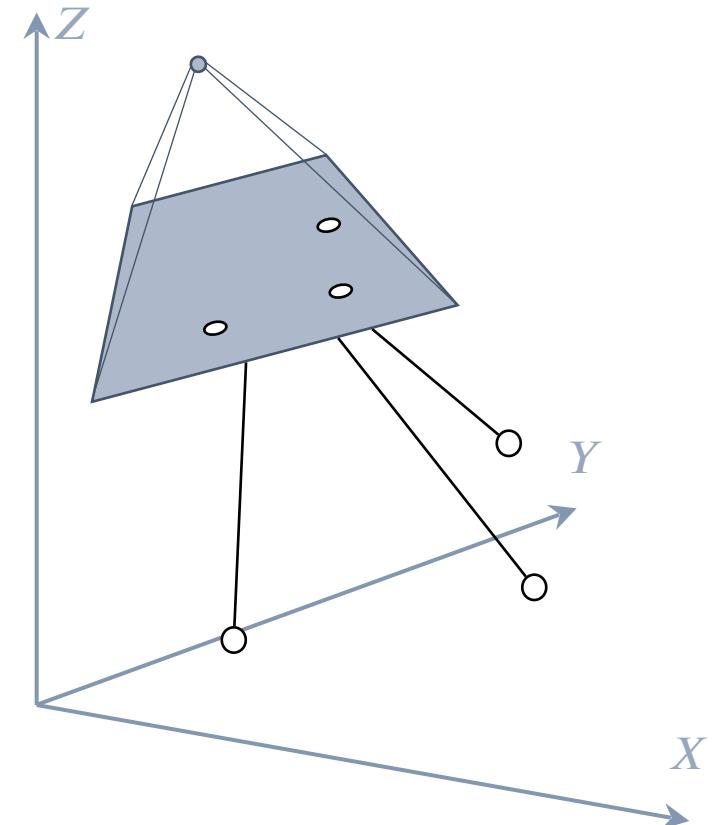
Multi-View Stereo Pipeline

- Identify best matching pair
 - Two-view stereo for relative pose
 - Reconstruct 3D points
 - Minimise reprojection error
- While still images to add
 - Select next best view
 - *Determine absolute pose*
 - Reconstruct more 3D points
 - Minimise reprojection error



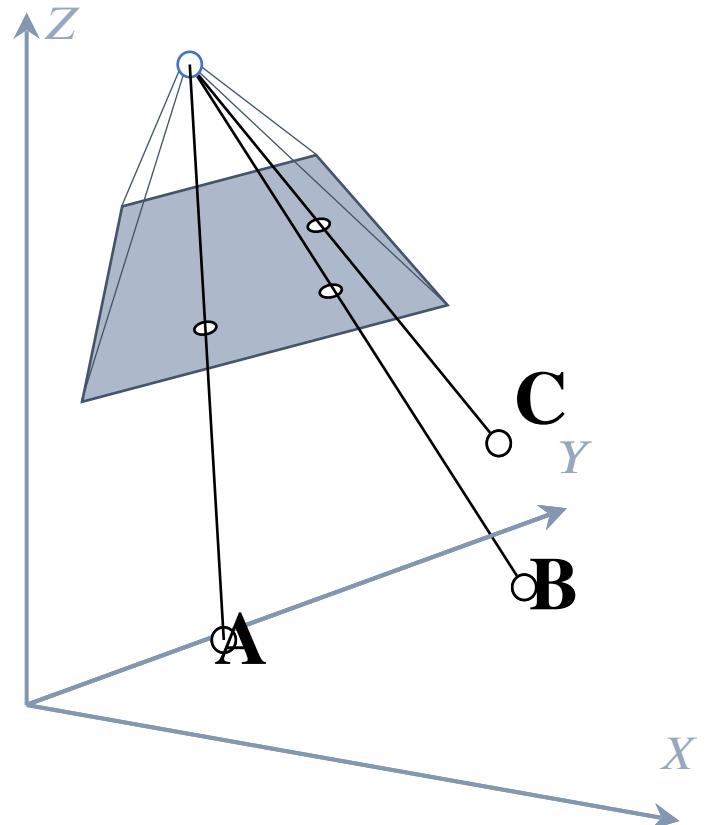
Absolute Pose Problem

- We have an image taken with a calibrated camera
- We've already established a coordinate frame
- Need to find pose (rotation and translation) of the camera when the image was taken



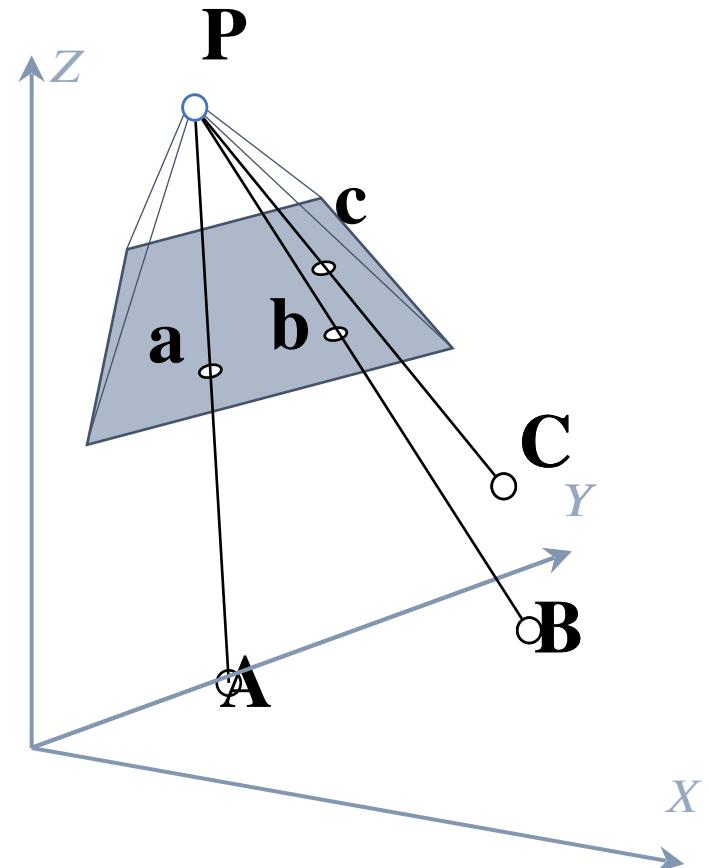
Perspective- n -Point Pose (PnP)

- Pose from 2D-3D matches
 - Have six unknowns (\mathbf{R} , \mathbf{t})
 - Each 2D-3D match $\mathbf{u}_i \leftrightarrow \mathbf{x}_i$ gives
$$k_i \mathbf{u}_i = \mathbf{K}[\mathbf{R} \quad \mathbf{t}] \mathbf{x}_i$$
- How many matches?
 - Each point gives 3 equations
 - Adds 1 unknown (k_i)
 - $6 + n$ unknowns, $3n$ equations,
 $n \geq 3$



Perspective-3-Point Pose (P3P)

- We have three 2D-3D matches
 - Call the 3D points **A**, **B**, and **C**
 - The matching image points are **a**, **b**, **c**
 - The (unknown) camera location is **P**
- Overview of solution
 - Use angles between \vec{PA} , \vec{PB} , and \vec{PC}
 - Find the distances $|\vec{PA}|$, $|\vec{PB}|$, $|\vec{PC}|$
 - This gives us equations for **P**



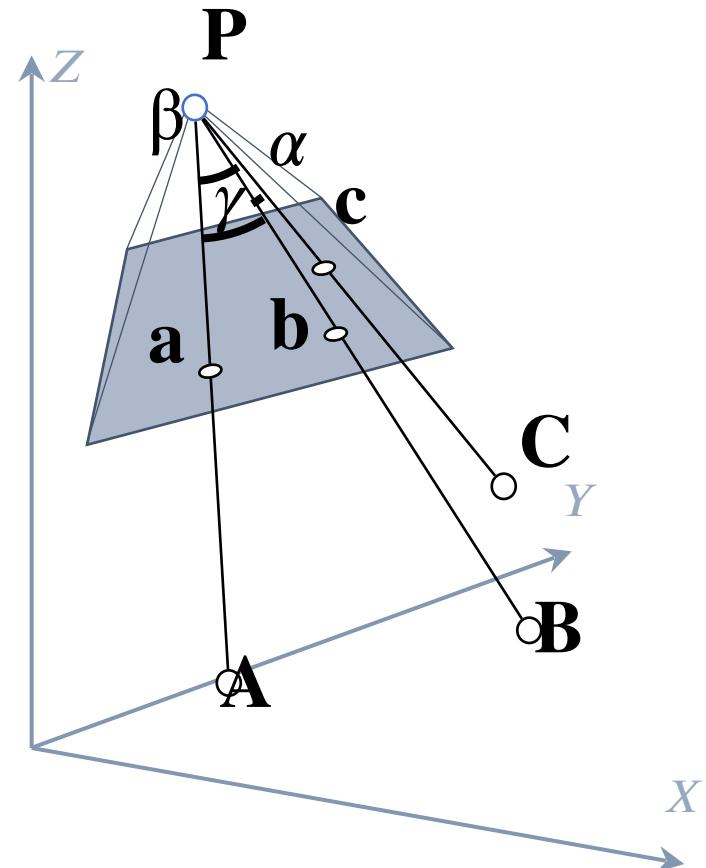
Perspective-3-Point Pose (P3P)

- Define the angles between rays
- $$\alpha = \angle BPC \quad \beta = \angle APC \quad \gamma = \angle APB$$

- Using dot products:

$$\cos\alpha = \frac{\mathbf{b} \cdot \mathbf{c}}{|\mathbf{b}| |\mathbf{c}|}$$

$$\cos\beta = \frac{\mathbf{a} \cdot \mathbf{c}}{|\mathbf{a}| |\mathbf{c}|} \quad \cos\gamma = \frac{\mathbf{a} \cdot \mathbf{b}}{|\mathbf{a}| |\mathbf{b}|}$$



Perspective-3-Point Pose (P3P)

- The law of cosines tells us that

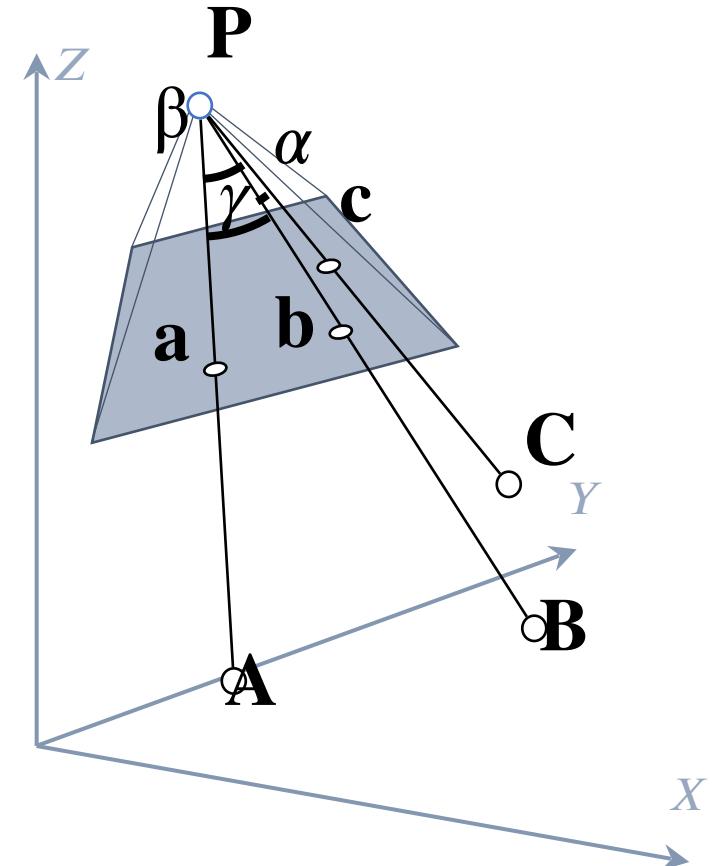
$$|\vec{AB}|^2 = |\vec{PA}|^2 + |\vec{PB}|^2 - 2|\vec{PA}||\vec{PB}|\cos\gamma$$

$$|\vec{AC}|^2 = |\vec{PA}|^2 + |\vec{PC}|^2 - 2|\vec{PA}||\vec{PC}|\cos\beta$$

$$|\vec{BC}|^2 = |\vec{PB}|^2 + |\vec{PC}|^2 - 2|\vec{PB}||\vec{PC}|\cos\alpha$$

- Solve for

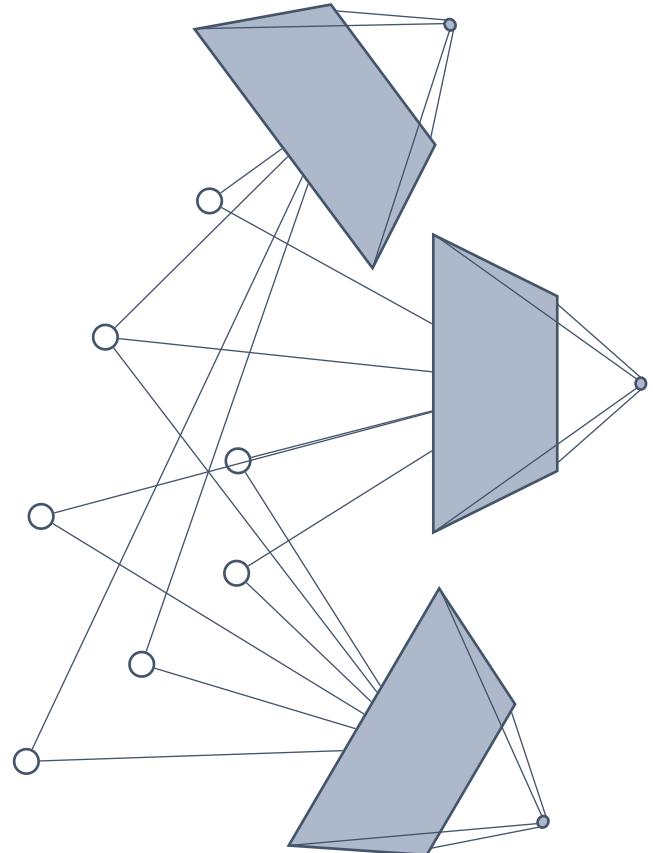
$$|\vec{PA}|, |\vec{PB}|, \text{ and } |\vec{PC}|$$



Bundle Adjustment

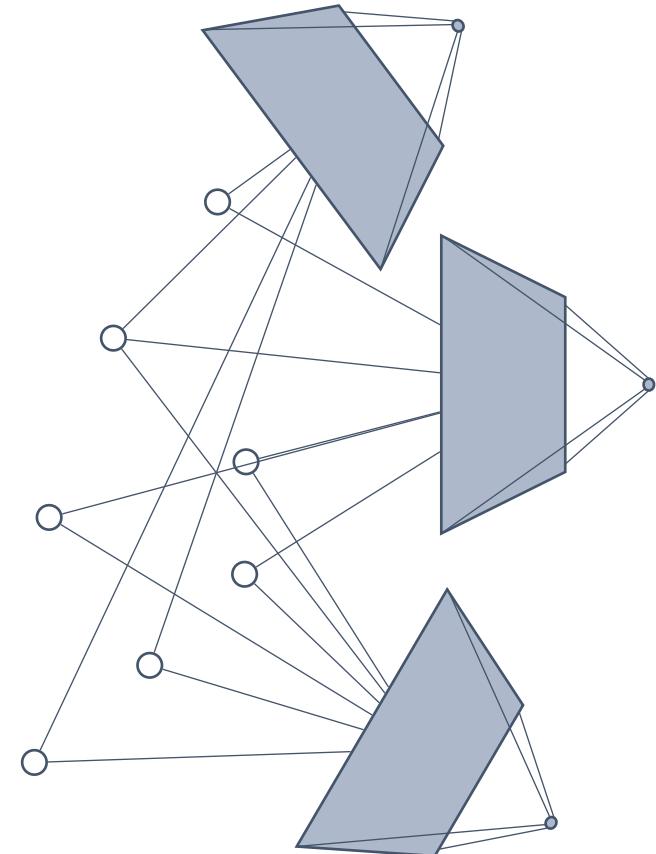
Multi-View Stereo Pipeline

- Identify best matching pair
 - Two-view stereo for relative pose
 - Reconstruct 3D points
 - Minimise reprojection error
- While still images to add
 - Select next best view
 - Determine absolute pose
 - Reconstruct more 3D points
 - *Minimise reprojection error*



Bundle Adjustment

- Minimising reprojection error for multi-view stereo
- Jointly refines:
 - Camera parameters (poses + intrinsics)
 - 3D point positions (scene structure)
- Goal: minimize reprojection error across all images
- Geometry forms **bundles** of rays for each camera
- **Adjusting** these gives an optimal reconstruction solution
- Large sparse non-linear least squares
- Levenberg-Marquardt is commonly used for this

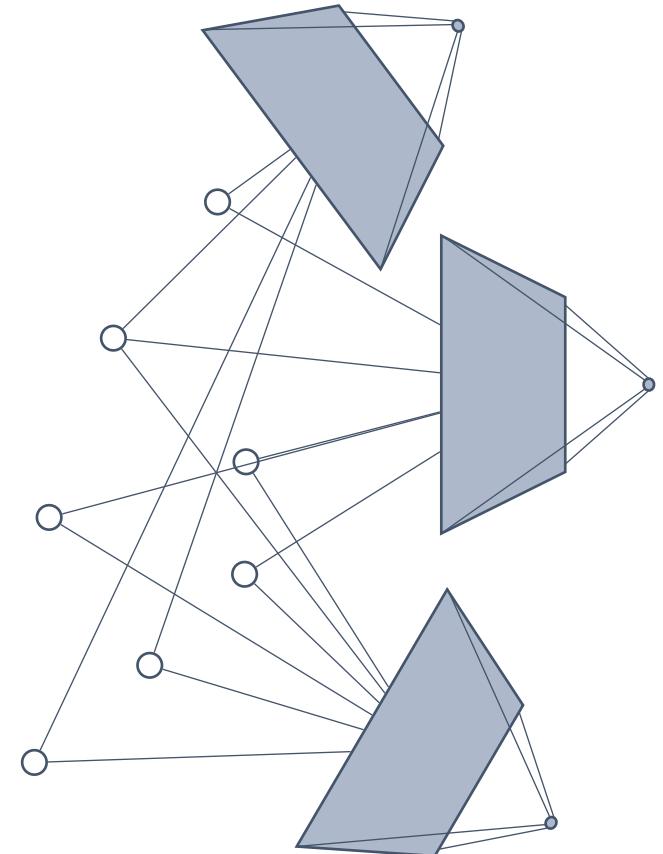


Reprojection Error

- If we have f features in n cameras, reprojection error is:

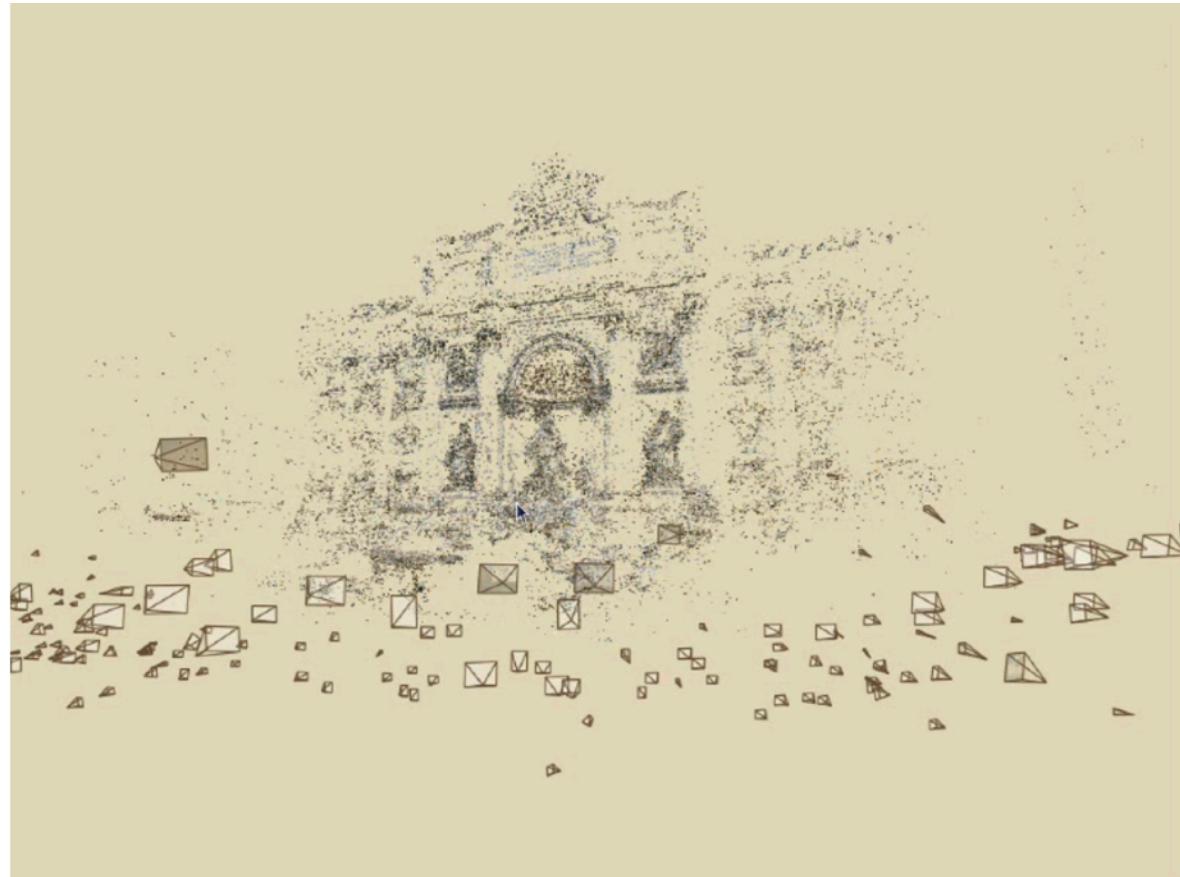
$$\sum_{i=1}^n \sum_{j=1}^f \left(\left(\mathbf{u}_{i,j} - \mathbf{K}_i [\mathbf{R}_i \quad \mathbf{t}_i] \mathbf{x}_j \right)^2 \right)$$

Measures Predictions



Bundle Adjustment

- Why is it Needed?
 - Individual pose or triangulation estimates are noisy
 - Errors accumulate across many views
- Bundle adjustment enforces global consistency



The end!